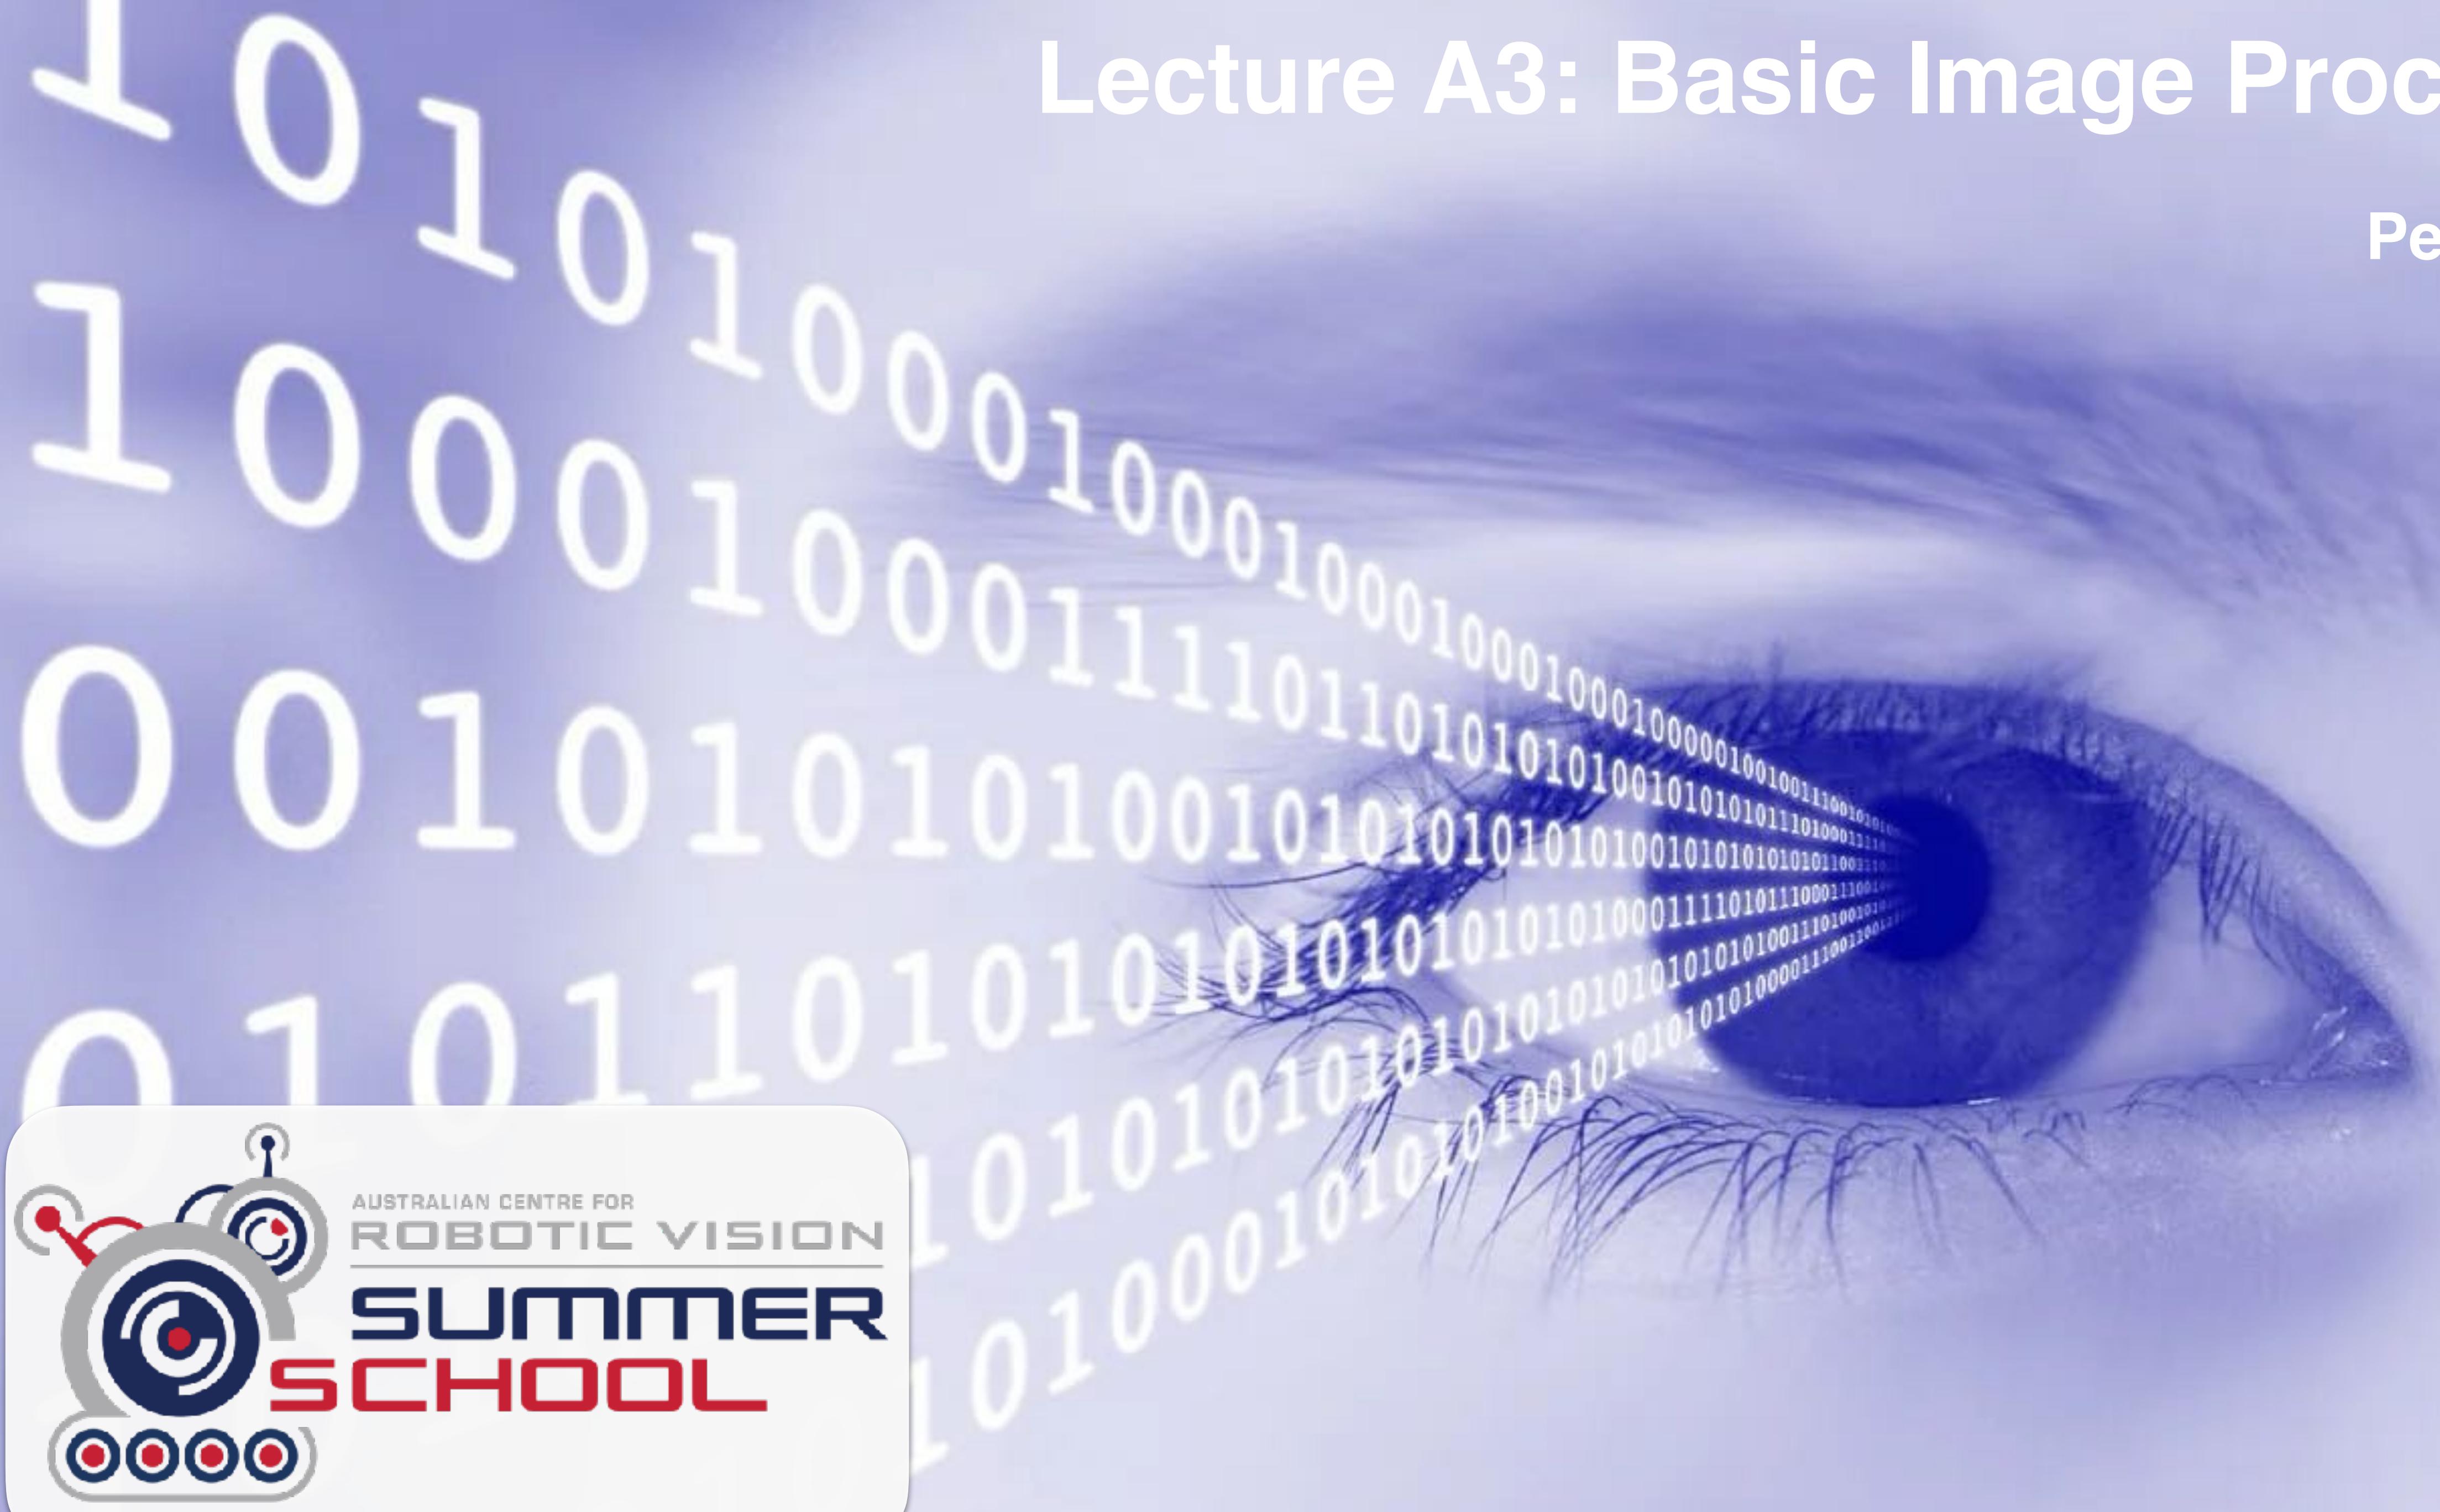
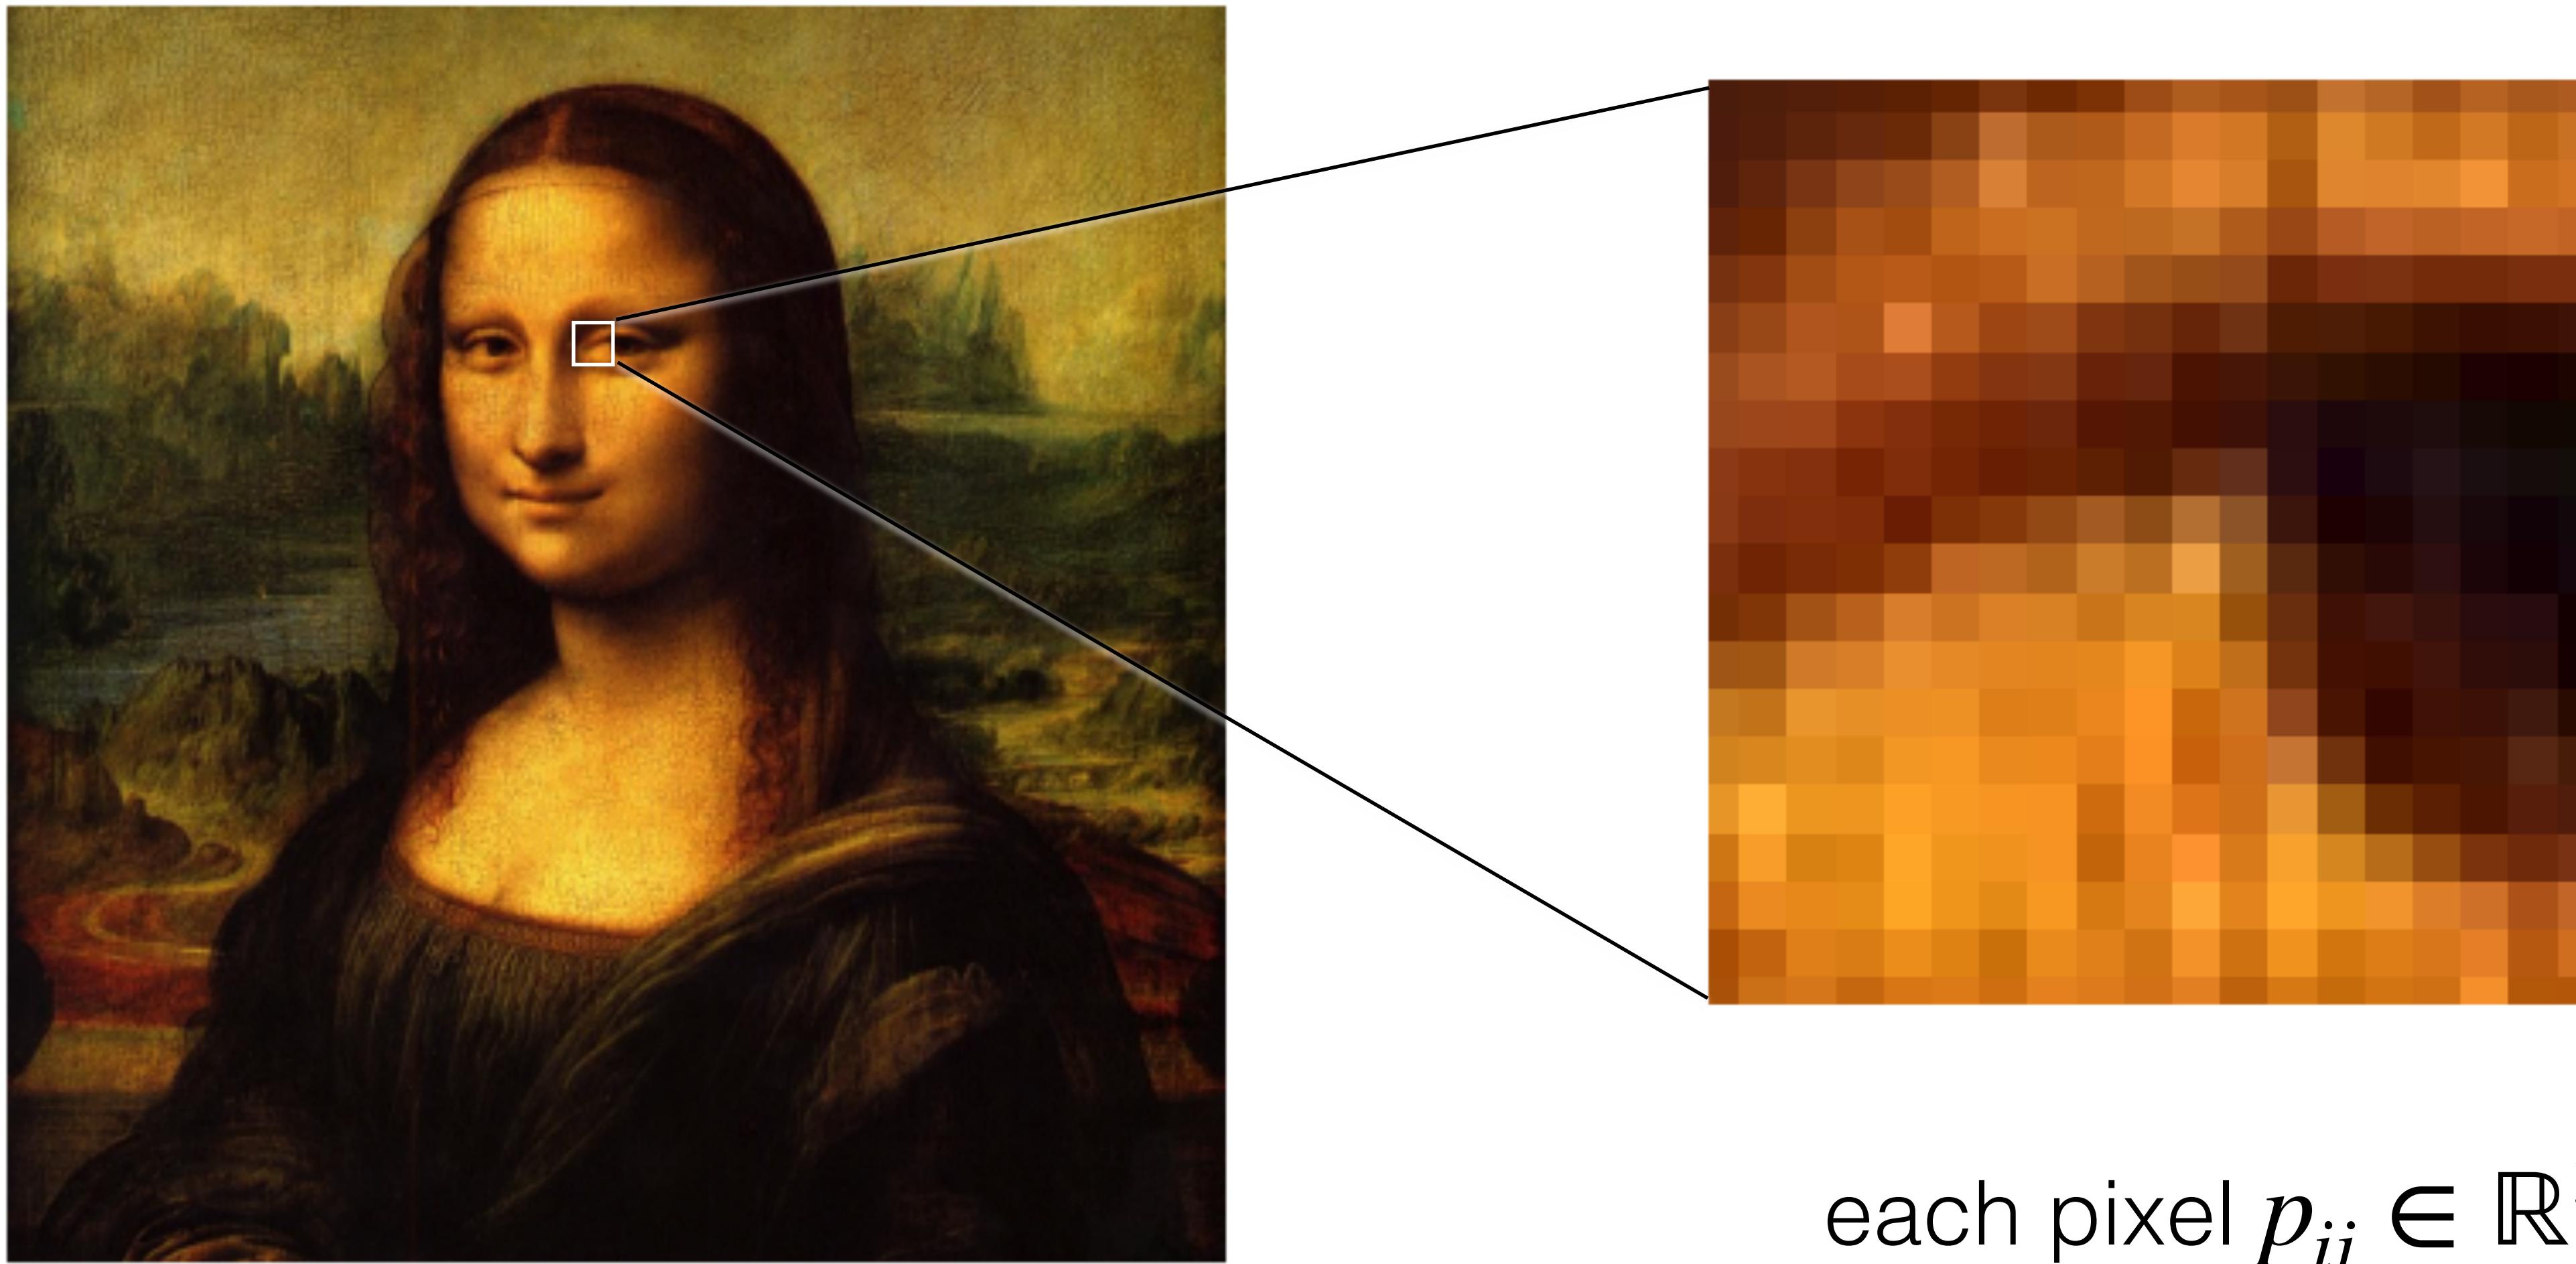


Lecture A3: Basic Image Processing

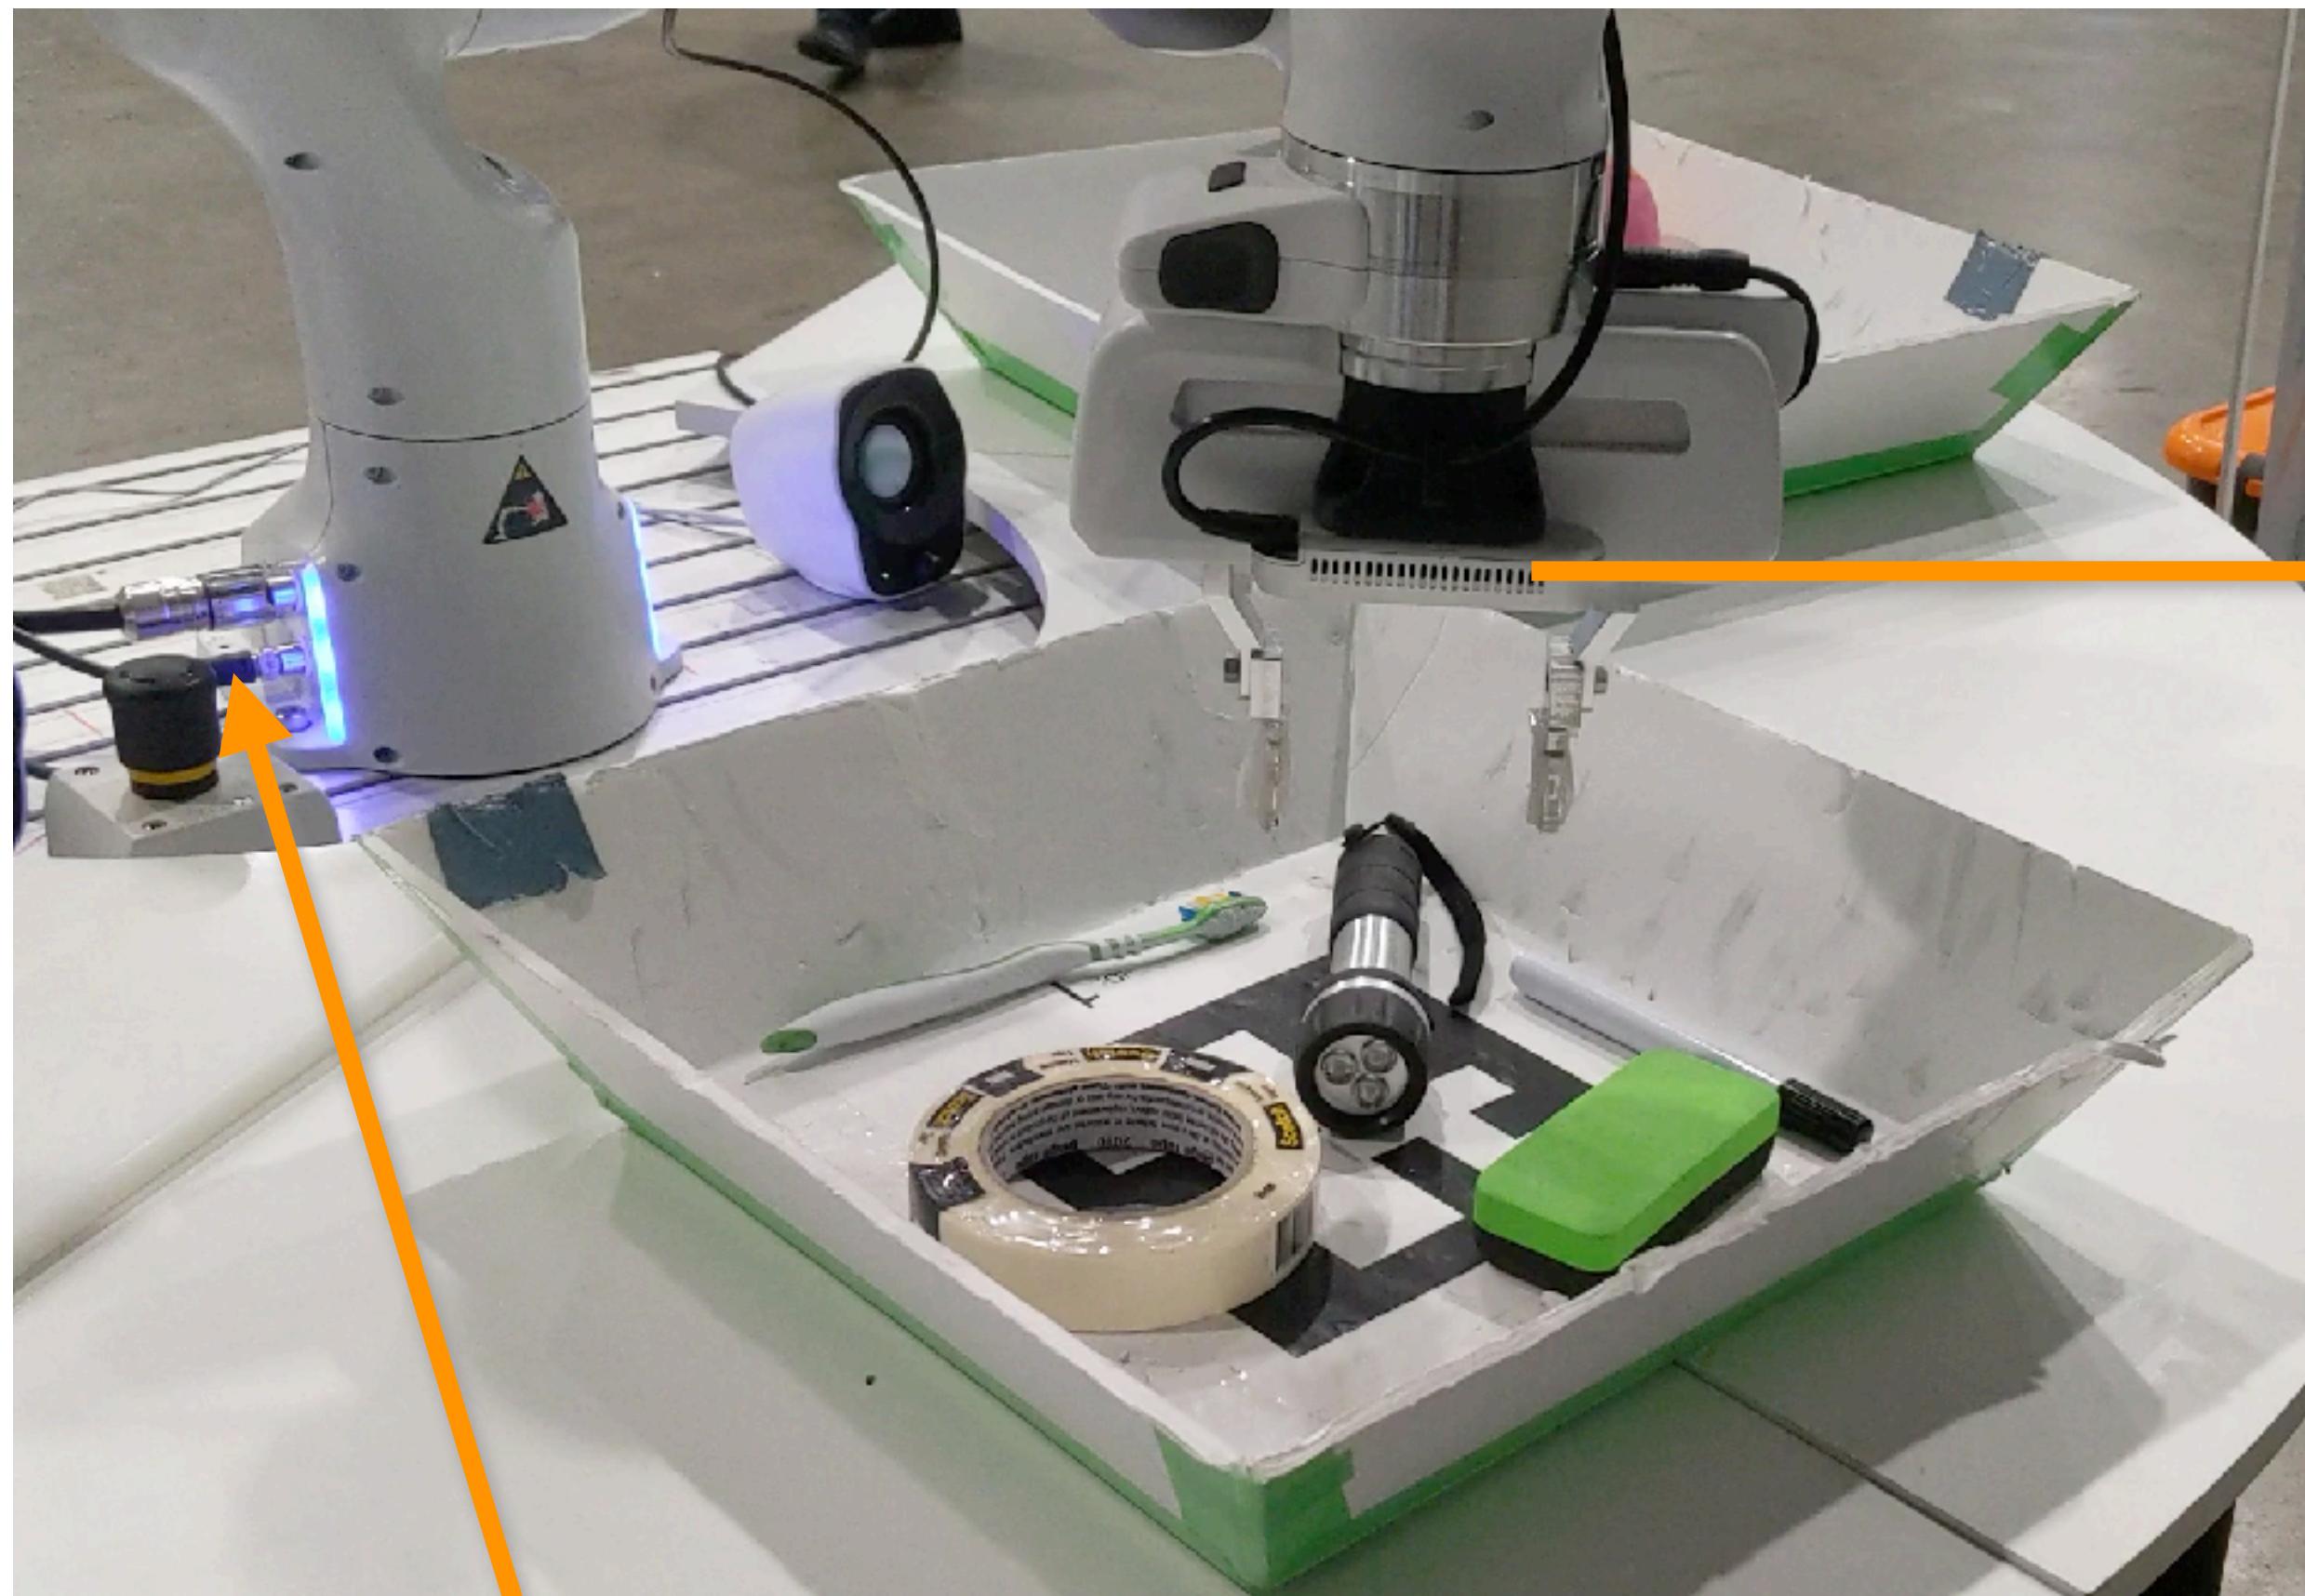
Peter Corke



A digital image comprises an array of square picture elements (pixels)



The data output rate of a camera and the requirements of a robot are fundamentally mismatched

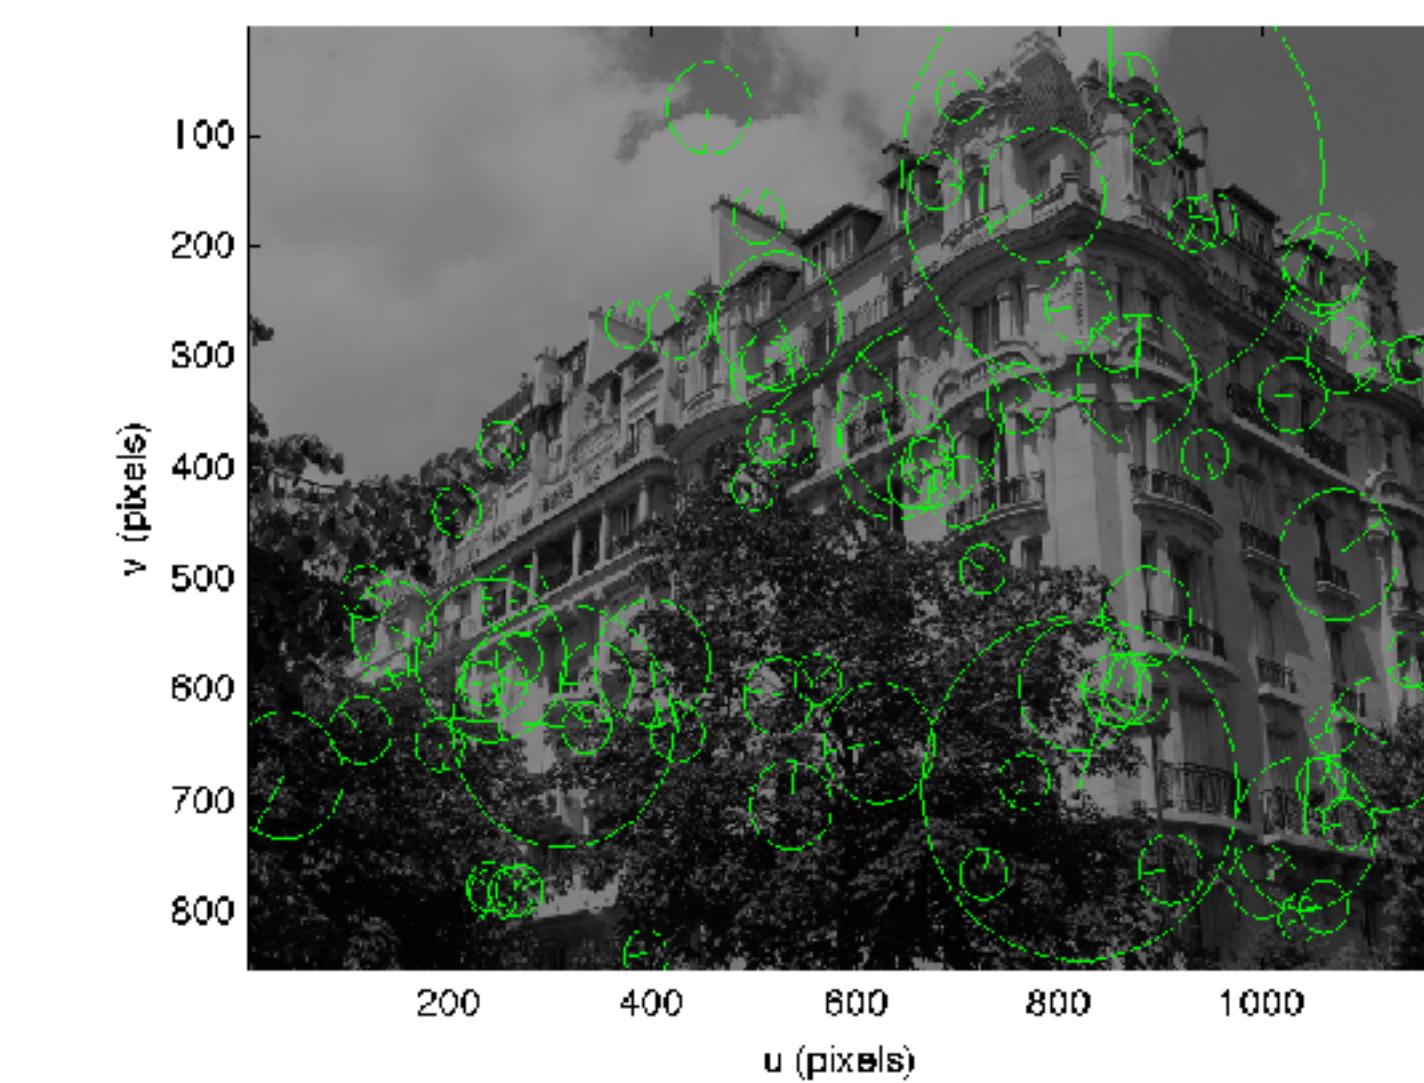
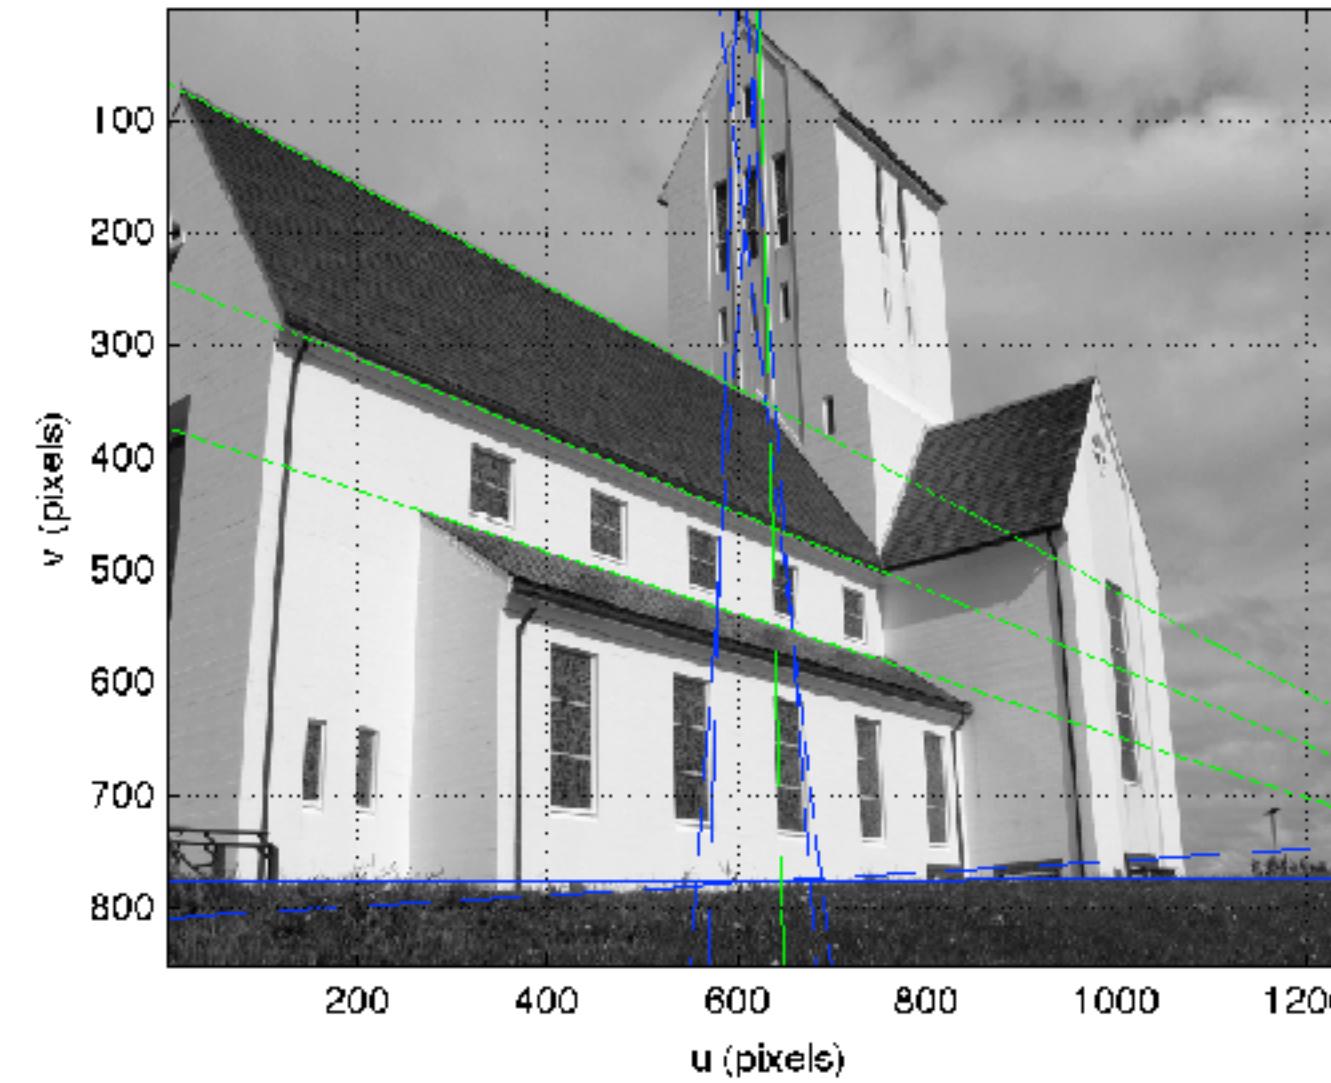
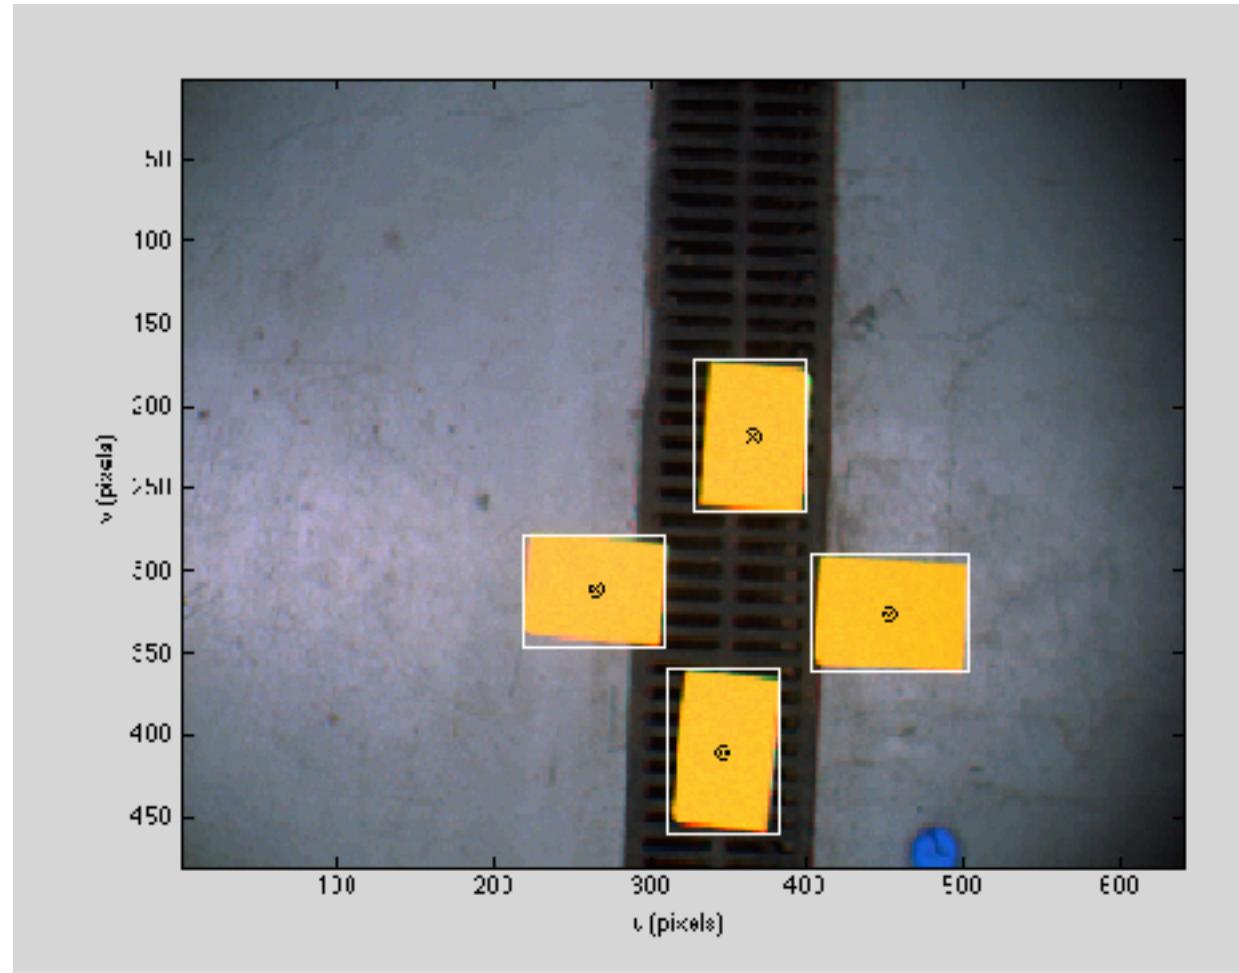


Robot input
~10 velocities at 10Hz
→ 400 byte/sec

Camera output
10 Mpix color image
at 30Hz
→ 900 Mbyte/sec

~ $\times 10^6$ mismatch

A classic approach to dealing with the data rate mismatch is to extract features from the image



Regions

- centroid coordinate
- size
- shape
- orientation
- etc.

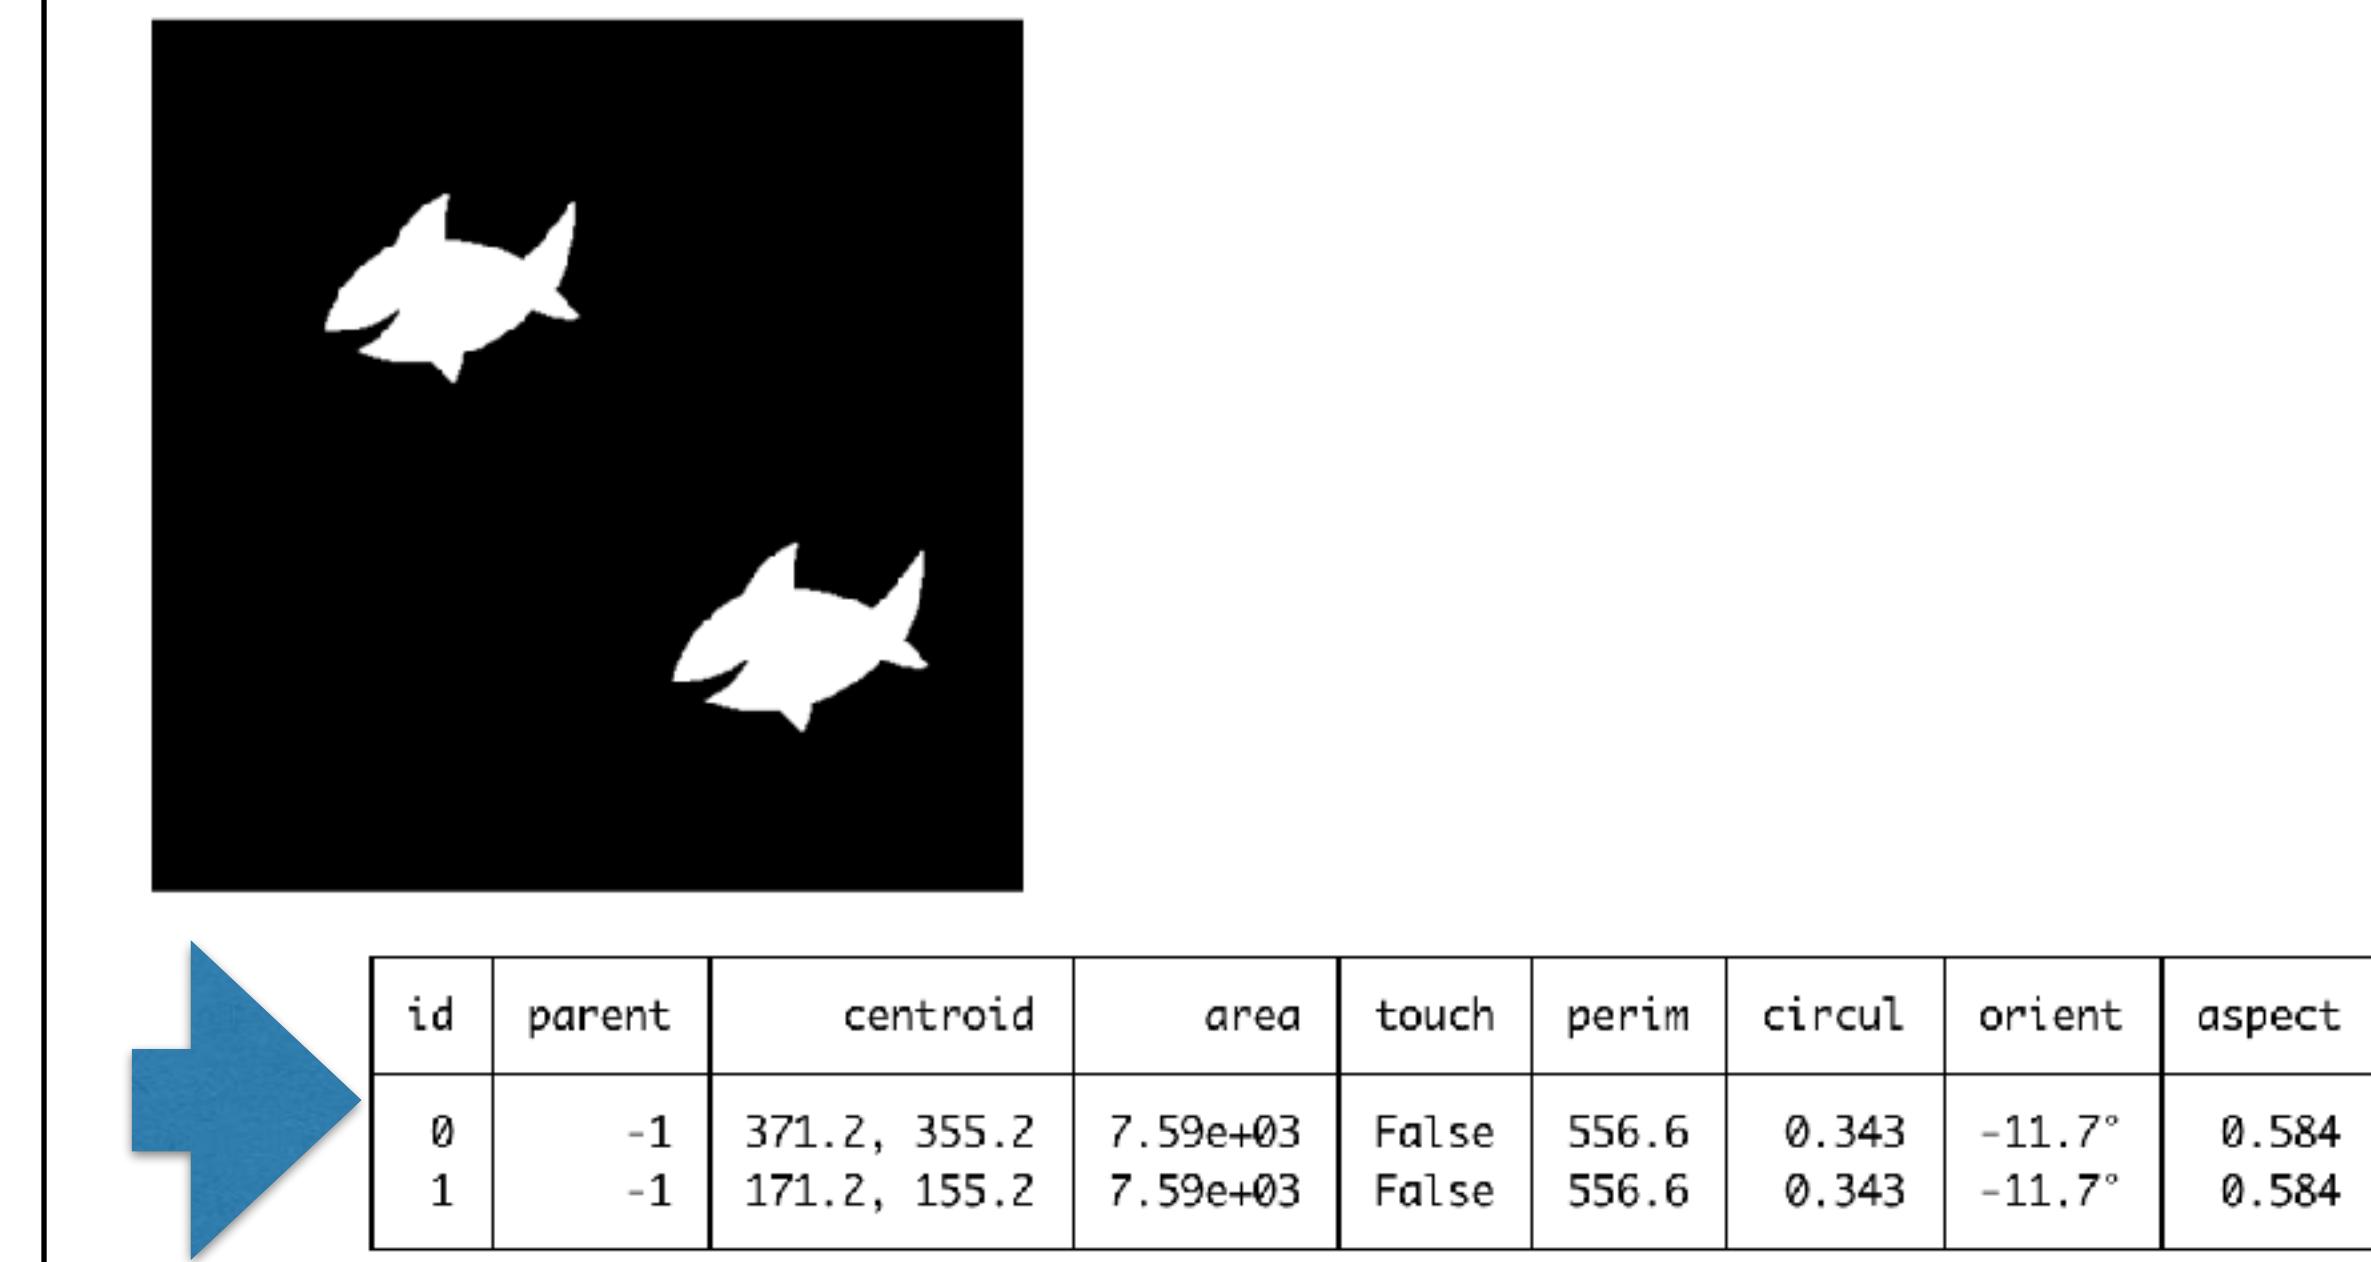
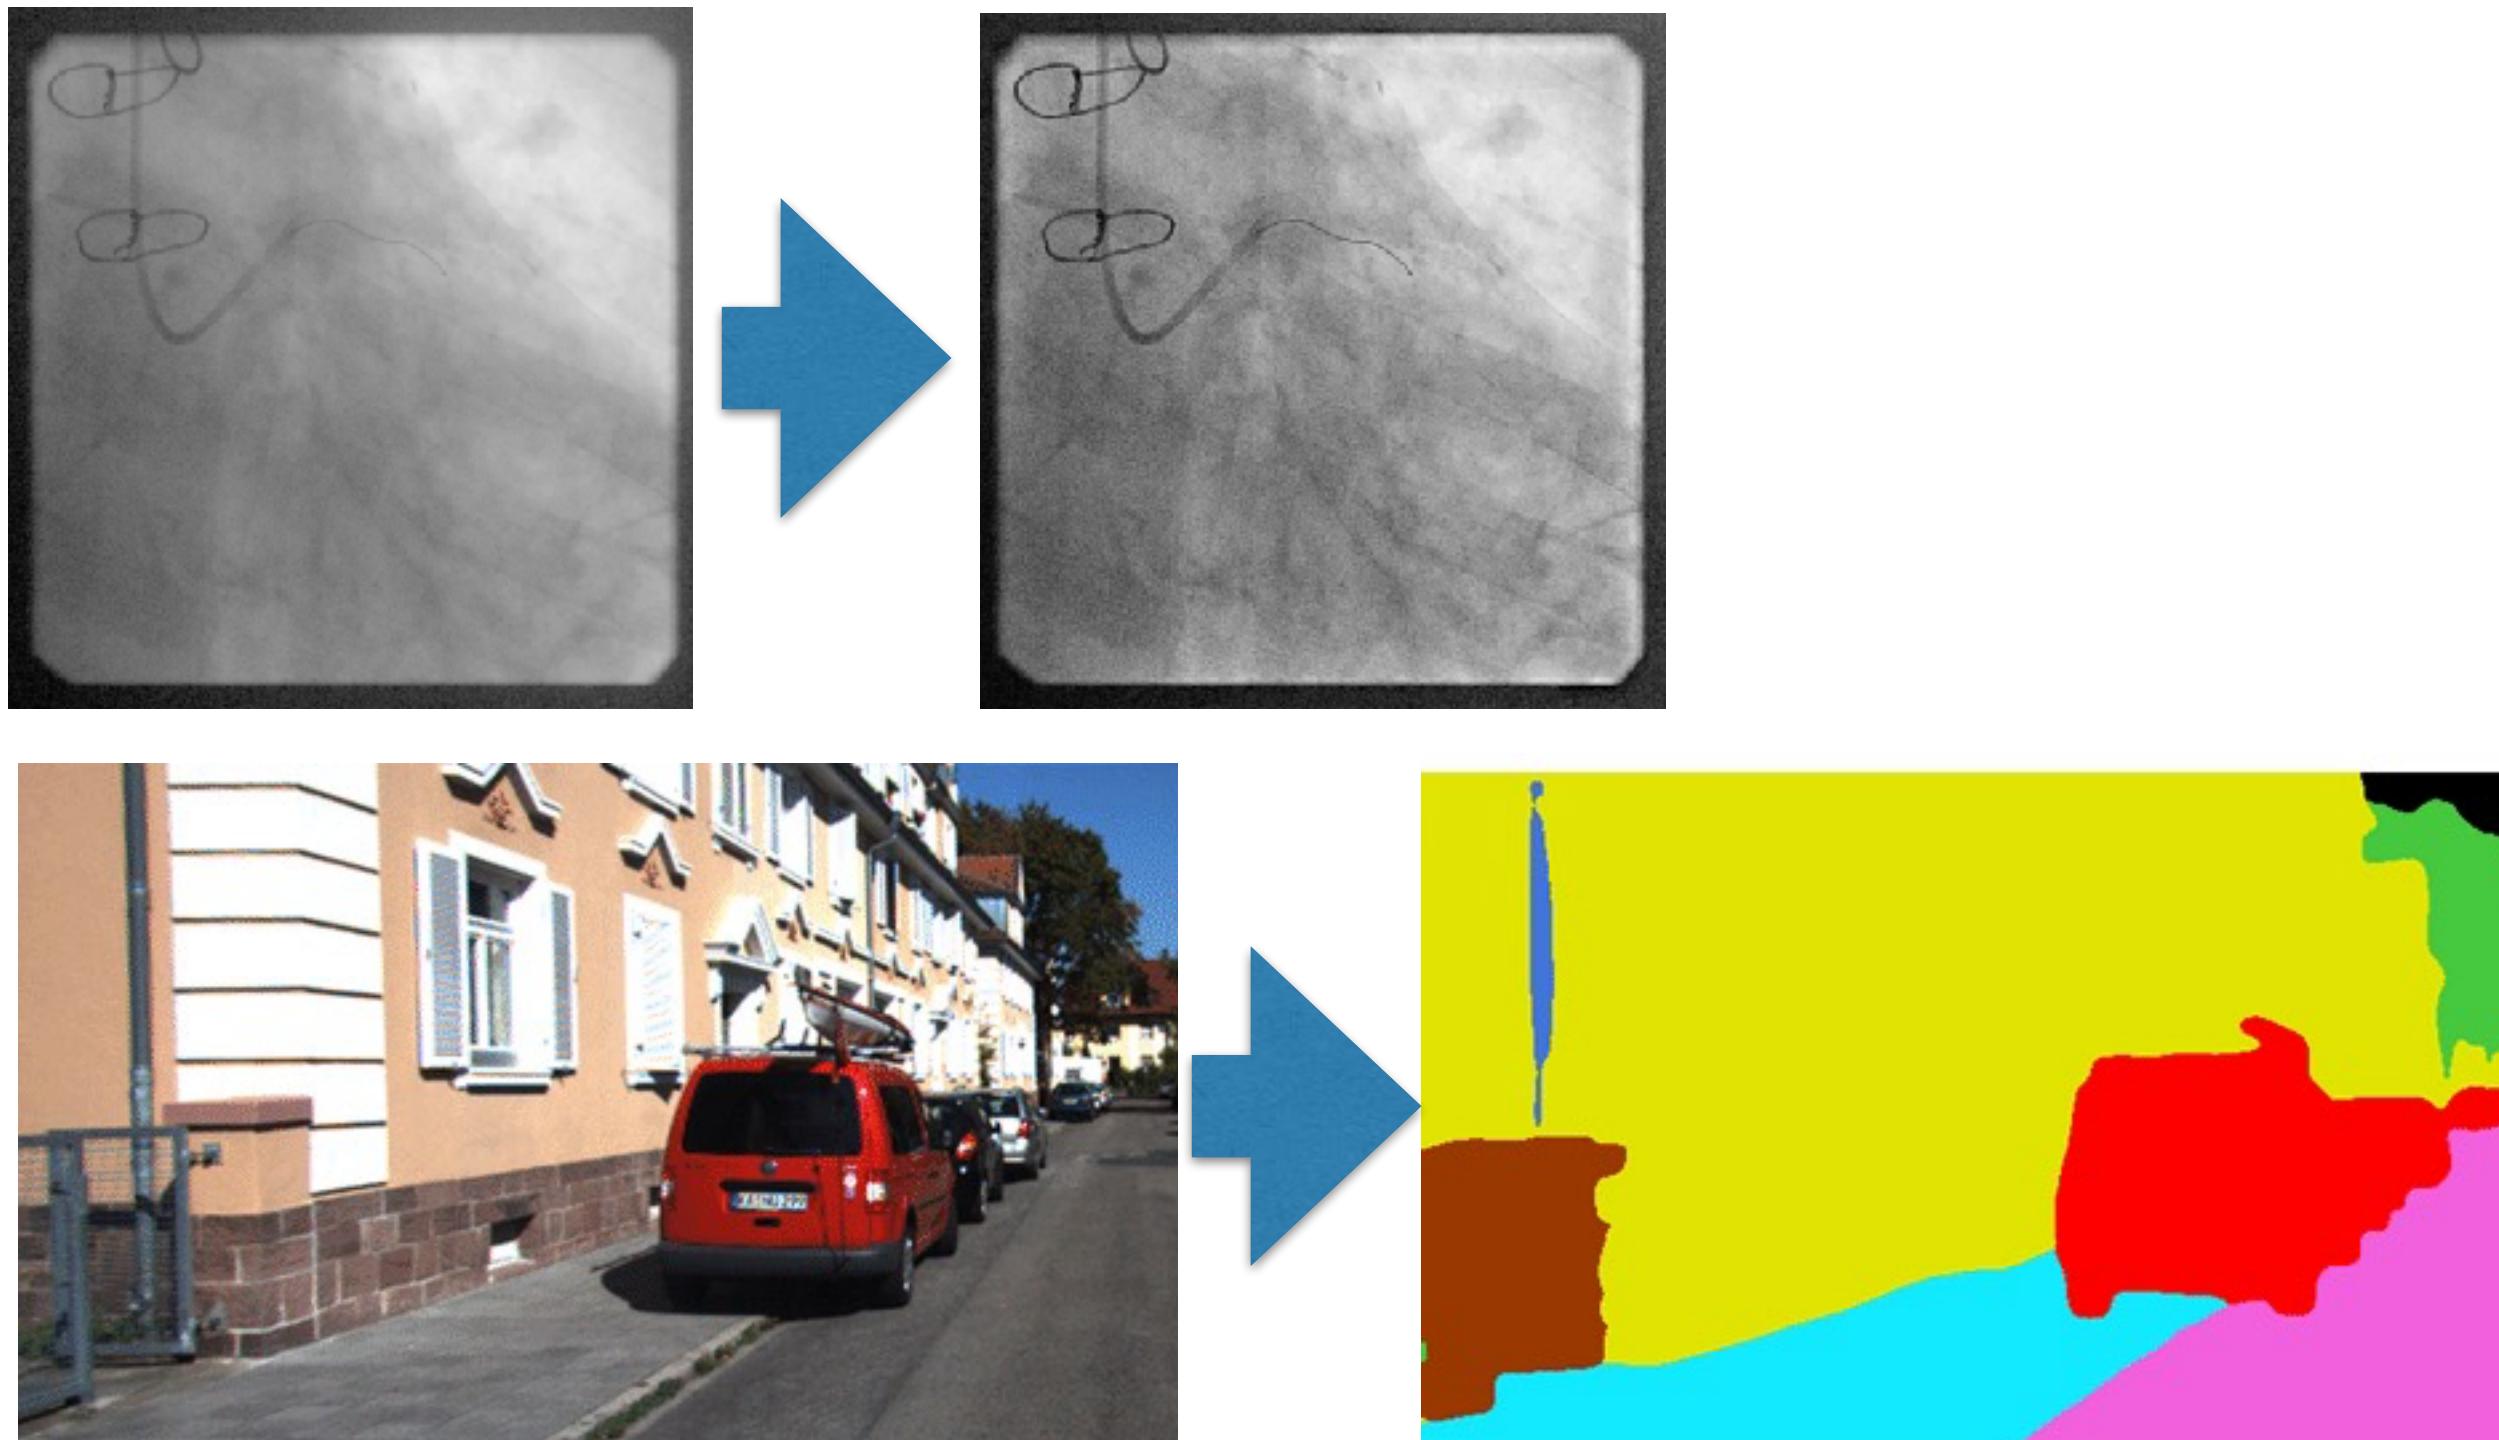
Lines

- equations of lines
- start/end of segments

Points

- coordinate
- scale

Image processing and computer vision are different things



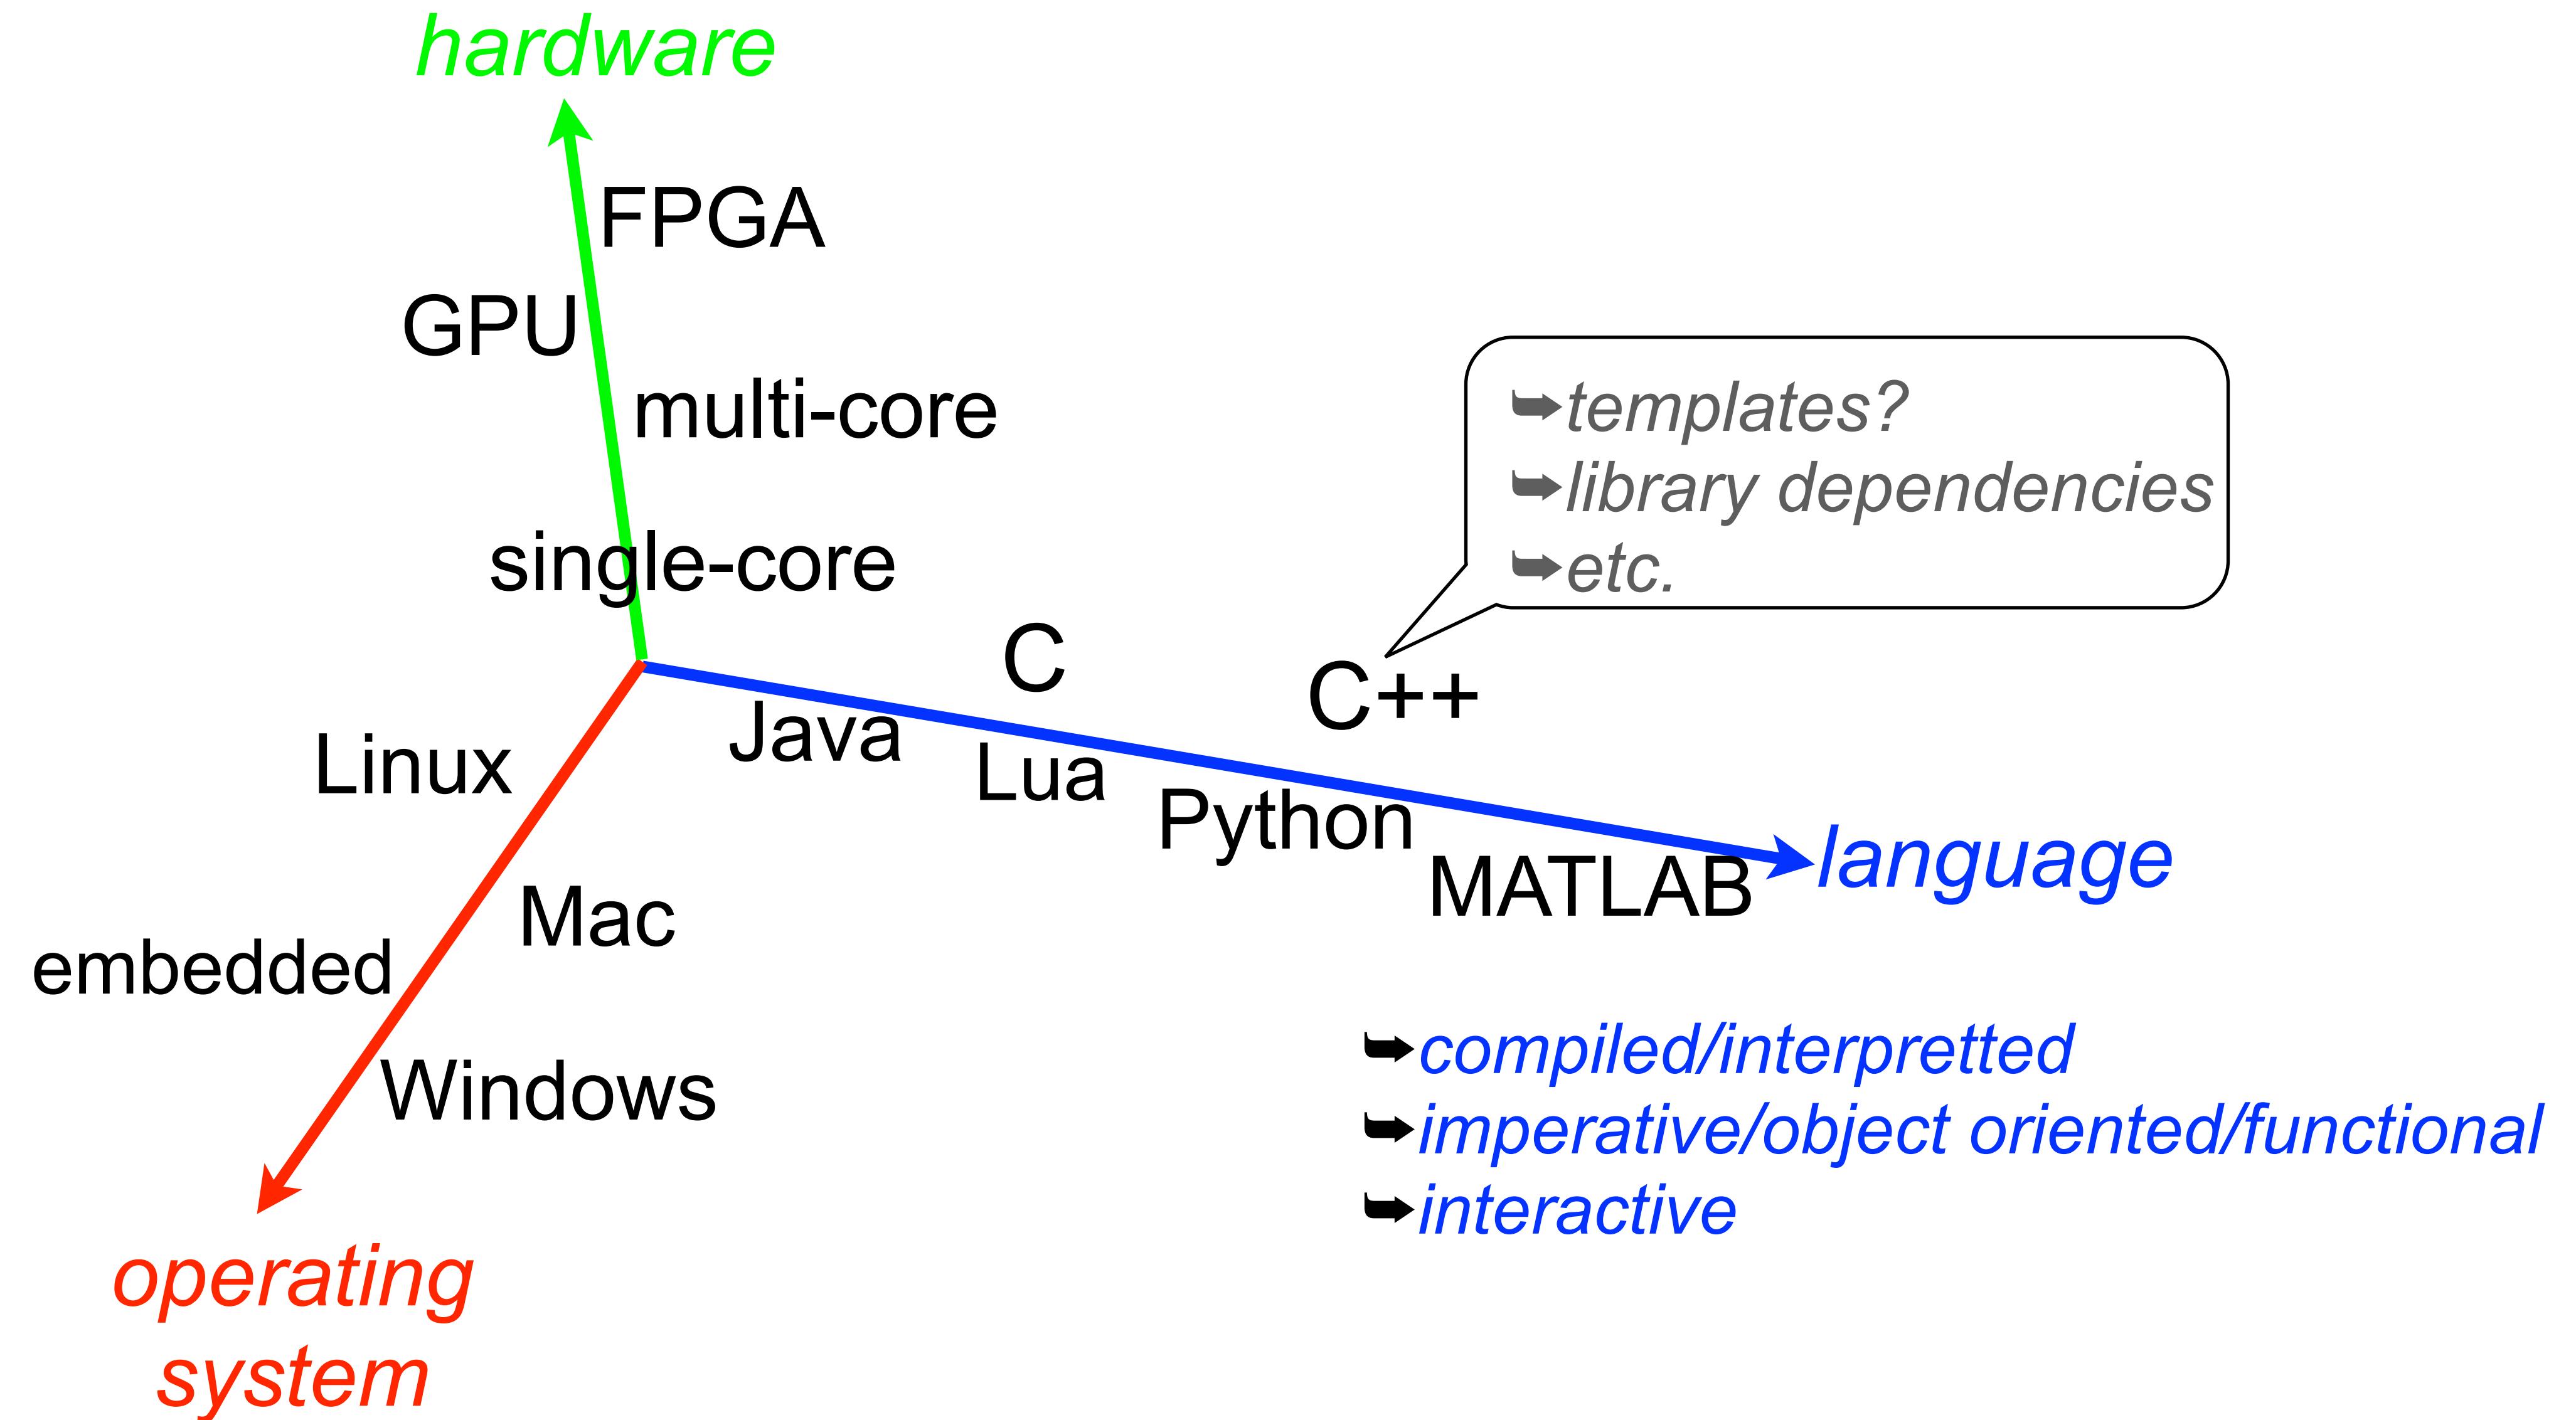
- **Image processing**
 - image → image
 - enhancement
 - semantic segmentation
- **Computer vision**
 - image → features
- **aka machine vision**

Technology for processing images

Over recent decades there have been hundreds of software implementations for computer vision (and image processing)

IMLAB IM Clmg VLFeat NeatVision
XVision VASARI pythonvision ViSP
Vista PIL OpenCV pixelvision
VXL vgg_ Gandalf CMVision Khoros
CVIPtools DIPlib ImageJ HIPS PBMplus
ImageLib TINA TargetJr Cpplma
ImageMagick

The design space has a number of axes



We have chosen to use Python and OpenCV

- Python
 - ✓ Common in teaching, research and industry
 - ✓ Runs on many platforms
 - ✓ With NumPy it eats matrices and vectors for breakfast
 - ✓ Interactive mode or programming mode
 - ✓ Great graphics
 - ✓ Object oriented
 - ✗ Slow
- OpenCV
 - ✓ Common in research and industry
 - ✓ Runs on many platforms
 - ✓ Has good Python wrappers
 - ✓ Highly optimised code
 - ✗ Not object oriented

For my part of the course we will be using a wrapped version of OpenCV - the Machine Vision Toolbox for Python

Machine Vision Toolbox for Python

collection PYTHONROBOTICS powered by spatial maths collection QUT Robotics

pypi package 1.0.0 python 3.9 | 3.10 | 3.11 | 3.12 powered by OpenCV 3.1.0 powered by Open3D 0.4.0

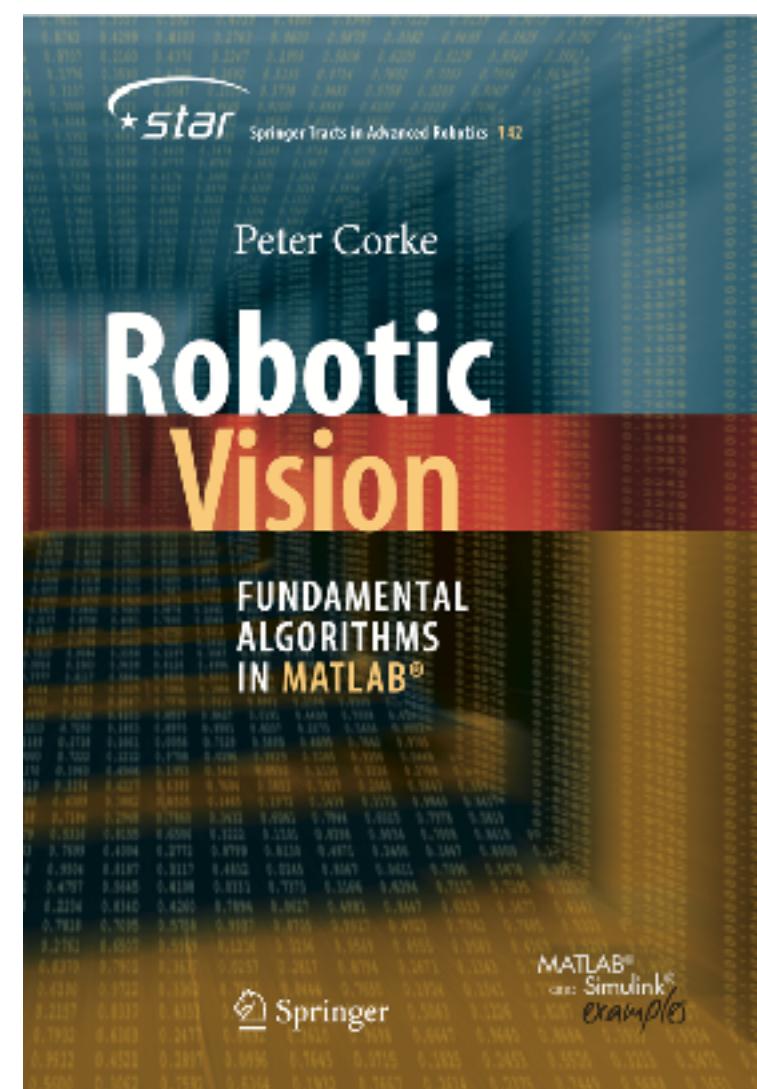
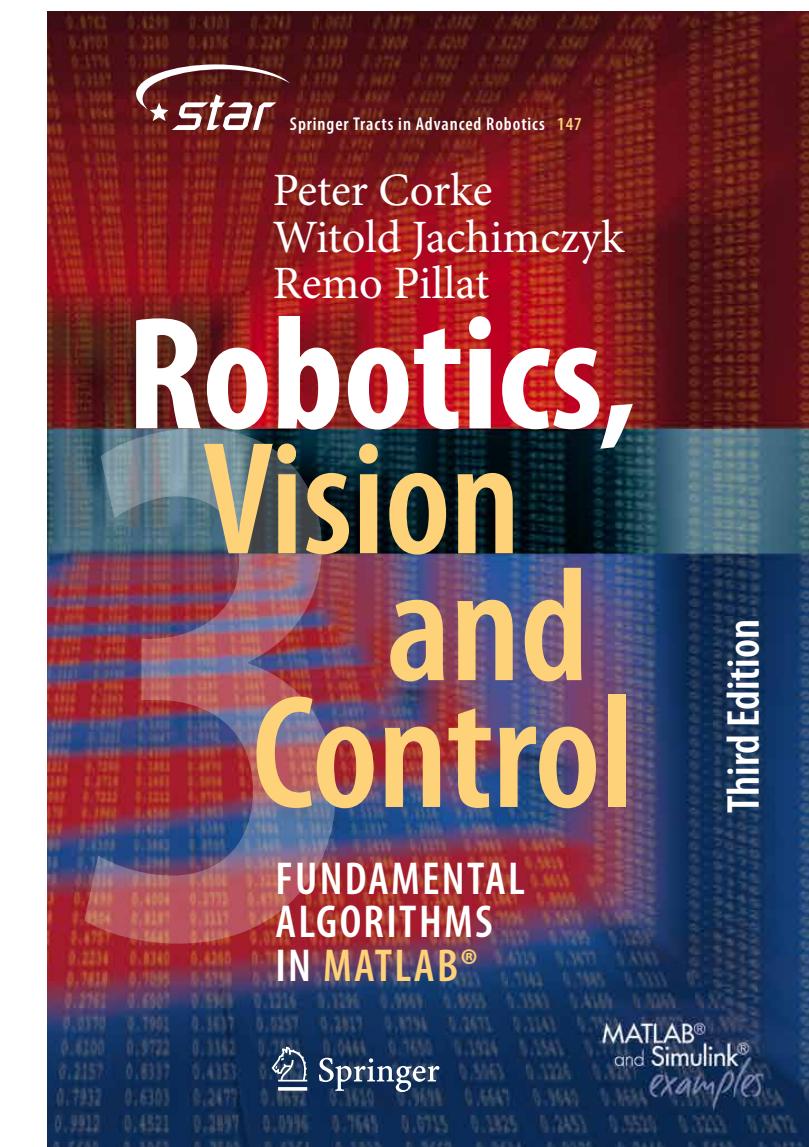
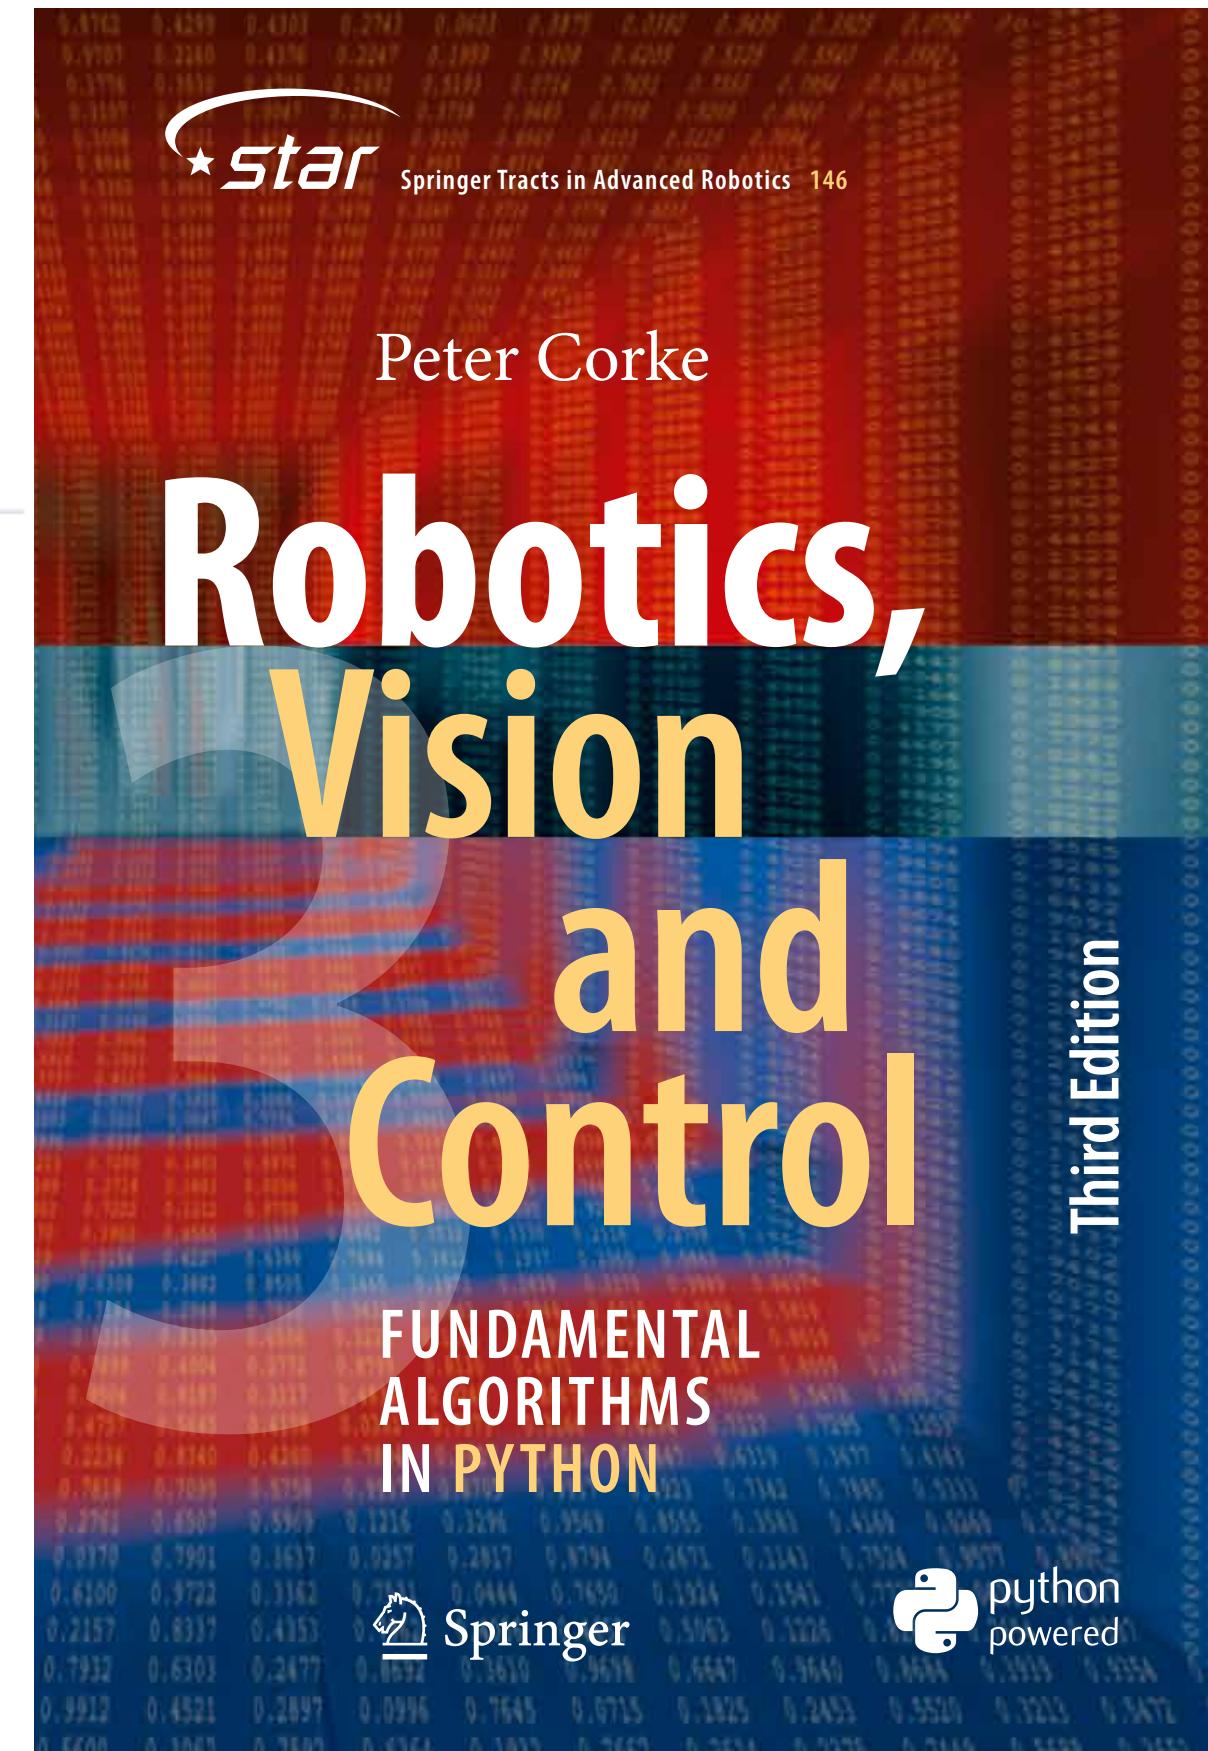
License MIT

mvtb-main passing codecov 32% downloads 122/week

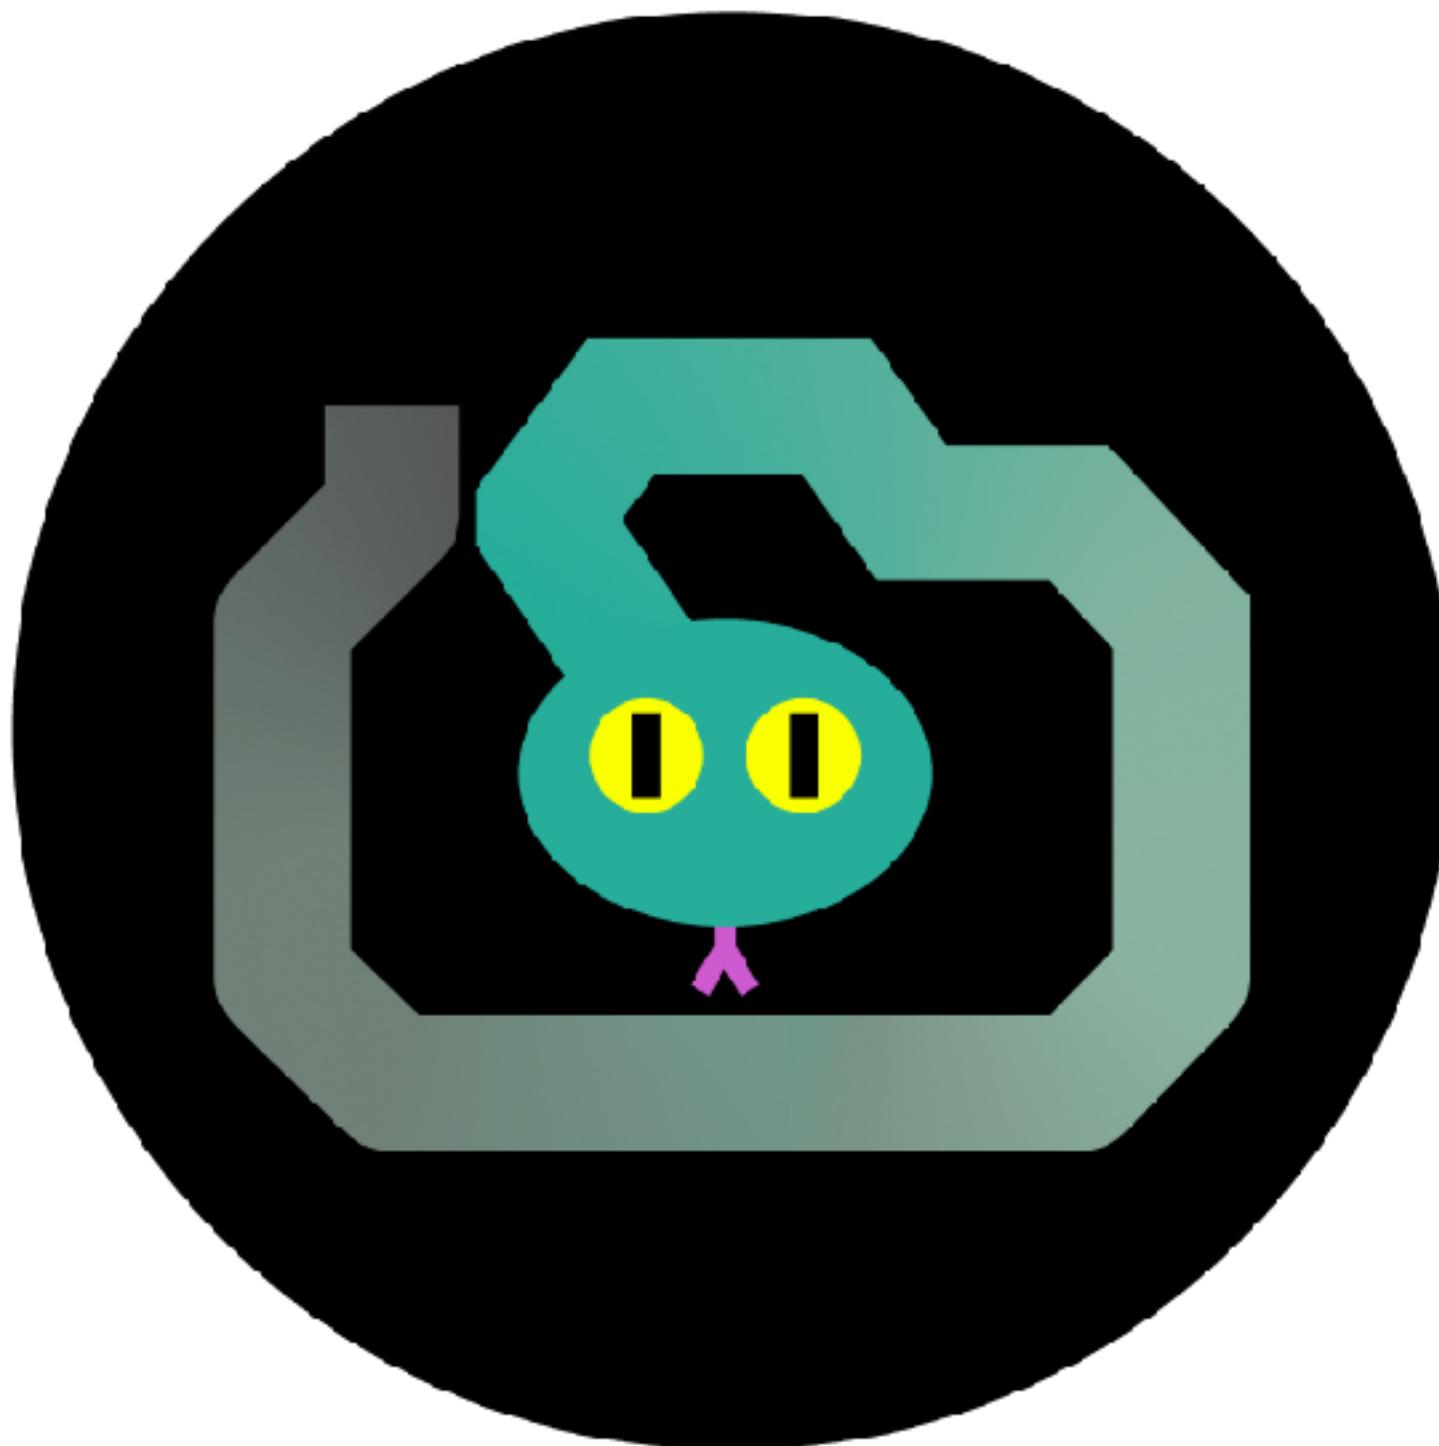


- [GitHub repository](#)
- [Documentation](#)
- [Examples and details](#)
- [Installation](#)
- [Examples and details](#)
- [Changelog](#)

A Python implementation of the [Machine Vision Toolbox for MATLAB®](#)



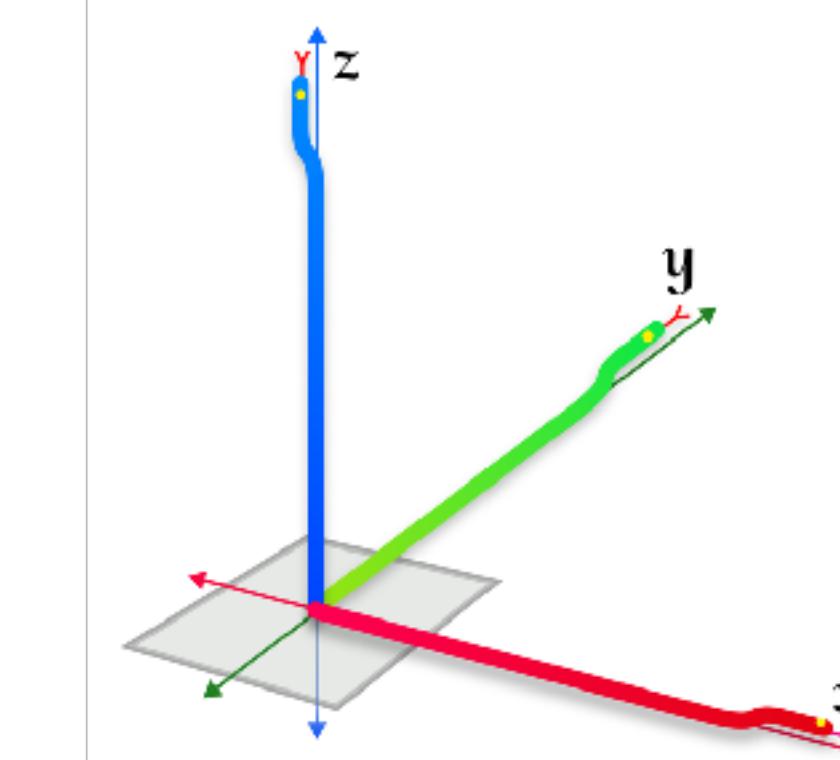
The install is easy



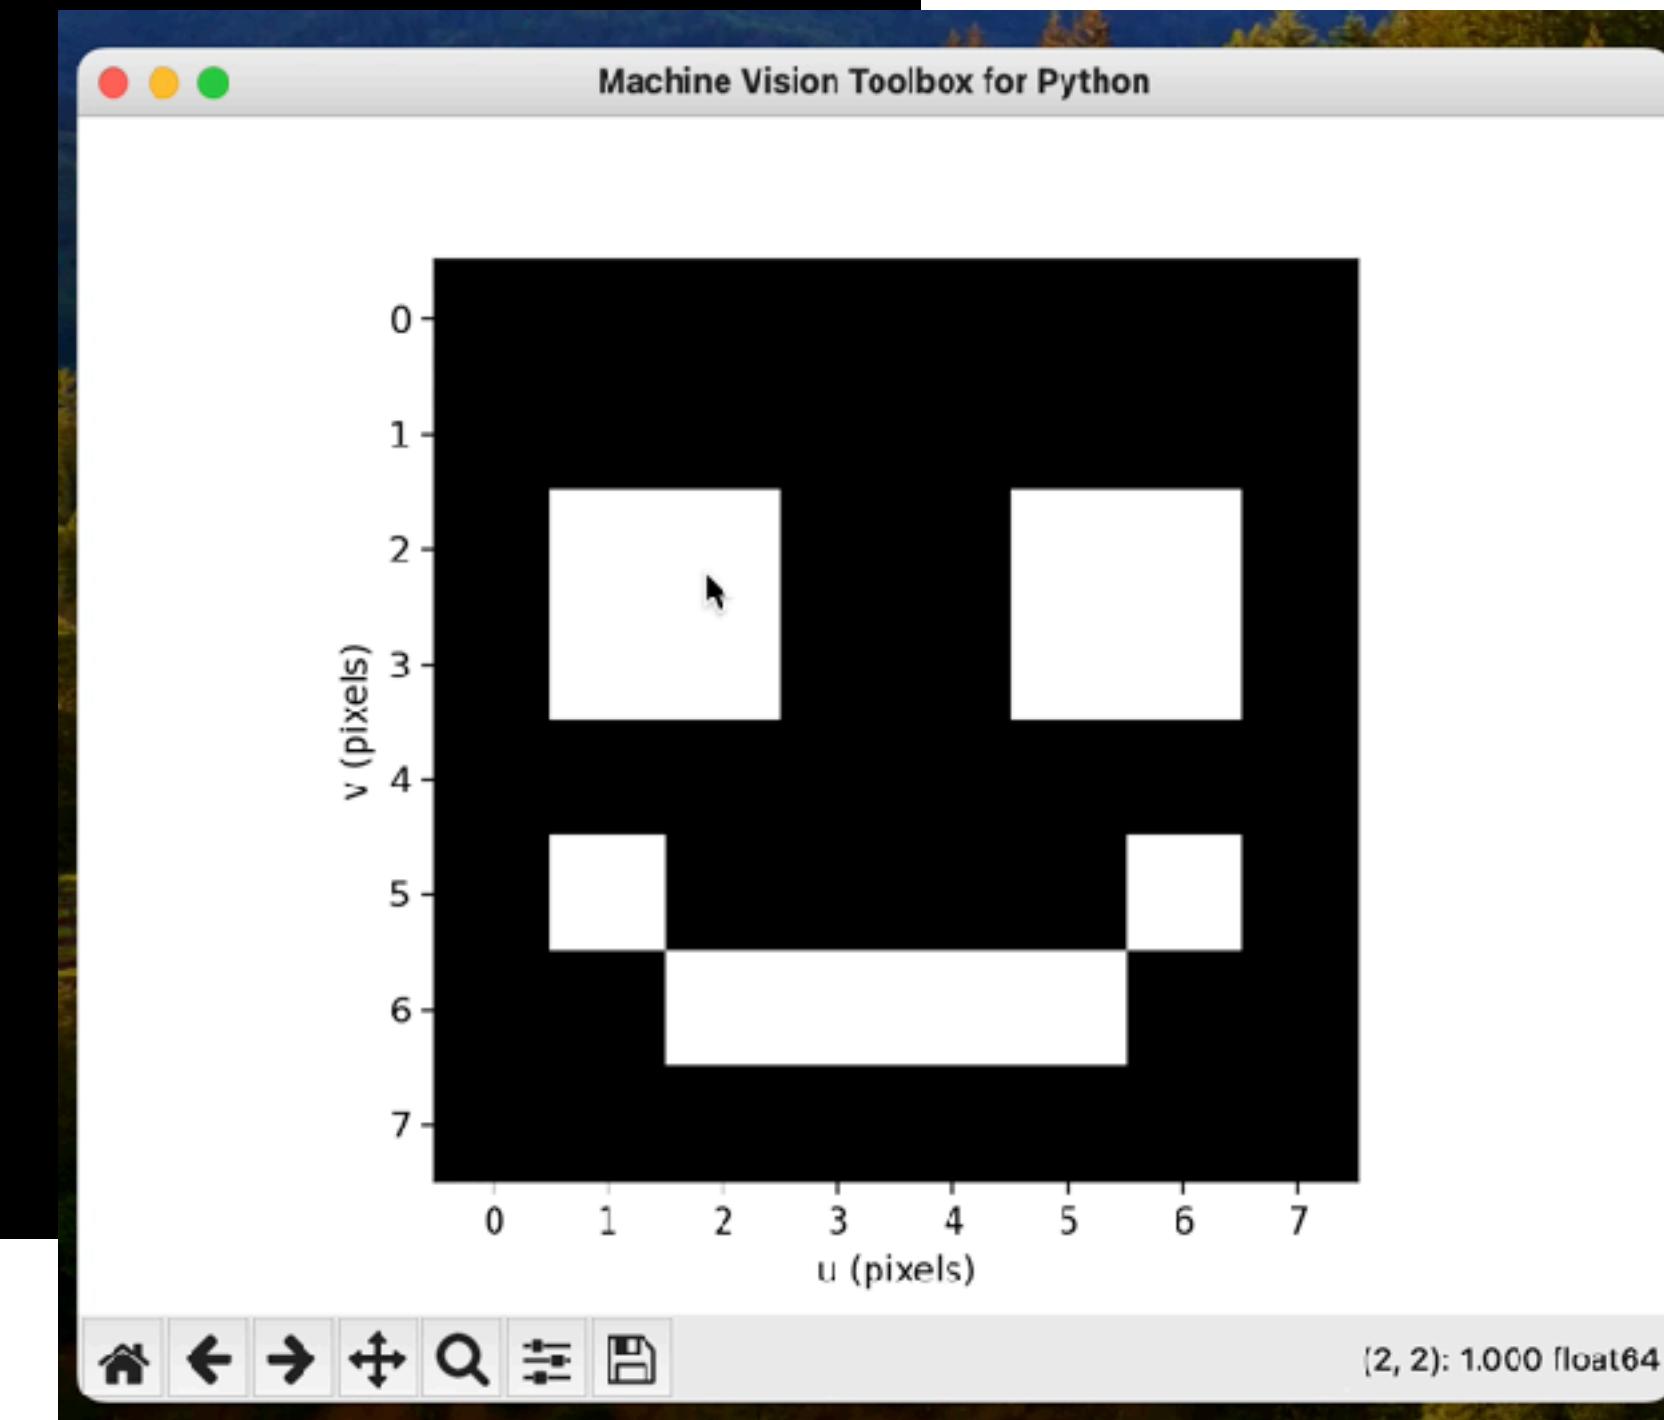
`pip install machinevision-toolbox-python`

and it also pulls in:

- `mvtb-data`
- `spatialmath-python`
- `OpenCV`
- `Open3D`



An image is simply a 2D array of numbers



An image is simply a 2D array of numbers

```
from machinevisiontoolbox import *

street, _ = iread("street.png")

print(street.shape)

(851, 1280)

print(street[125:140, 295:310])

[196 195 197 198 198 199 197 199 199 201 199 197 197 200 201]
[194 193 195 197 198 199 198 198 200 202 201 196 195 199 198]
[194 195 194 194 196 195 197 200 199 198 200 198 199 200 200]
[194 202 199 196 194 195 198 198 197 196 200 195 198 198 199]
[ 91 130 170 200 205 204 198 194 195 197 198 192 198 199 200]
[ 29  31  48  89 122 156 191 202 199 198 198 194 196 199 199]
[ 30  32  31  27  30  41  67 115 161 183 194 194 196 201 200]
[ 25  27  27  28  30  29  26  26  40  70 110 170 196 195 198]
[ 21  22  21  22  23  24  27  29  28  27  78 169 200 196 197]
[ 18  19  20  20  22  23  23  23  28  30 101 190 195 194 198]
[ 17  18  19  19  20  22  22  24  29  32  91 181 195 196 197]
[ 15  17  18  19  20  23  24  27  29  34 102 183 195 197 197]
[ 15  18  19  18  19  22  23  26  33  40 101 190 197 196 198]
[ 15  17  17  18  19  22  23  26  32  41 103 183 192 197 197]
[ 16  17  17  17  20  21  23  26  35  42 106 197 194 193 196]]
```

```
idisp(street, block=True)
```



There are advantages in making the images into Python objects

```
from machinevisiontoolbox import *
street, _ = iread("street.png")
print(street.shape)
(851, 1280)
print(street[125:140, 295:310])
[[196 195 197 ...
idisp(street, block=True)
```

```
from machinevisiontoolbox import *
street = Image.Read("street.png")
print(street)
Image: 1280 x 851 (uint8) [.../images/street.png]
print(street.shape)
(851, 1280)
print(street[125:140, 295:310].A)
[[196 195 197 ...
street.disp(block=True)
```

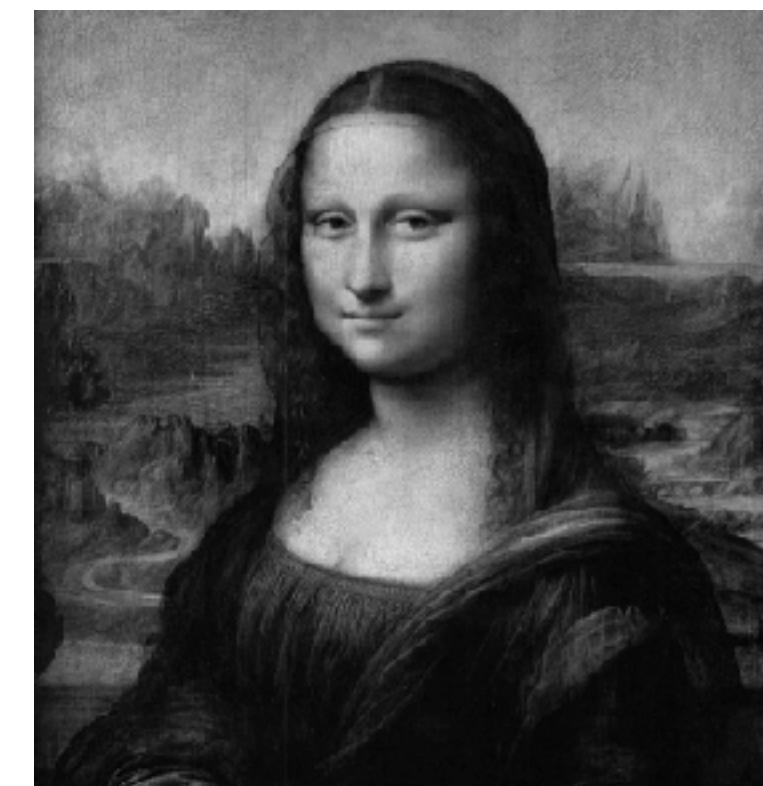
Python awesomeness (language, tools, support) plus the speed of OpenCV

There are lots of ways to acquire an image to work with

```
from machinevisiontoolbox import Image
```

```
mona = Image.Read("monalisa.png")  
mona.disp()
```

read from a file

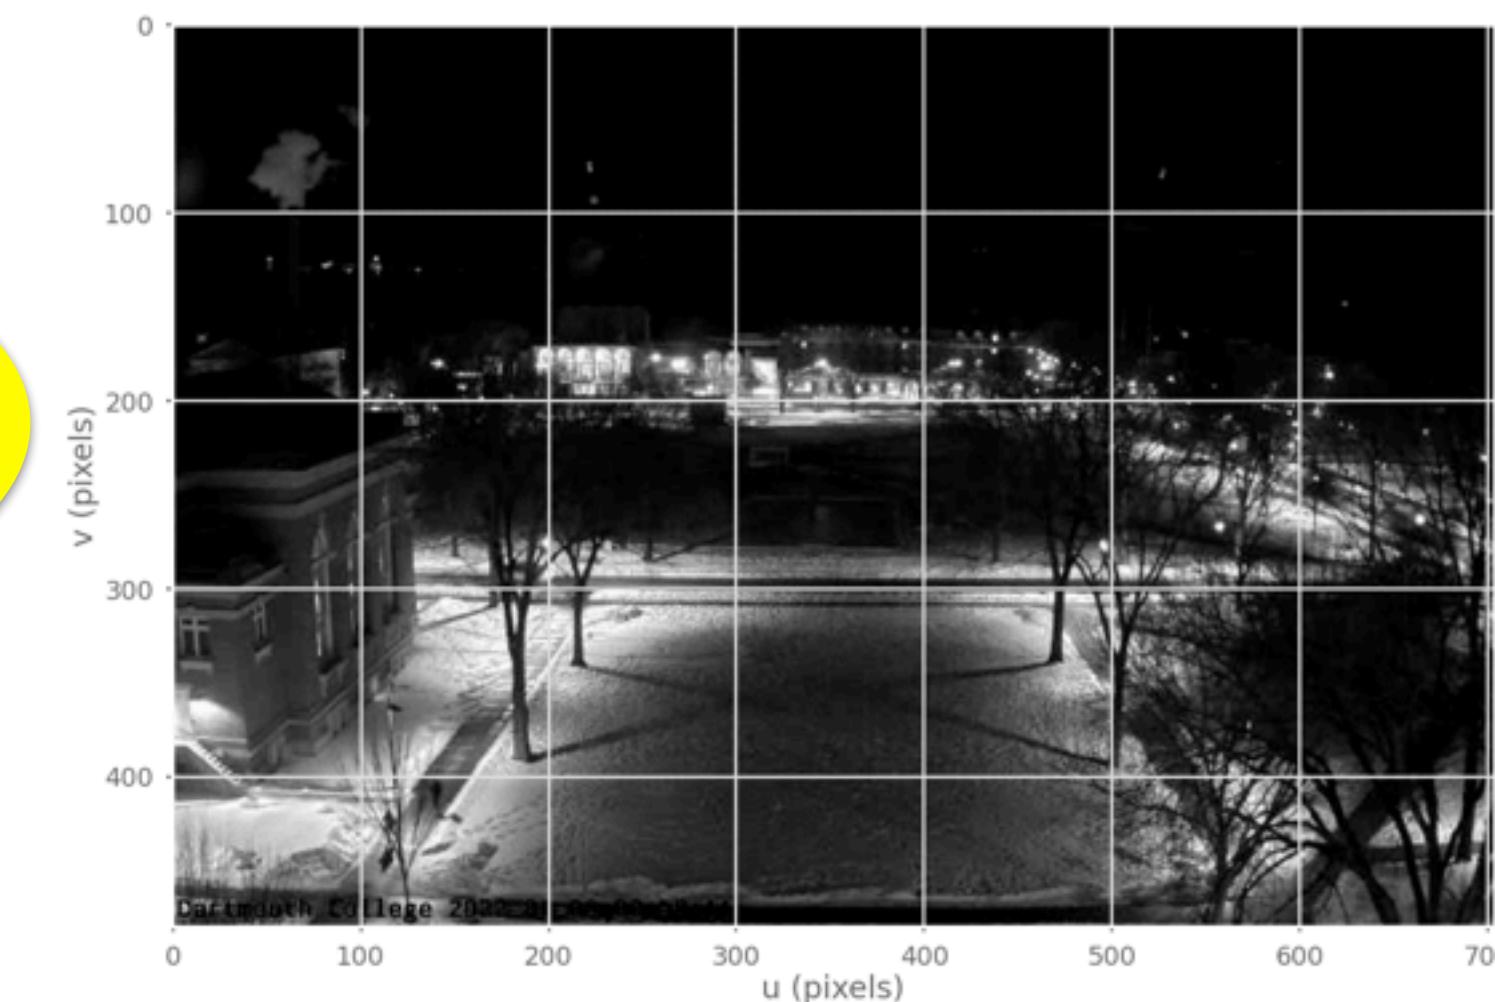


```
camera = VideoCamera()  
camera.grab().disp()
```

read from
an attached camera
via OpenCV

```
porjus = WebCam("https://uk.jokkmokk.jp/photo/nr4/latest.jpg")  
porjus.grab().disp()
```

read from
somebody else's
camera



```
im = Image.Random(20).disp()
```

create it using
code

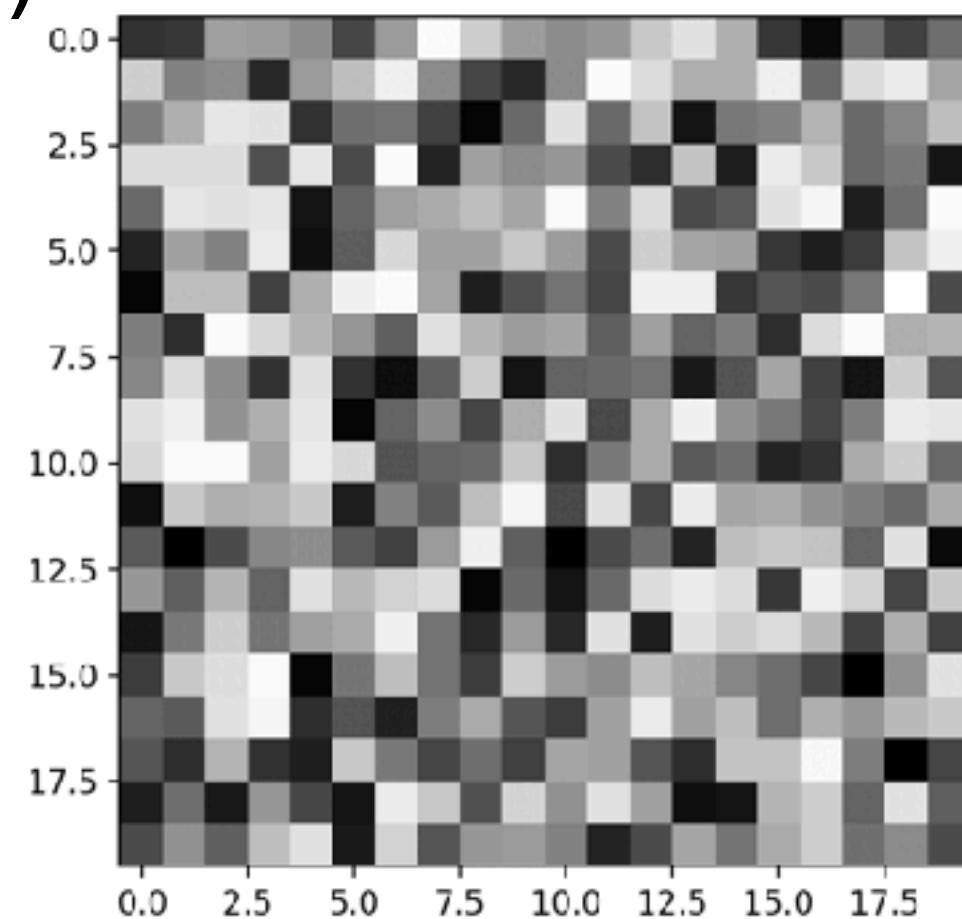
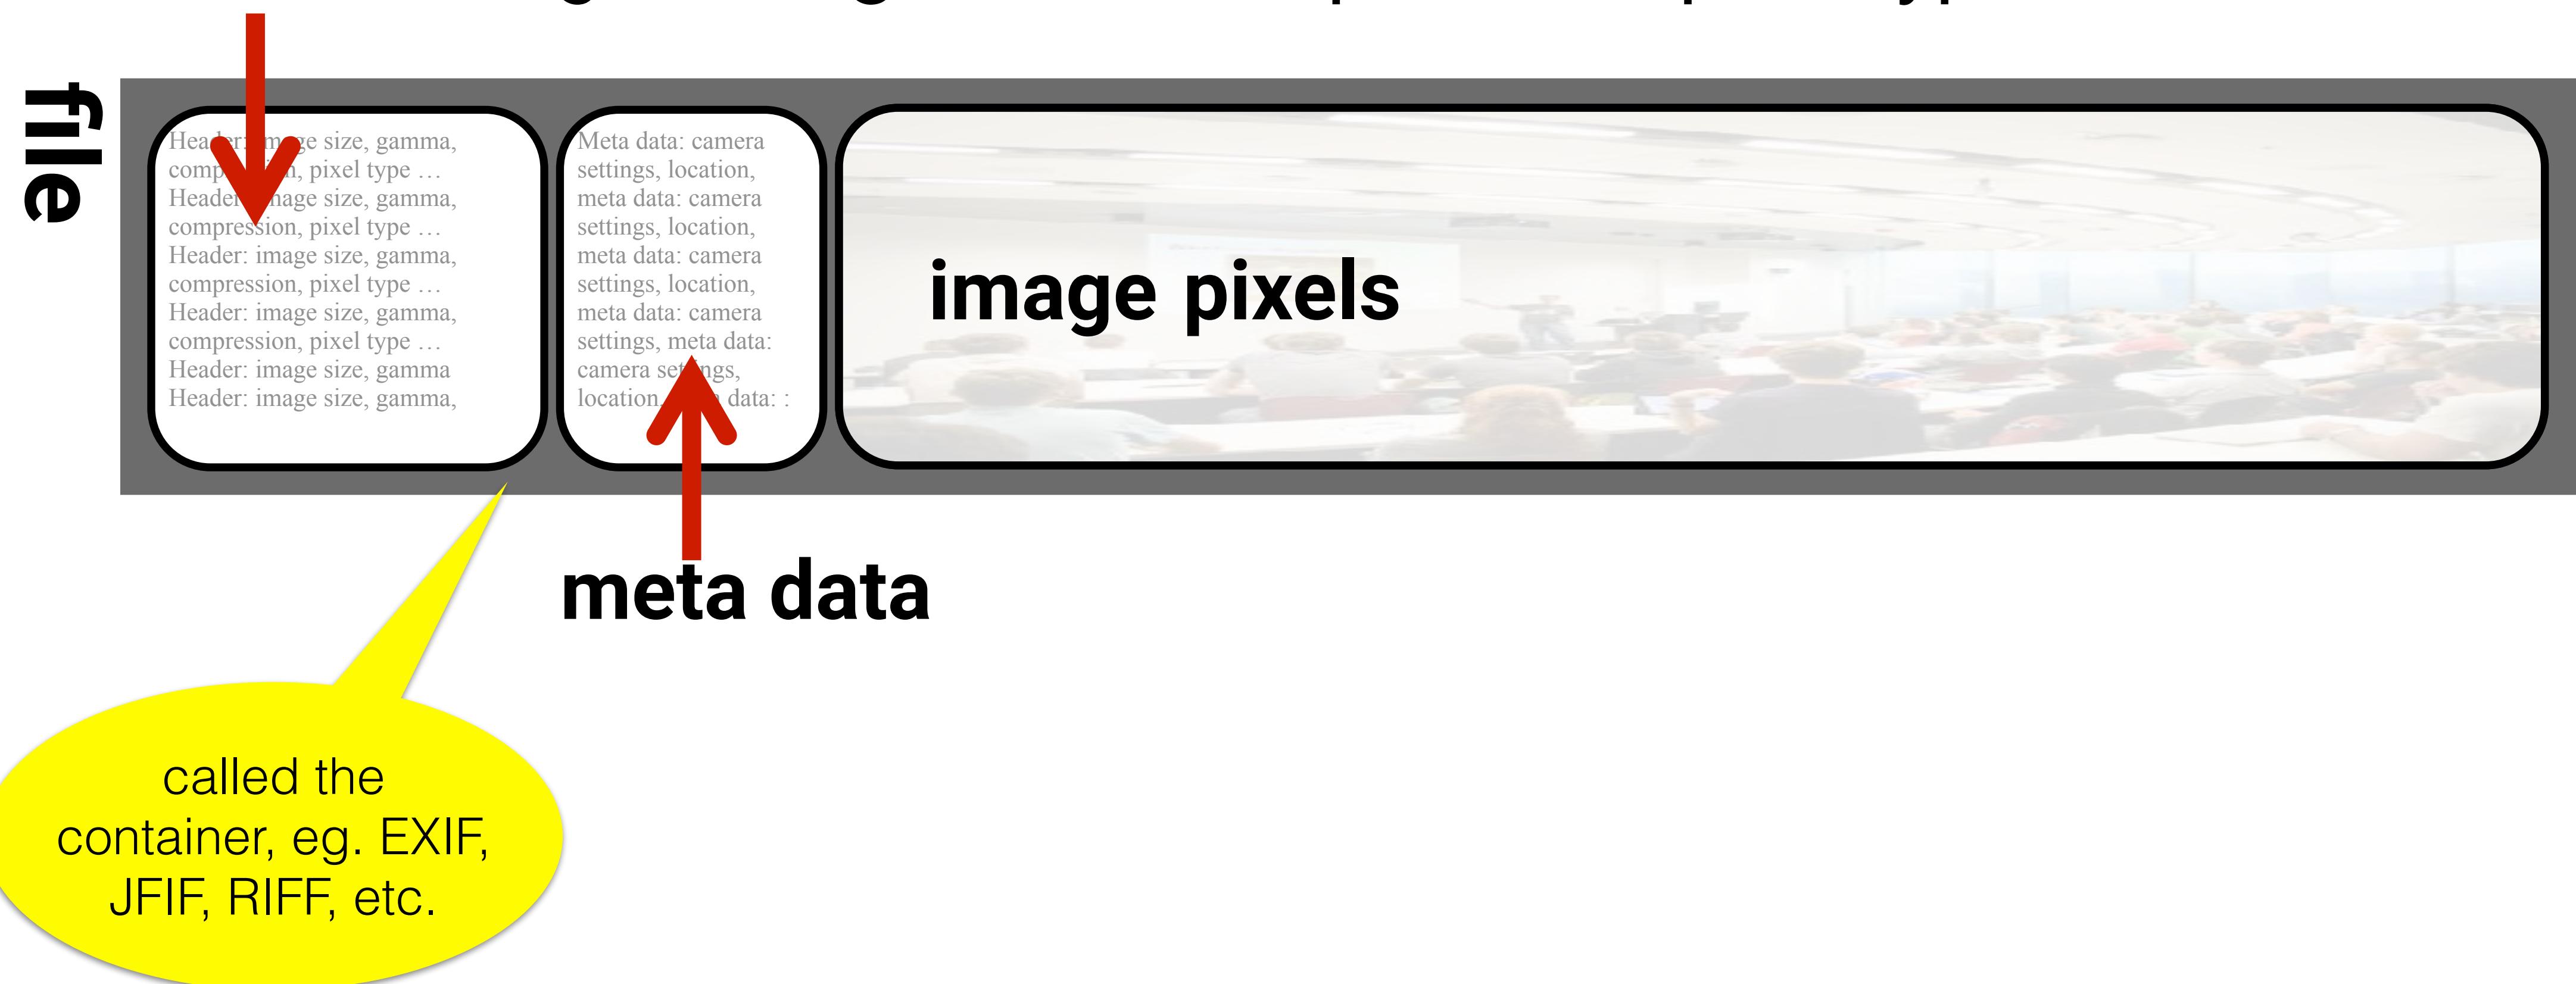


Image file formats

header: image size, gamma, compression, pixel type...



Display non-linearity

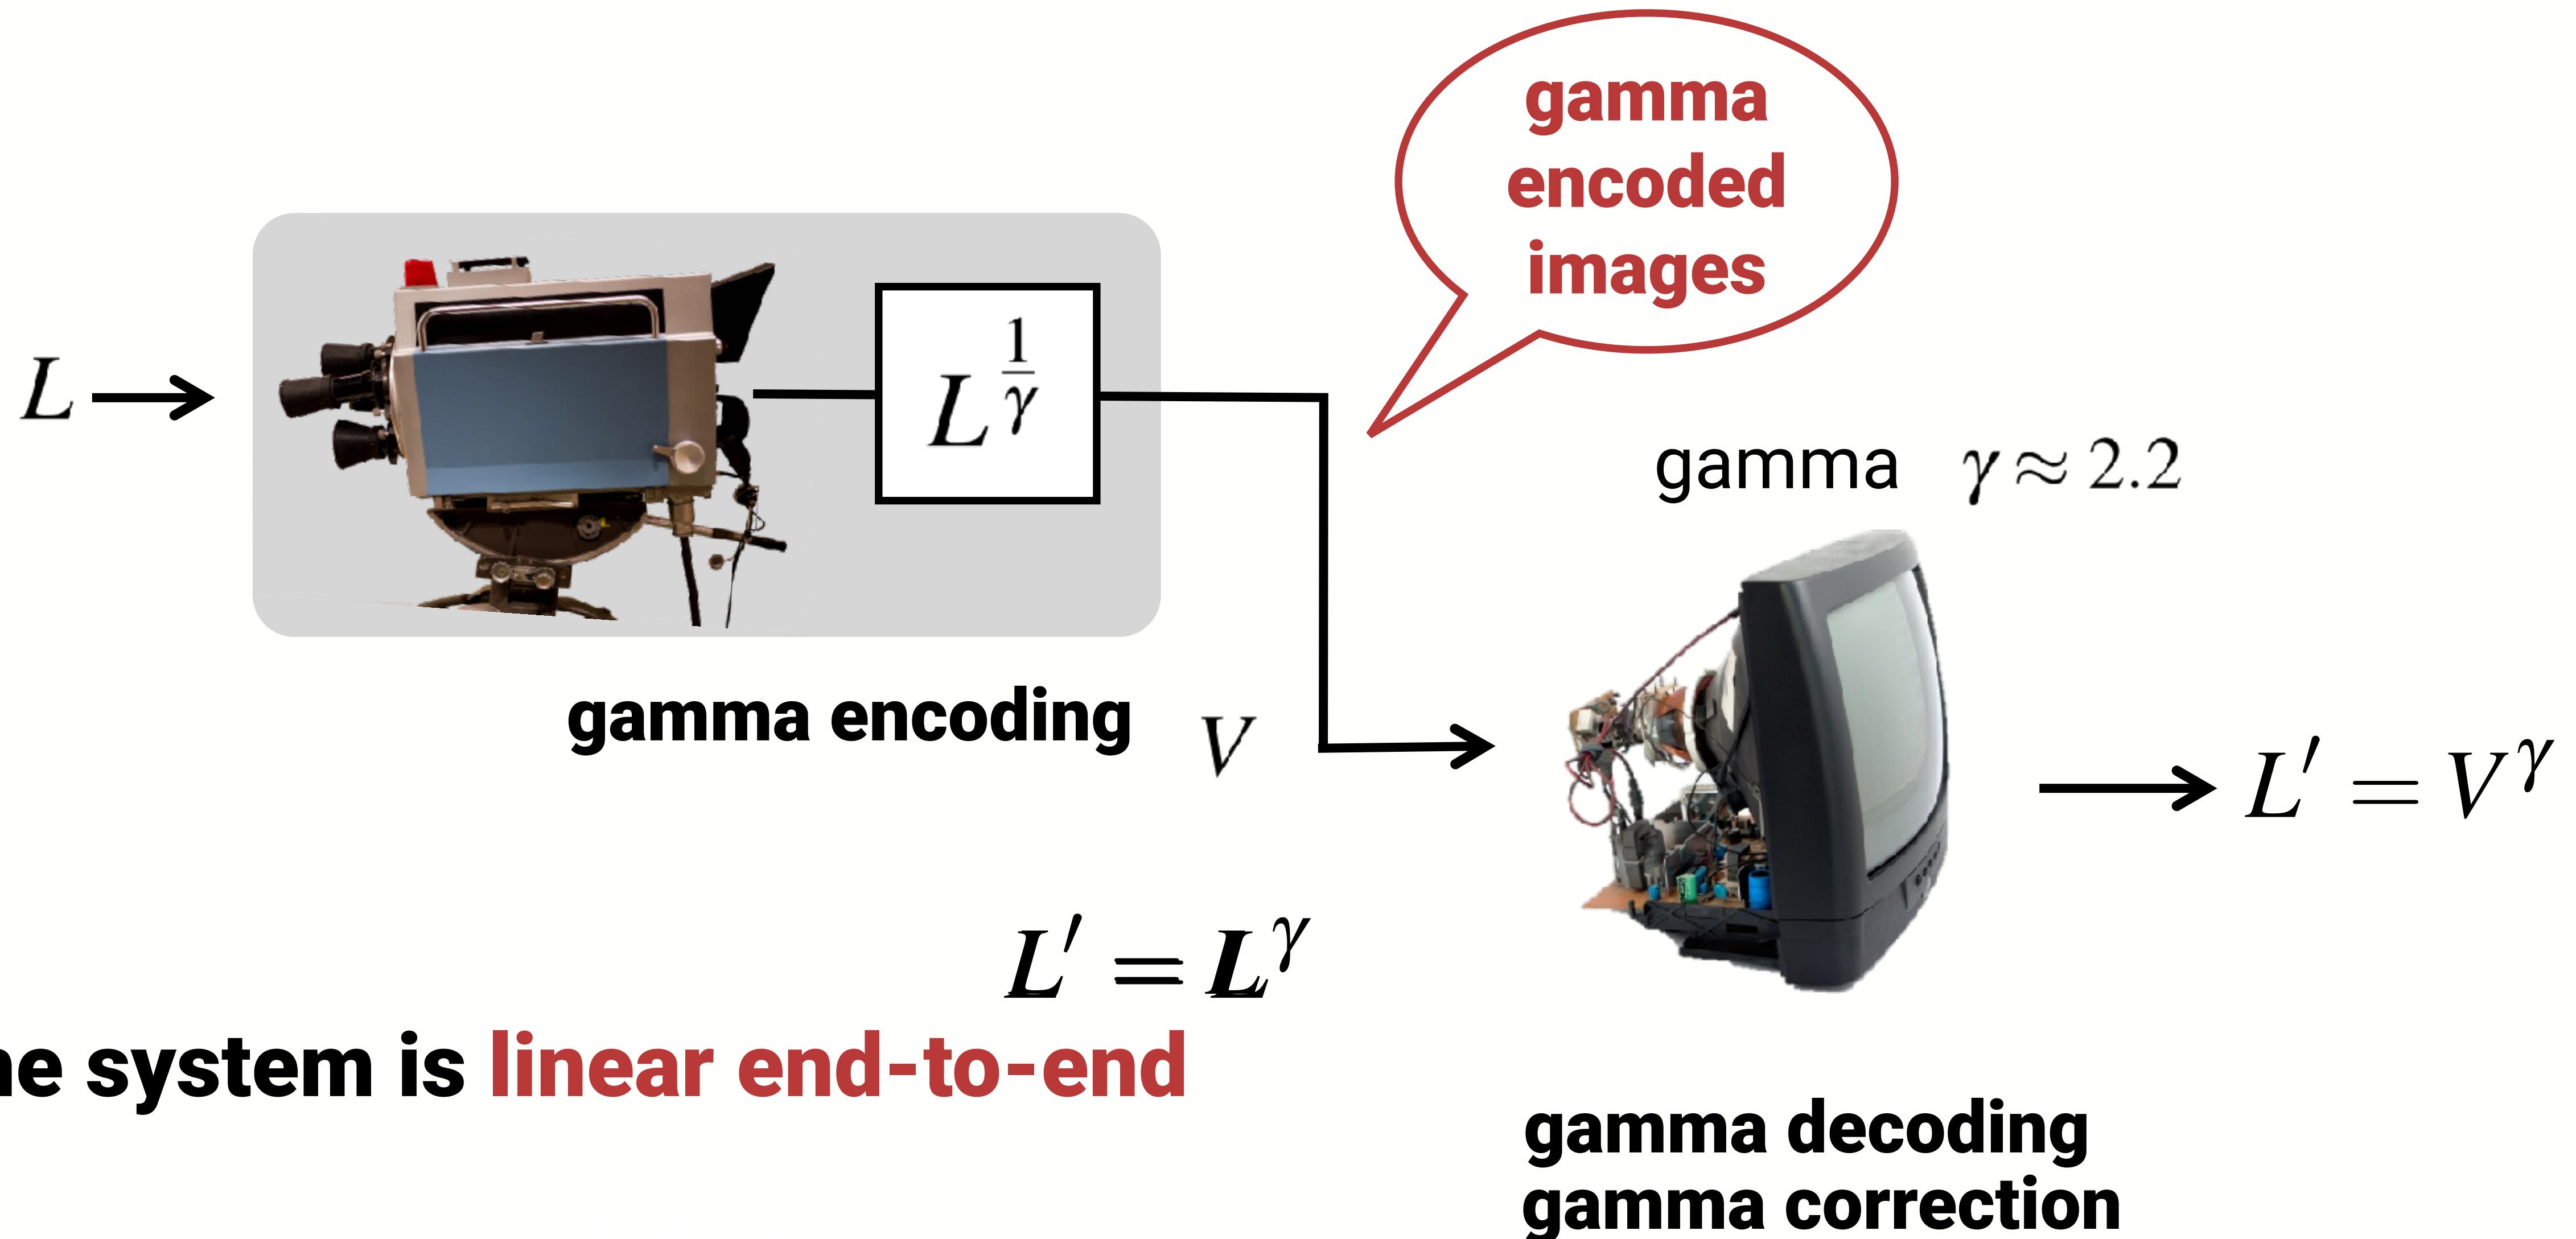
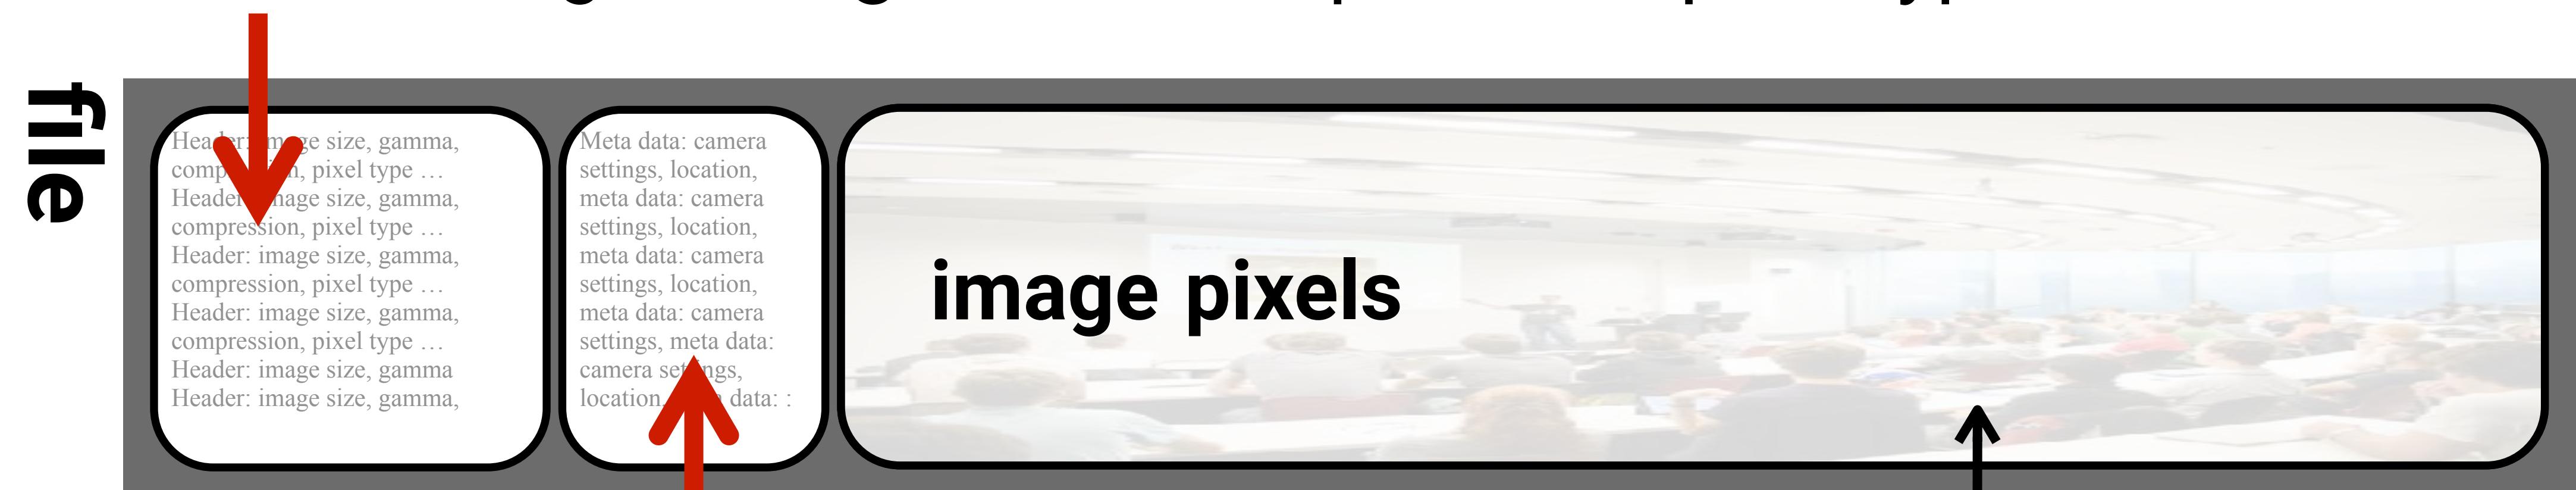


Image file formats

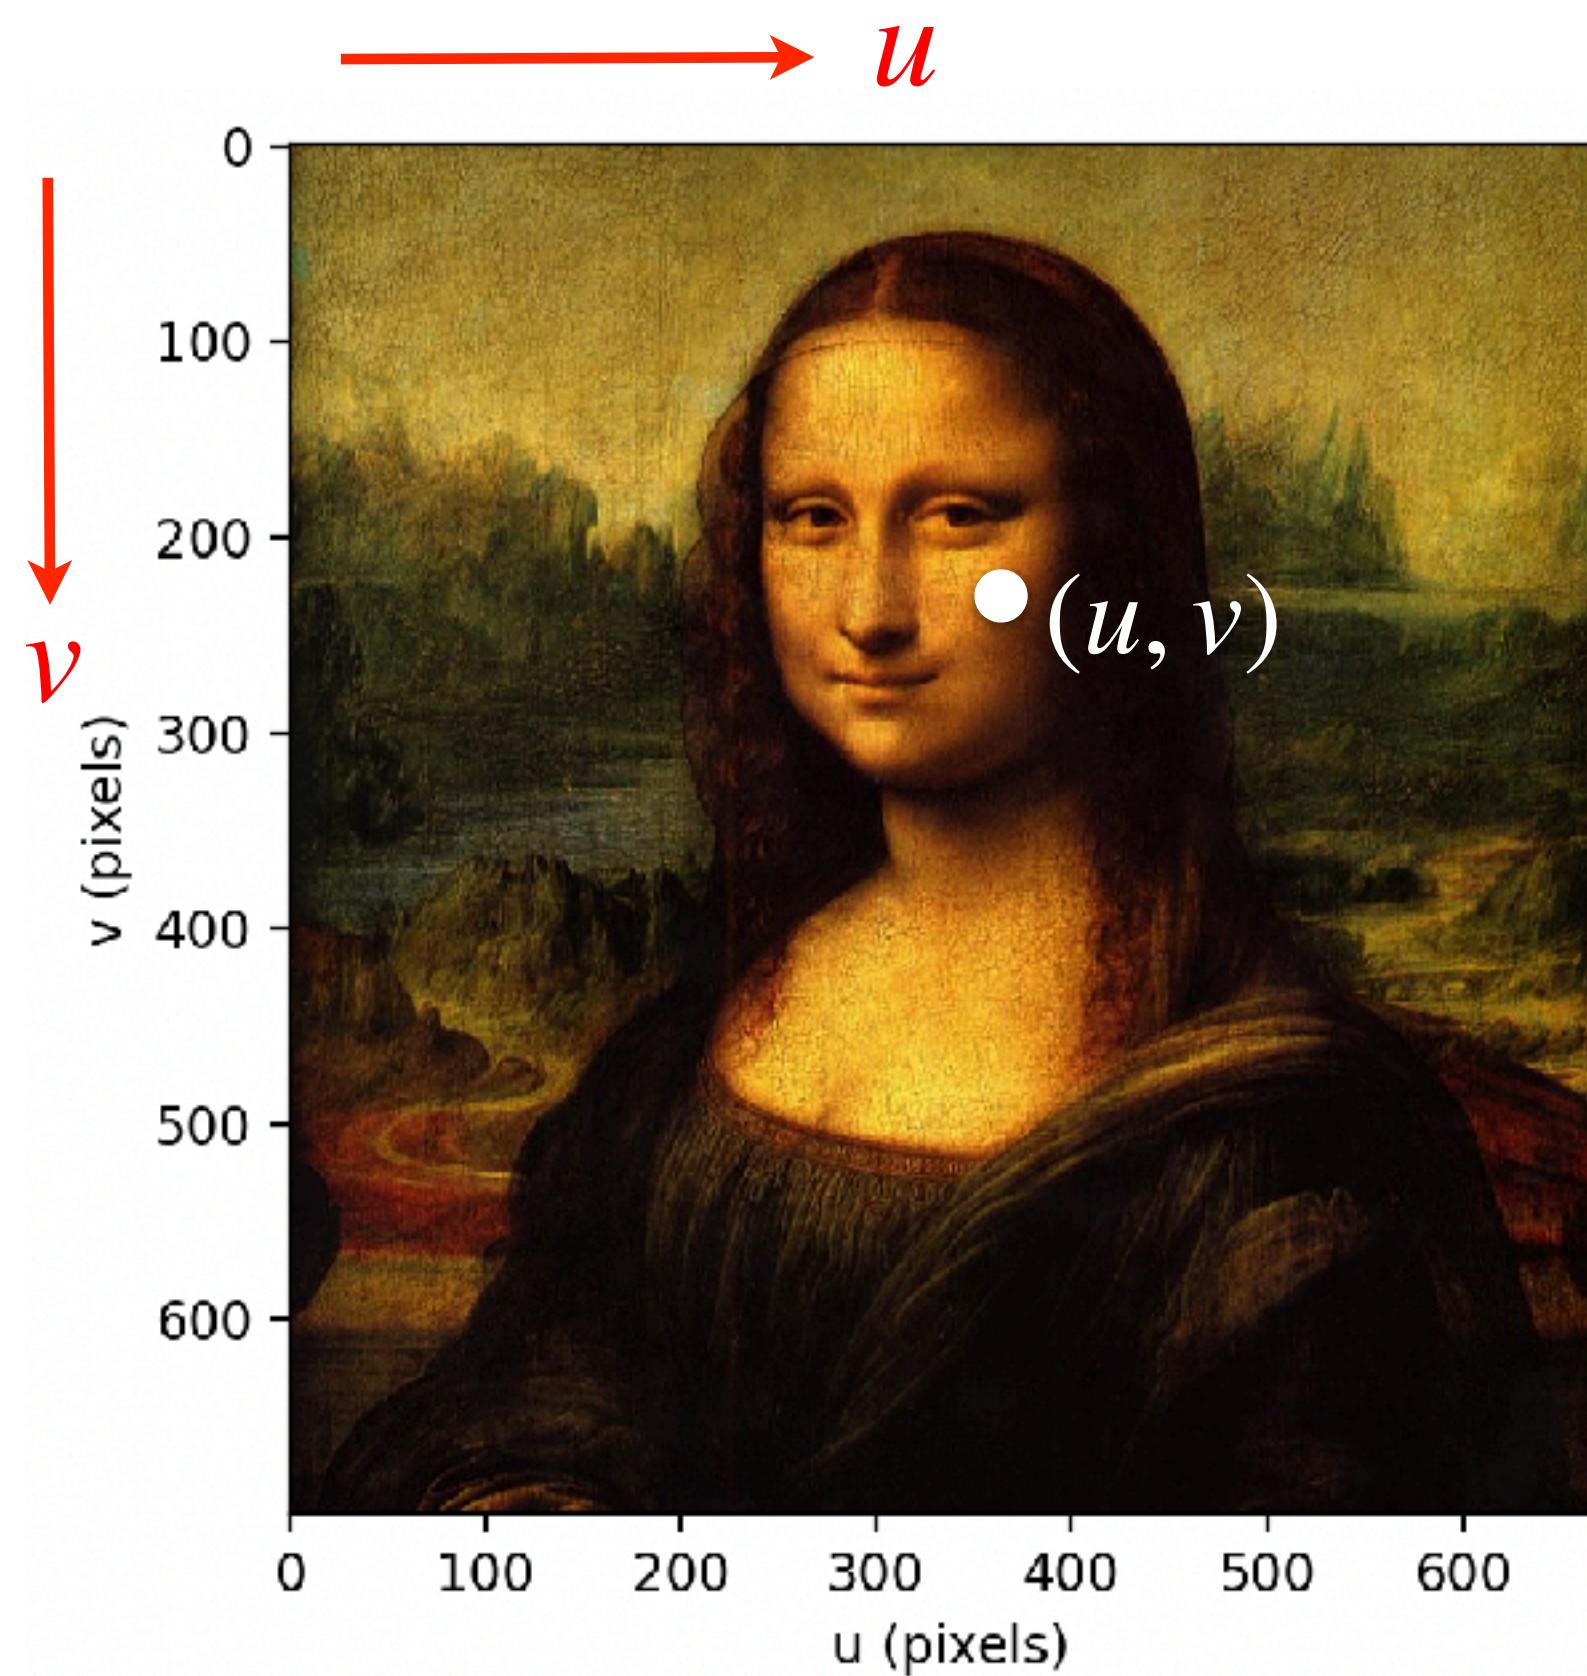
header: image size, gamma, compression, pixel type...



meta data



Beware that indexing into an image Numpy array is reversed compared to the way we normally think about coordinates



```
row=v          column=u  
              ↓           ↗  
=> mona, _ = imread("monalisa.png")  
=> mona[200,100]  
array([111, 98, 39], dtype=uint8)
```

DANGER: Reversed compared to graph coordinates!

```
>>> mona = Image.Read("monalisa.png")  
>>> mona[100,200]  
array([111, 98, 39], dtype=uint8)
```

Images can have tens of millions of pixels, so we want to use a compact representation of each pixel value. uint8 is common

single pixel
“images”

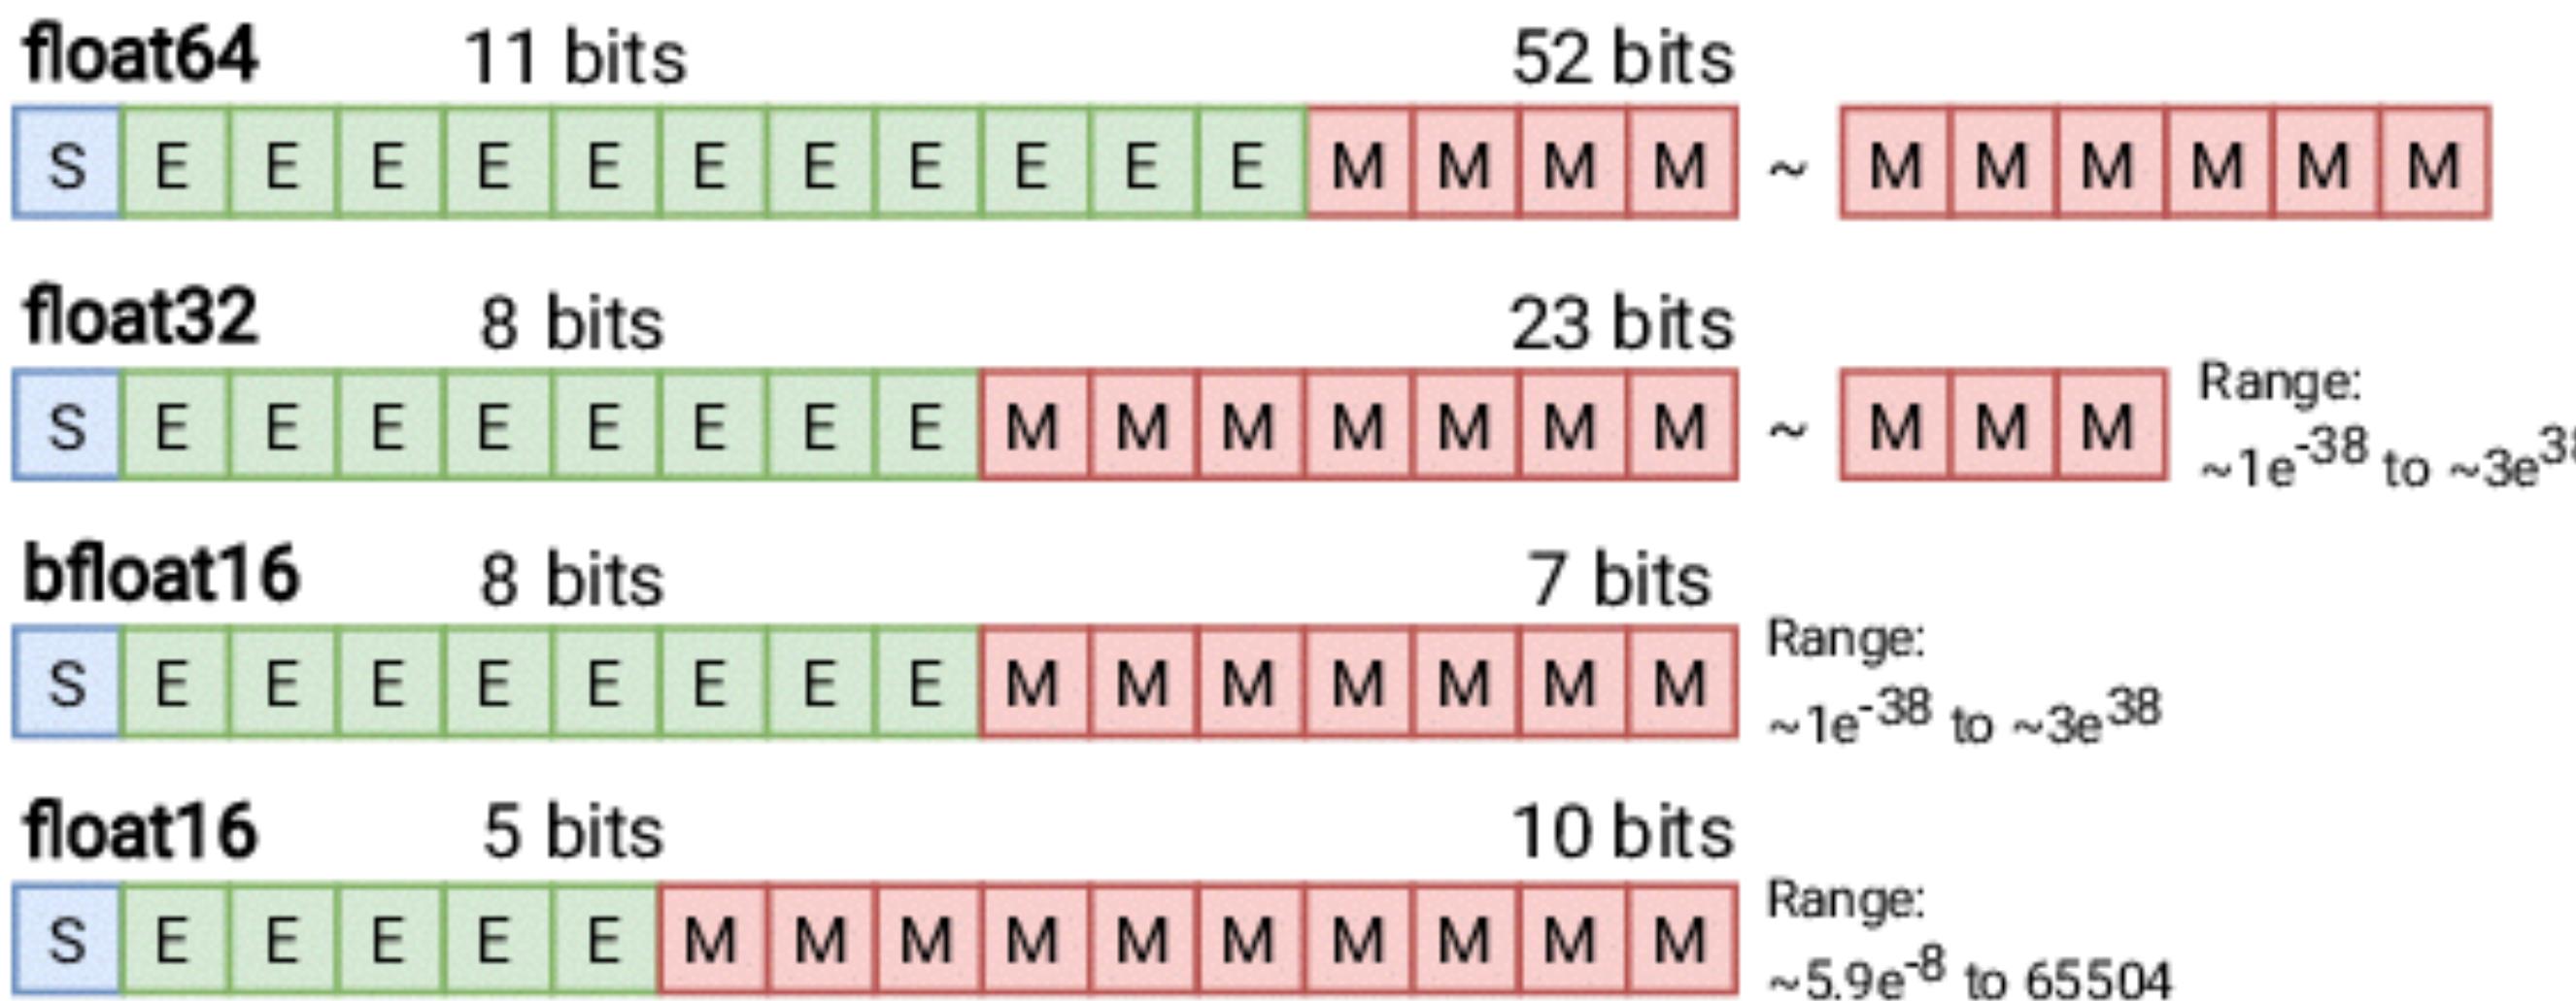
- uint8 = unsigned integer 8 bits
 - range 0 to 255
- each pixel requires 1 byte
- limited dynamic range
- saturation effects

```
>>> a = np.array(5, "uint8")
>>> b = np.array(10, "uint8")
>>> a + b
15
>>> import cv2
>>> cv2.add(a,b)
array([[15]], dtype=uint8)

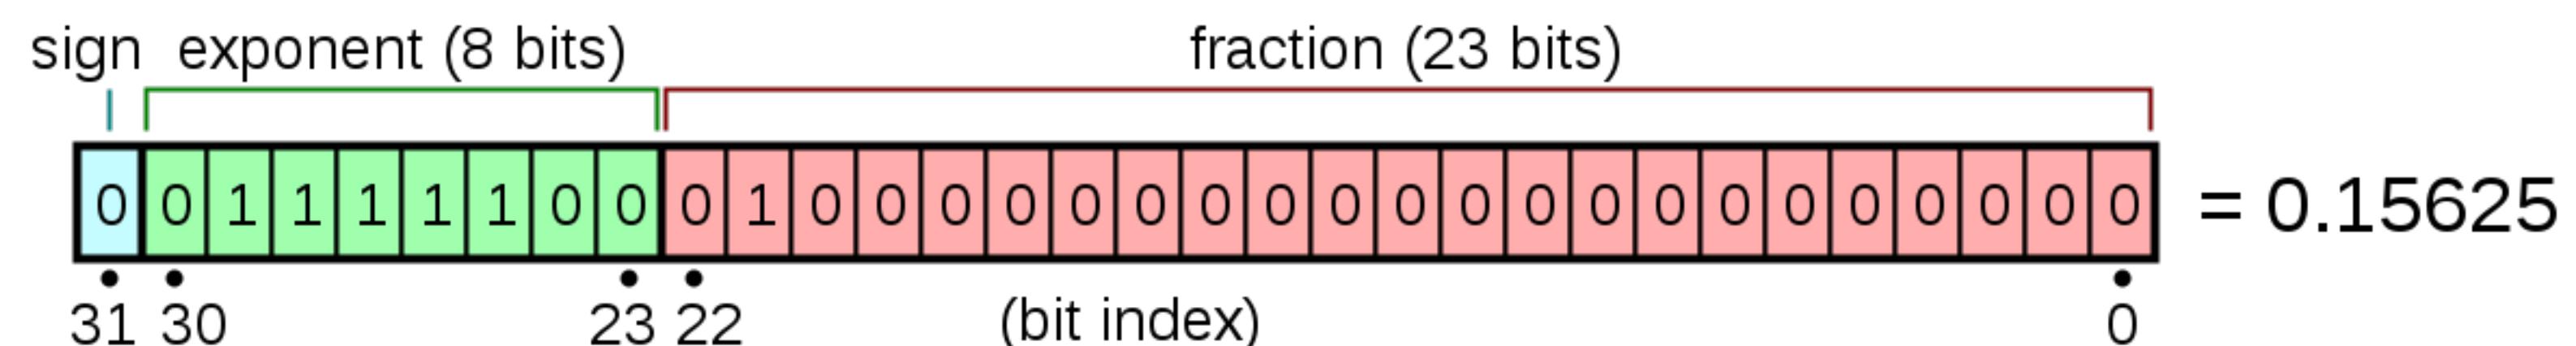
>>> a - b
251
>>> cv2.subtract(a,b)
array([[0]], dtype=uint8)

>>> c = np.array(250, "uint8")
>>> b + c
4
>>> cv2.add(b,c)
array([[255]], dtype=uint8)
```

Sometimes we need fractional values so we choose an IEEE 754 standard floating point number. They come in 3 sizes: 16, 32 or 64 bits.



$$p = M \cdot 2^{(E-E_0)}, M \in [0.5, 1)$$



- greyscale values typically in the range [0, 1]

Image processing

Histograms

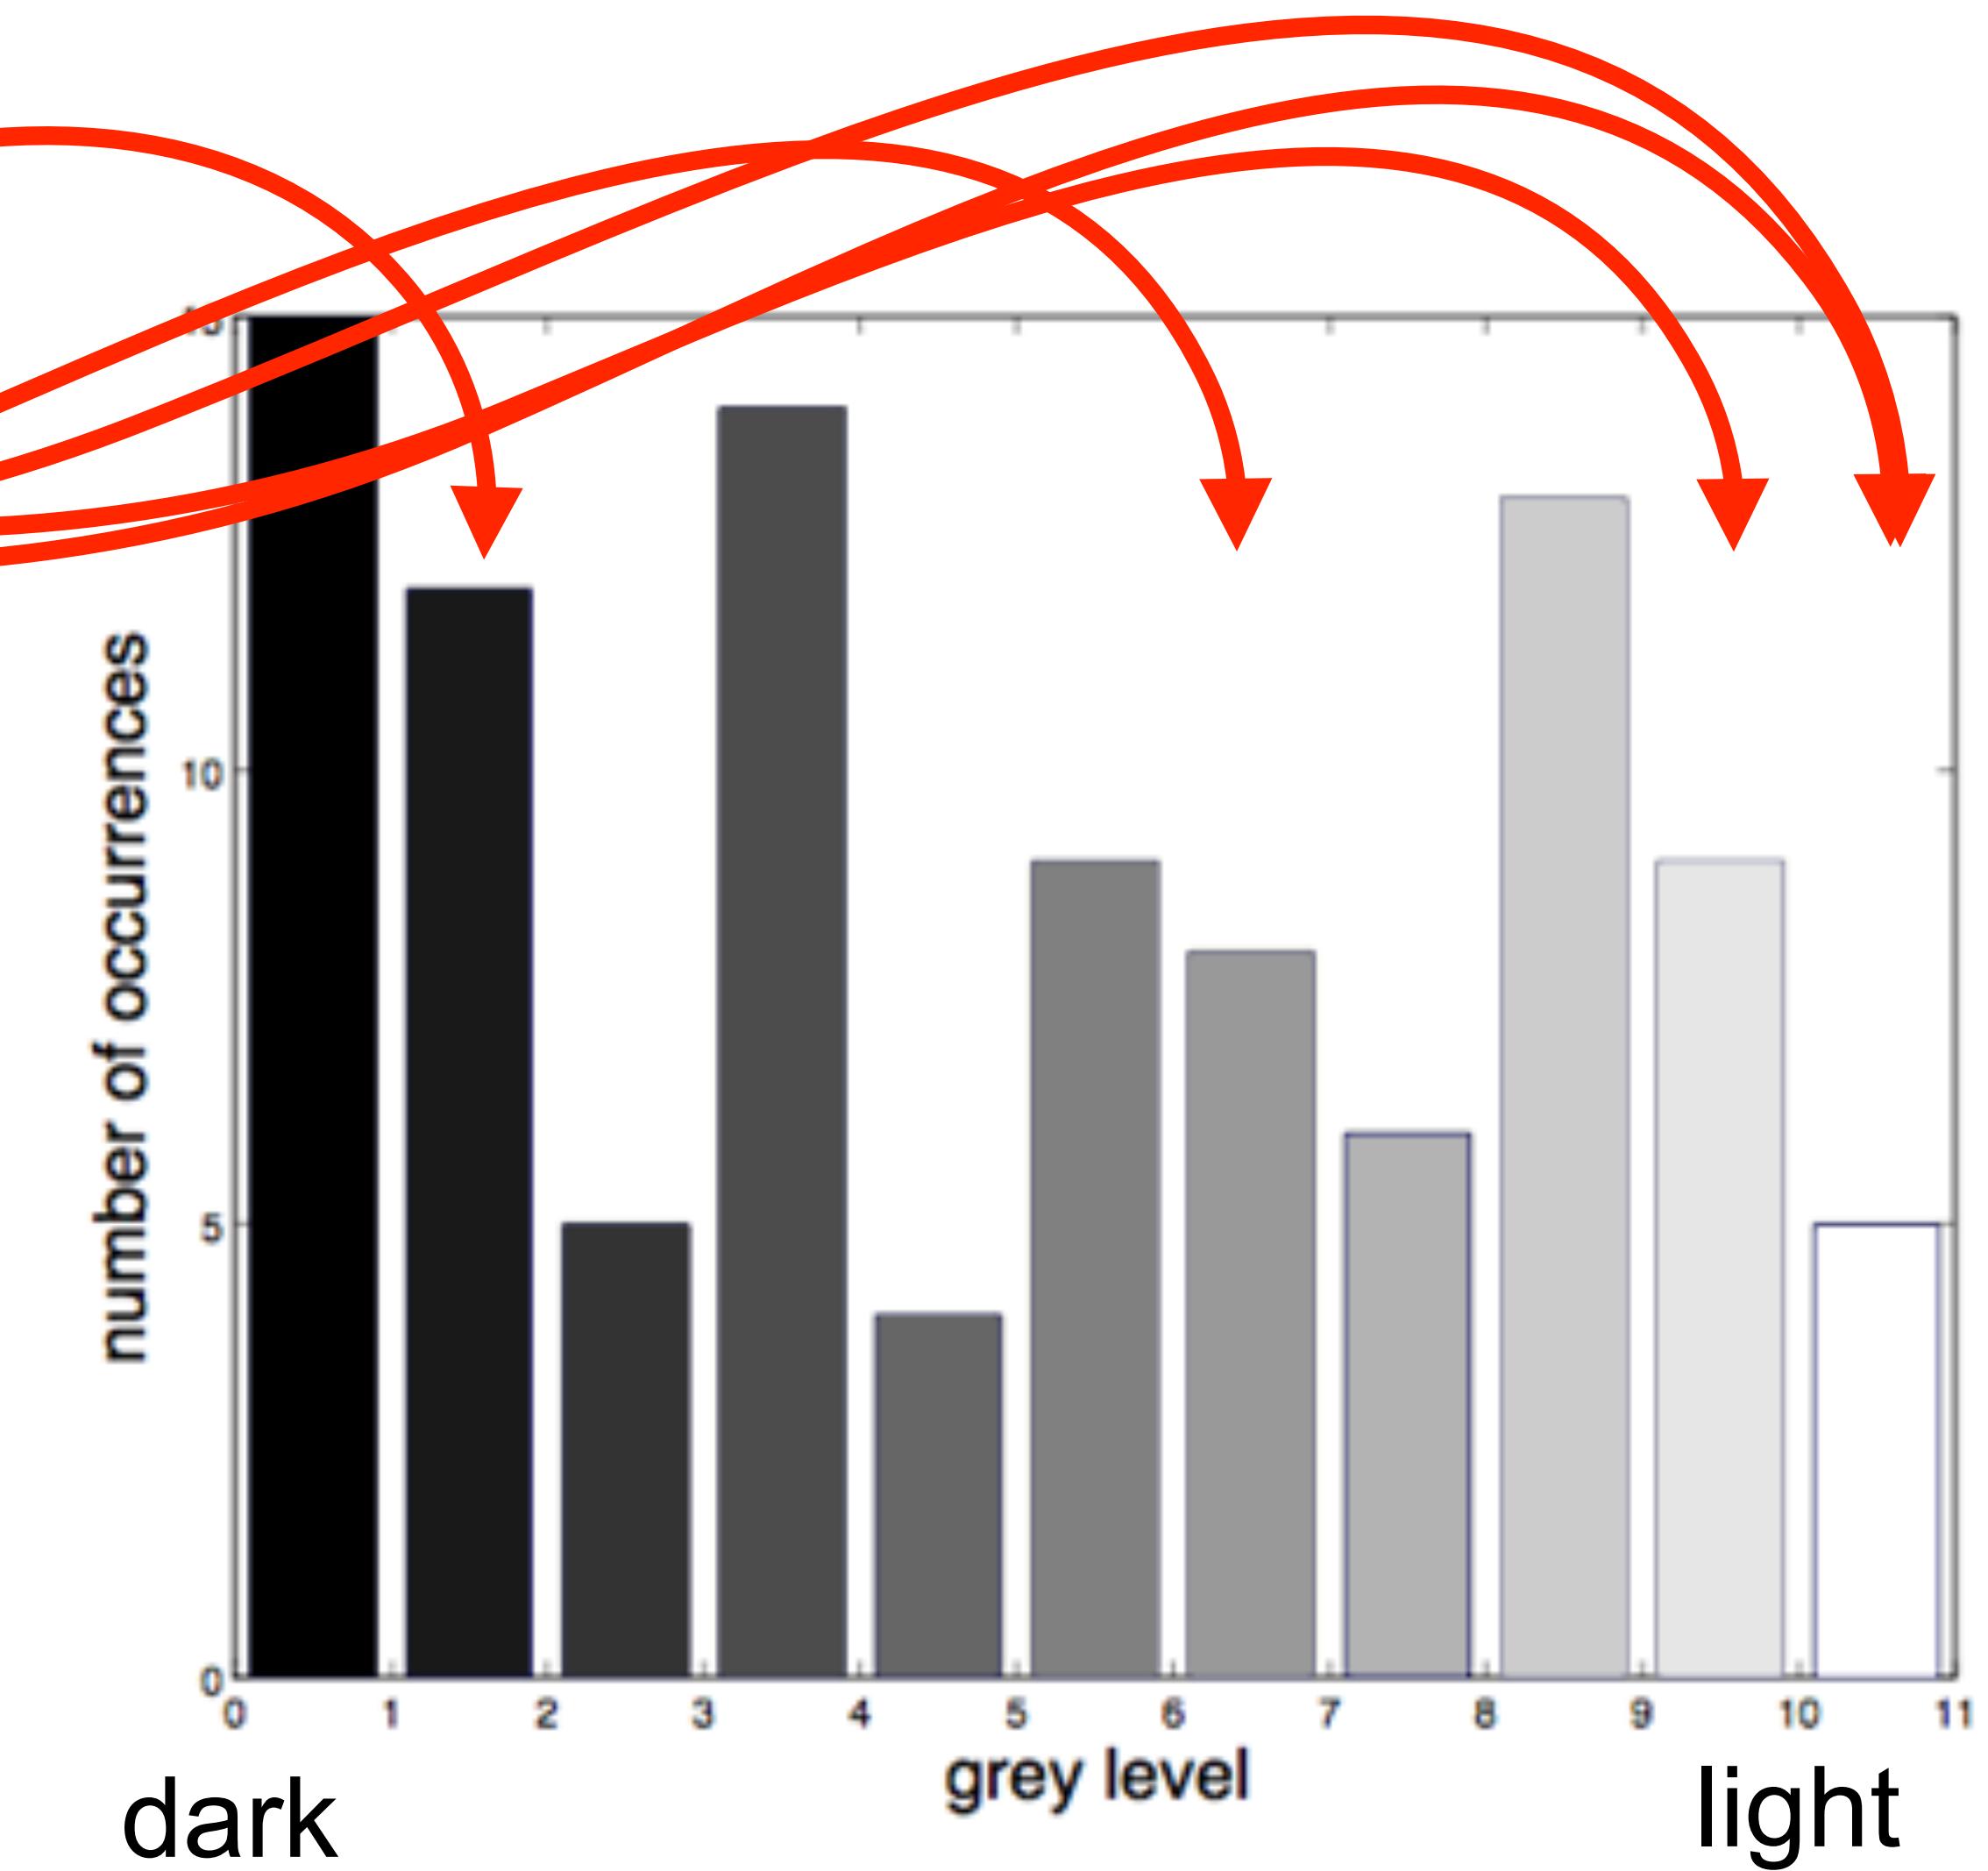
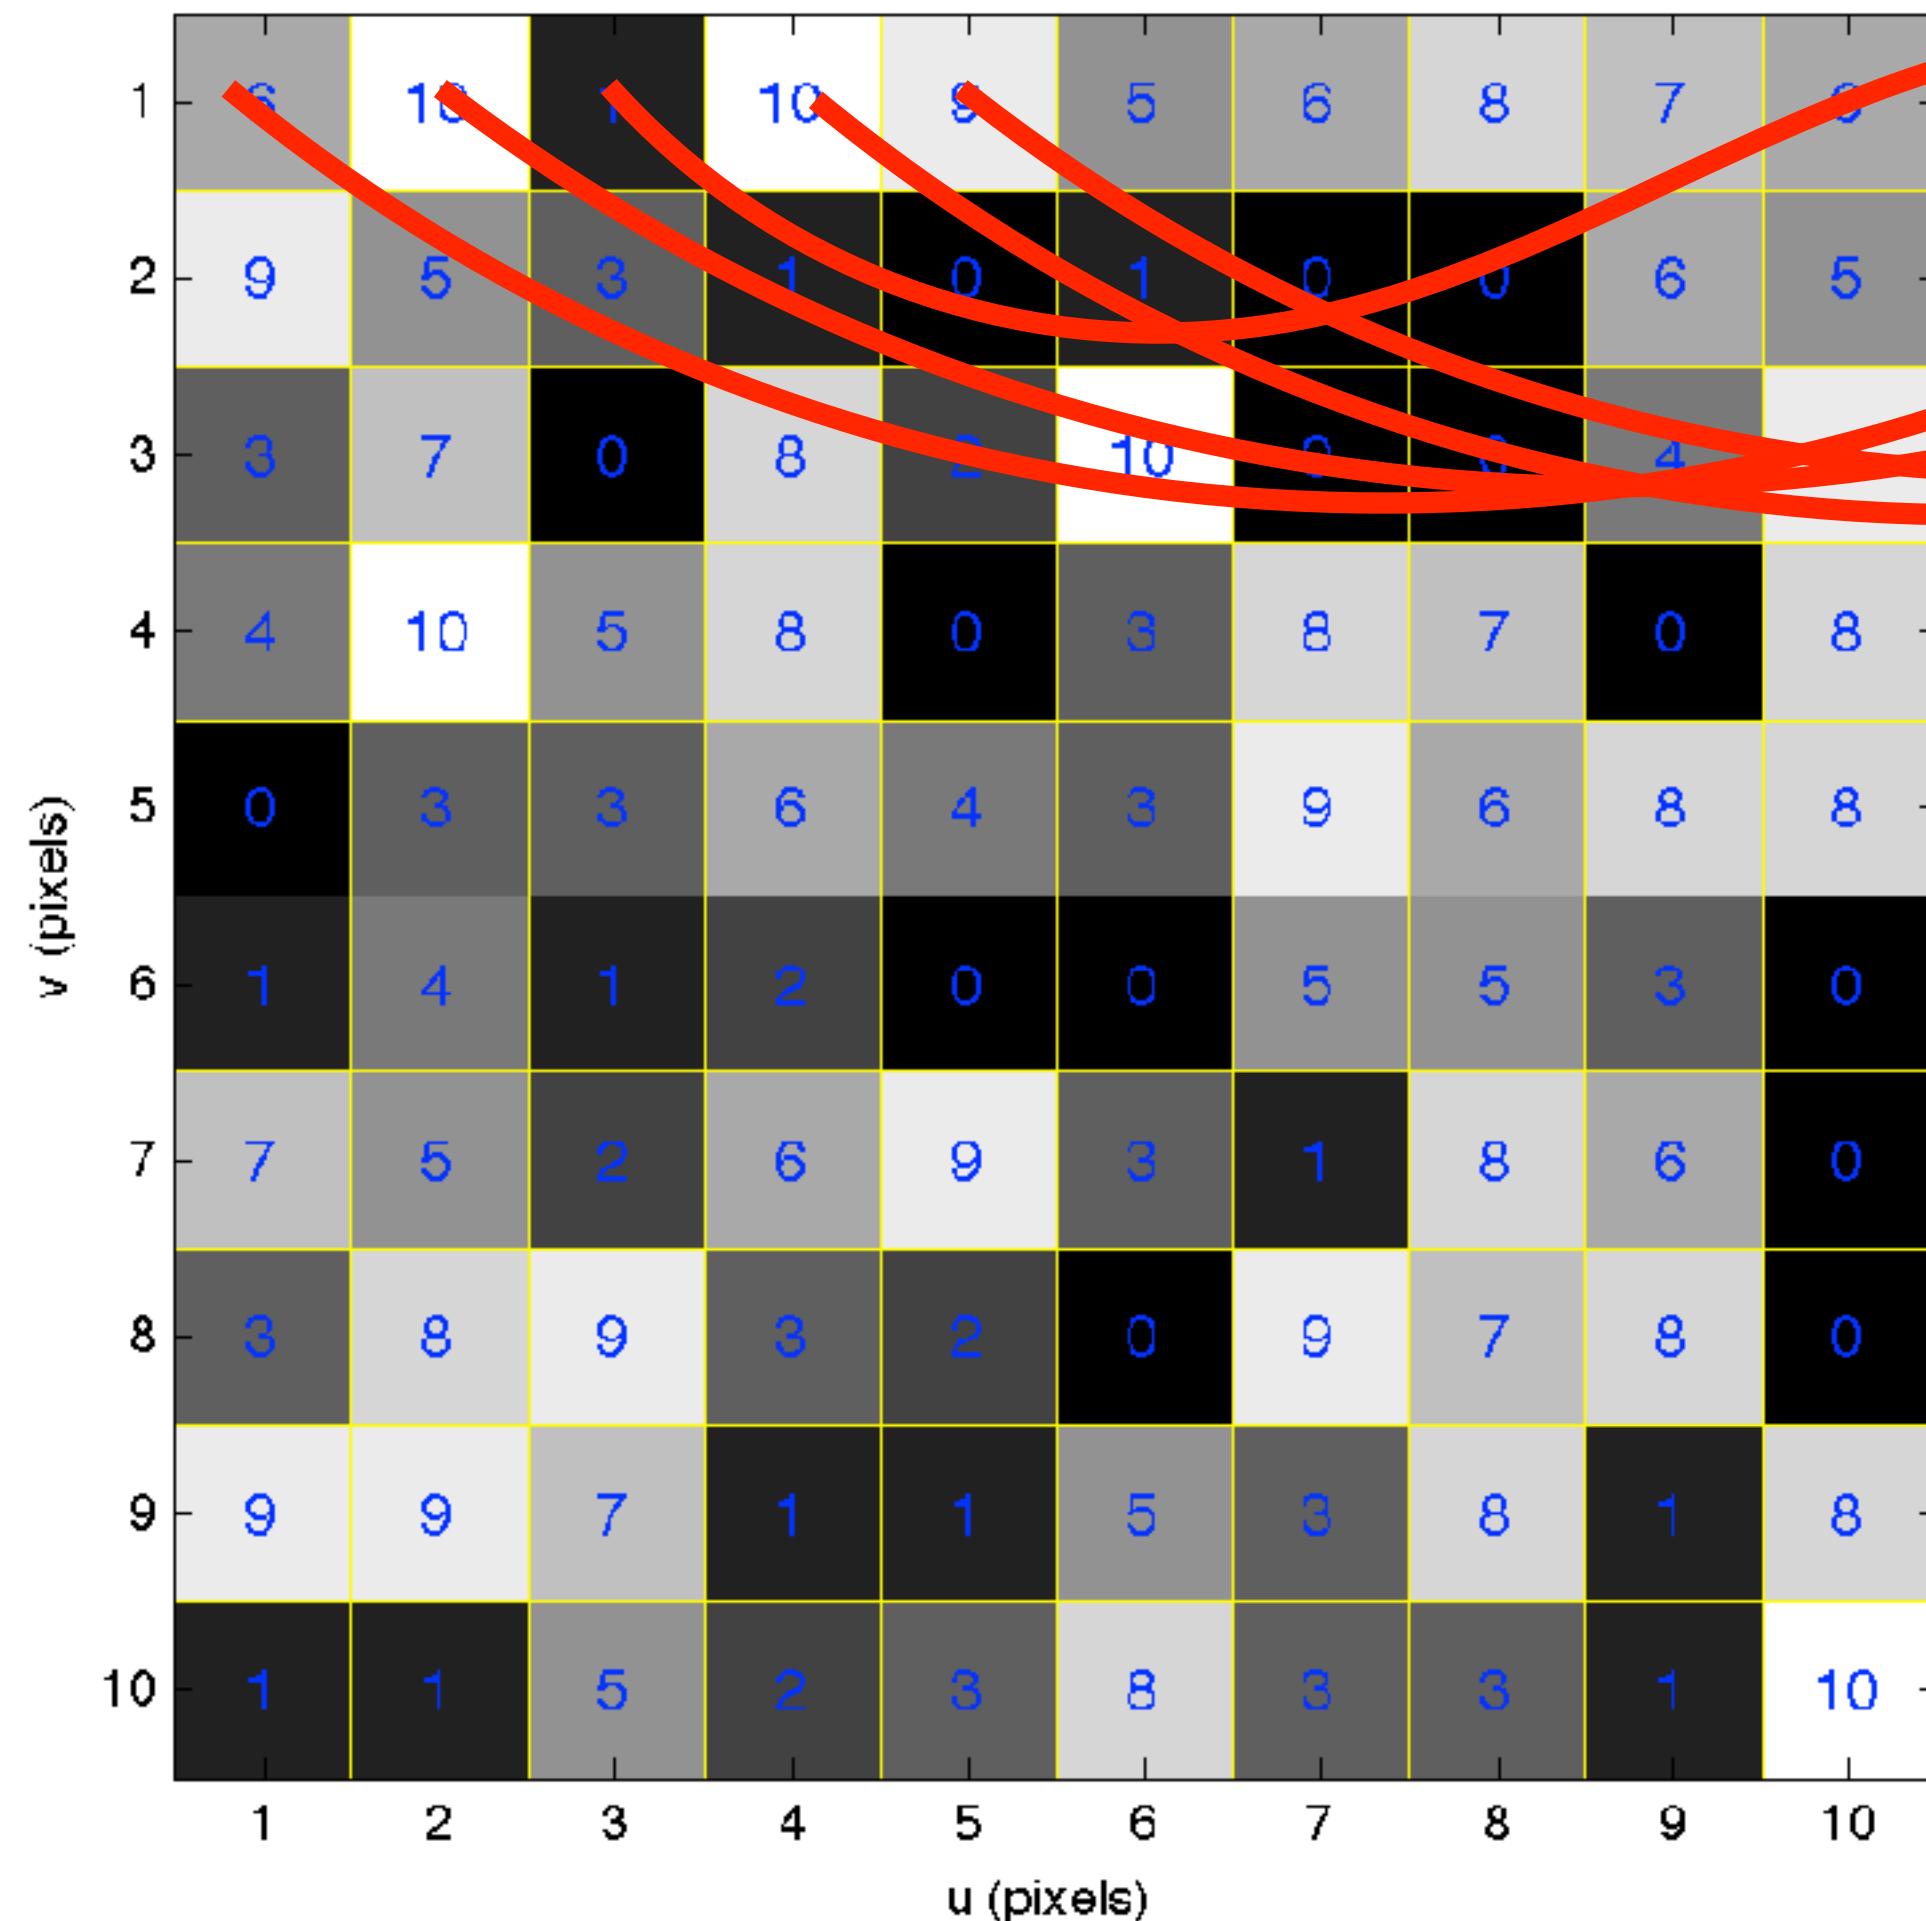
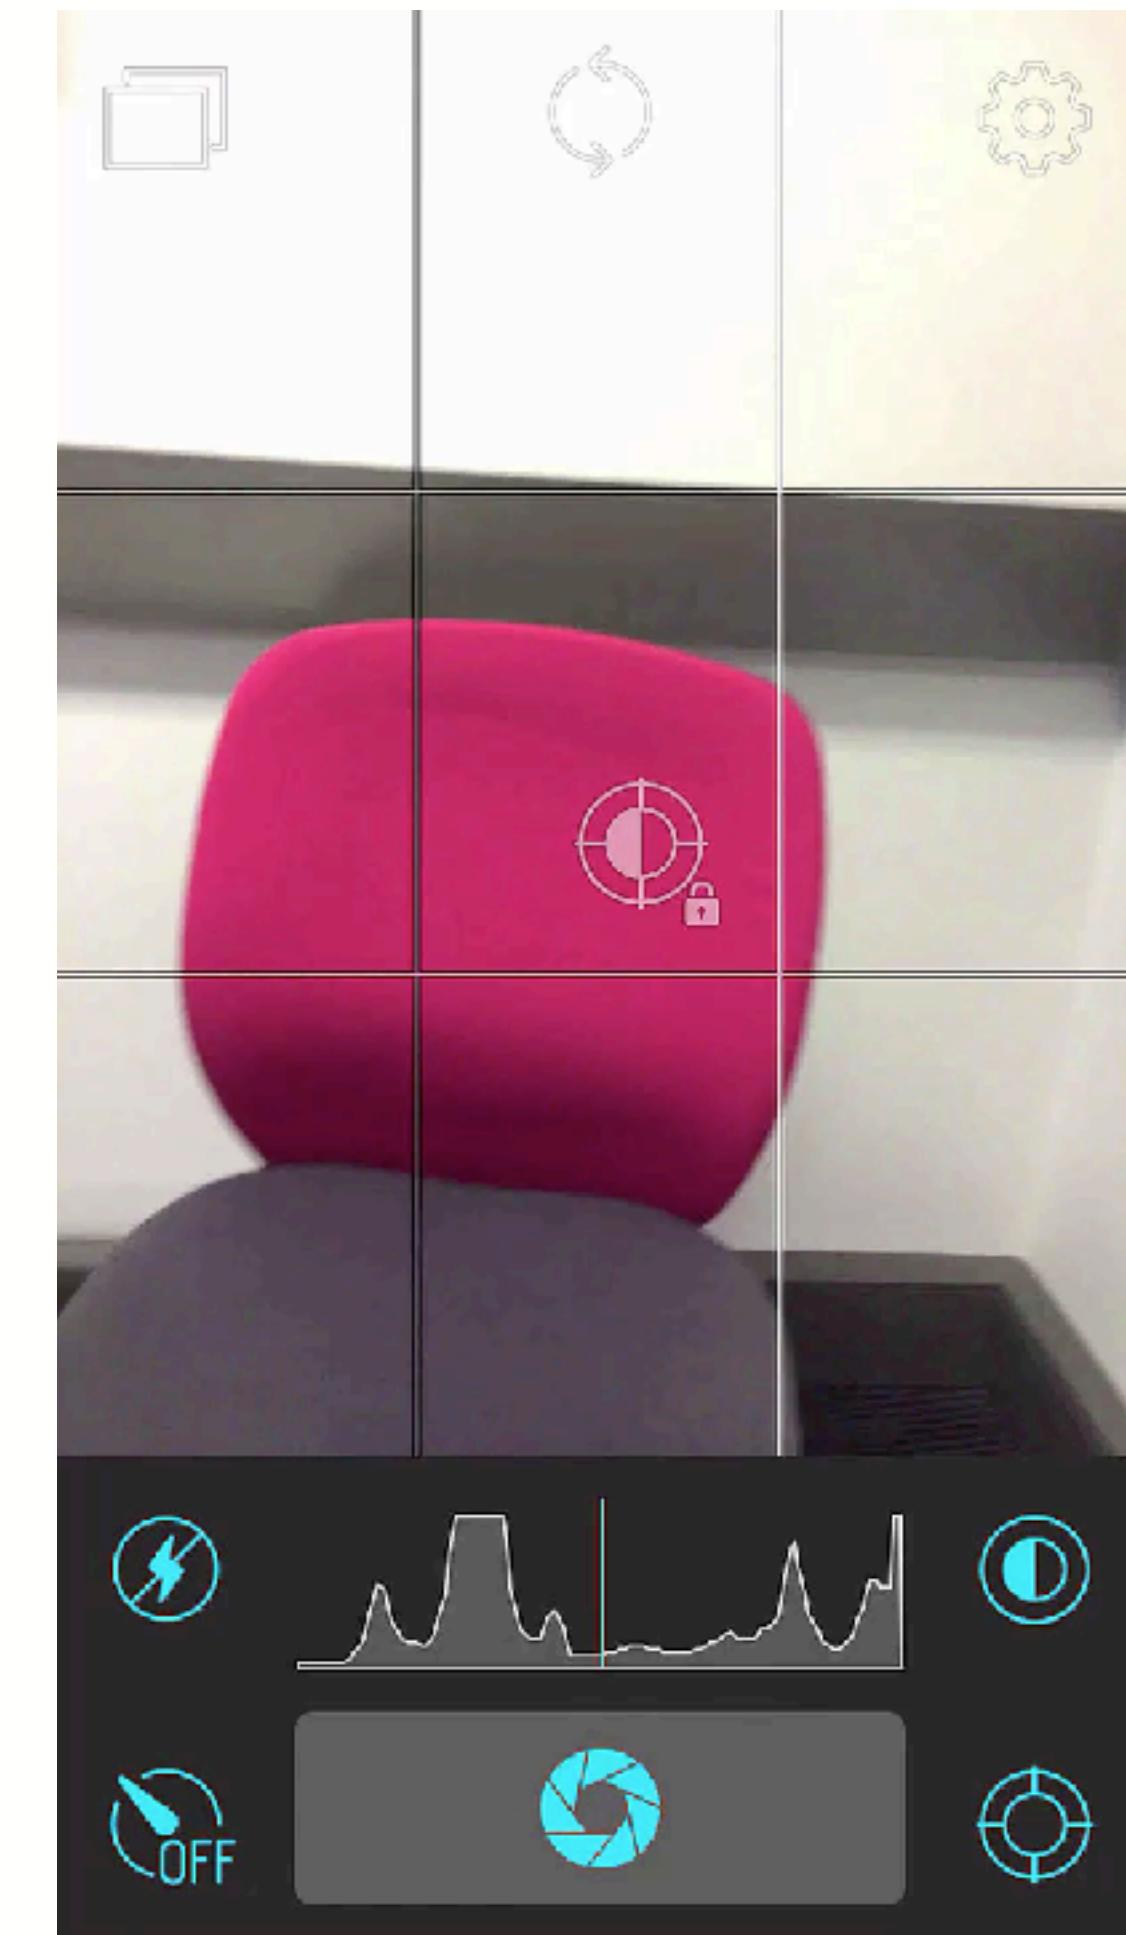
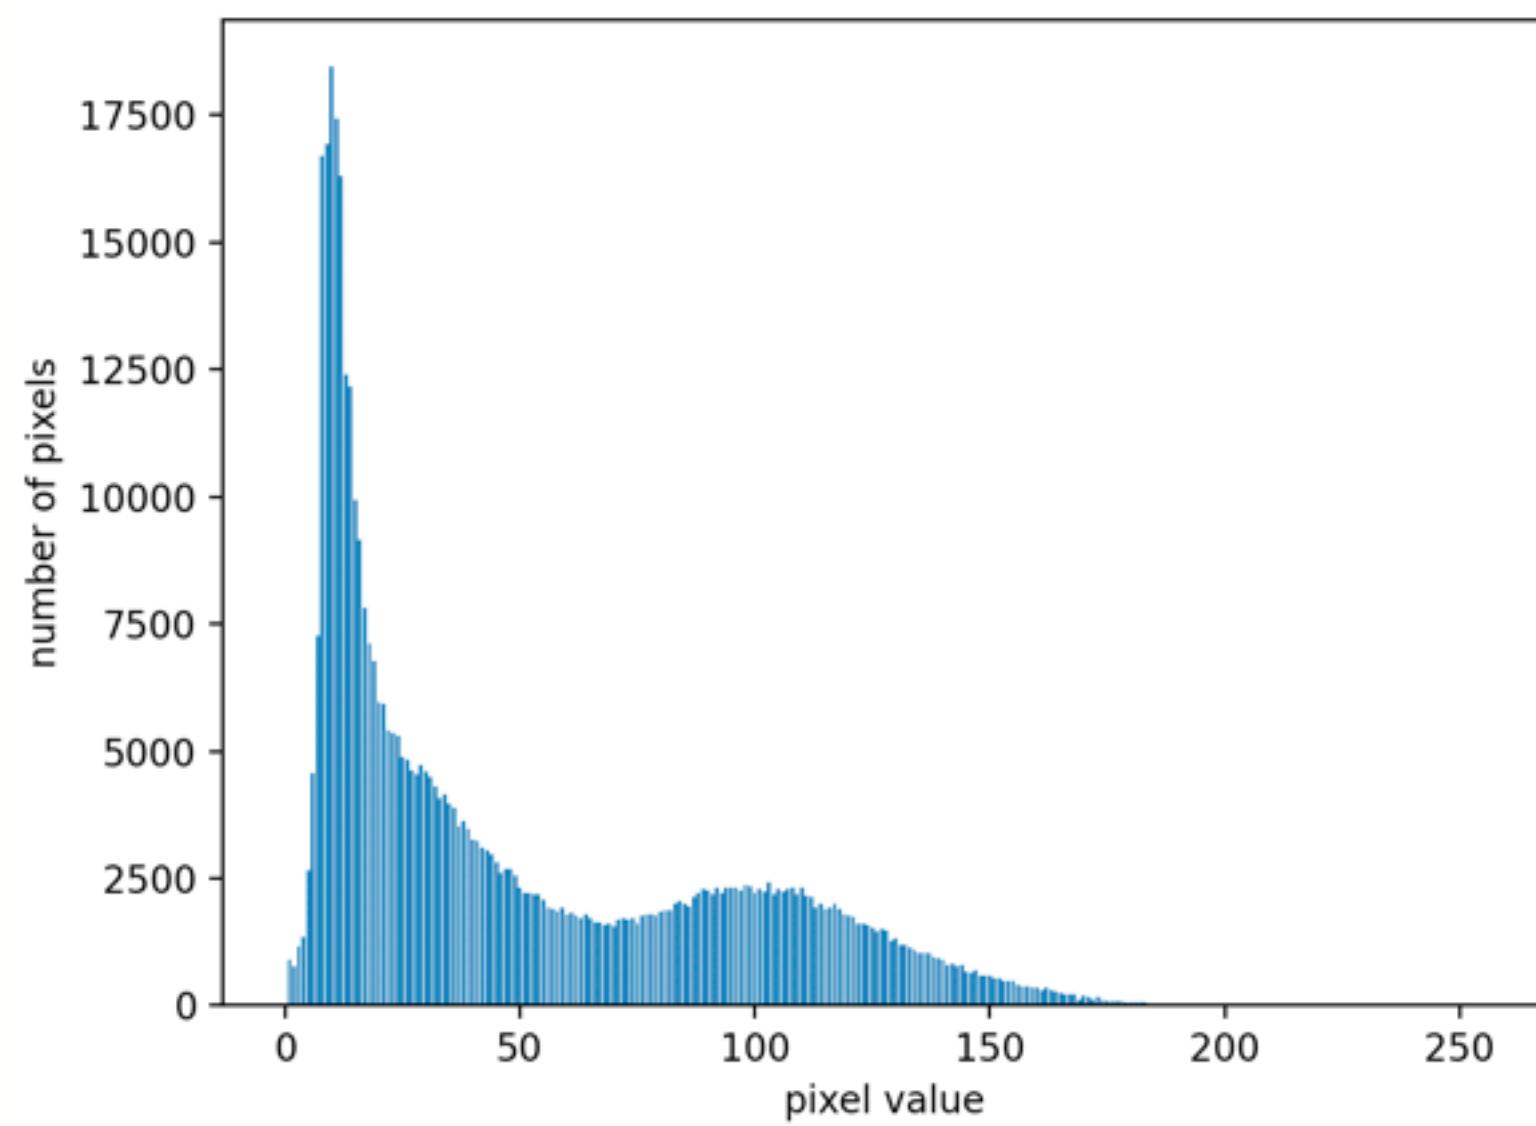
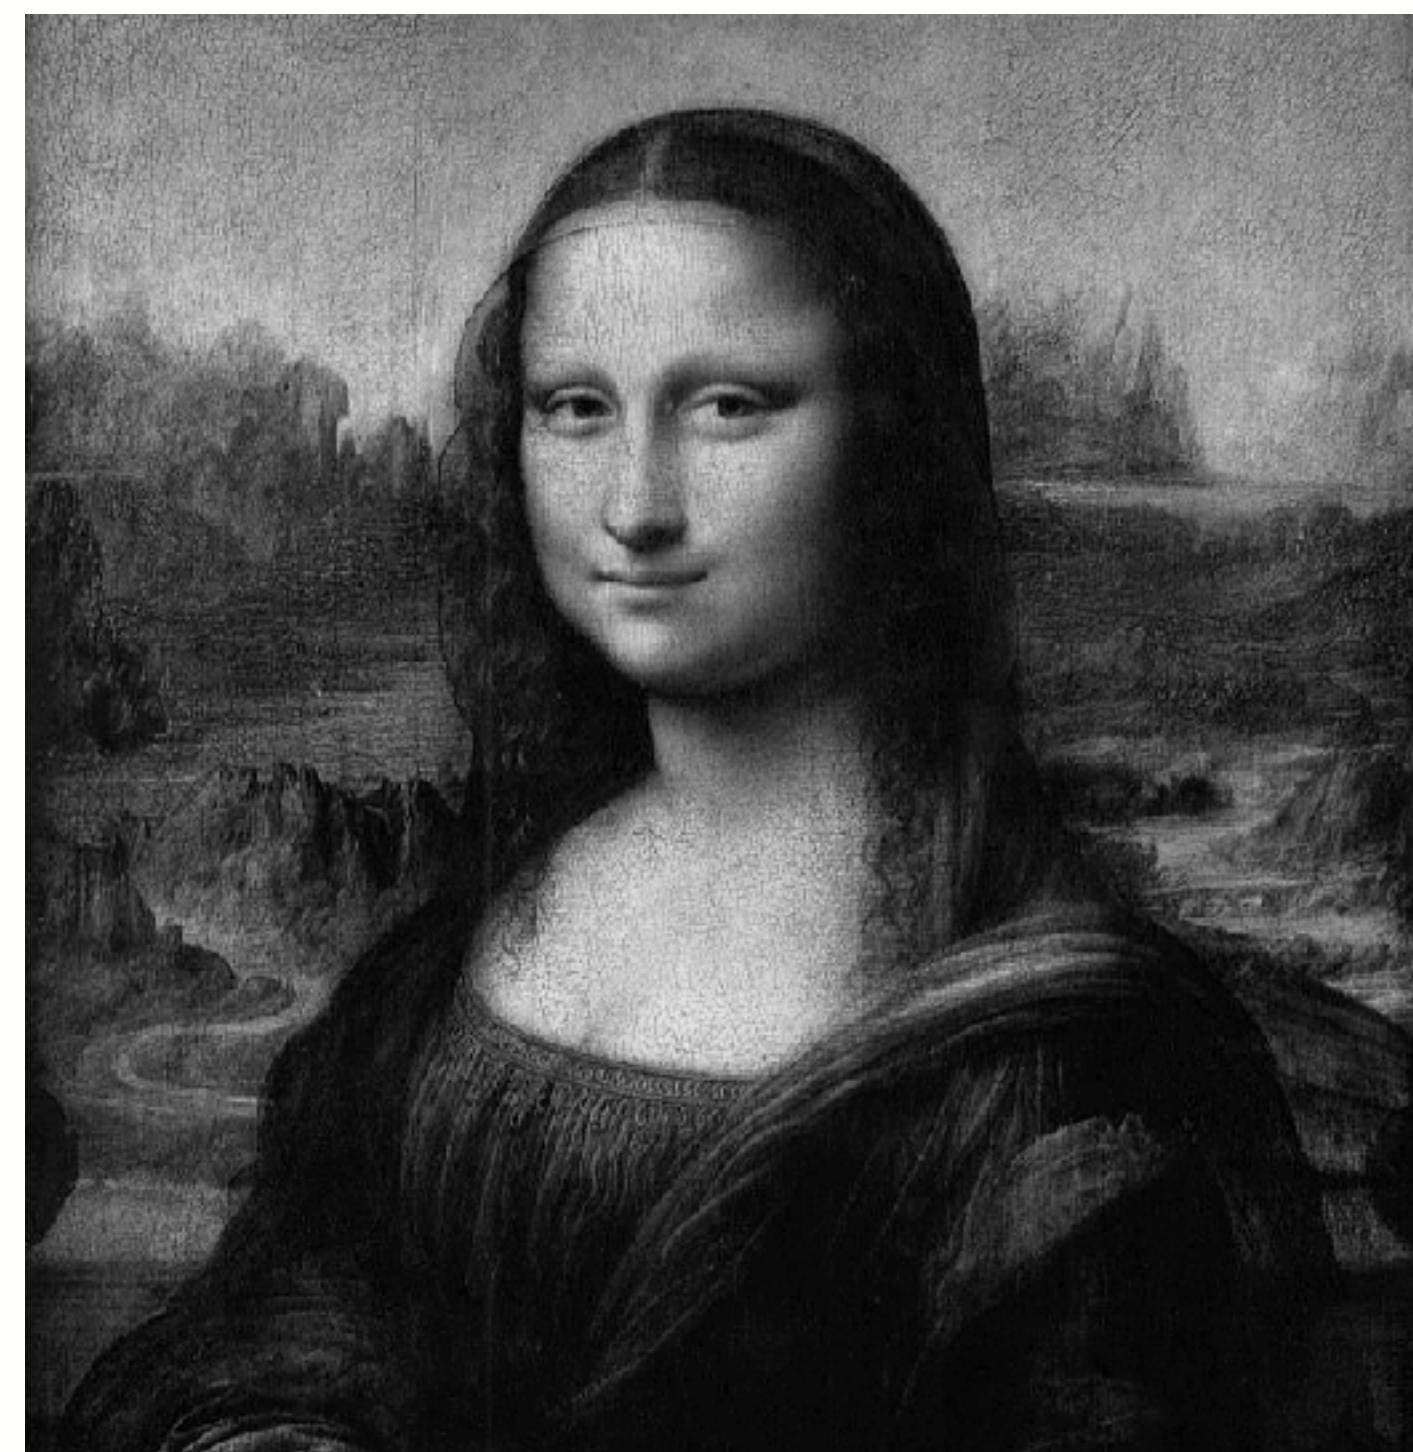
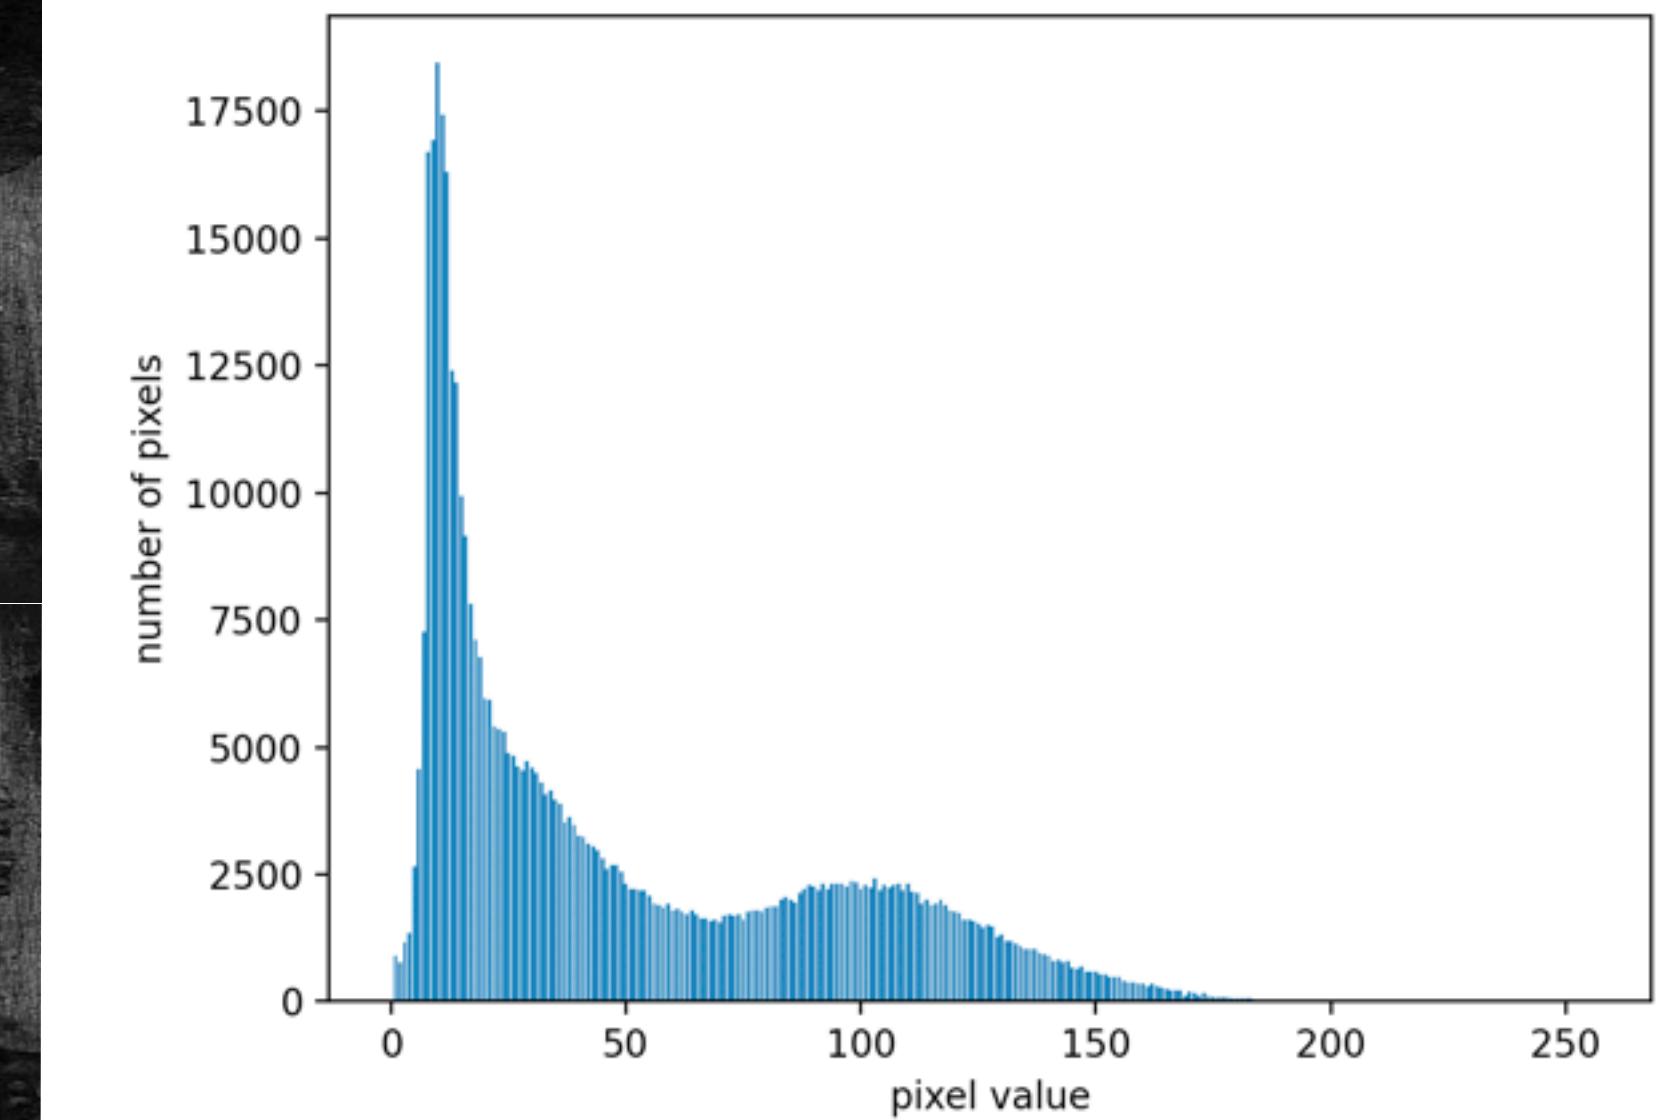


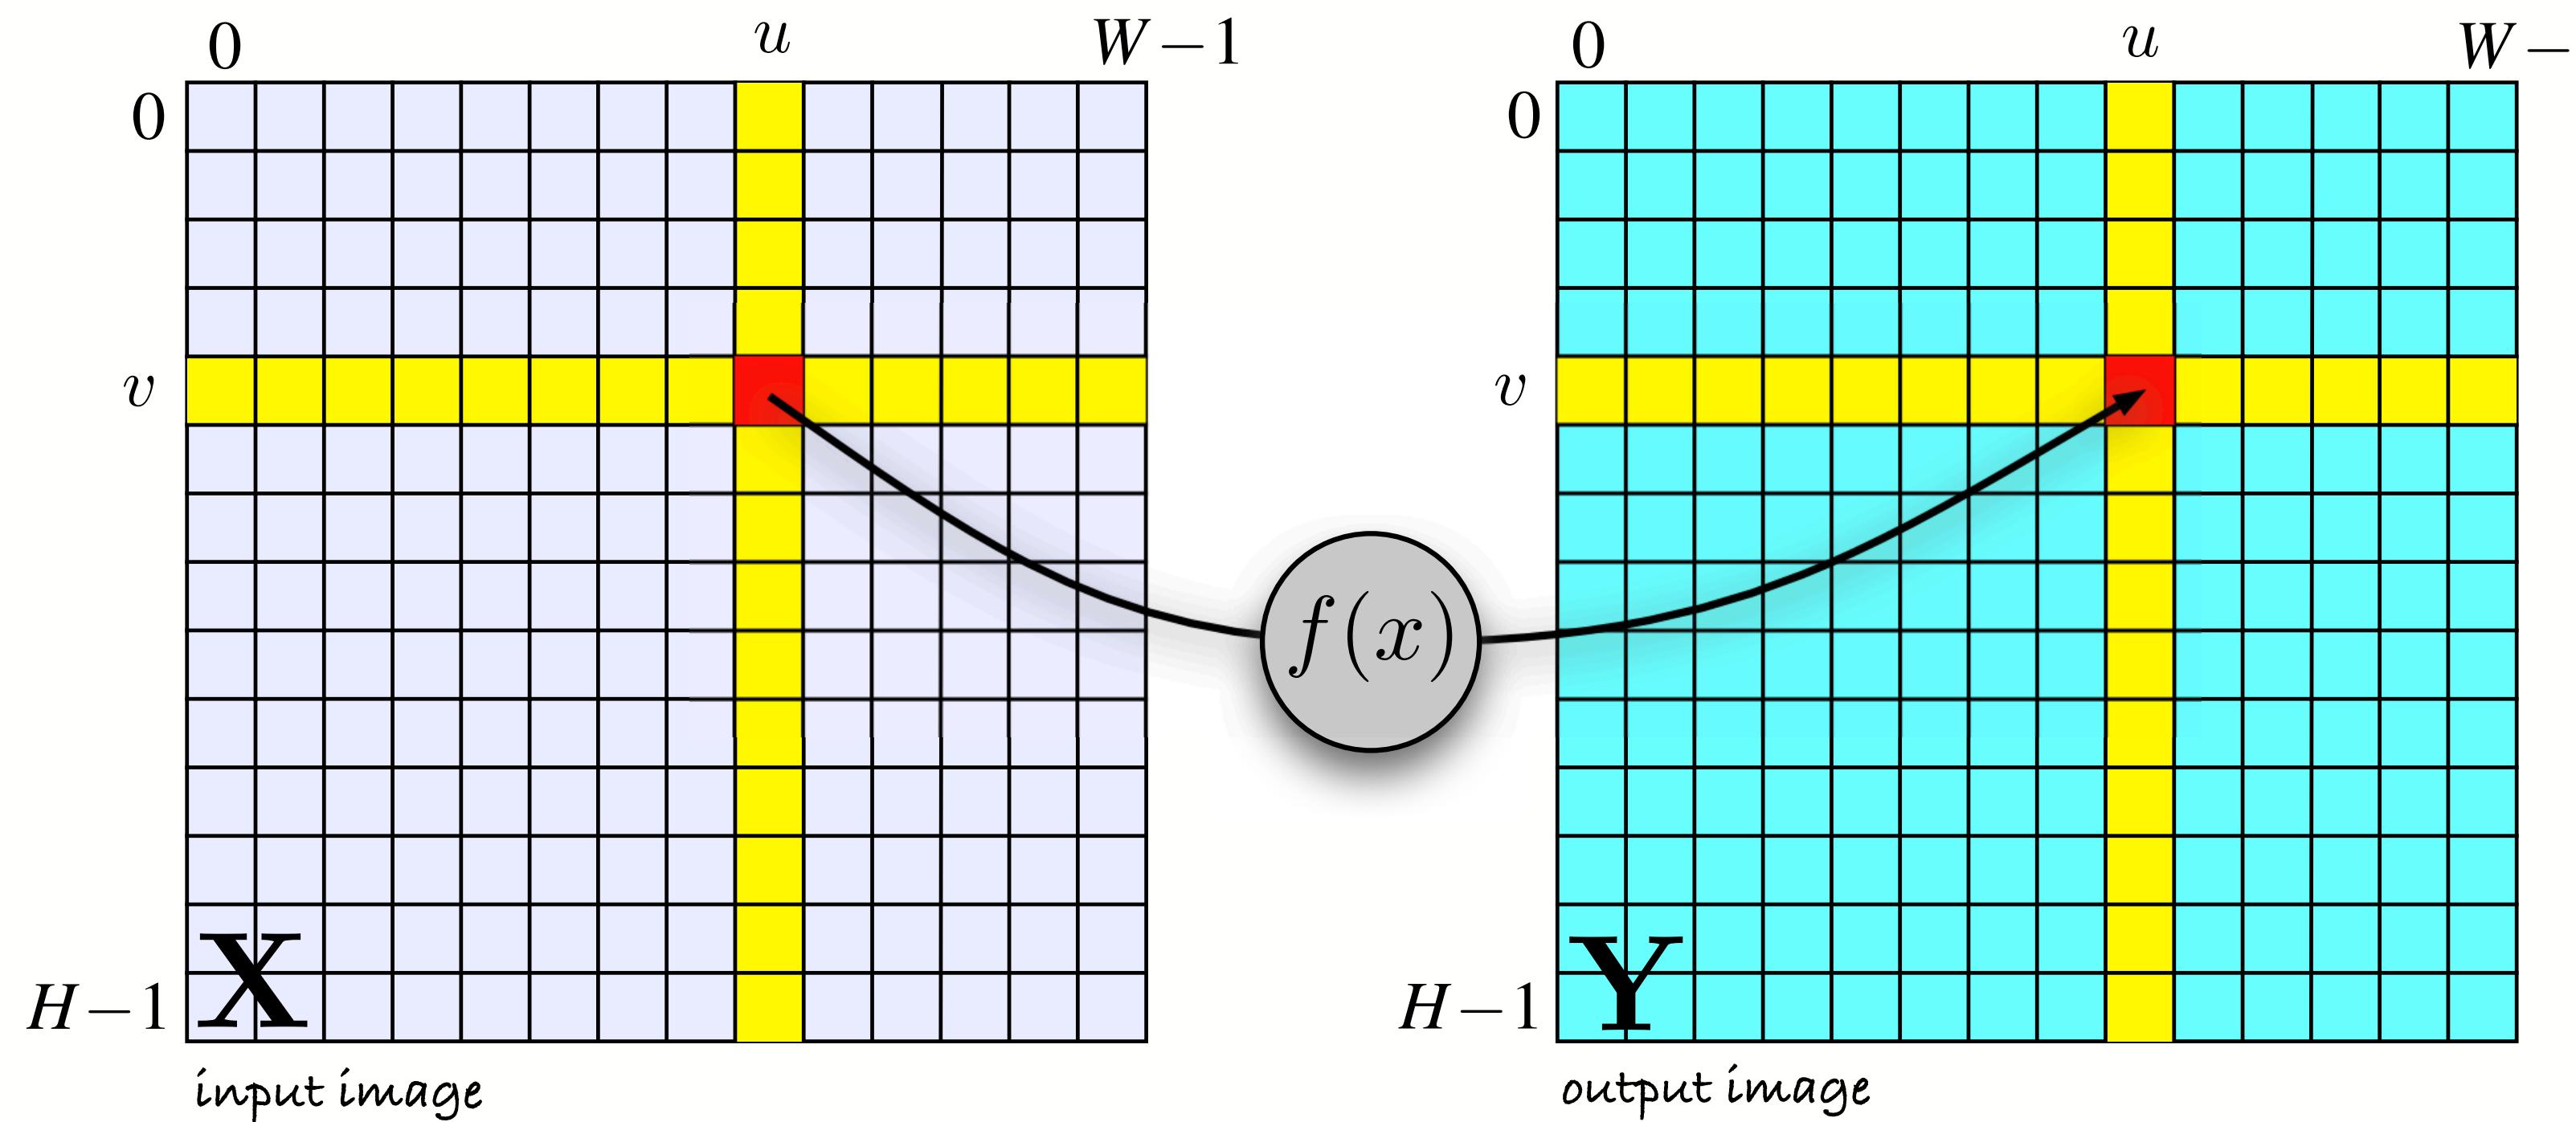
Image histograms



The histogram is invariant to rotation or even complete spatial rearrangement of the pixels

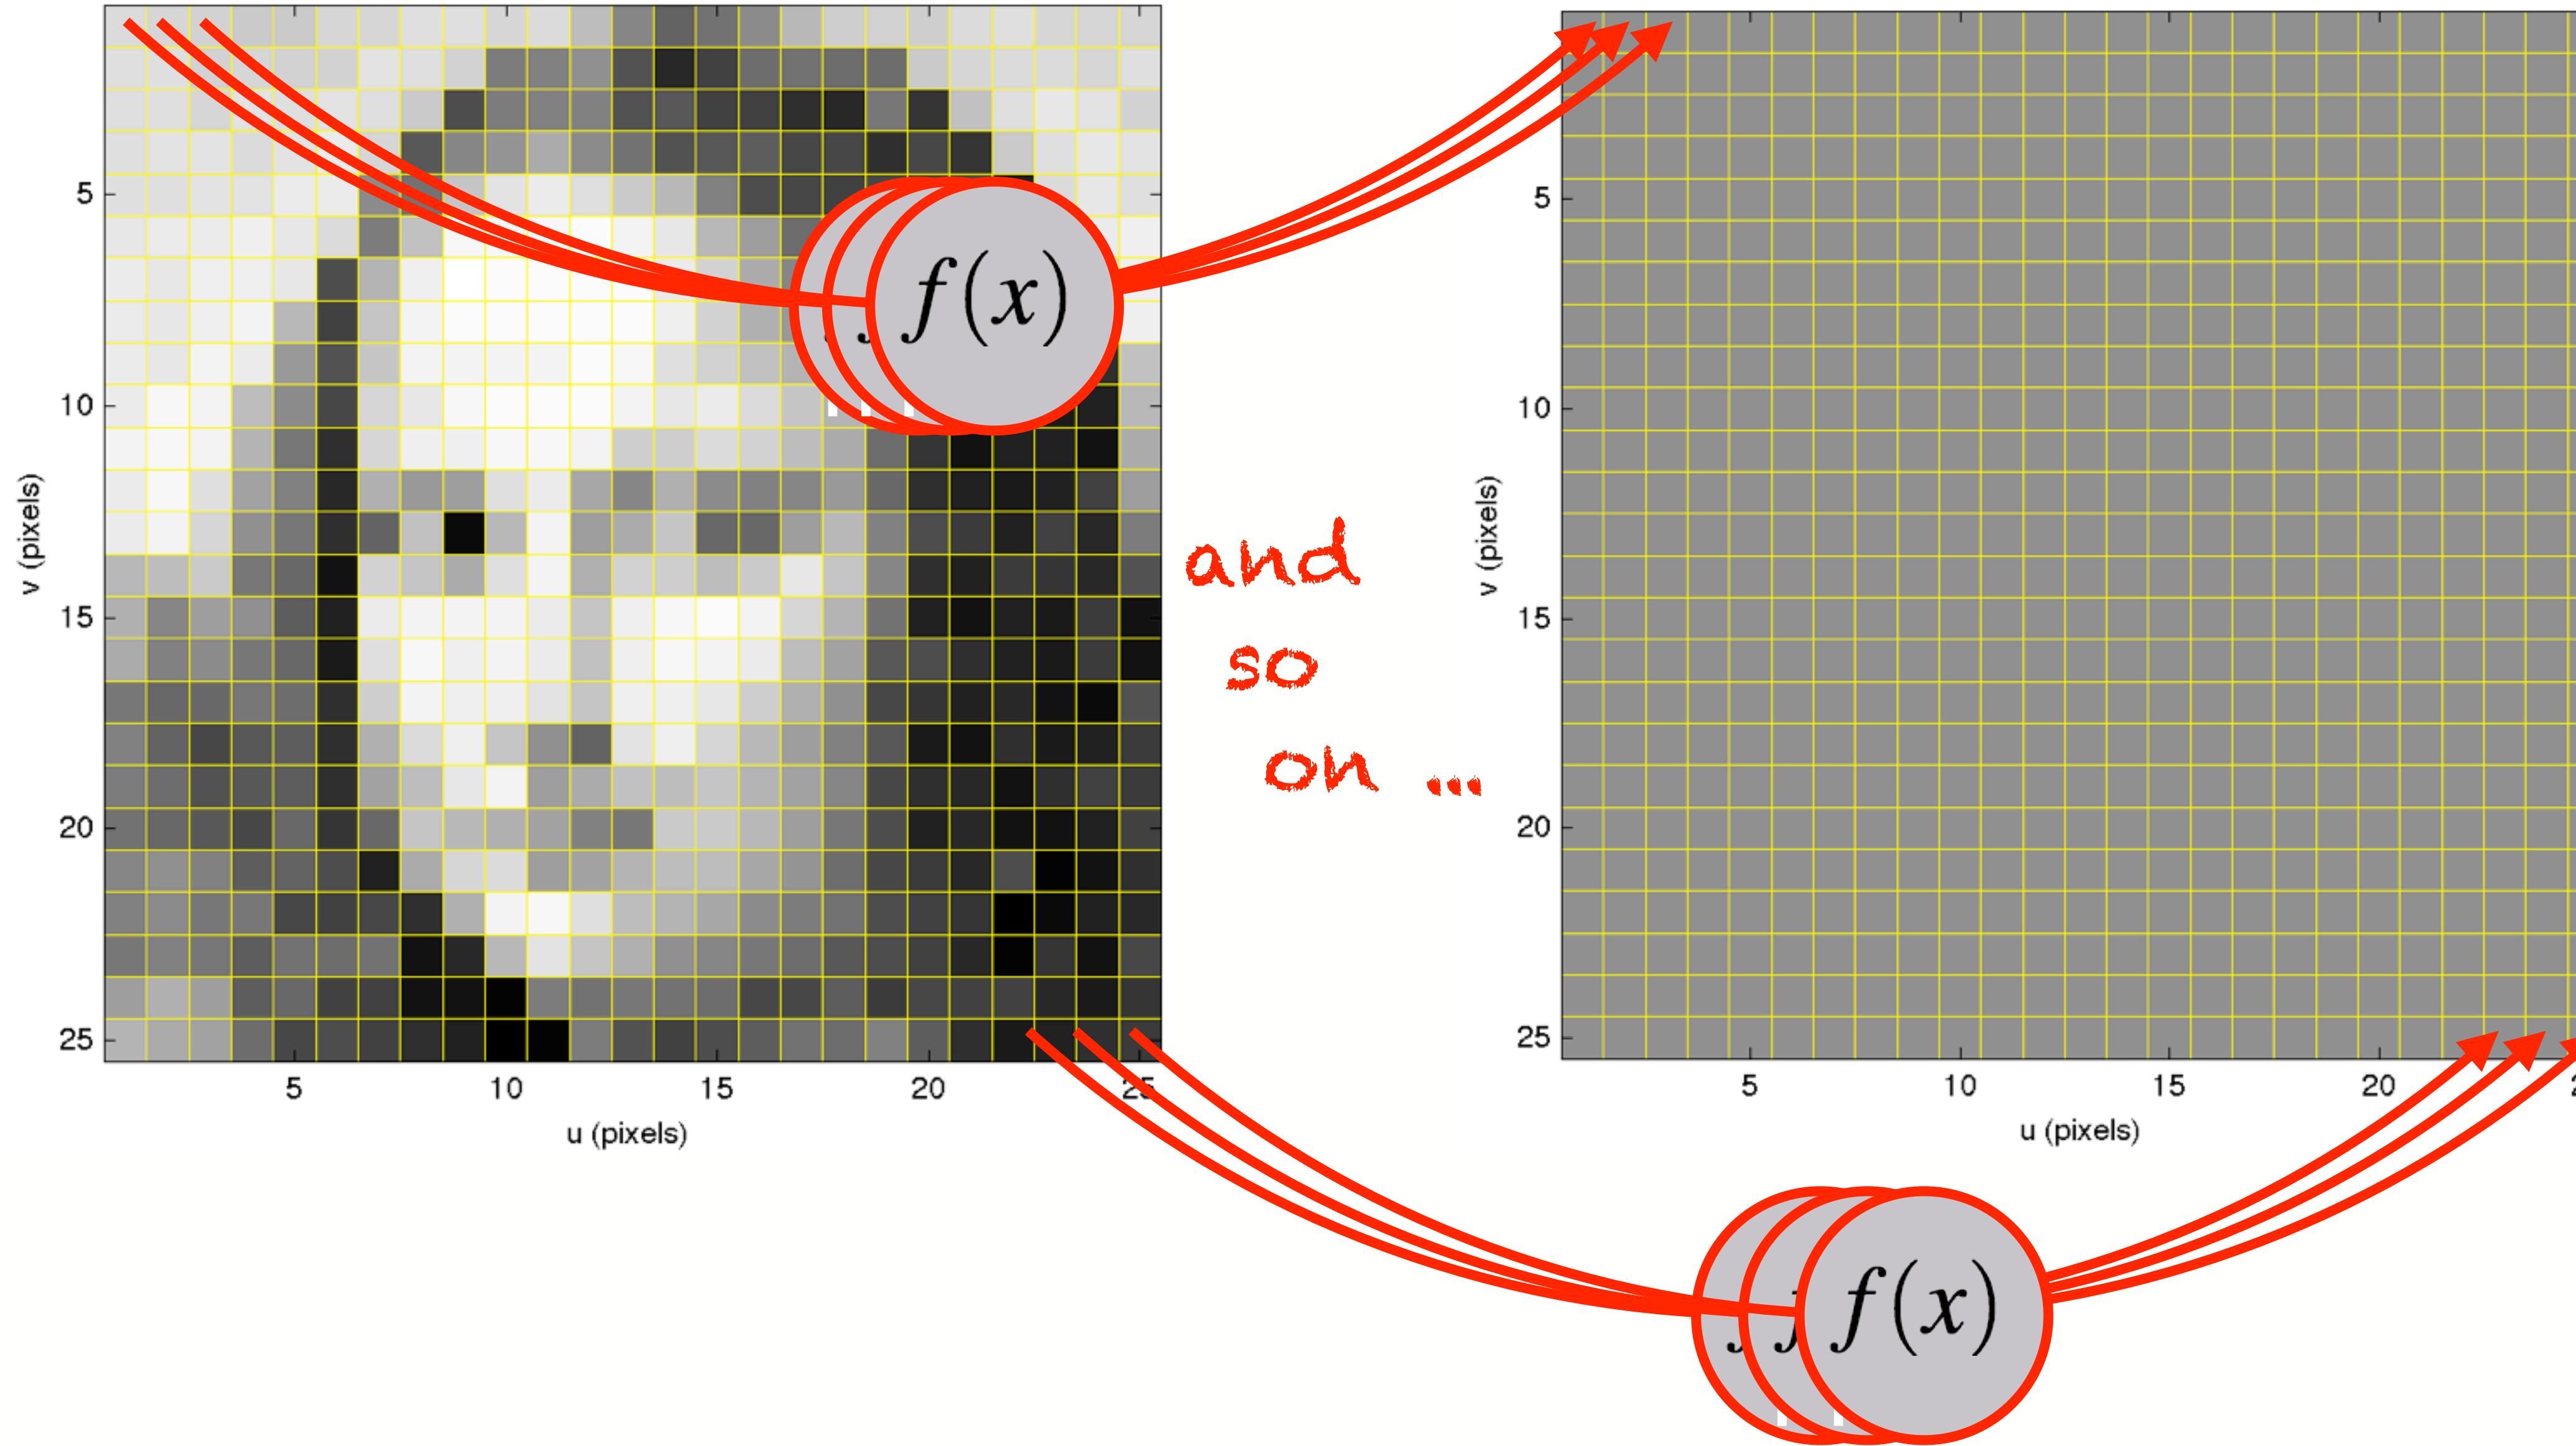


Monadic processing



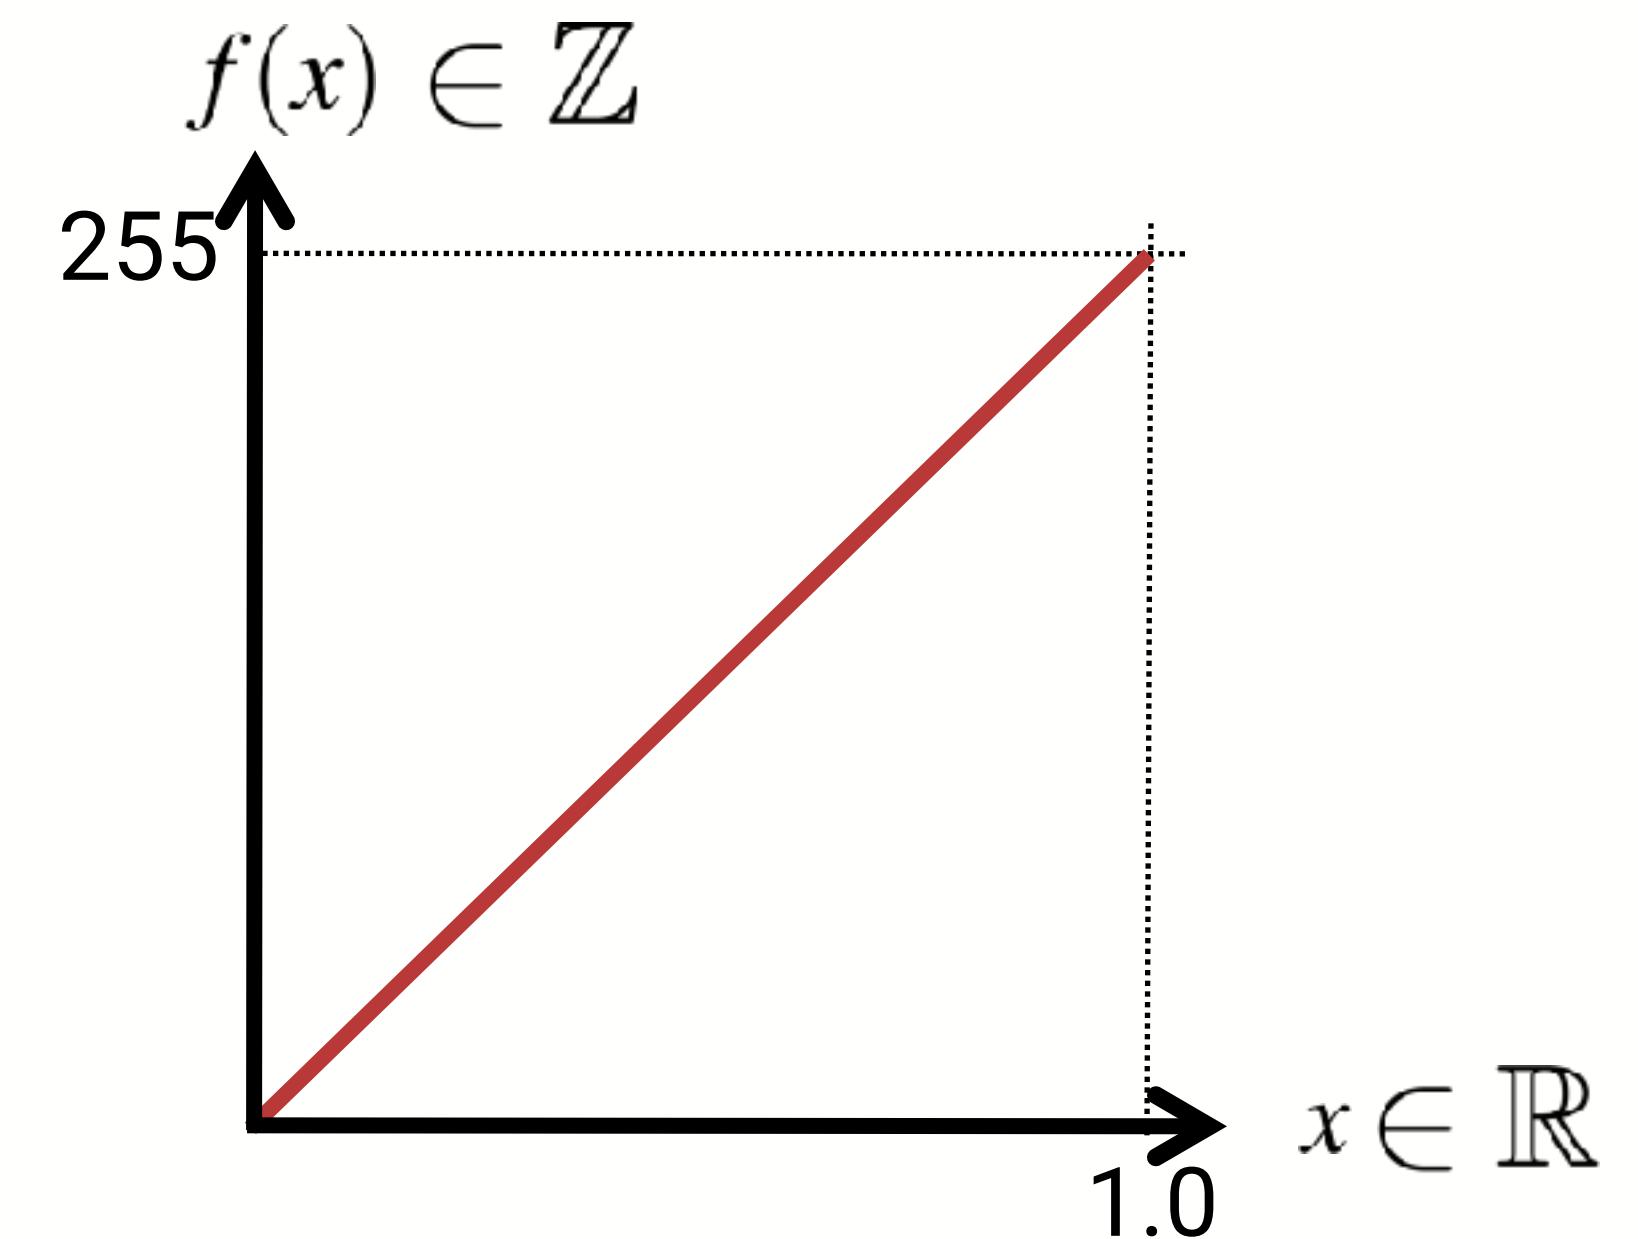
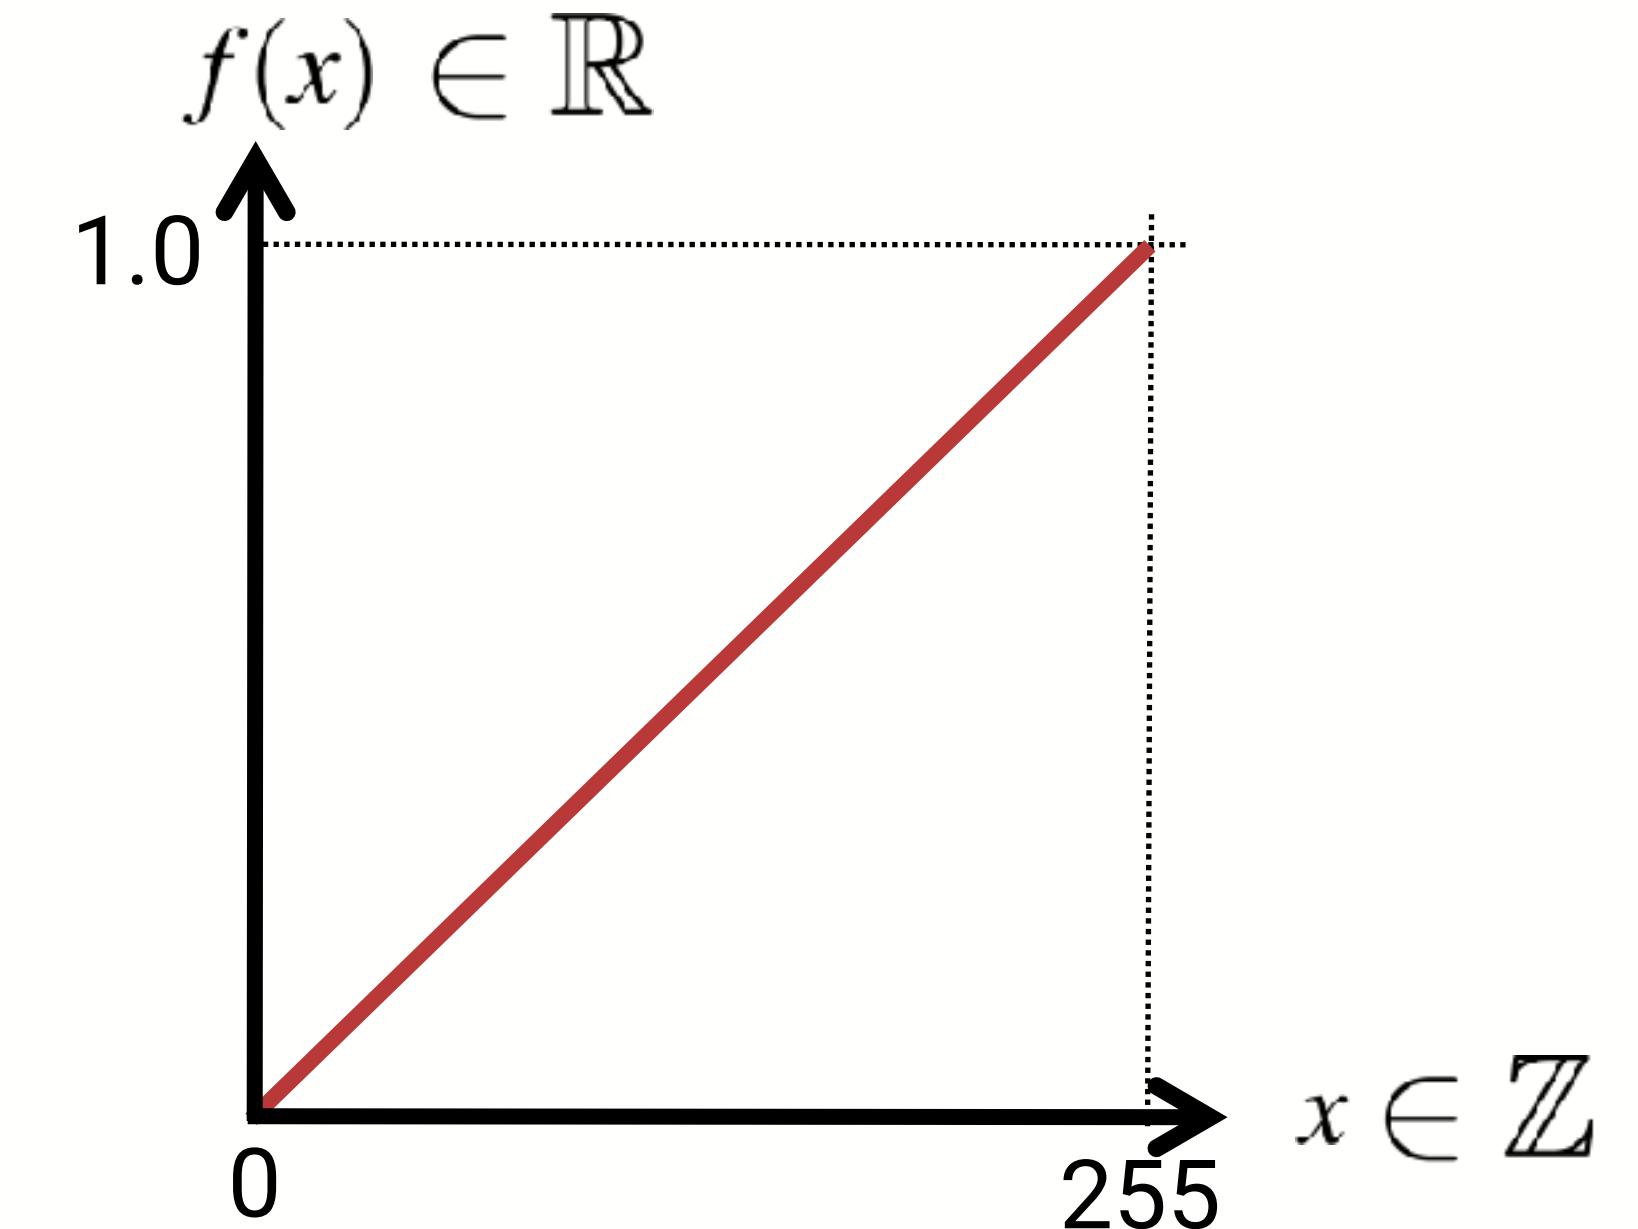
- Each output pixel is a function of corresponding input pixel
- The function is **the same** for all pixels

Monadic image processing



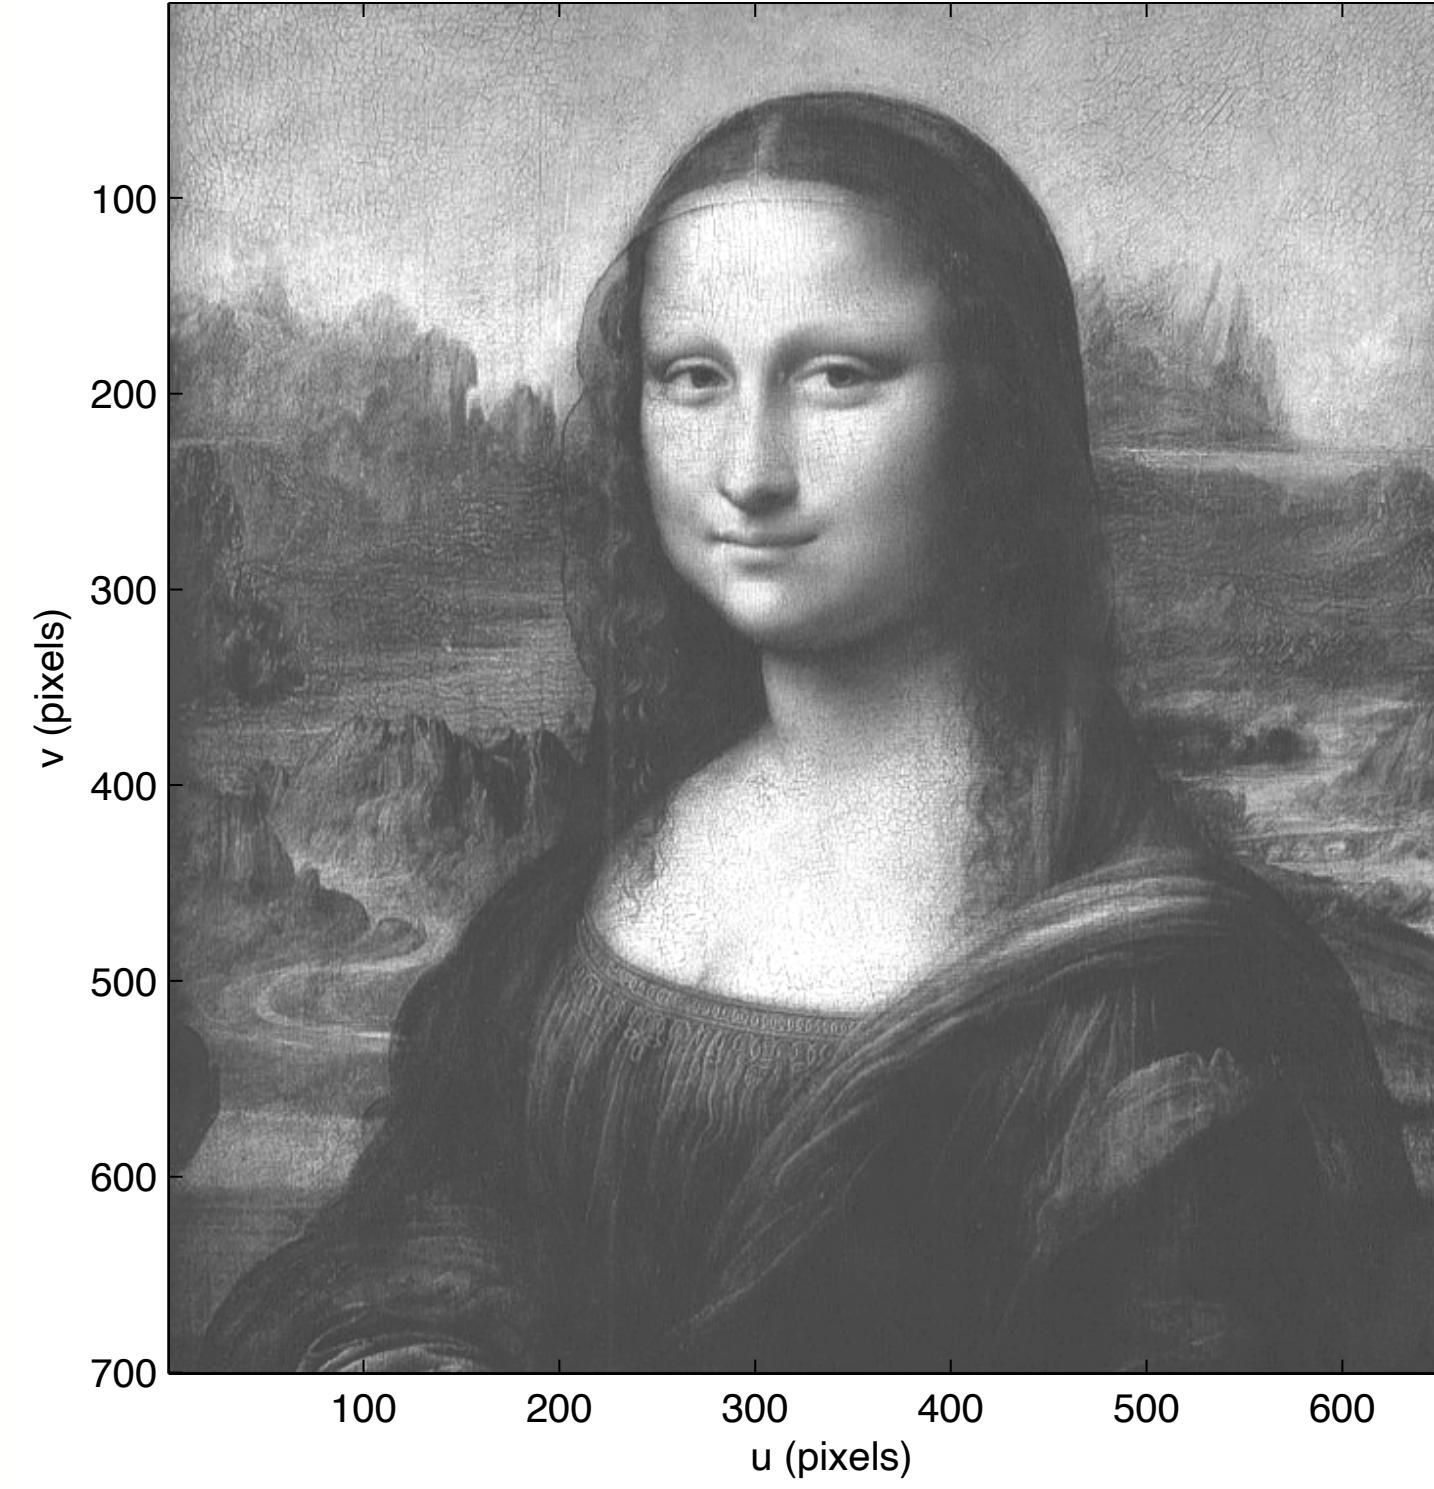
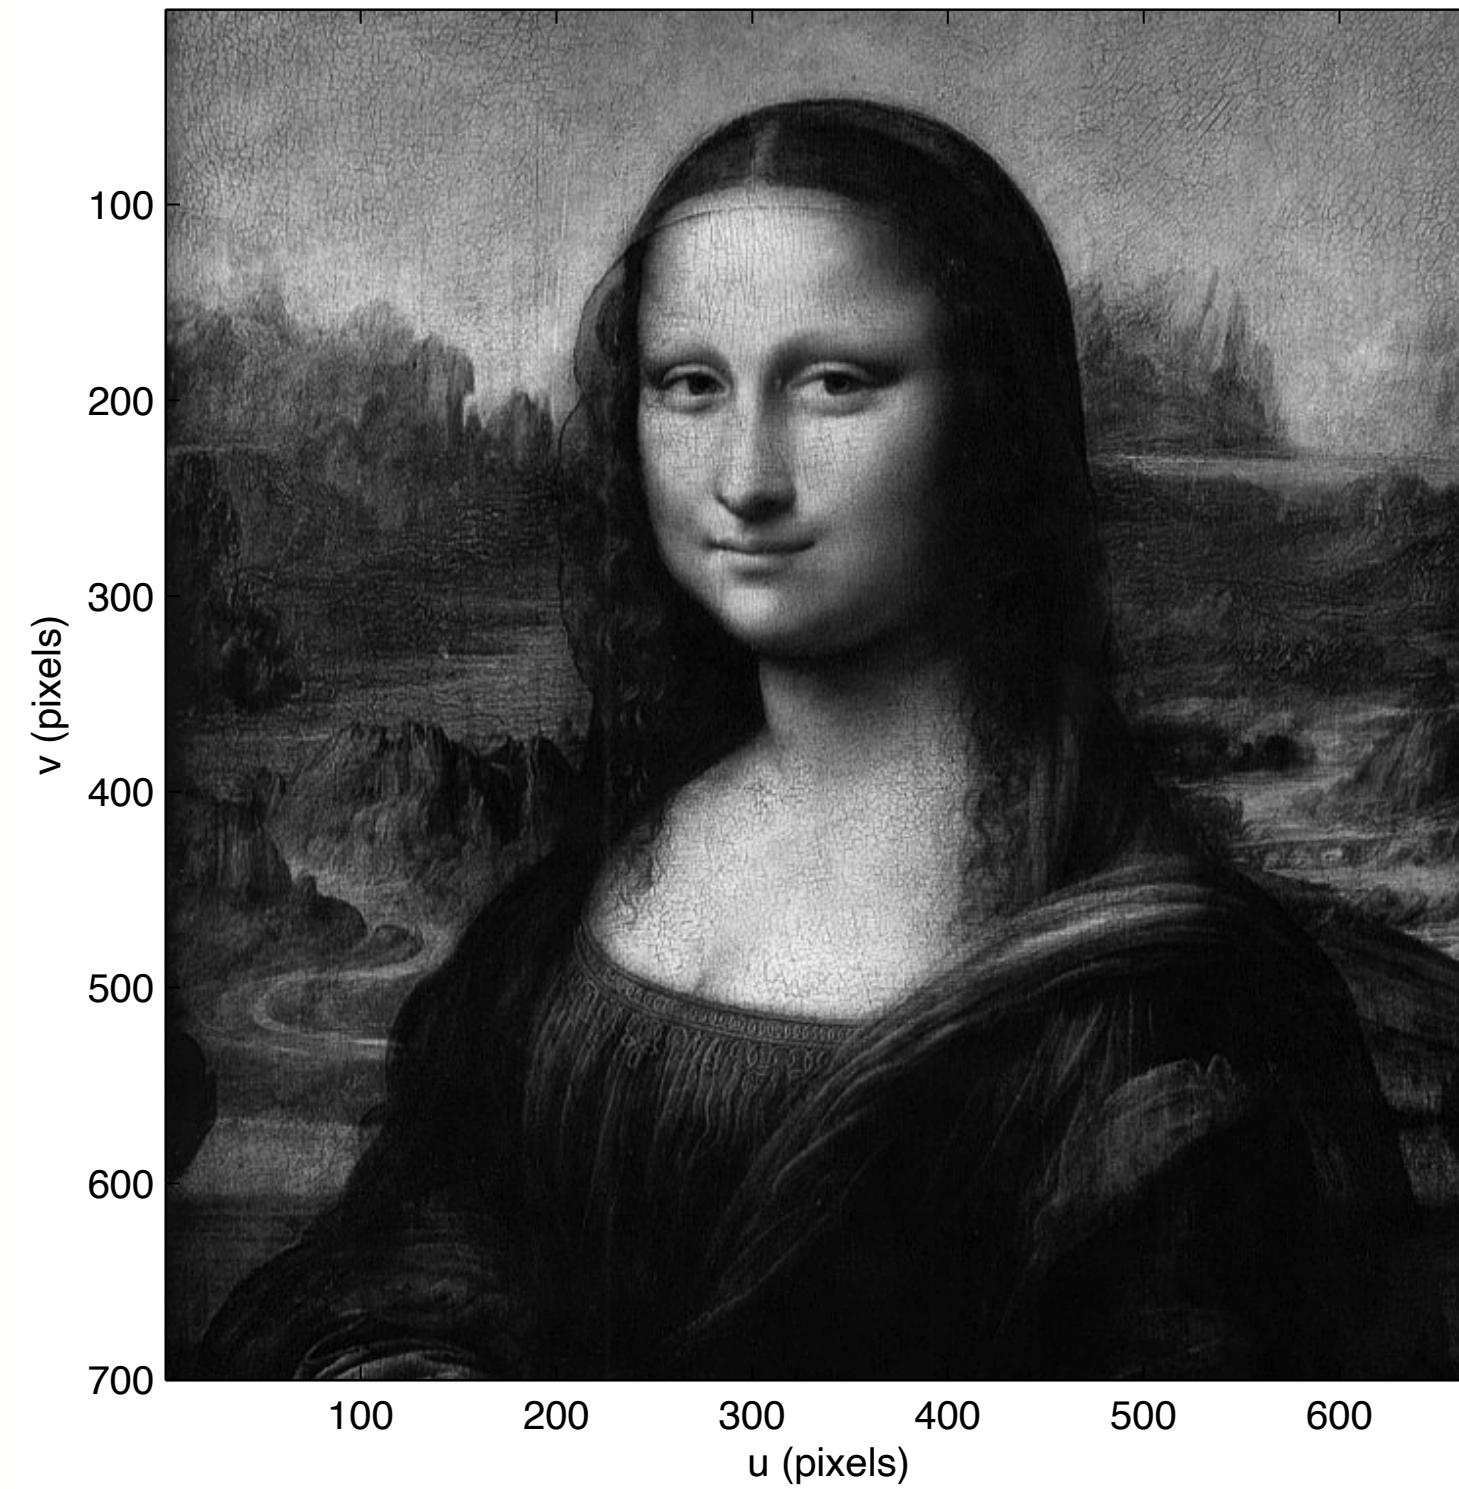
Changing pixel data type

- $\text{uint8} \rightarrow \text{float32}$
 - pixels in the range $[0, 1 \dots 255]$ mapped to $[0.0 \dots 1.0]$
- $\text{float32} \rightarrow \text{uint8}$
 - pixels in the range $[0.0 \dots 1.0]$ mapped to $[0, 1 \dots 255]$

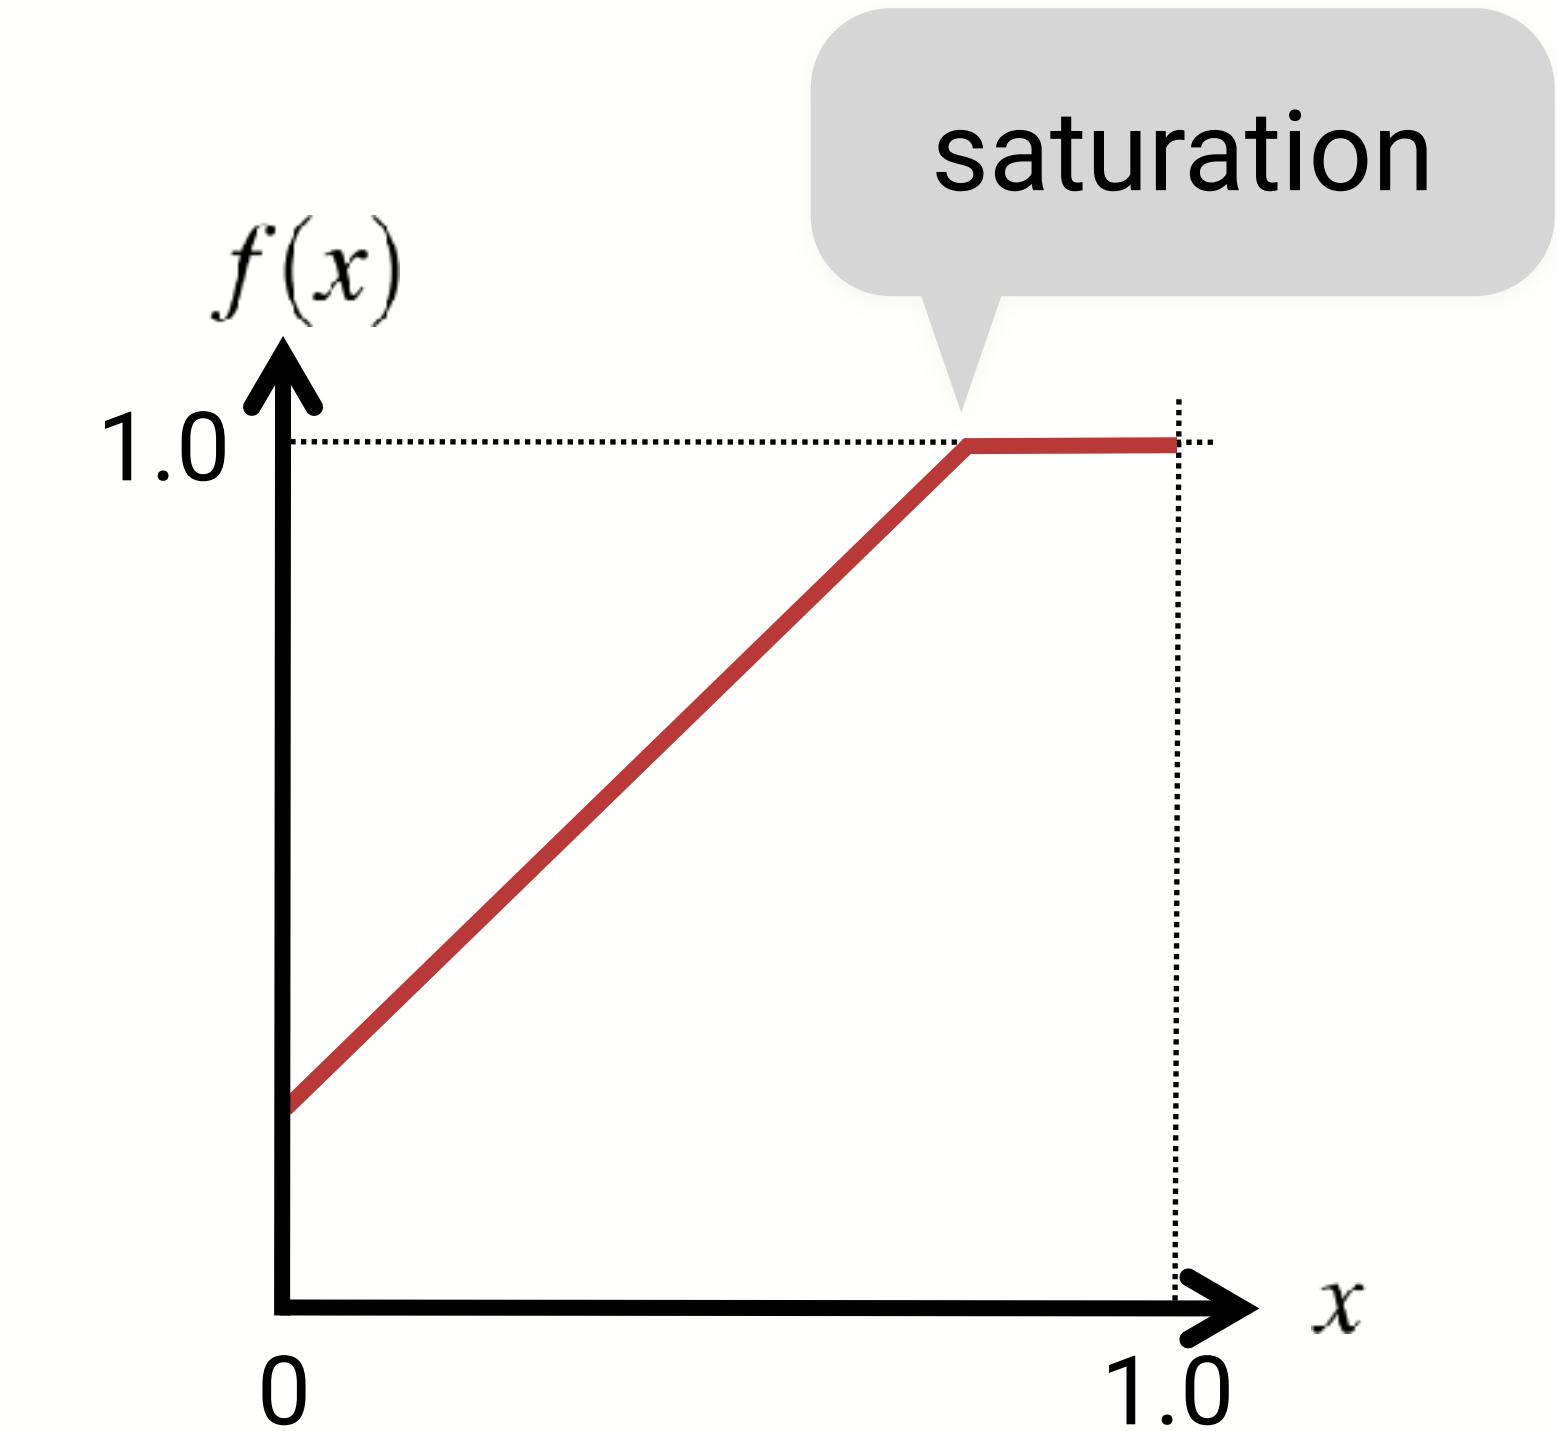


Changing brightness

- by offsetting the grey values
→ $f(x) = x + 0.25$

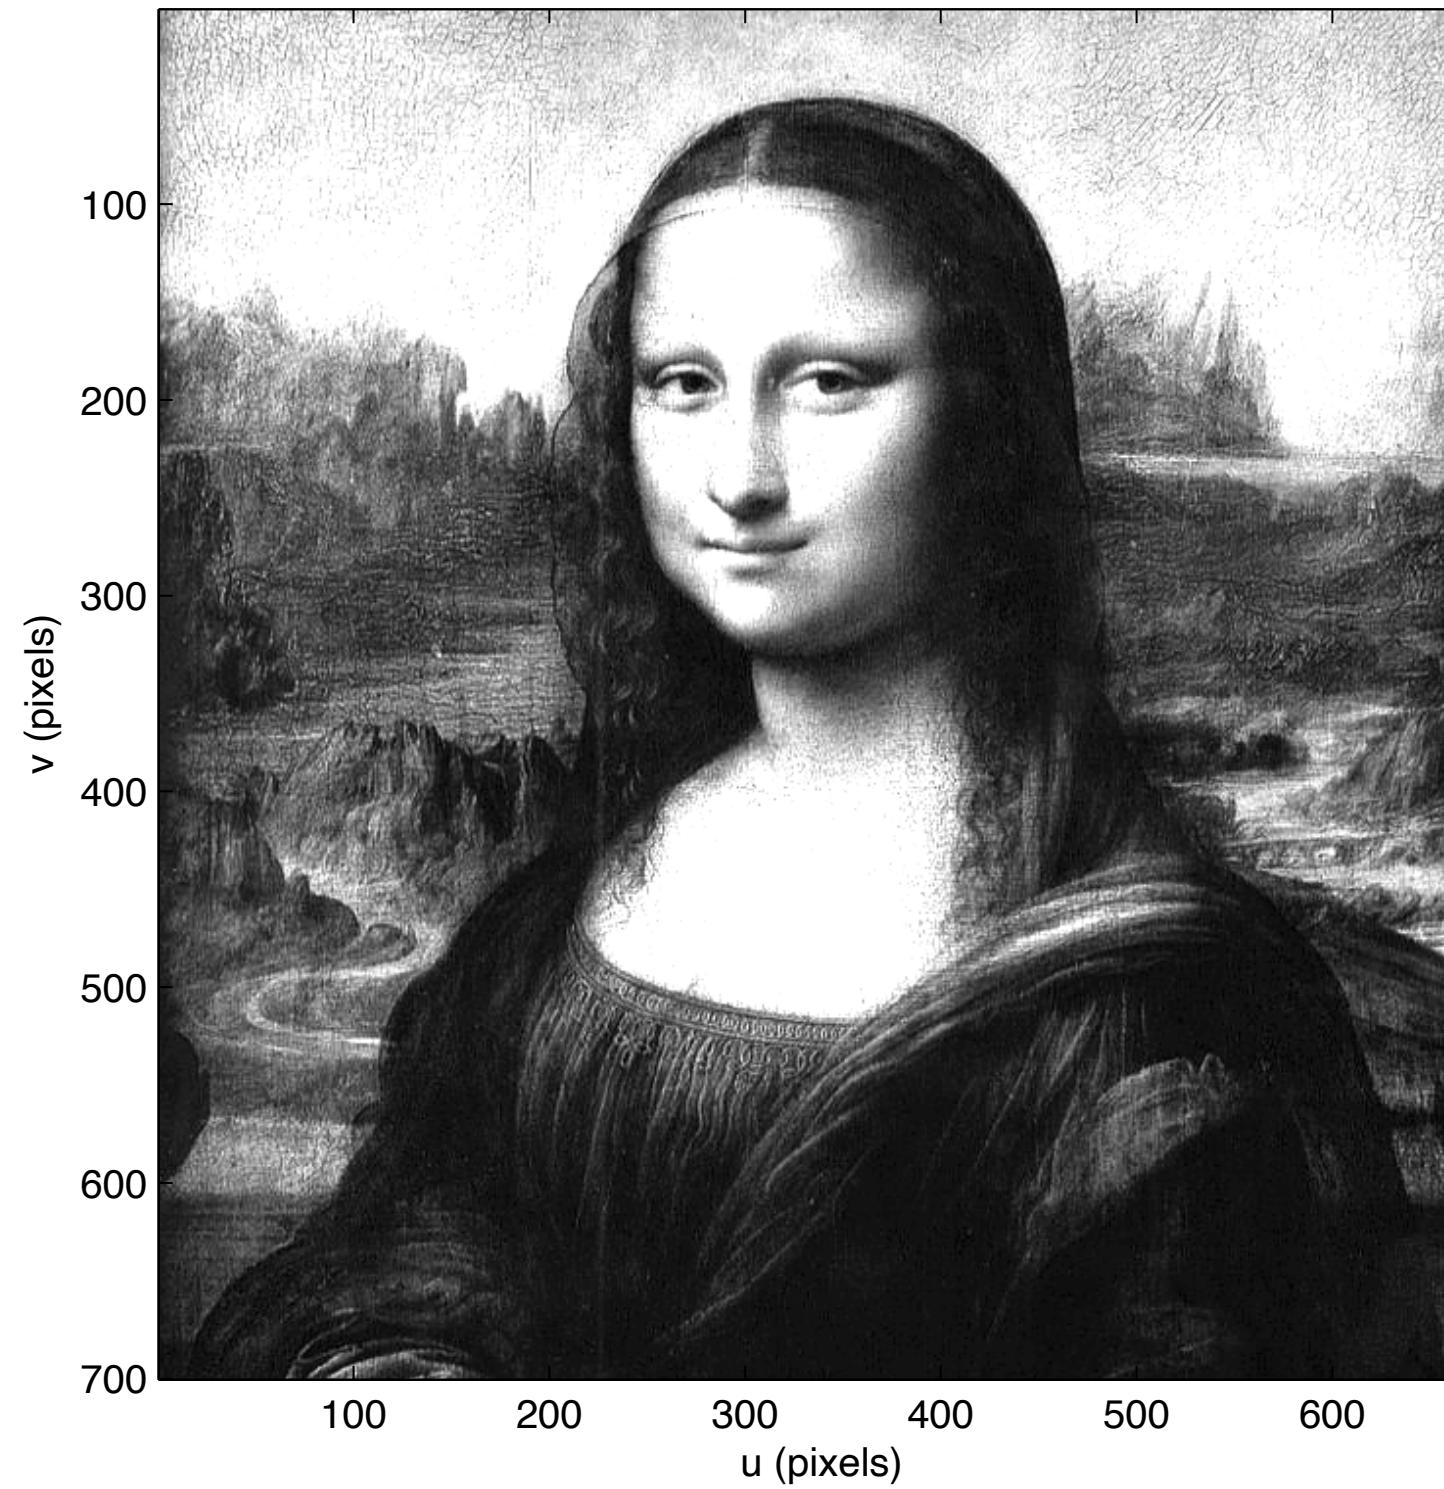
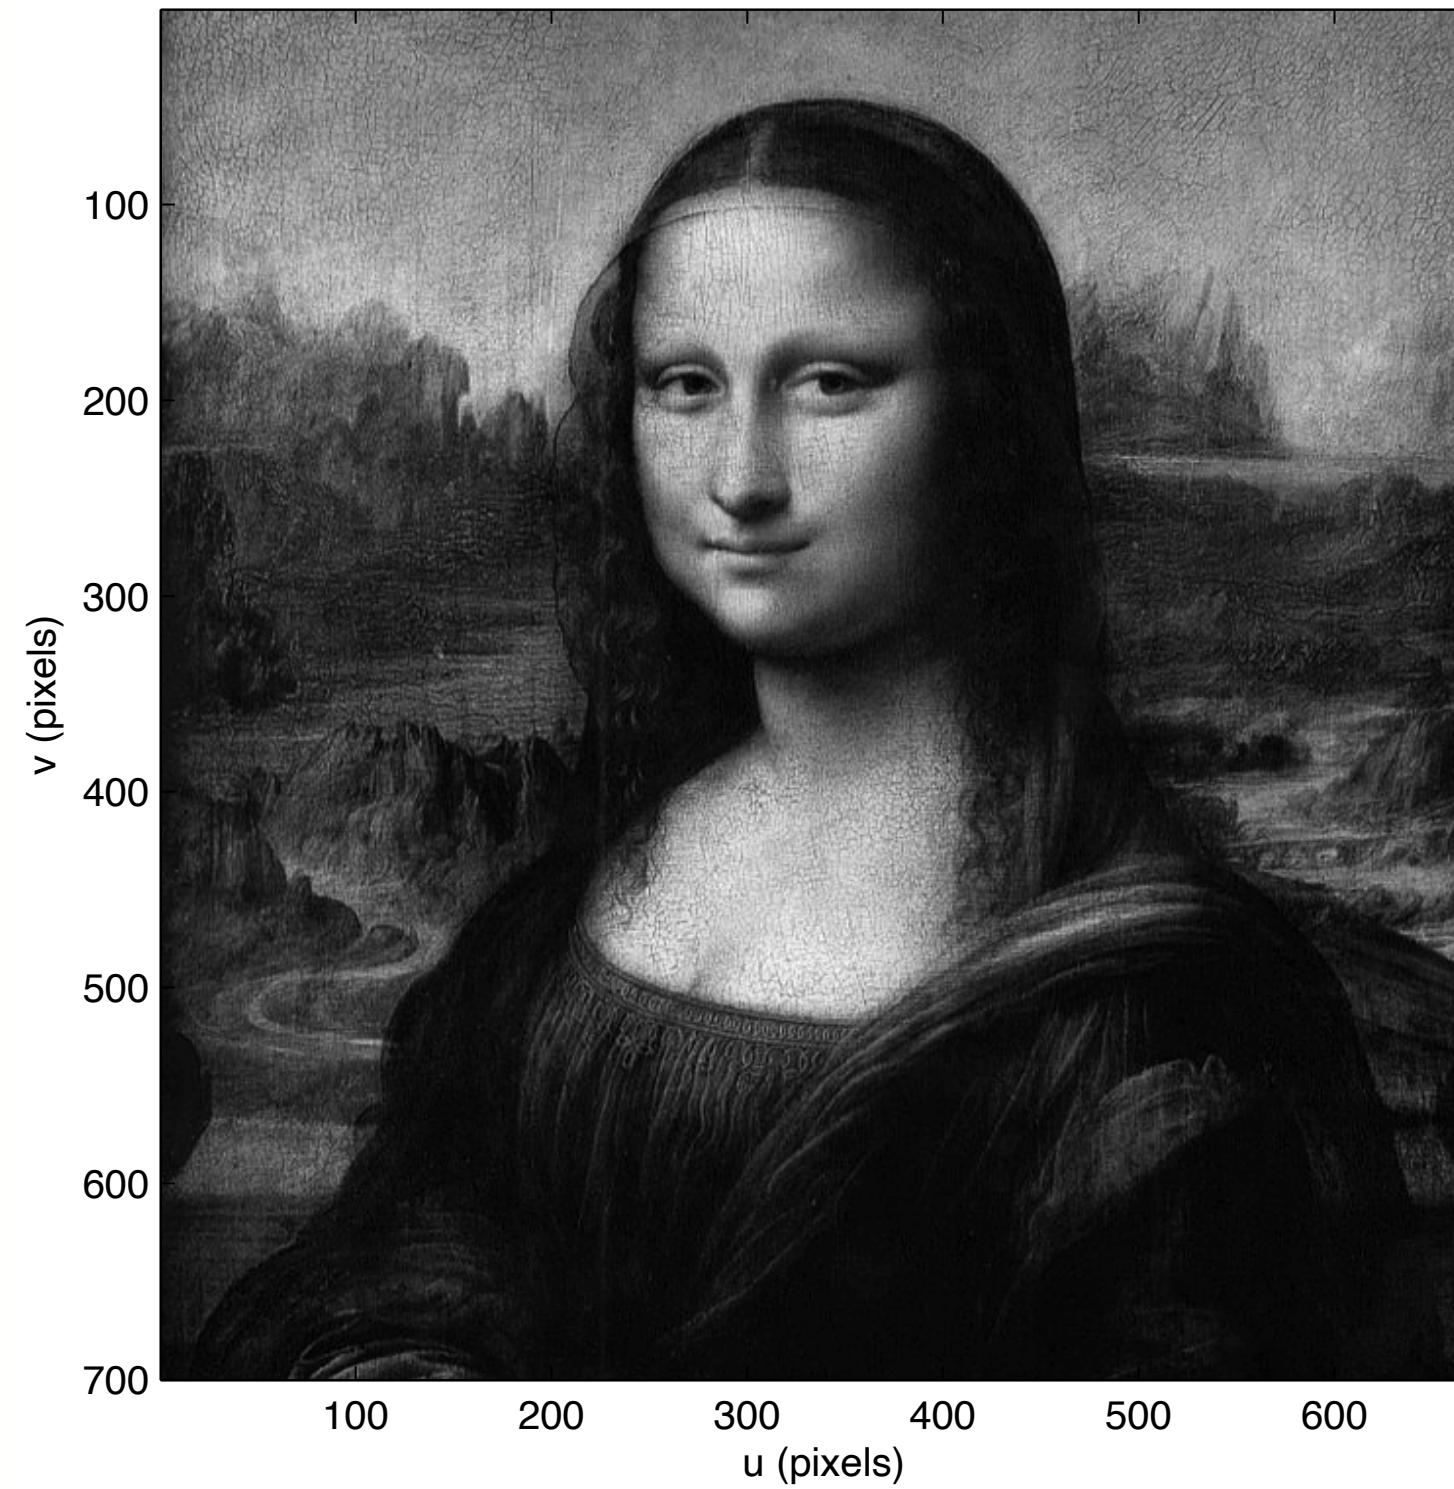


Mona Lisa c.1503–1505
Da Vinci, L.

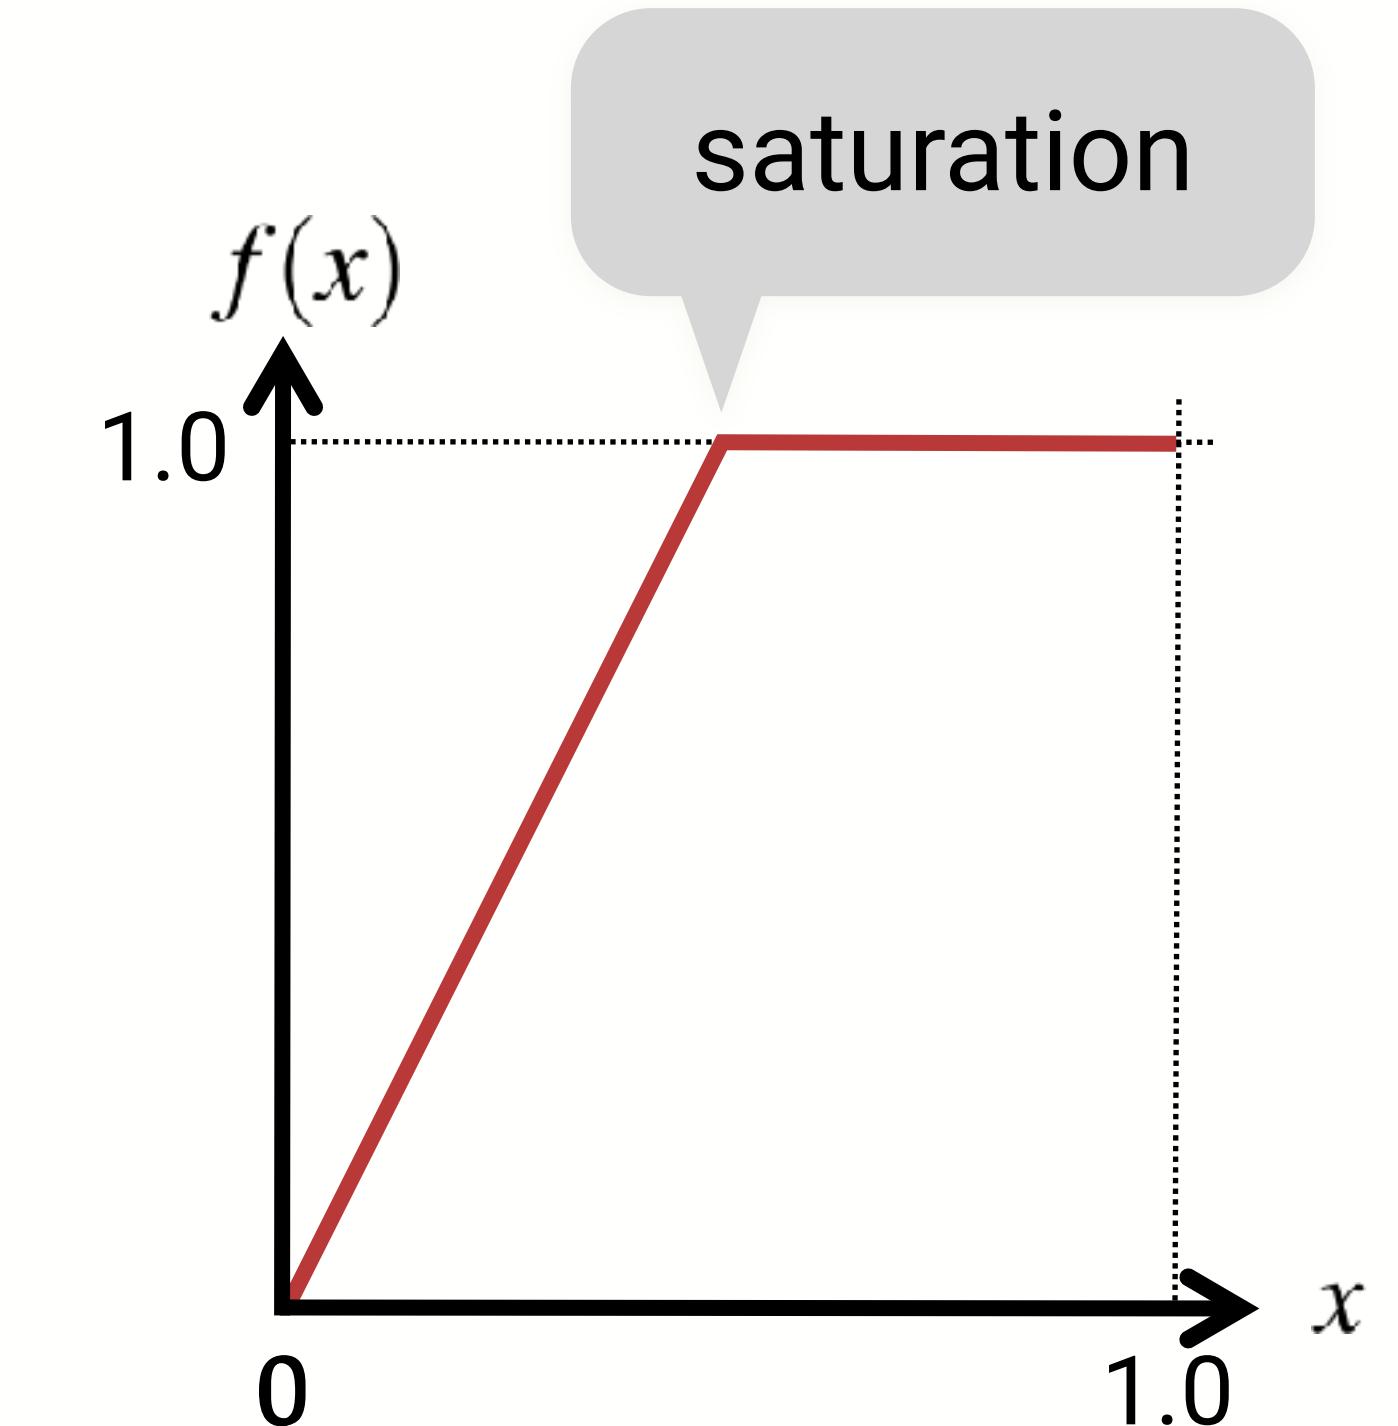


Changing contrast

- by scaling the grey values
→ $f(x) = 2x$

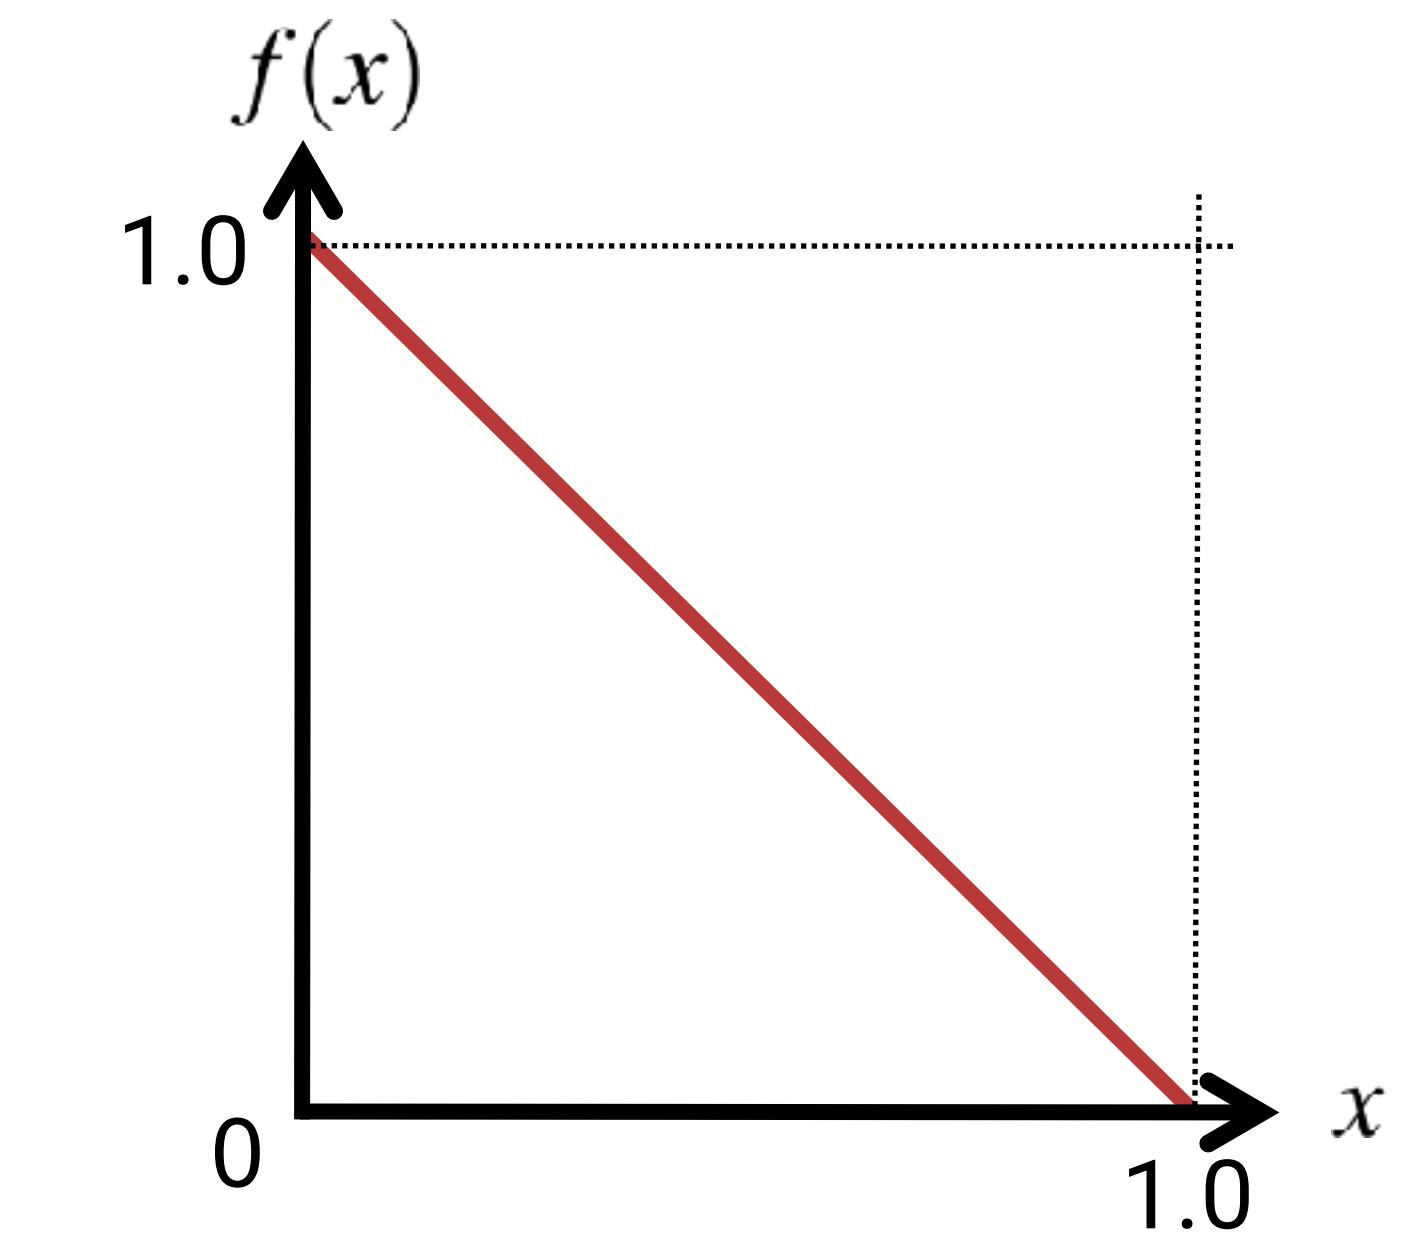
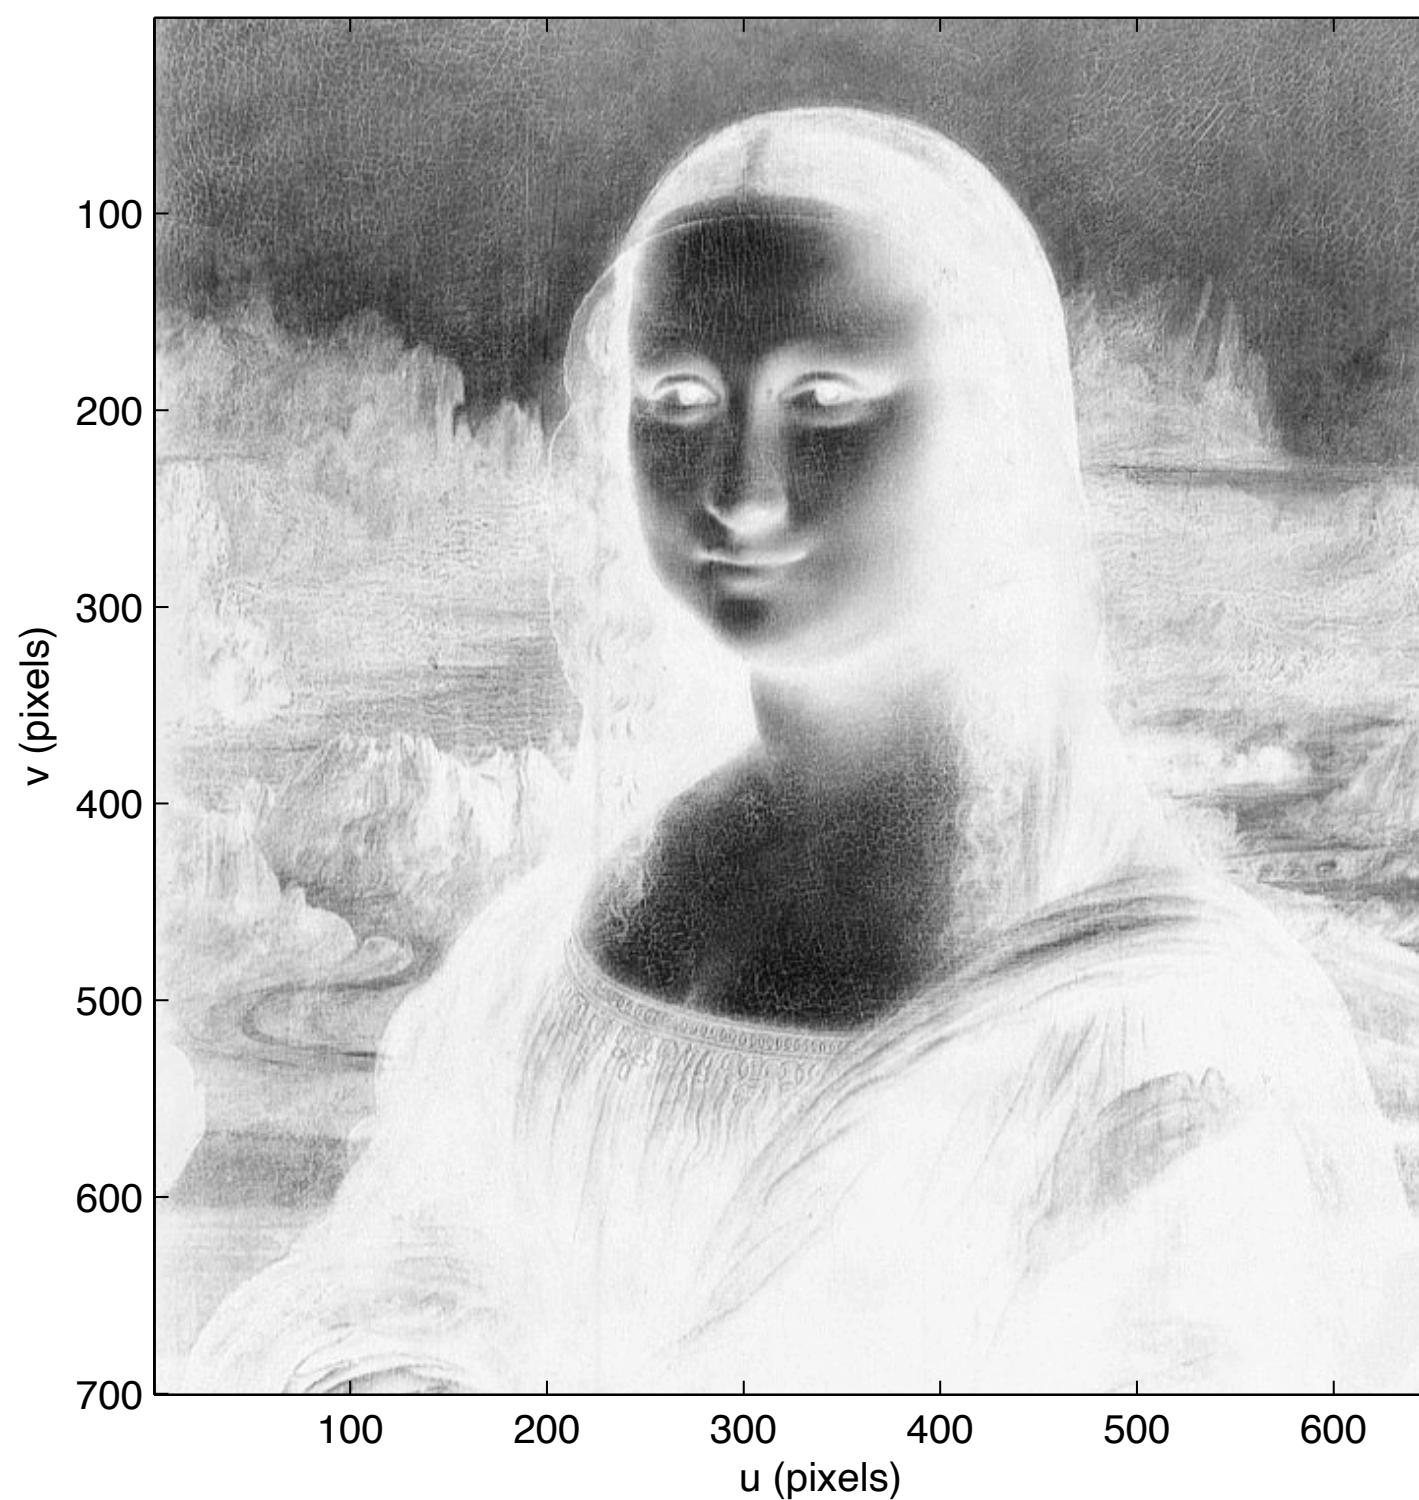
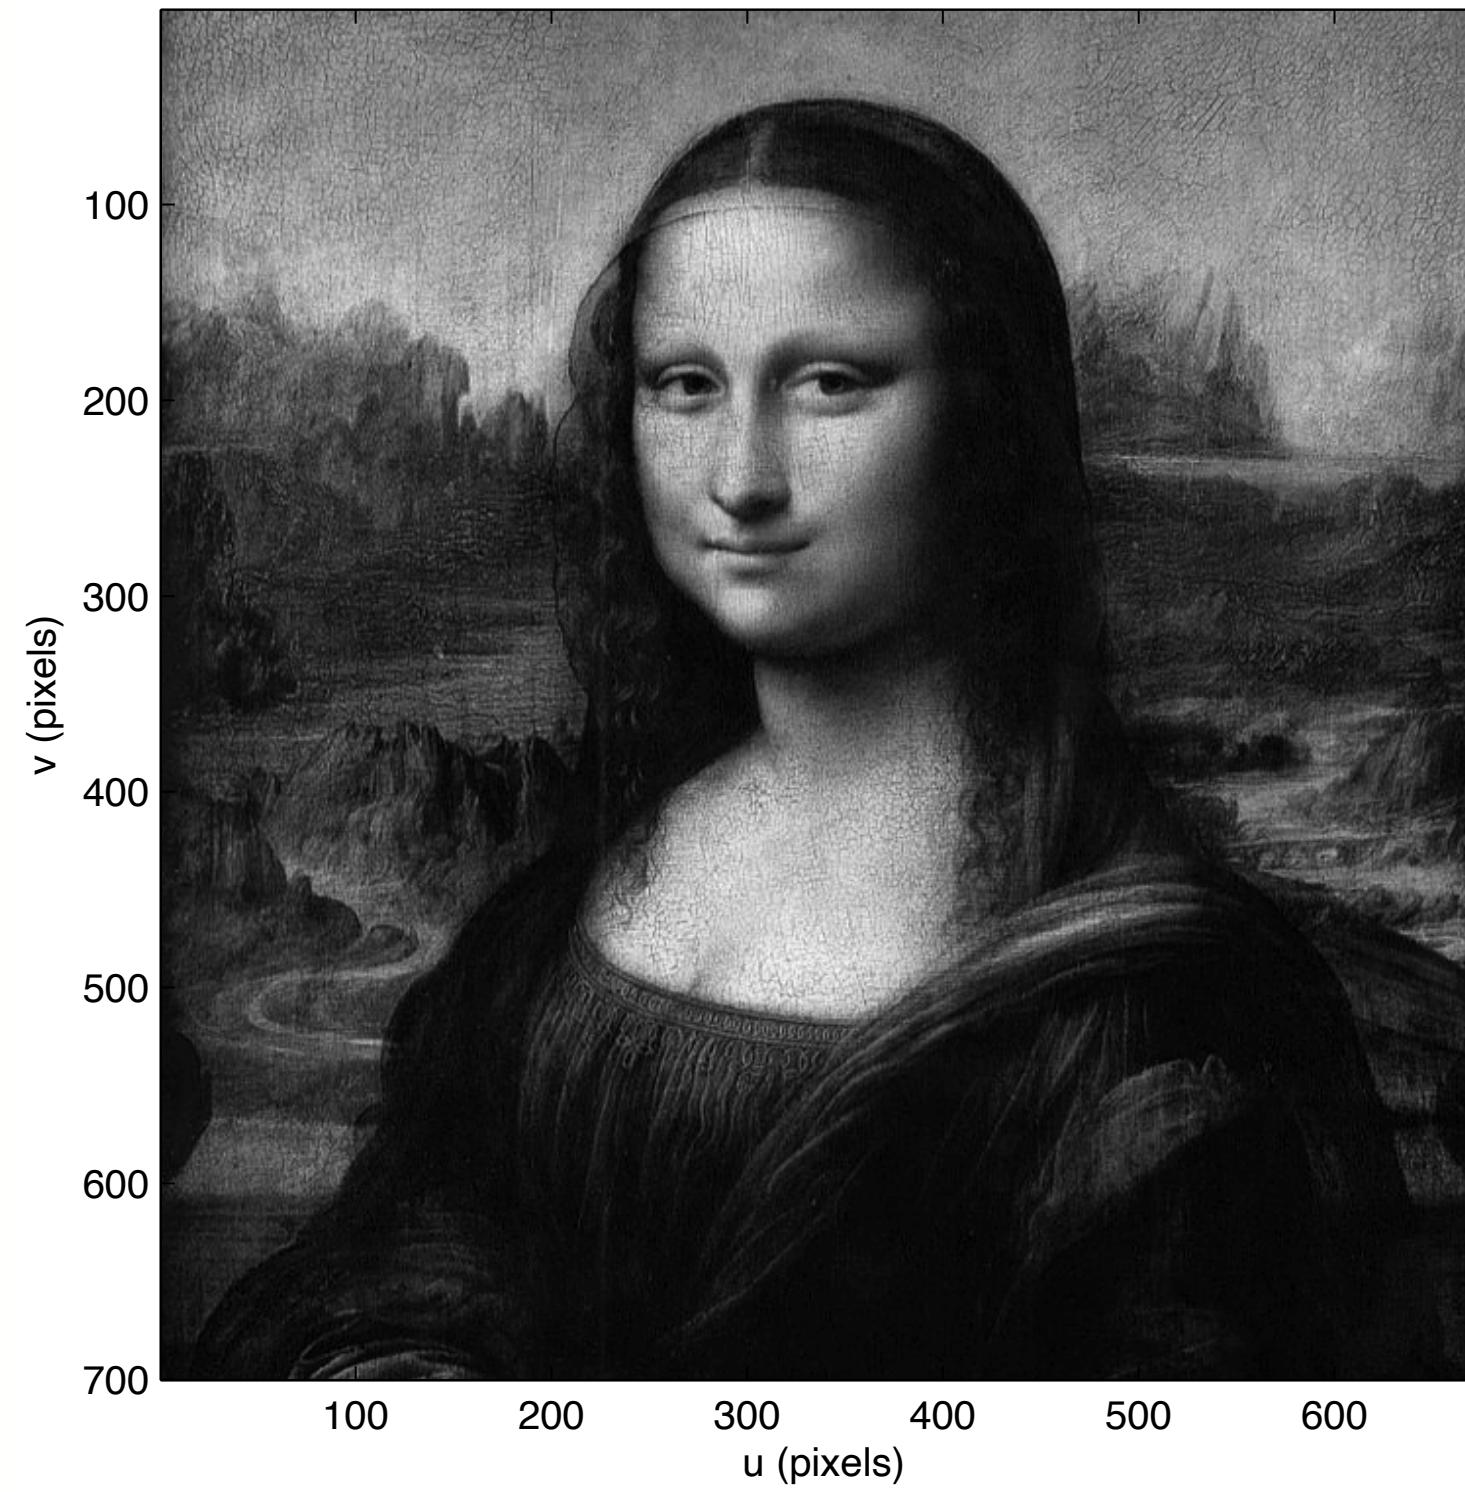


Mona Lisa c.1503–1505
Da Vinci, L.



Negative image

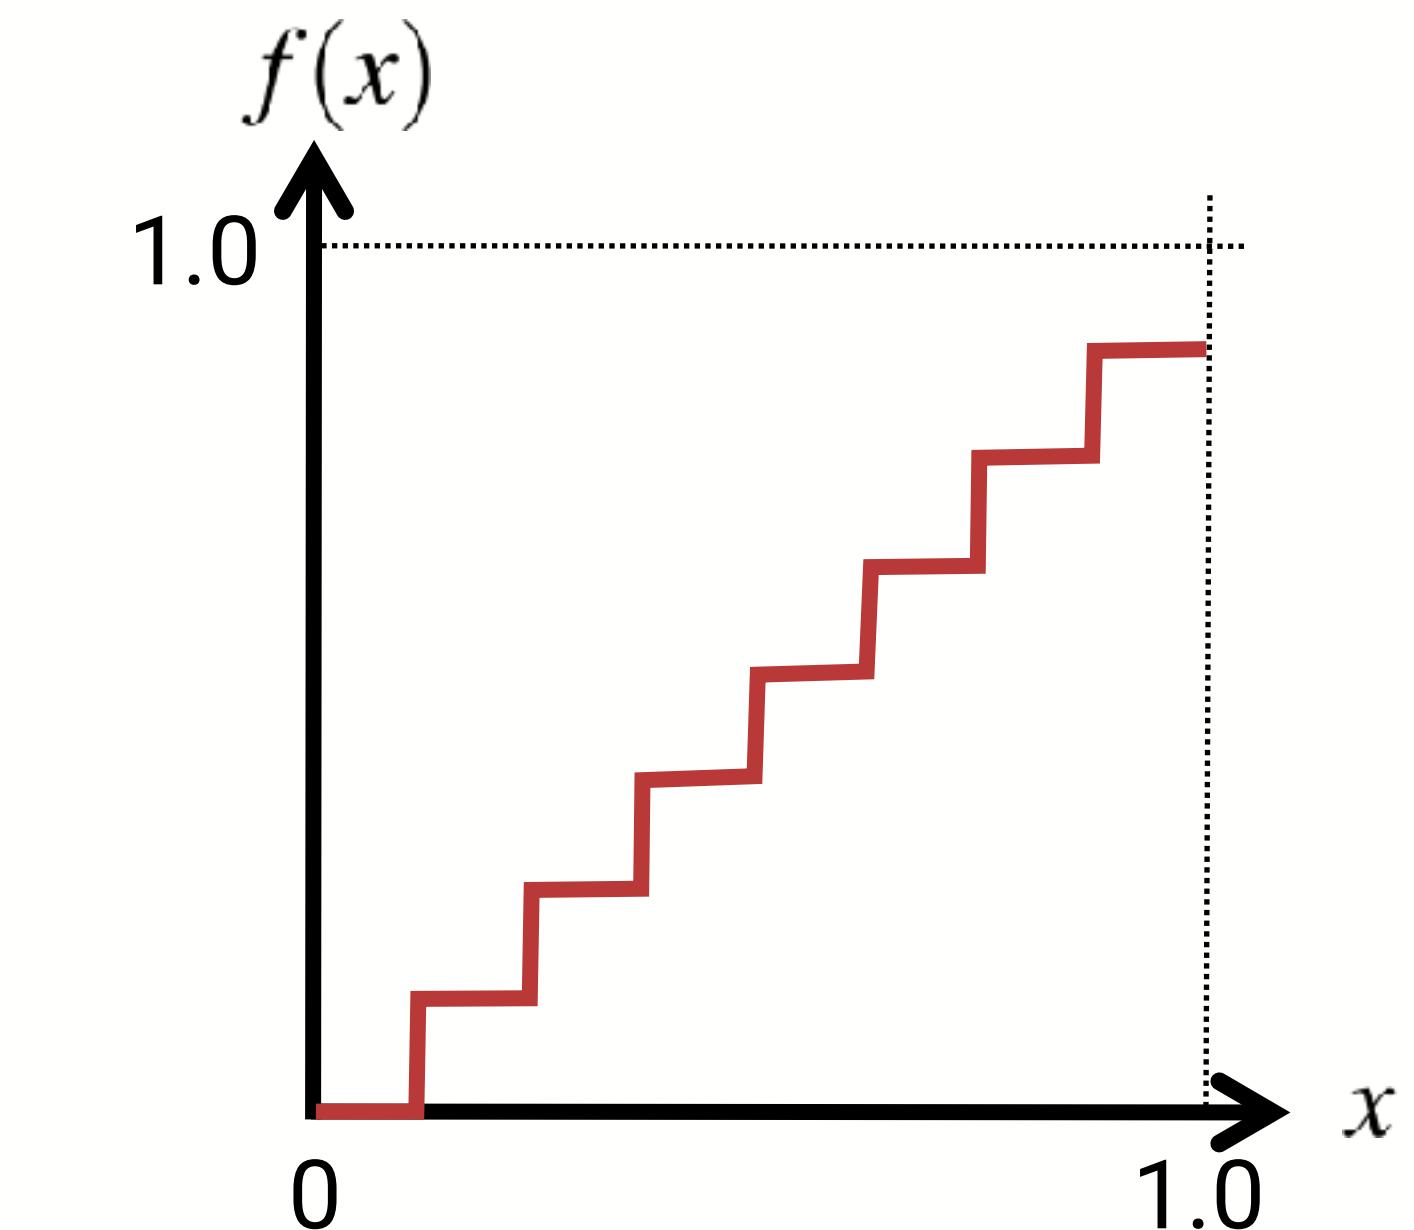
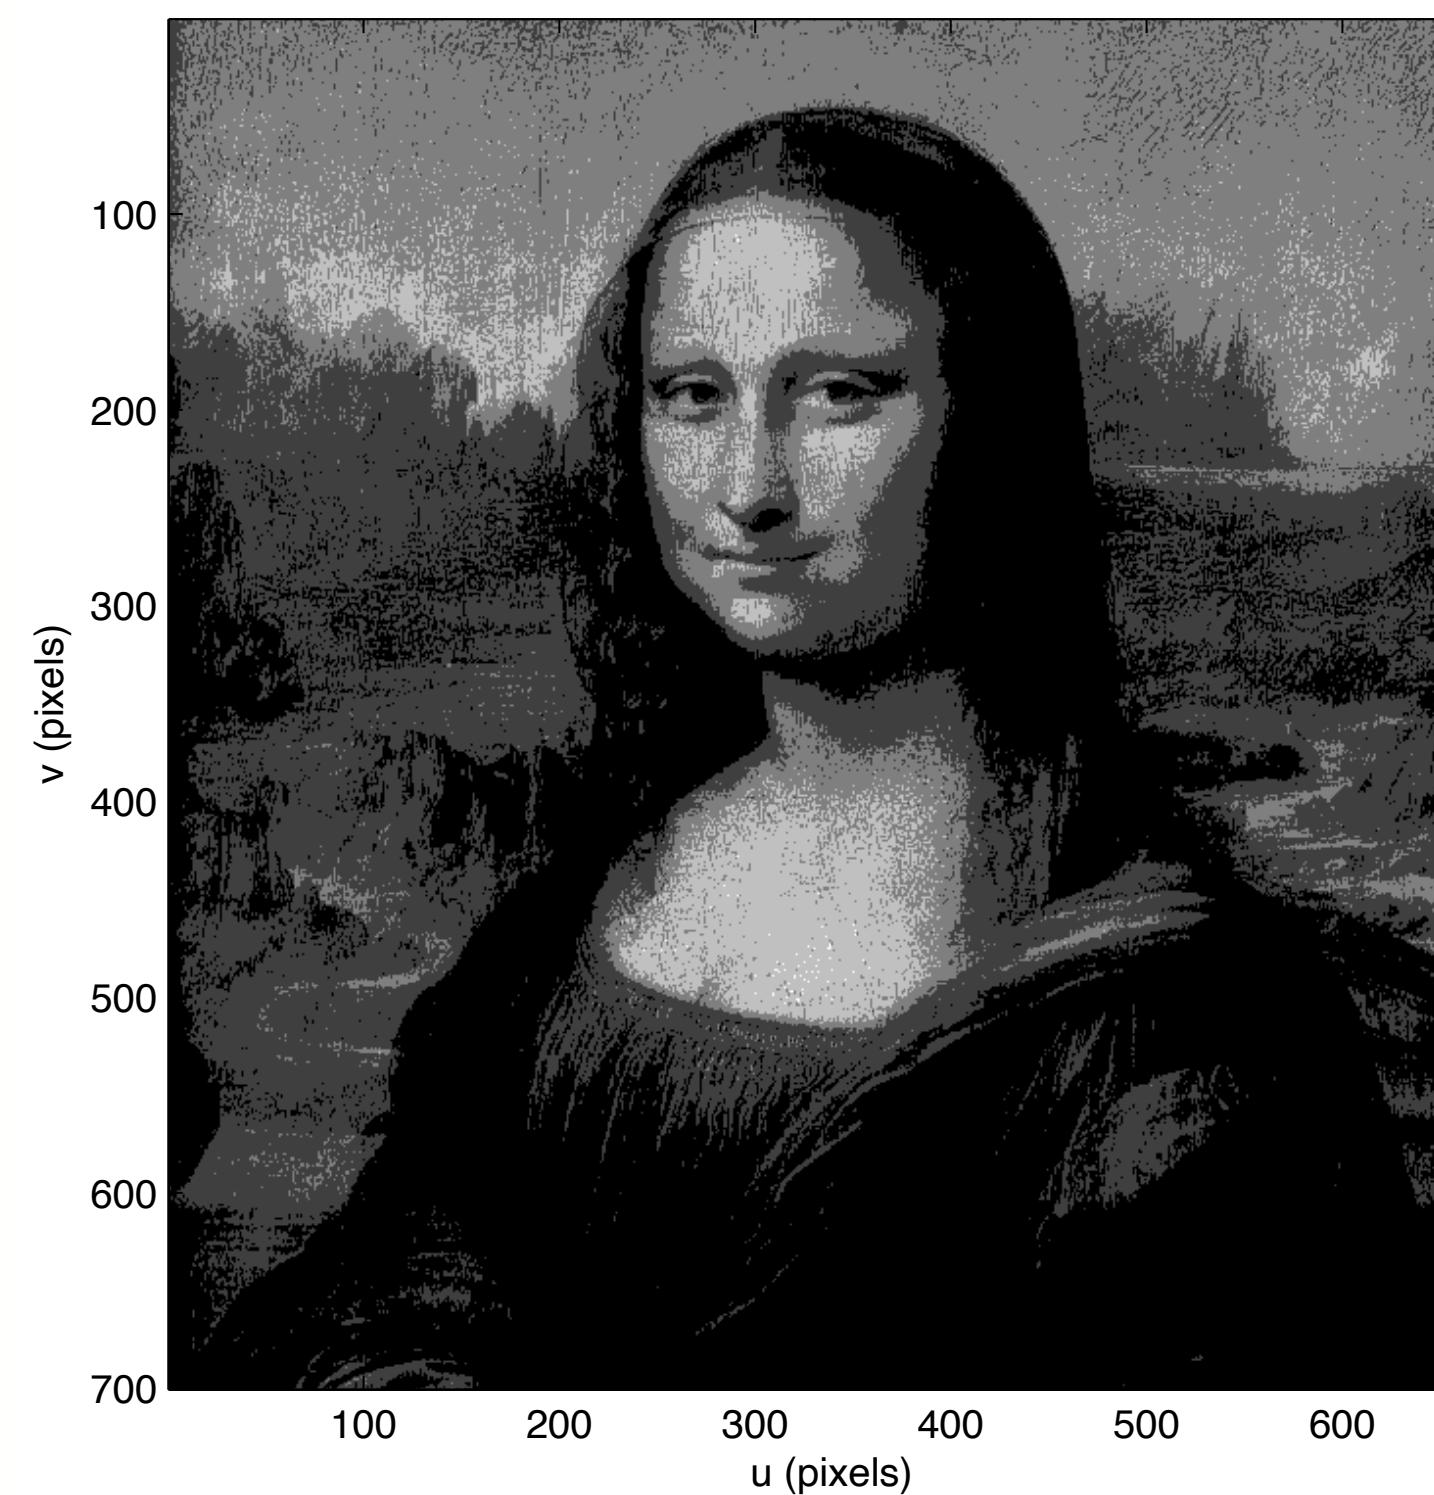
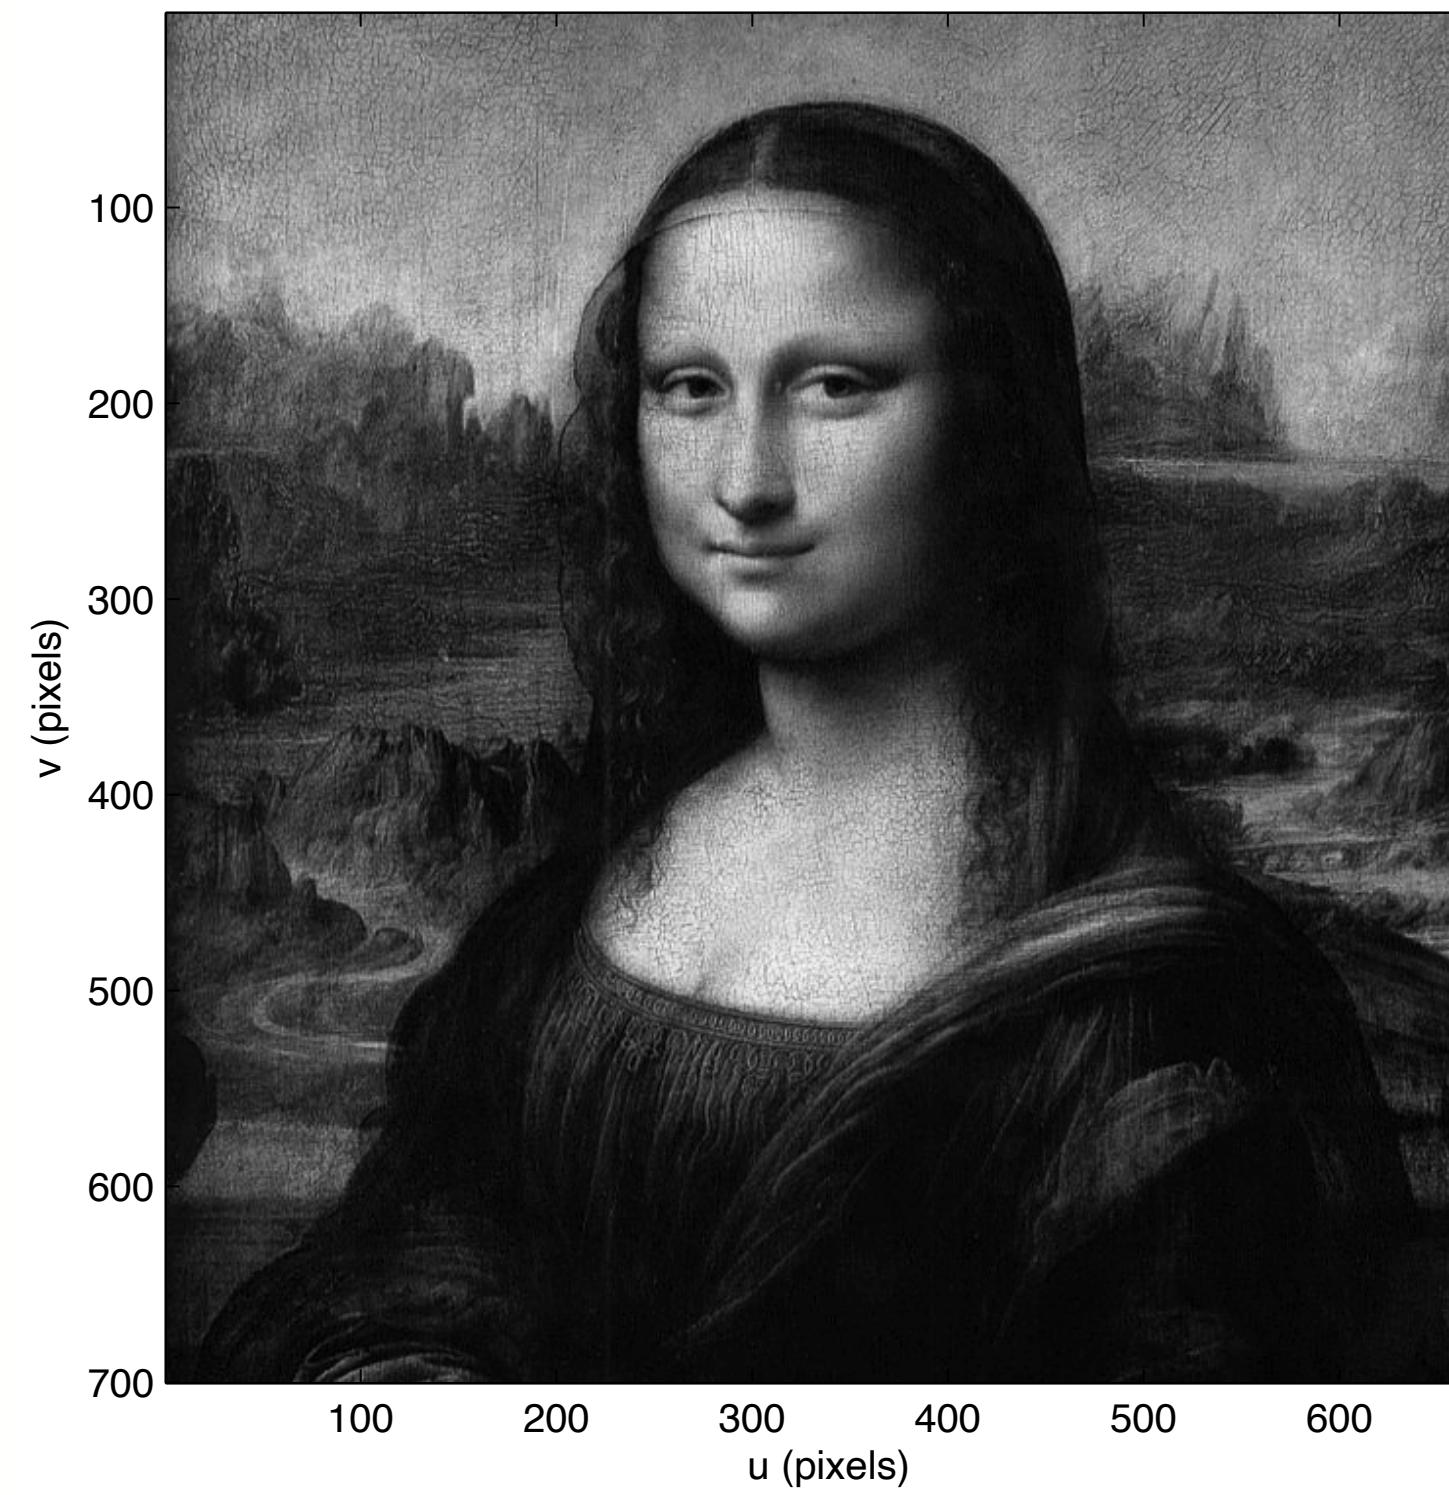
- $f(x) = 1-x$



Mona Lisa c.1503–1505
Da Vinci, L.

Posterisation

- Only N unique grey levels in the output image
→ $f(x) = N \text{ floor}(x/N)$

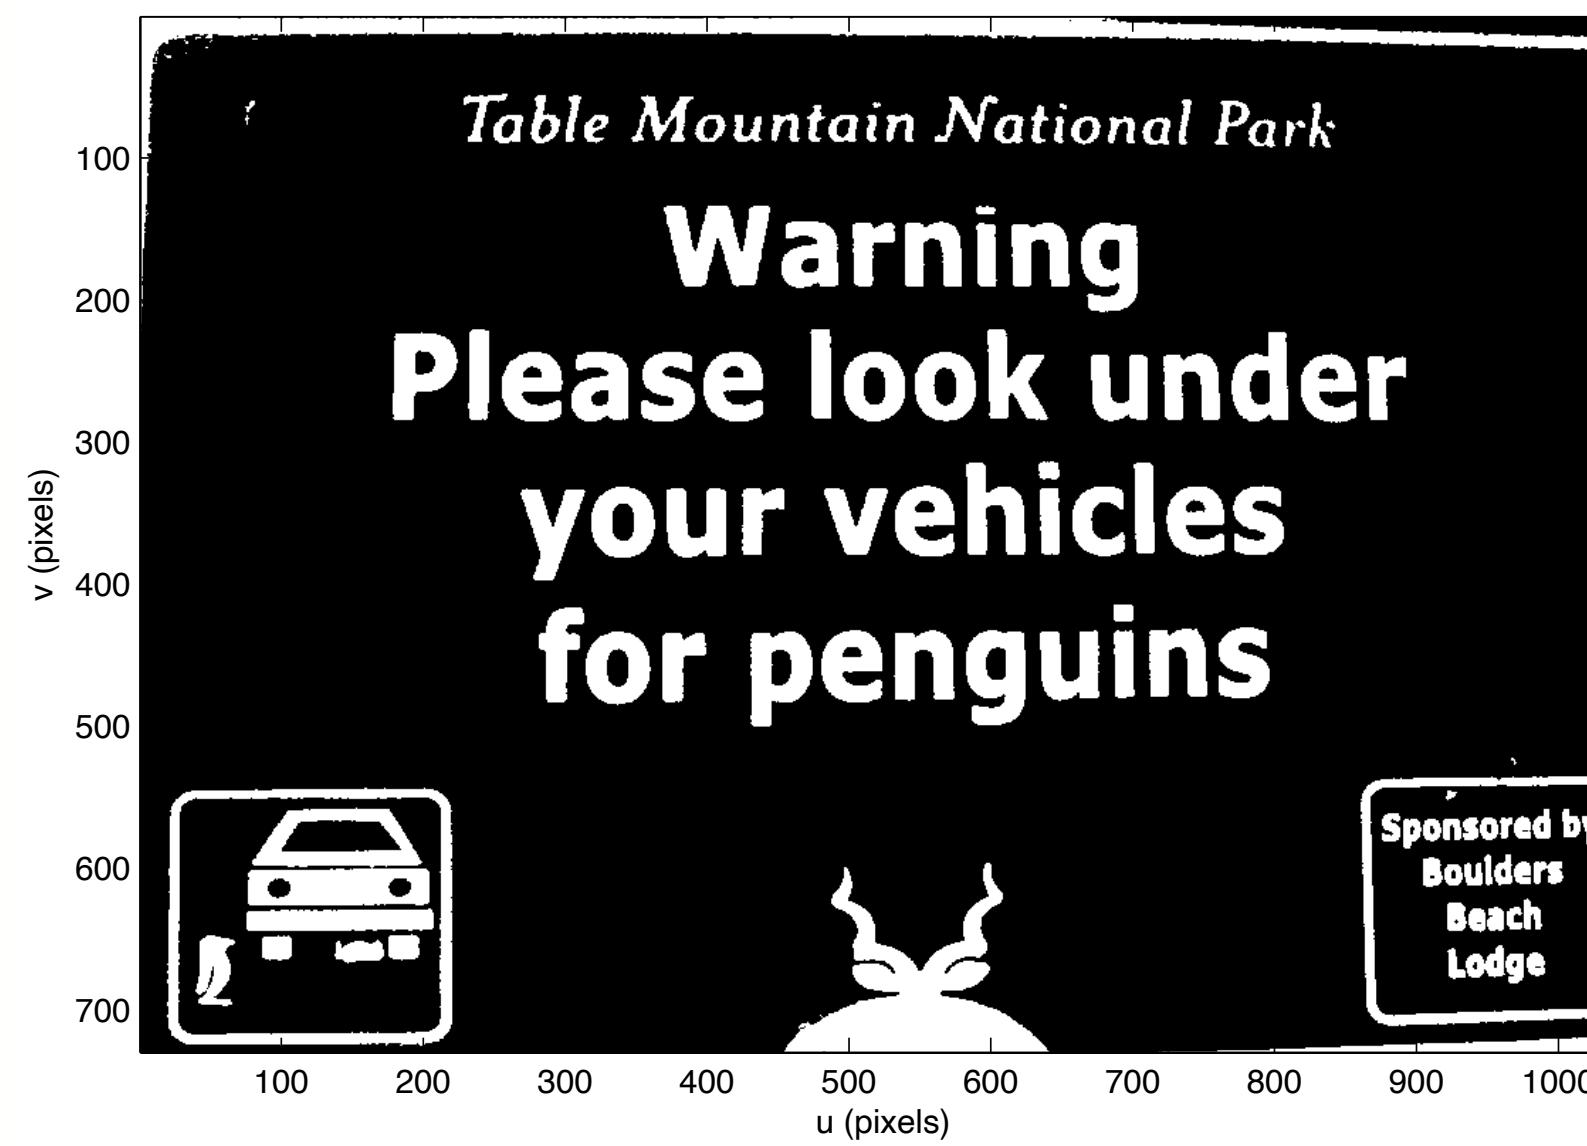
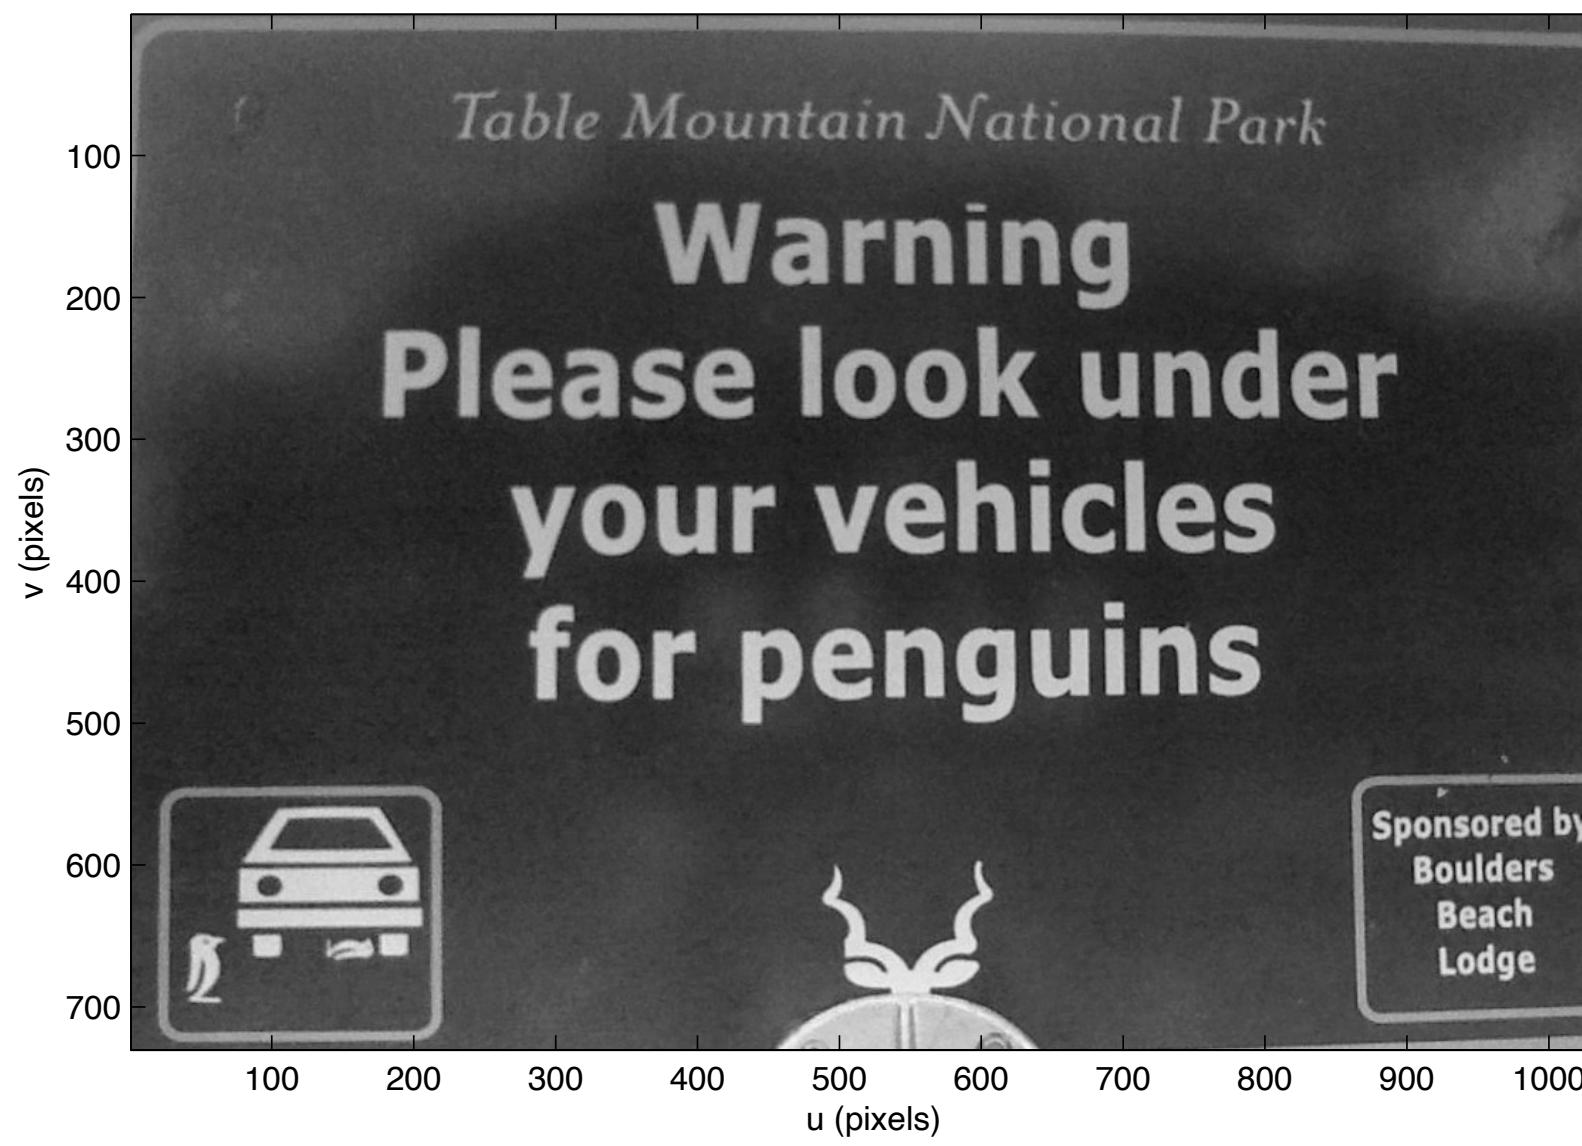
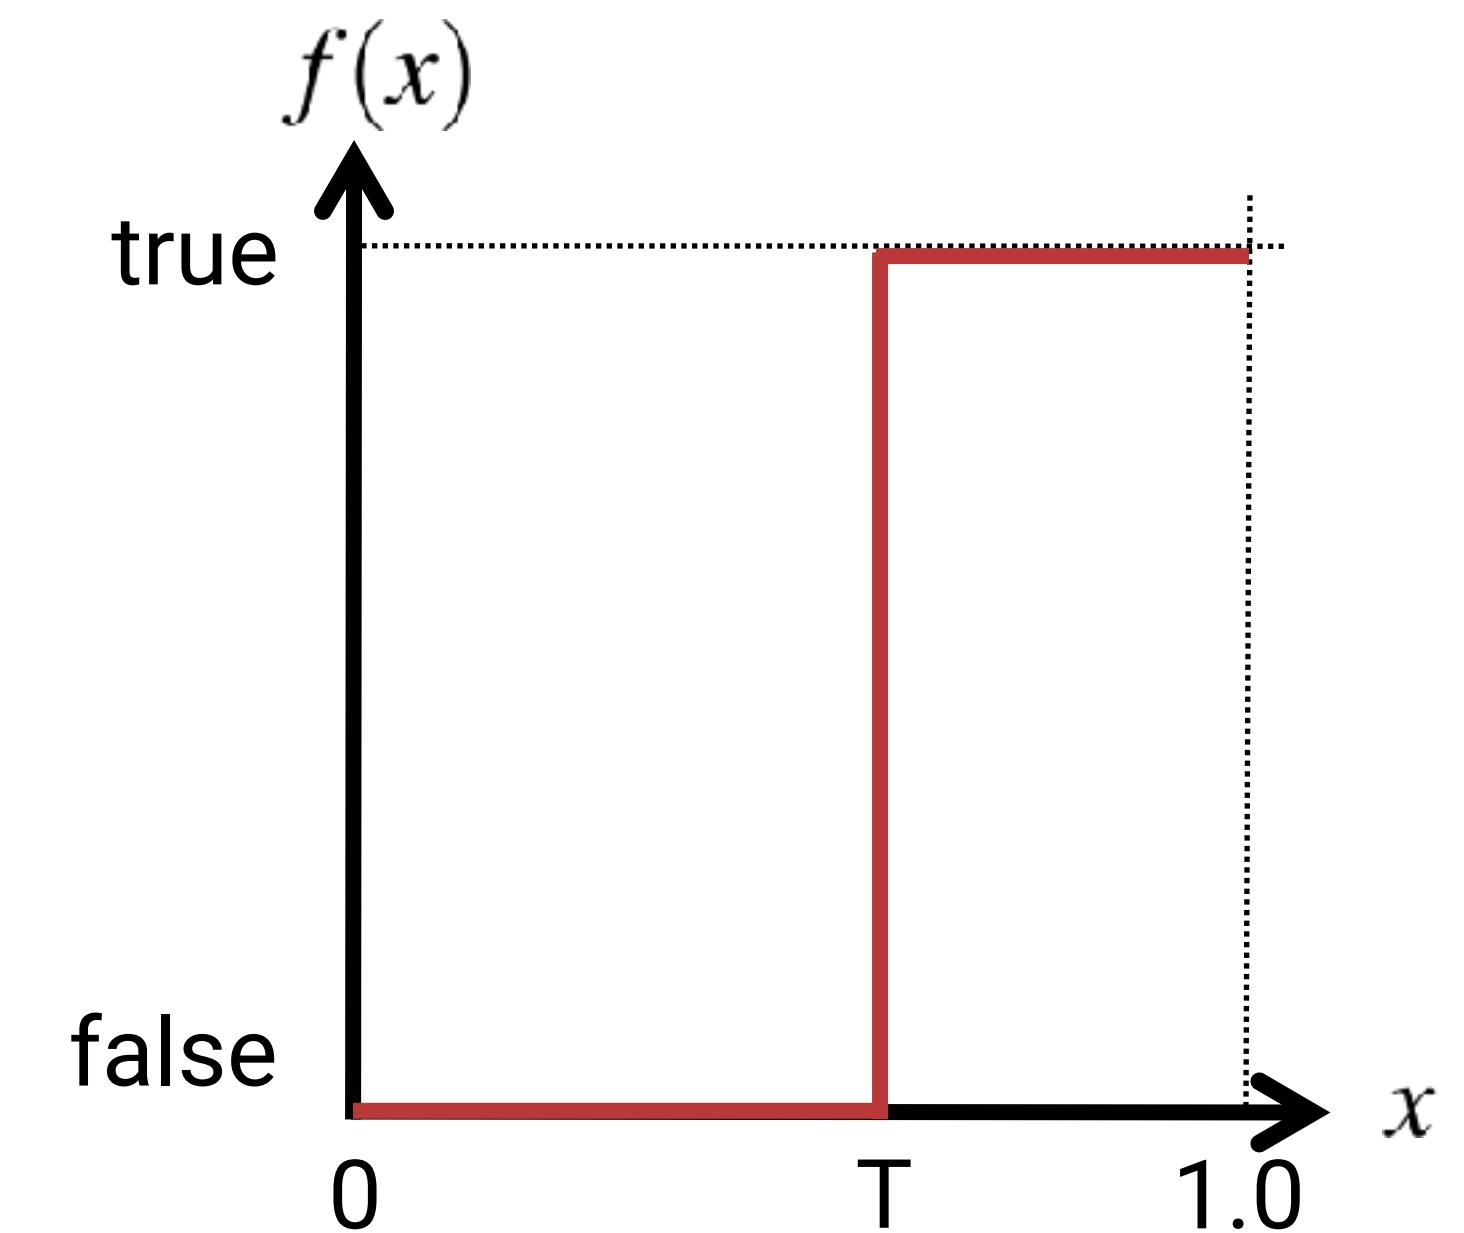


Mona Lisa c.1503–1505
Da Vinci, L.

N=4

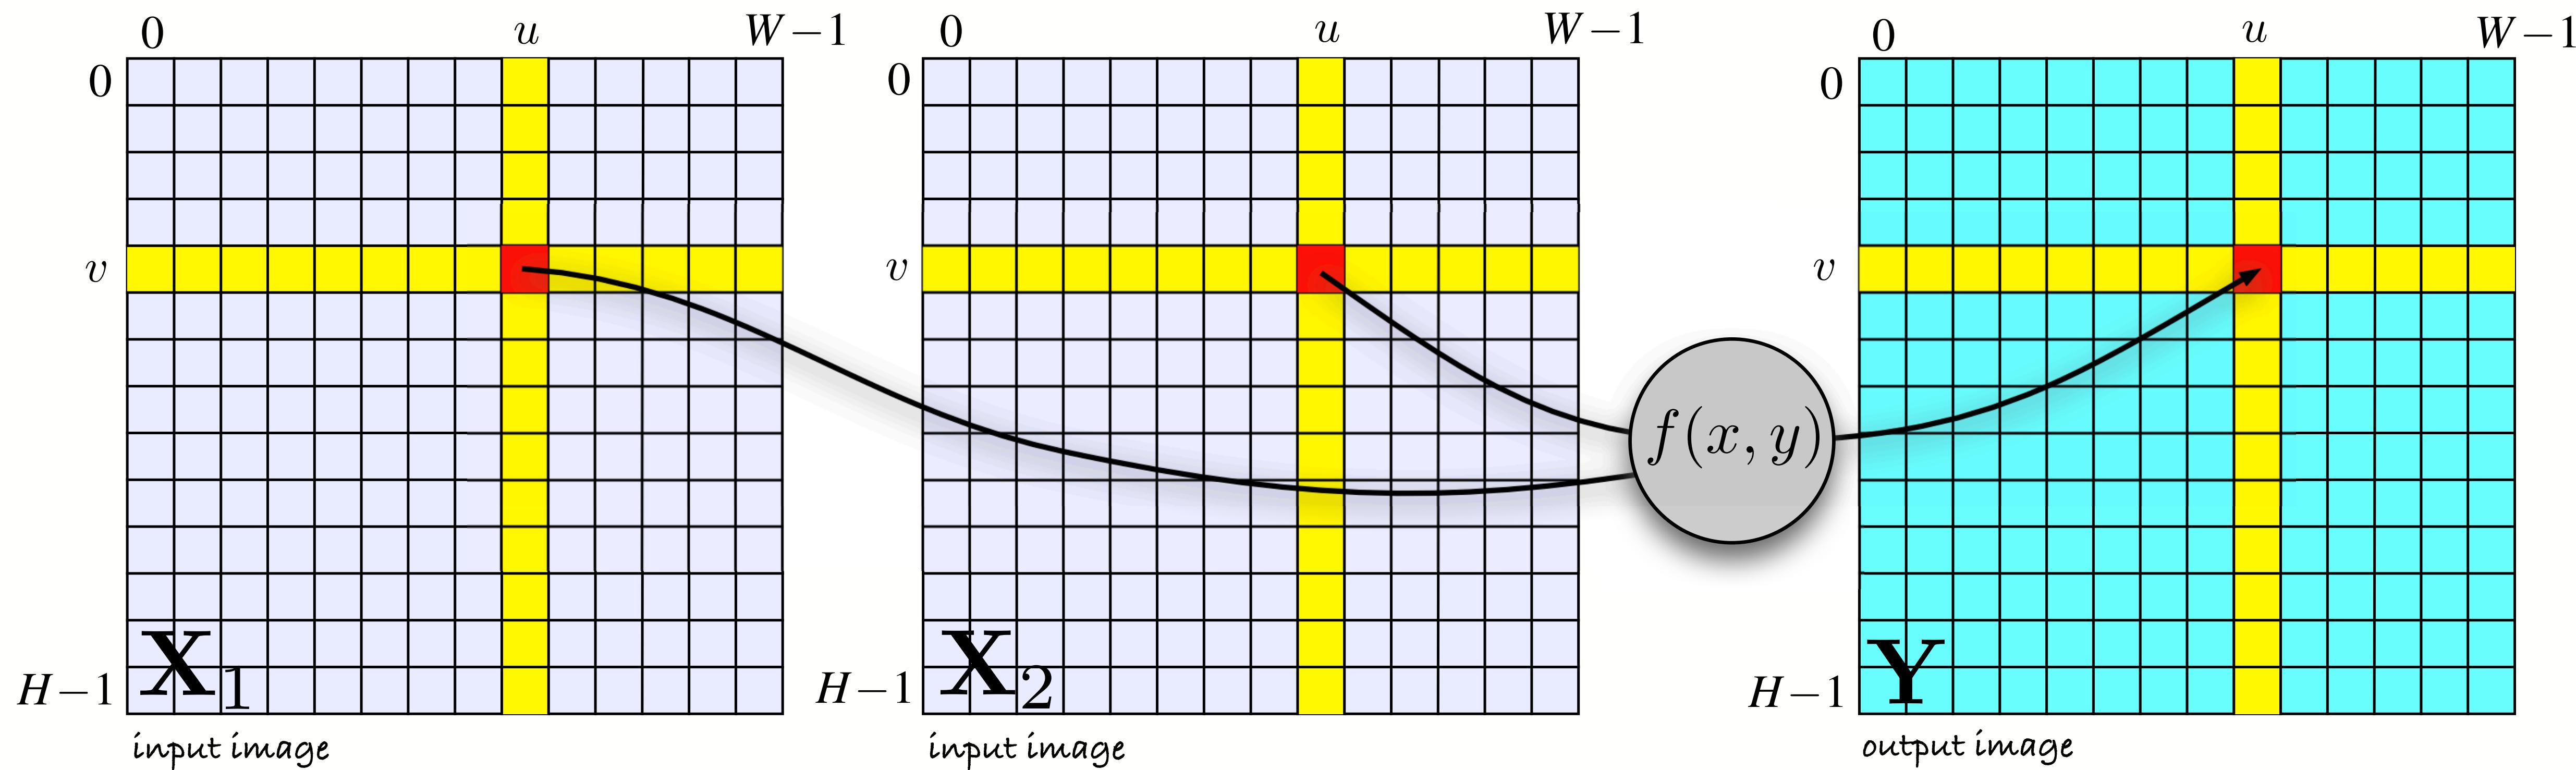
Thresholding

- The resulting image has only two values:
 - false (black) if $x < T$
 - true (white) if $x \geq T$



$T = 0.4$

Diadic operations



- Each output pixel is a function of two corresponding input pixels
- The function is **the same** for all pixels

Diadic image processing

- What is $f(x,y)$?

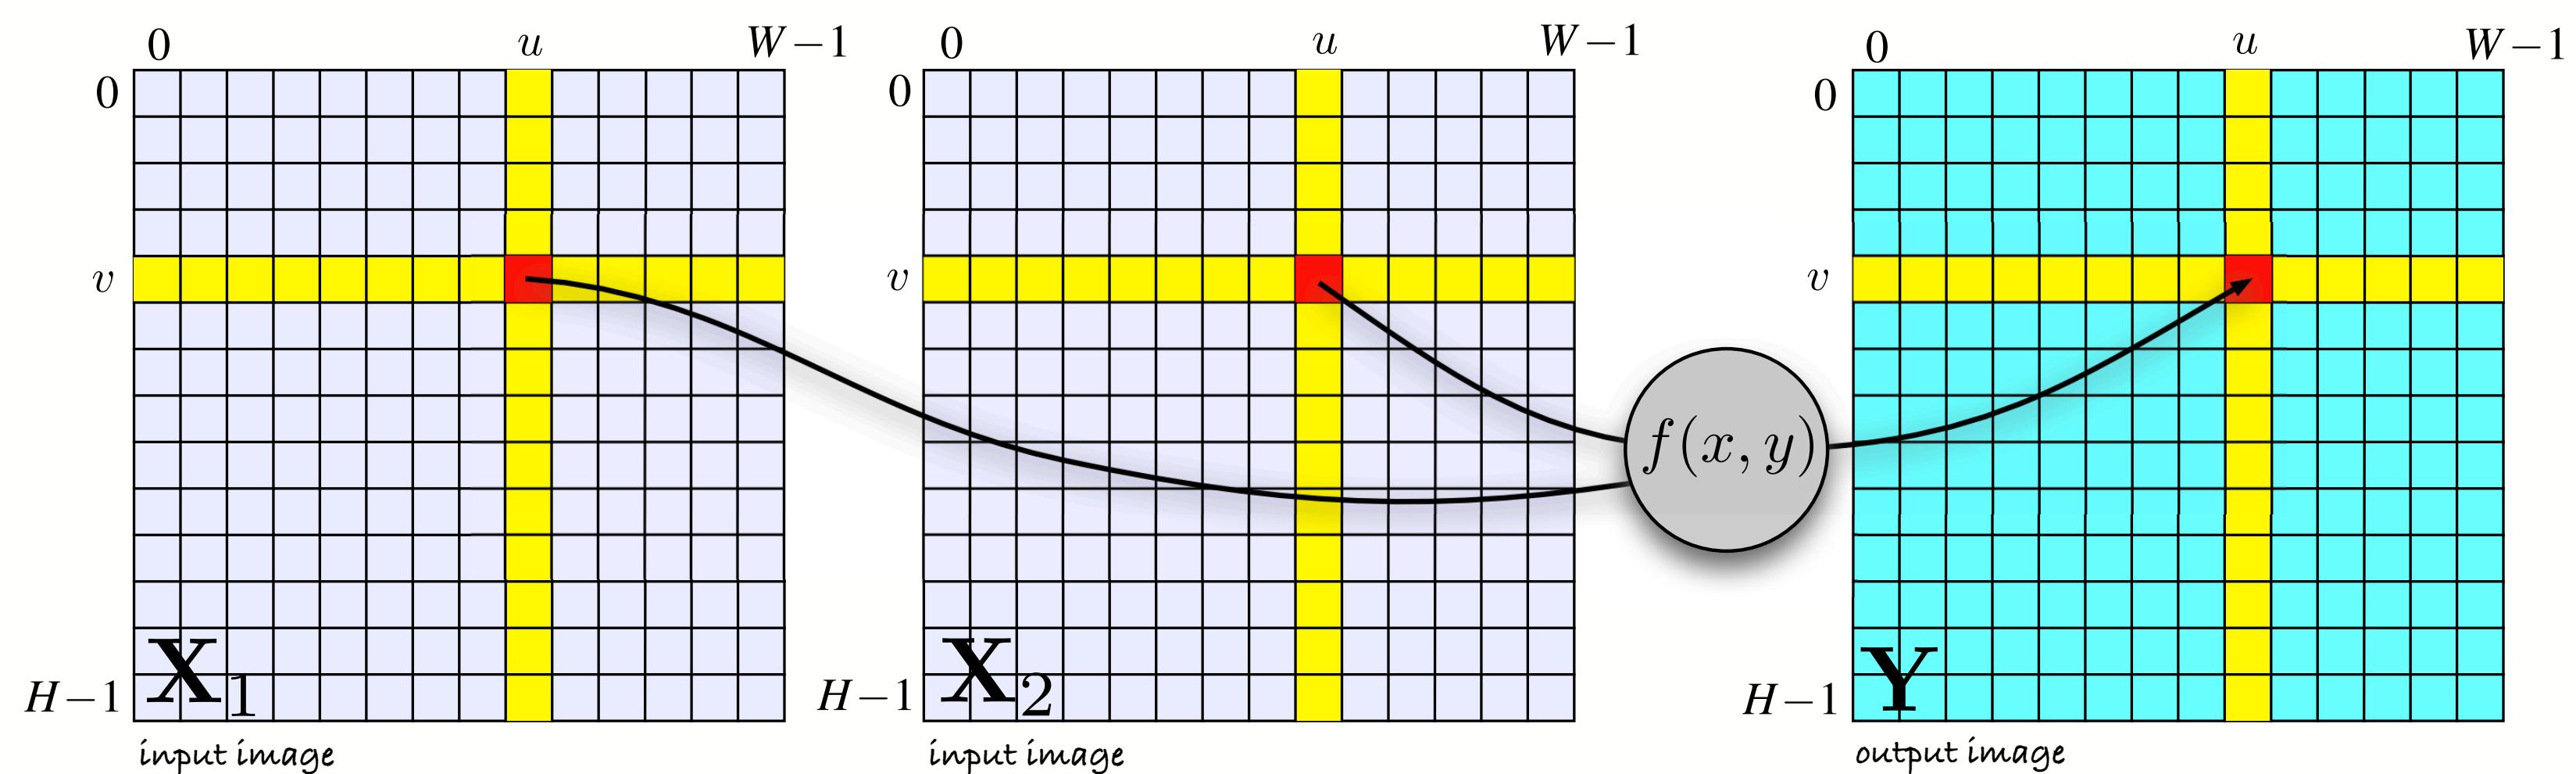
→ arithmetic

- $x+y$
- $x-y$
- x^*y
- etc.

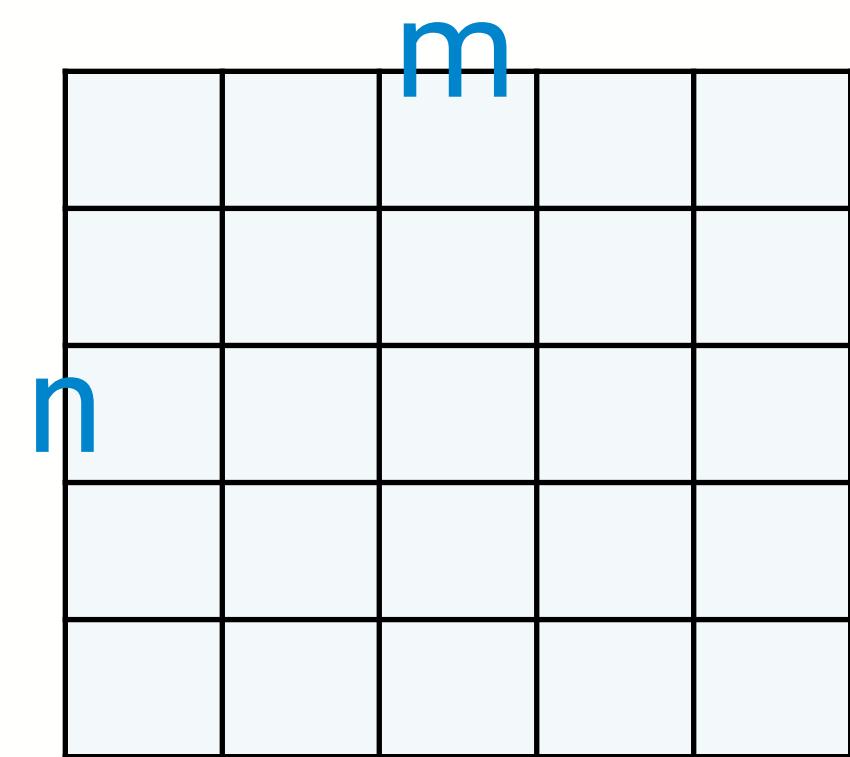
→ $\sqrt{x^2 + y^2}$

→ relational

- $x > y$
- $x < y$
- etc.

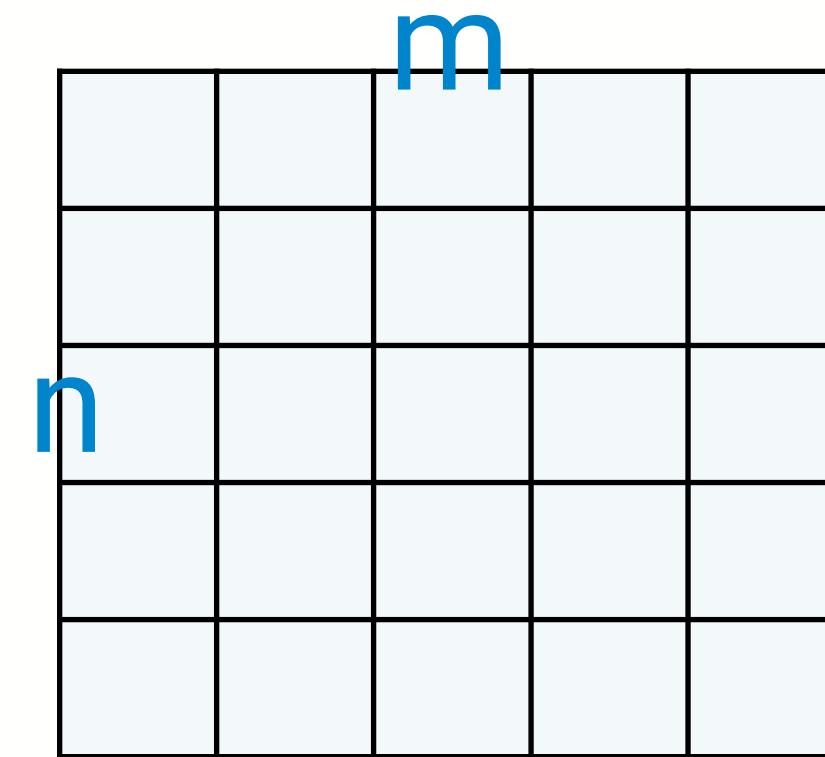


NumPy array operations

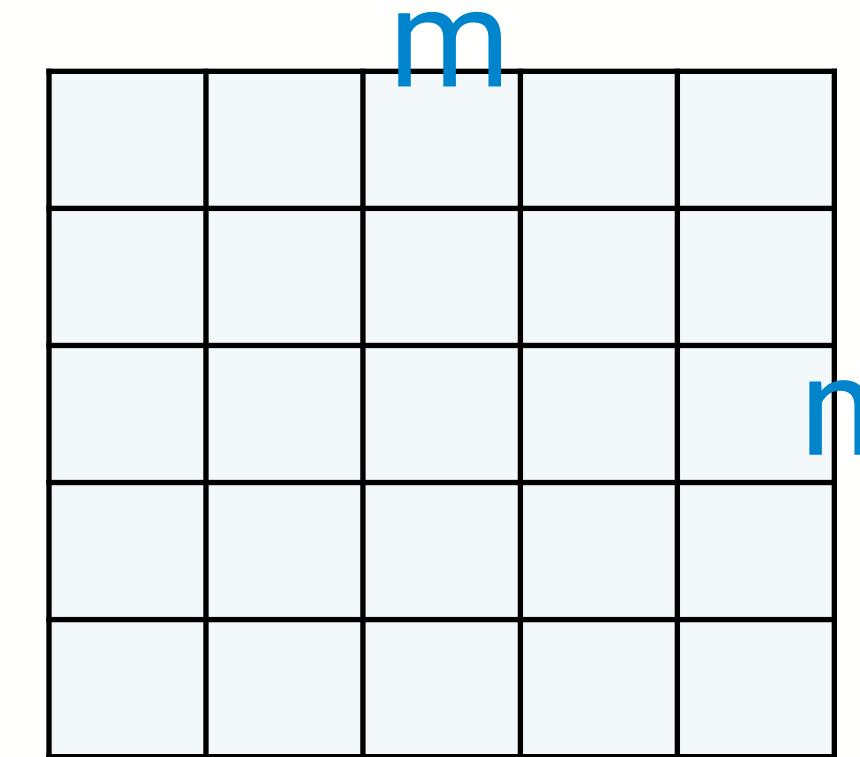


A vertical column of operators:

- +
-
- *
- /
- >
- \geq
- <
- \leq
- \equiv
- \neq



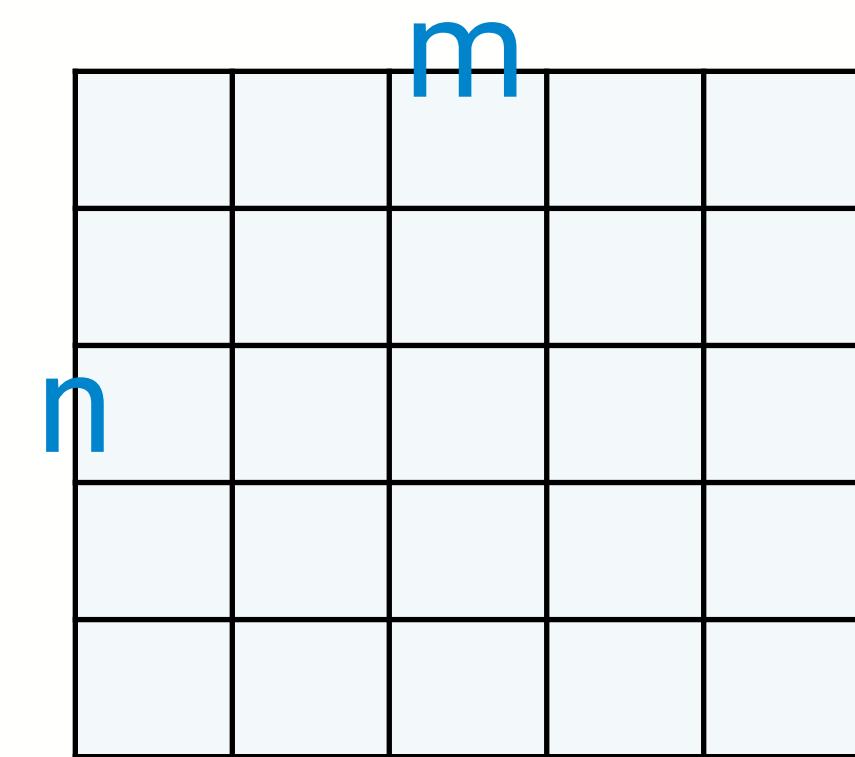
=



Both operands must have the same size

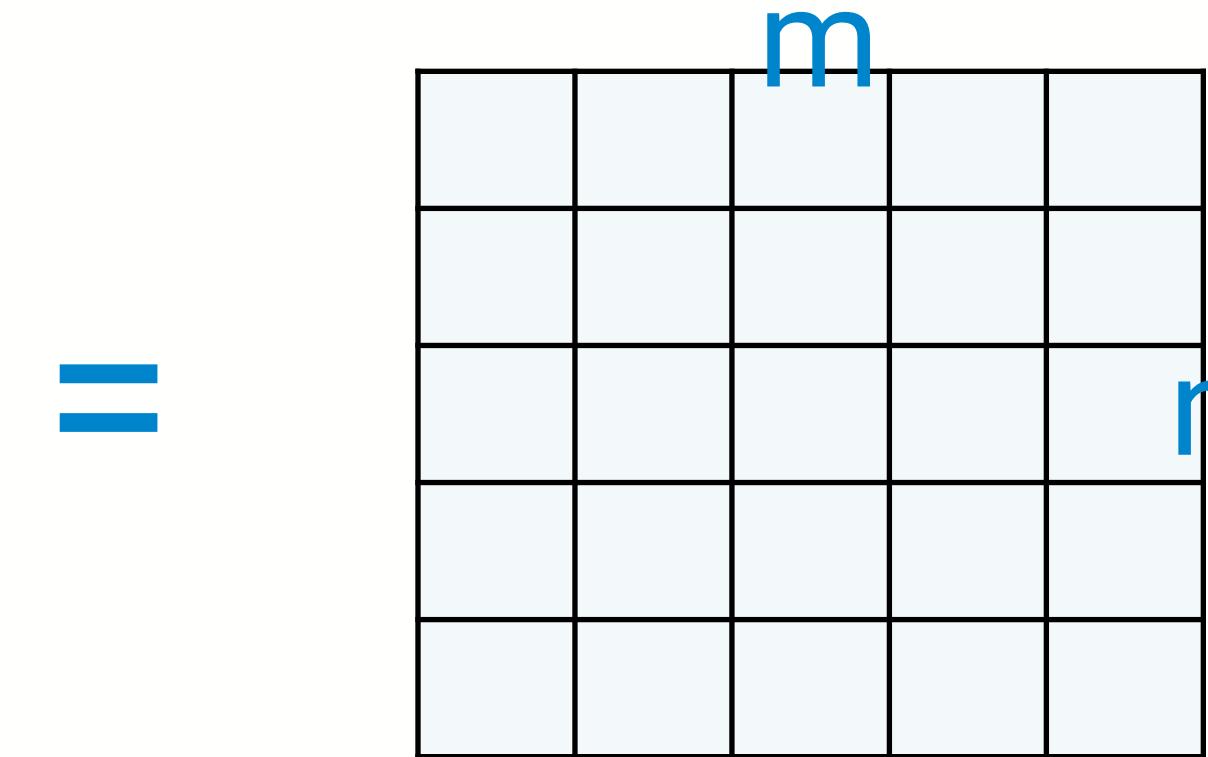
The output is the same size as well

NumPy array operations with broadcasting



+ >
- >=
* <
/ <=
==
!=

scalar



scalar

+ >
- >=
* <
/ <=
==
!=

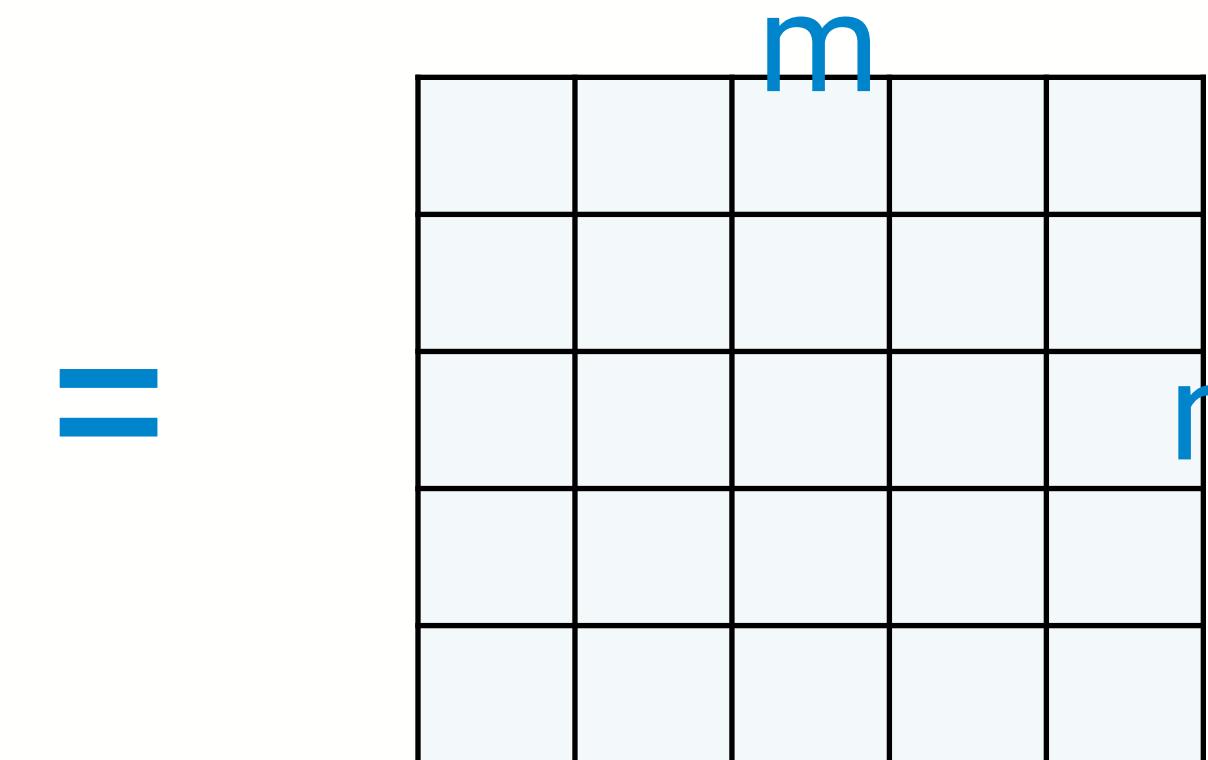
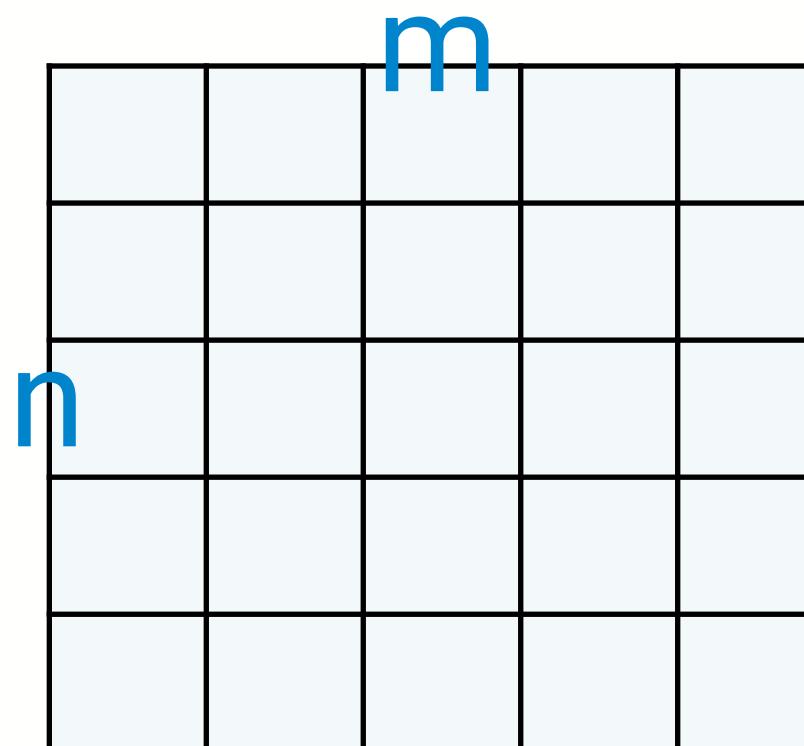
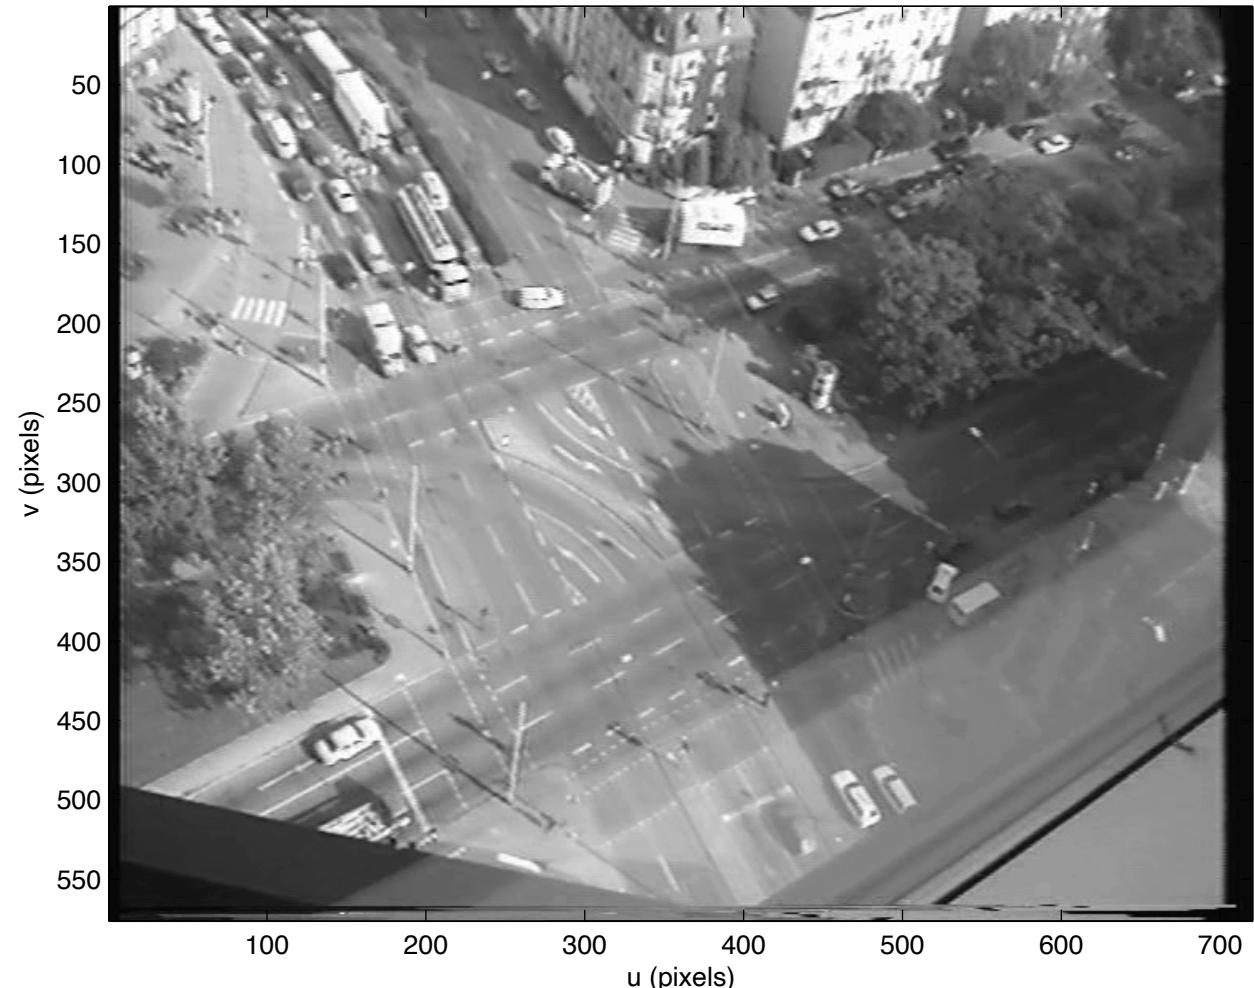
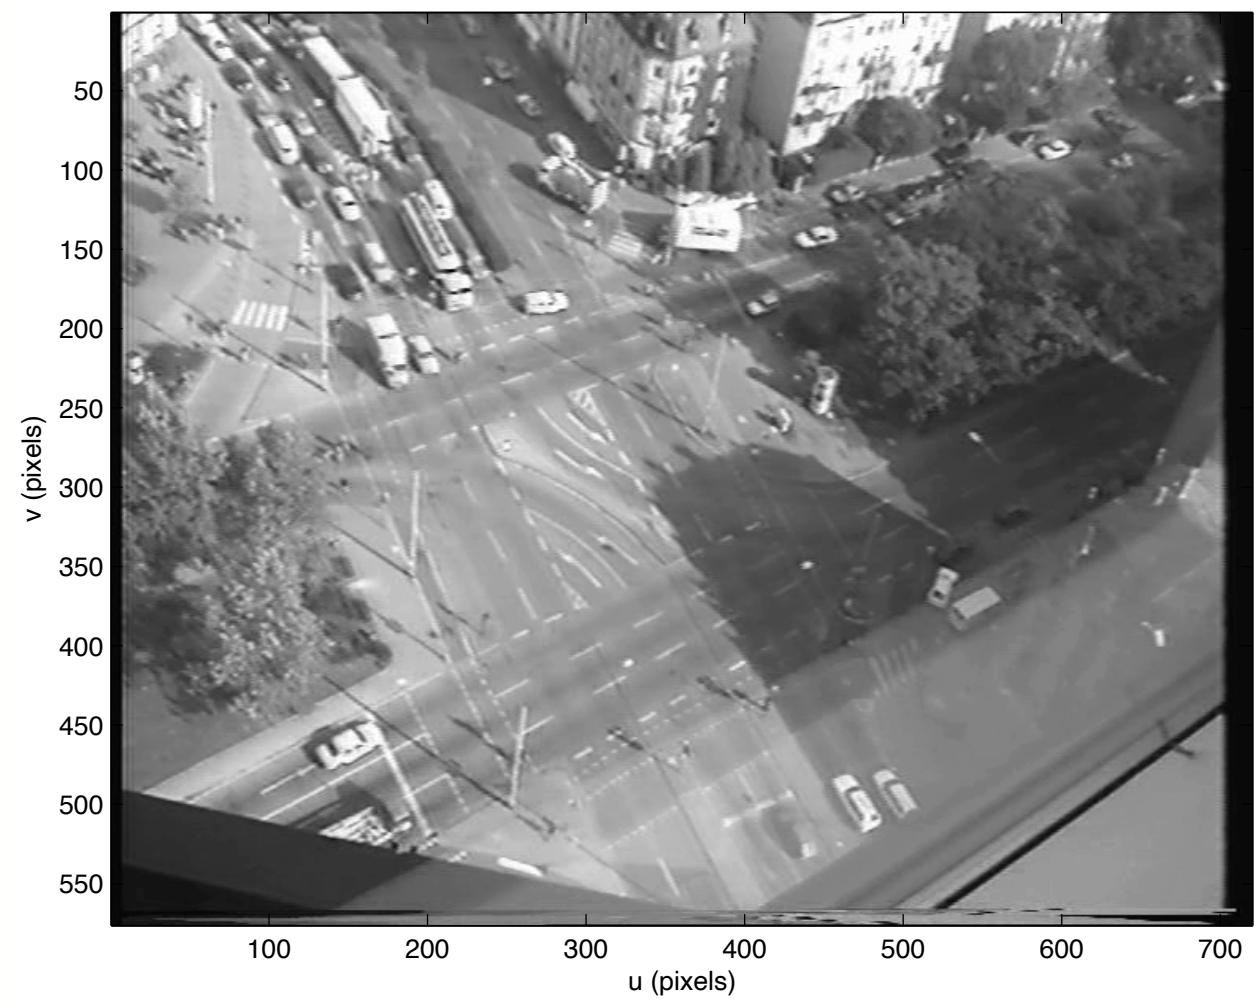


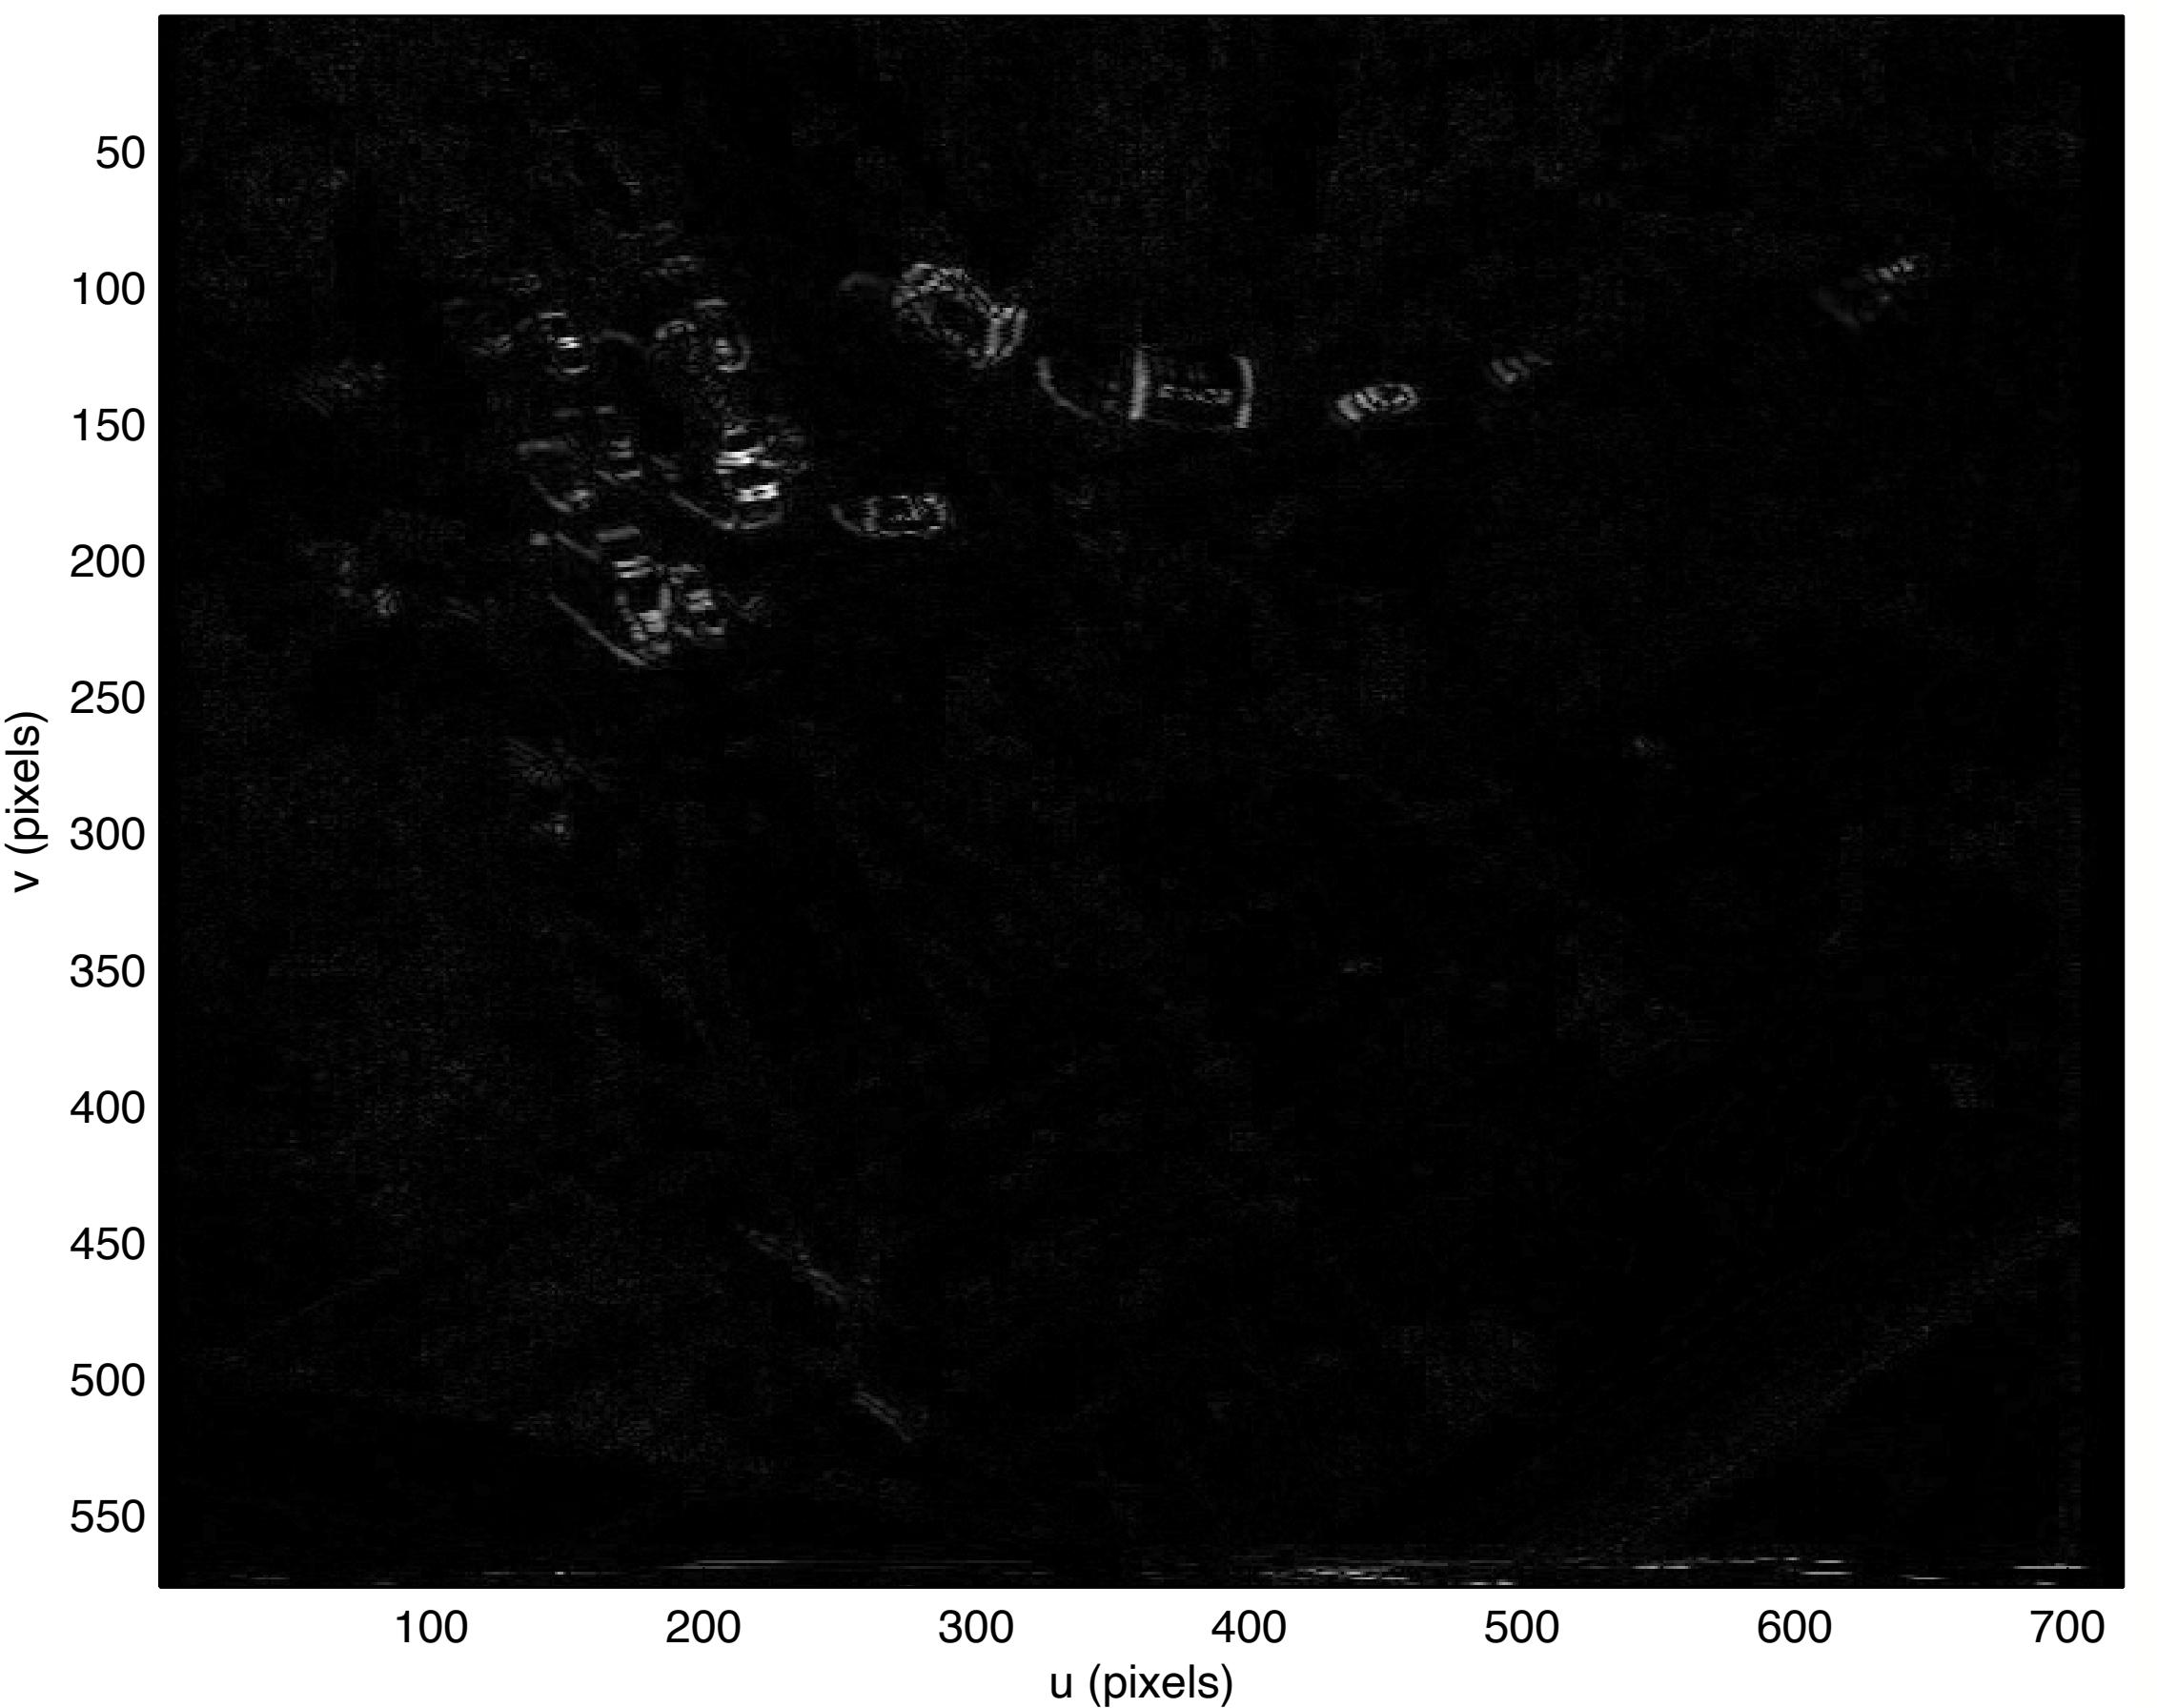
Image subtraction



frame 1



frame 4



- Pixel by pixel difference highlights image change
- Be careful when subtracting uint8 types



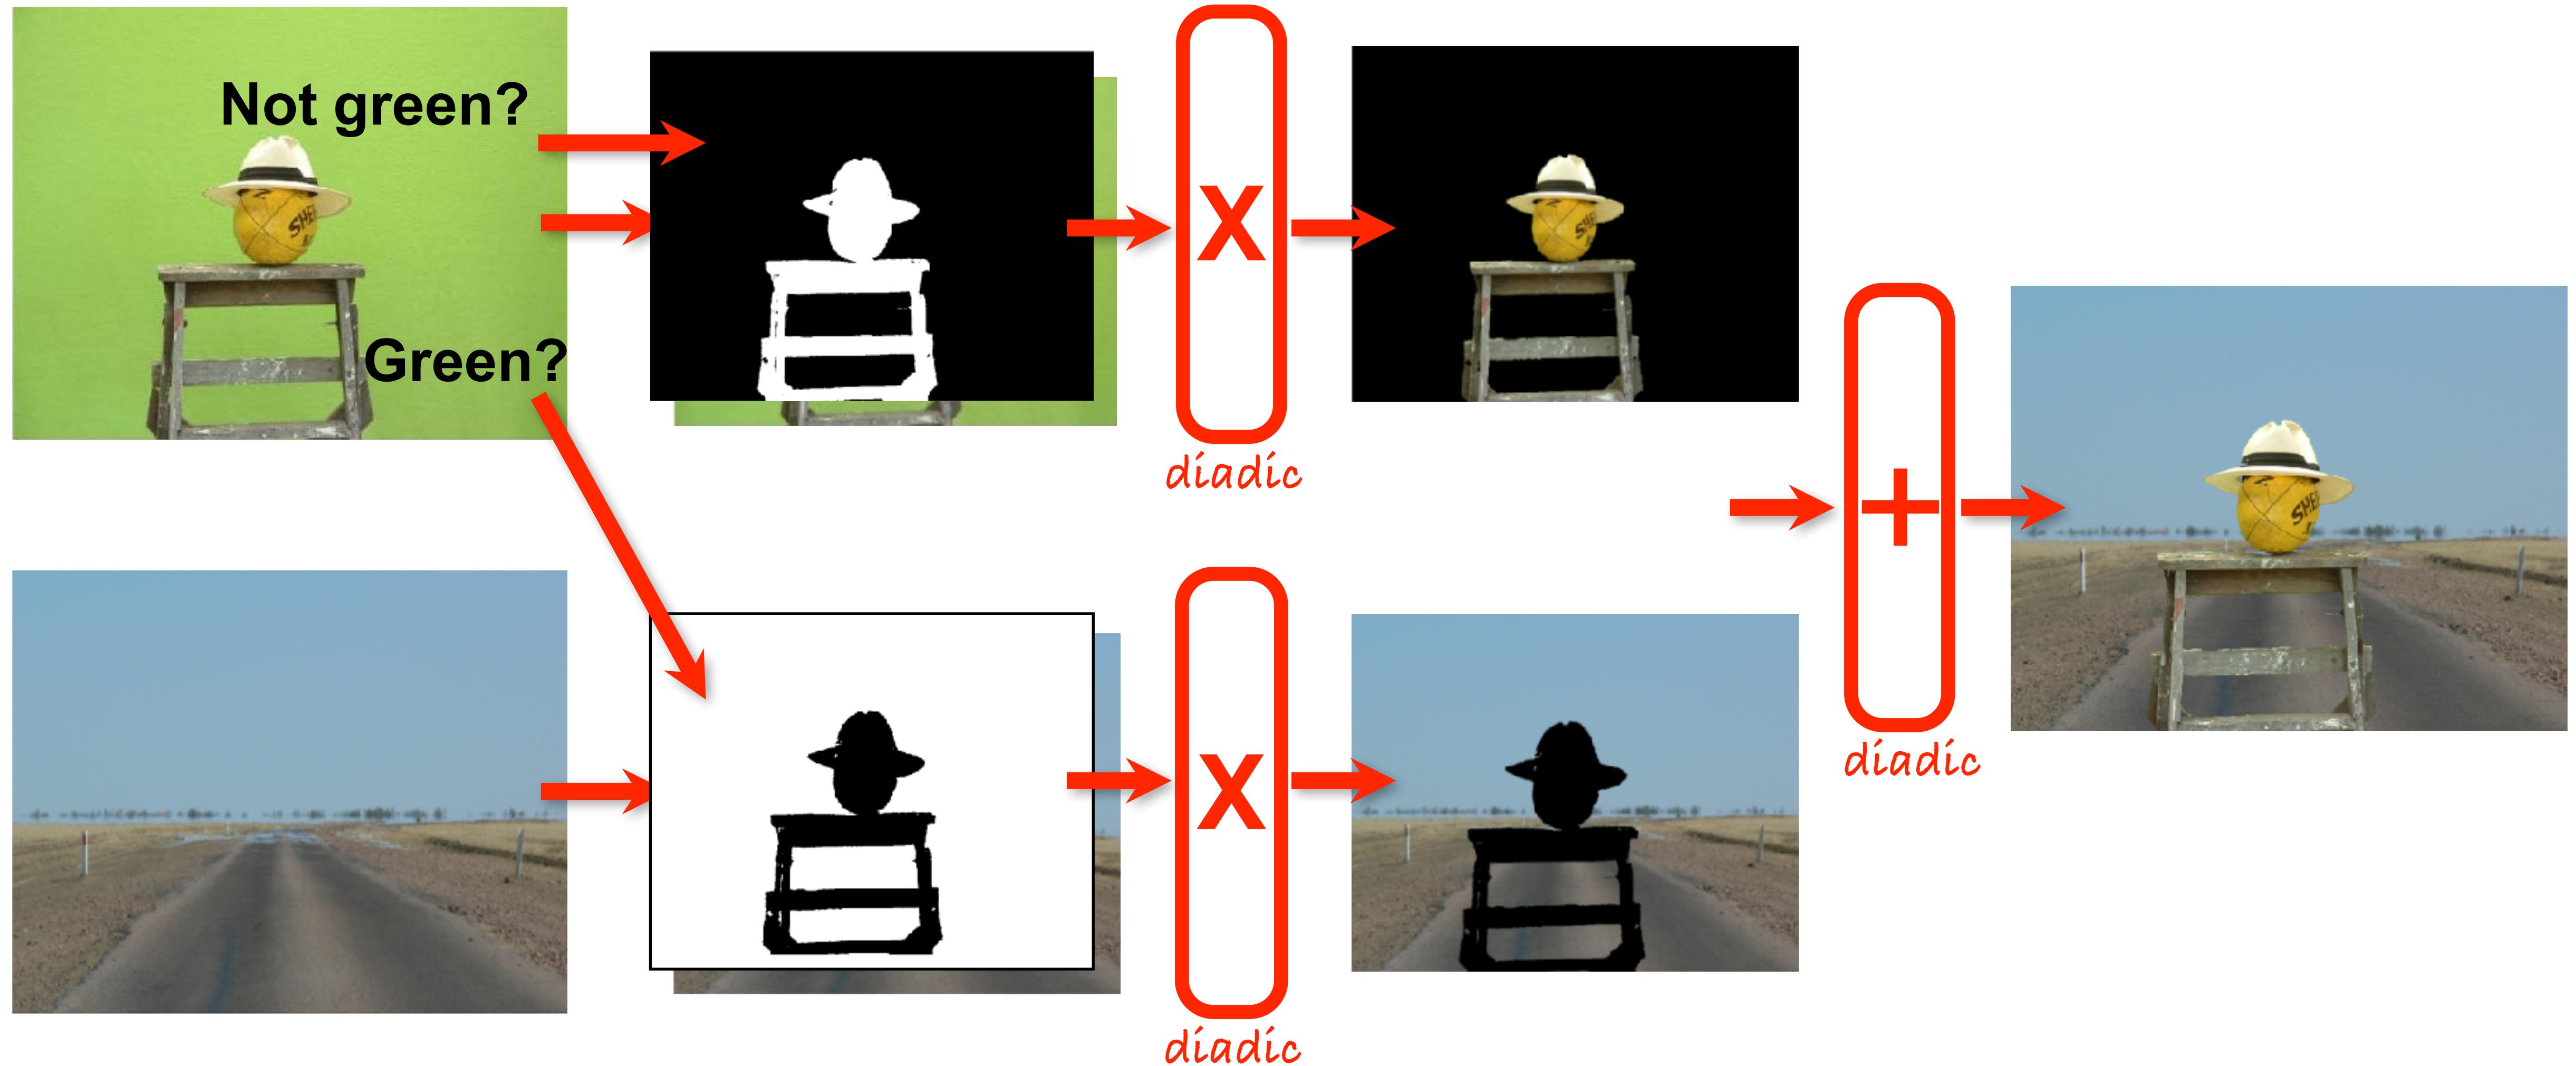
Green screen effect



- Merge two images to create a composite image
- Green pixels are substituted for background pixels



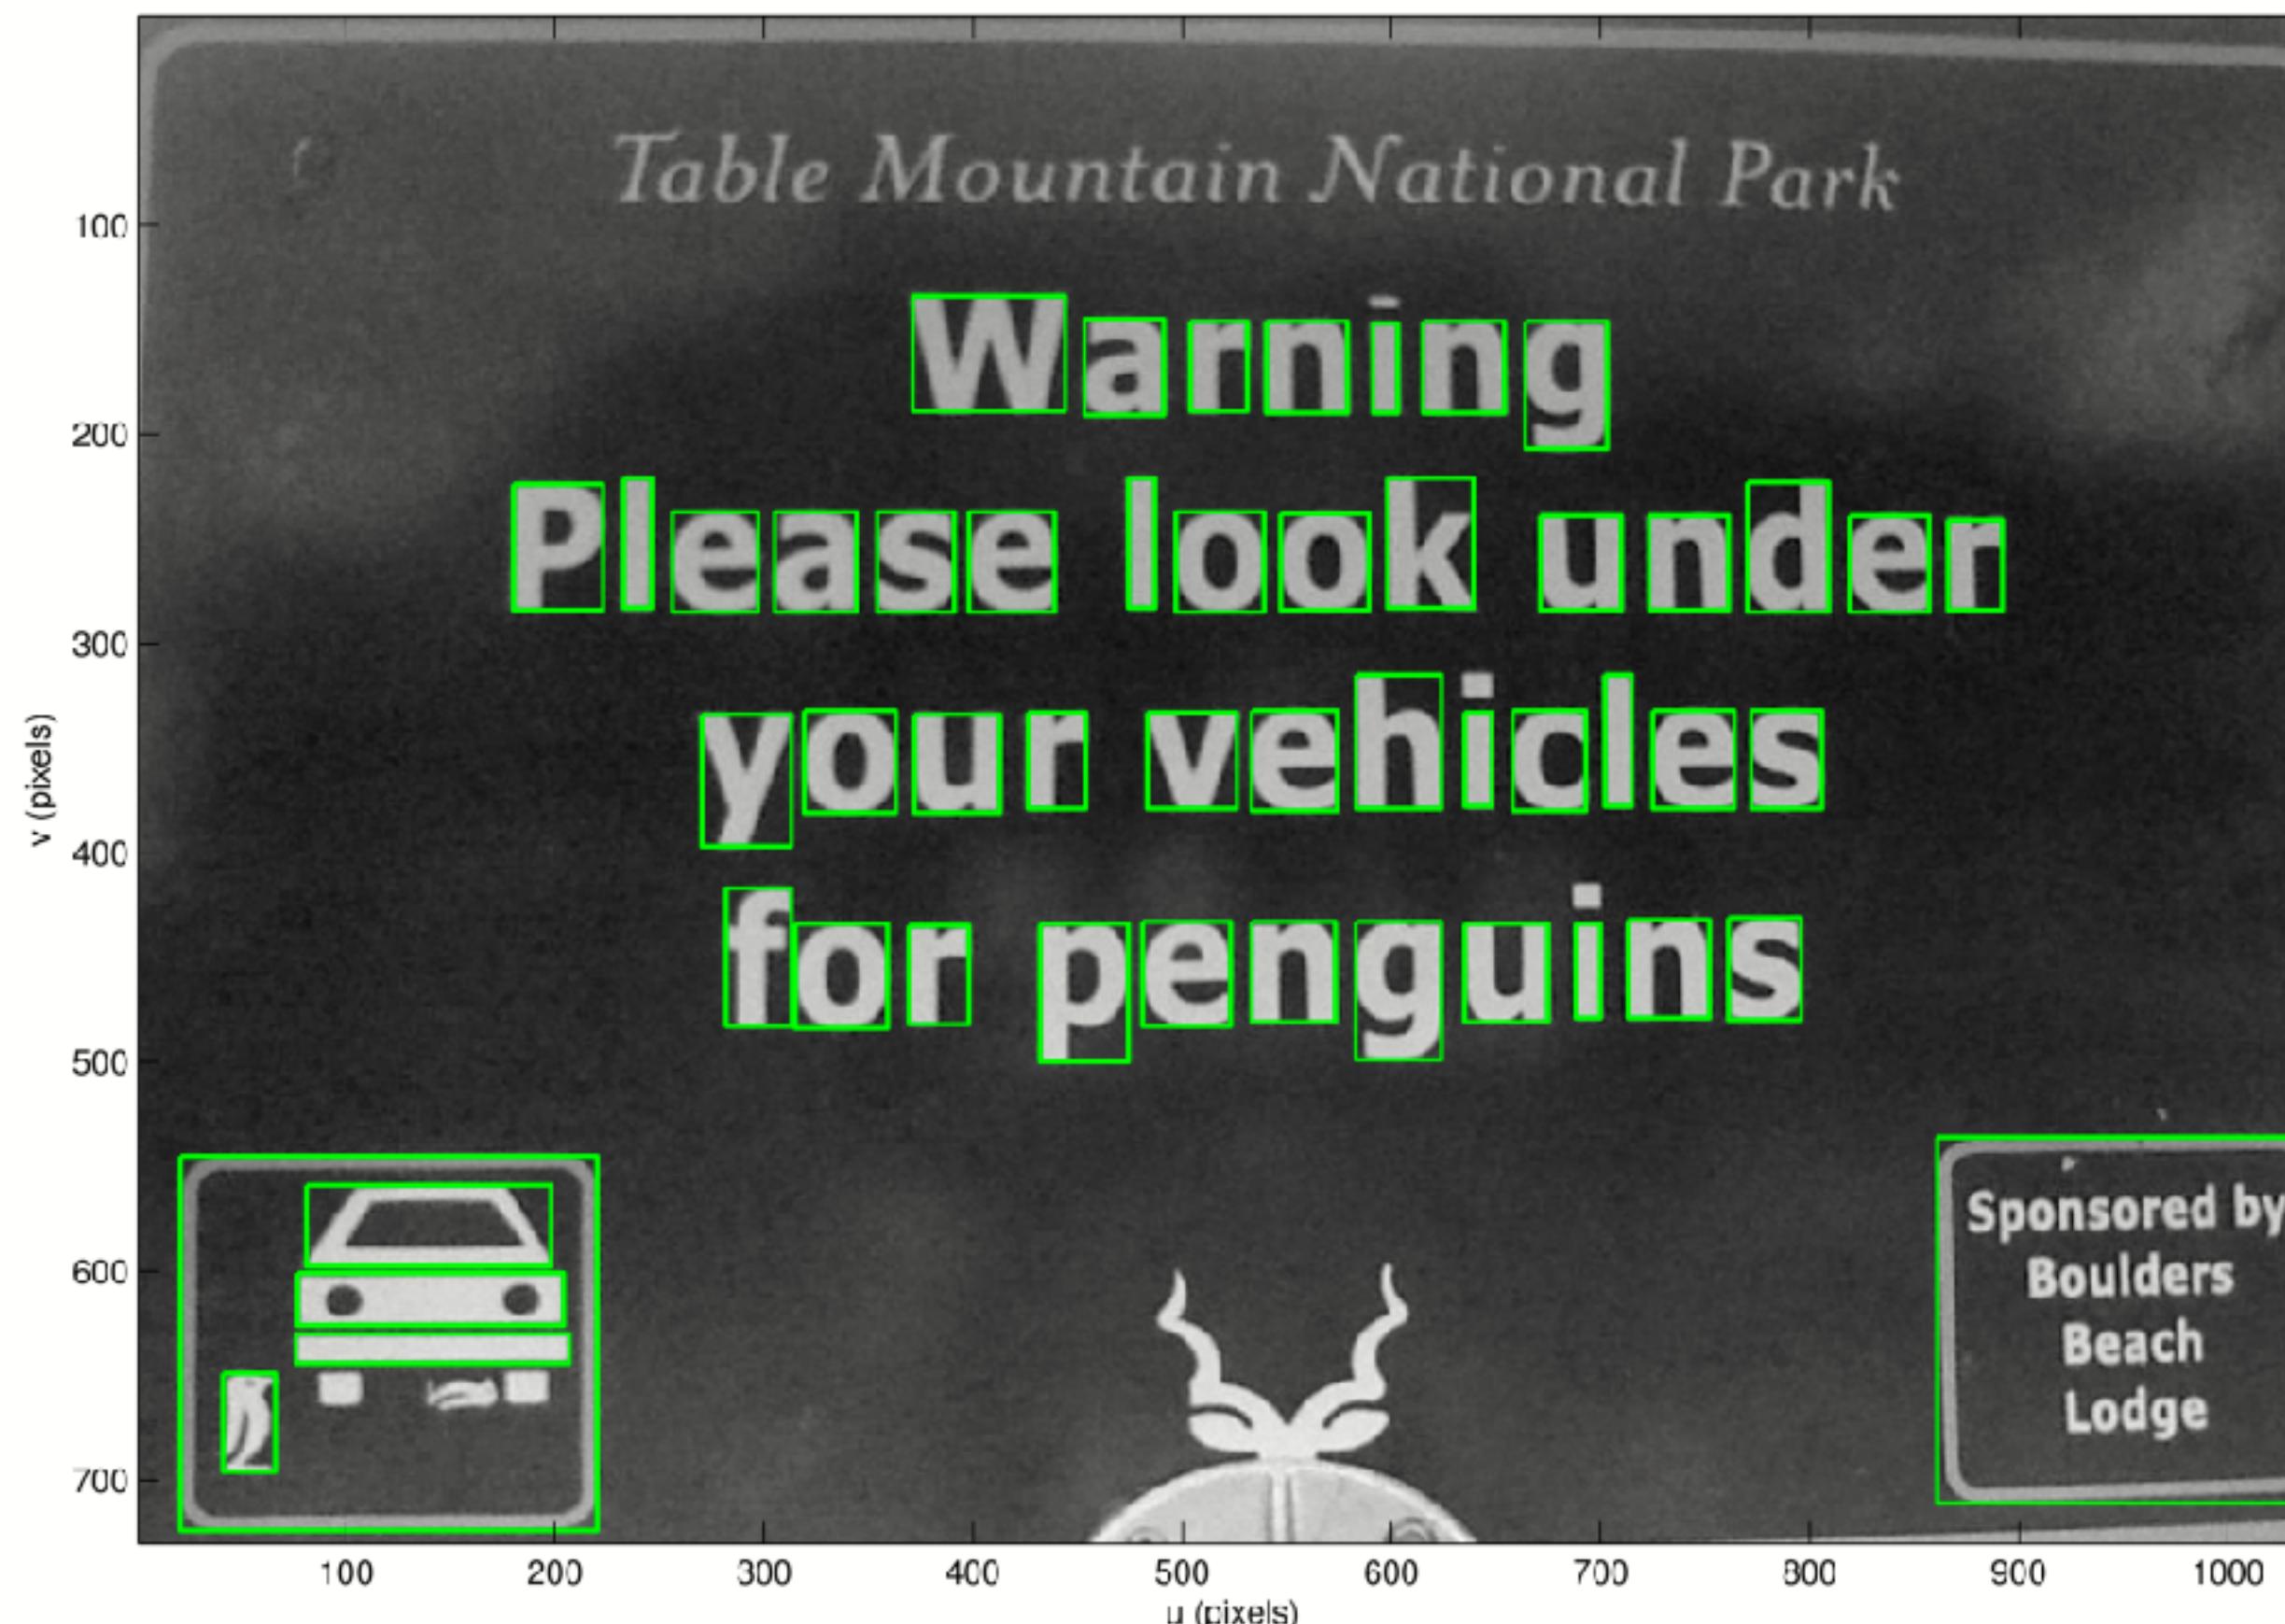
Green screen effect



Section 2

Binary blobs

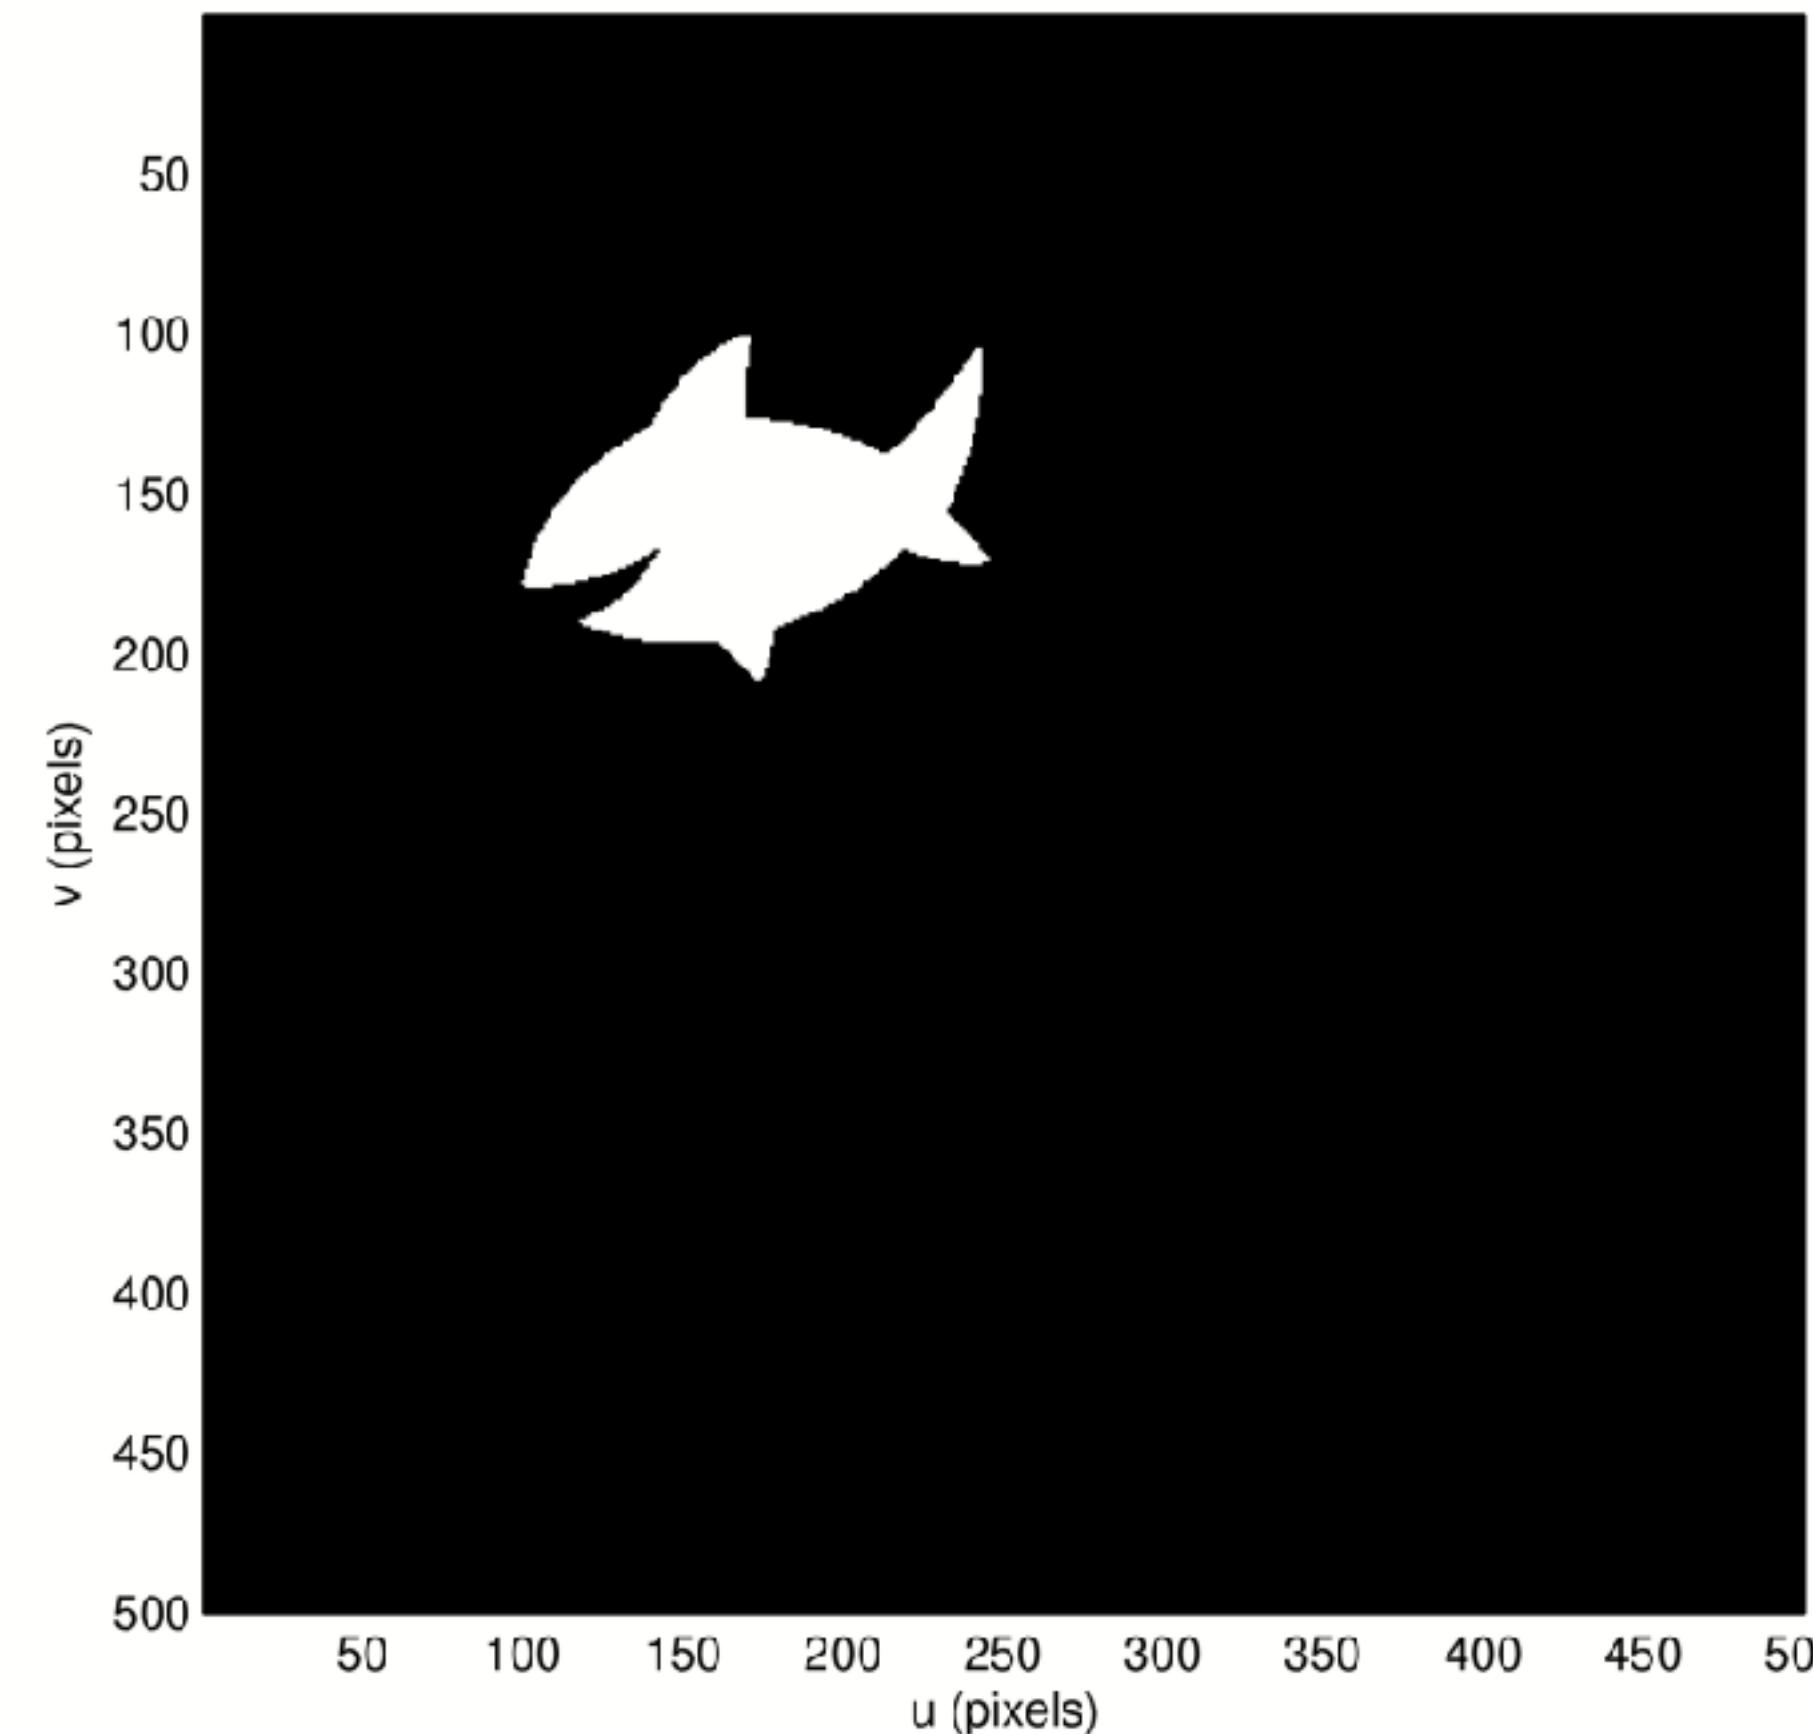
Region features



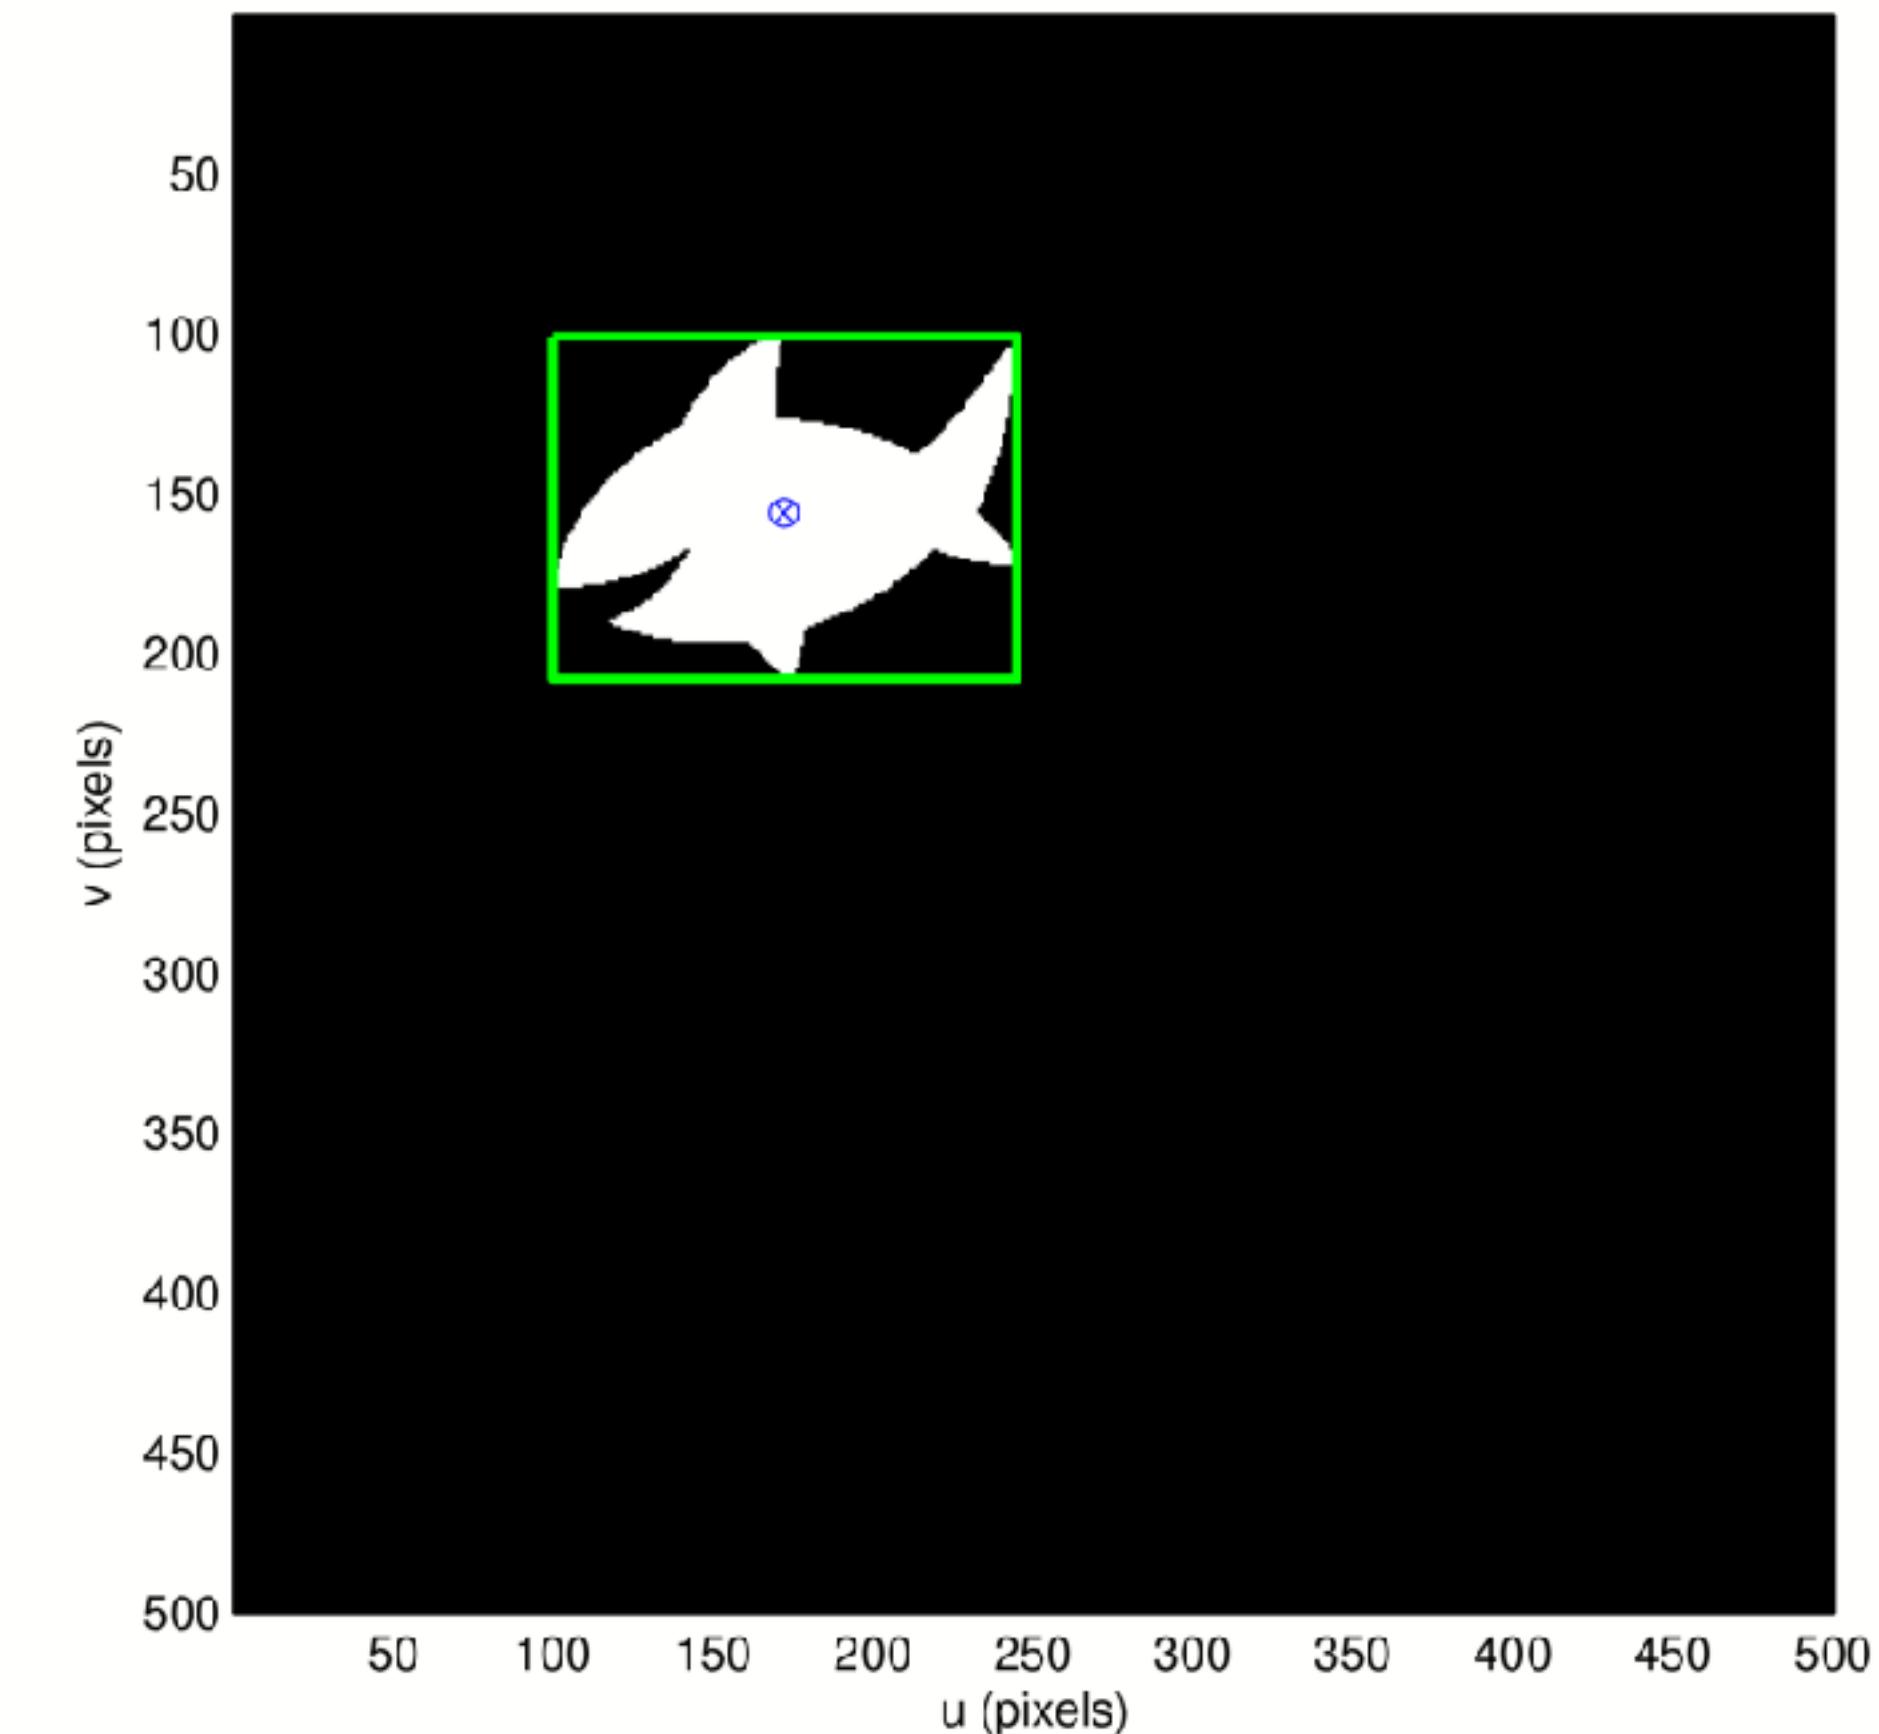
Where are the symbols on this sign?

Finding regions

<http://sweetclipart.com>



a bunch of pixels



a region

Moments

The p - q^{th} moment of an image is given by $m_{pq} = \sum_{(u,v) \in I} u^p v^q I[u, v]$

- where $(p+q)$ is the **order** of the moment

The zeroth ($p=q=0$) moment of an image is simply $m_{00} = \sum_{(u,v) \in I} I[u, v]$

- which is the number of one pixels in the image
- the **area** of the region

First moments

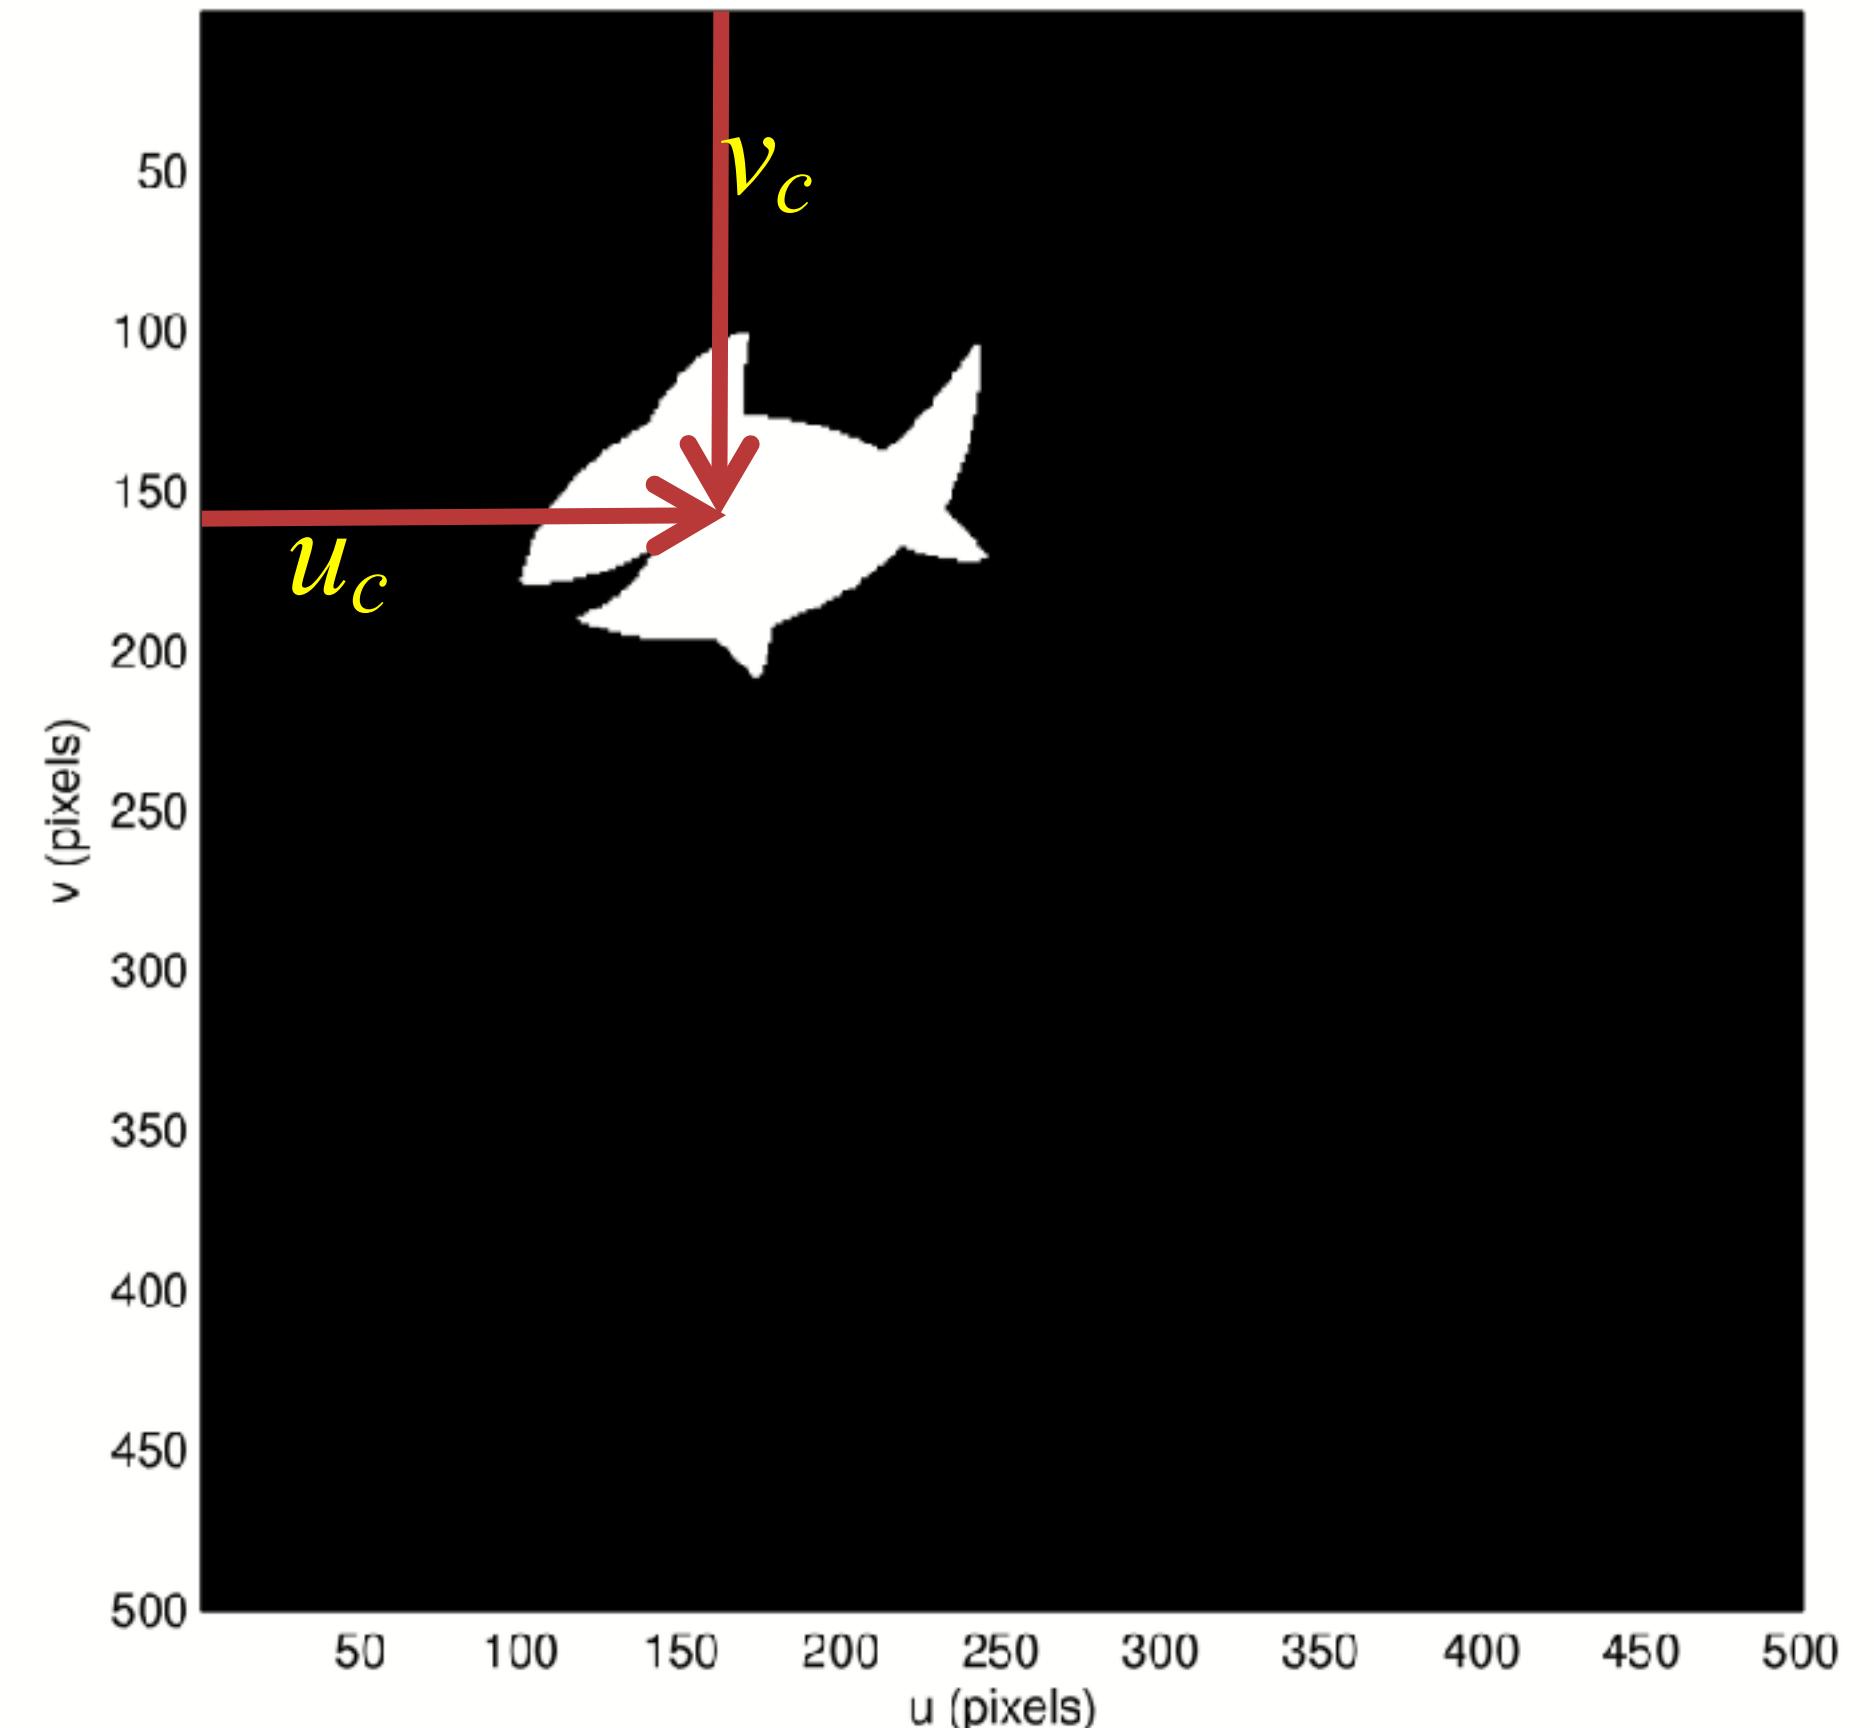
$$m_{10} = \sum_{(u,v) \in I} u I[u, v]$$

$$m_{01} = \sum_{(u,v) \in I} v I[u, v]$$

The **centroid**, or **centre** of mass, is given by

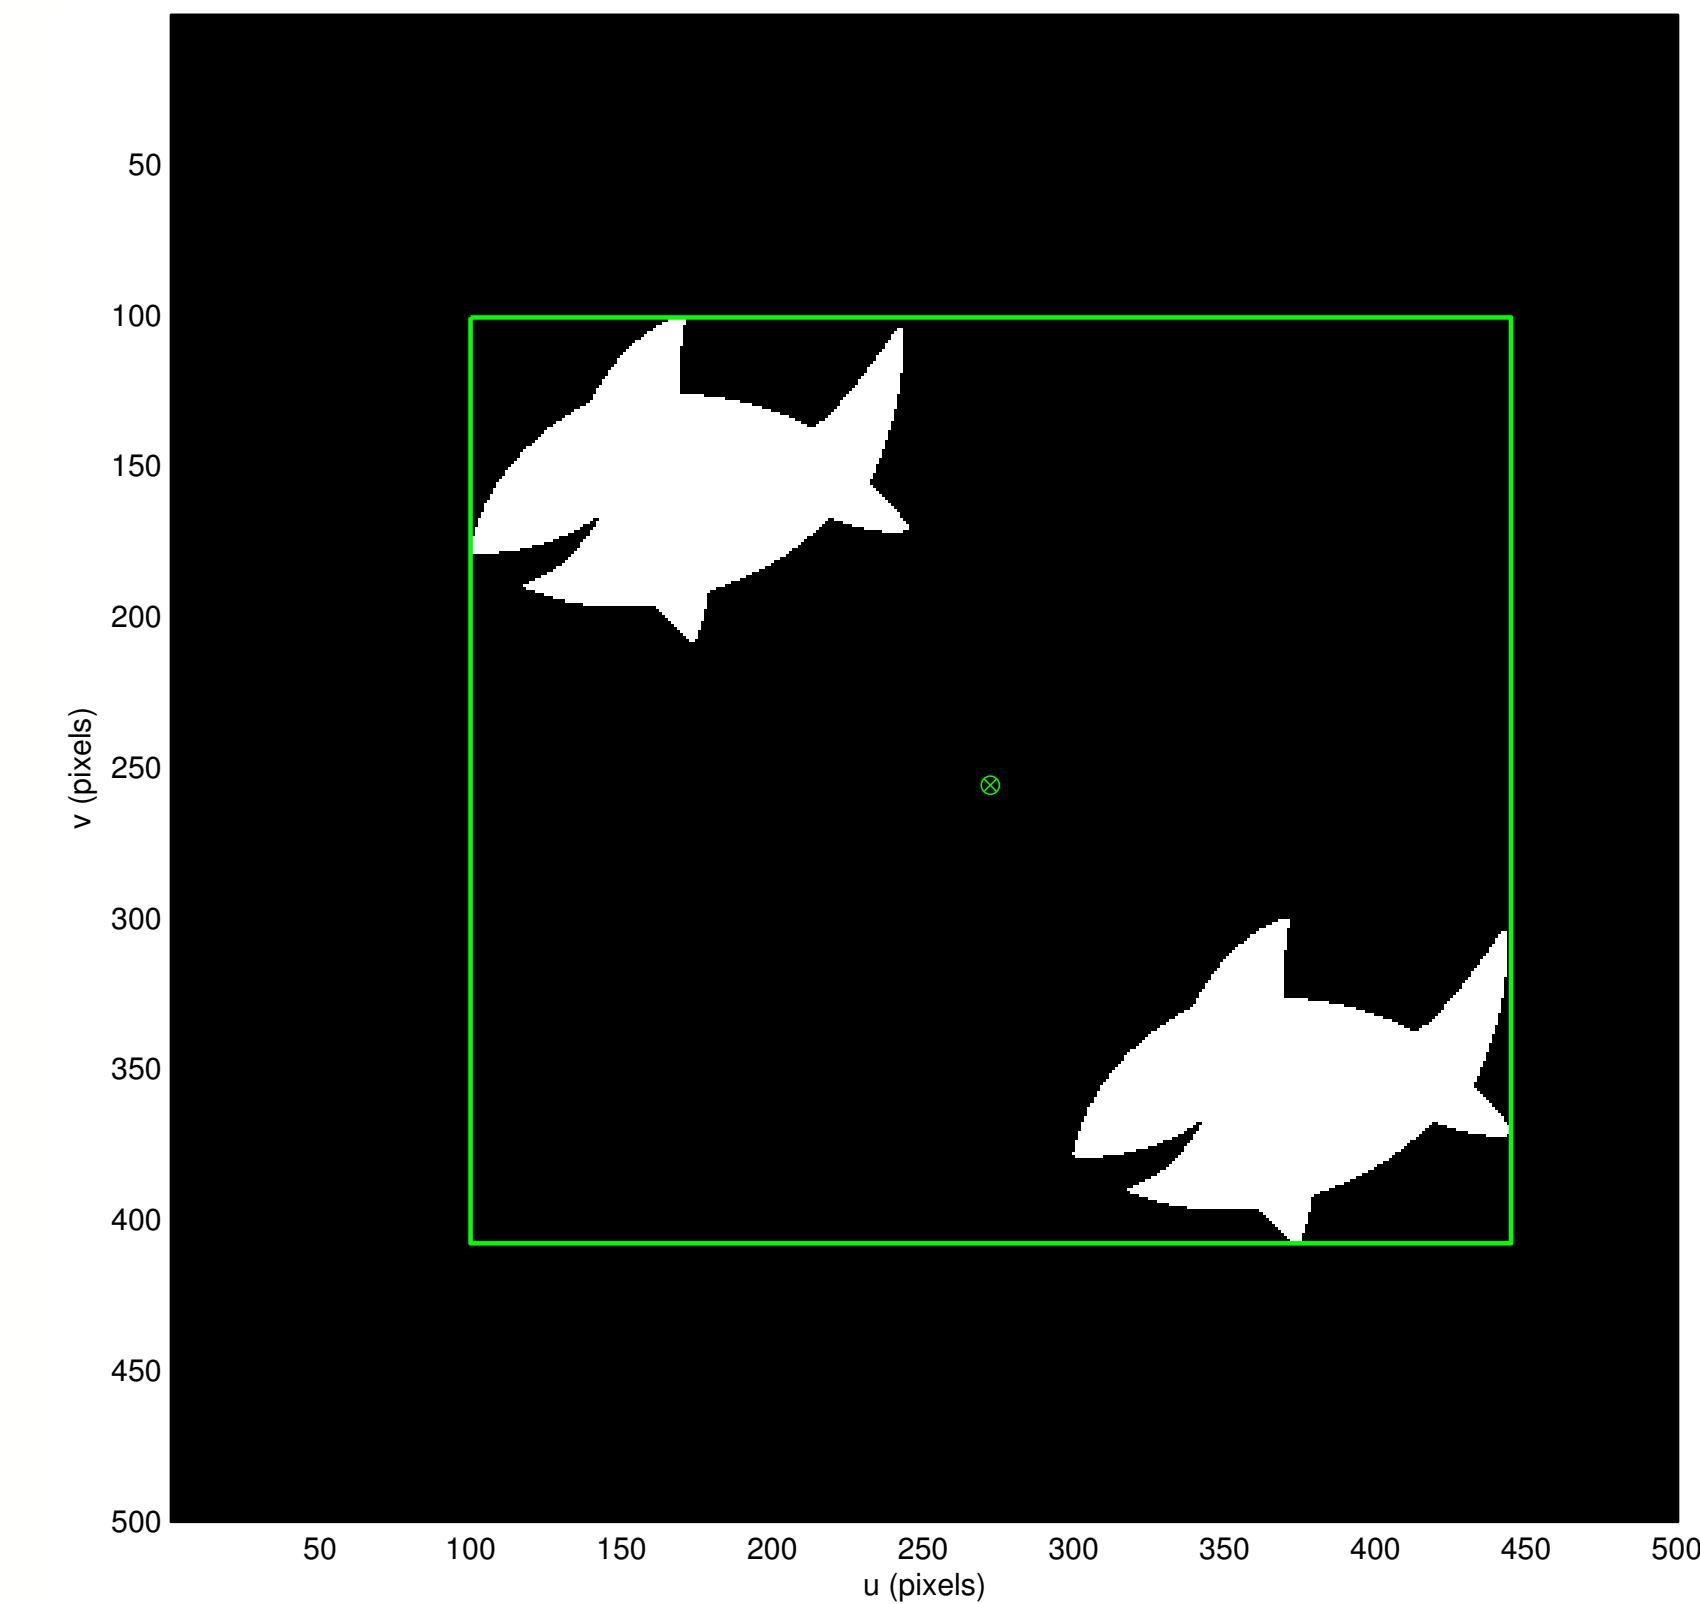
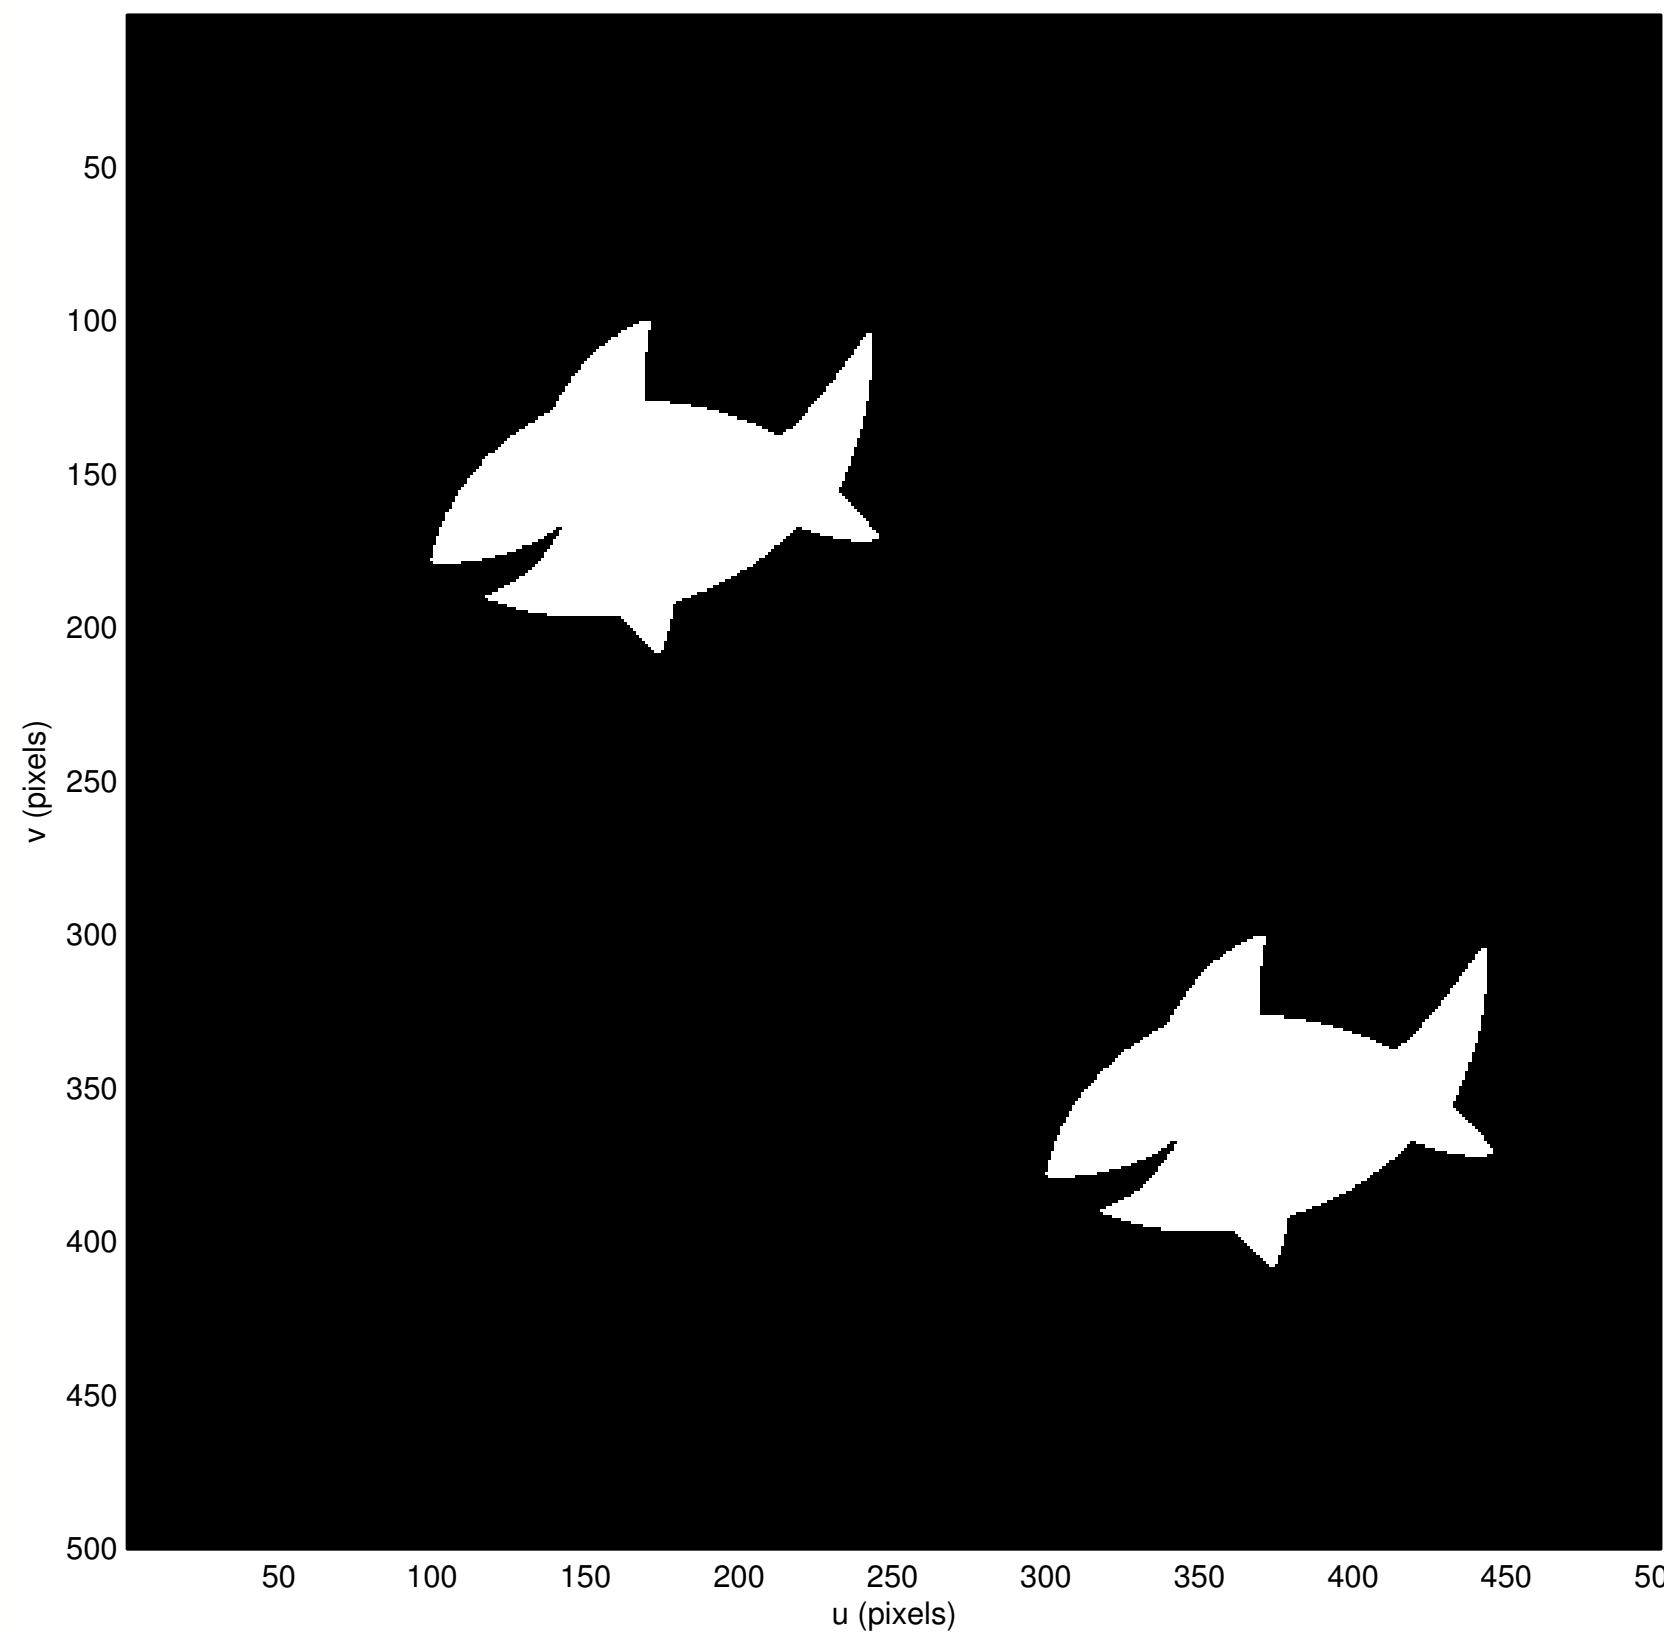
$$u_c = \frac{m_{10}}{m_{00}}, \quad v_c = \frac{m_{01}}{m_{00}}$$

<http://sweetclipart.com>



Multiple blobs

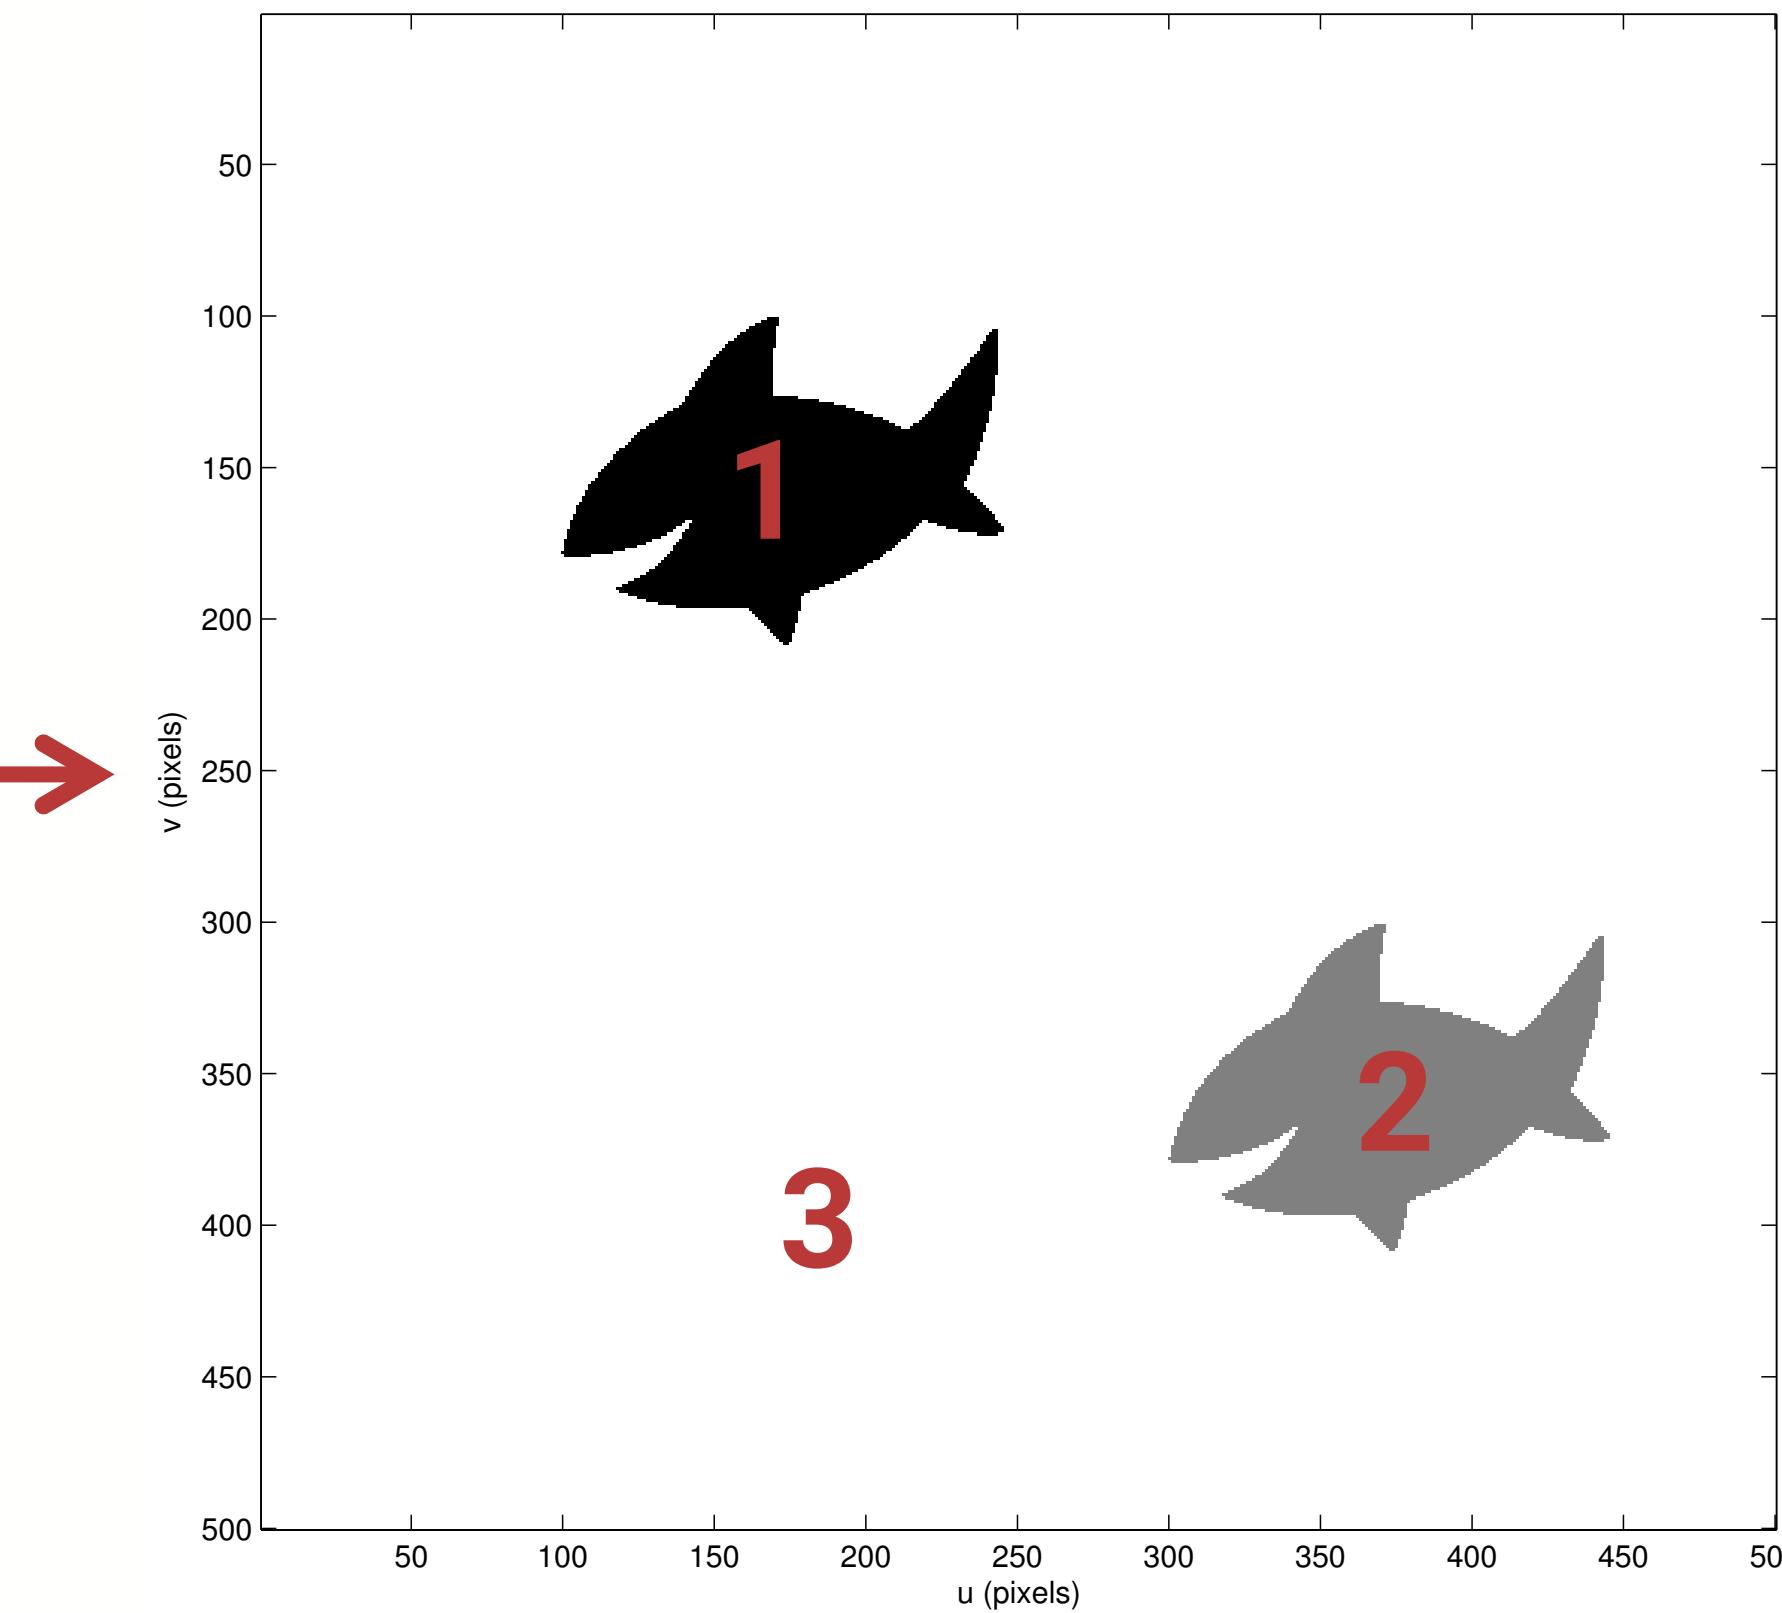
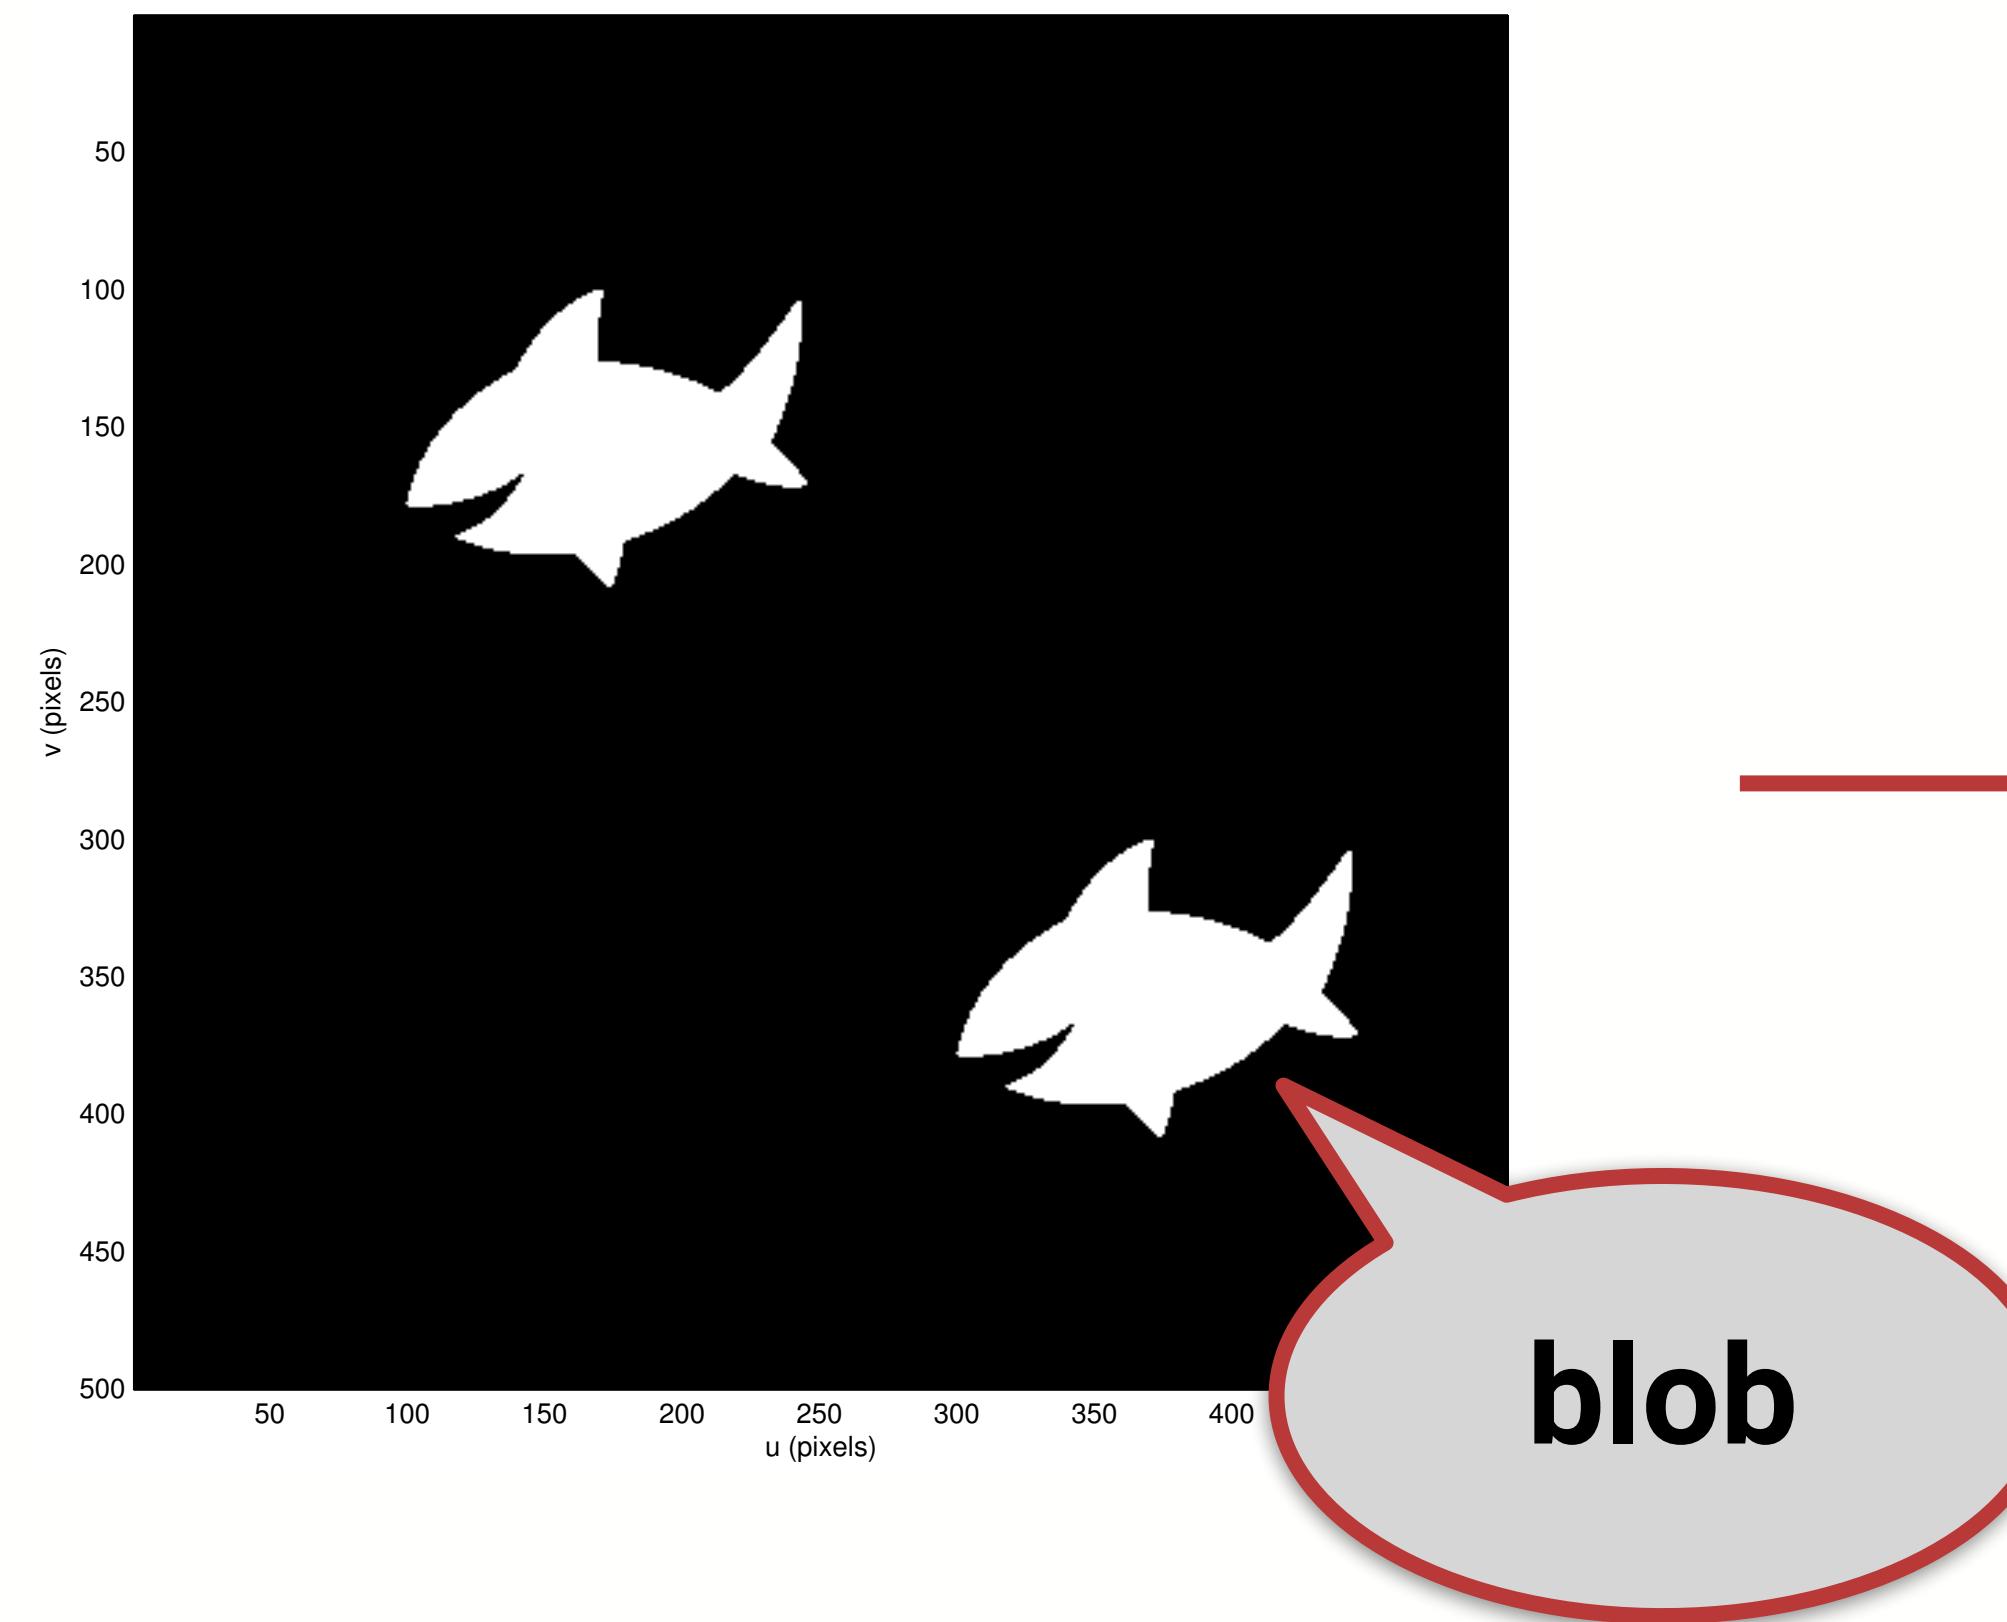
<http://sweetclipart.com>



- We can see two distinct regions
- but the moment computation has no sense of this

Transform the image

<http://sweetclipart.com>

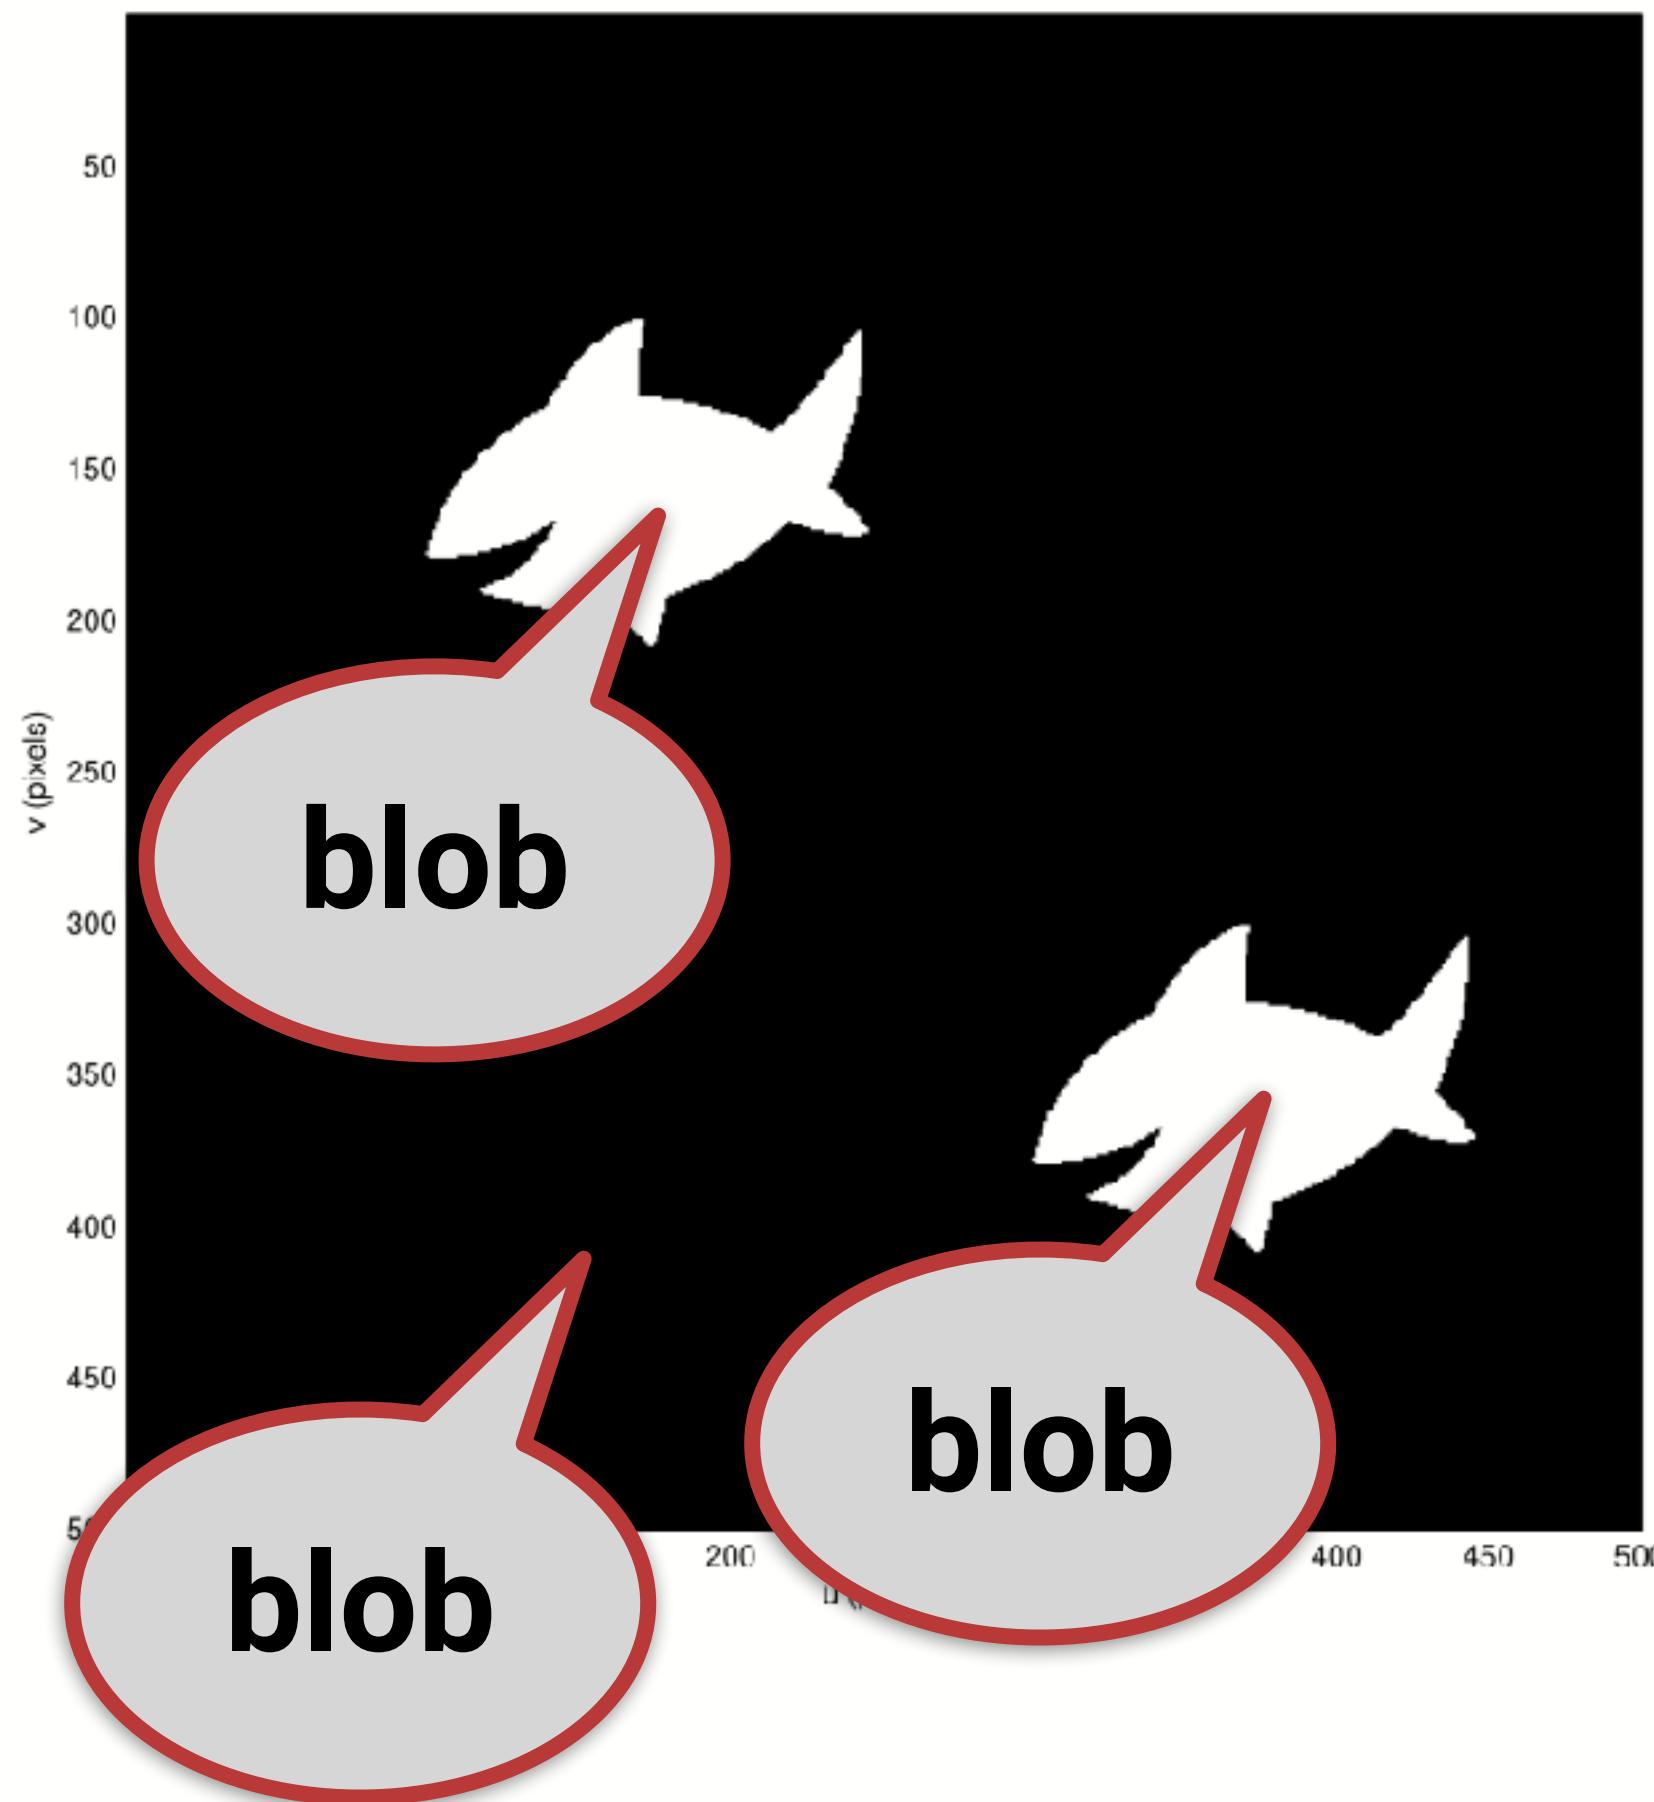


- We apply **connectivity analysis**
- Each pixel in the output image (value 1 to 3) indicates the **blob number** of the corresponding input pixel

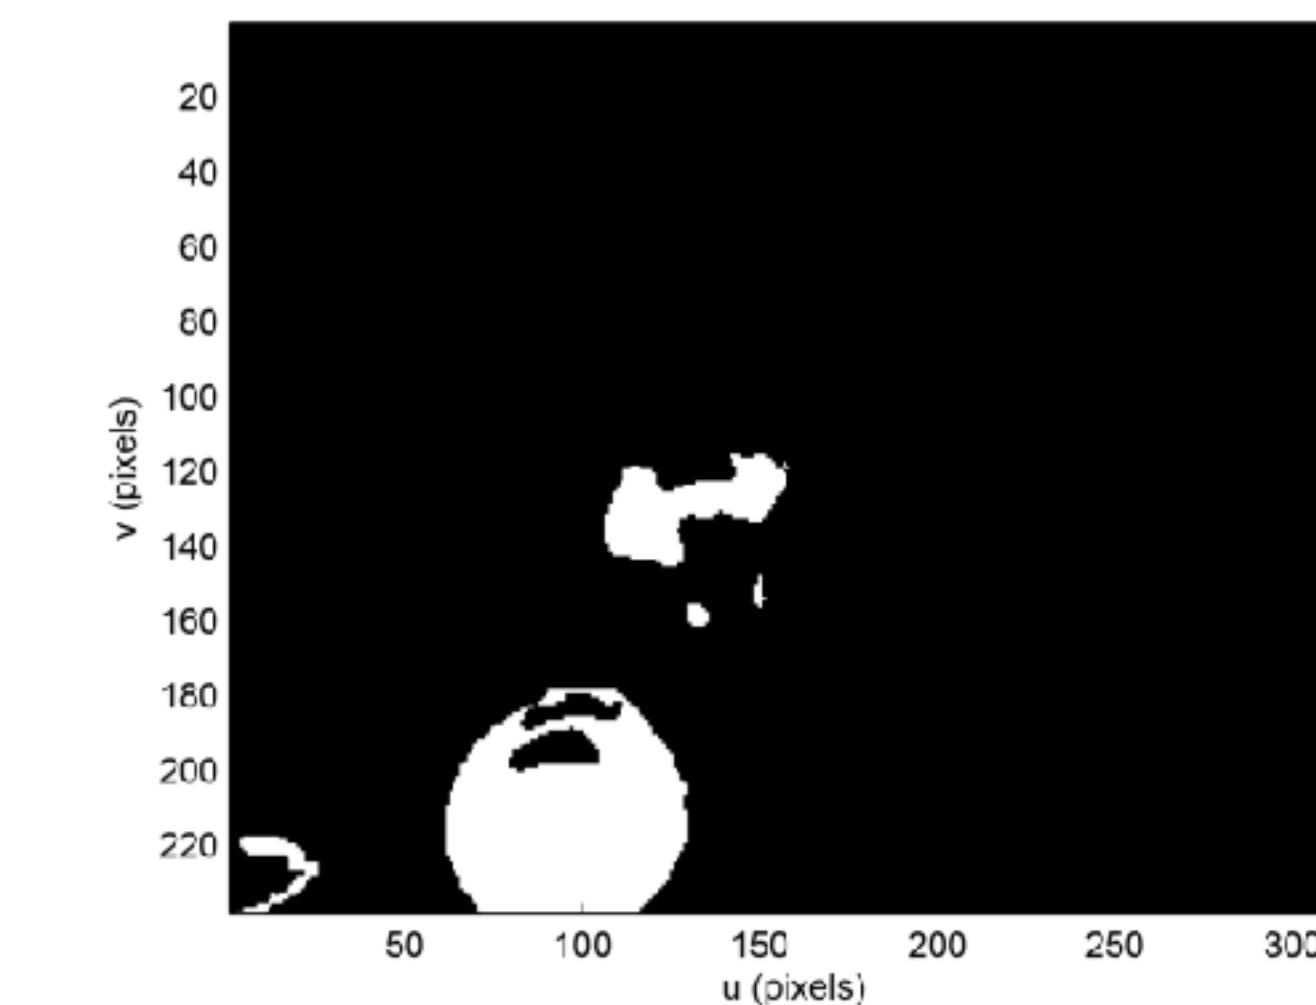
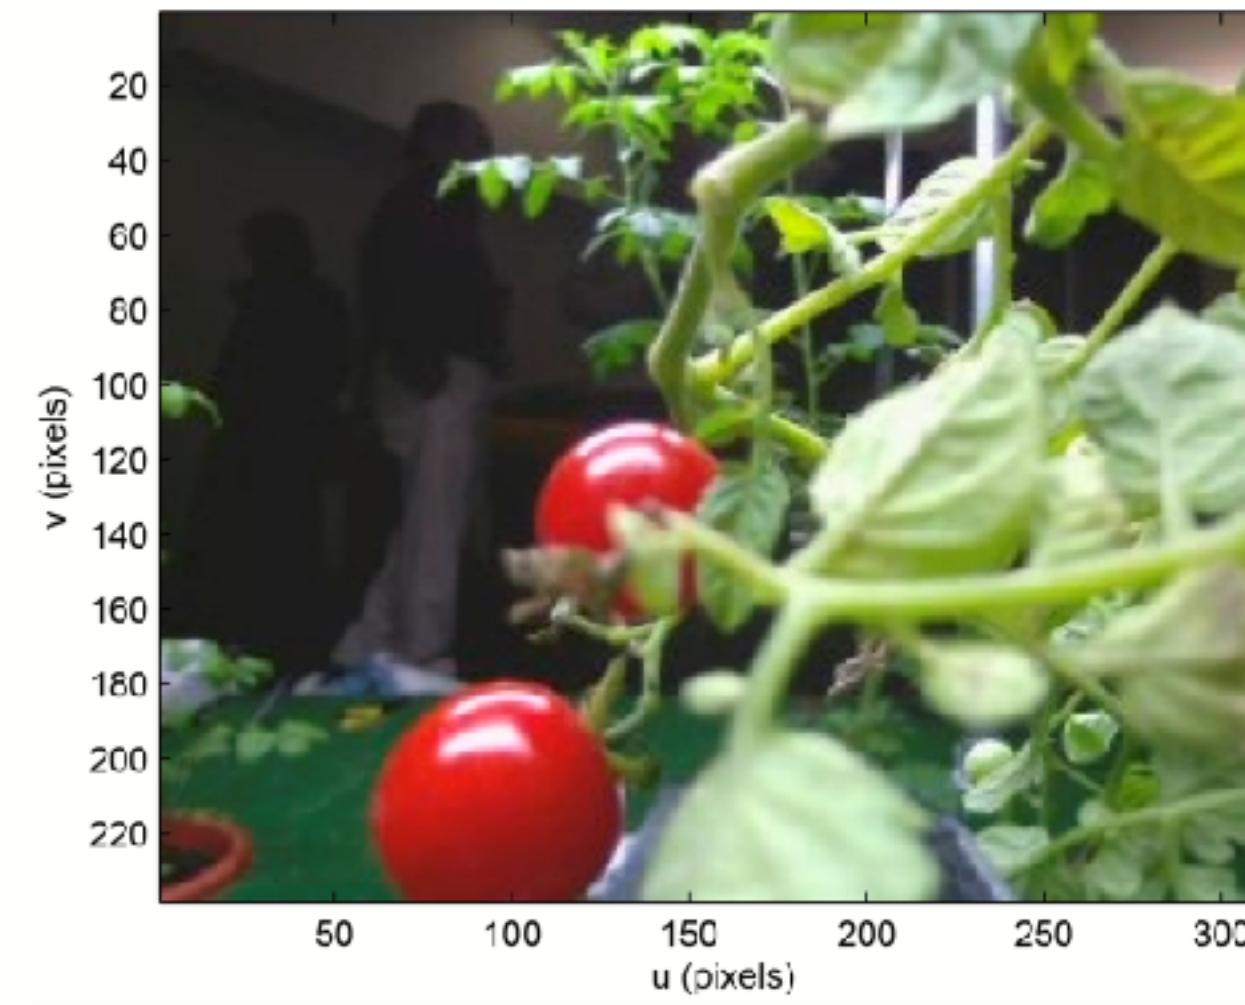
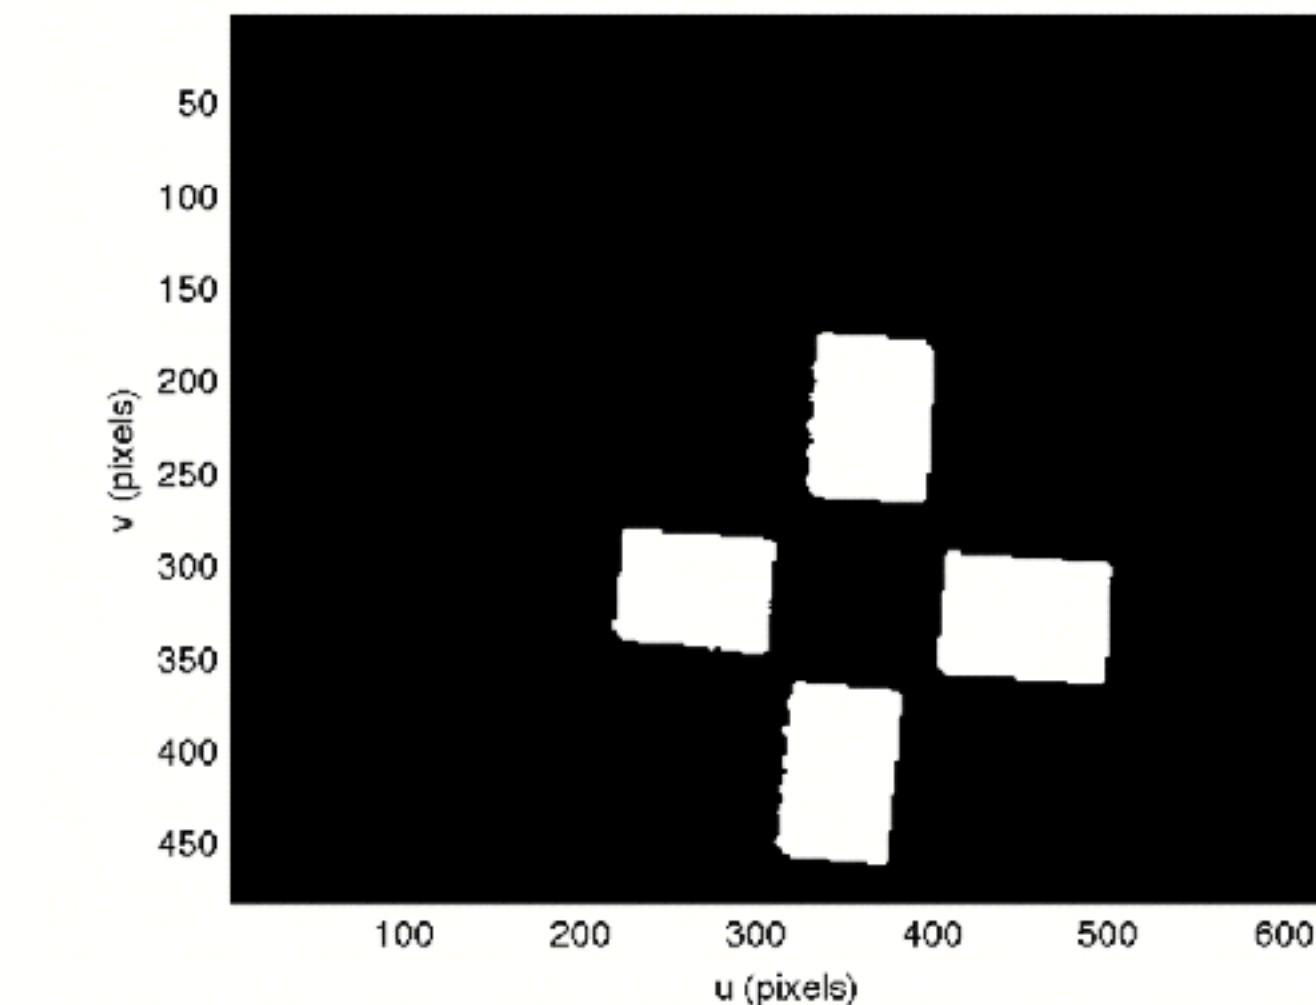
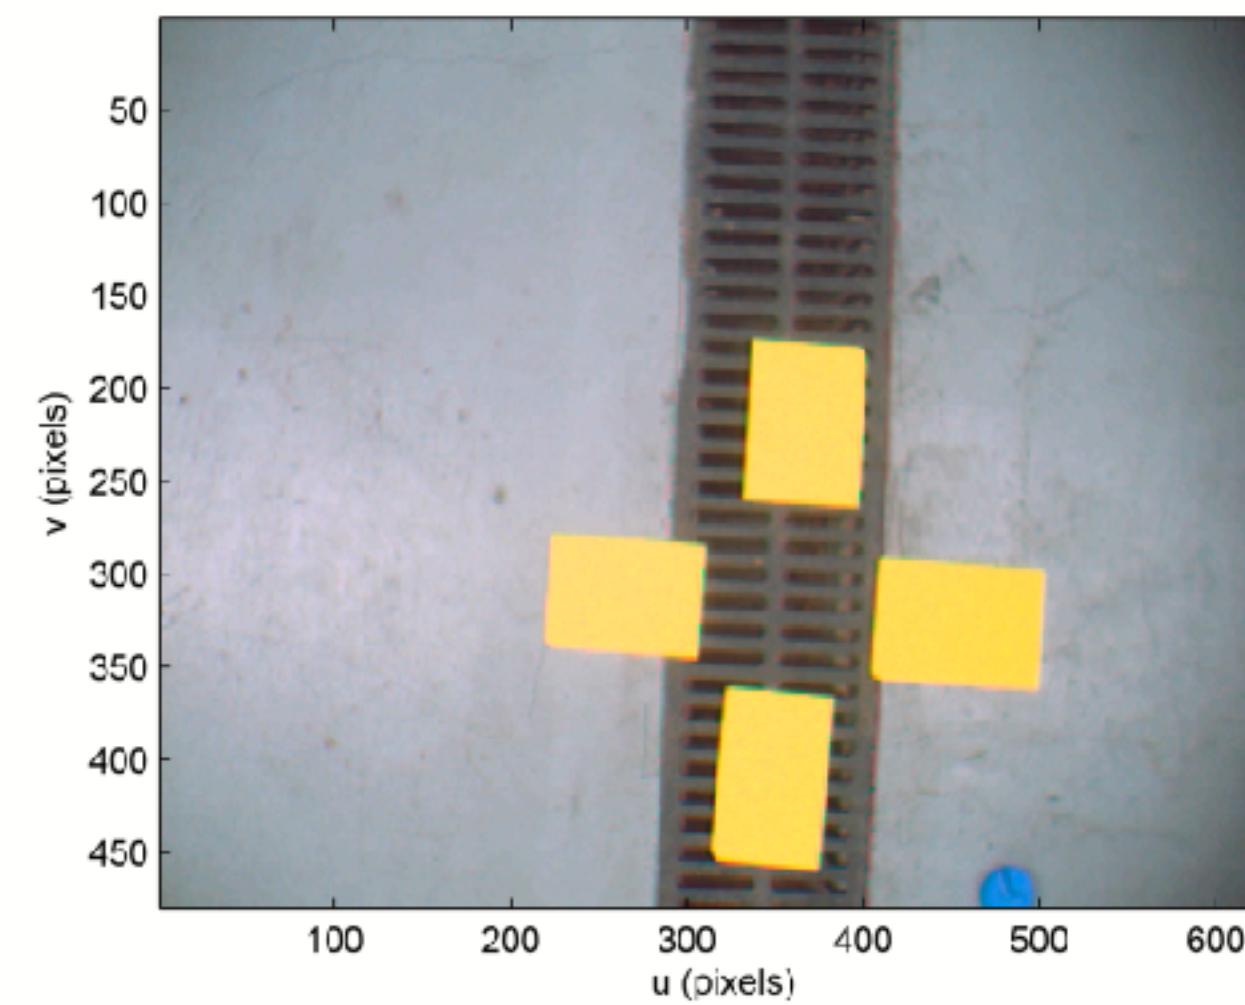
What is a blob?

- aka region, connected component
- a set of **contiguous** pixels of the same color (value)

<http://sweetclipart.com>

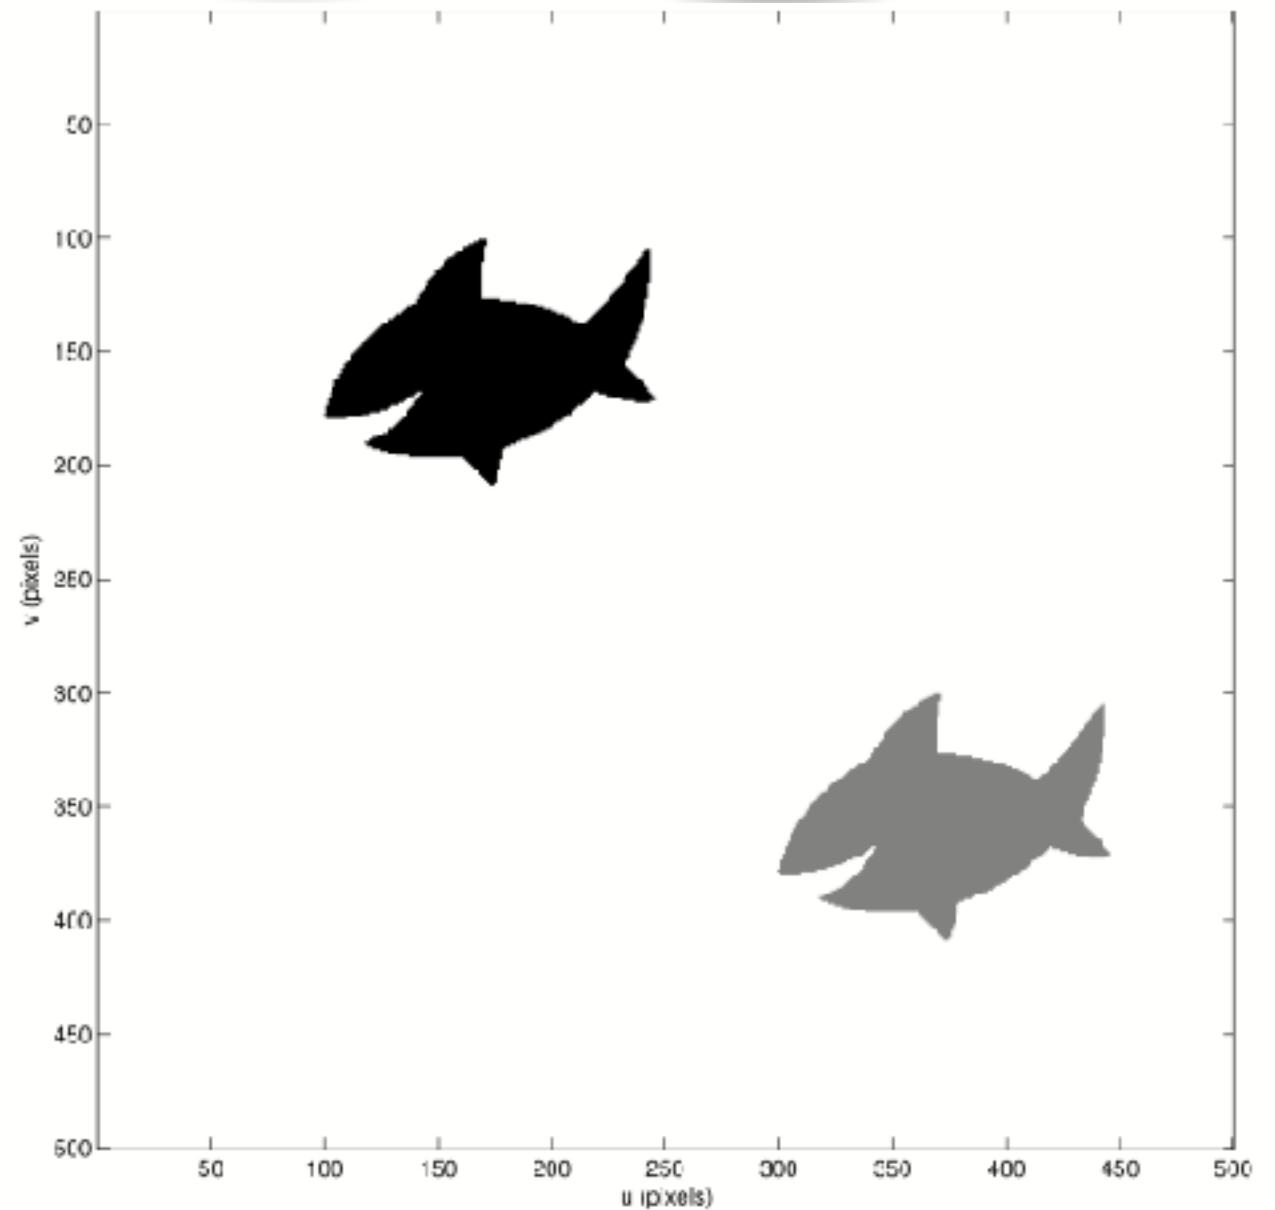
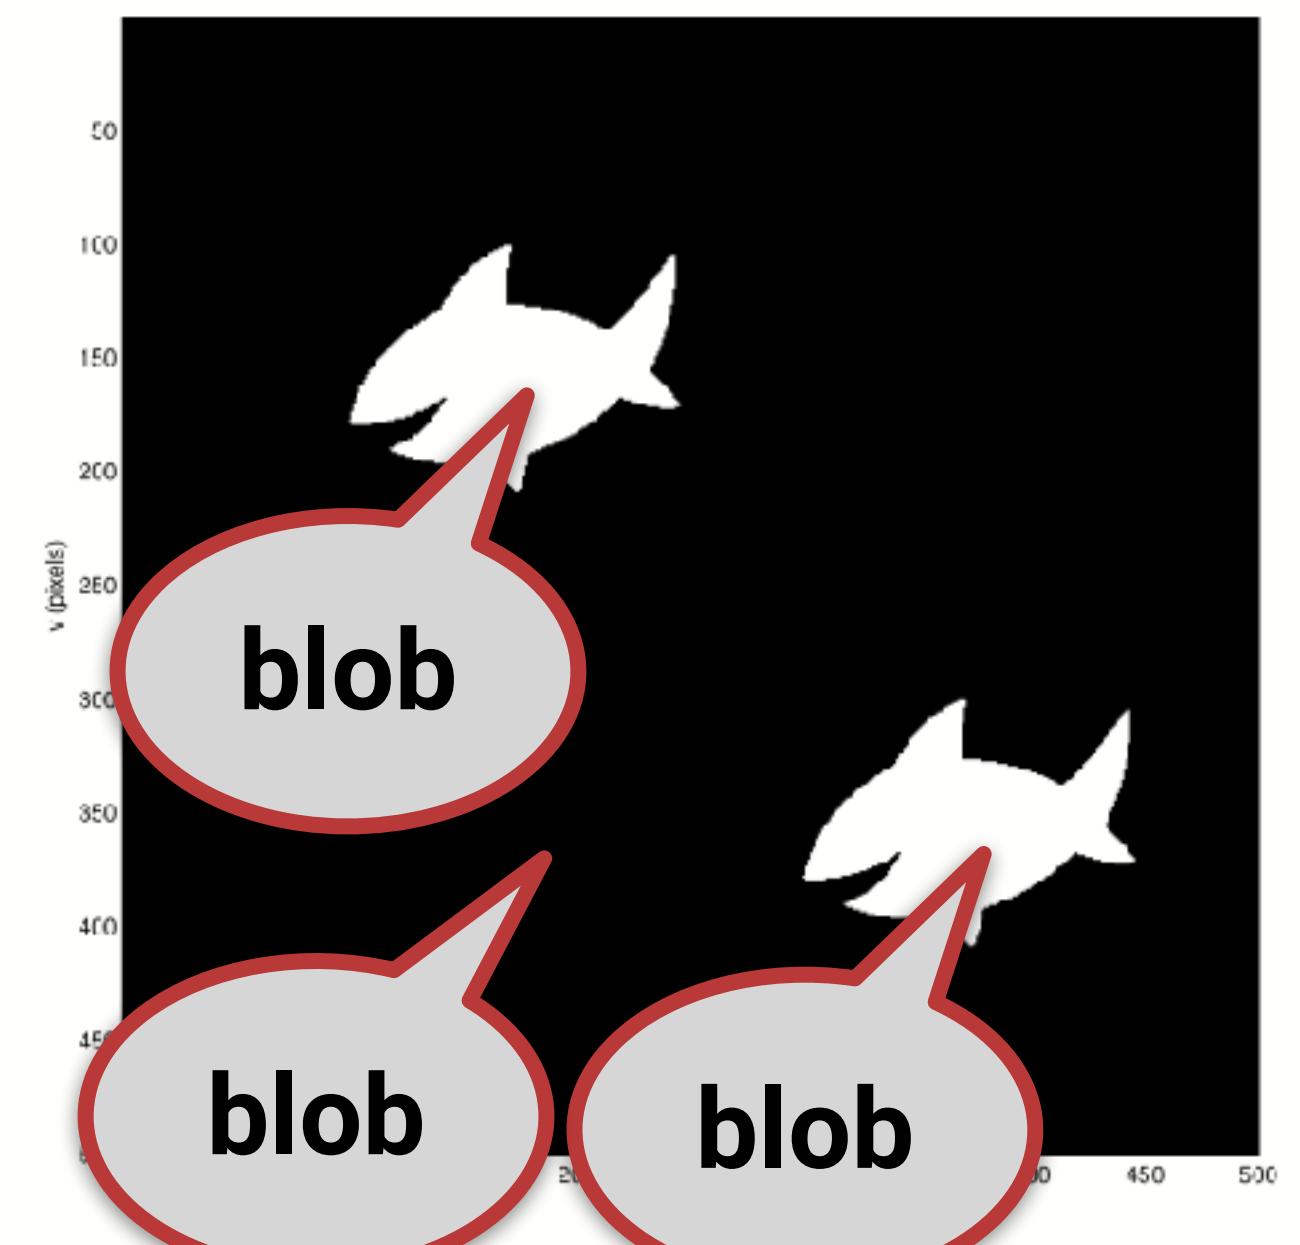


Pixel classification



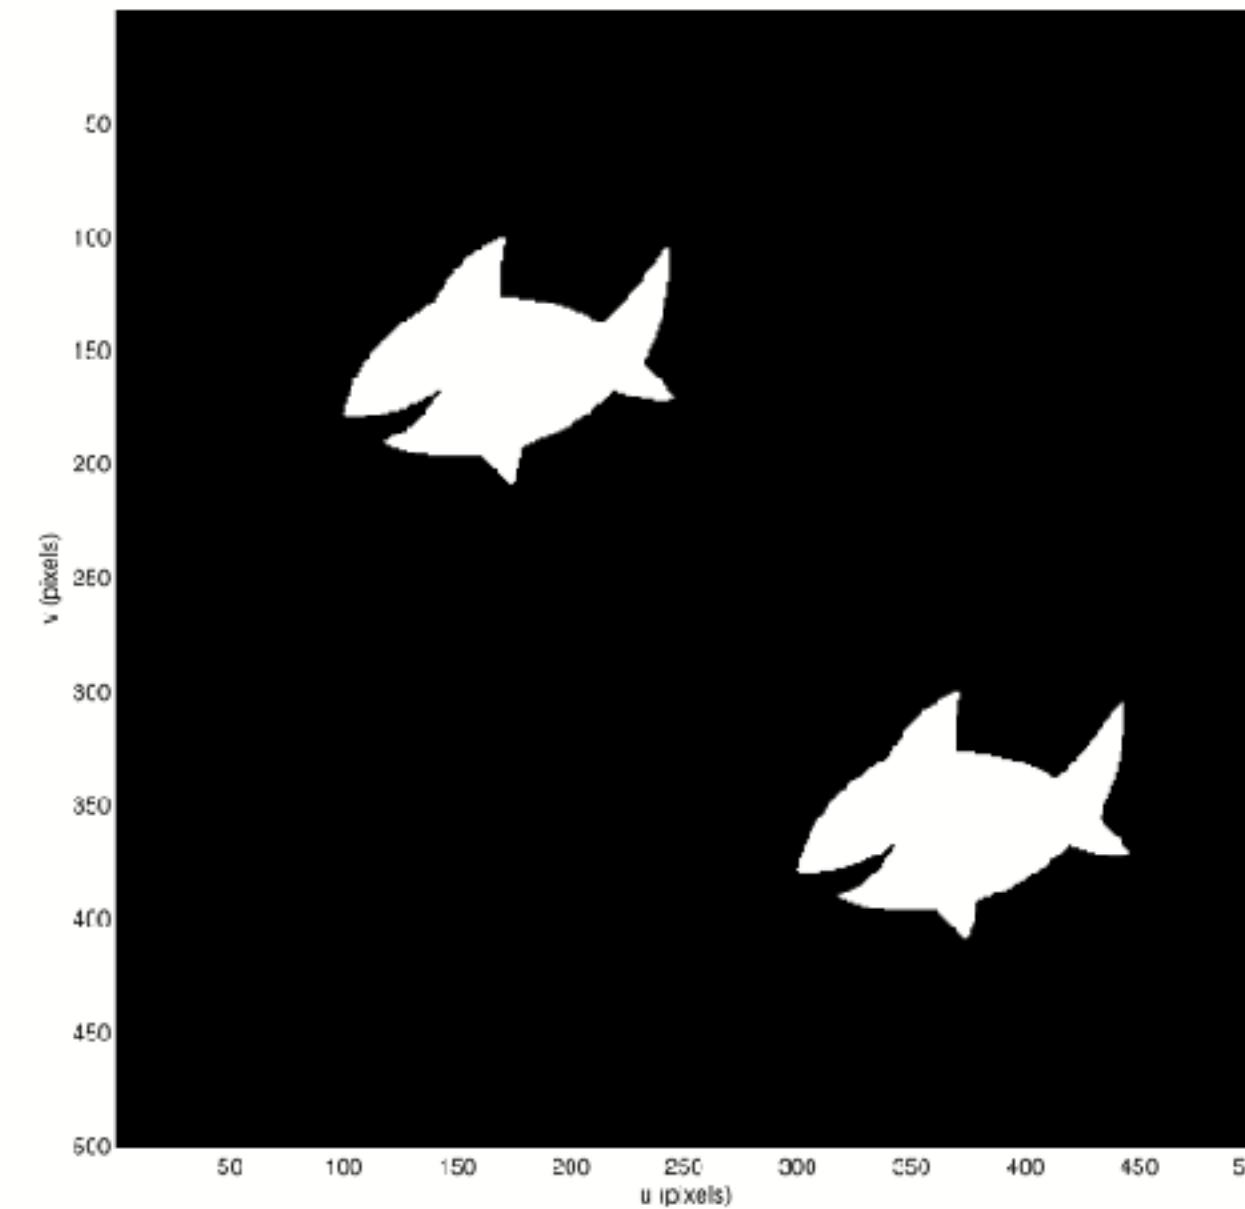
What is a blob?

- aka region, connected component
- a set of **contiguous** pixels of the **same color** (value)
- the label assigned to a pixel indicates which set (1...N) it belongs to
- every pixel has the same label as its N, S, E or W neighbour of the same color
- the process has many names:
 - connectivity/blob analysis,
 - connected component analysis,
 - blob/region/image labelling,
 - blob/region coloring

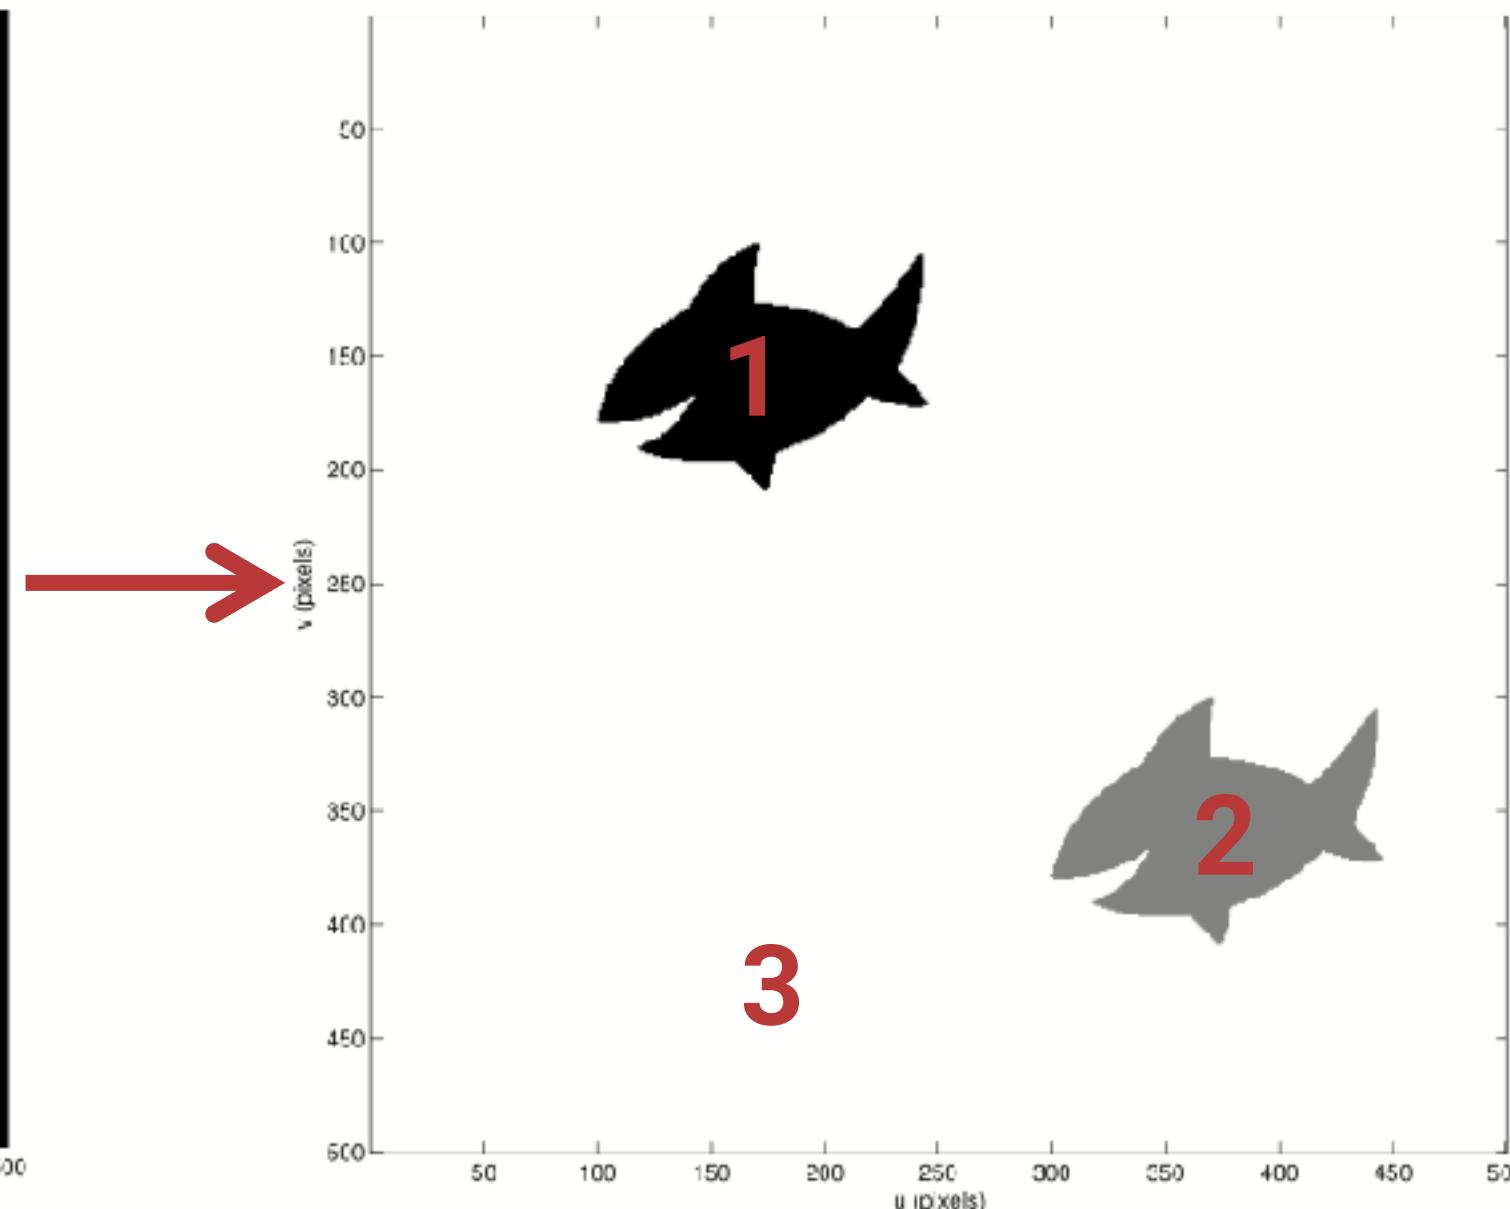


Select individual blobs

<http://sweetclipart.com>

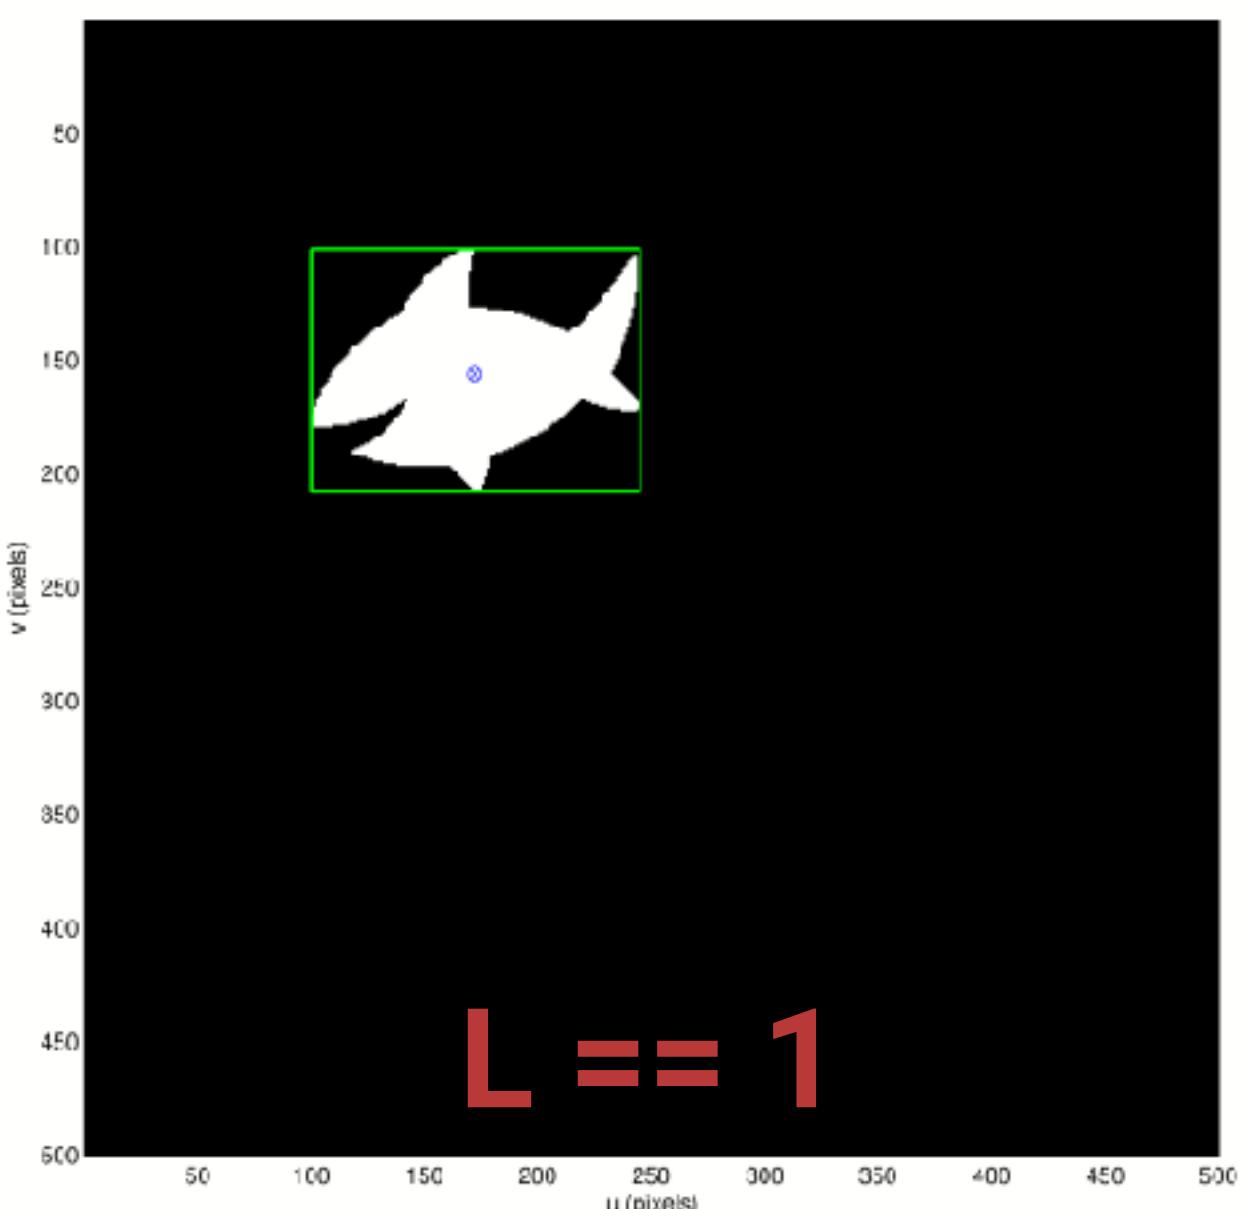


Original binary image

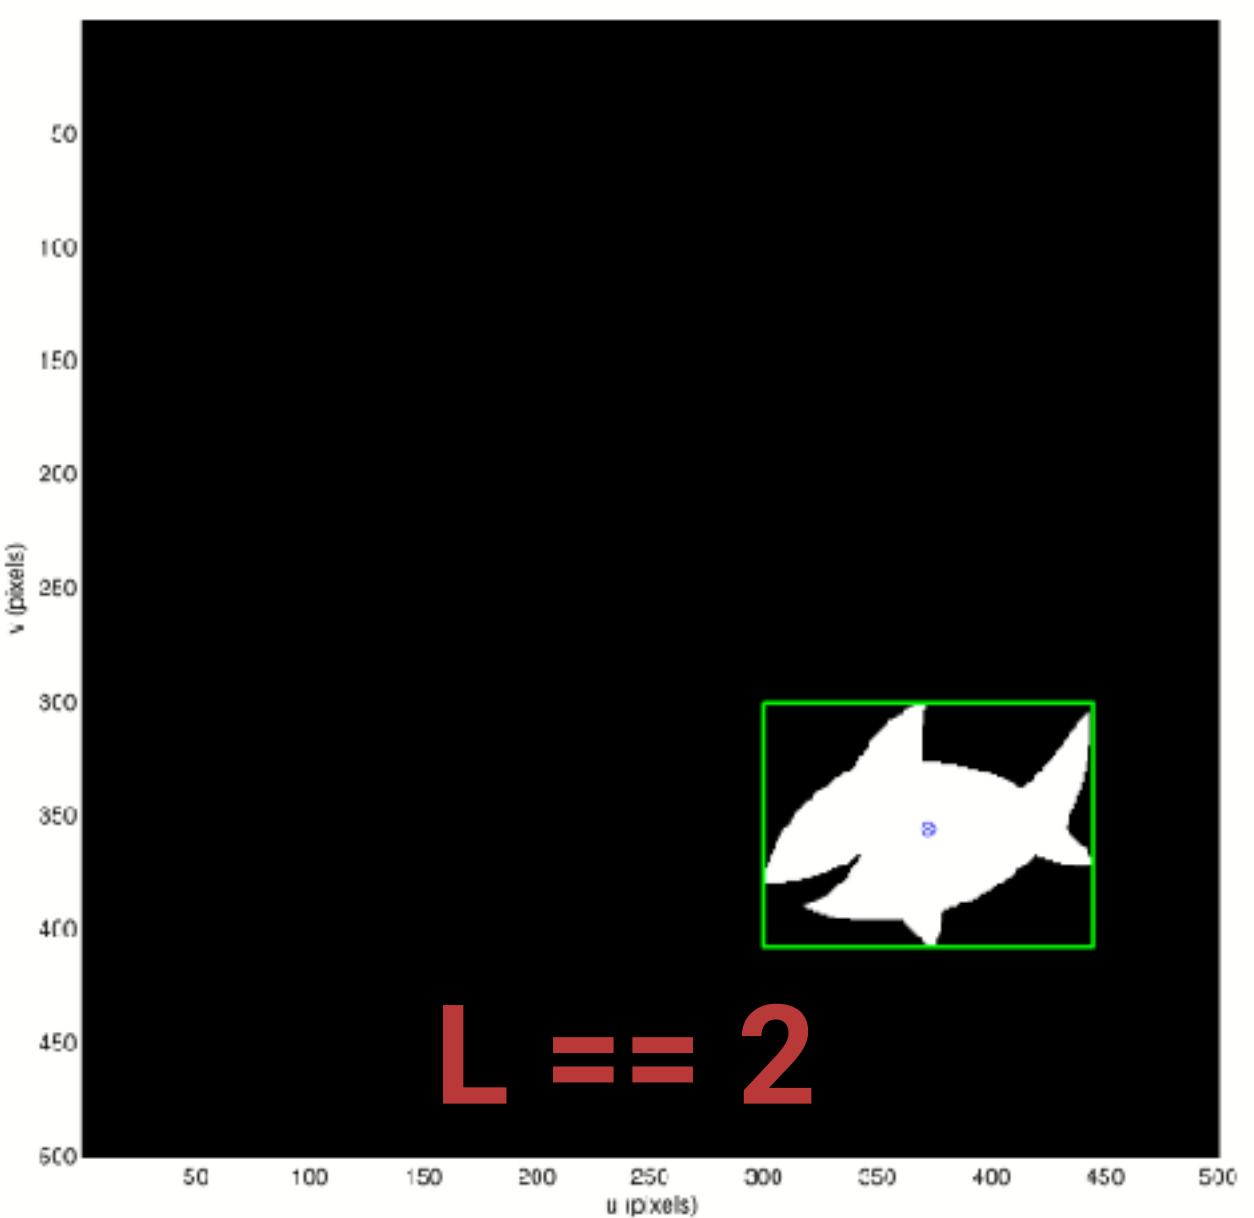


L = label image

Individual blob binary image

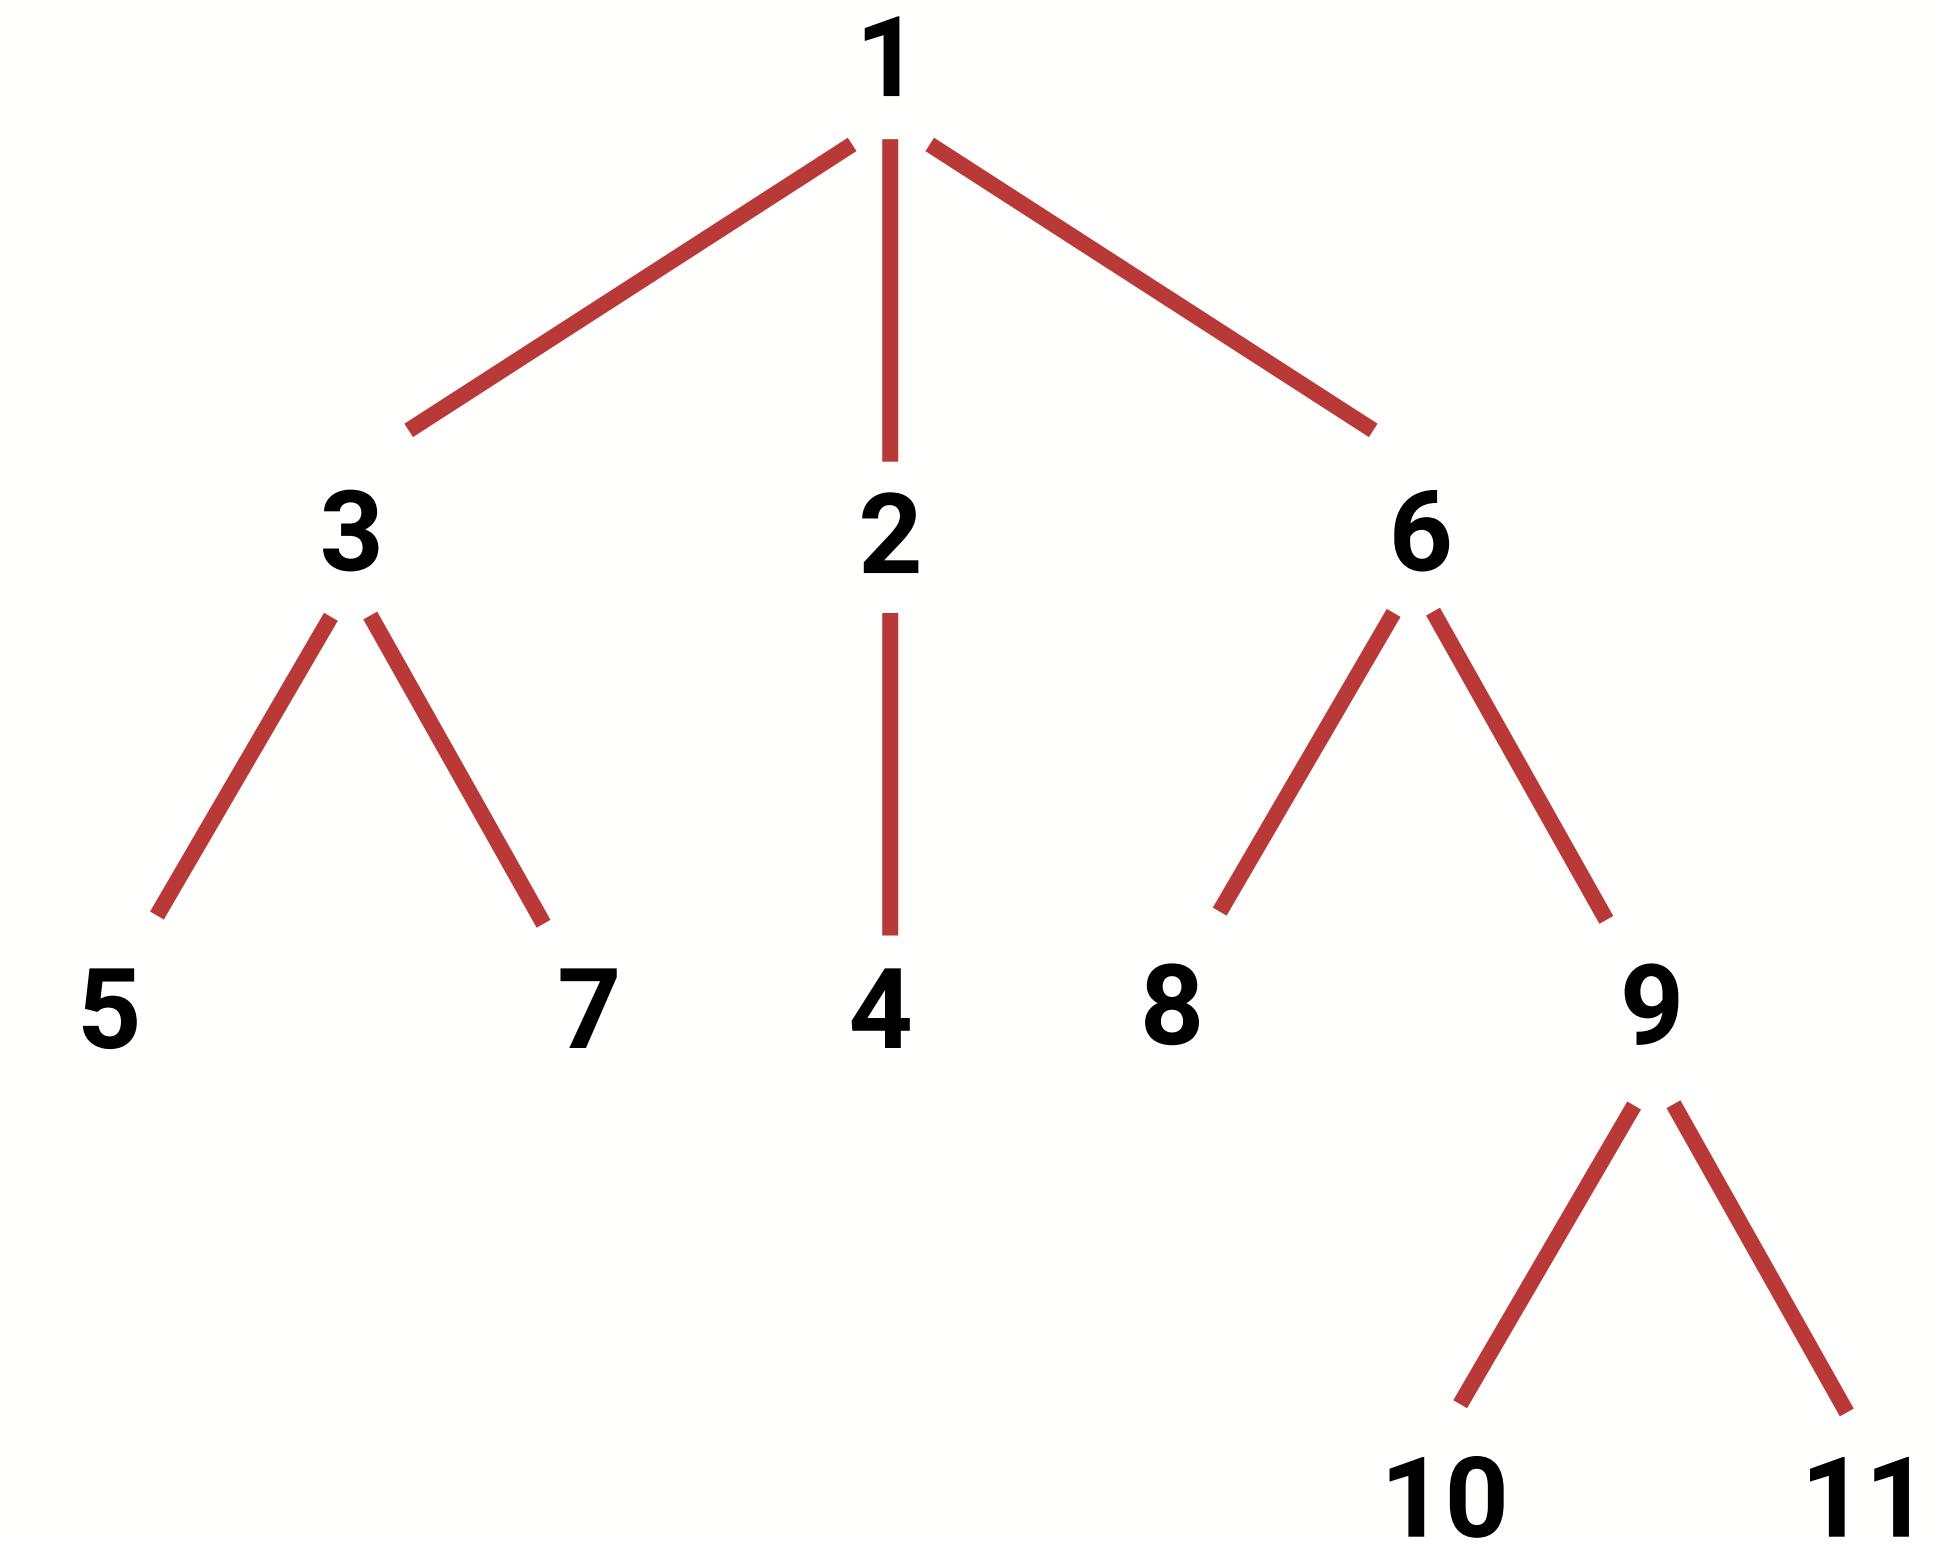
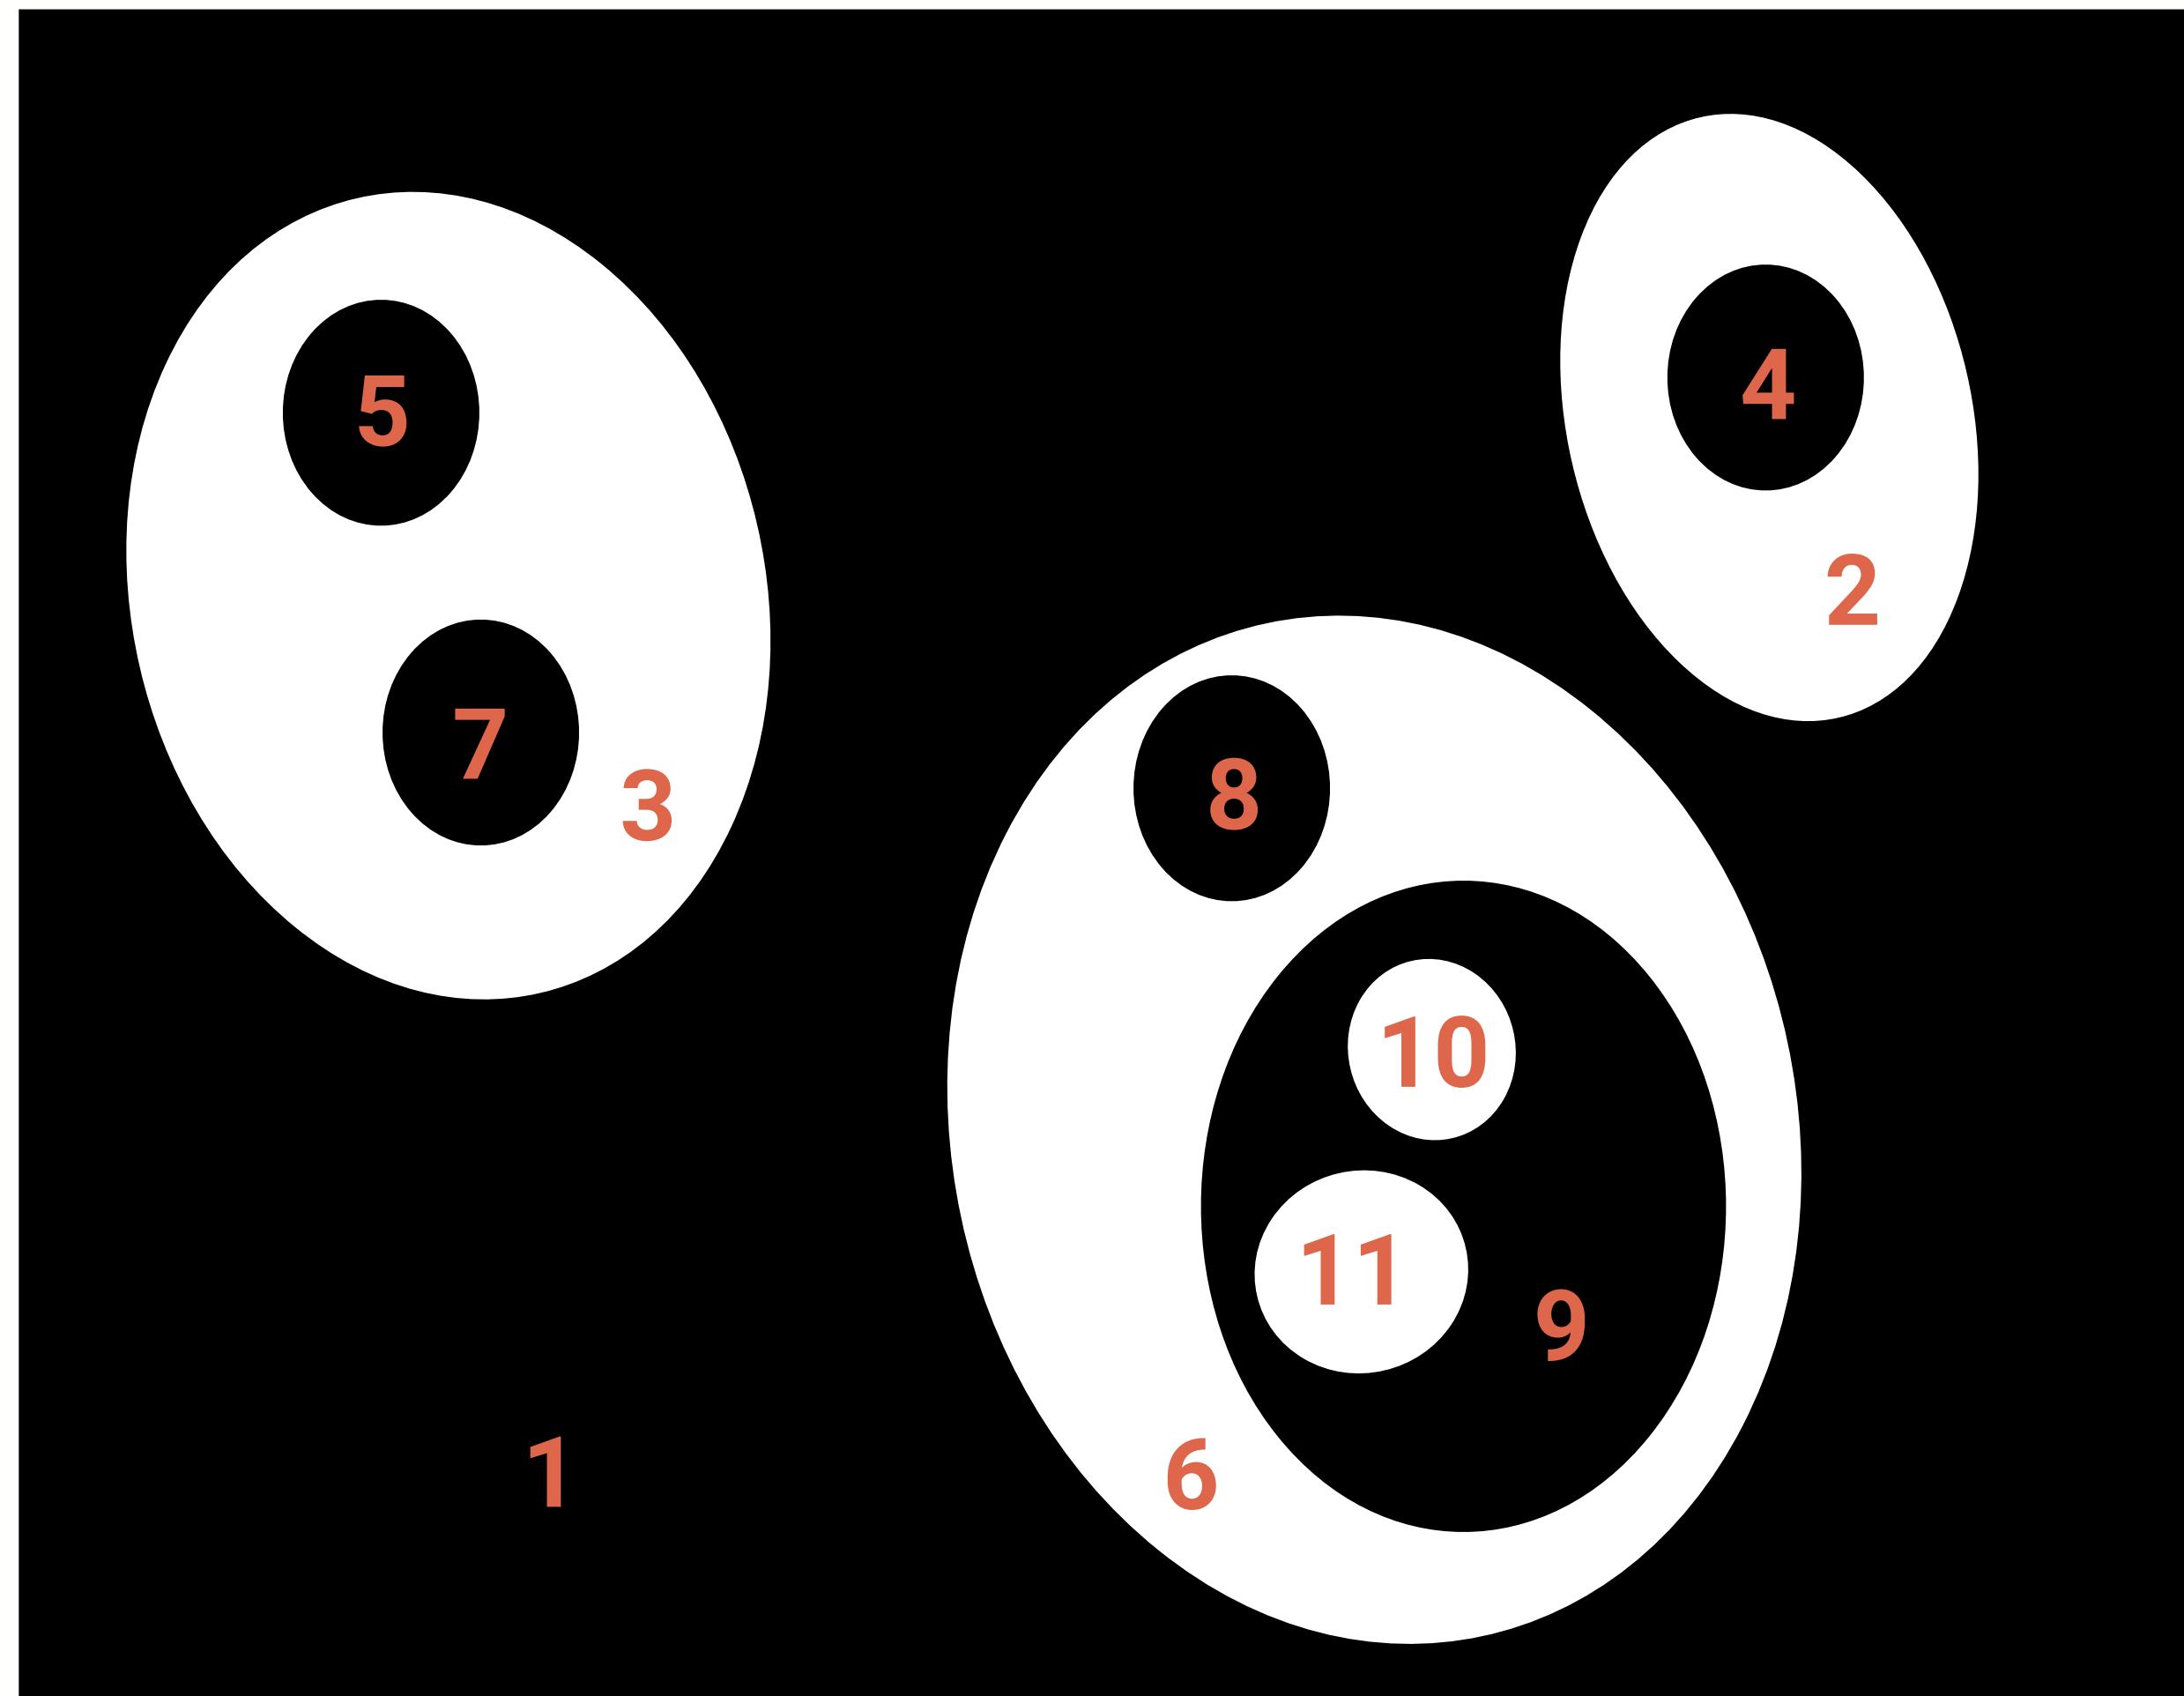


$L == 1$



$L == 2$

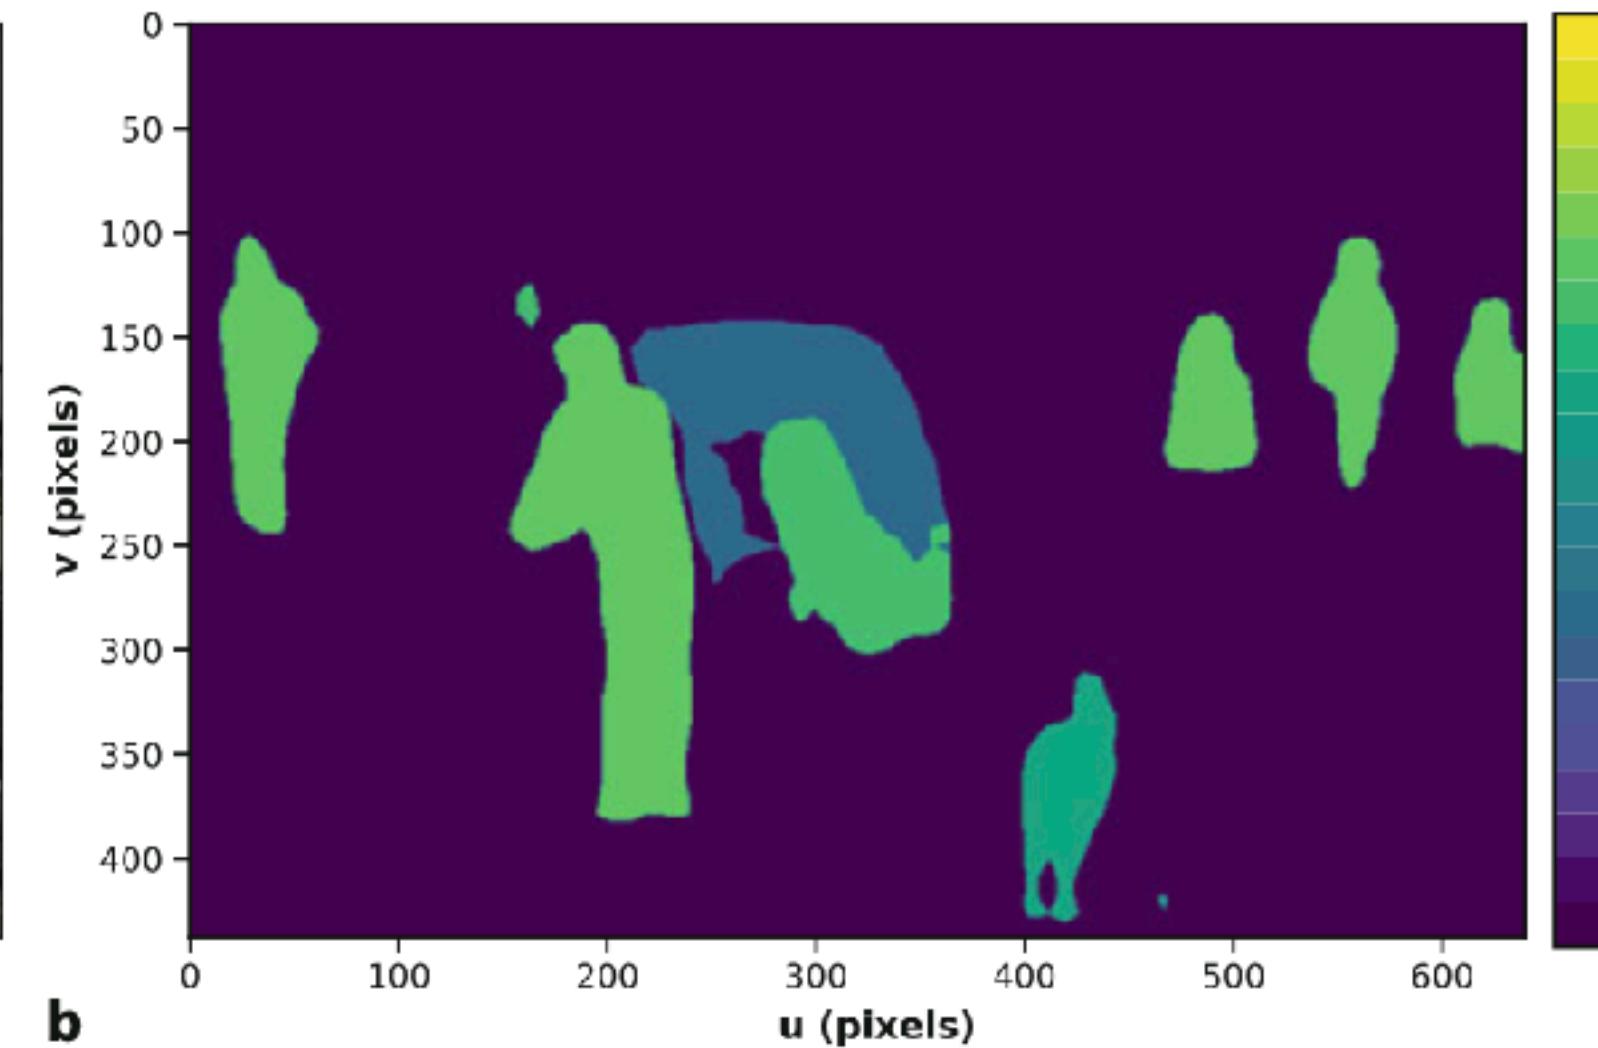
Blob hierarchy



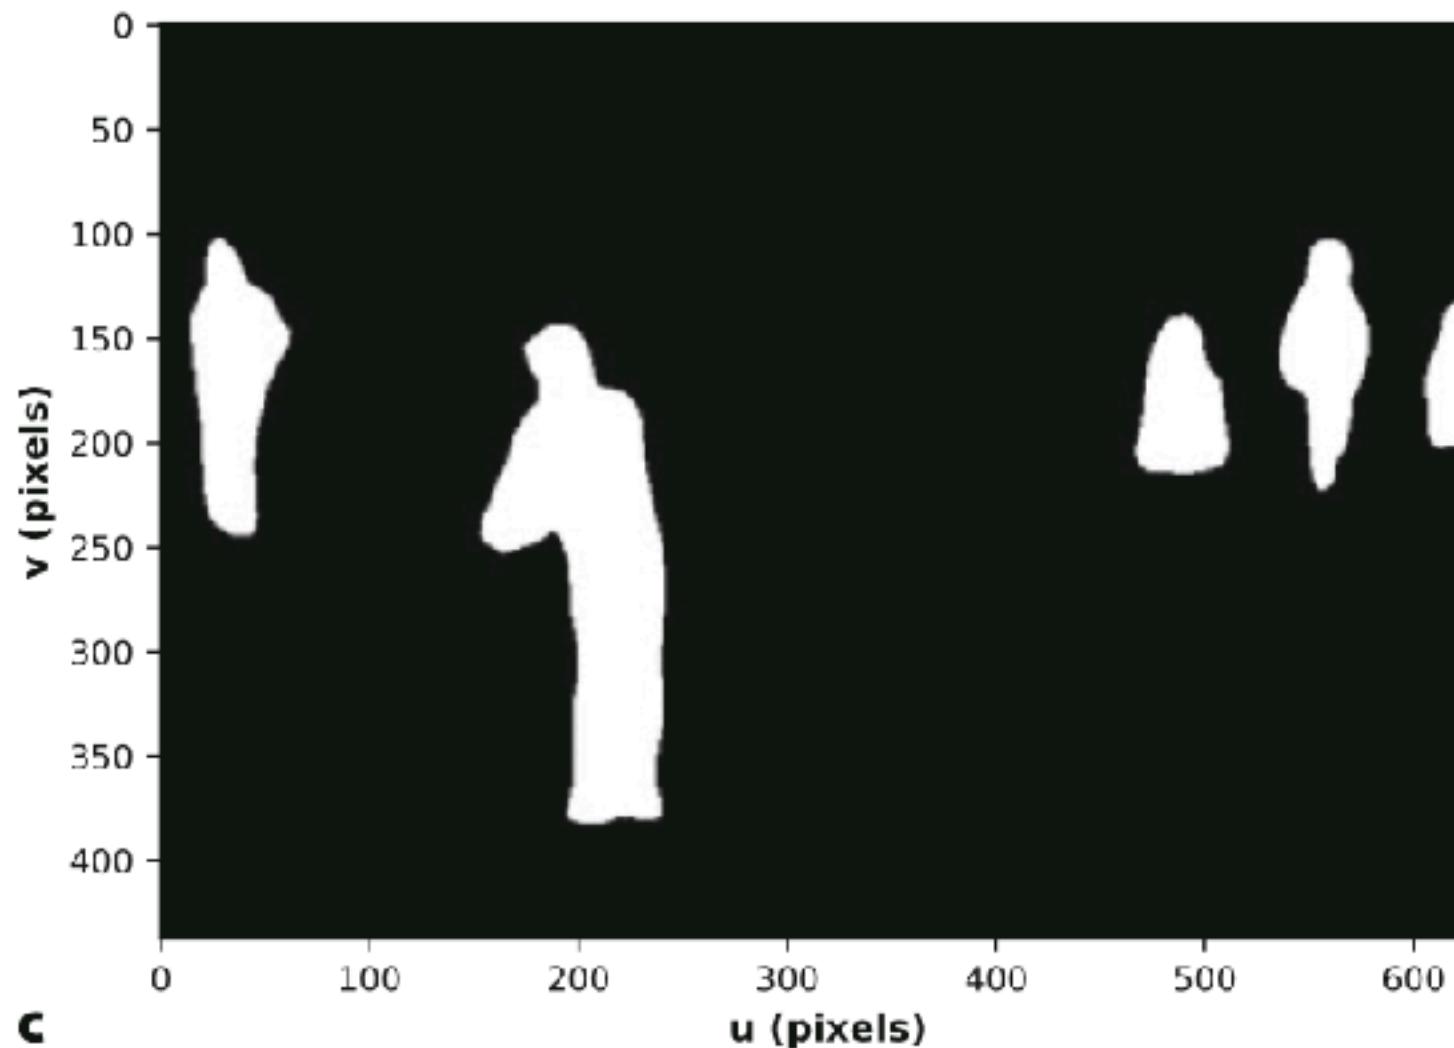
Deep networks allow us to classify pixels based on their semantic meaning (semantic segmentation) even if they have a range of brightness and color



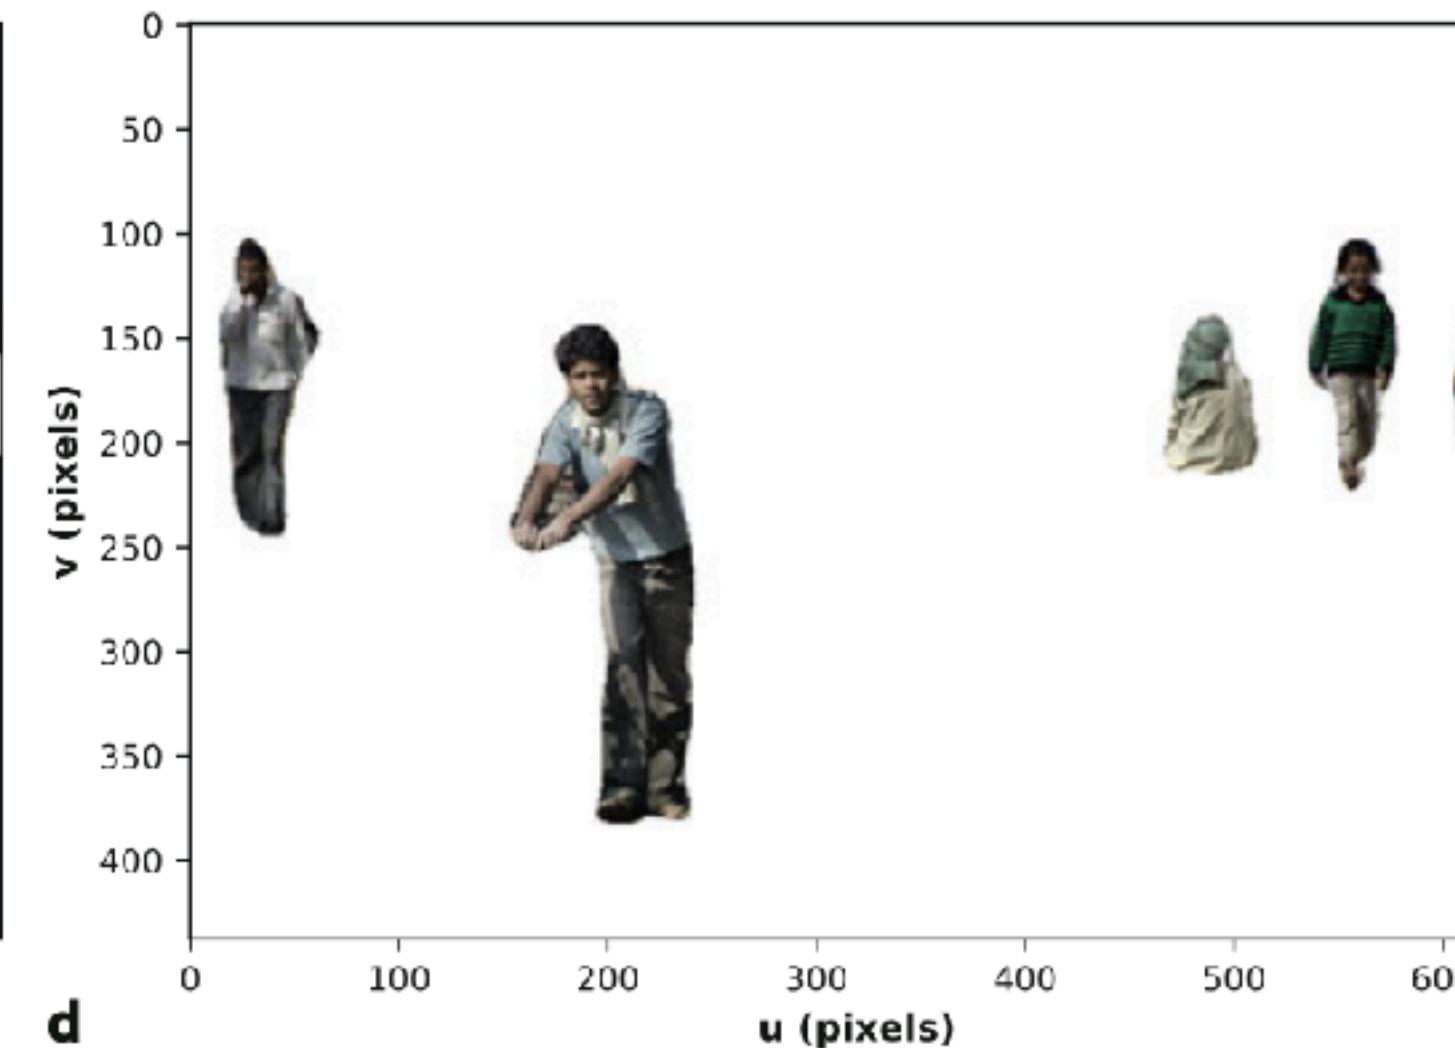
a



b

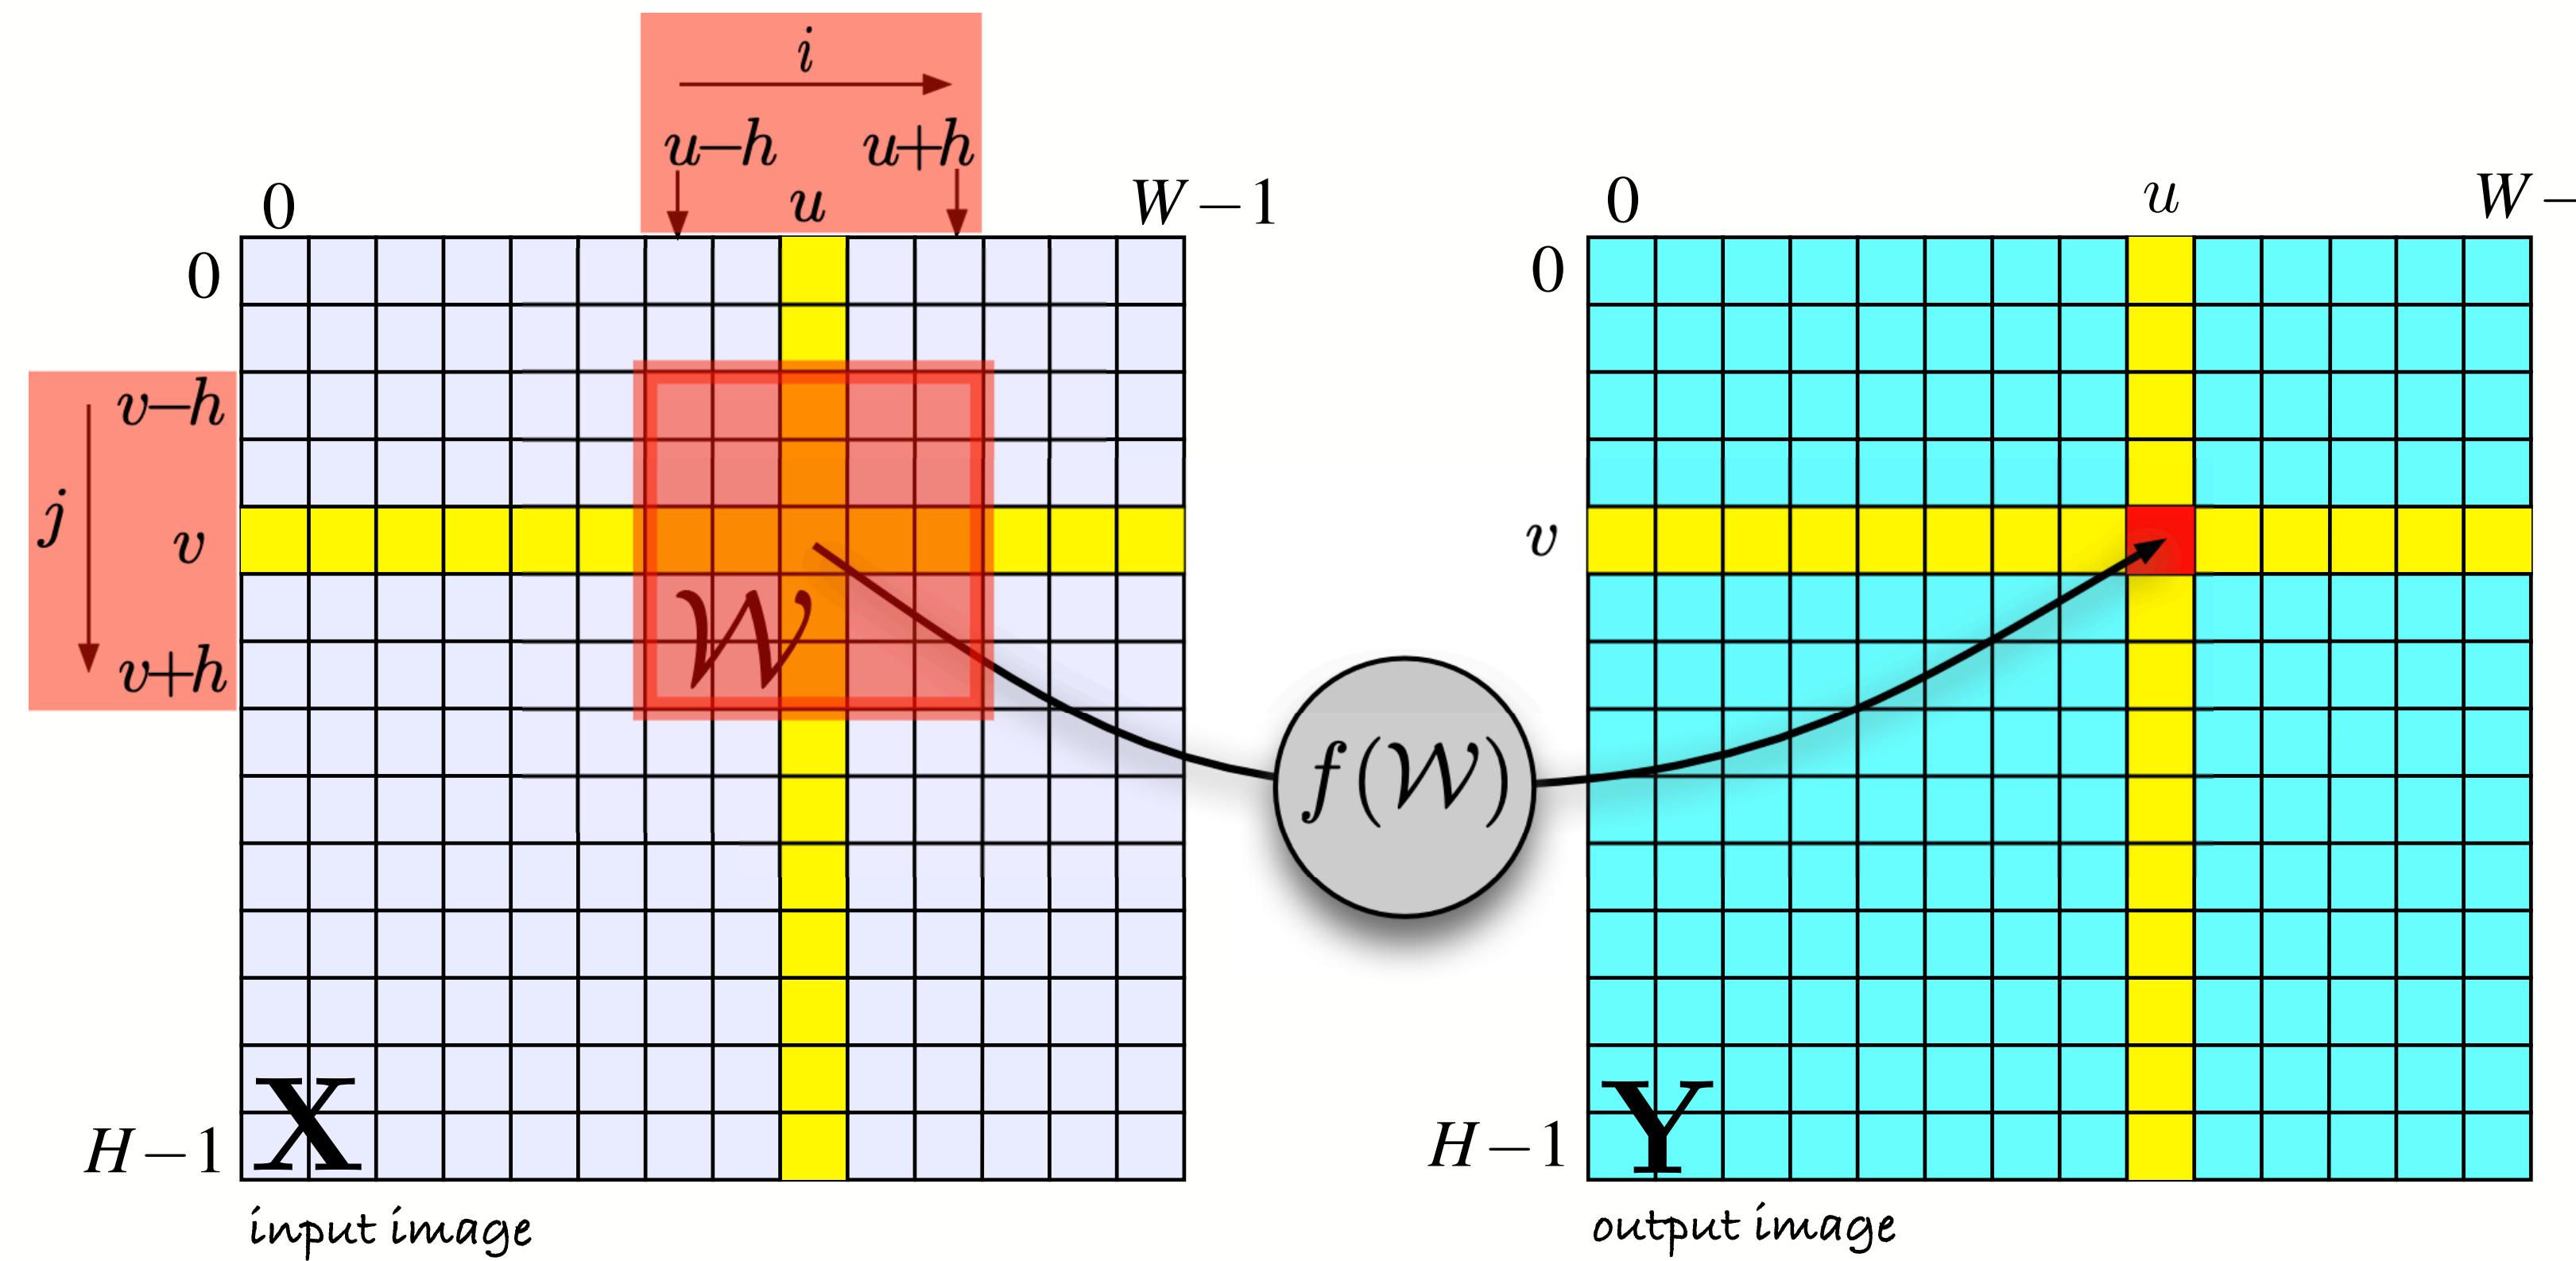


c



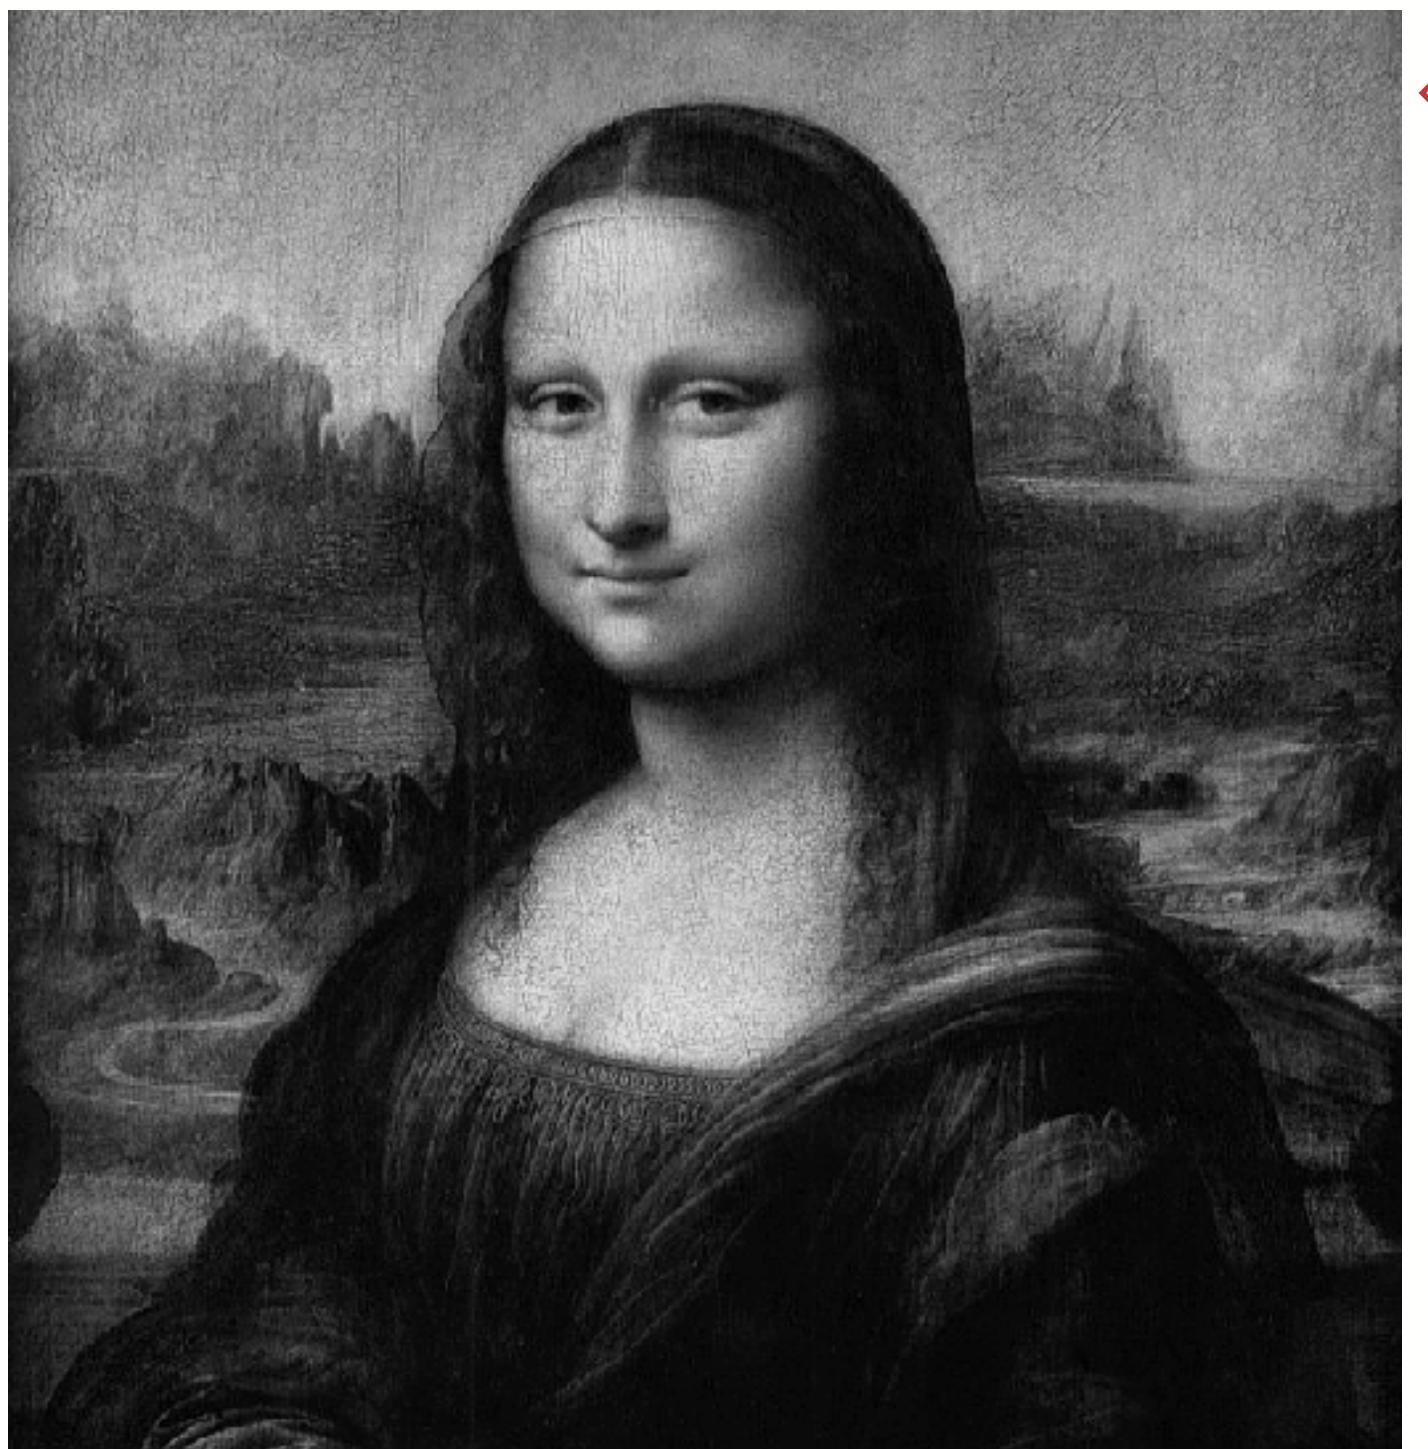
d

Spatial operators



- The function can capture something about the **uniformity** or **variation** over the local pixel **window** \mathcal{W}

$f(\cdot)$ could be an average



↙Original image

Average each
pixel over a
7x7 window⇒



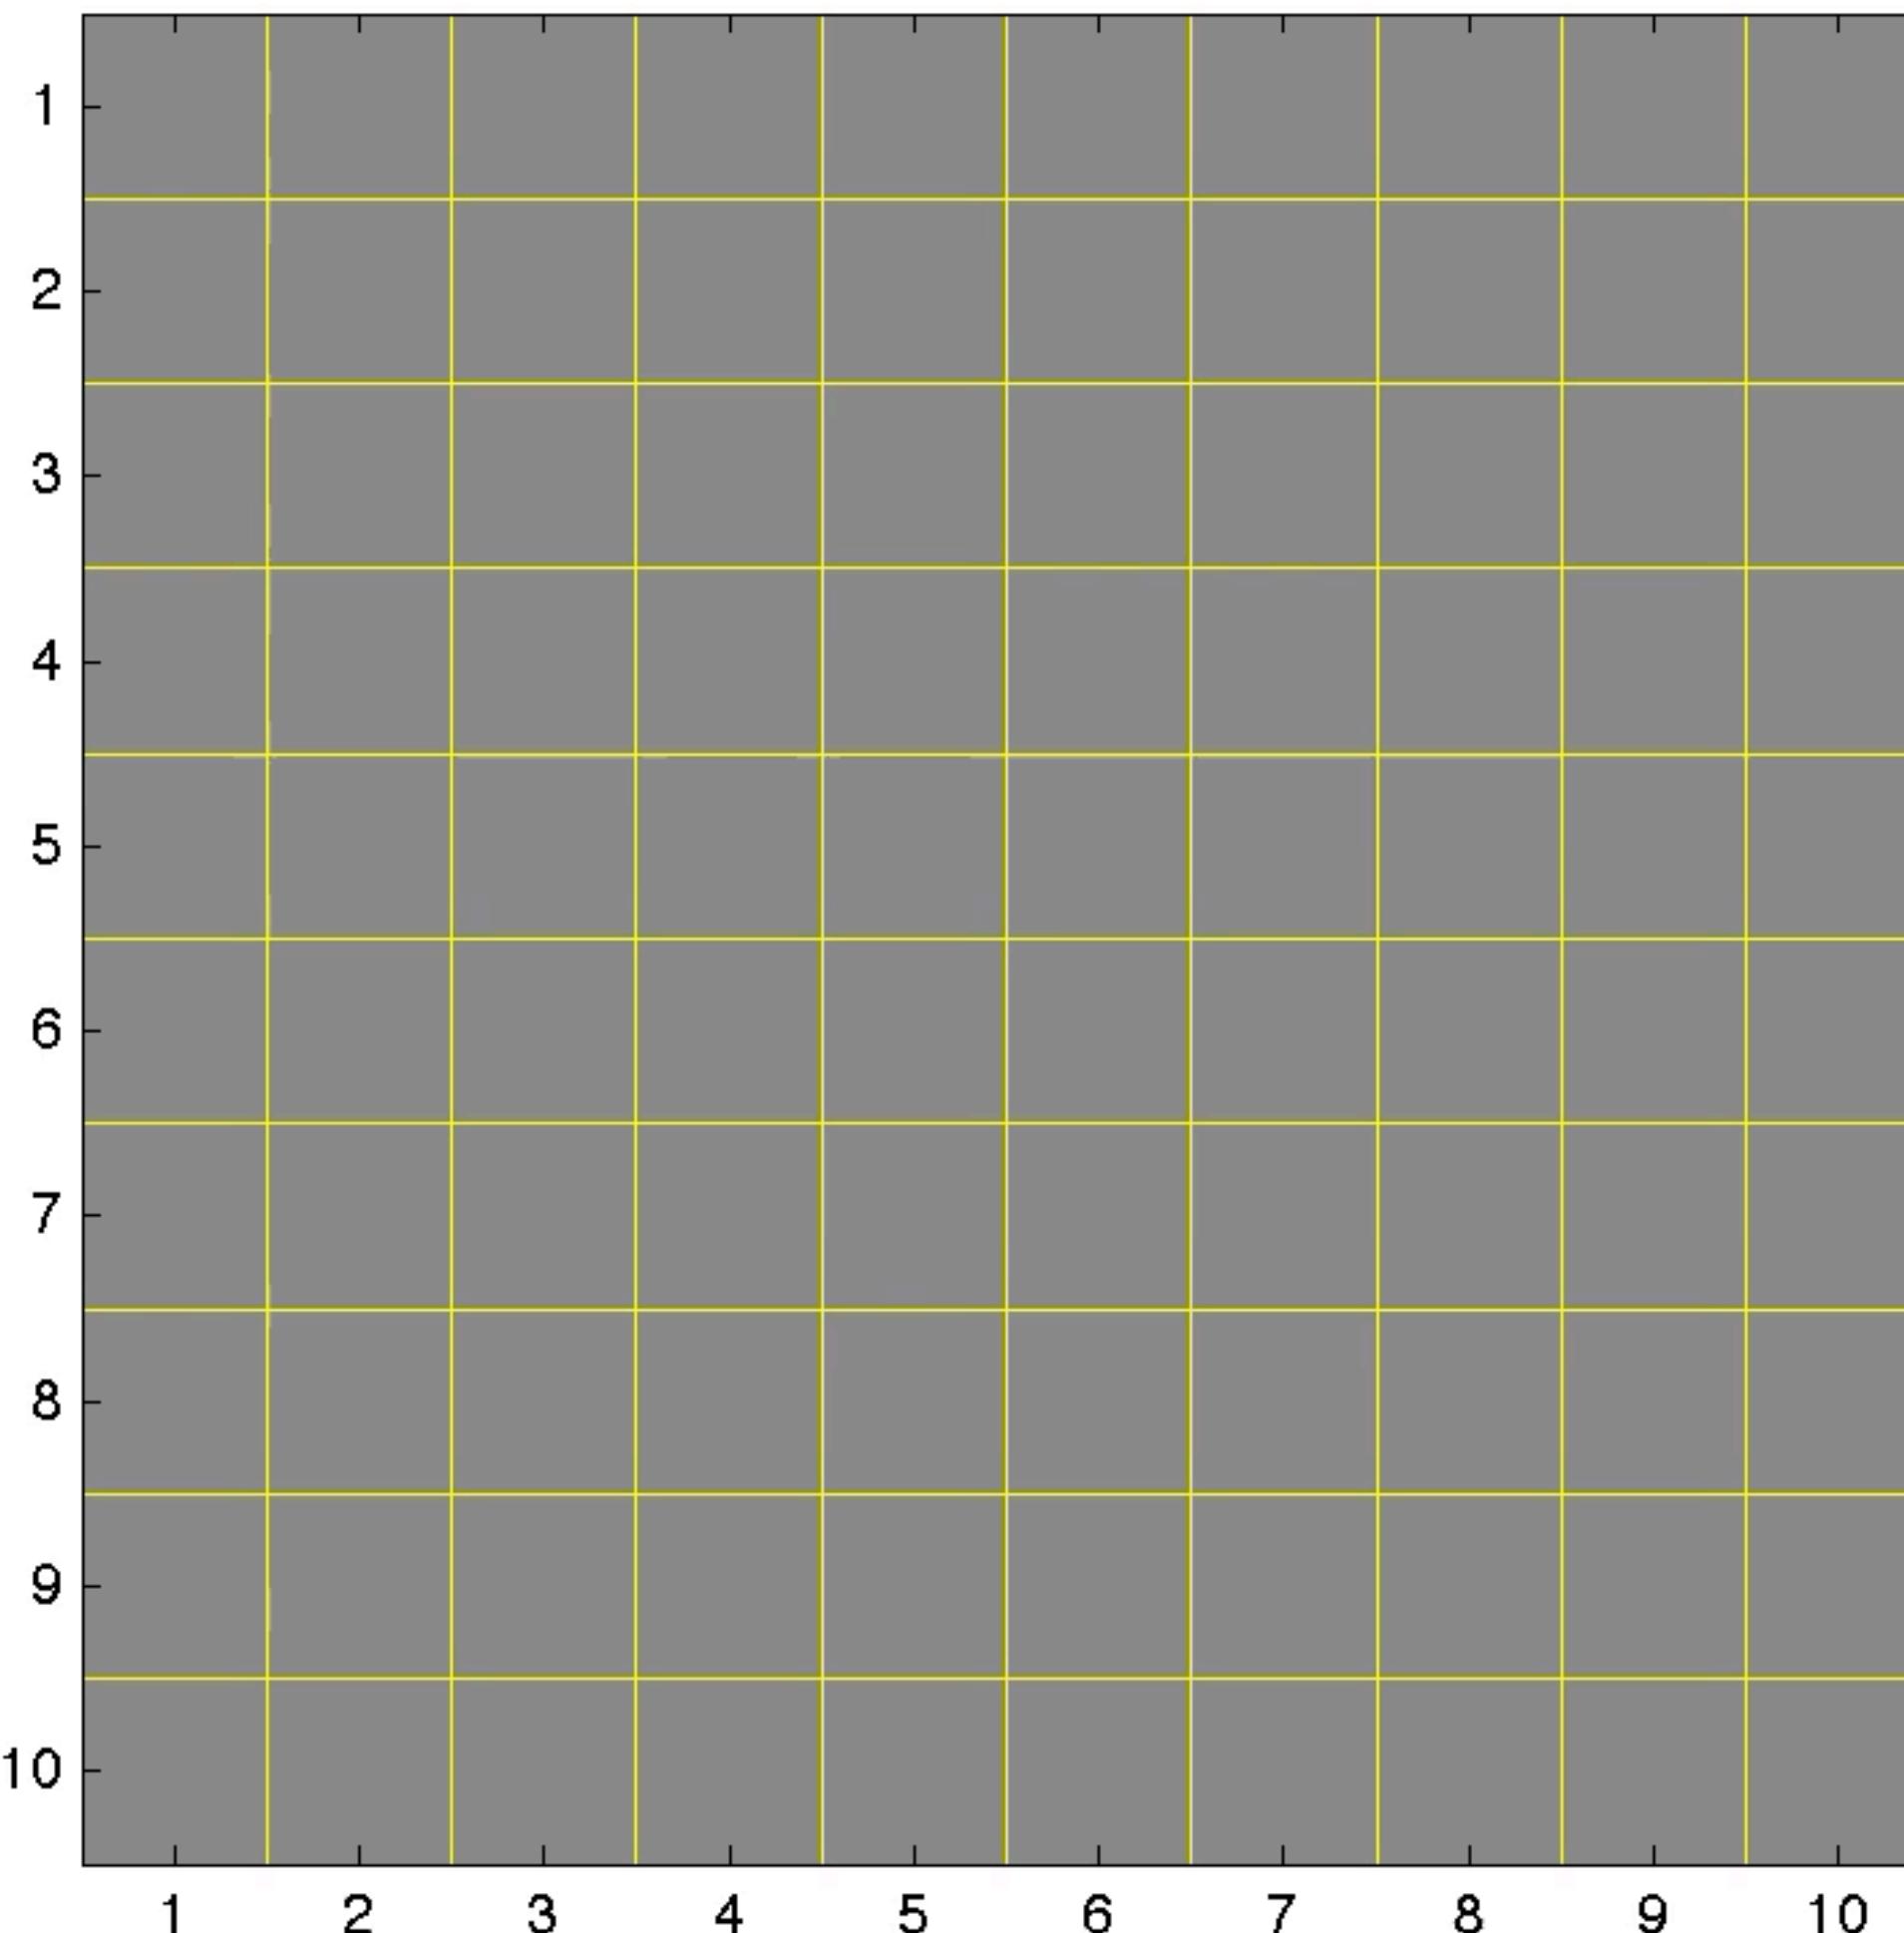
- The average over the window
 - can reduce noise
 - reduces the resolution

Moving window process

Input image

1	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30
2	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30
3	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30
4	0.30	0.30	0.30	0.80	0.80	0.80	0.30	0.30	0.30	0.30
5	0.30	0.30	0.30	0.80	0.80	0.80	0.30	0.30	0.30	0.30
6	0.30	0.30	0.30	0.80	0.80	0.80	0.30	0.30	0.30	0.30
7	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30
8	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30
9	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30
10	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30	0.30

Output image



Effect of window size



Original image



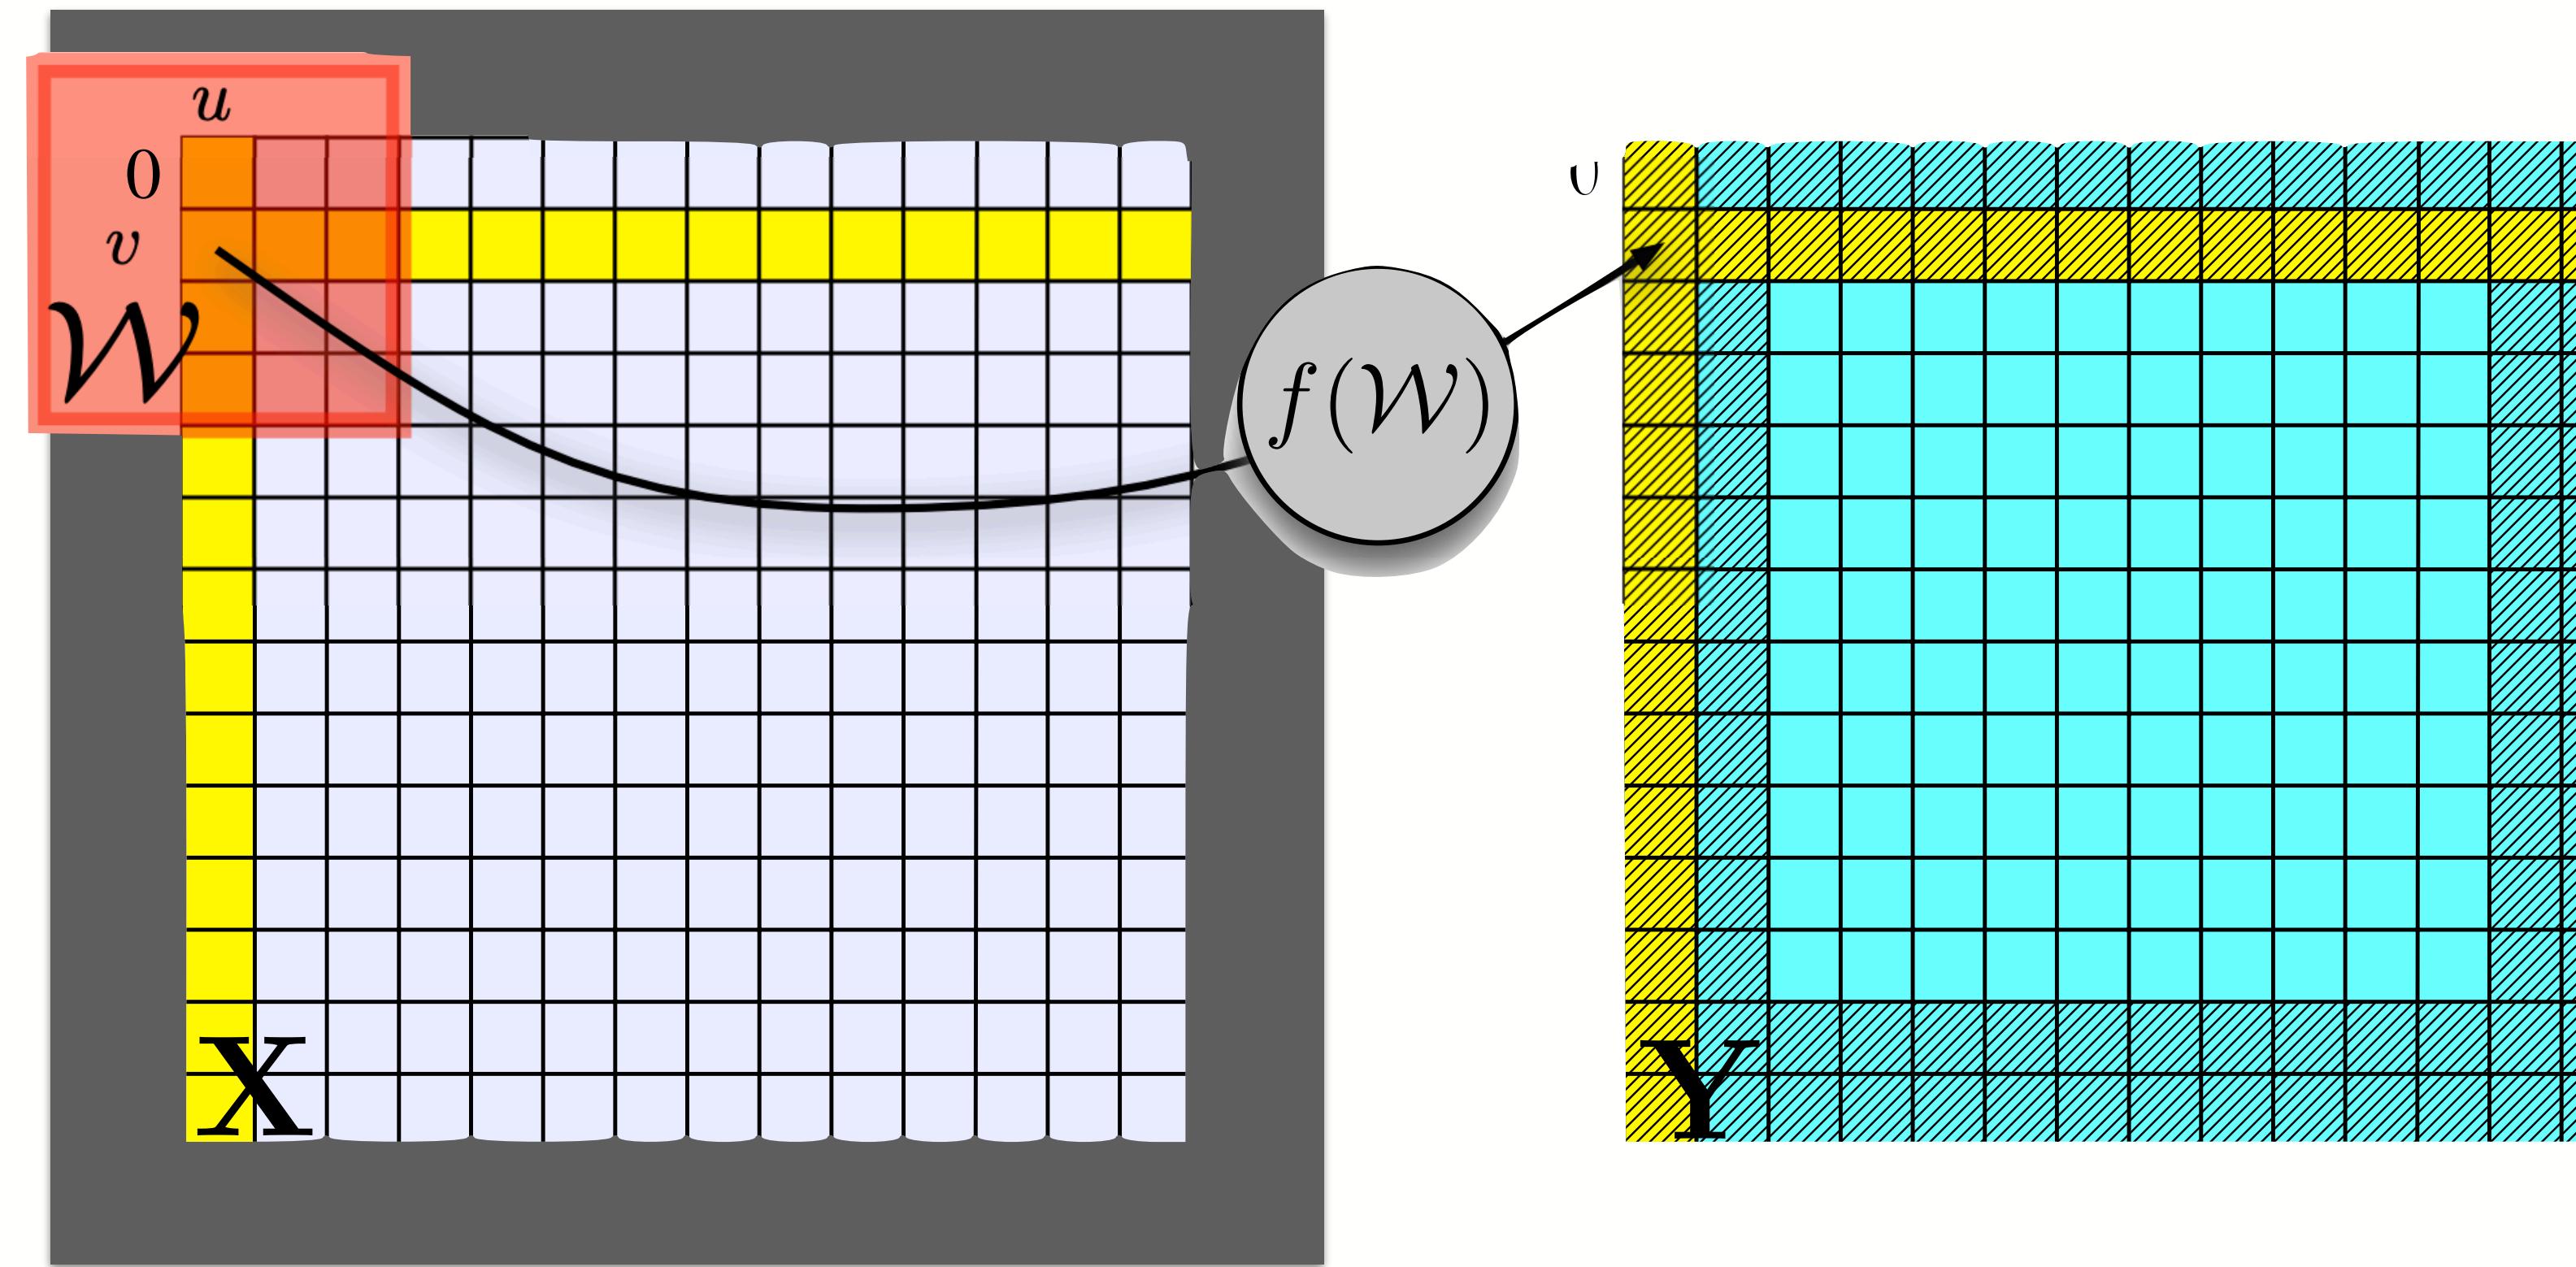
Average each pixel over a
7x7 window



Average each pixel over a
21x 21 window

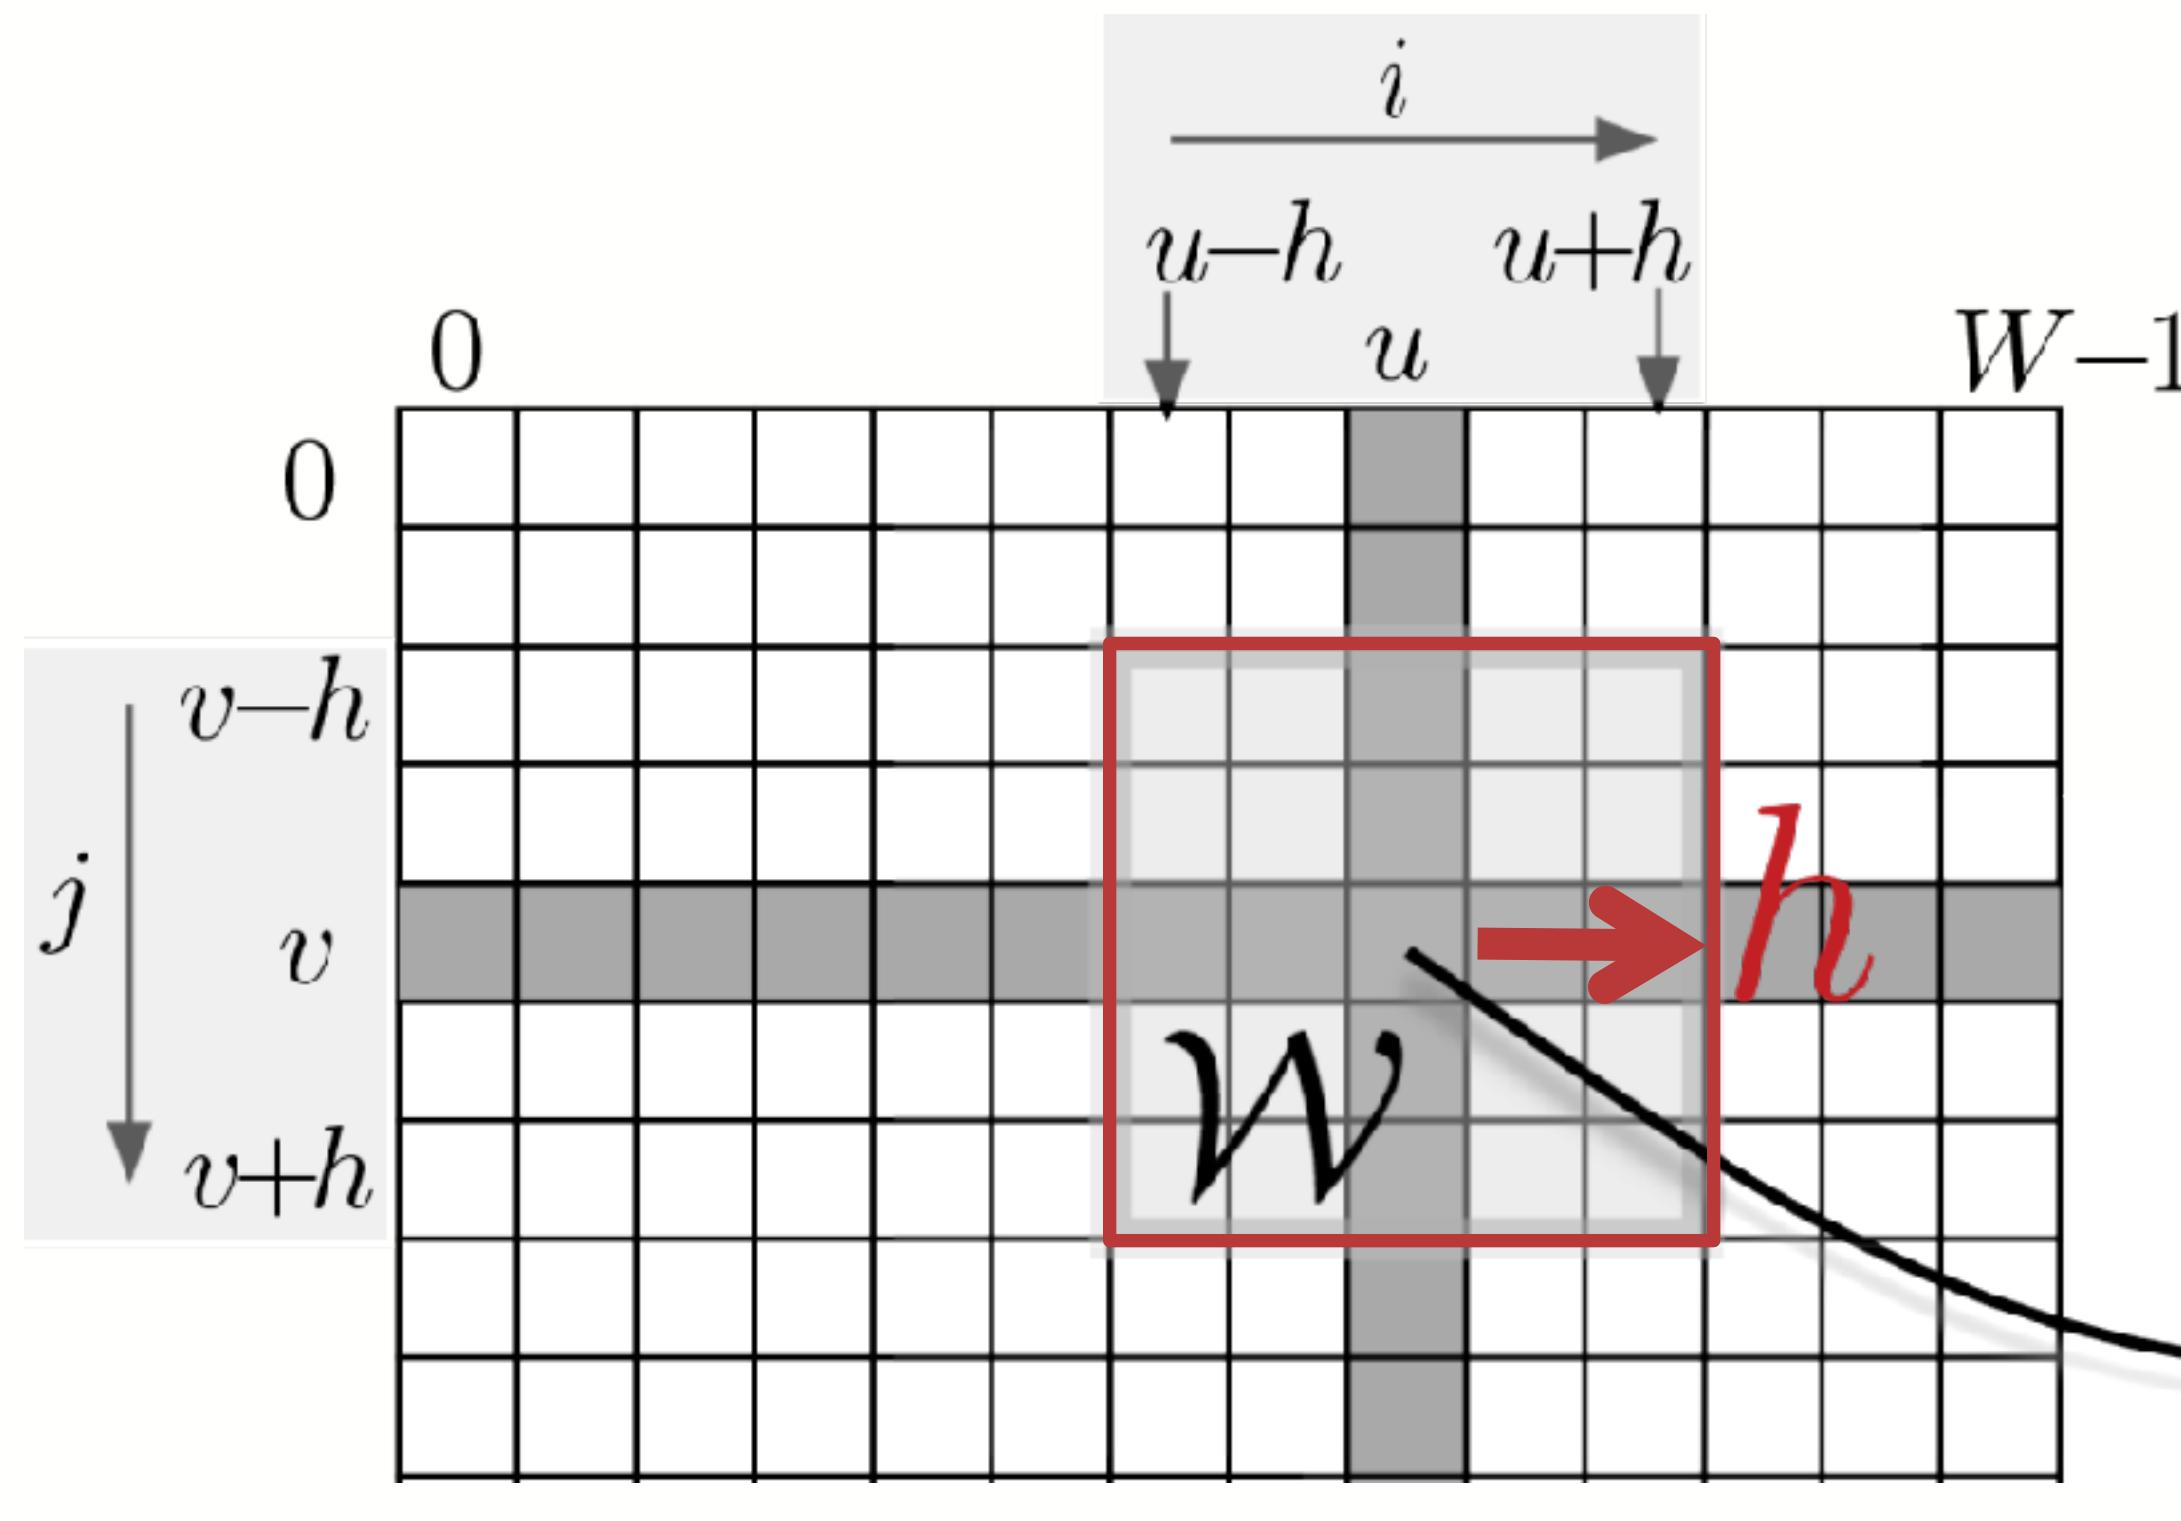
- Features smaller than the window size will be strongly attenuated
 - we say the image is blurry, fuzzy, lower resolution etc.

The edge problem



- Some solutions:
 - don't compute the output value when the window "falls off the edge"
 - assume the image is surrounded by zero pixels
 - assume the edge pixels are replicated outward

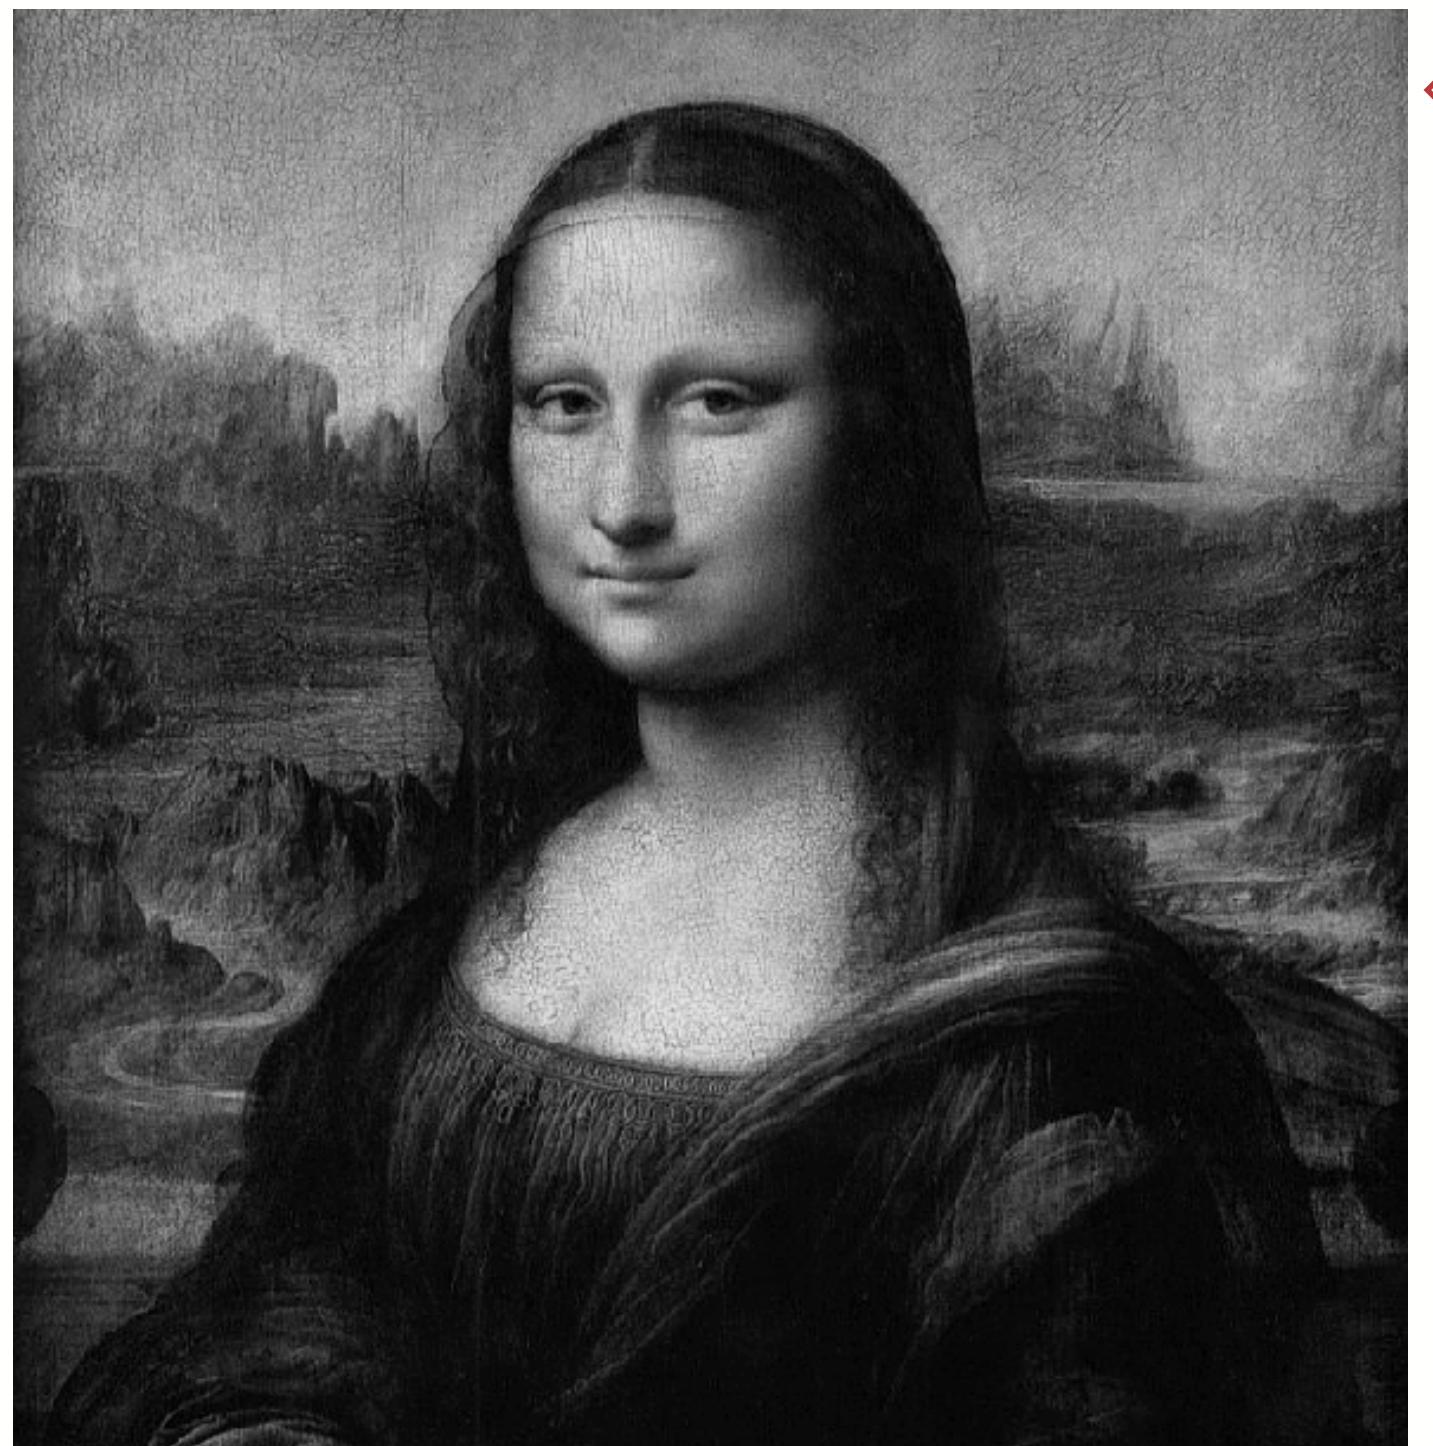
Window size is always odd



- The window is:
 - square
 - always centred on the input pixel
 - edge is integer h pixels from the centre
- The window width is $2h + 1$
 - always odd

With kind permission of Springer Science+Business Media

Artefacts of averaging



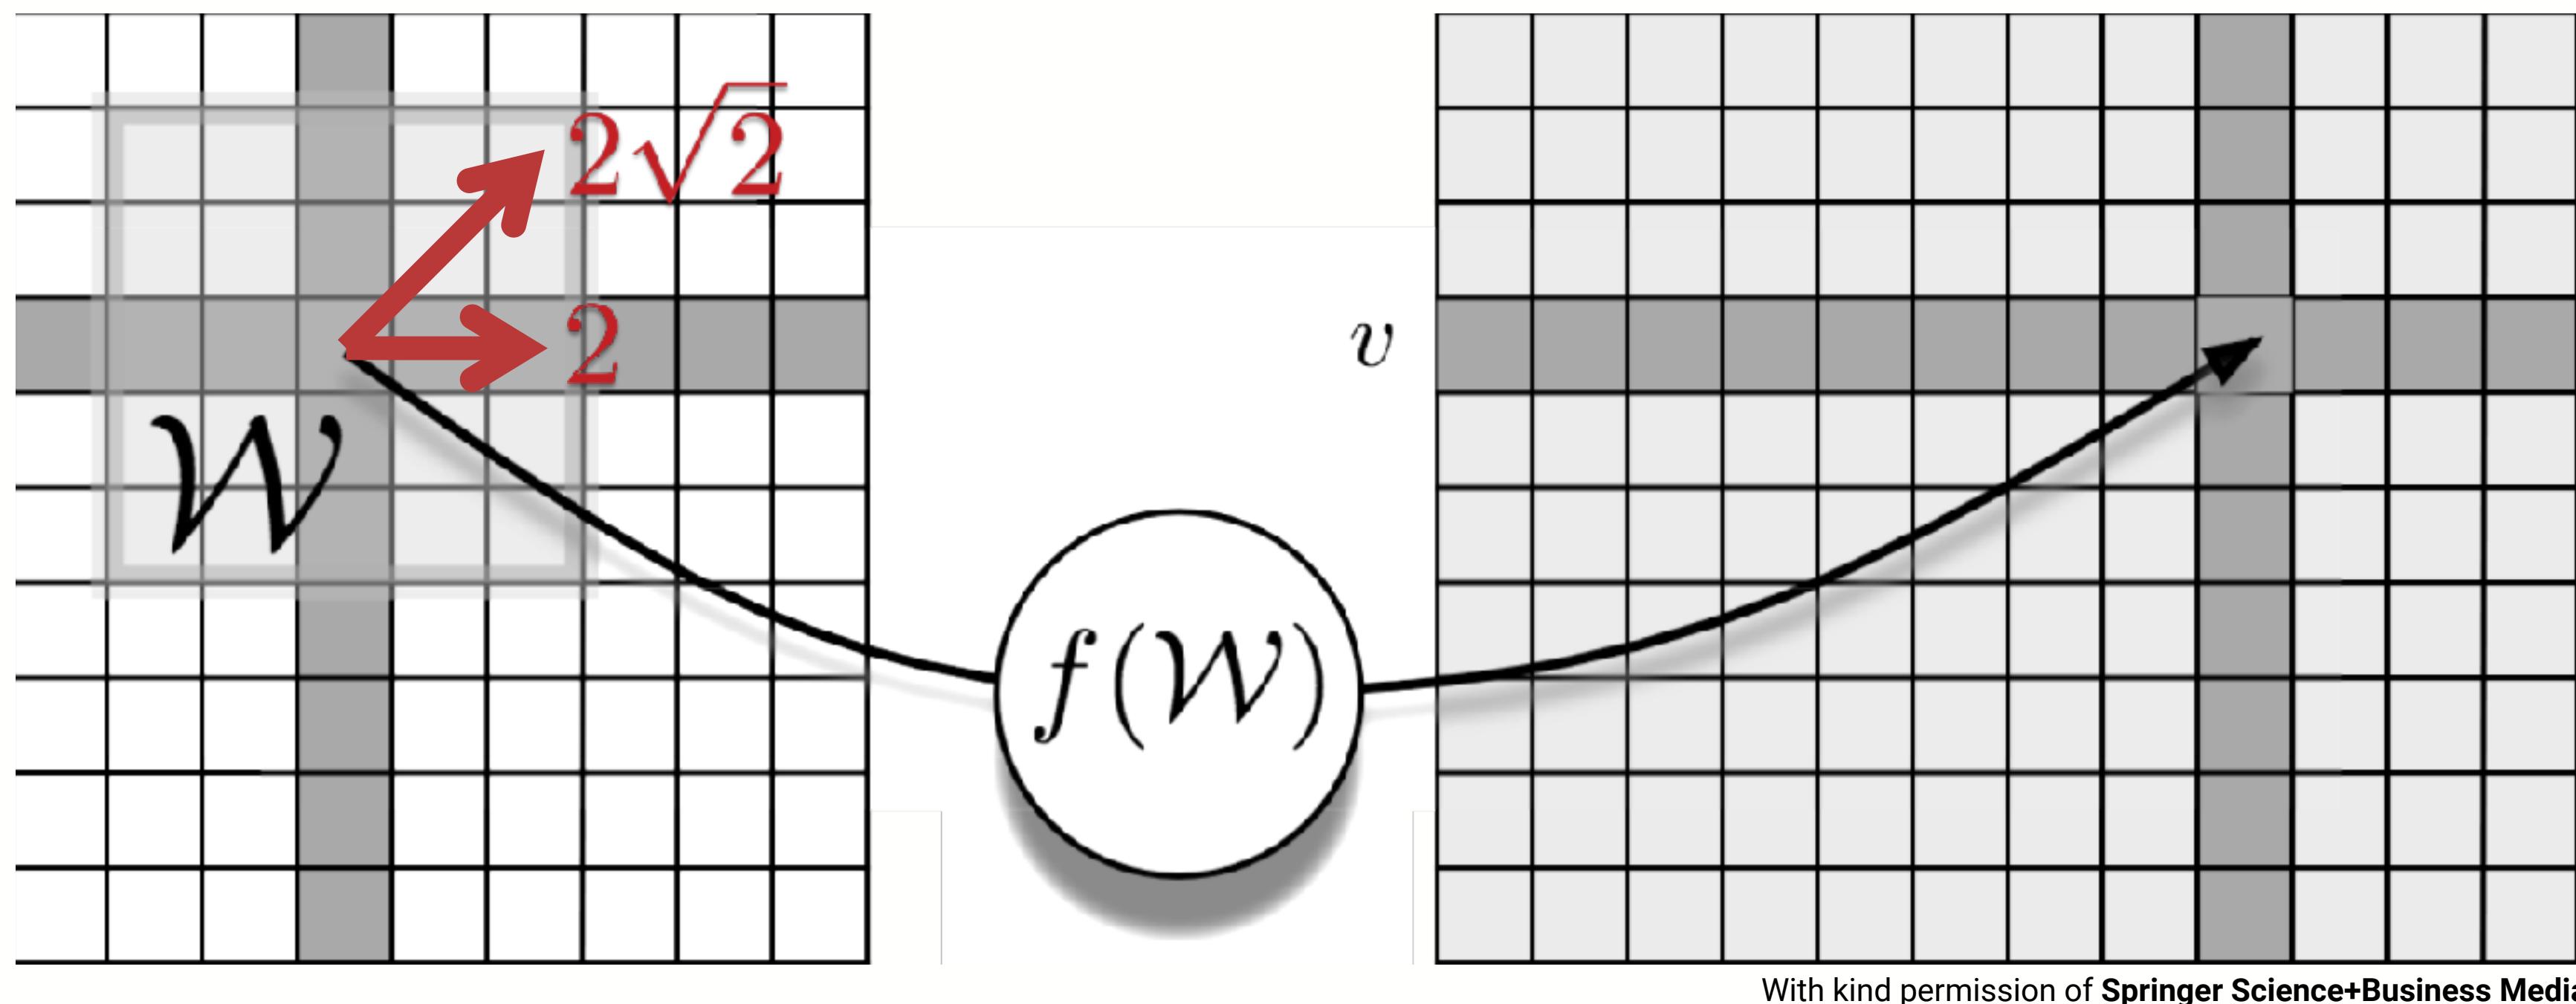
⬅Original image

Average each
pixel over a
 21×21 window ➡



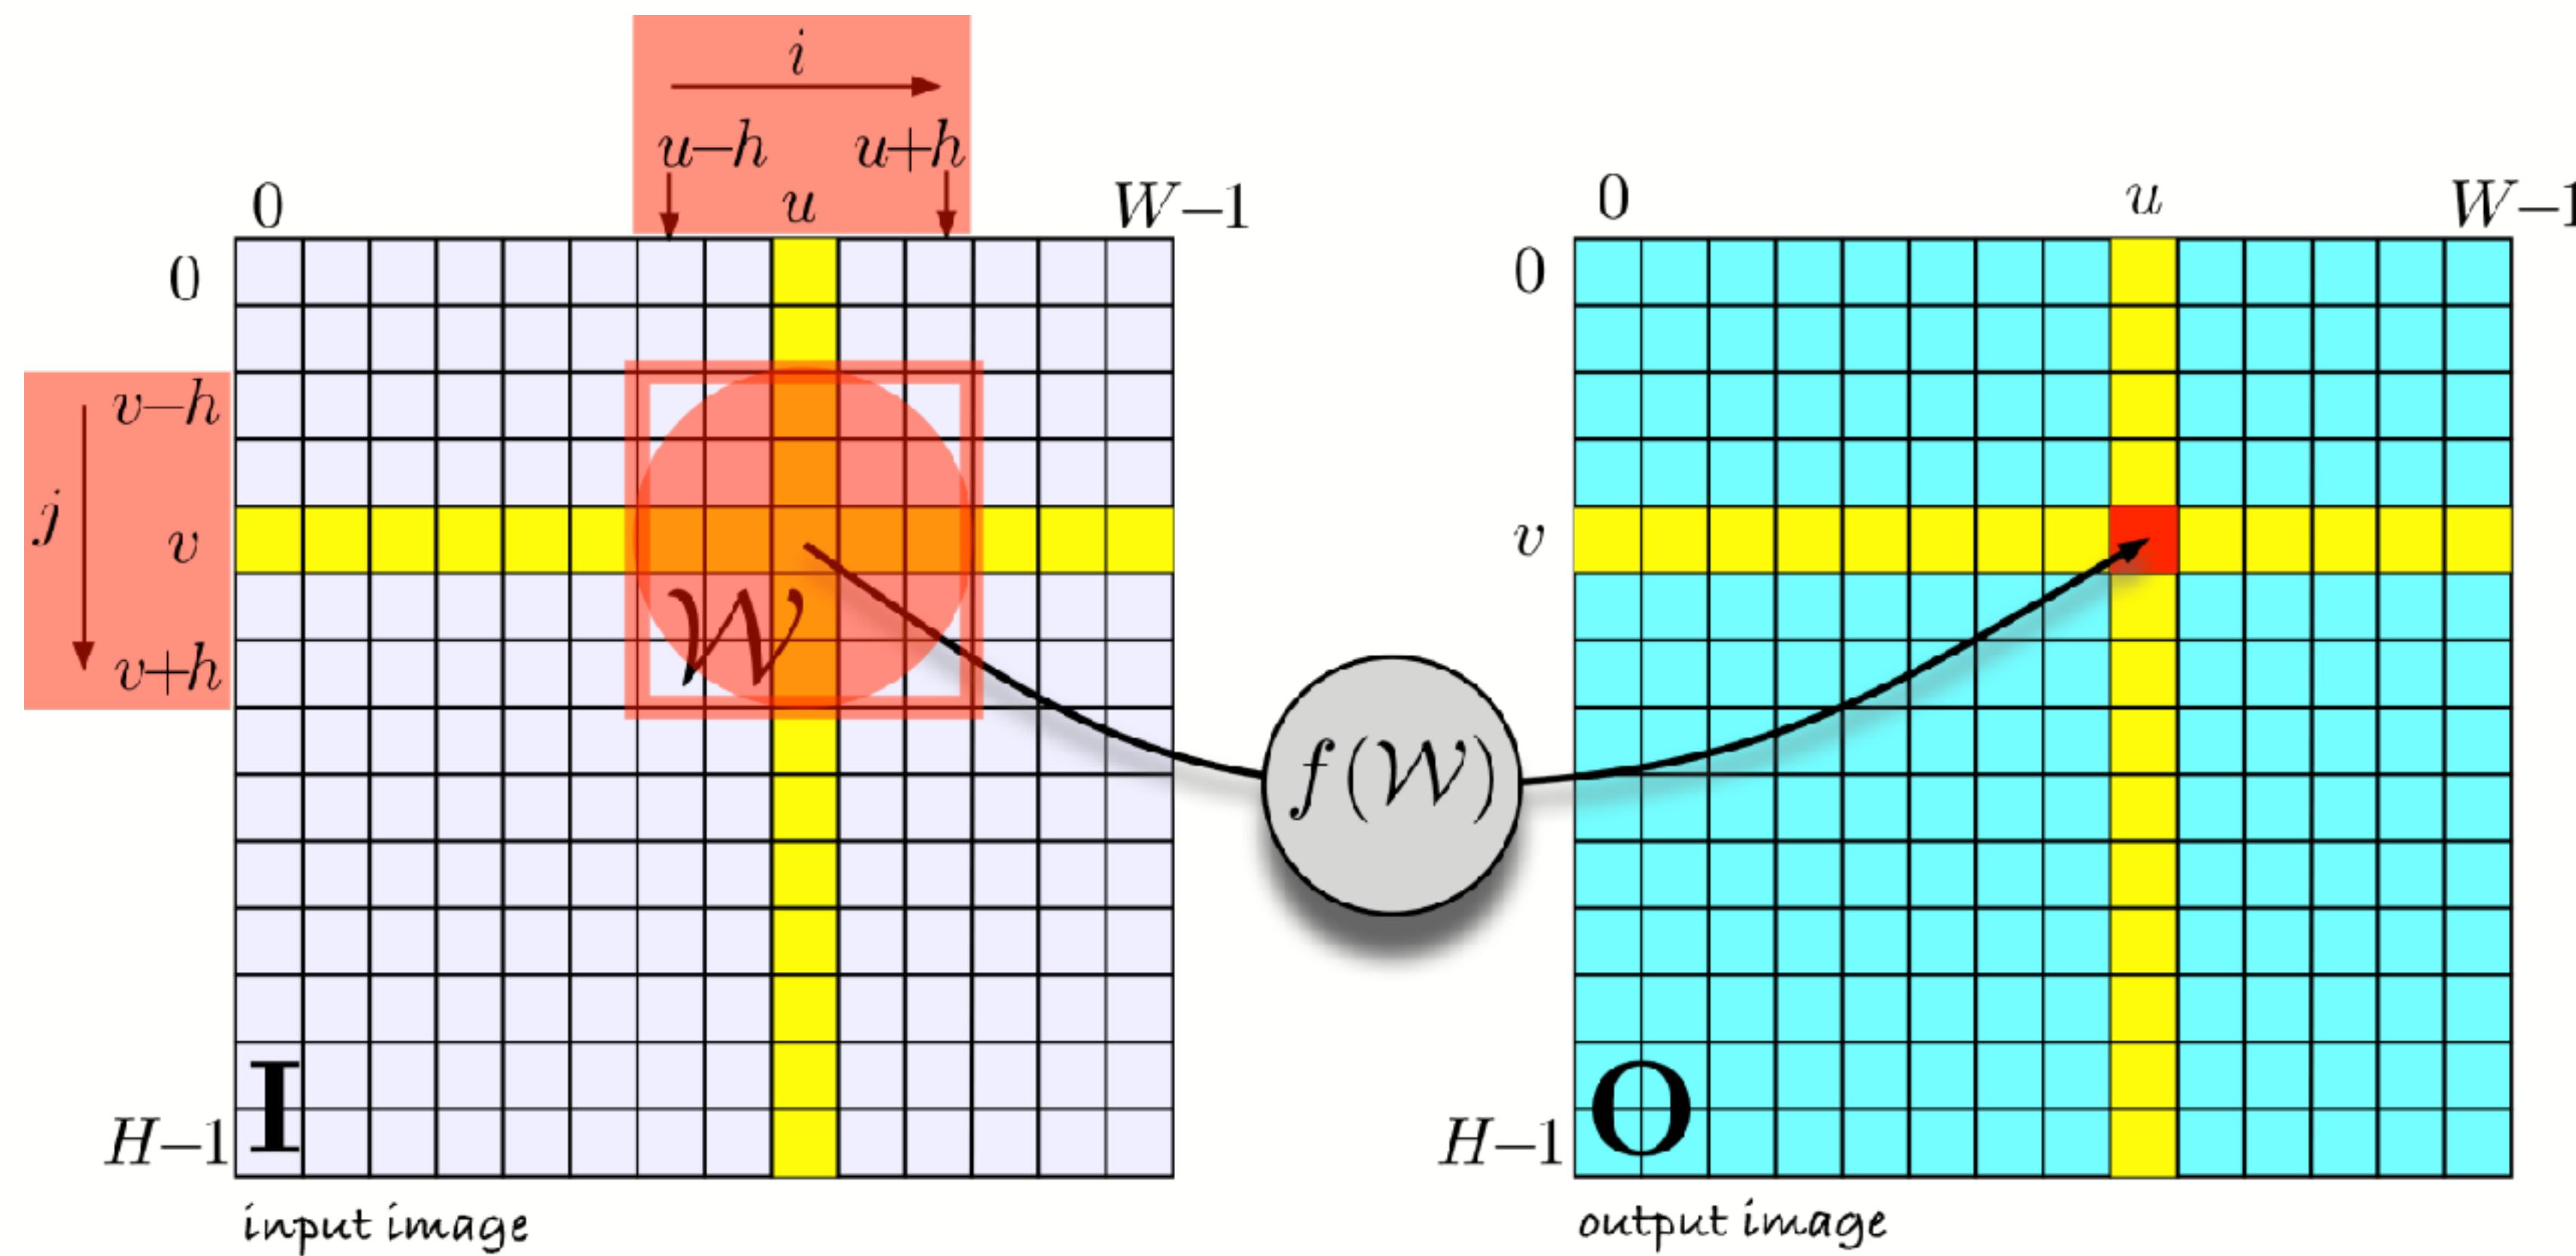
- Can lead to *ringing*
 - faint vertical & horizontal lines are introduced

Averaging over a square is not isotropic



- Not all values used in the average are the same distance away
 - Undue influence by distant values

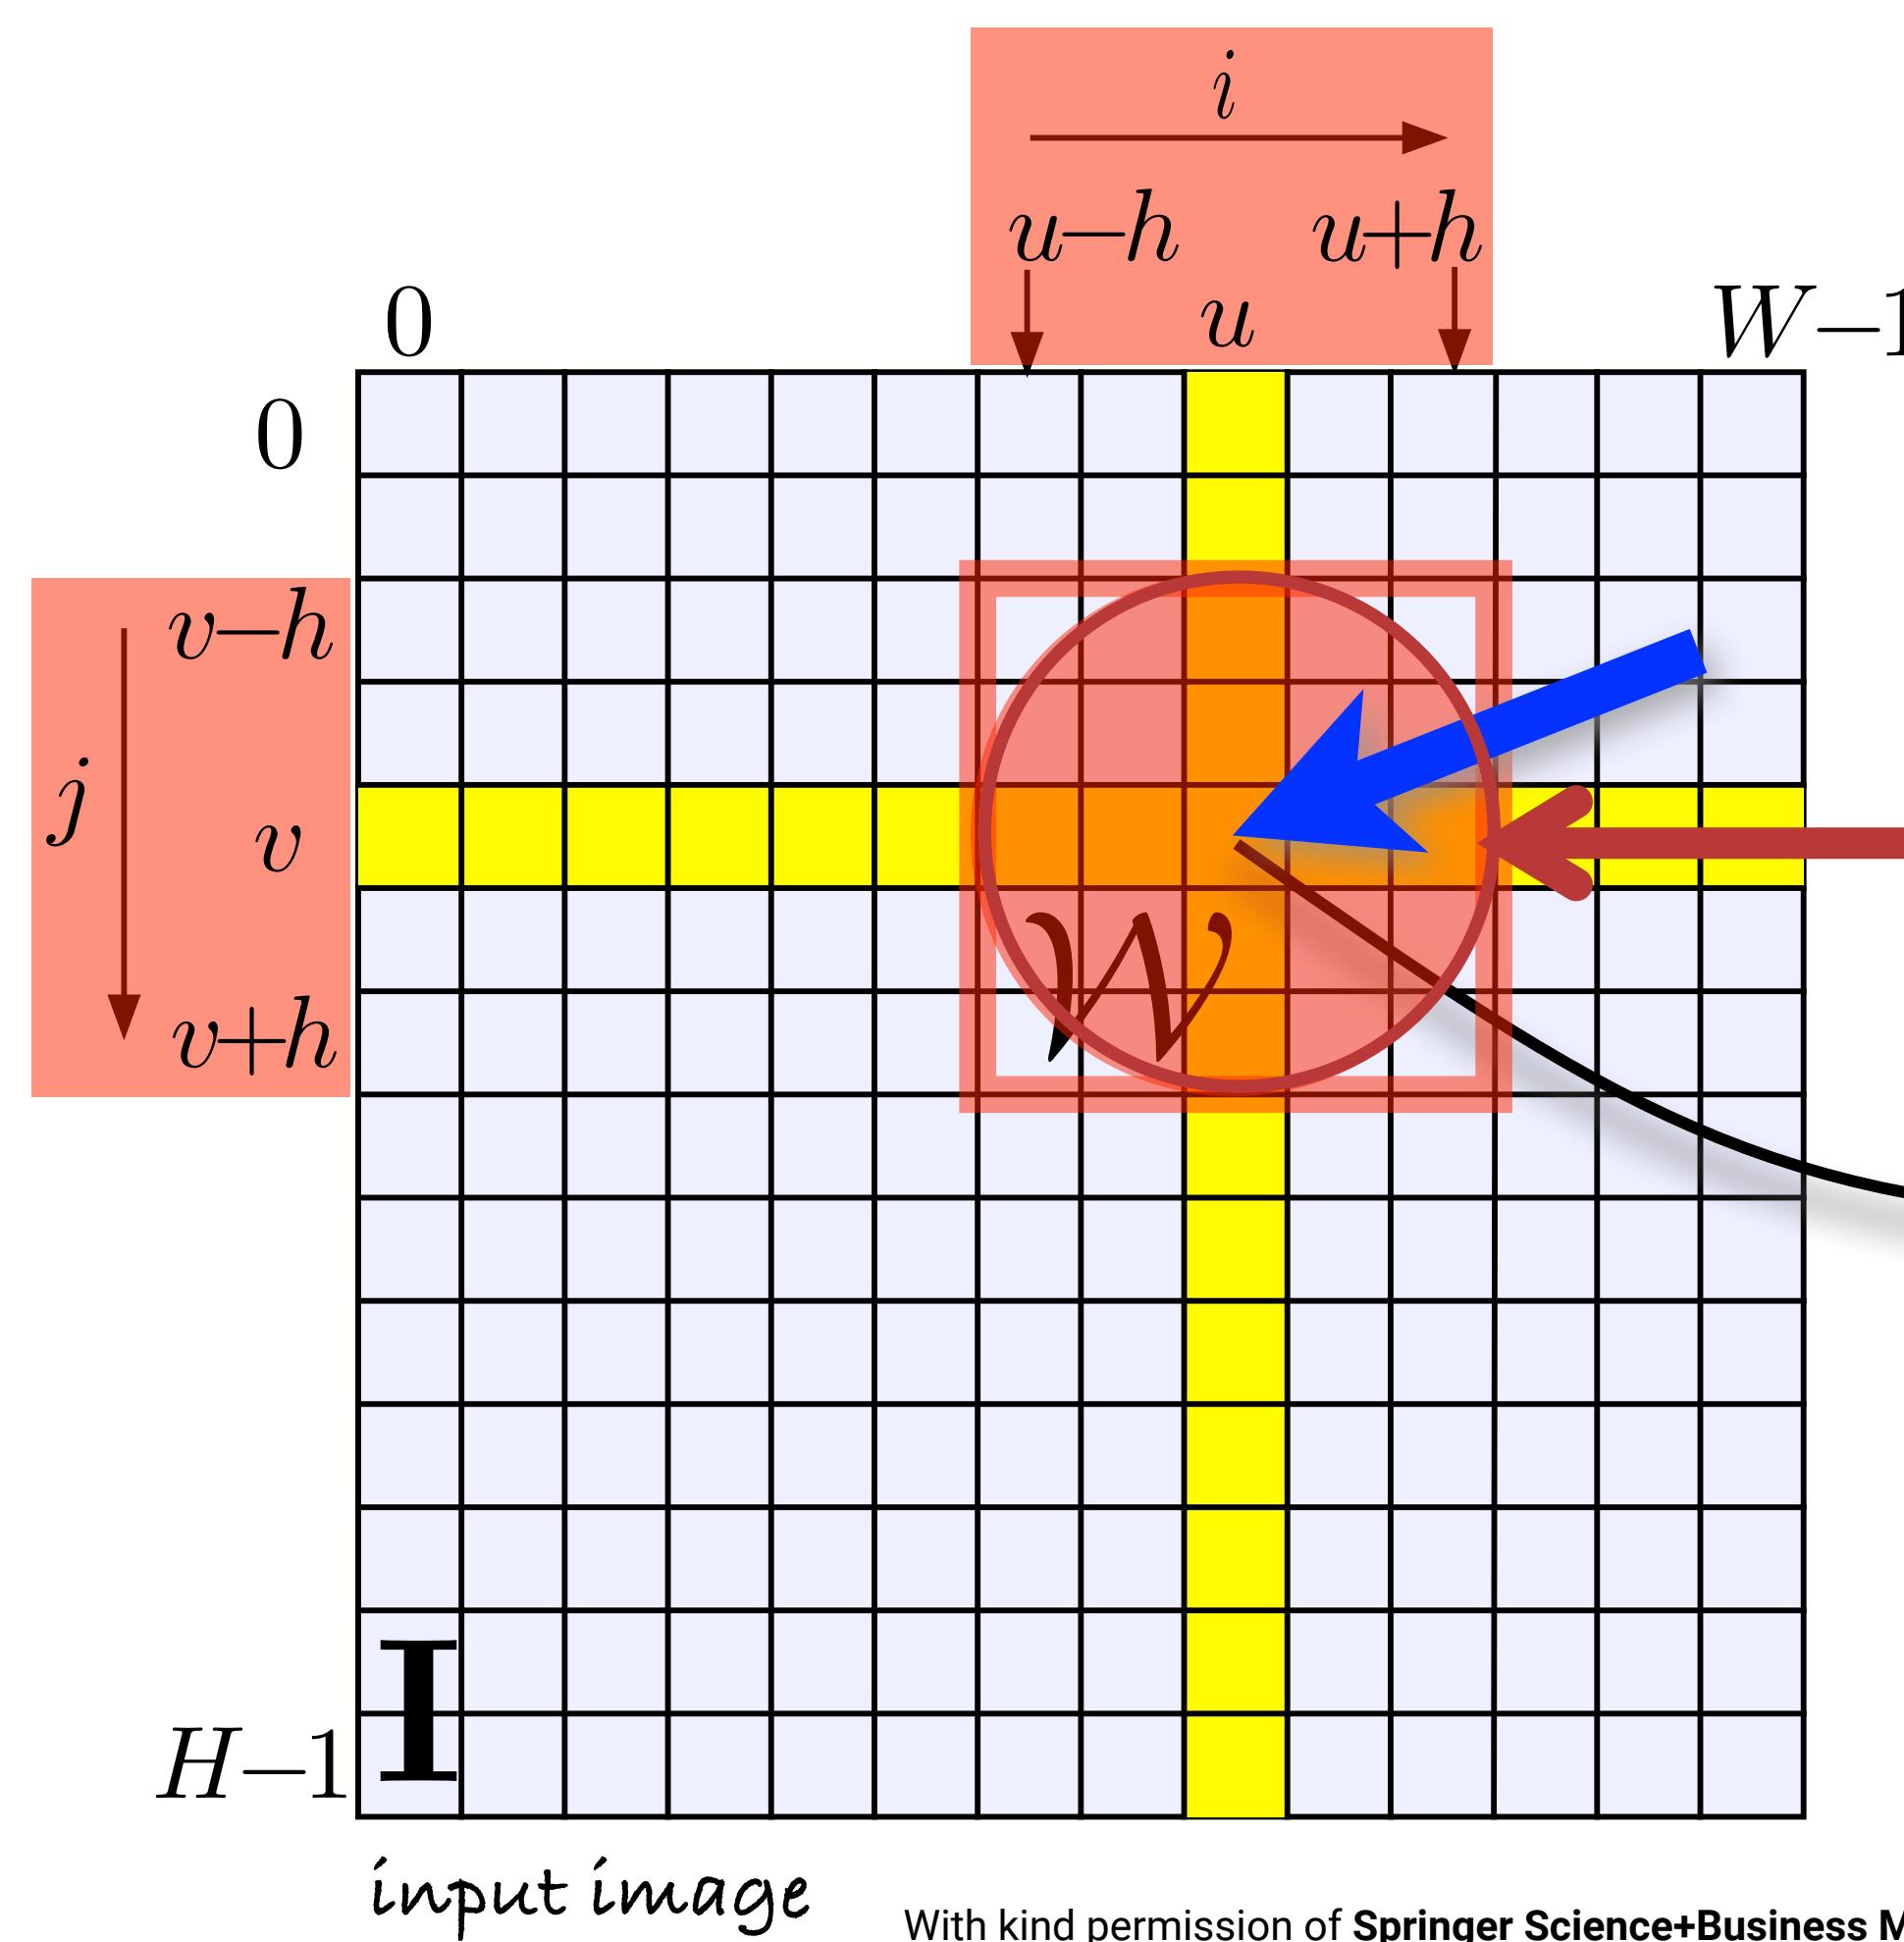
Going isotropic



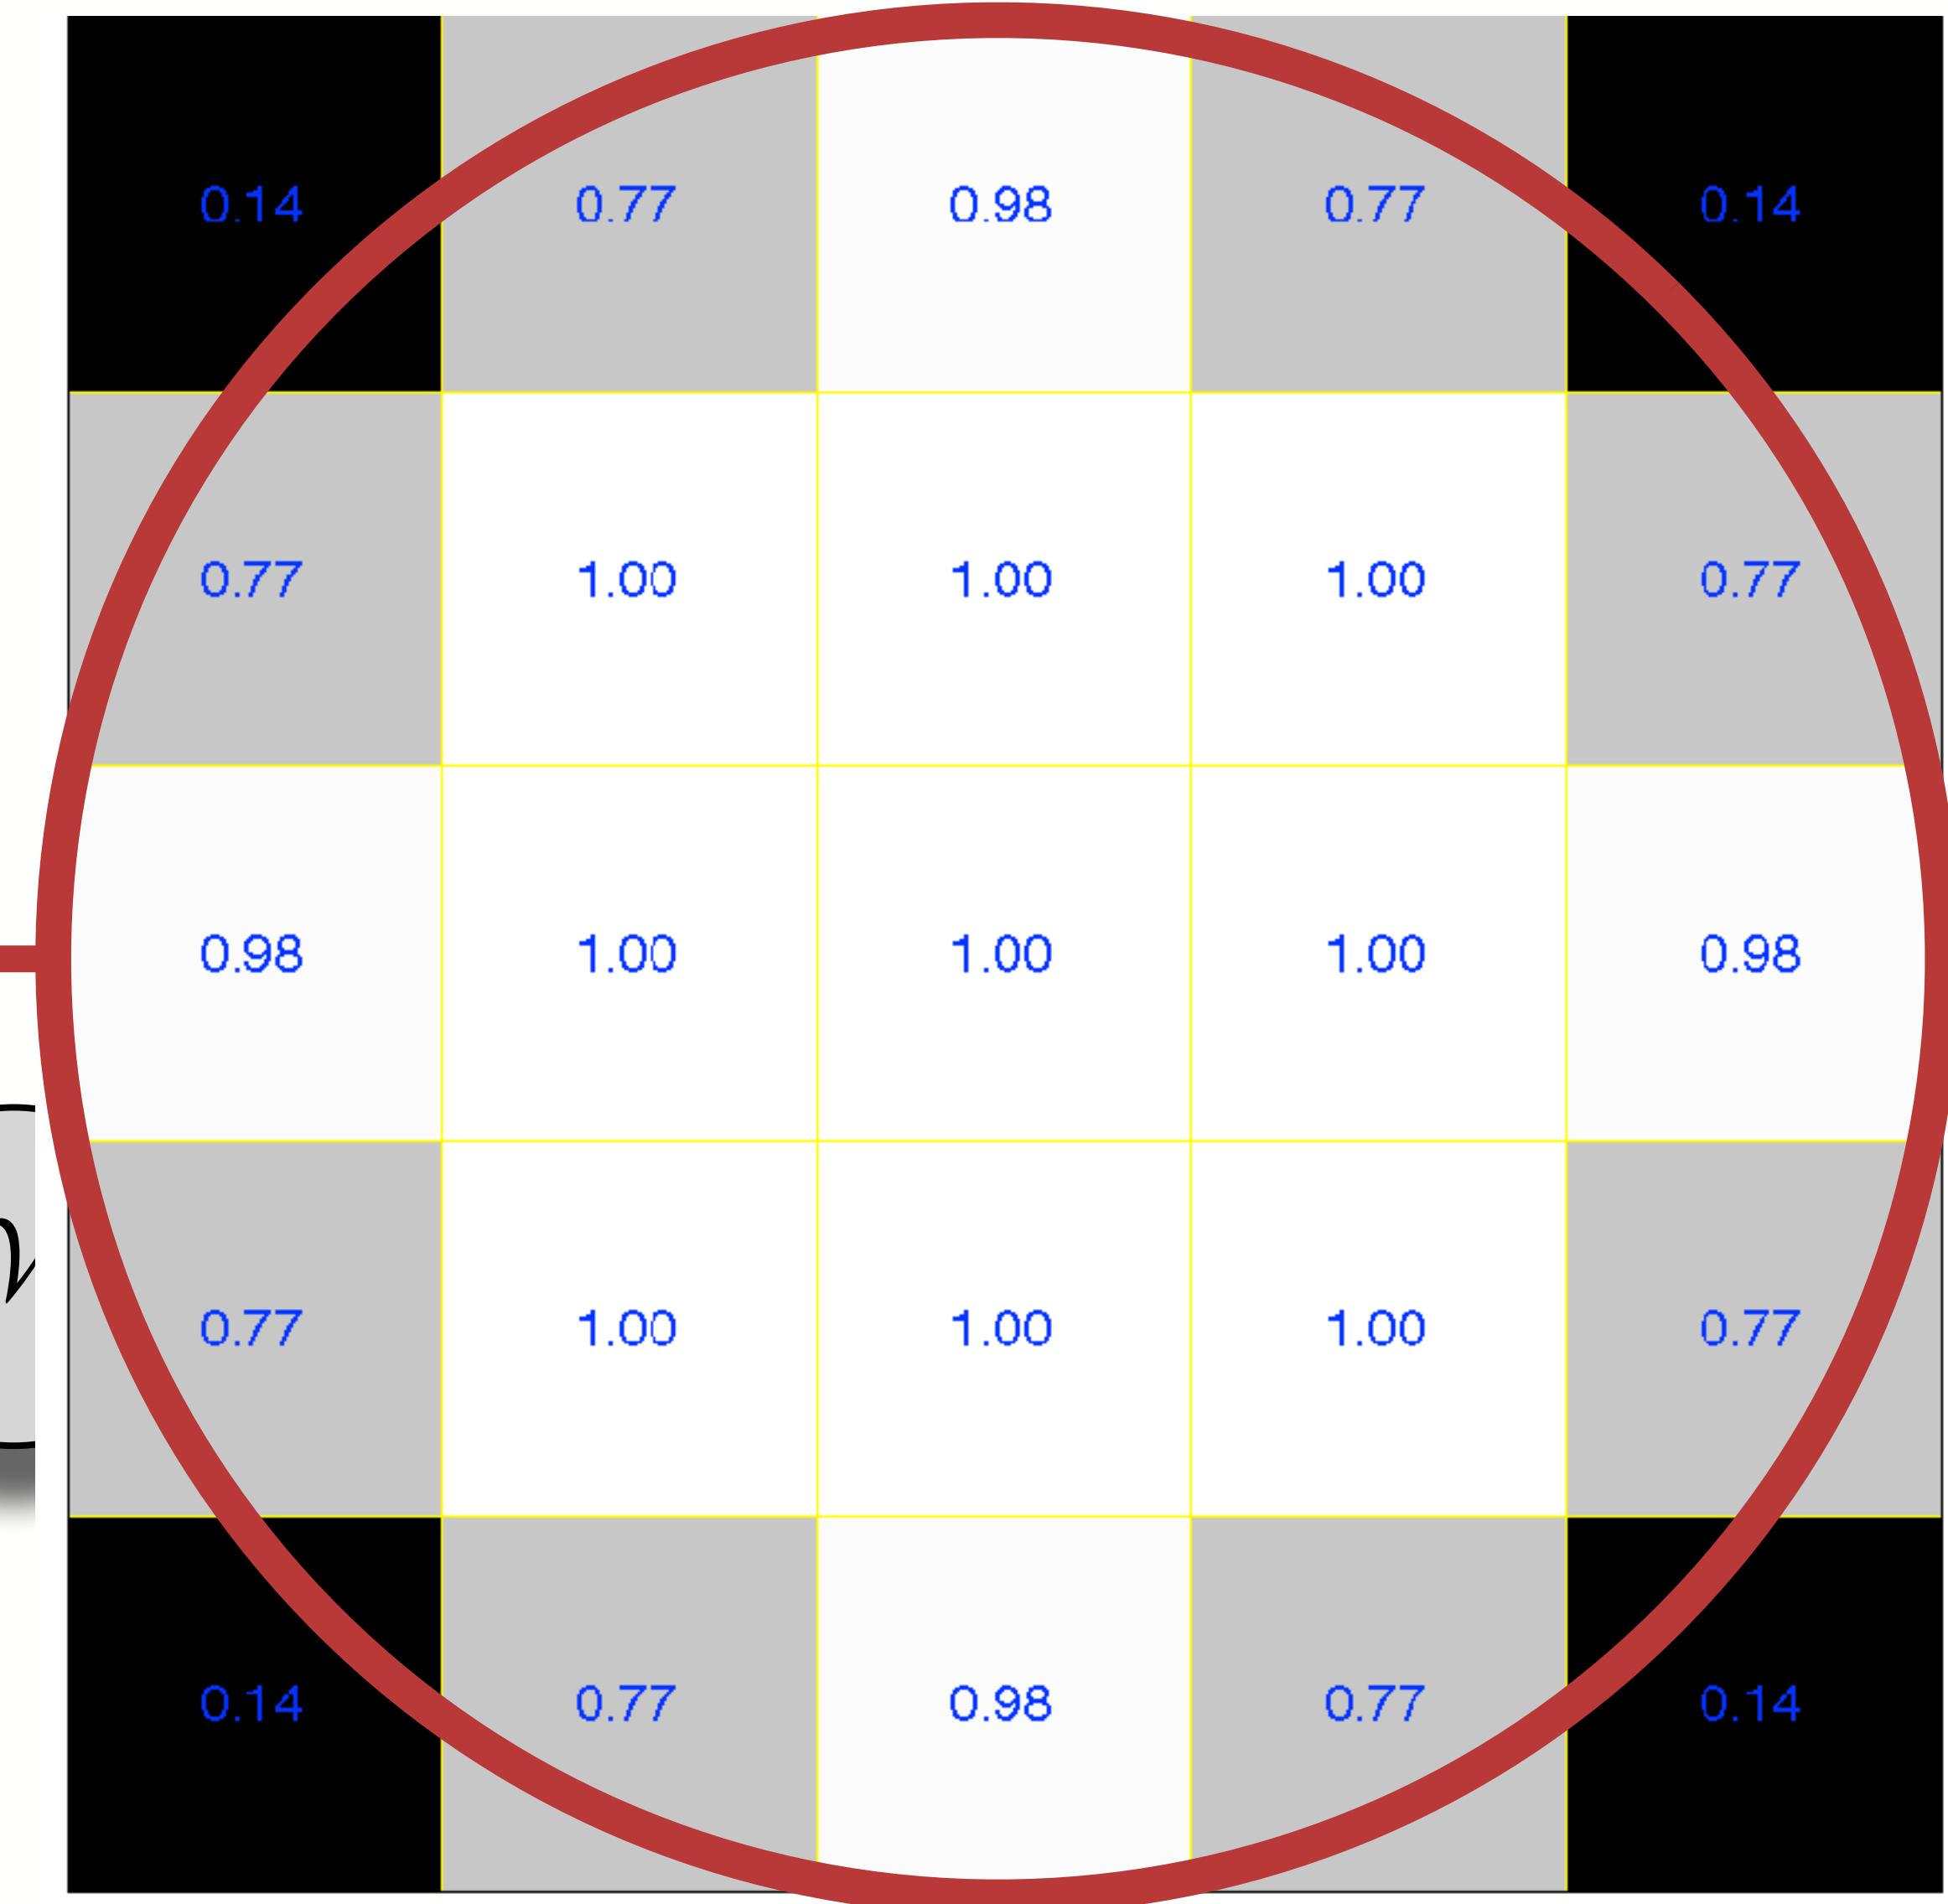
With kind permission of Springer Science+Business Media

- Ideally we'd like to extract a circular region
 - but that would involve taking fractions of pixels

Apply a weighting

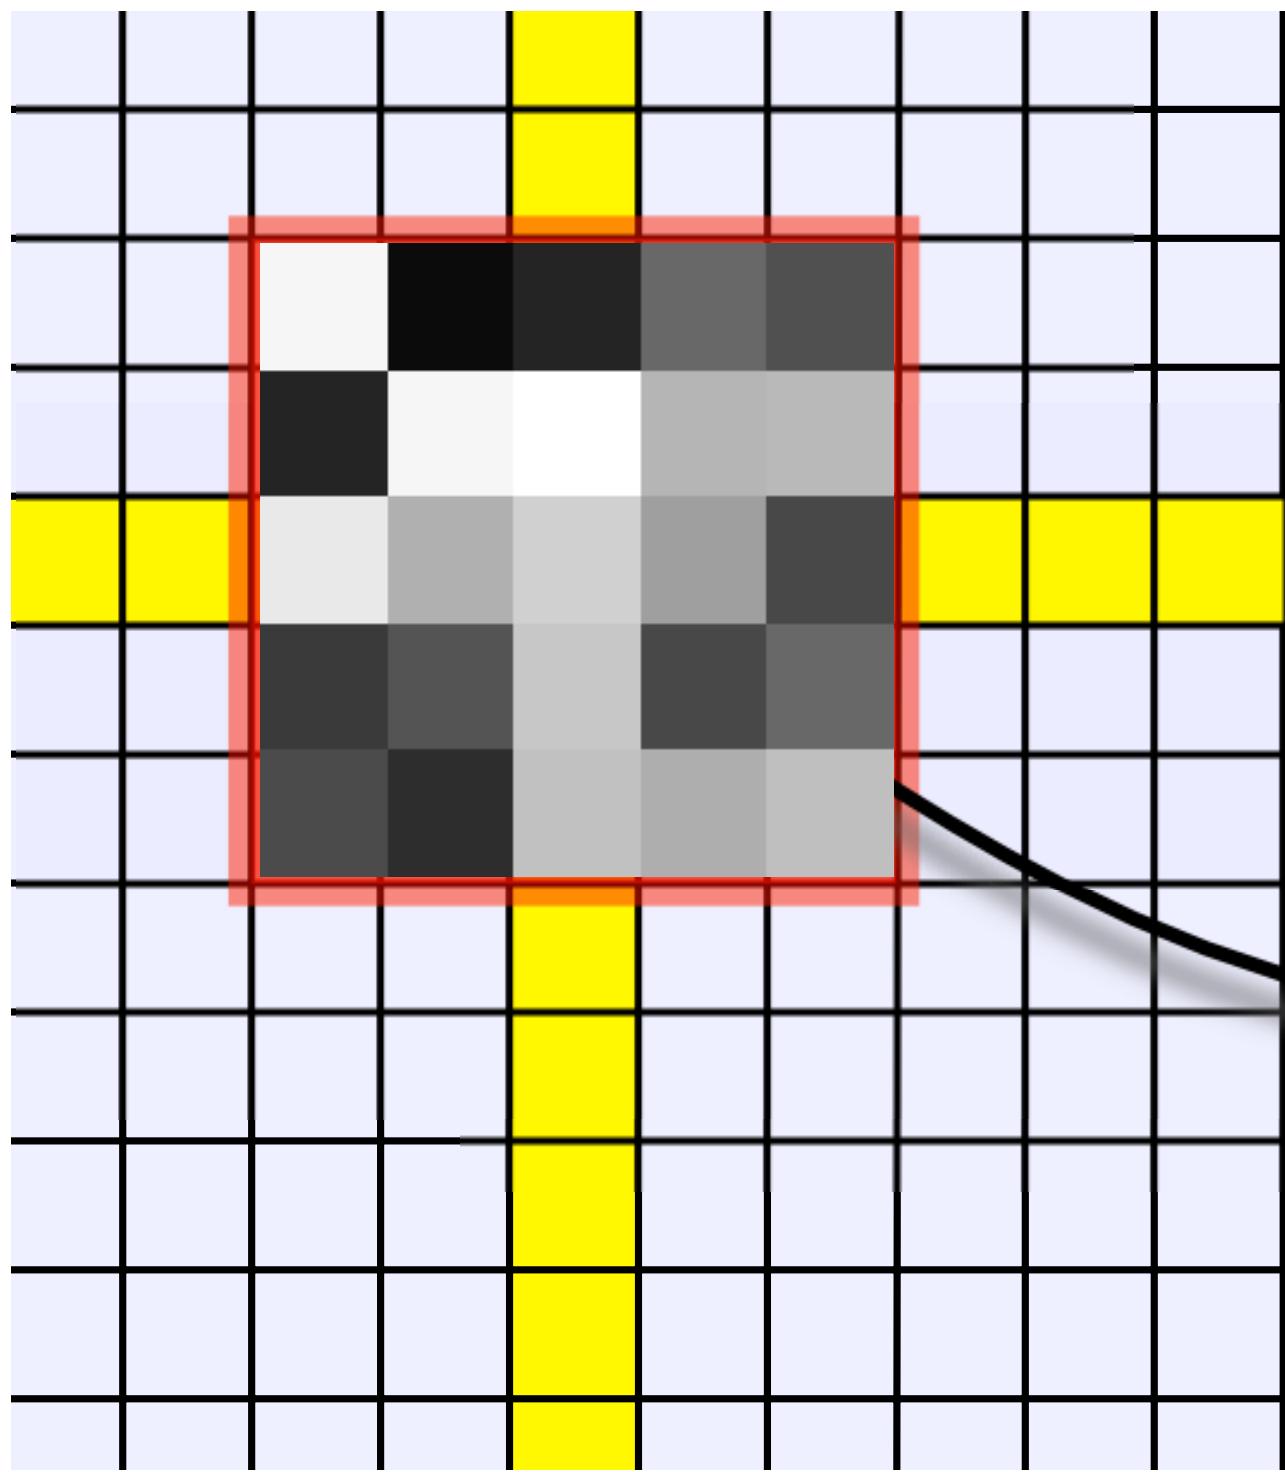


With kind permission of Springer Science+Business Media



- Circle of radius 2.5 pixels

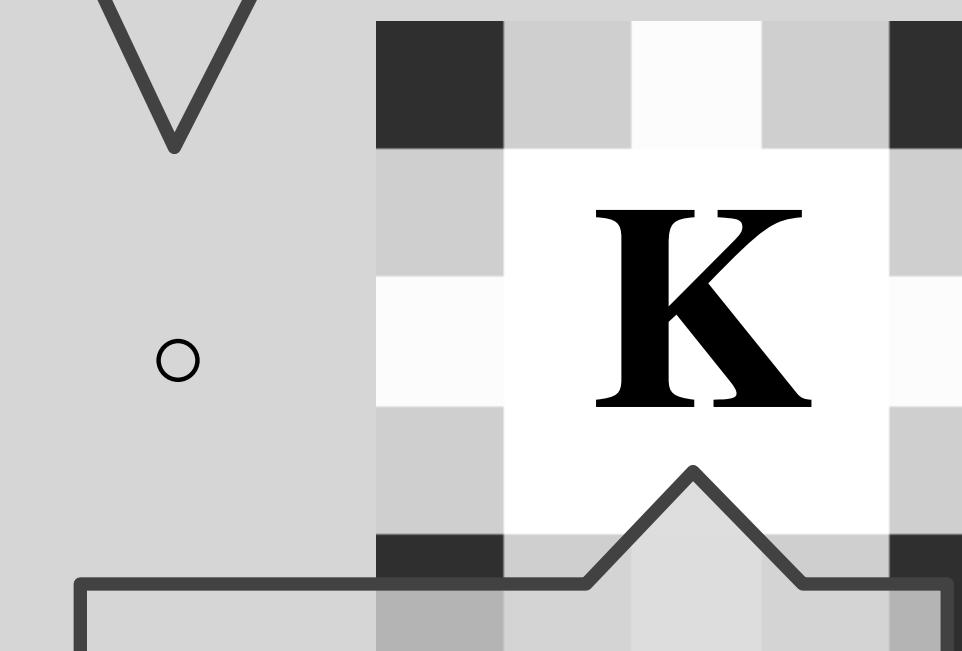
Weighted kernel



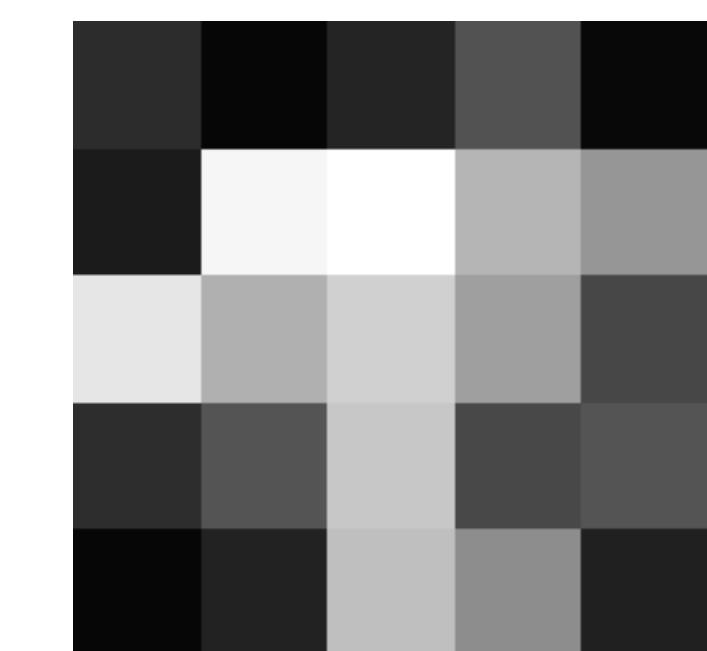
$f(\mathcal{W})$

$$f(\mathcal{W}) = \sum_{\mathcal{W}} (\mathcal{W} \circ K) =$$

Hadamard
product



Kernel



Scaling the kernel

$$\mathbf{K} = \begin{pmatrix} 0.14 & 0.77 & 0.98 & 0.77 & 0.14 \\ 0.77 & 1.0 & 1.0 & 1.0 & 0.77 \\ 0.98 & 1.0 & 1.0 & 1.0 & 0.98 \\ 0.77 & 1.0 & 1.0 & 1.0 & 0.77 \\ 0.14 & 0.77 & 0.98 & 0.77 & 0.14 \end{pmatrix}$$

Total value is 19.6

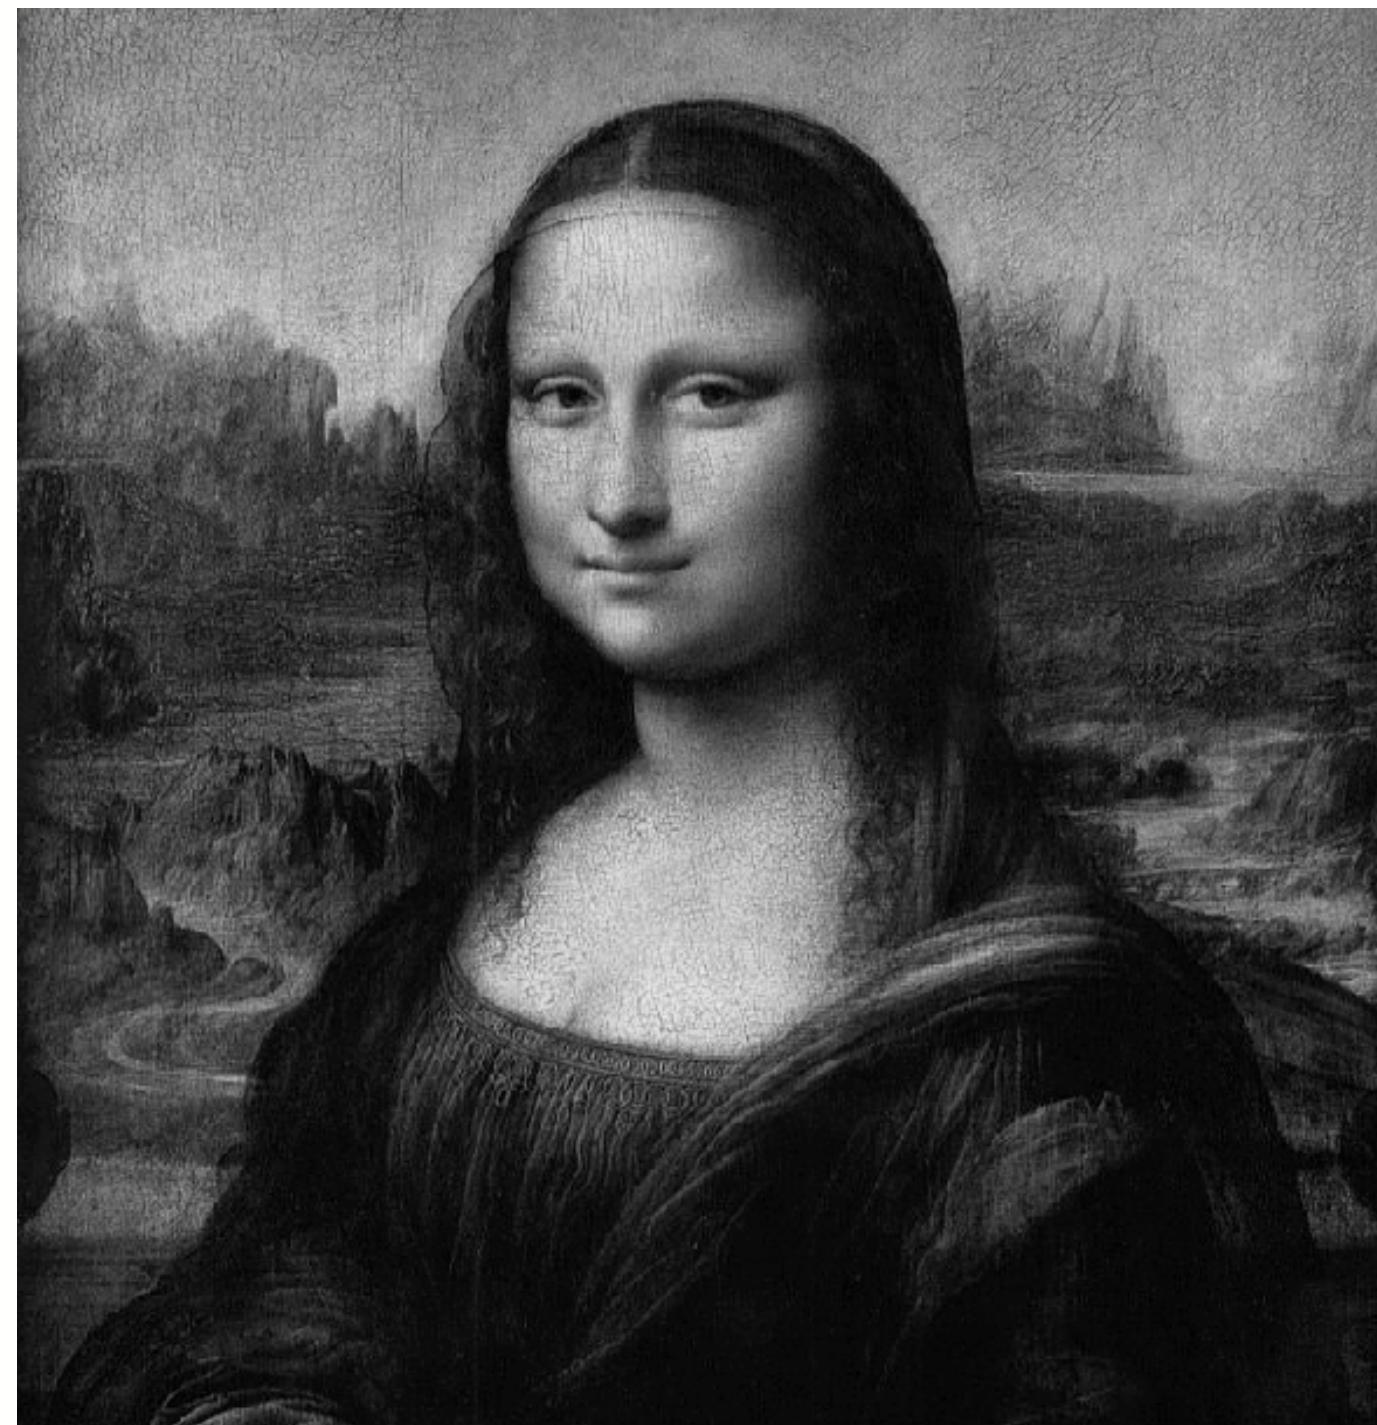
- The scale factor is $S = \sum |\mathbf{K}_{i,j}|$

Scaling the kernel

$$\mathbf{K} = \frac{1}{19.6} \begin{pmatrix} 0.14 & 0.77 & 0.98 & 0.77 & 0.14 \\ 0.77 & 1.0 & 1.0 & 1.0 & 0.77 \\ 0.98 & 1.0 & 1.0 & 1.0 & 0.98 \\ 0.77 & 1.0 & 1.0 & 1.0 & 0.77 \\ 0.14 & 0.77 & 0.98 & 0.77 & 0.14 \end{pmatrix}$$

- The scale factor is $S = \sum |\mathbf{K}_{i,j}|$
- Typically make $S = 1$ to keep grey levels the same as the input image

Simple averaging is also a kernel



Original image

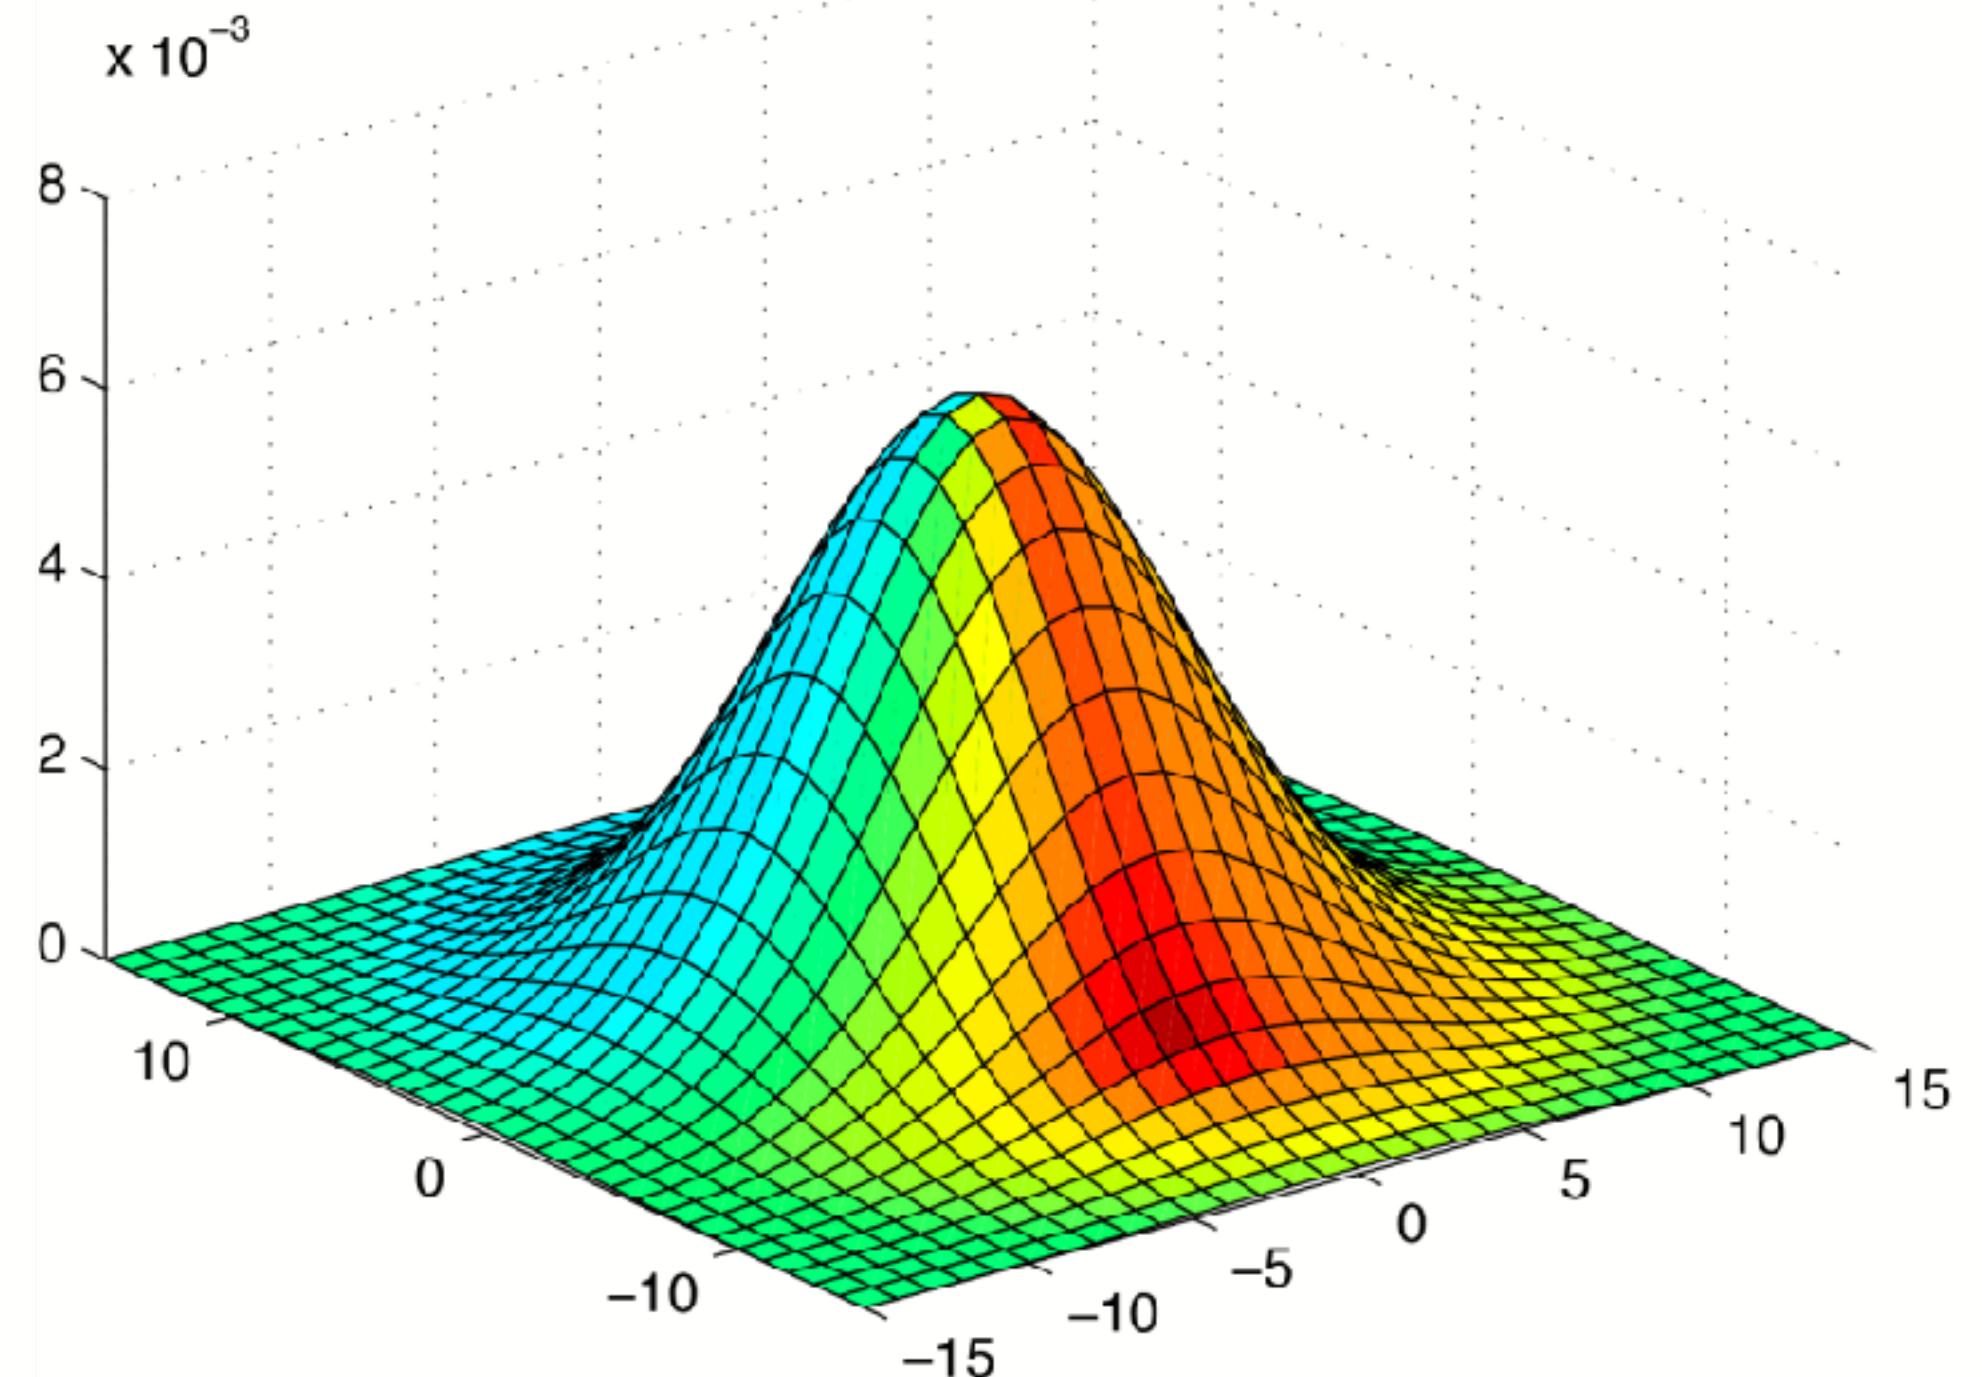


Average each pixel over a 7x7 window

$$\mathbf{K} = \frac{1}{49} \begin{pmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{pmatrix}$$

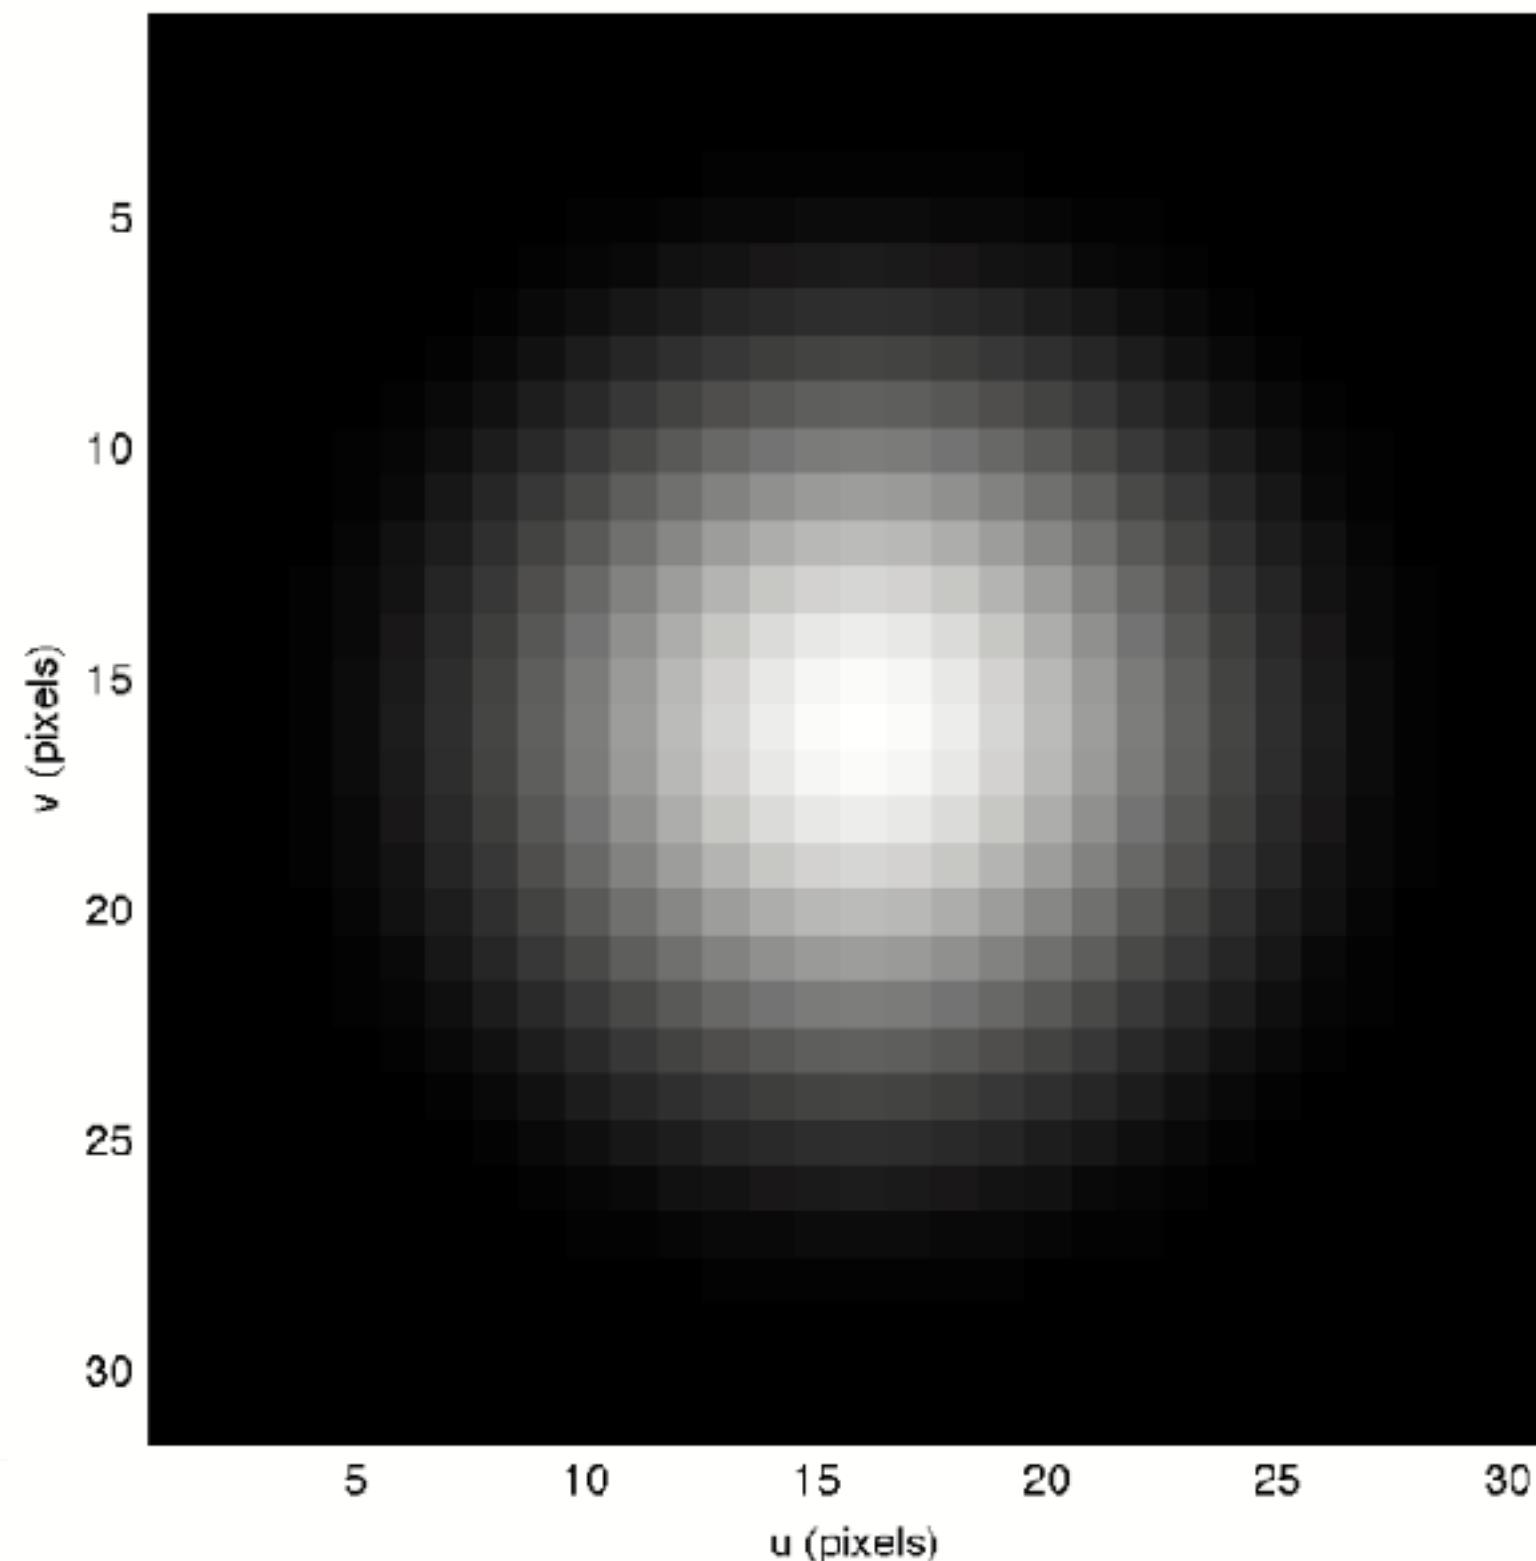
```
>>> out = mona.convolve(np.ones((7,7)) / 49)
```

Gaussian kernel



With kind permission of Springer Science+Business Media

$$G(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



- Isotropic
- Decreasing weight away from centre

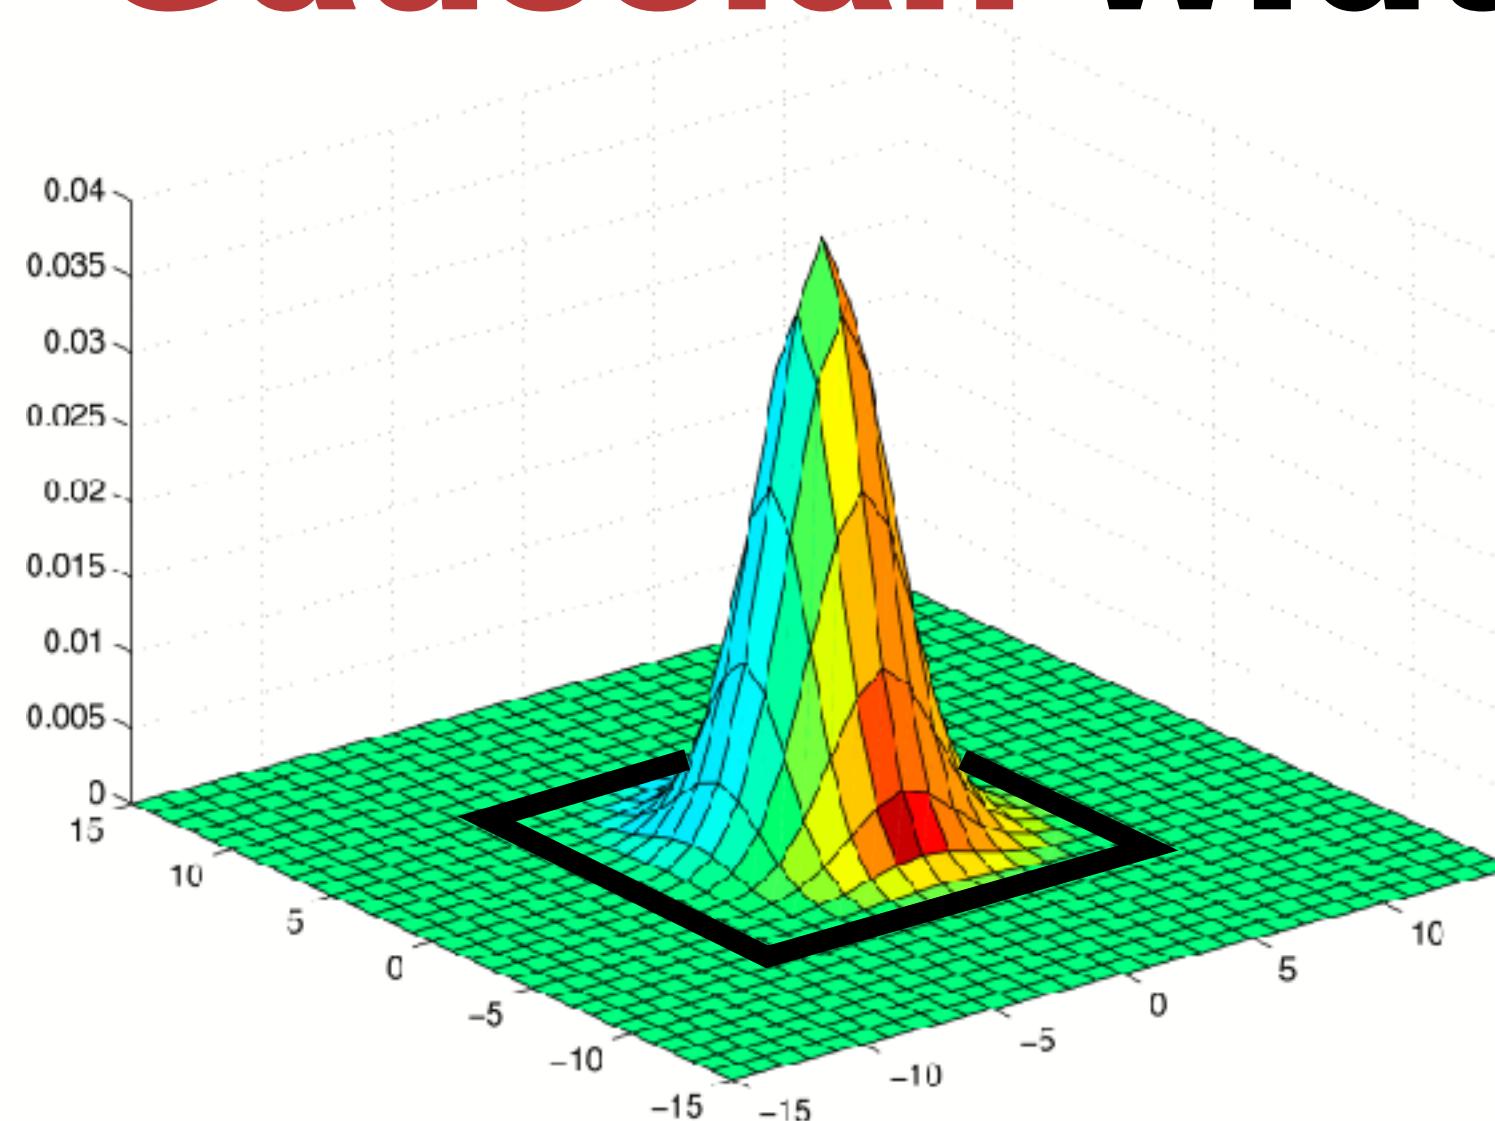


Carl-Friedrich Gauss
1777–1855

By Gottlieb Biermann A. Wittmann (photo)
[Public domain], via Wikimedia Commons

Gaussian width

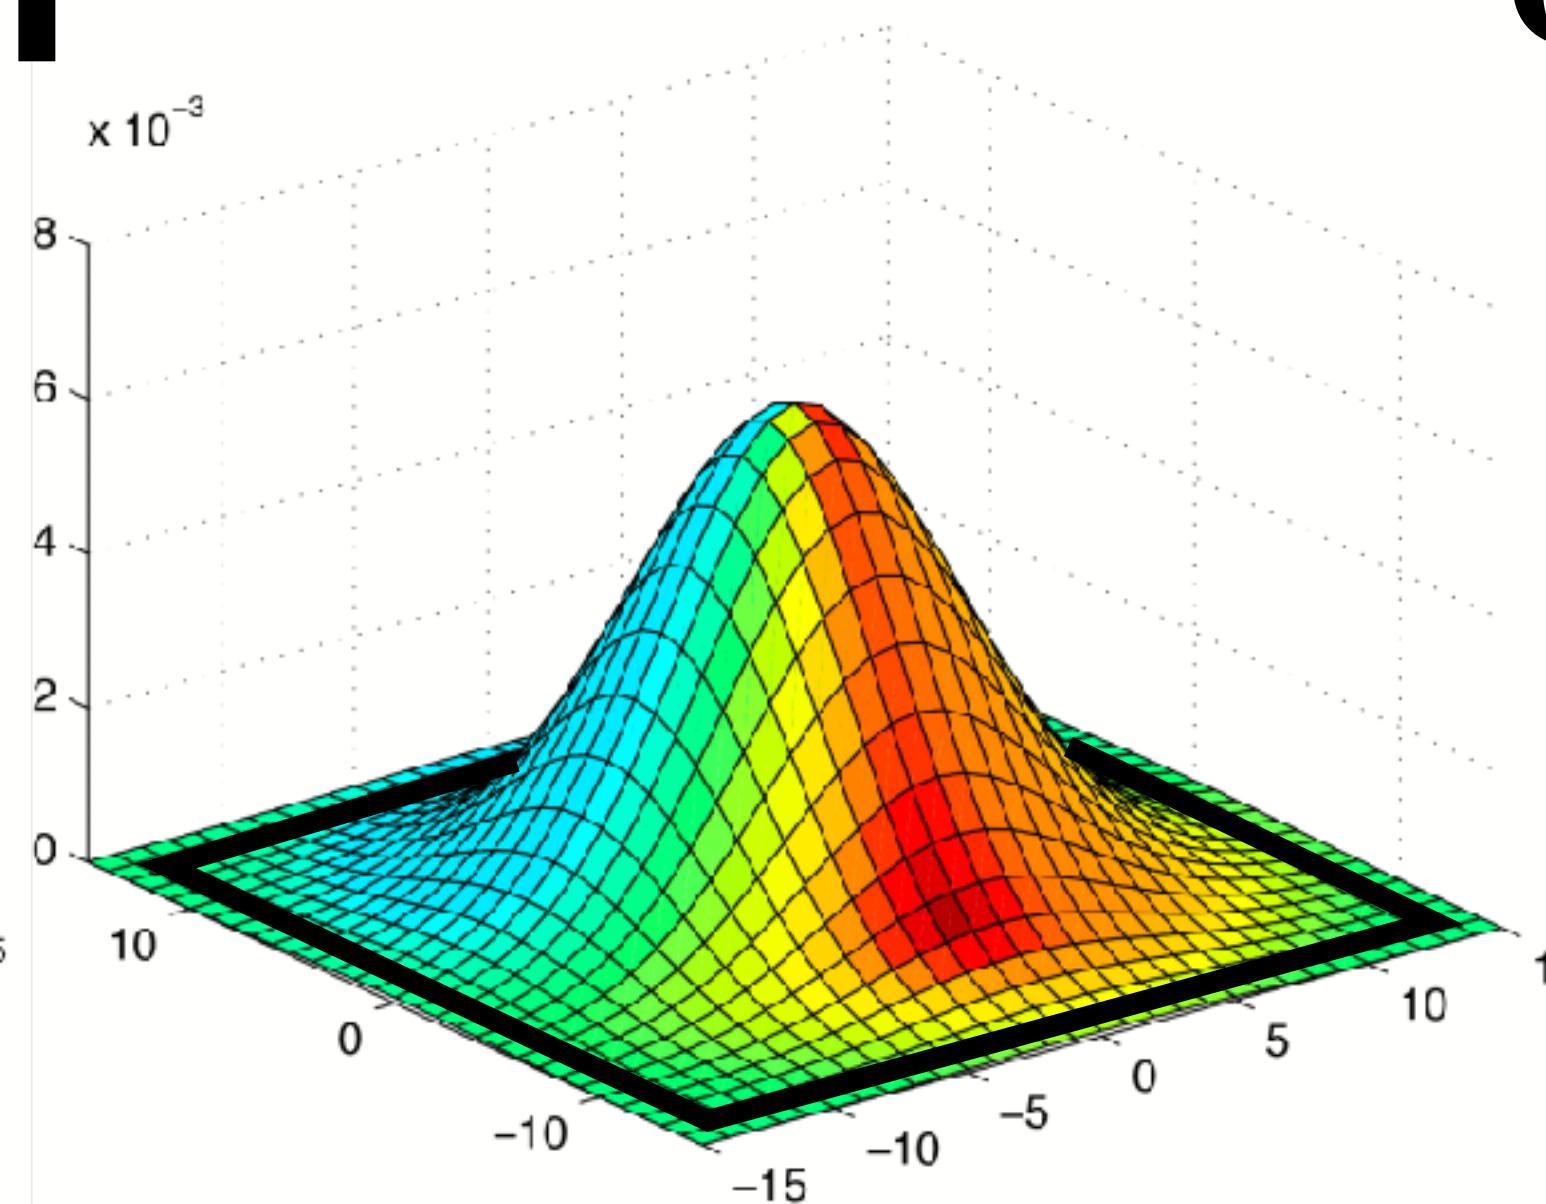
$$G(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



With kind permission of Springer Science+Business Media

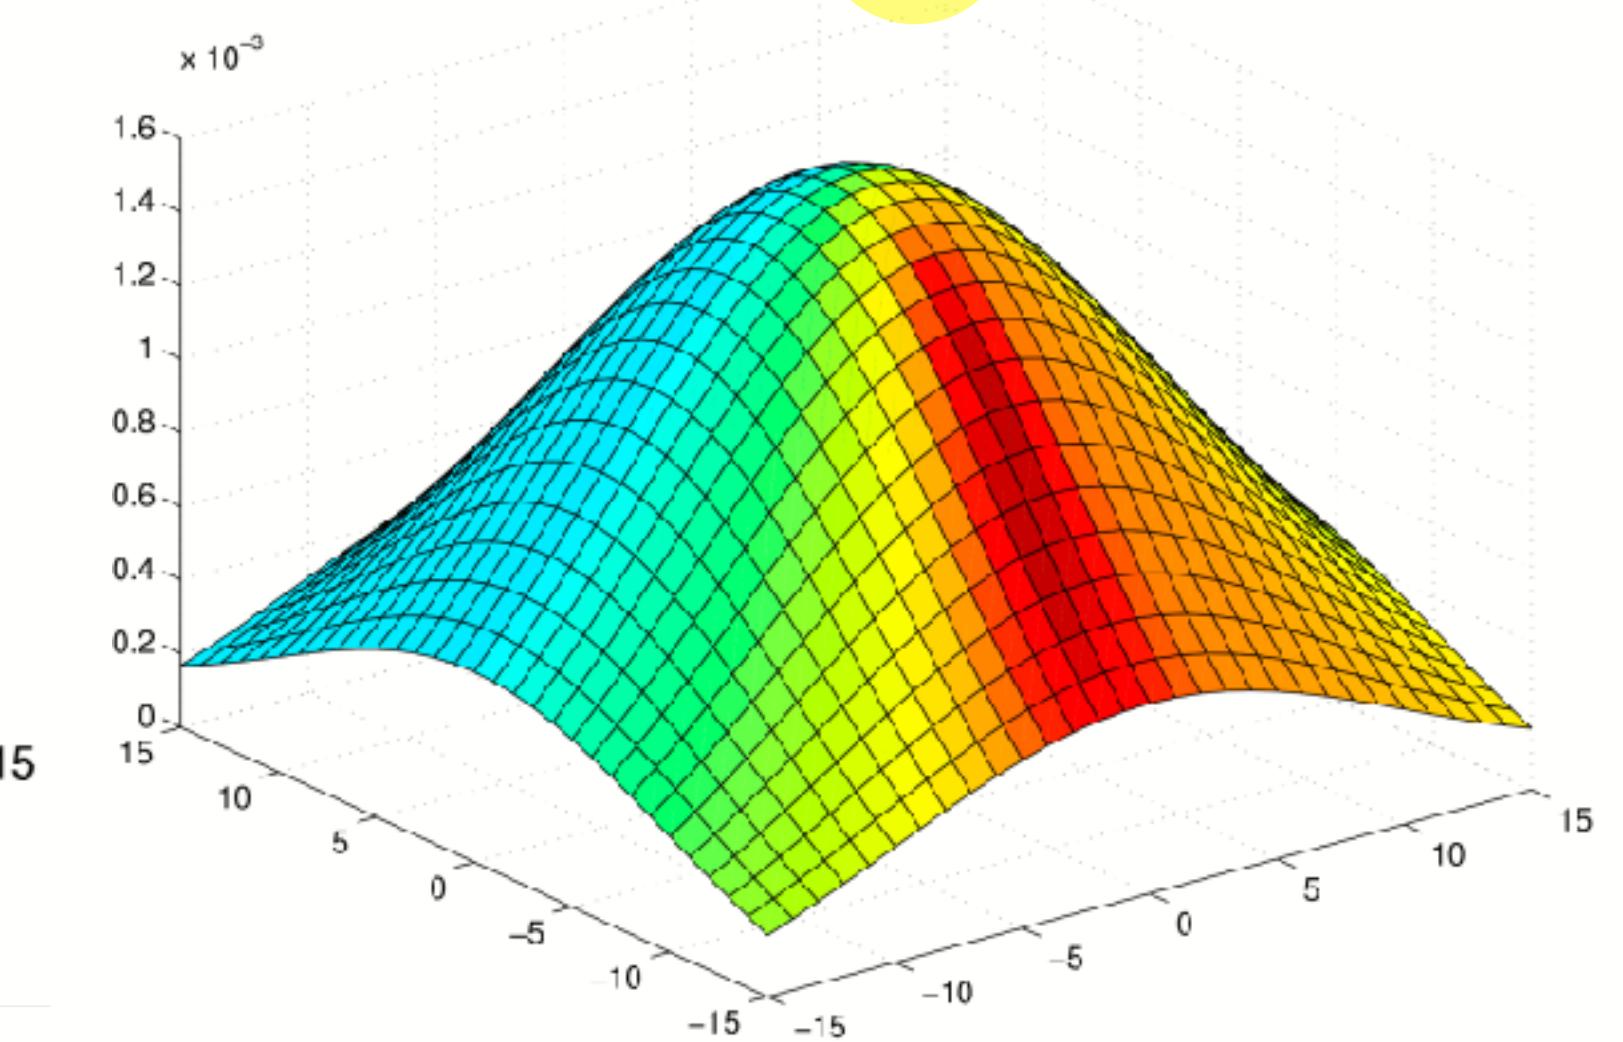
$$\sigma = 2$$

`Kernel.Gauss(2, 15)`



$$\sigma = 5$$

`Kernel.Gauss(5, 15)`

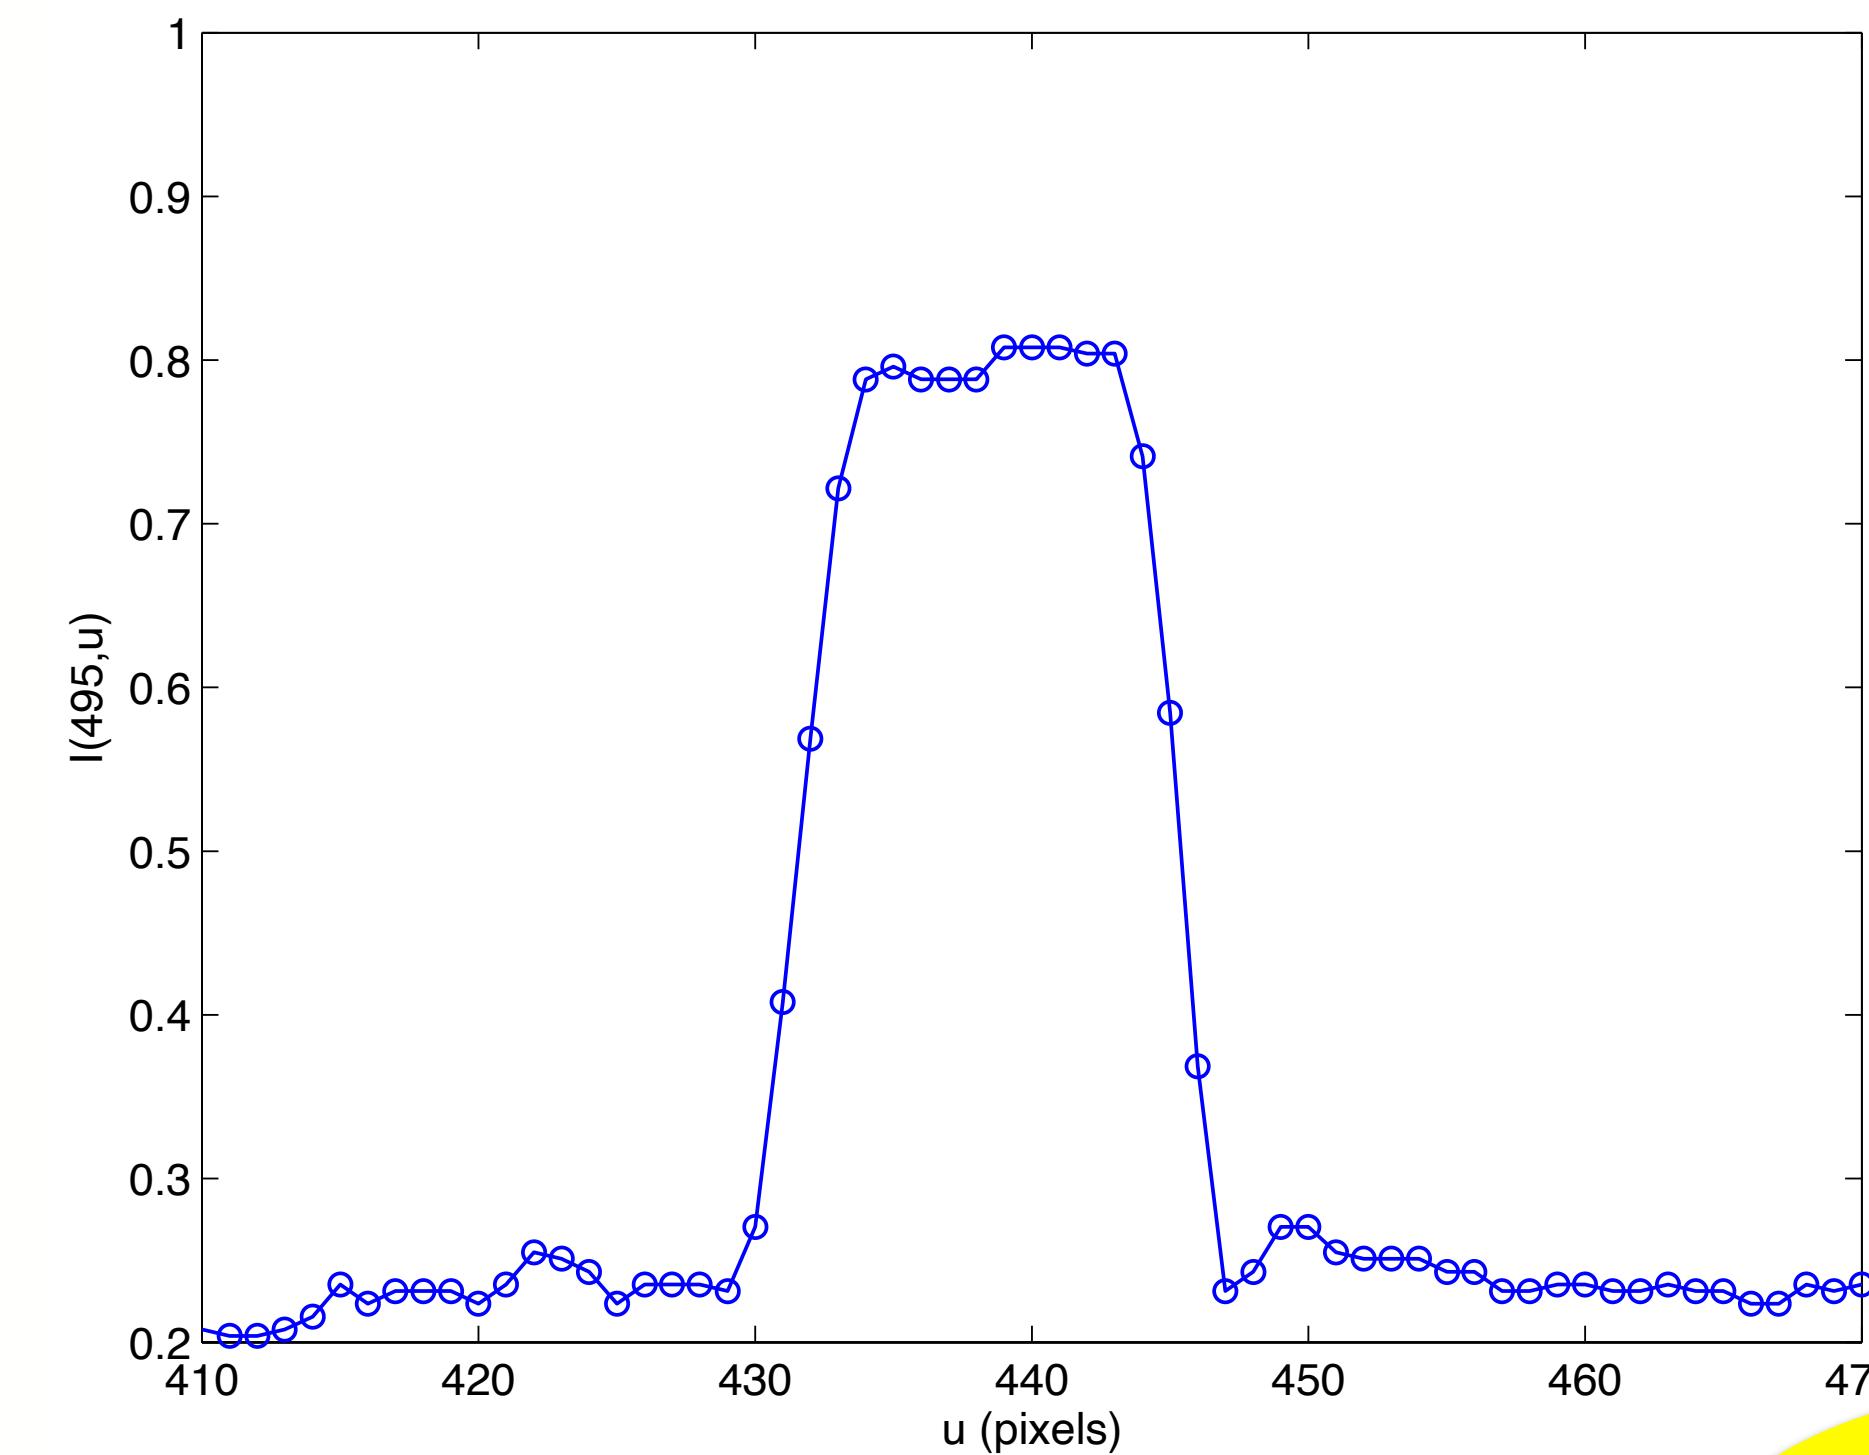
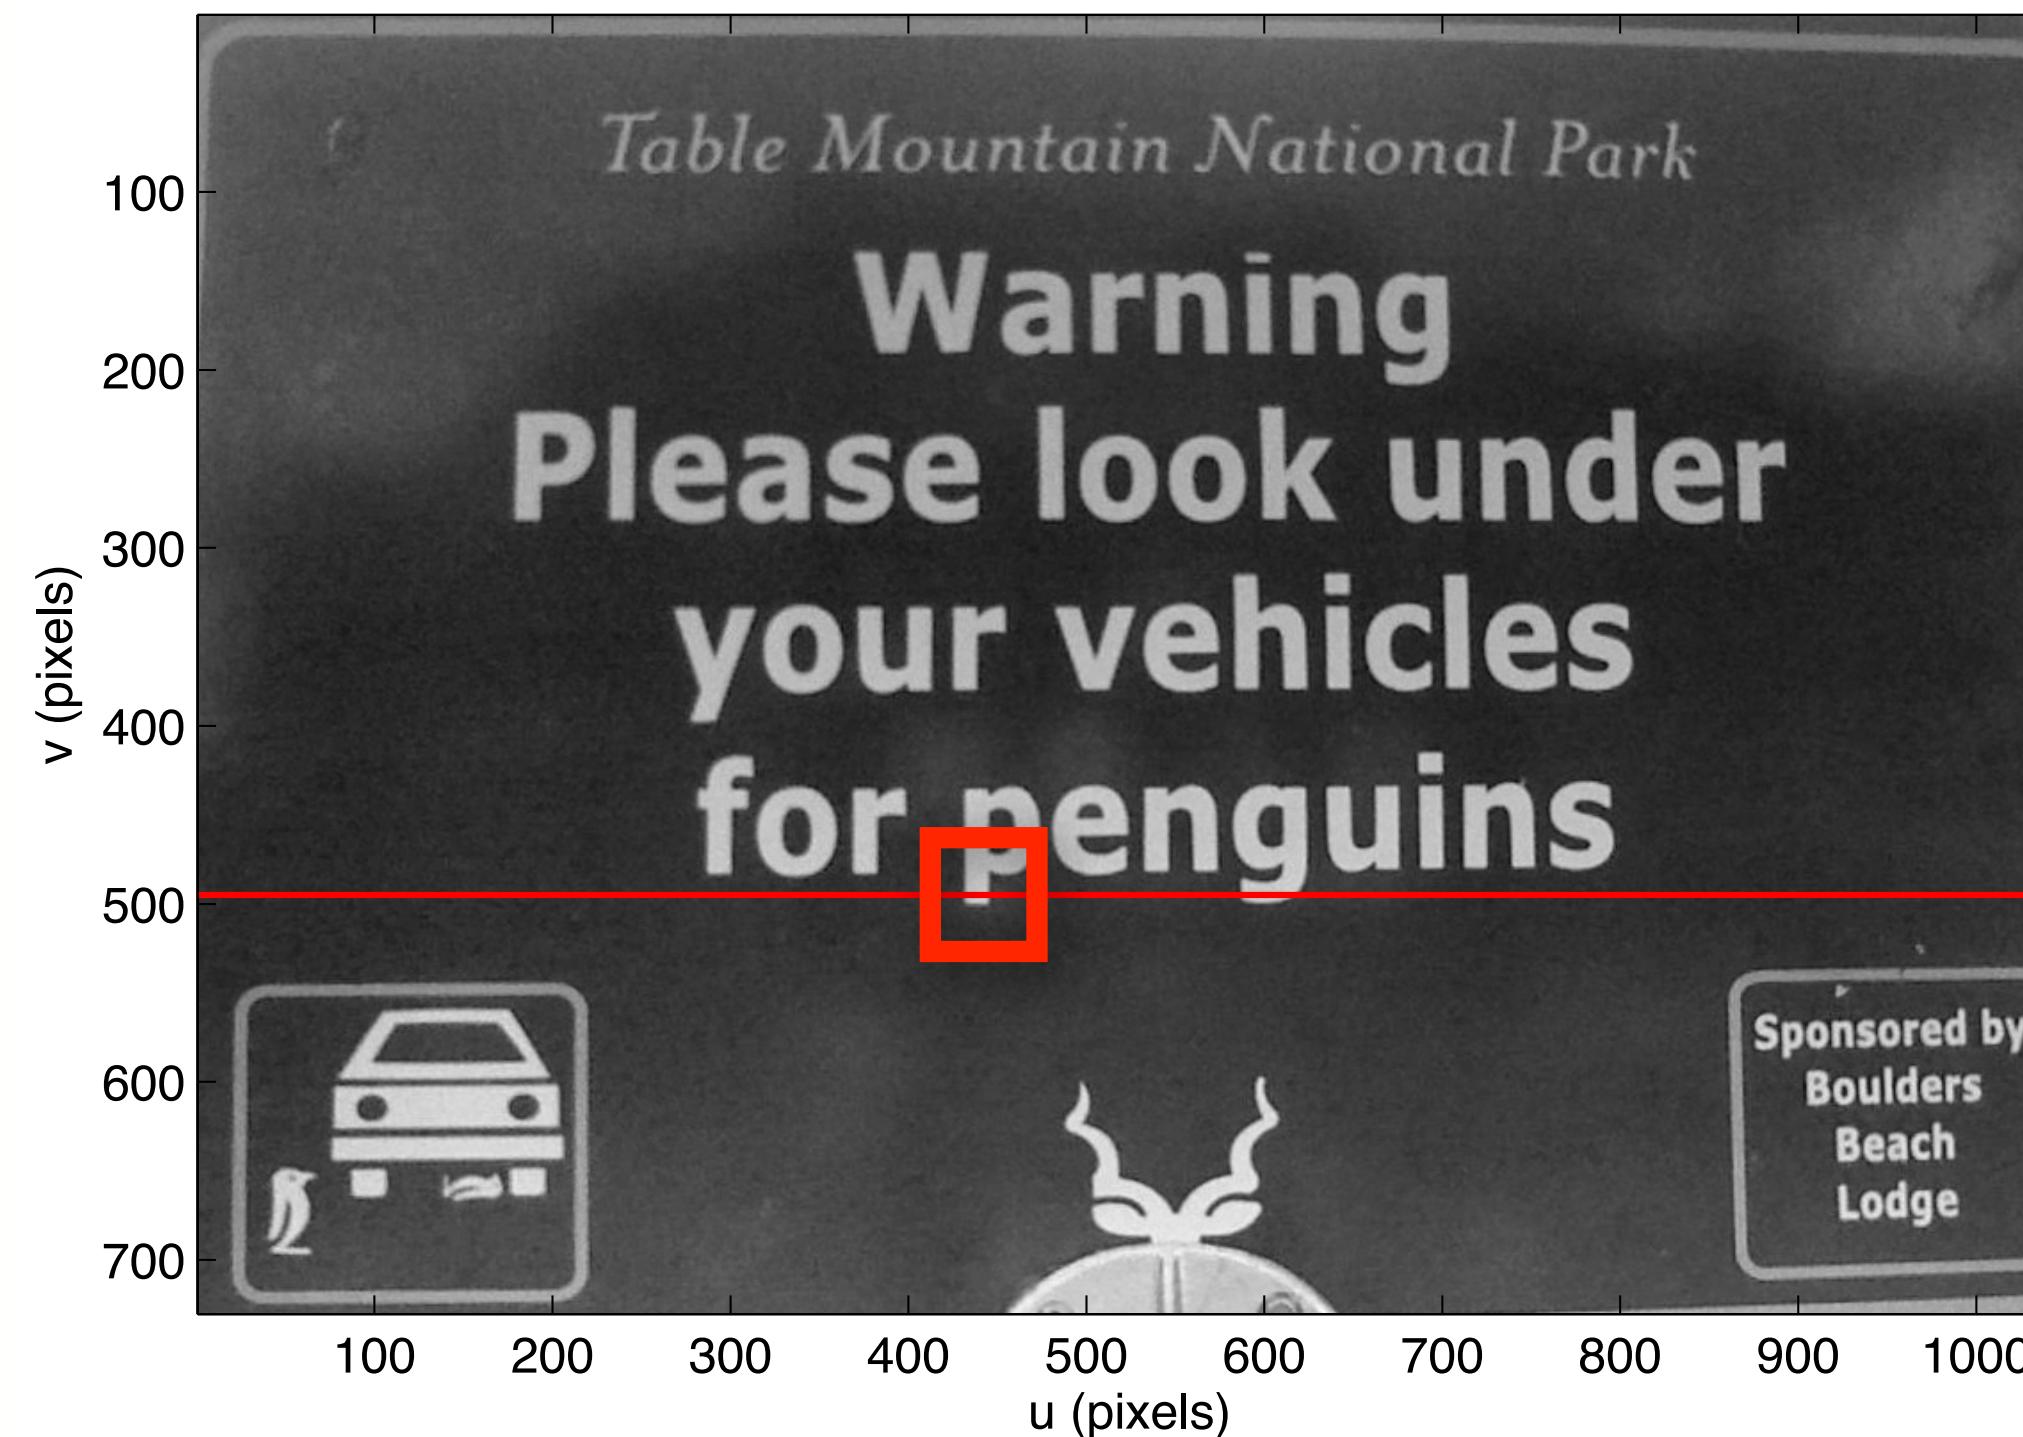


$$\sigma = 10$$

`Kernel.Gauss(10, 15)`

- Choose the size of the square kernel to fit the Gaussian
- Rule of thumb $h=3\sigma$

The intensity function



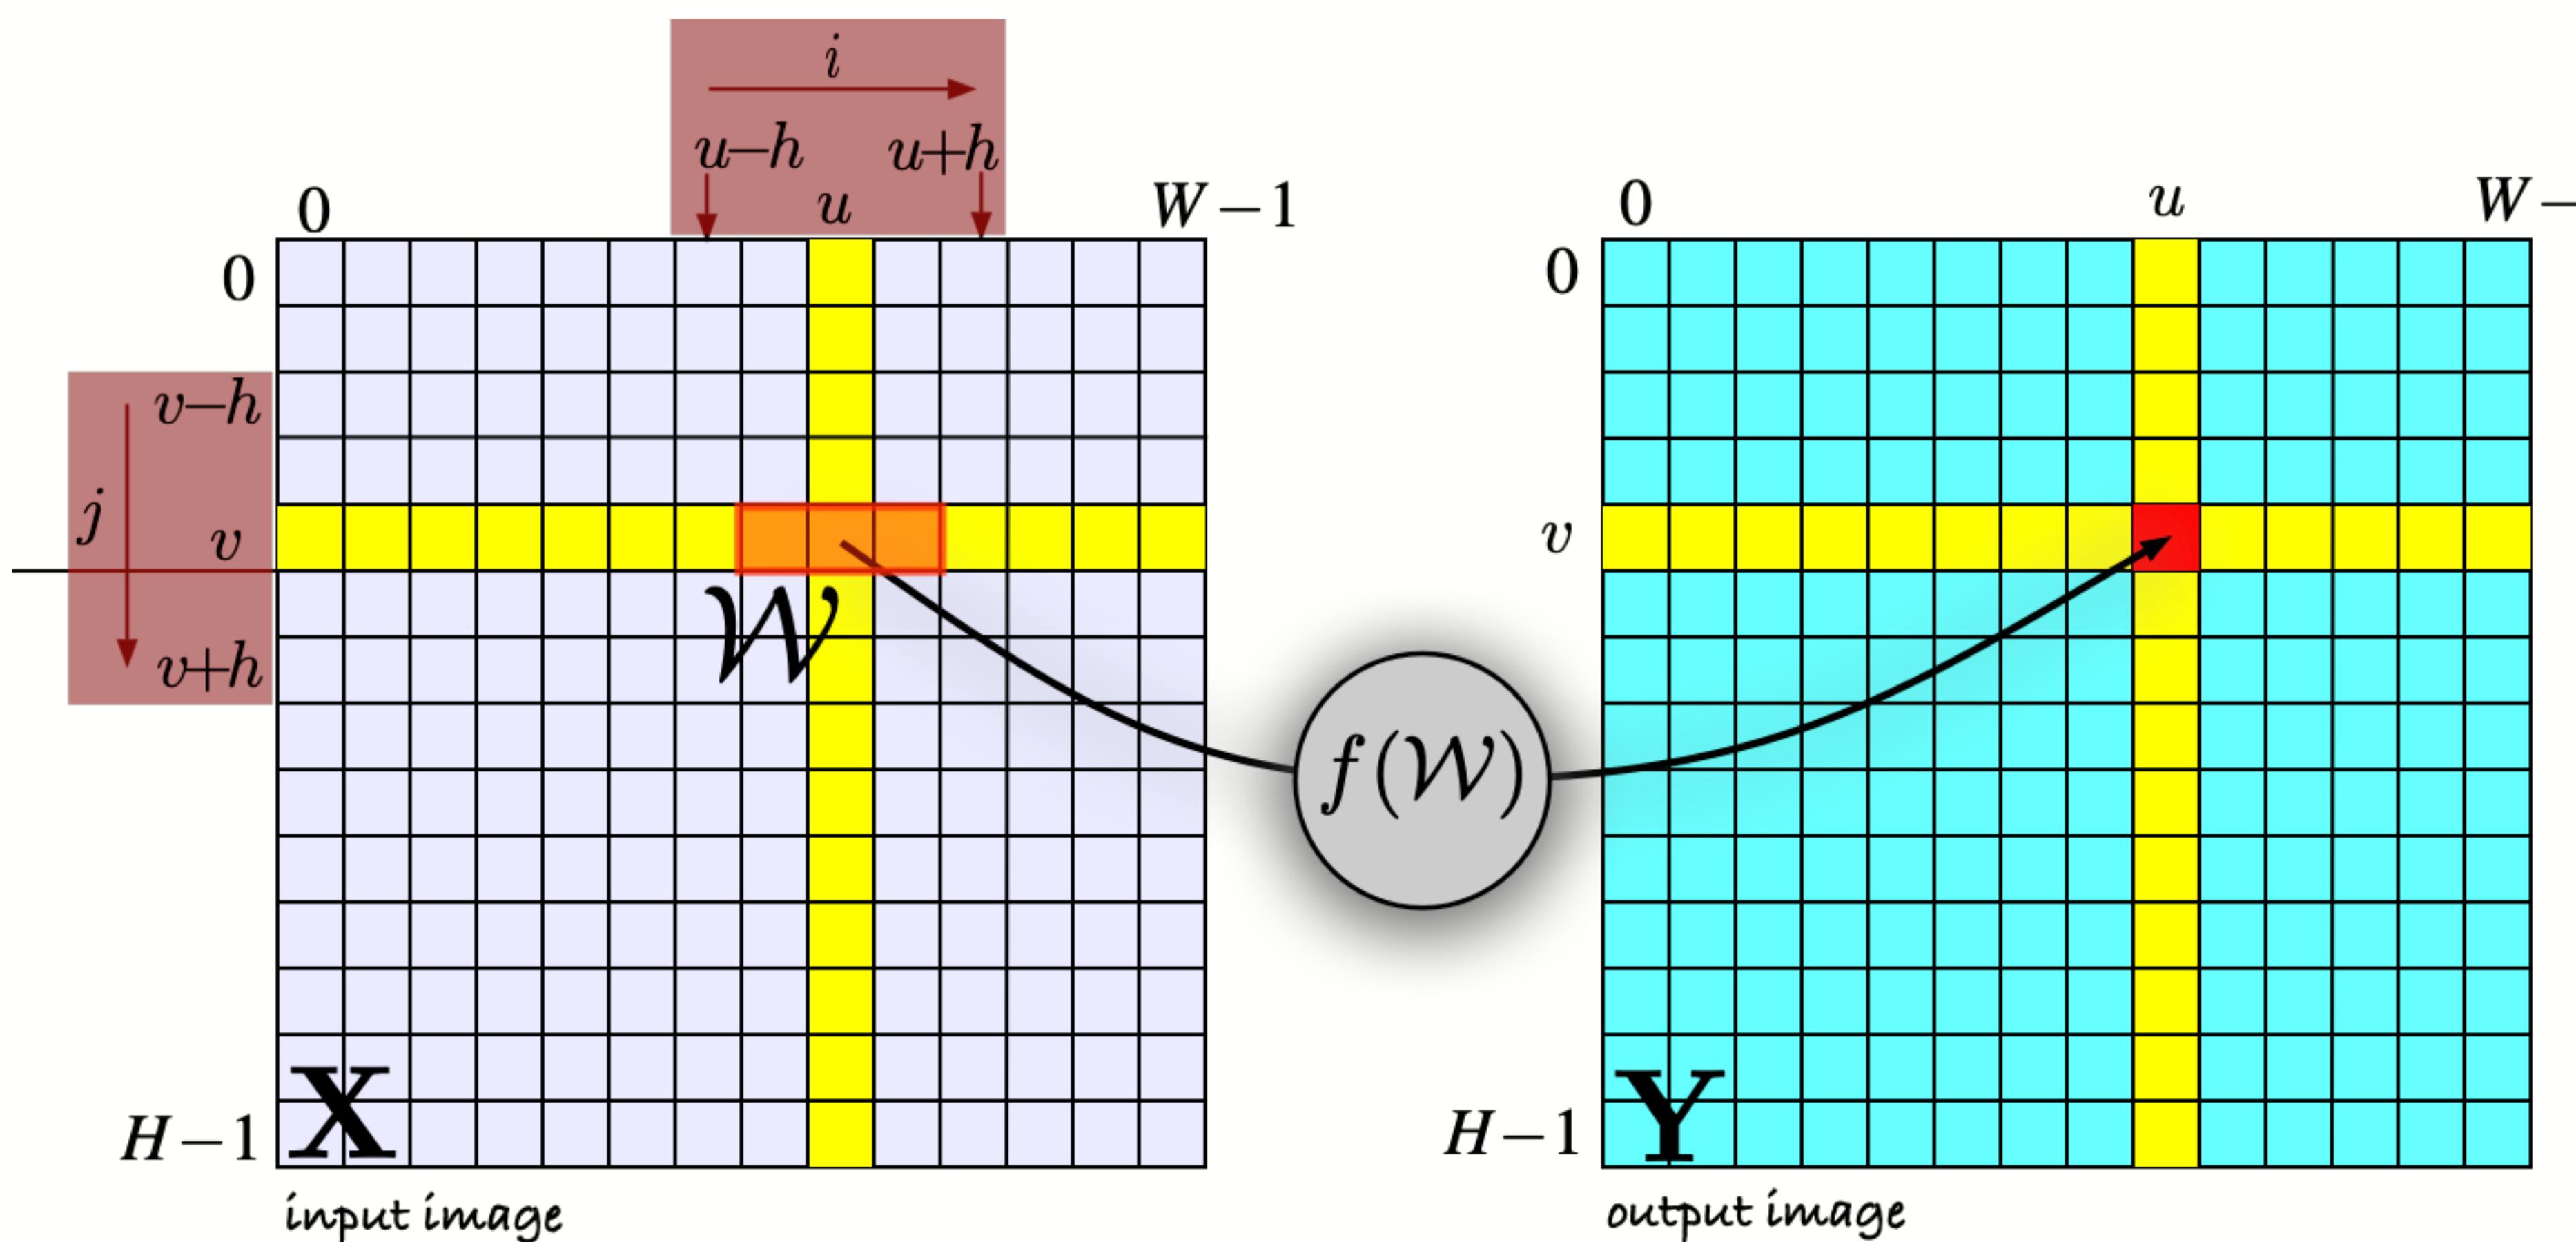
- Horizontal gradient can be approximated:

$$\frac{dI}{du} \approx I(u+1) - I(u)$$

$$\left. \frac{dI}{du} \right|_u \approx \frac{I(u+1) - I(u-1)}{2}$$

symmetrical
difference about
 u

Expressed as correlation with a kernel



With kind permission of Springer Science+Business Media

$$\frac{d\mathbf{X}}{du} \Big|_u \approx \frac{-\mathbf{X}(u-1) + \mathbf{X}(u+1)}{2}$$

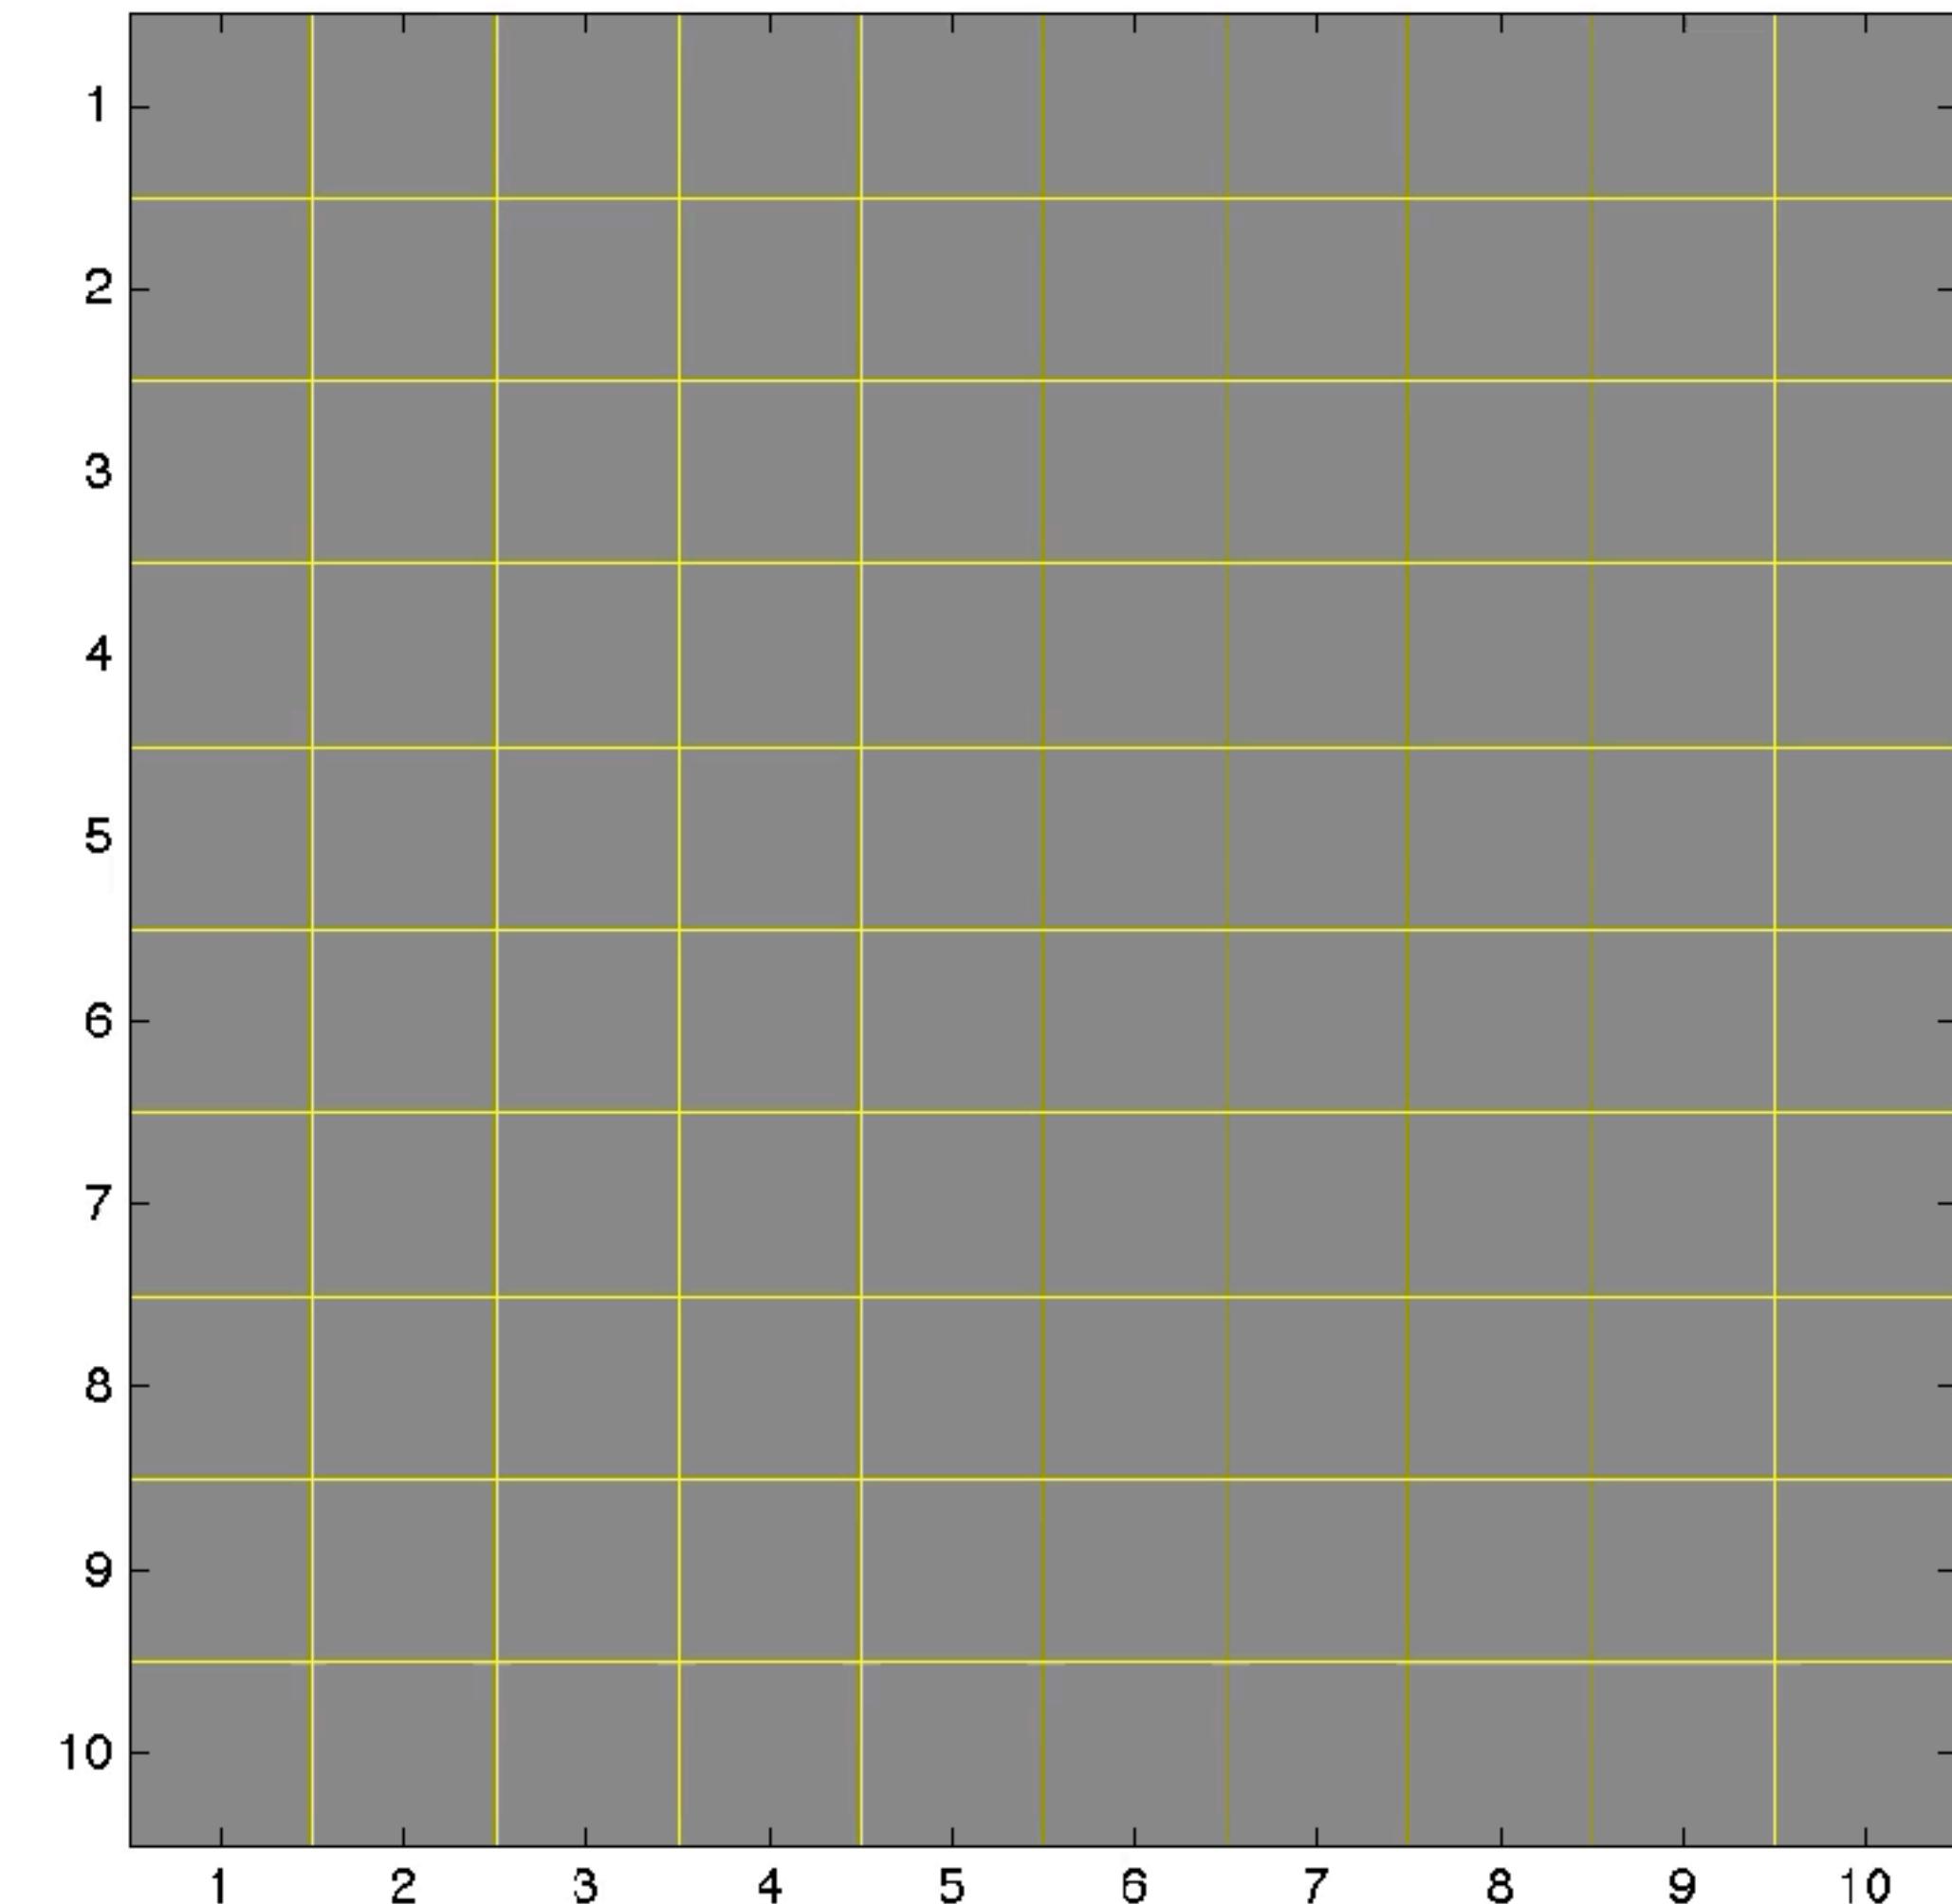
$$\mathbf{K} = \frac{1}{2} \begin{pmatrix} 1 & 0 & -1 \end{pmatrix}$$

Convolution for edge detection

Input image

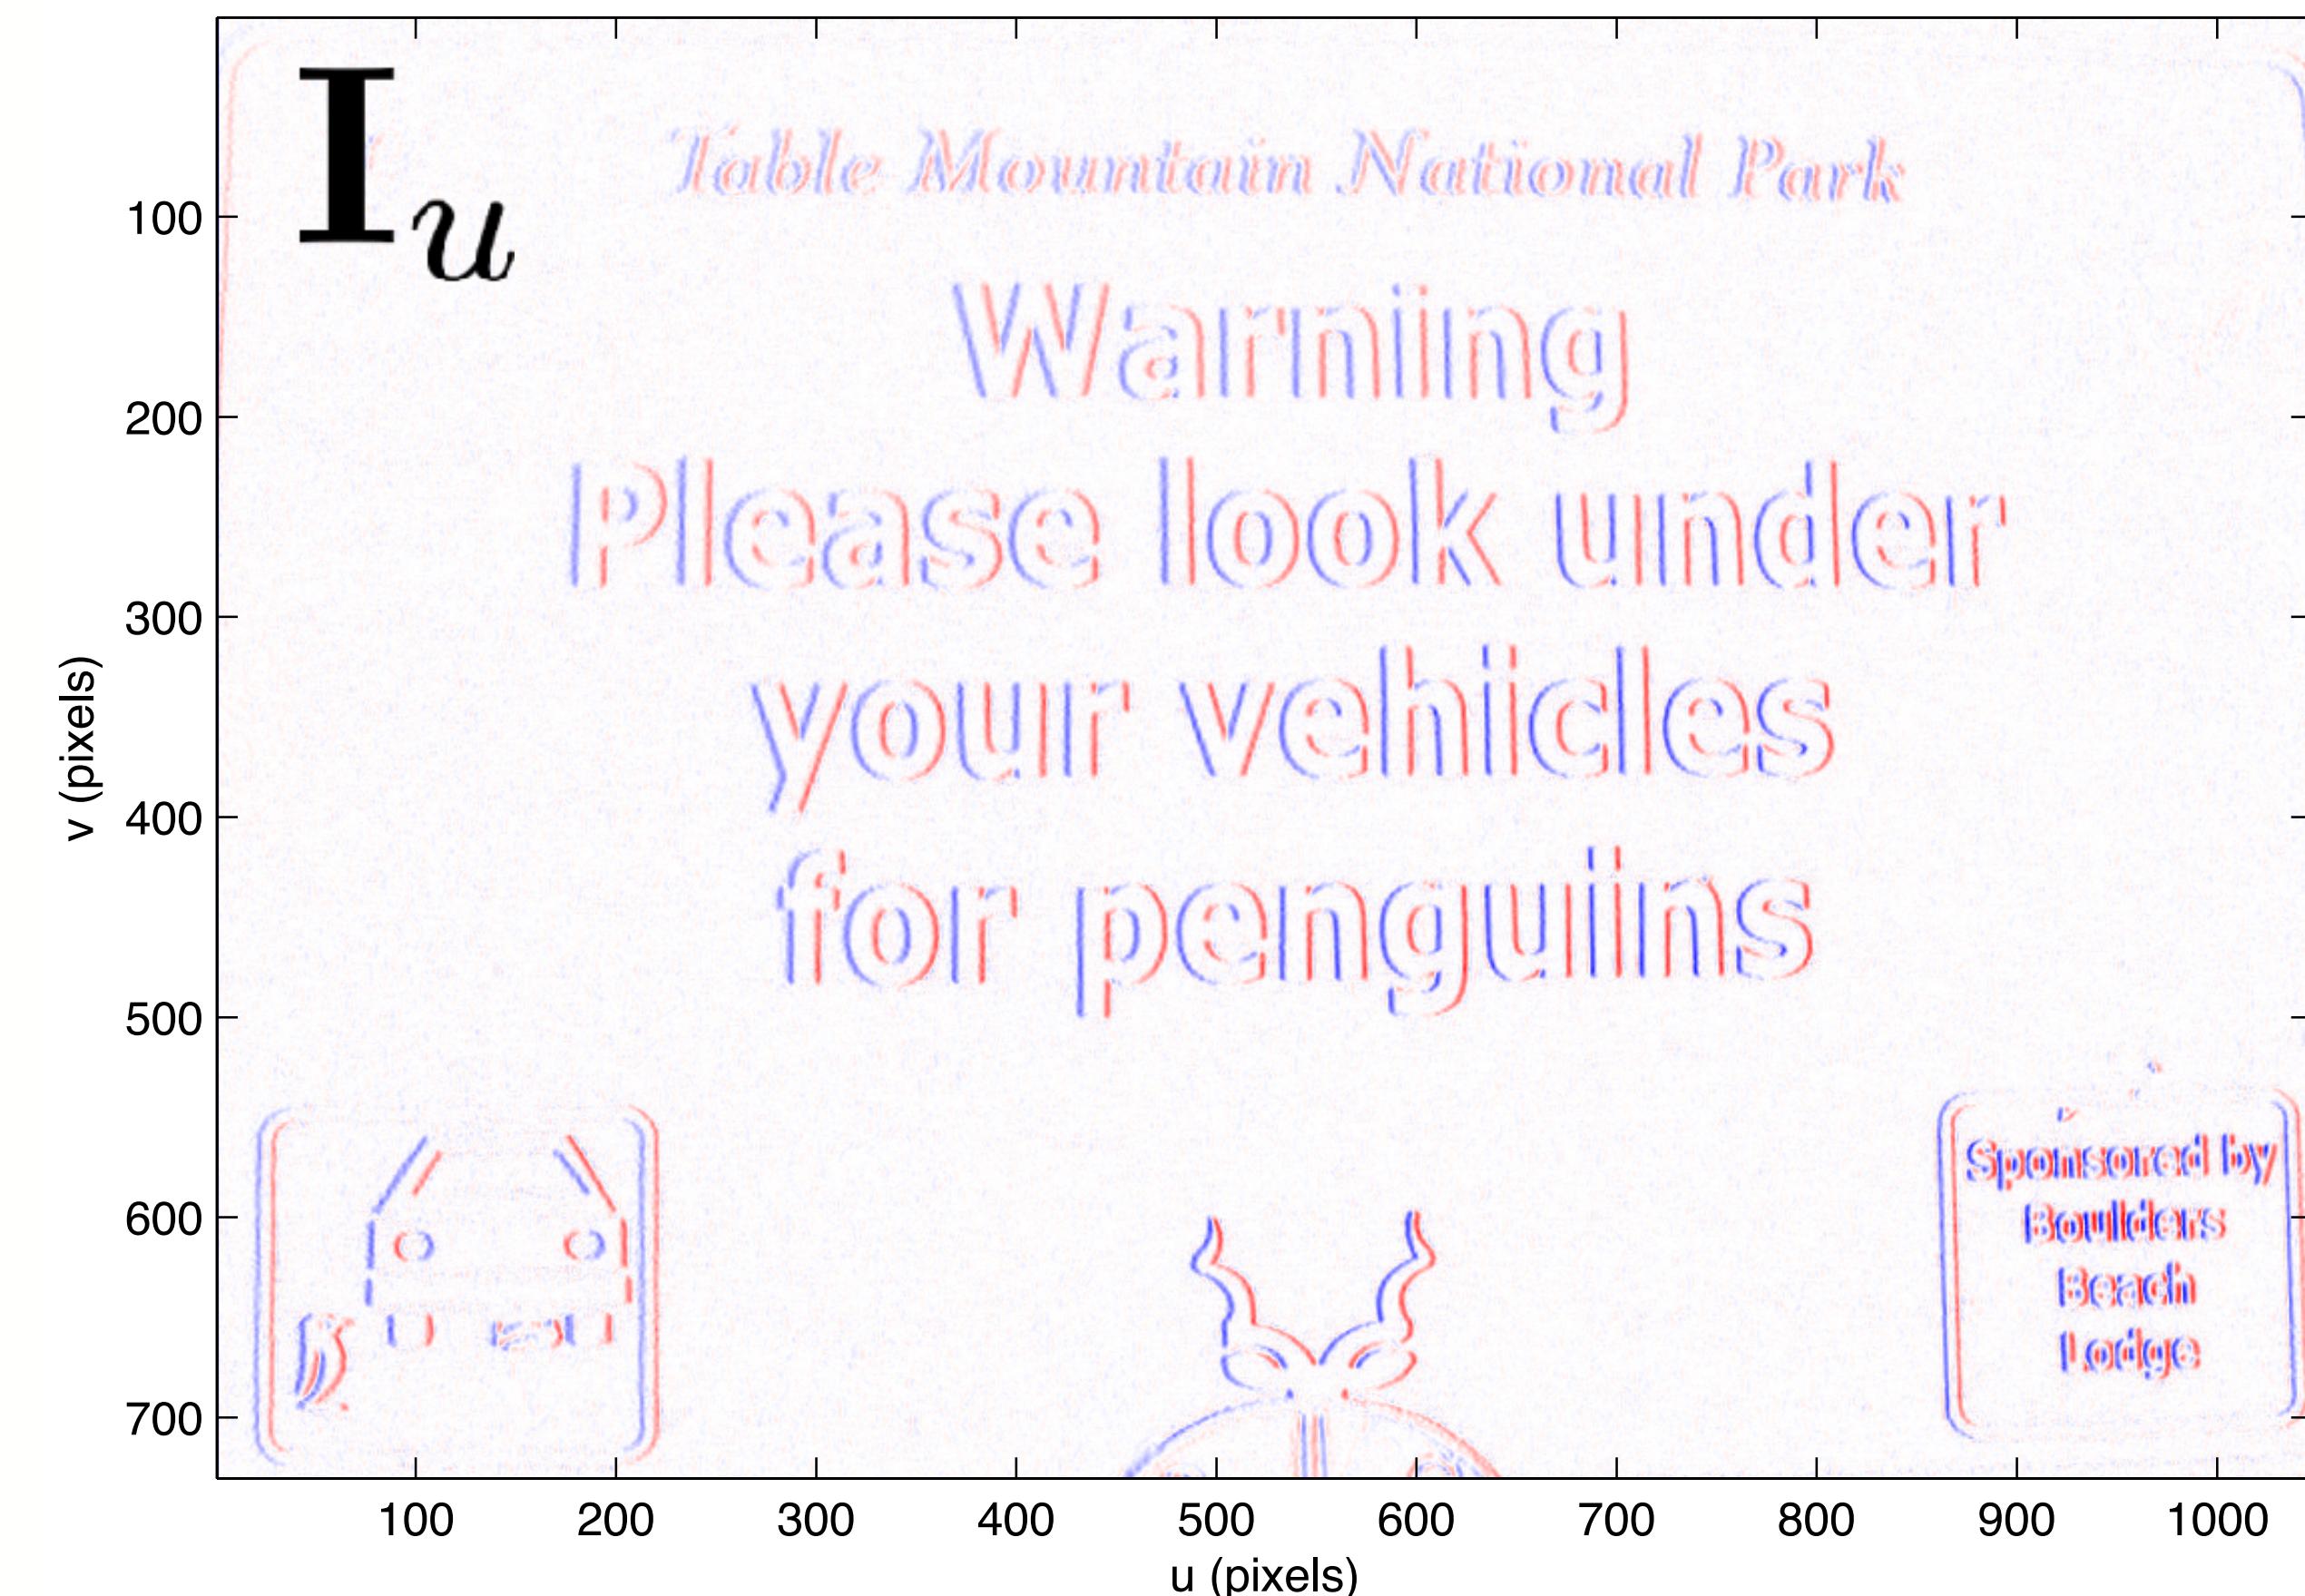
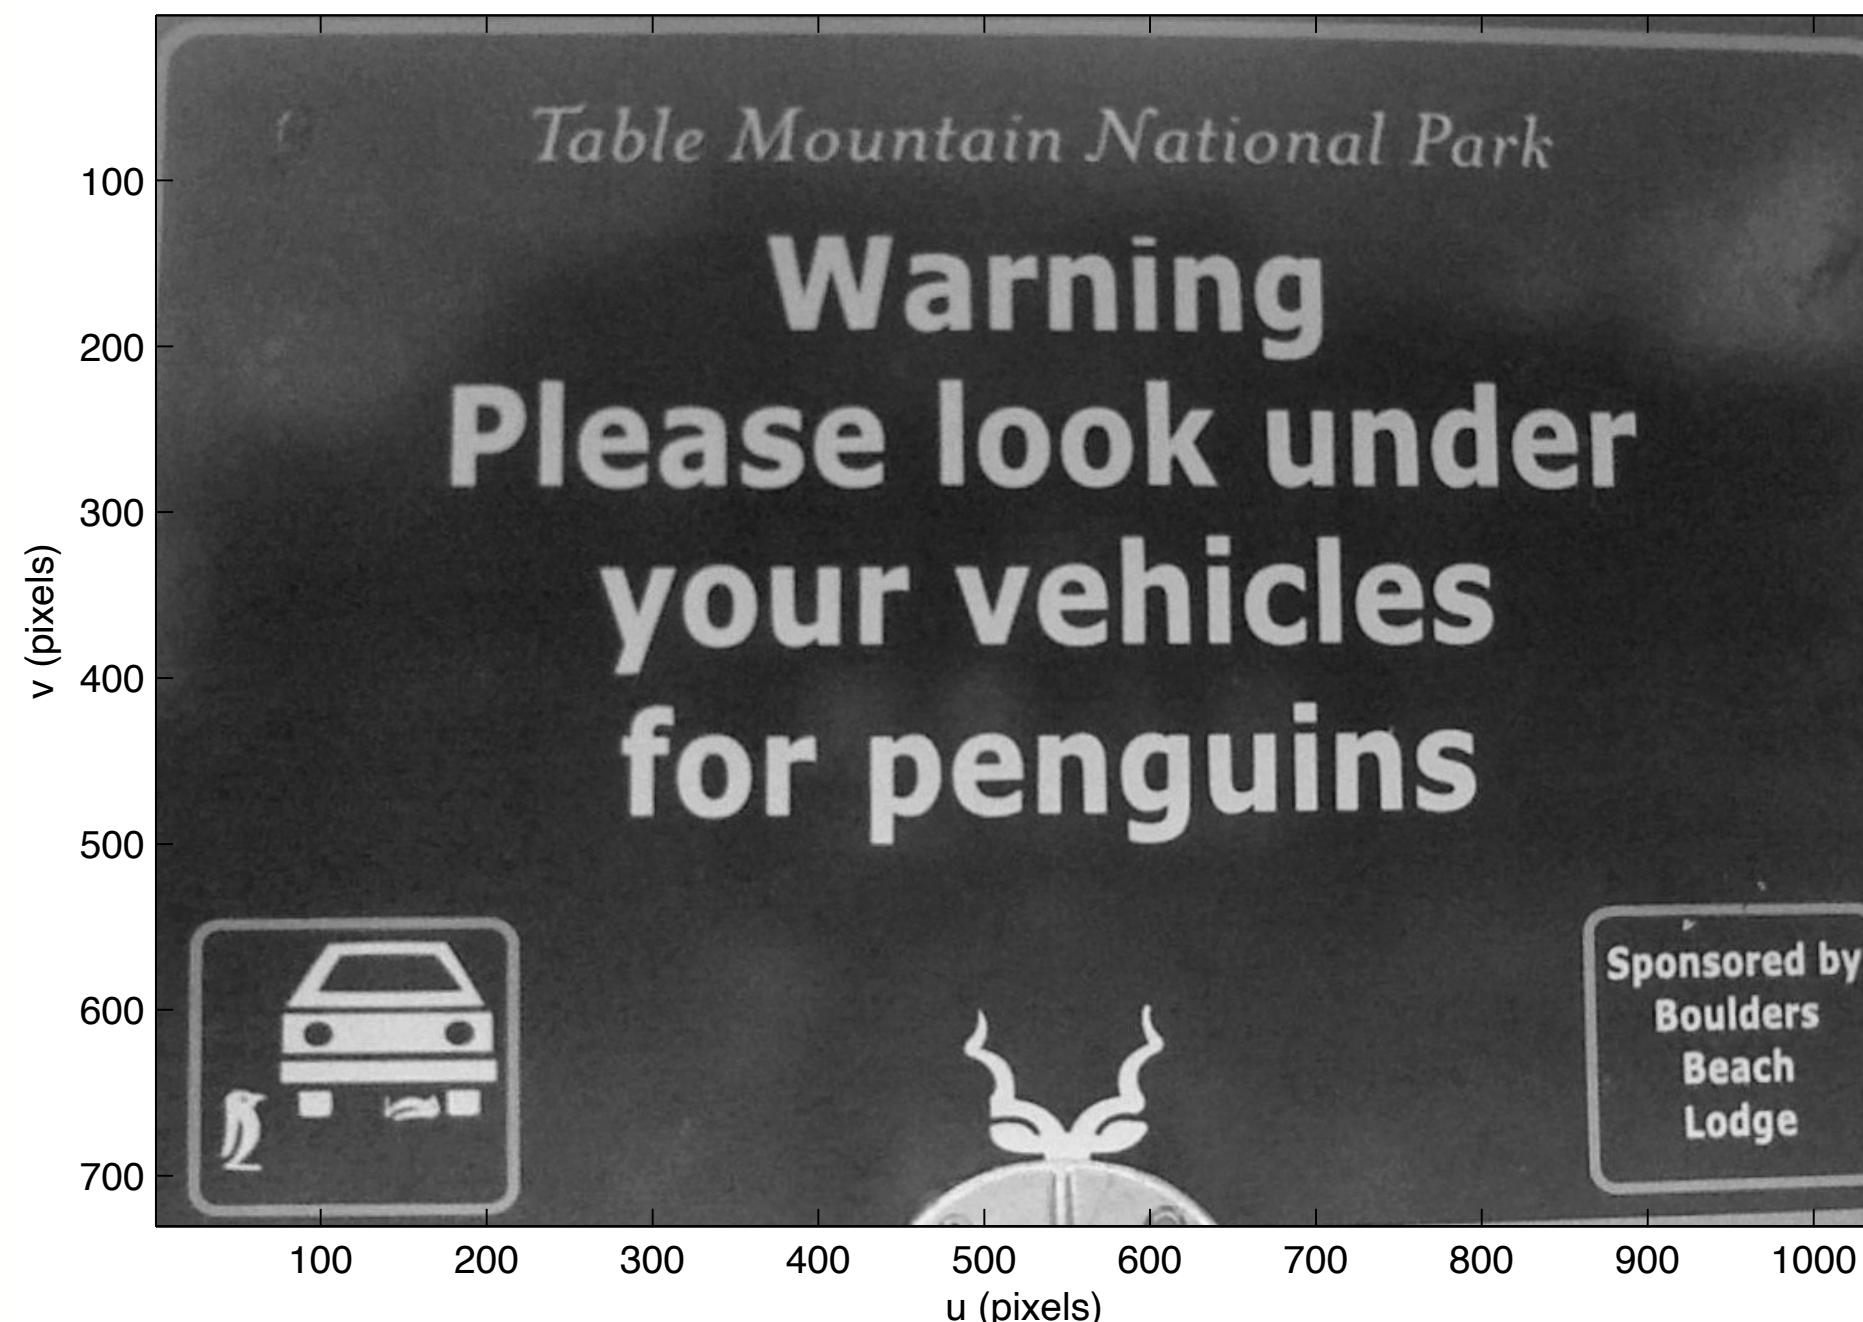
1	0.15	0.21	0.28	0.35	0.41	0.65	0.72	0.79	0.85	0.92
2	0.16	0.23	0.30	0.36	0.43	0.67	0.74	0.80	0.87	0.93
3	0.18	0.25	0.31	0.38	0.44	0.69	0.75	0.82	0.88	0.95
4	0.20	0.26	0.33	0.39	0.46	0.70	0.77	0.84	0.90	0.97
5	0.21	0.28	0.35	0.41	0.48	0.72	0.79	0.85	0.92	0.98
6	0.23	0.30	0.36	0.43	0.49	0.74	0.80	0.87	0.93	1.00
7	0.25	0.31	0.38	0.44	0.51	0.28	0.25	0.21	0.18	0.15
8	0.26	0.33	0.39	0.46	0.53	0.28	0.25	0.21	0.18	0.15
9	0.28	0.35	0.41	0.48	0.54	0.28	0.25	0.21	0.18	0.15
10	0.30	0.36	0.43	0.49	0.56	0.28	0.25	0.21	0.18	0.15

Output image



$$K = \frac{1}{2} \begin{pmatrix} 1 & 0 & -1 \end{pmatrix}$$

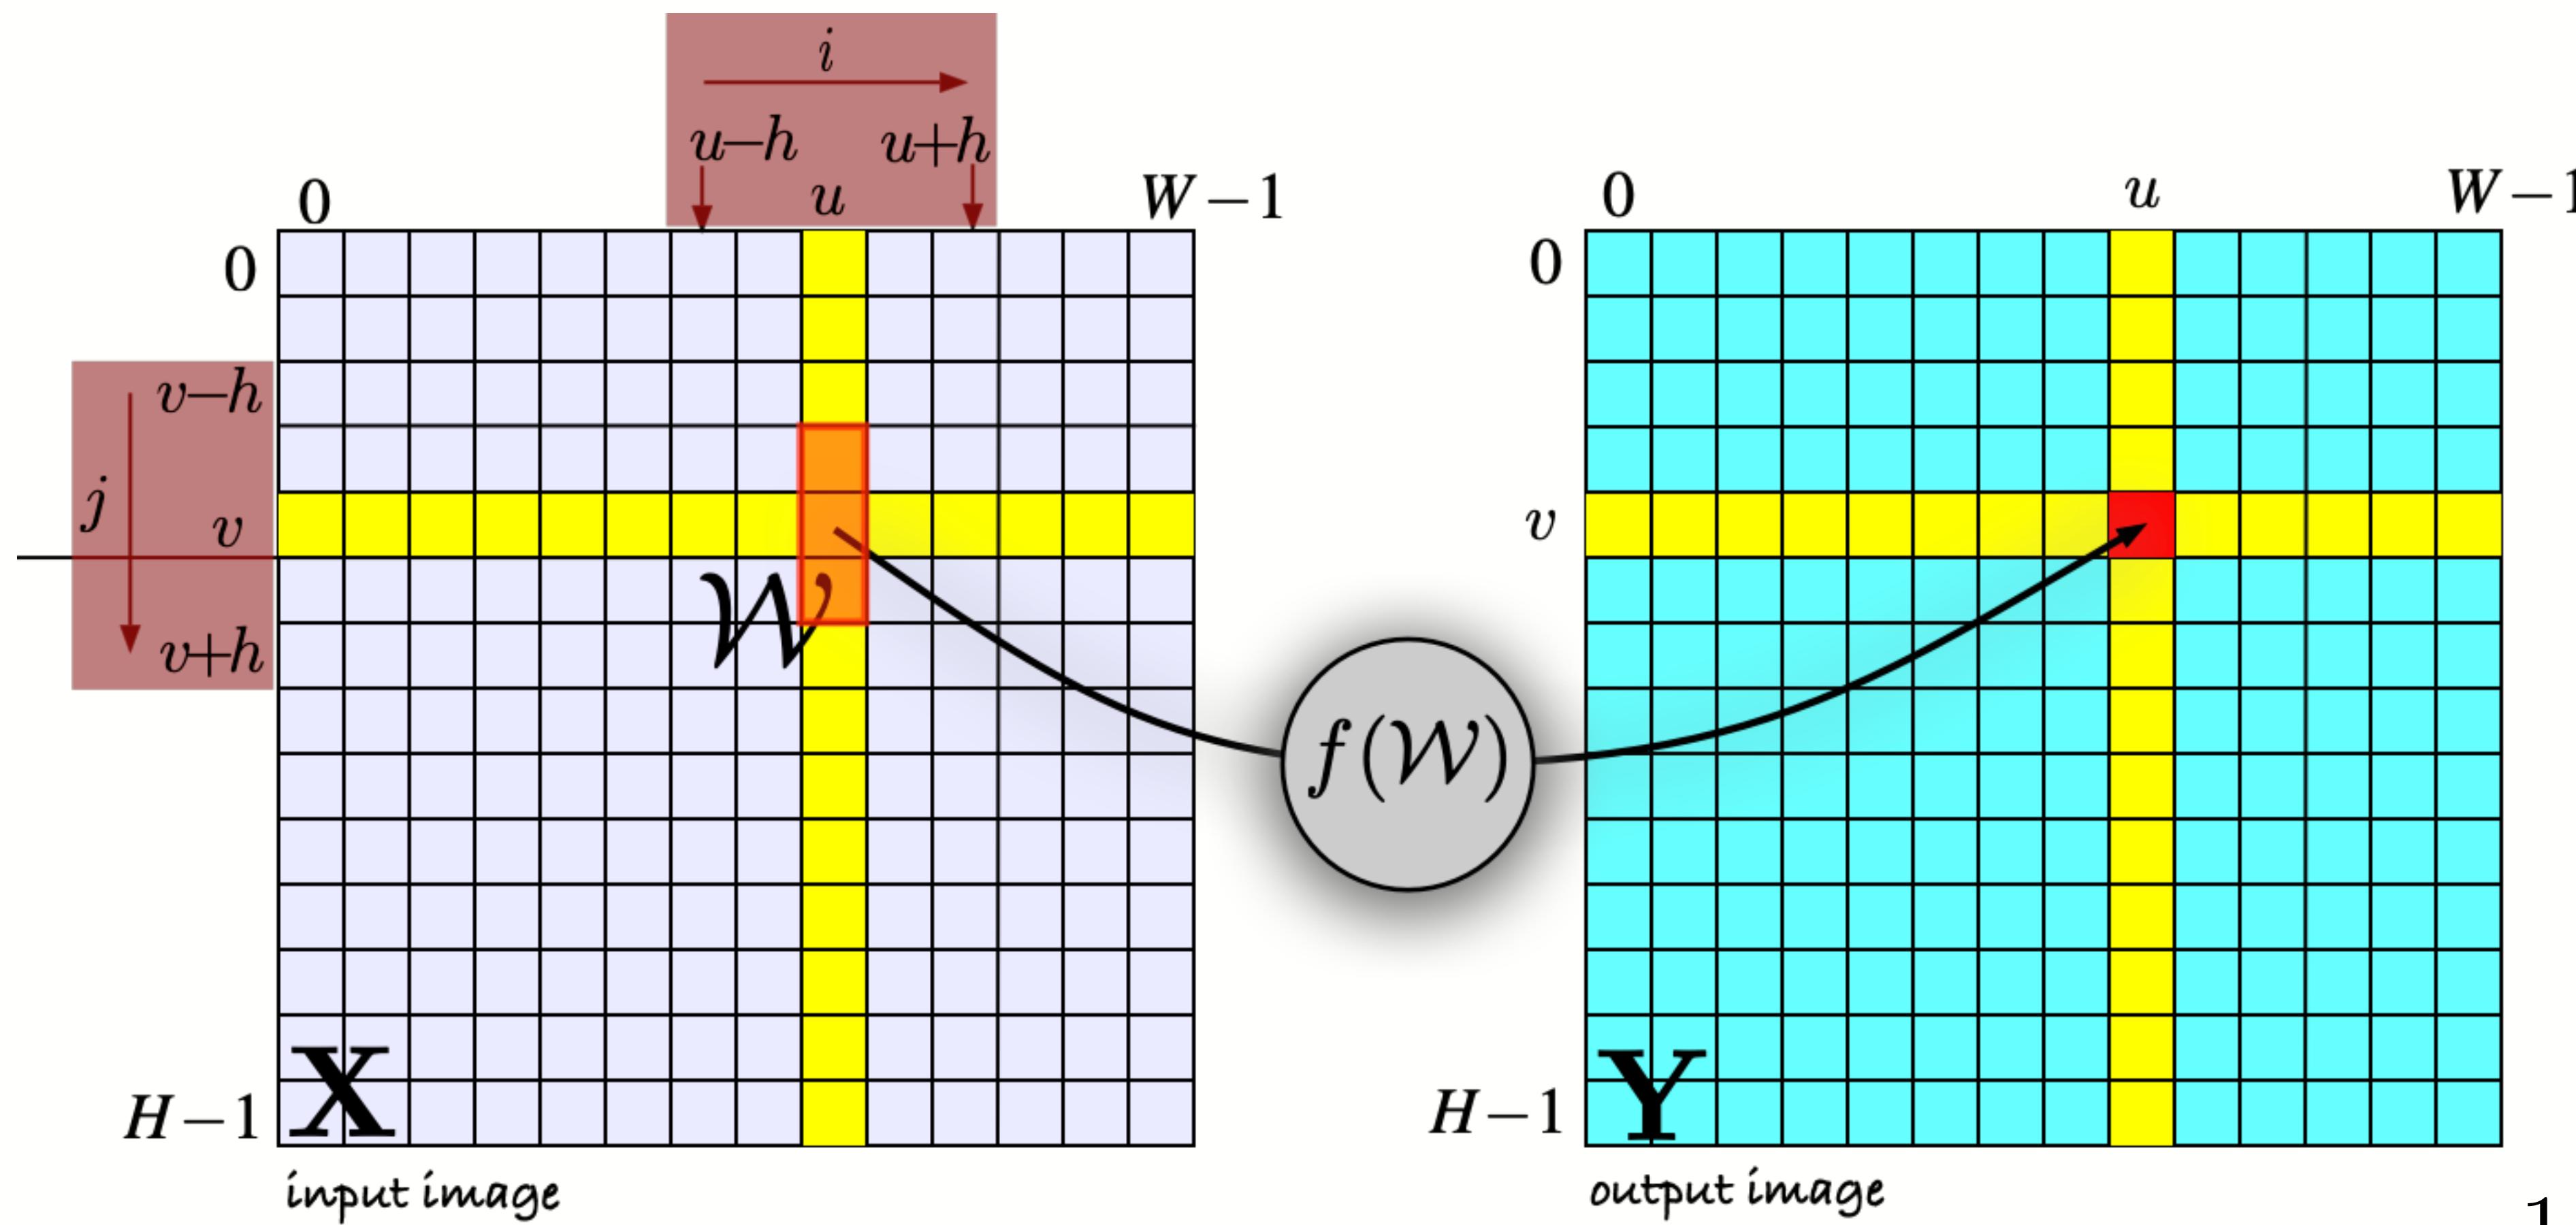
Horizontal gradient image



$$K = \frac{1}{2} \begin{pmatrix} 1 & 0 & -1 \end{pmatrix}$$

```
>>> K = np.c_[1, 0, -1] / 2
>>> im.convolve(K)
```

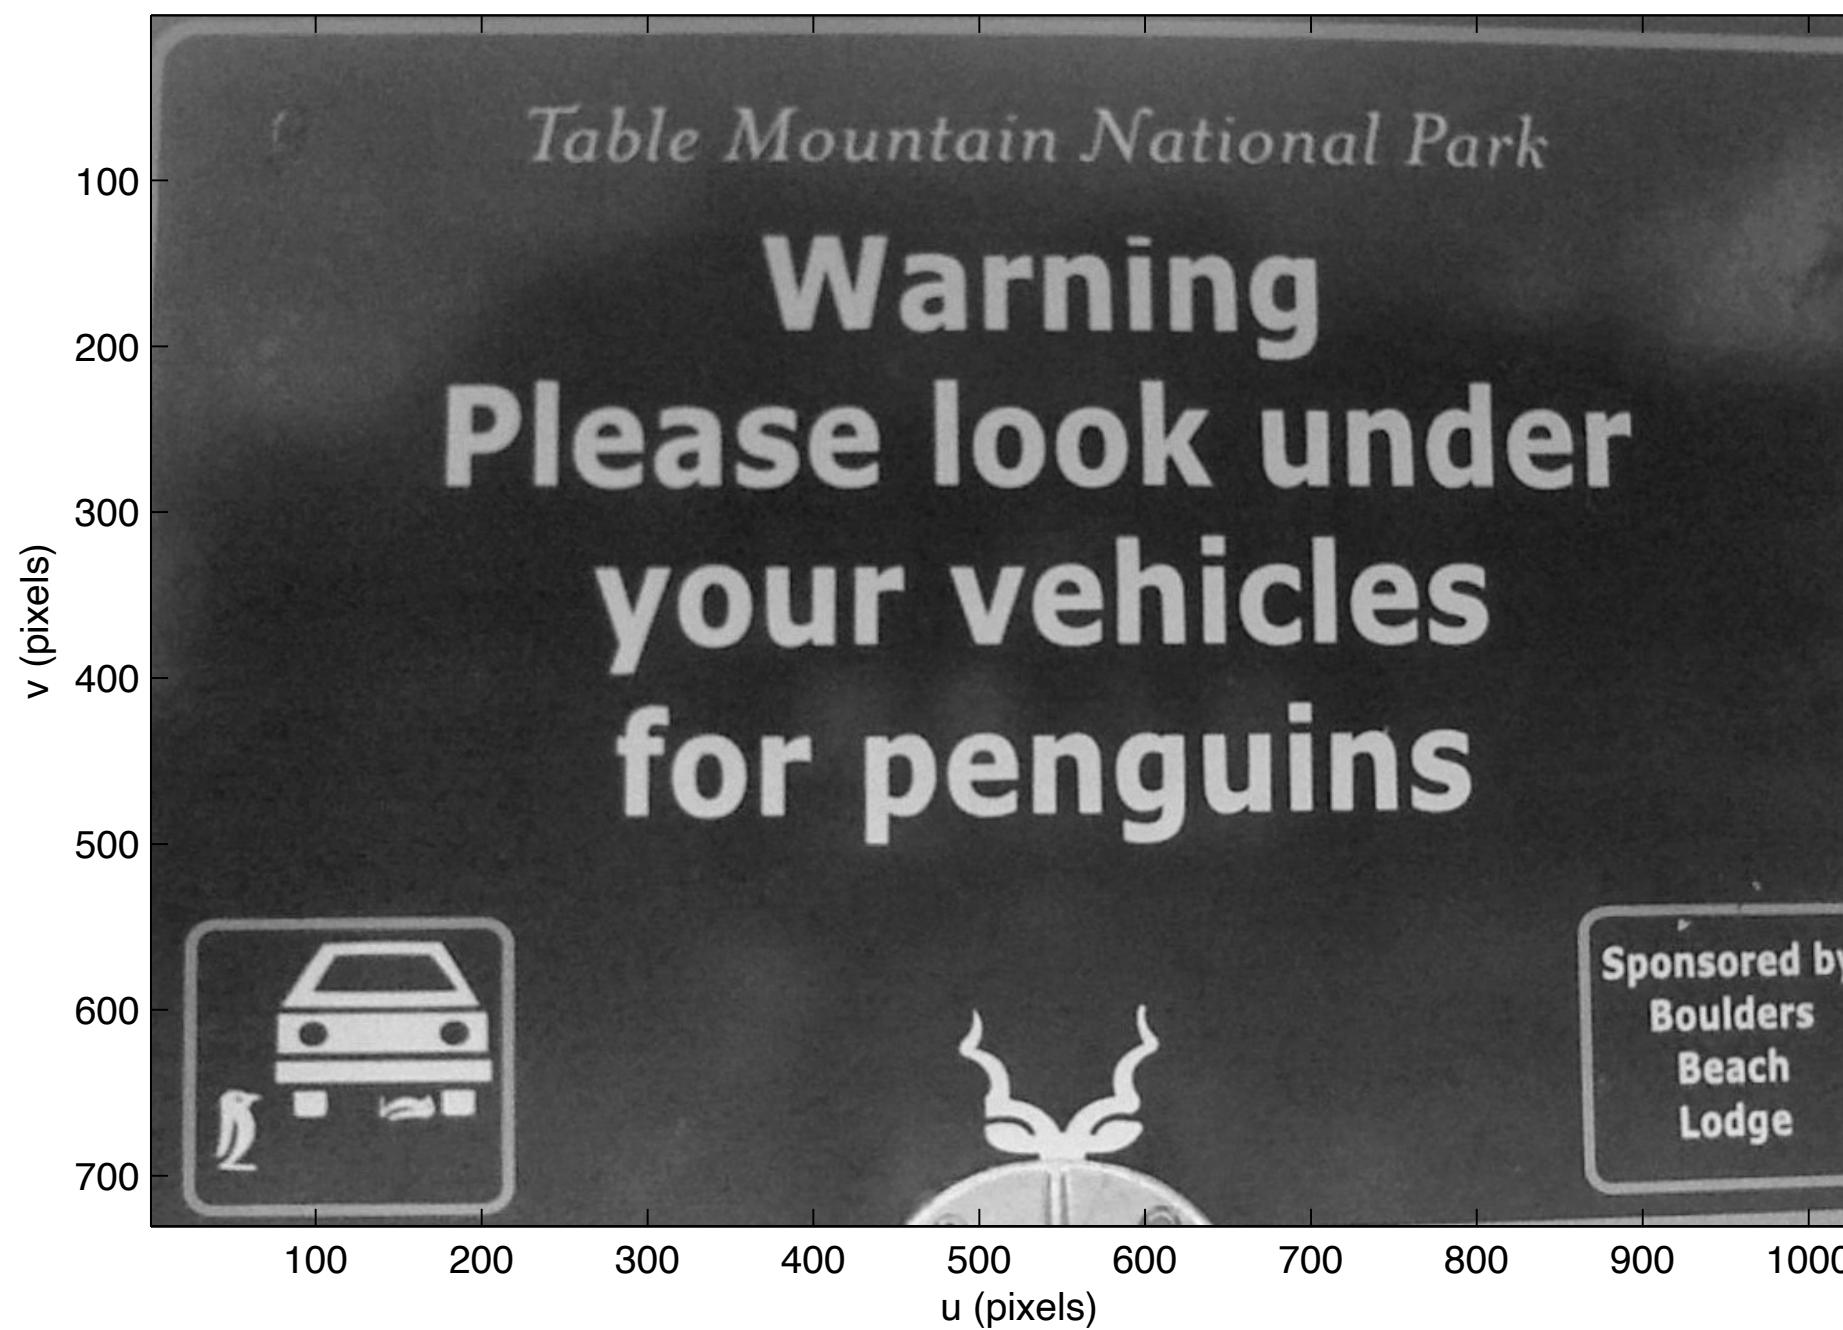
Vertical gradient as correlation with a kernel



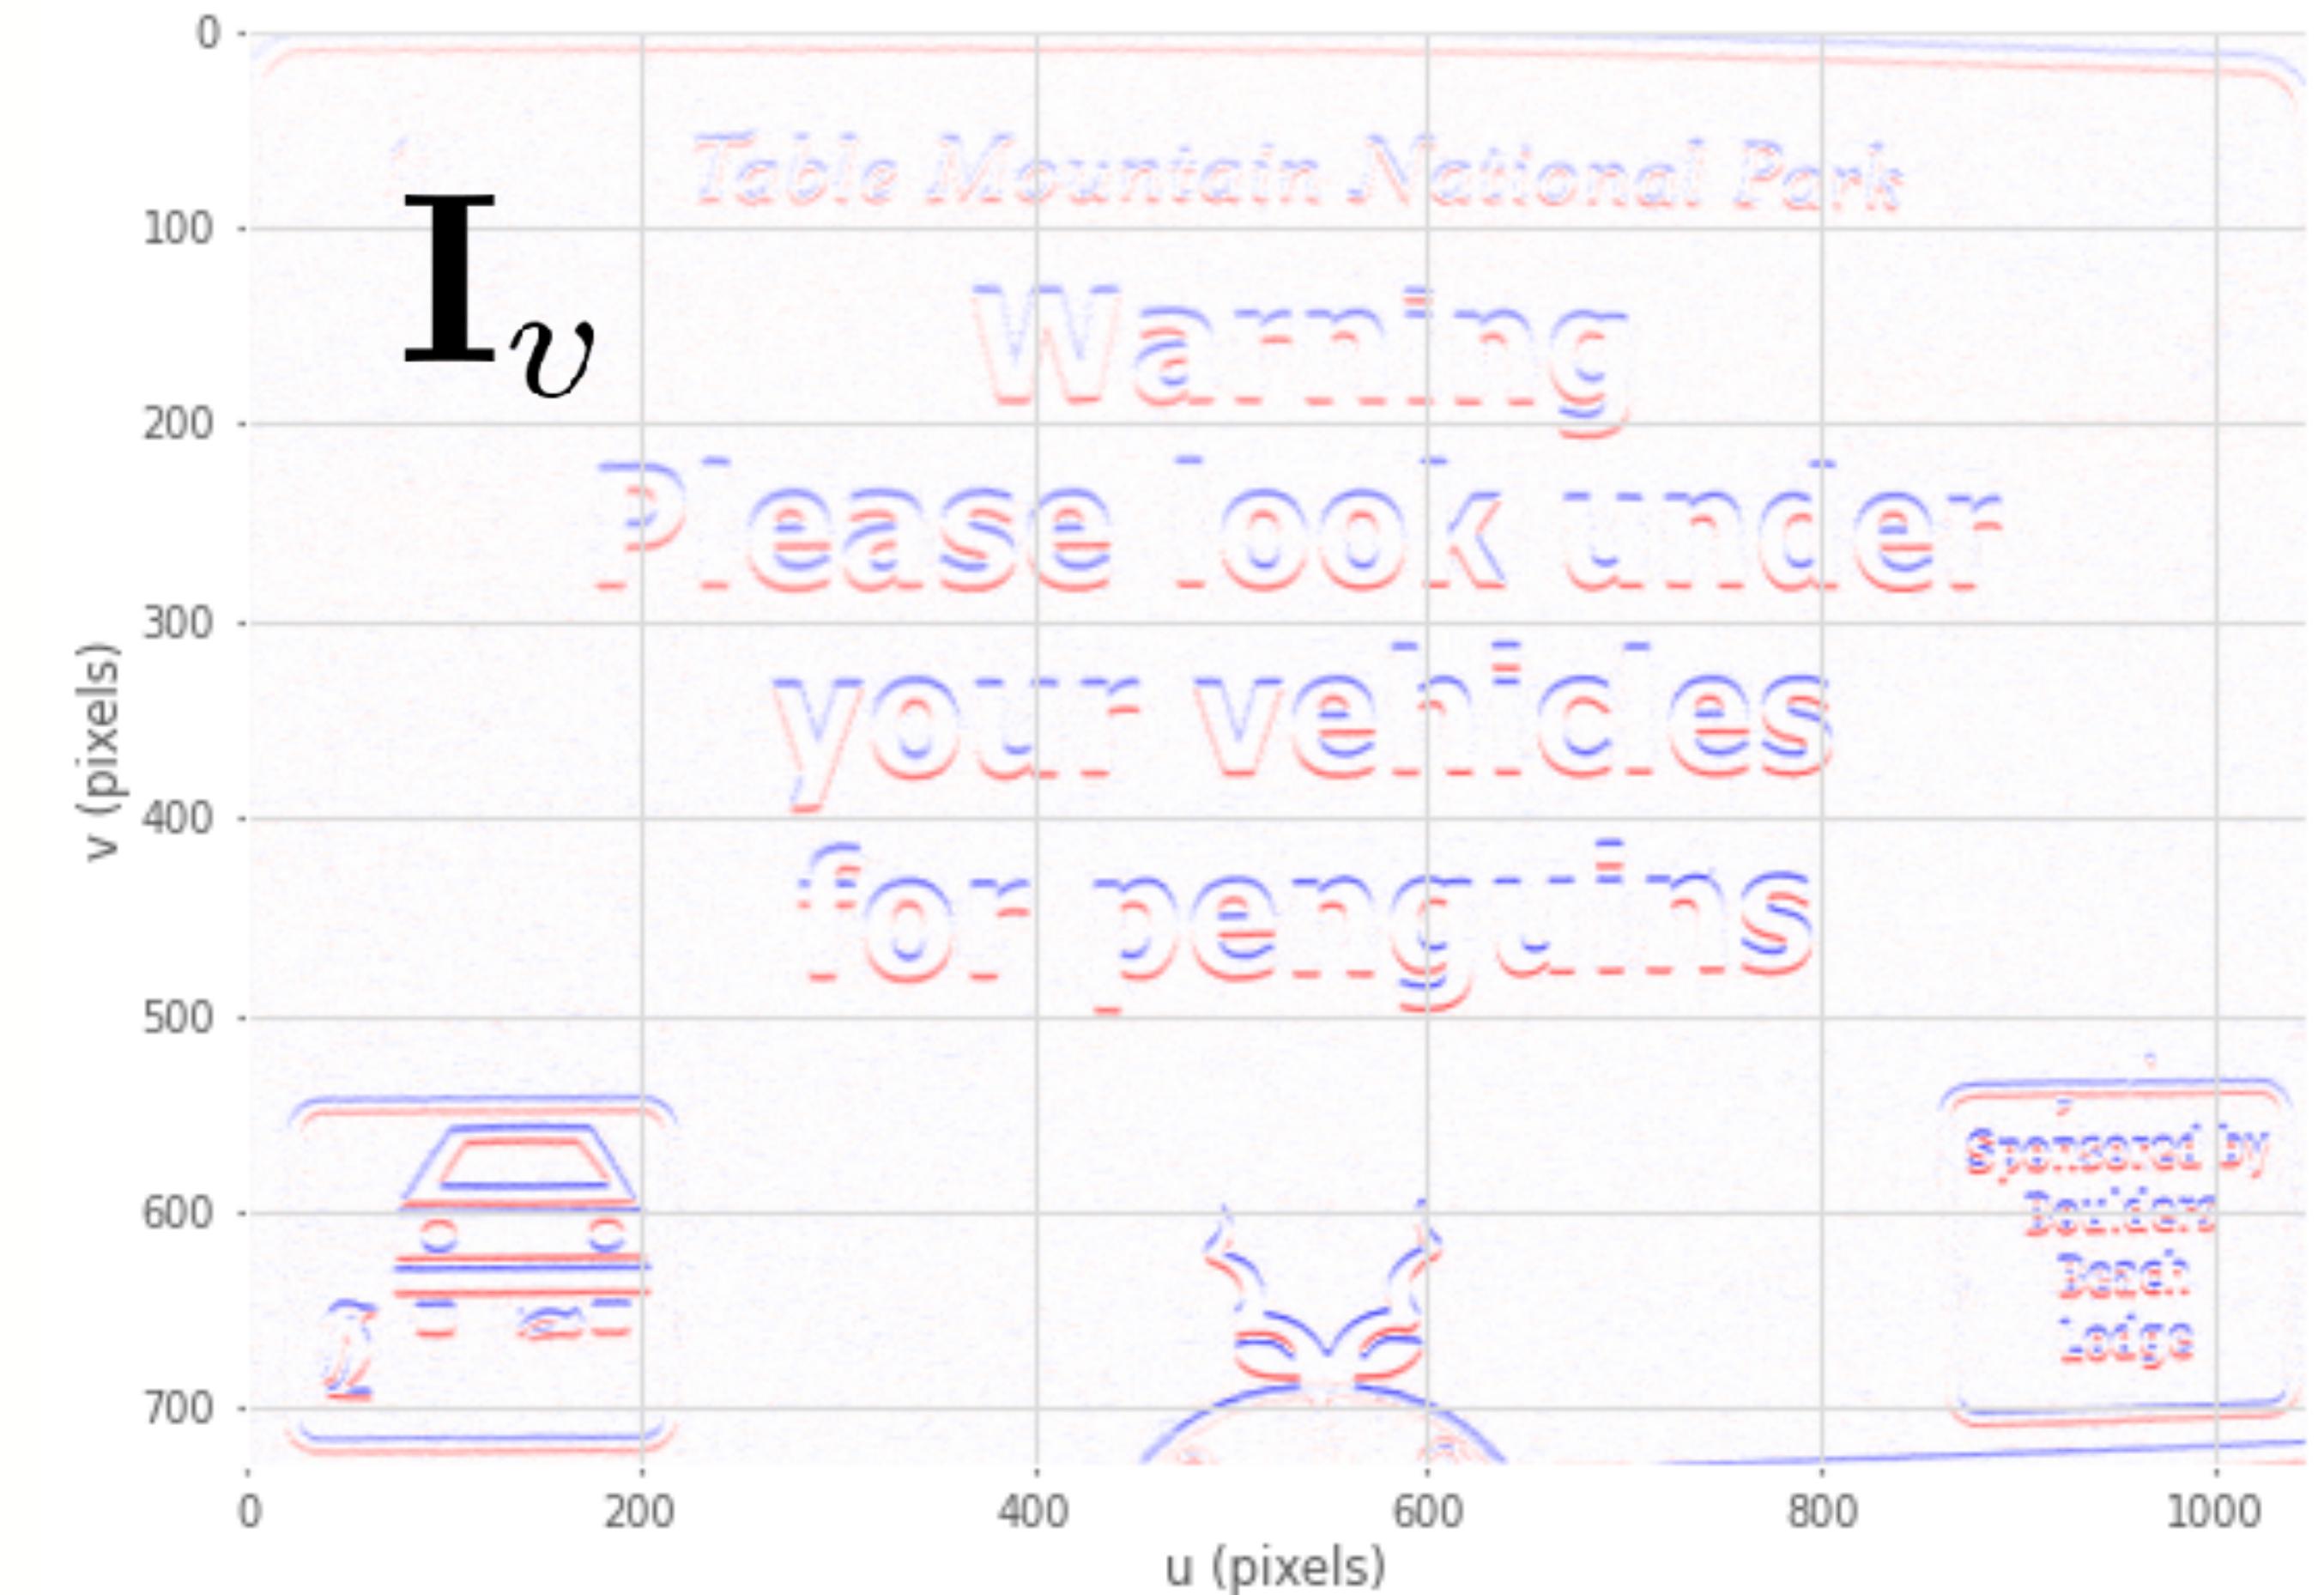
With kind permission of Springer Science+Business Media

$$\mathbf{K} = \frac{1}{2} \begin{pmatrix} 1 \\ 0 \\ -1 \end{pmatrix}$$

Vertical gradient image, use the transposed kernel



$$K = \frac{1}{2} \begin{pmatrix} -1 \\ 0 \\ 1 \end{pmatrix}$$



```
>>> im.convolve(K.T)
```

2D convolution

- Correlation is closely related to convolution

$$\mathbf{O}[u, v] = \sum_{(i,j) \in \mathcal{W}} \mathbf{I}[u - i, v - j] \mathbf{K}[i, j], \forall (u, v) \in \mathbf{I}$$

- Convolution is the same as correlation if the kernel is symmetric**

- smoothing kernels are always symmetric
 - derivative kernels are anti-symmetric
- Often written in operator form $\mathbf{O} = \mathbf{K} \otimes \mathbf{I}$

negative signs

Properties of convolution

- Commutative

$$A \otimes B = B \otimes A$$

- Associative

$$A \otimes B \otimes C = (A \otimes B) \otimes C = A \otimes (B \otimes C)$$

- Distributive

$$A \otimes (\alpha B) = \alpha(A \otimes B)$$

- and Linear

$$A \otimes (B + C) = A \otimes B + A \otimes C$$

Advantages of associativity

- Convolving an image with a Gaussian kernel twice

$$\mathbf{I} = \mathbf{G} \otimes (\mathbf{G} \otimes \mathbf{I})$$

- Is the same as convolving the image with a kernel that is the Gaussian convolved with itself

$$\begin{aligned}\mathbf{I} &= (\mathbf{G} \otimes \mathbf{G}) \otimes \mathbf{I} \\ &= \mathbf{G}' \otimes \mathbf{I}\end{aligned}$$

Properties of Gaussians

$$\mathbf{G}(\sigma_1) \otimes \mathbf{G}(\sigma_2) = \mathbf{G}(\sqrt{\sigma_1^2 + \sigma_2^2})$$

$$\mathbf{G}(\sigma) \otimes \mathbf{G}(\sigma) = \mathbf{G}(\sqrt{2}\sigma)$$

Derivative with smoothing

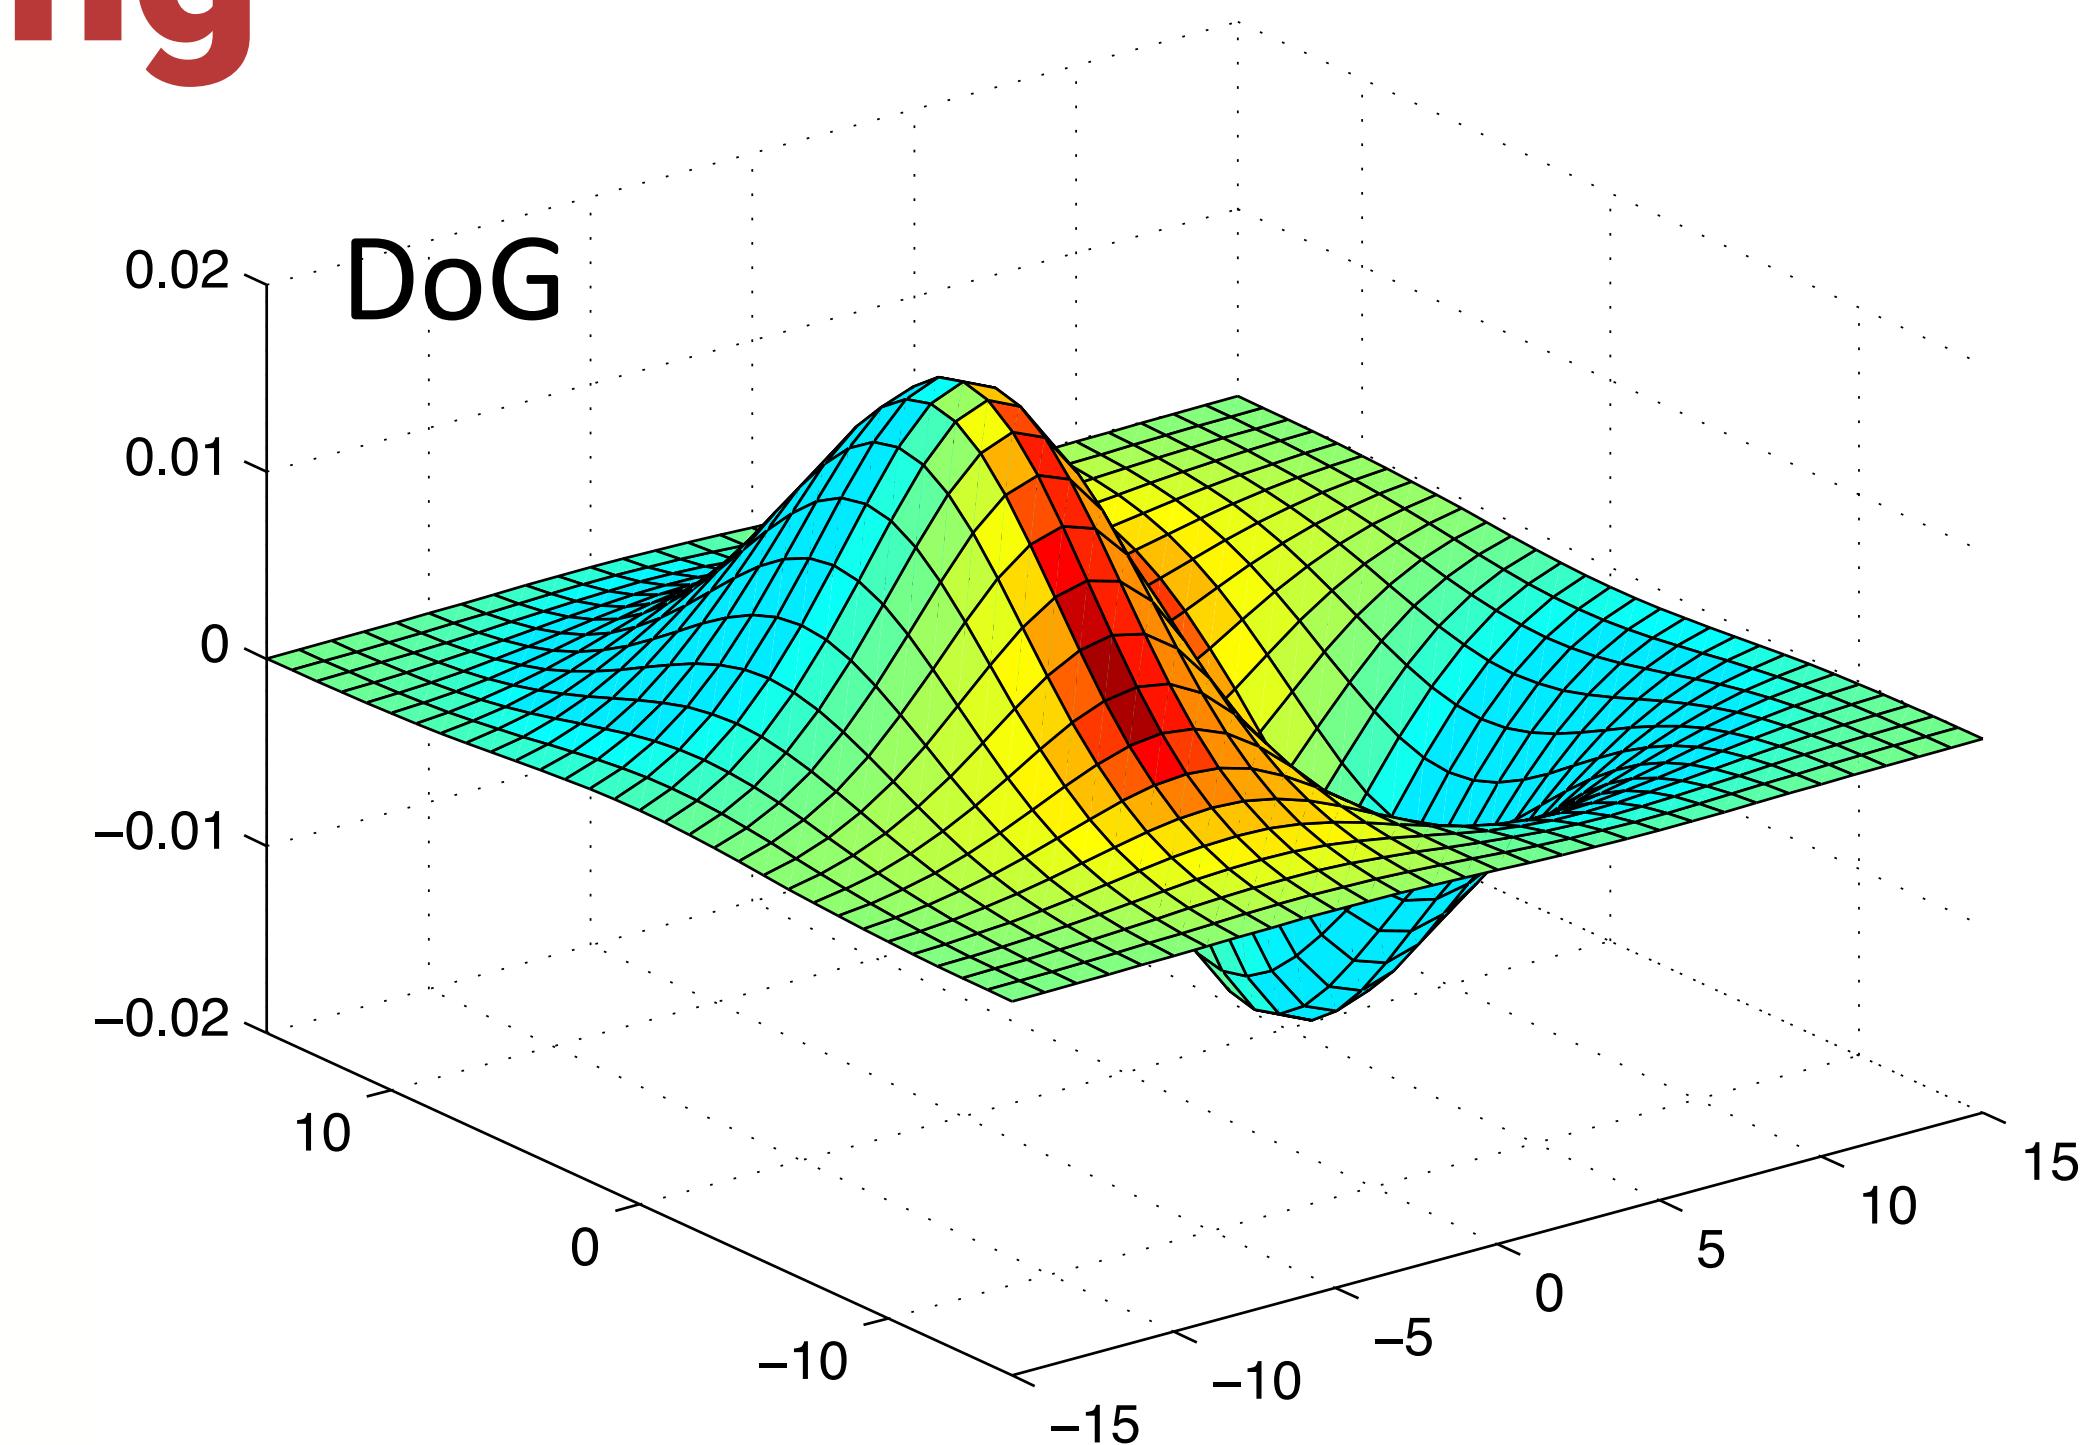
With kind permission of Springer Science+Business Media

- Convolve image with the derivative kernel \mathbf{D}

$$\nabla \mathbf{I} = \mathbf{D} \otimes \mathbf{I}$$

- If we smooth the image first then

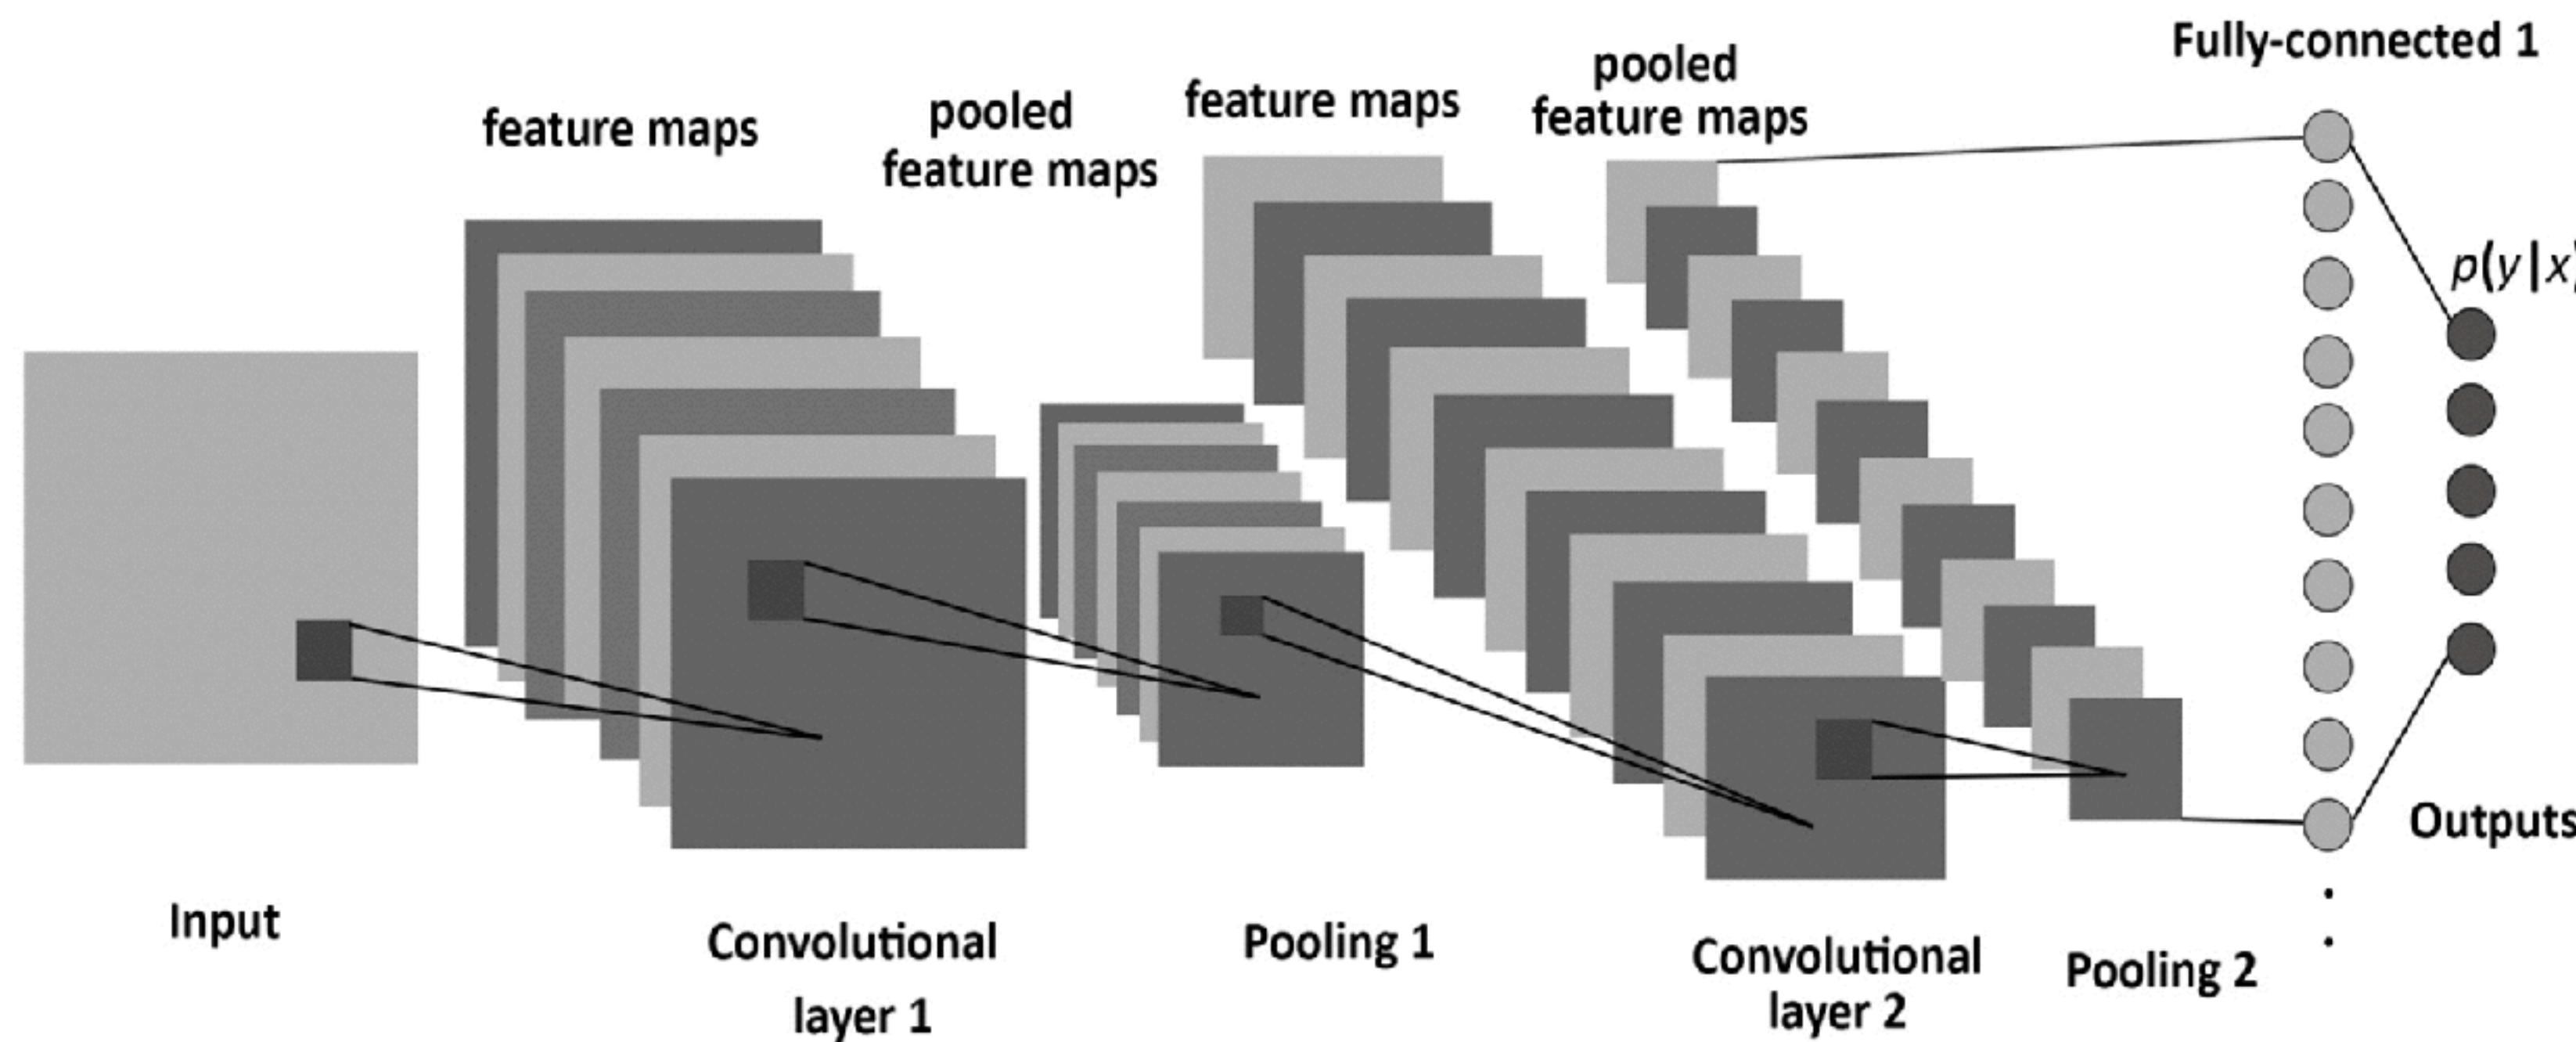
$$\nabla \mathbf{I} = \mathbf{D} \otimes (\mathbf{G}(\sigma) \otimes \mathbf{I}) = \underbrace{(\mathbf{D} \otimes \mathbf{G}(\sigma))}_{DoG} \otimes \mathbf{I}$$



$$\mathbf{G}_u(u, v) = -\frac{u}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

```
>>> Image.kdgauss(5, 15)
```

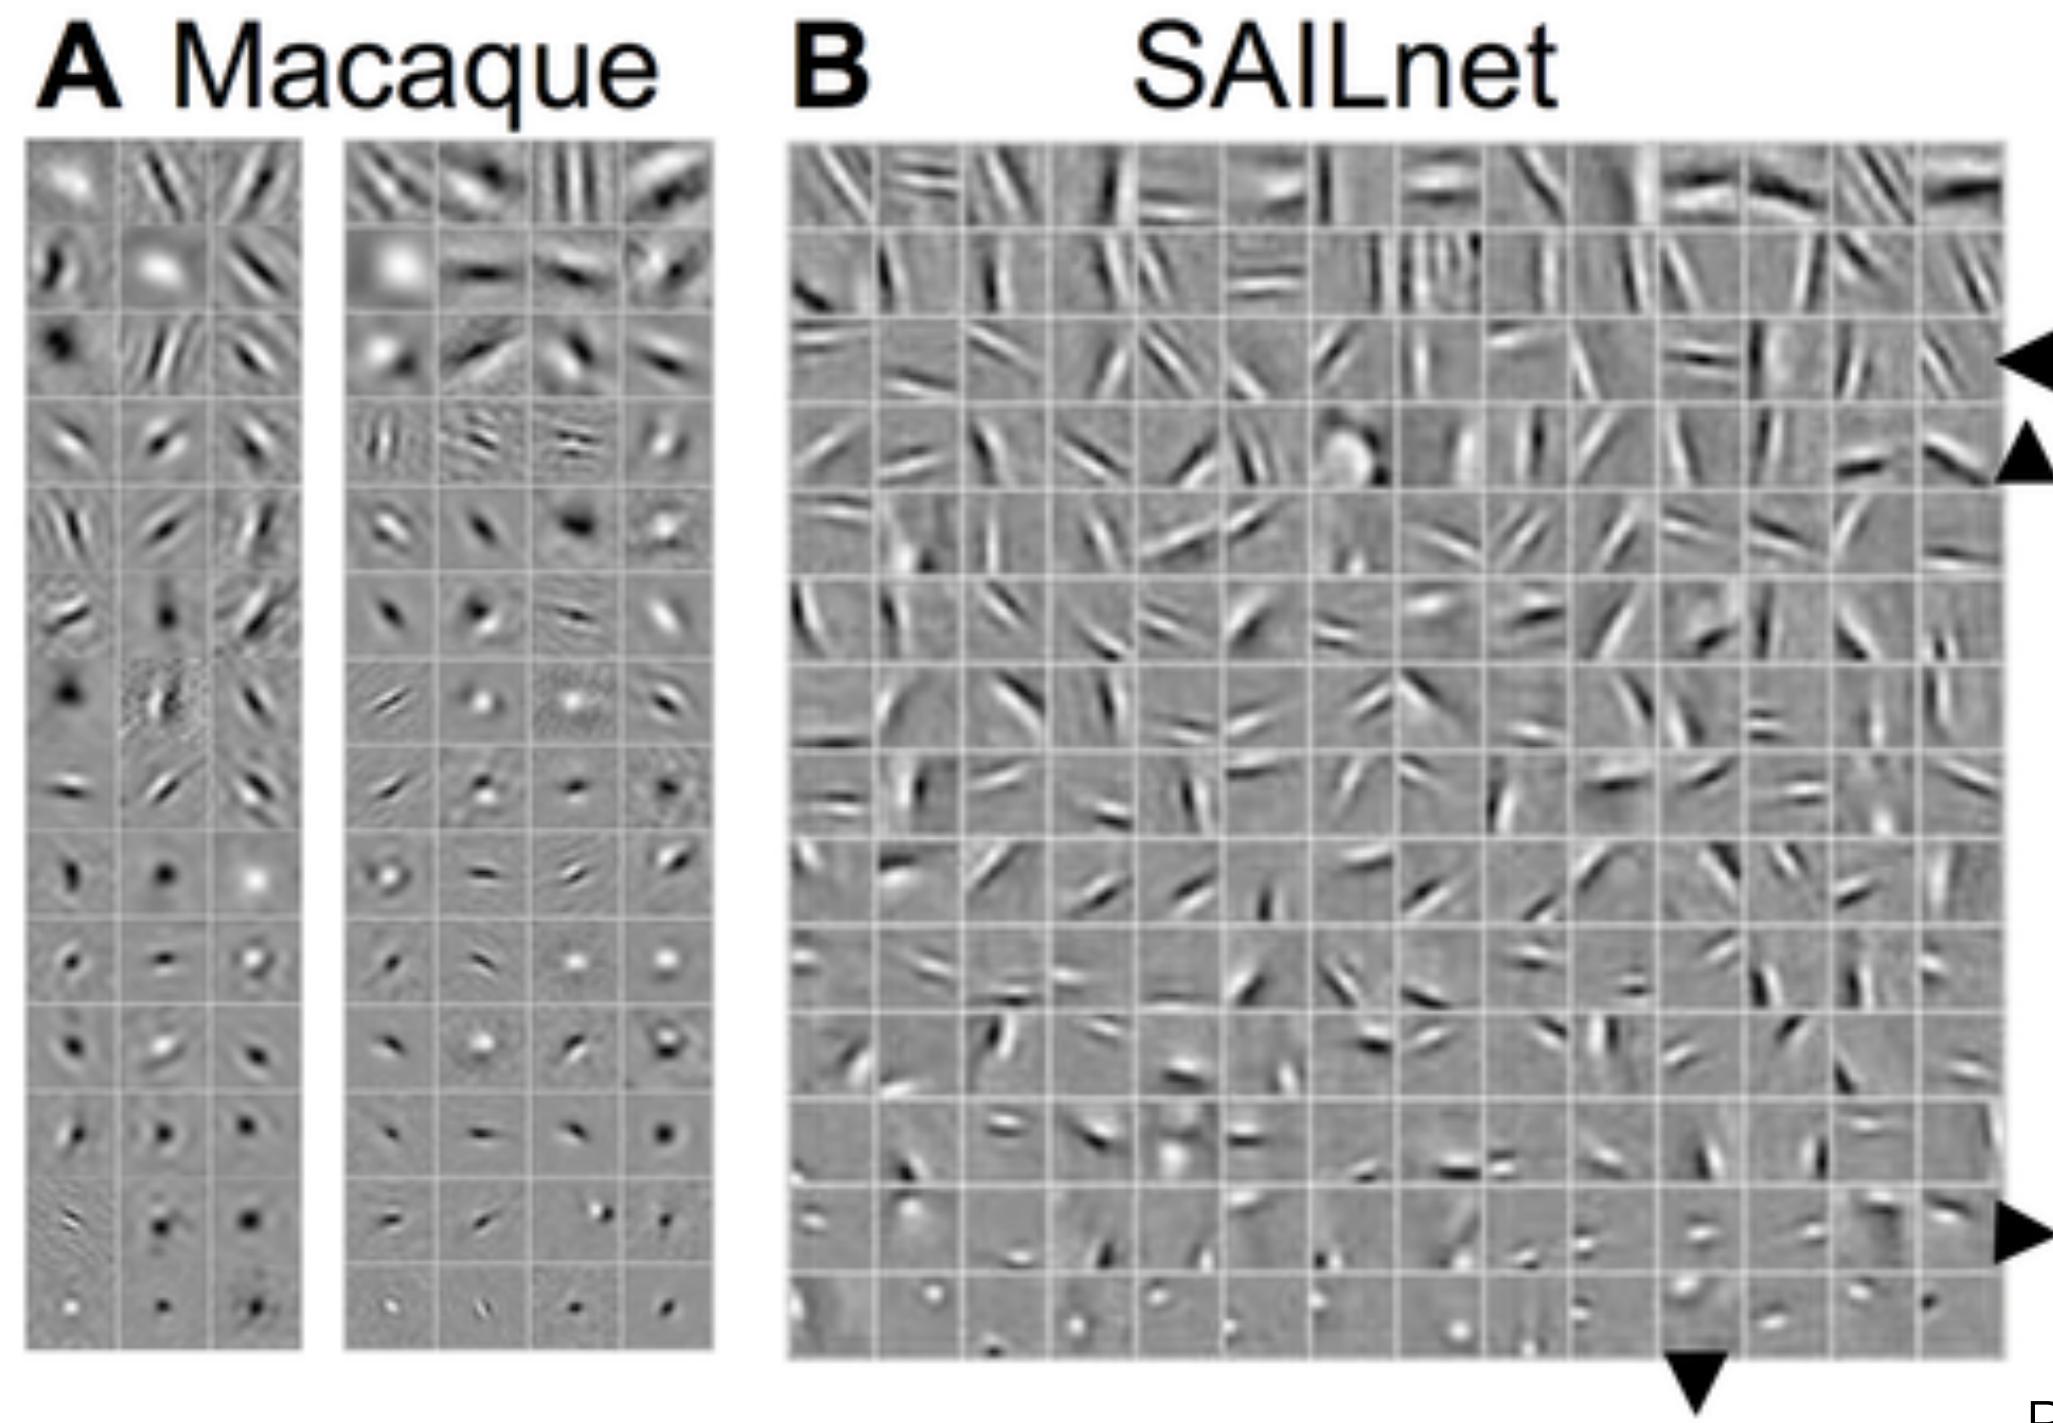
Convolutional neural networks



Entropy 2017, 19(6), 242; <https://doi.org/10.3390/e19060242>

- The input image is convolved with various learnt kernels to form each feature map:
 - a non-linear activation function (often ReLU which is $\max(0, x)$) is applied
- Pooling is average or max over the window
 - the window moves in steps greater than 1 to reduce resolution

Convolutional neural networks

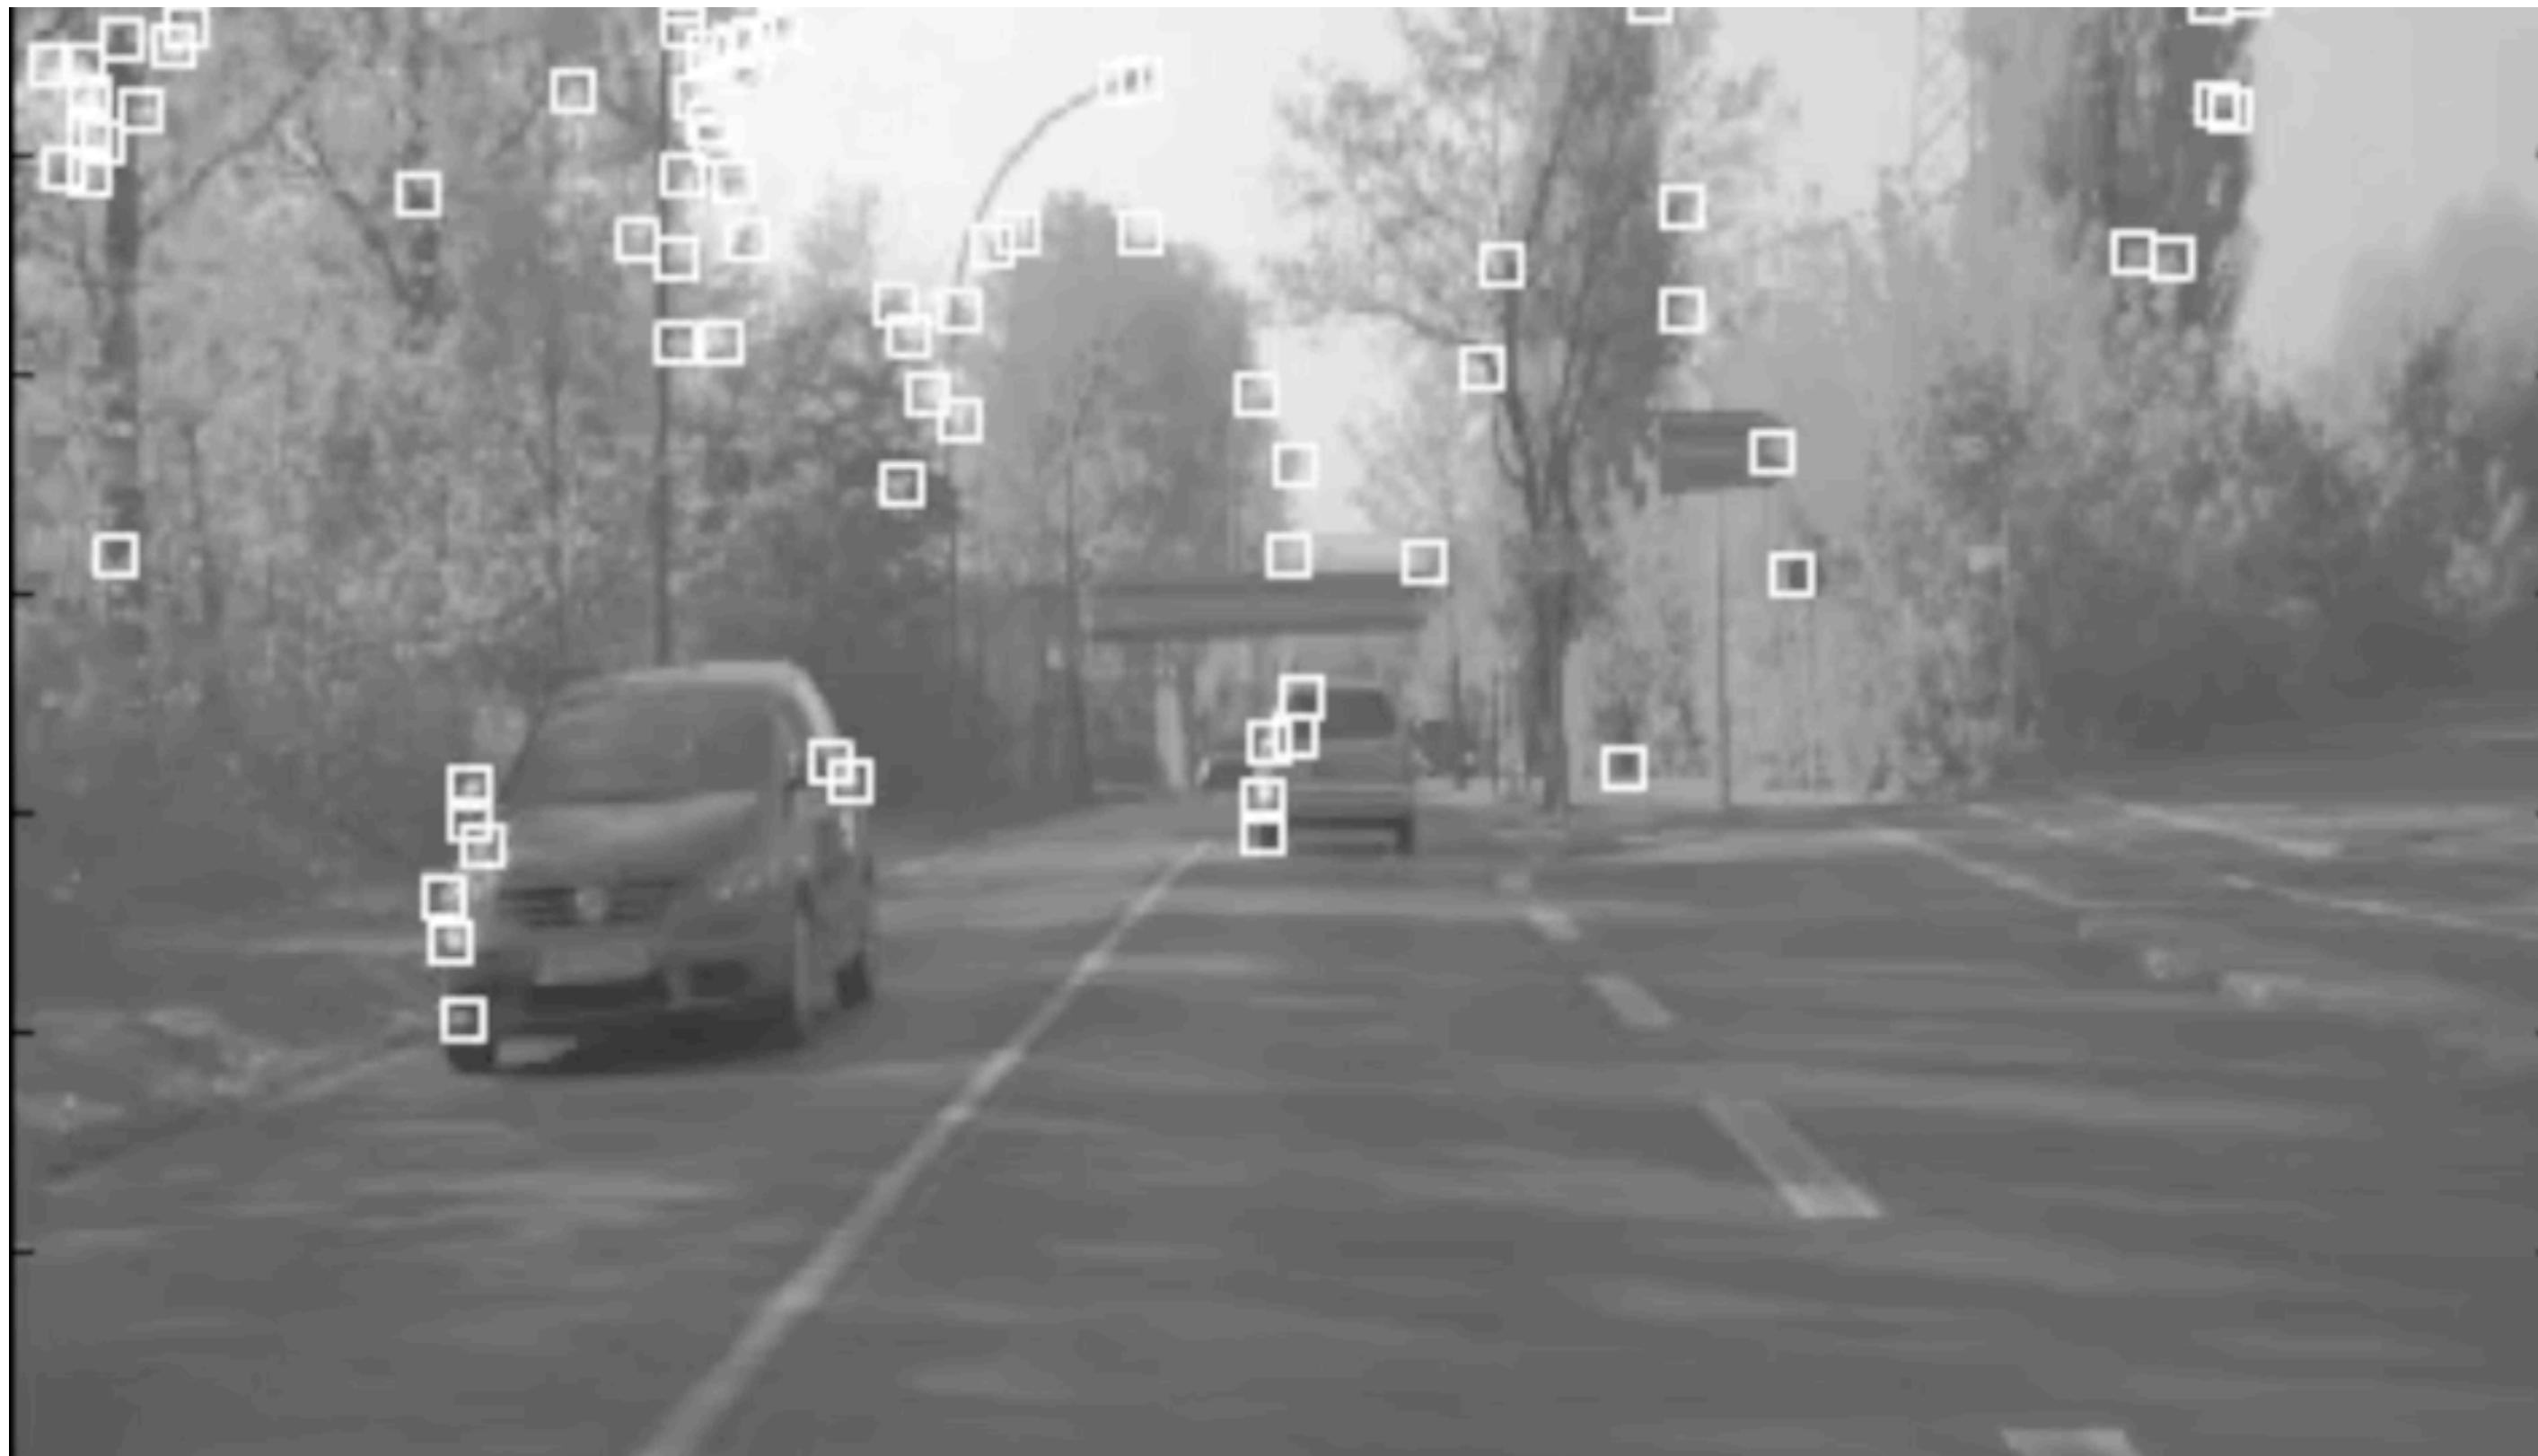


PLoS computational biology. 7. e1002250. 10.1371/journal.pcbi.1002250

- The kernels are tuned to respond to edges (at various angles) and simple light/dark patterns
- The first-level feature maps have filters remarkably similar to those of primate visual cortex (V1)

Image point features

We often want to find the same points in another image



- between different views captured at the same time
- across consecutive views

Finding corresponding points

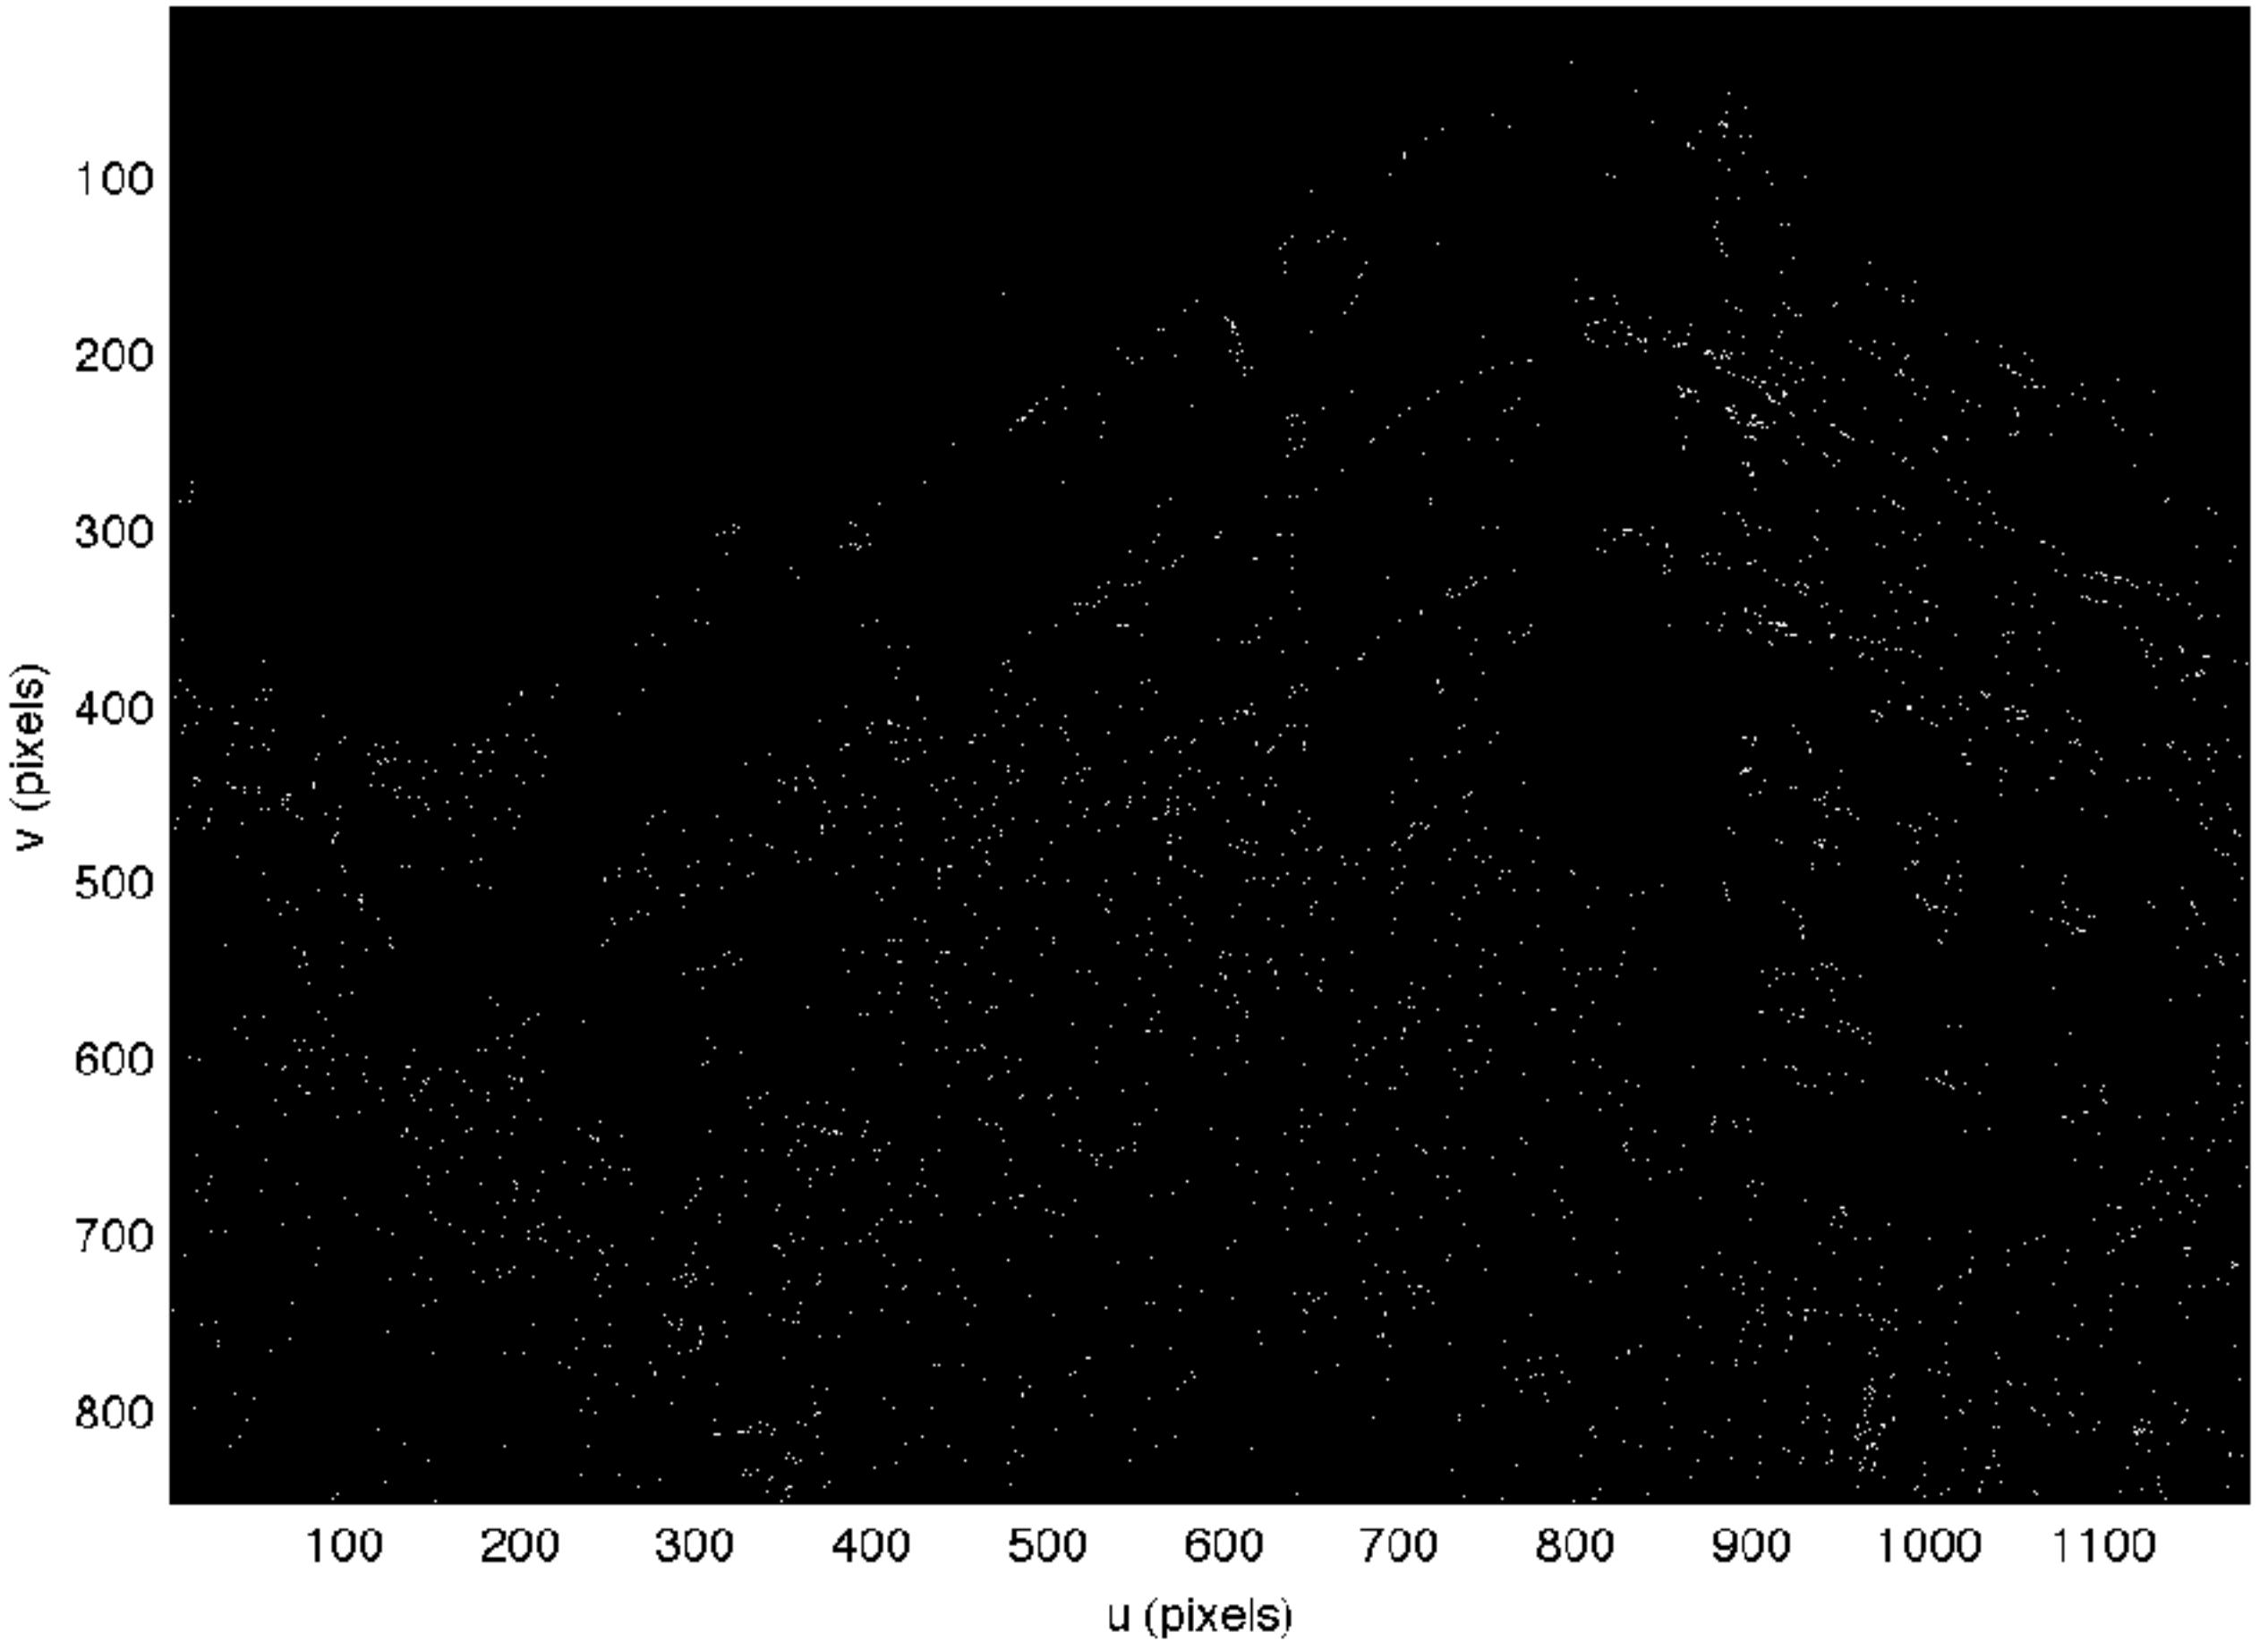


(818,164)

the pixel value is **51**



(619,184)



- All the pixels with a value of 51 in second view
 - 6017 points (0.6% of the image)

Finding corresponding points



(818,164)

pixel value is **51**



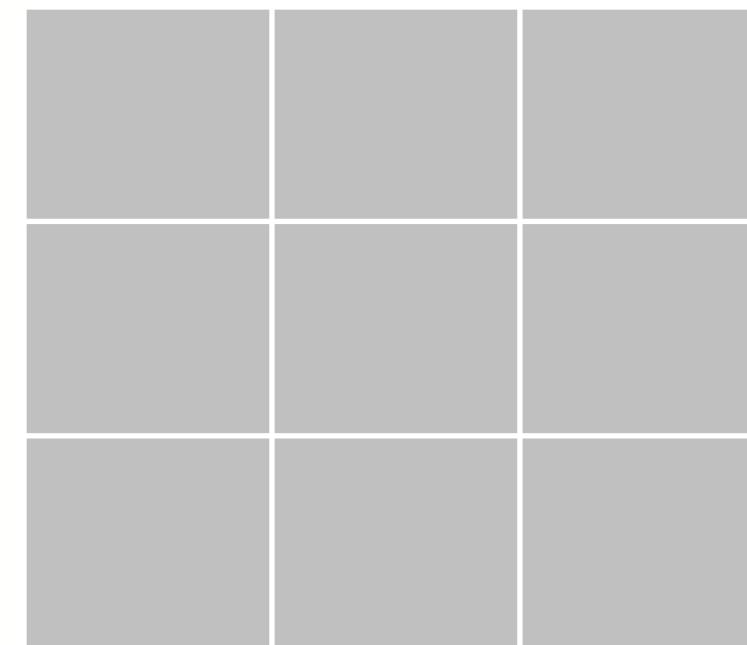
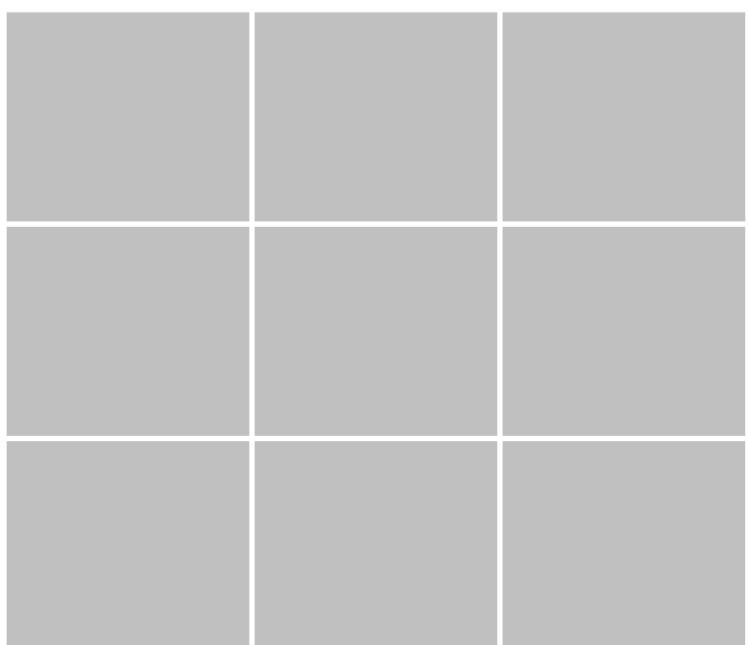
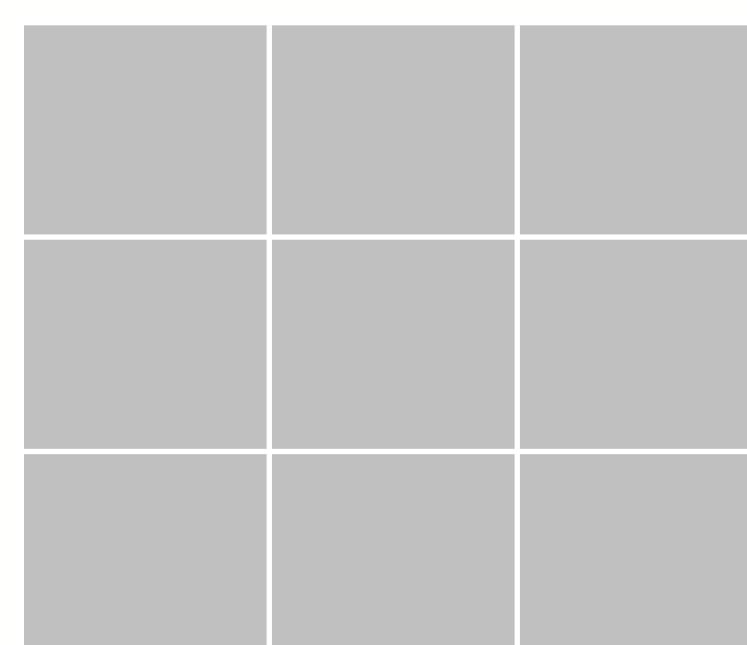
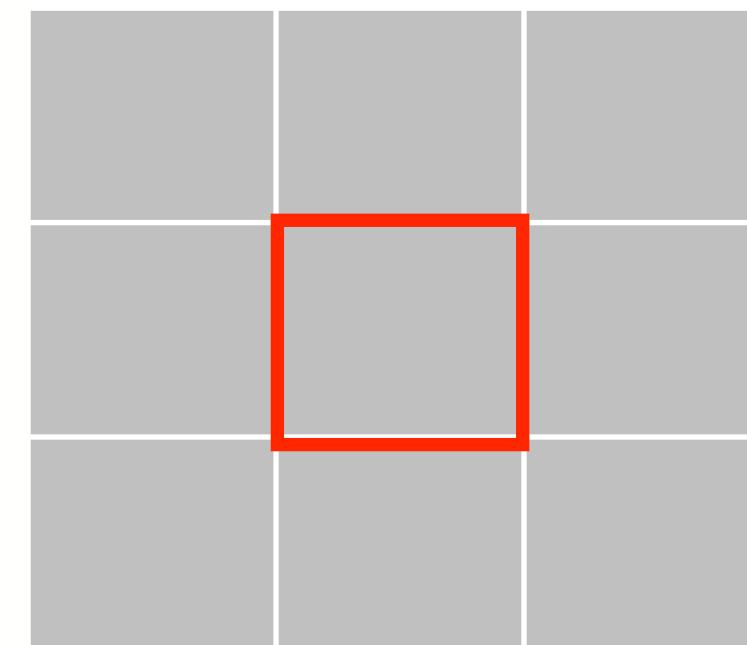
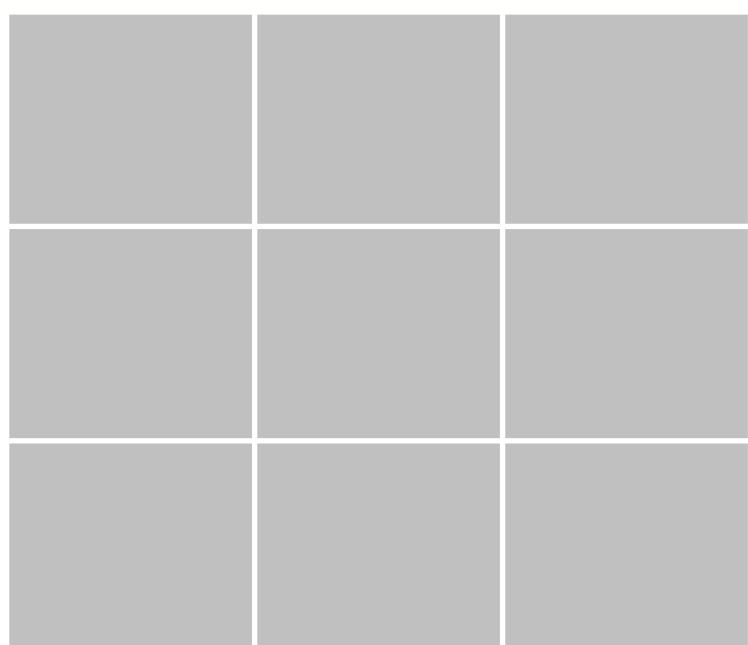
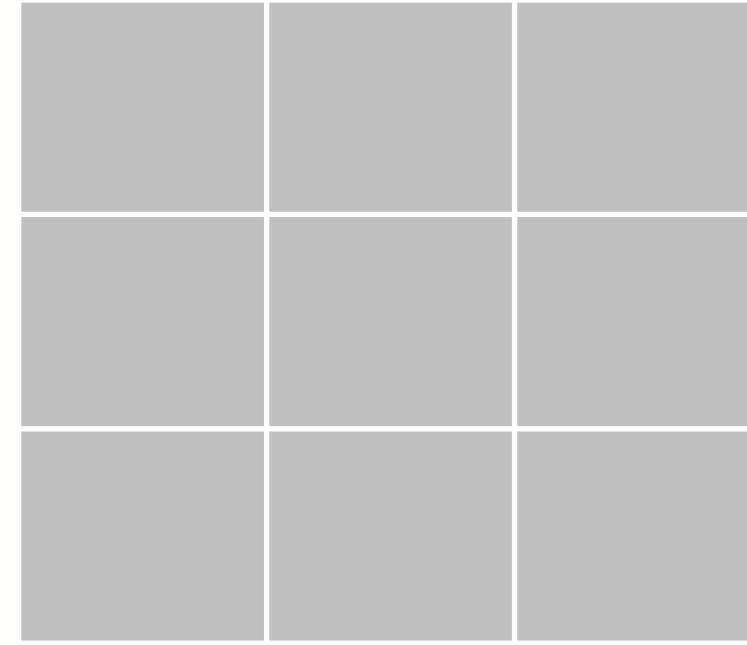
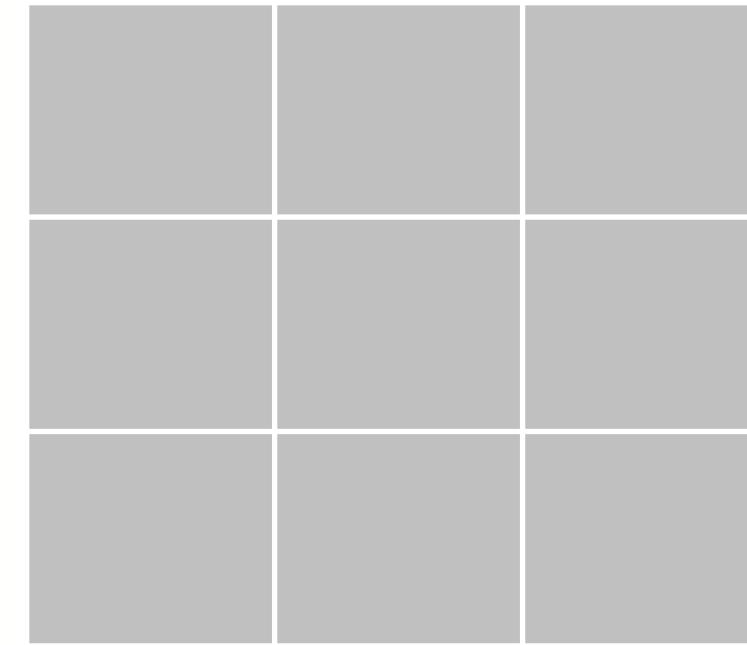
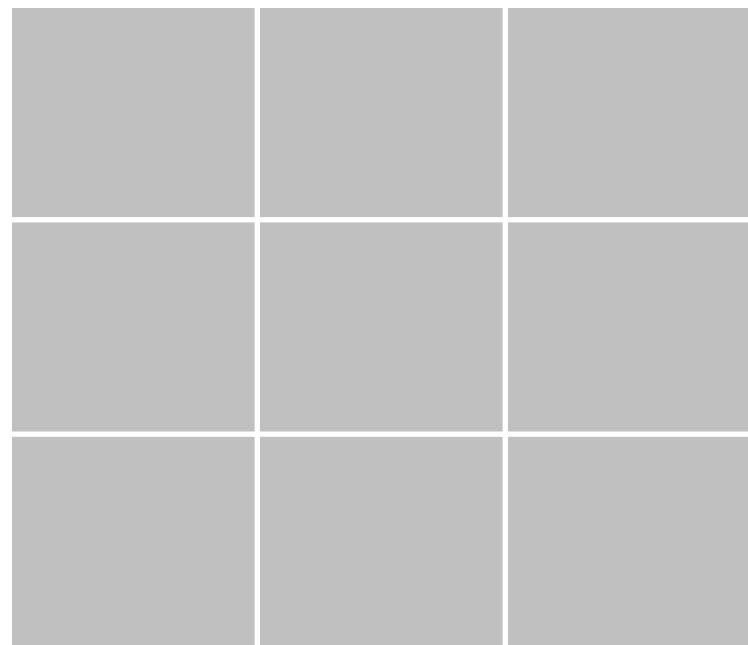
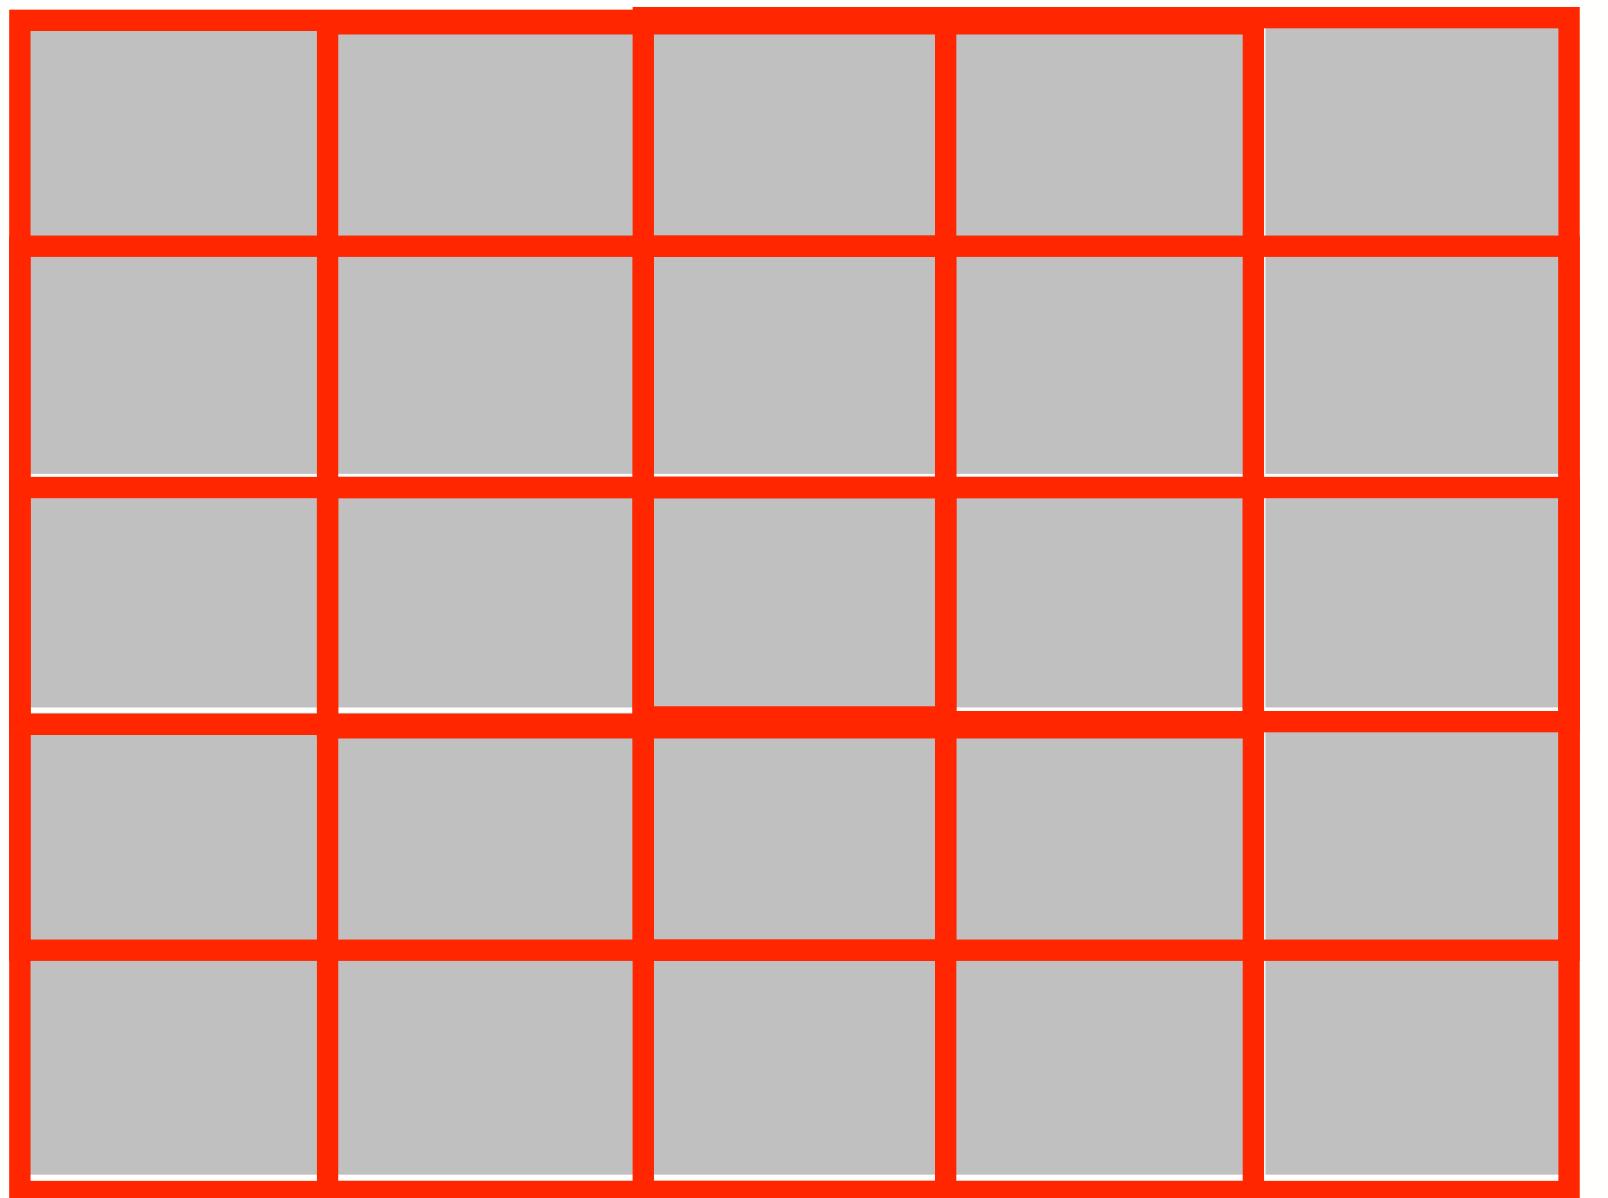
(619,184)

pixel value is **47**

Finding corresponding points

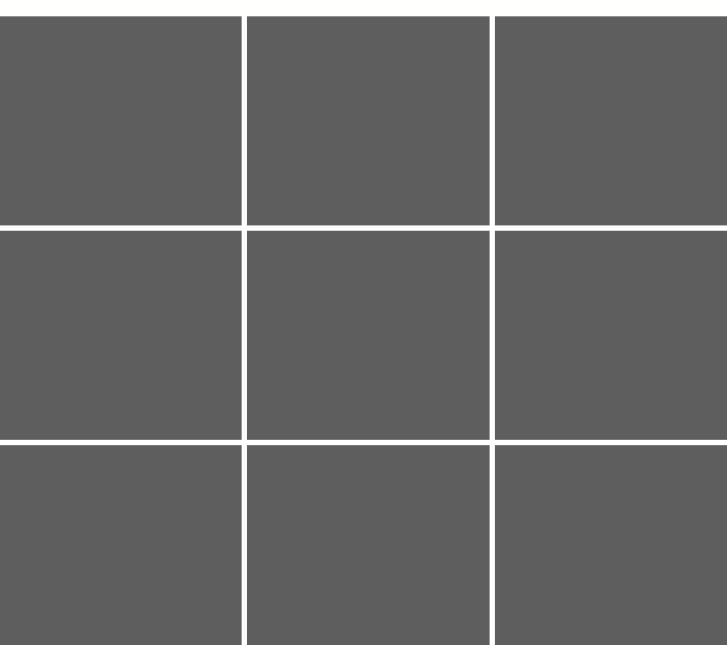
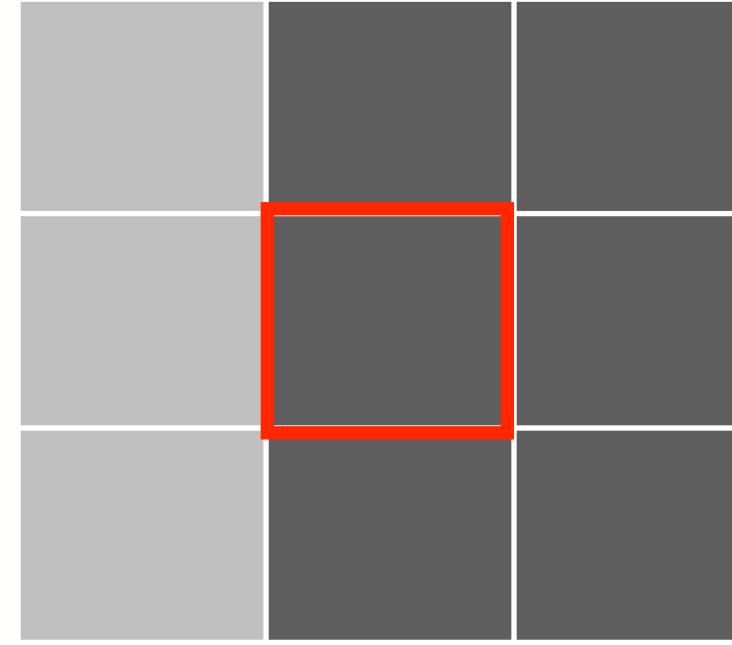
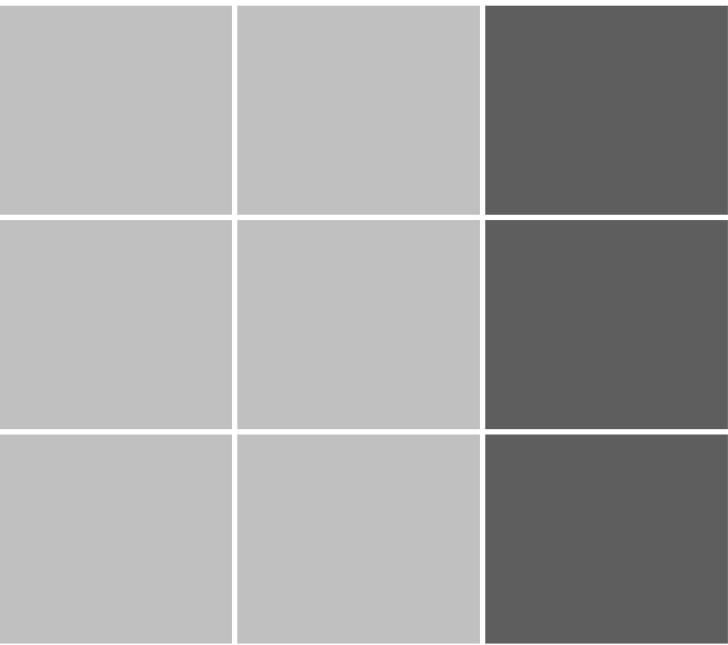
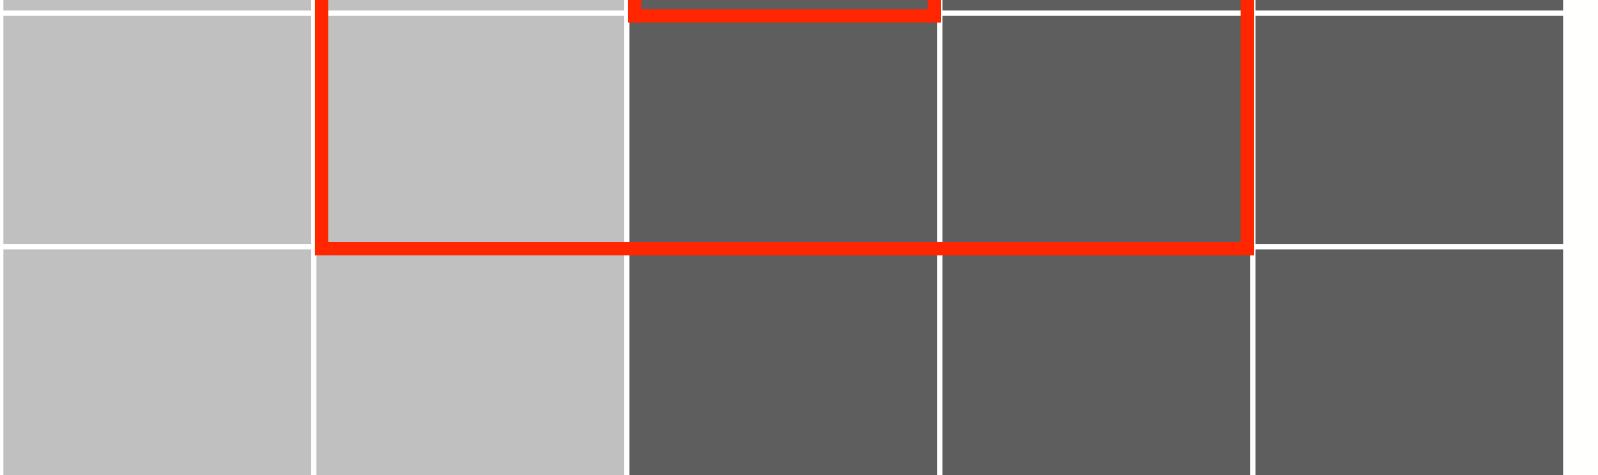
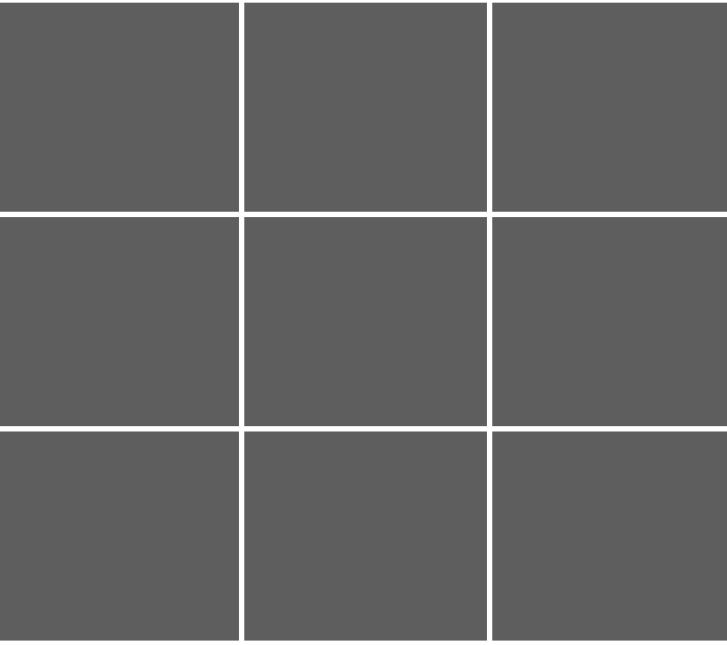
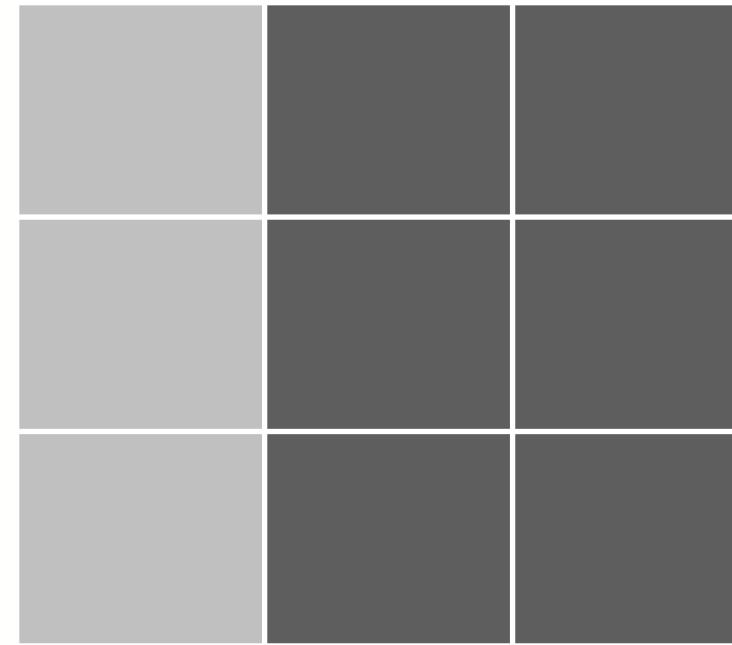
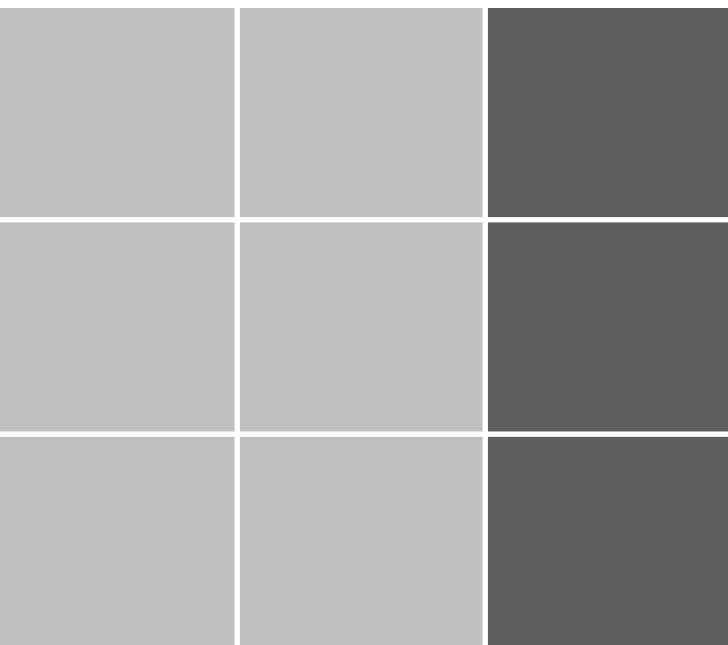
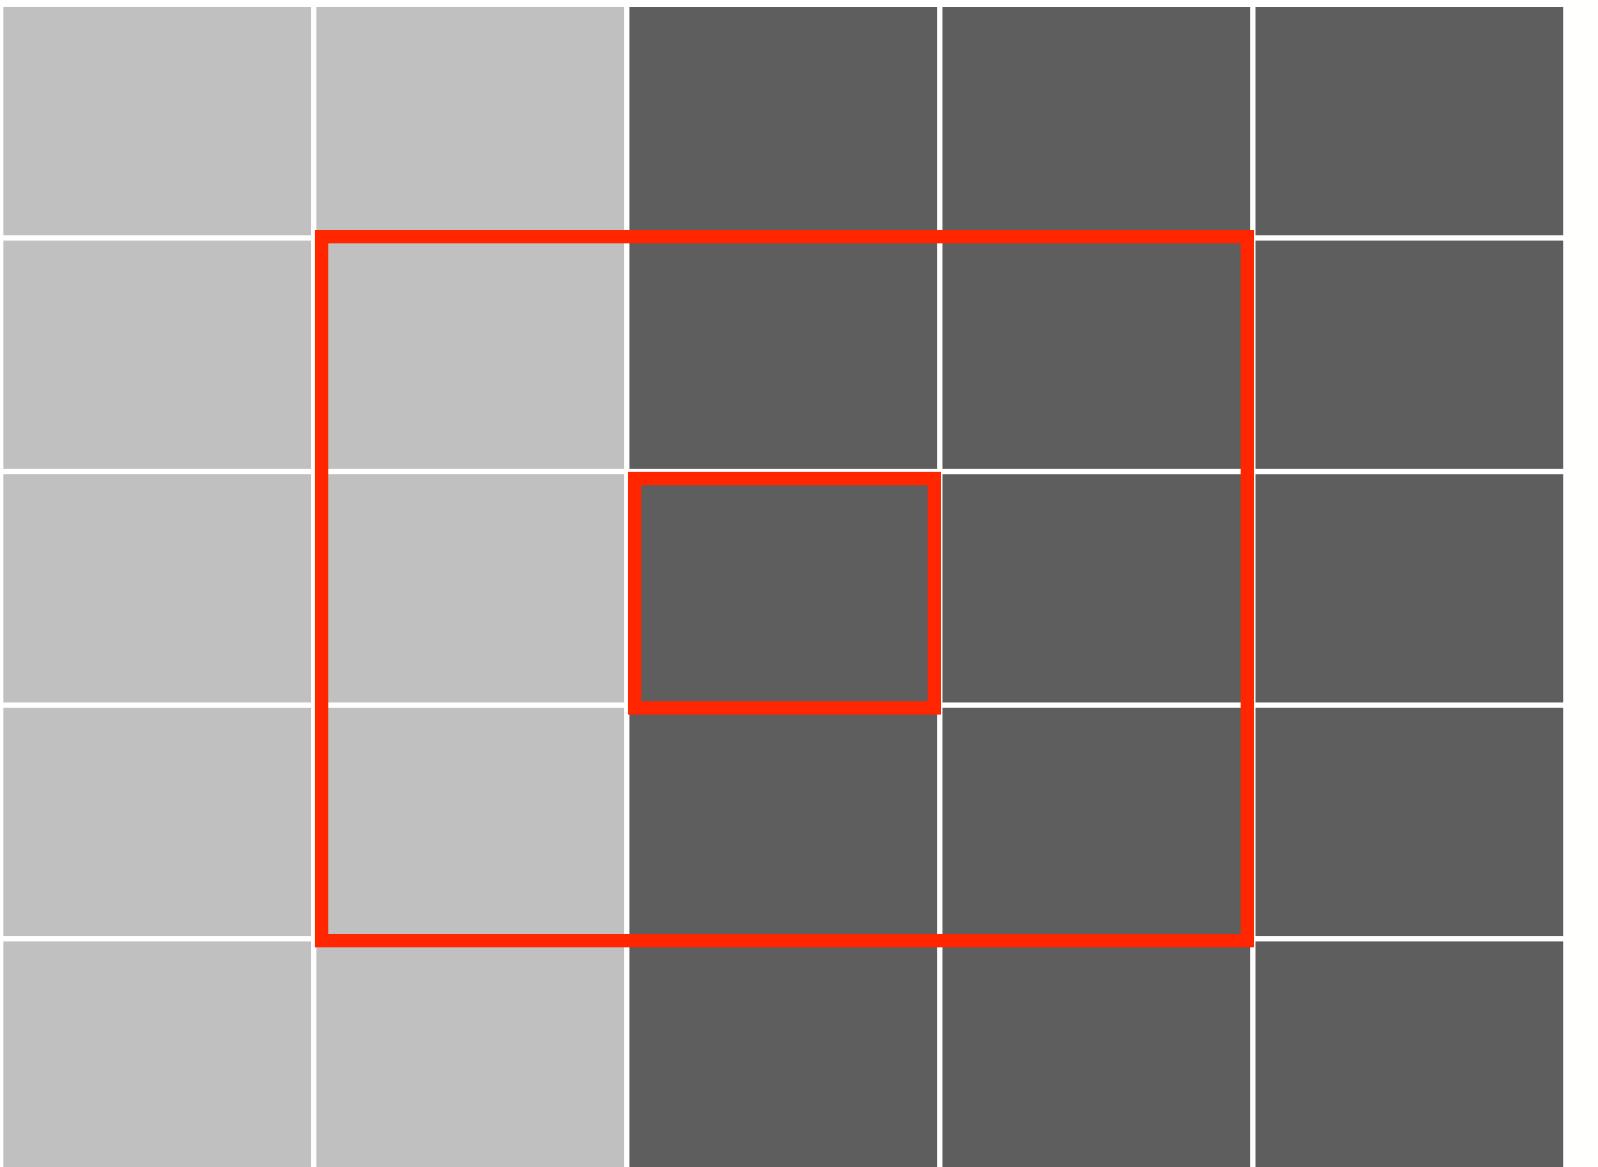


Corner detector intuition



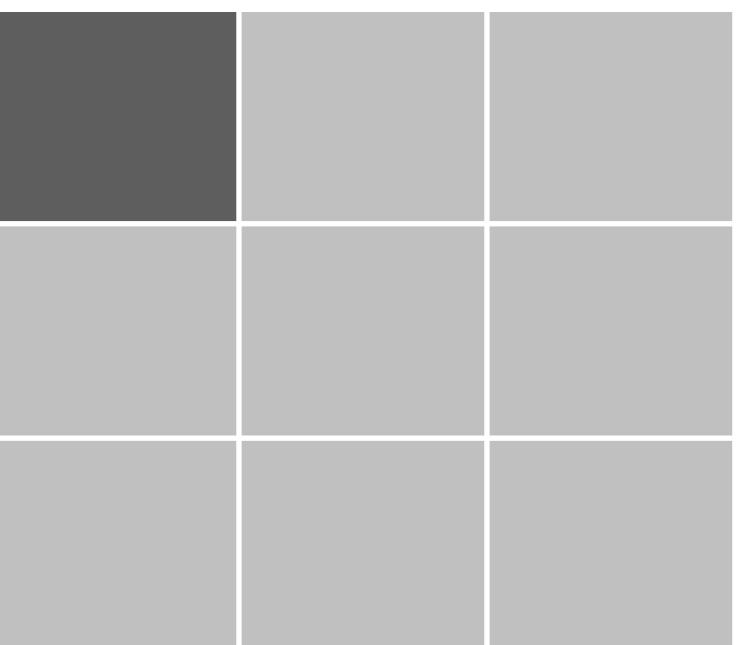
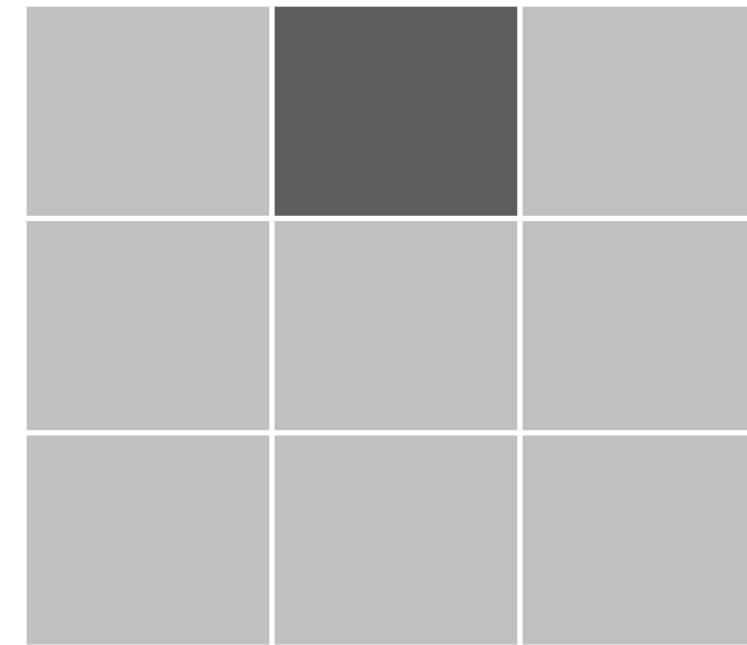
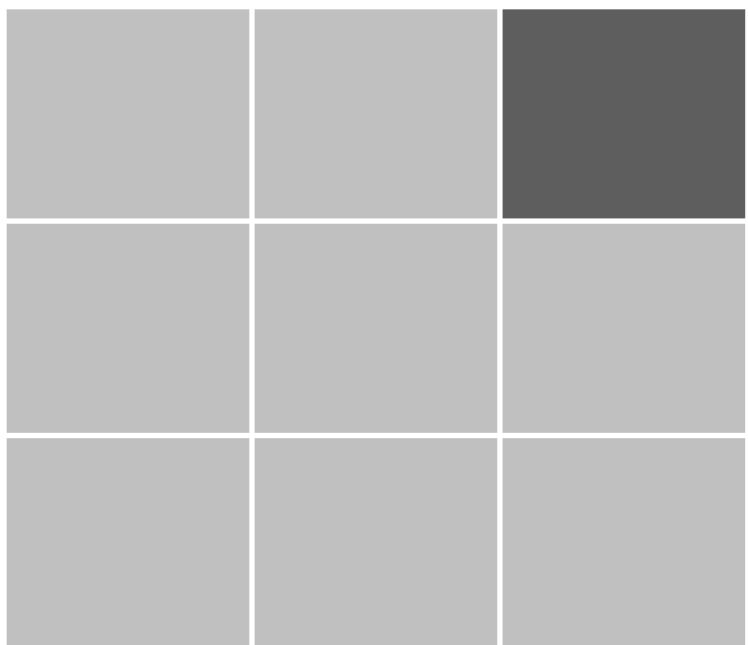
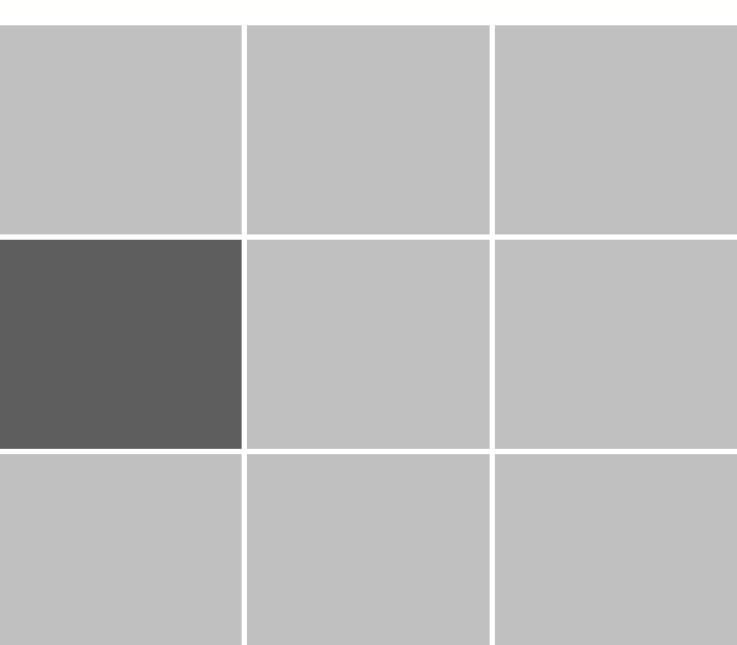
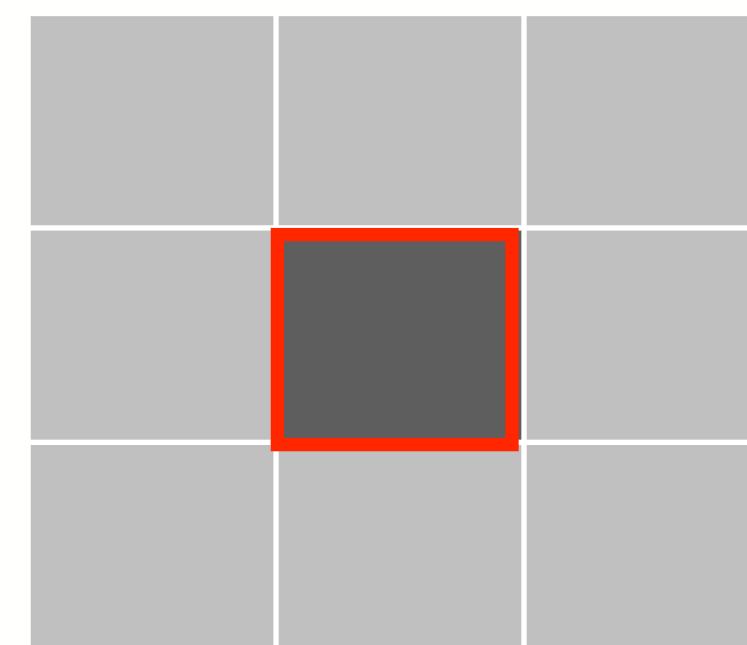
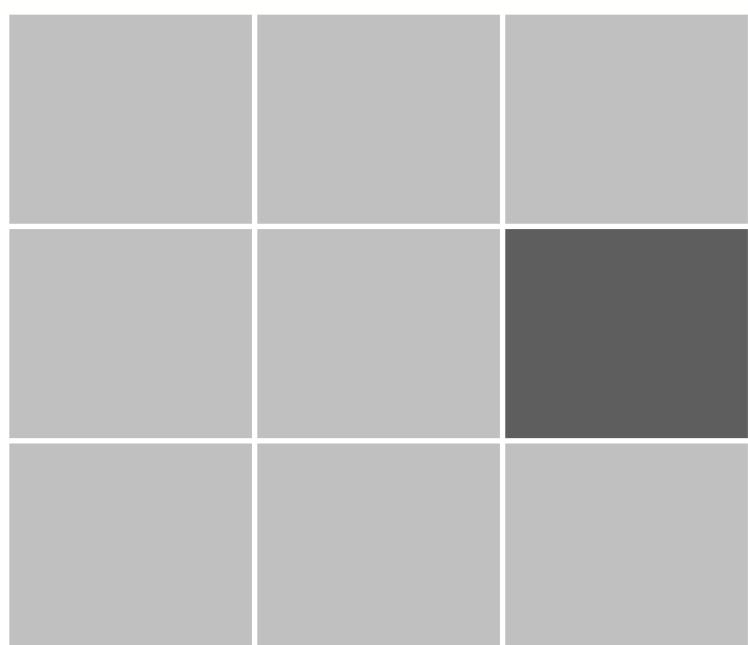
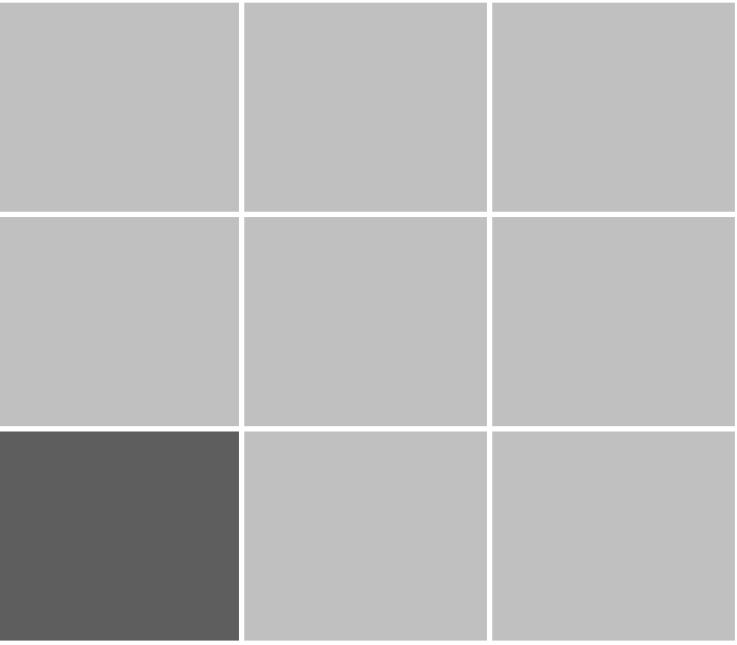
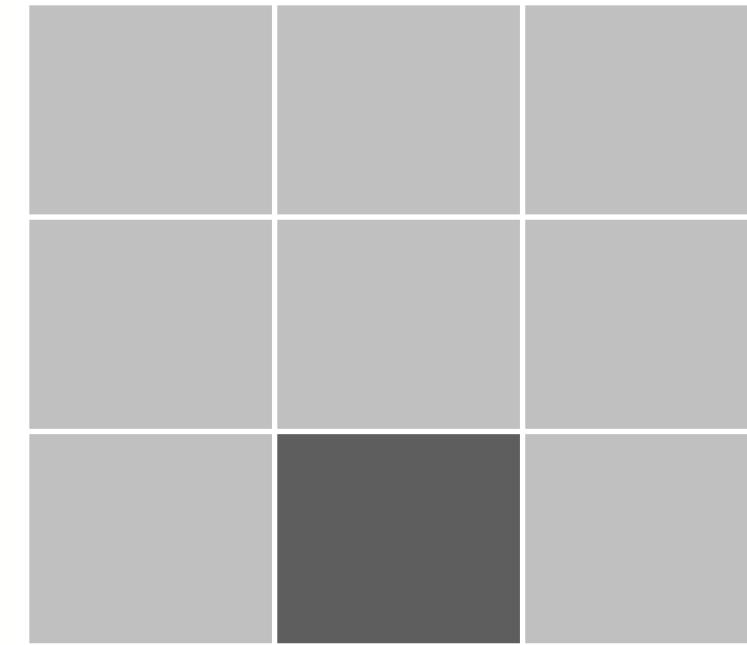
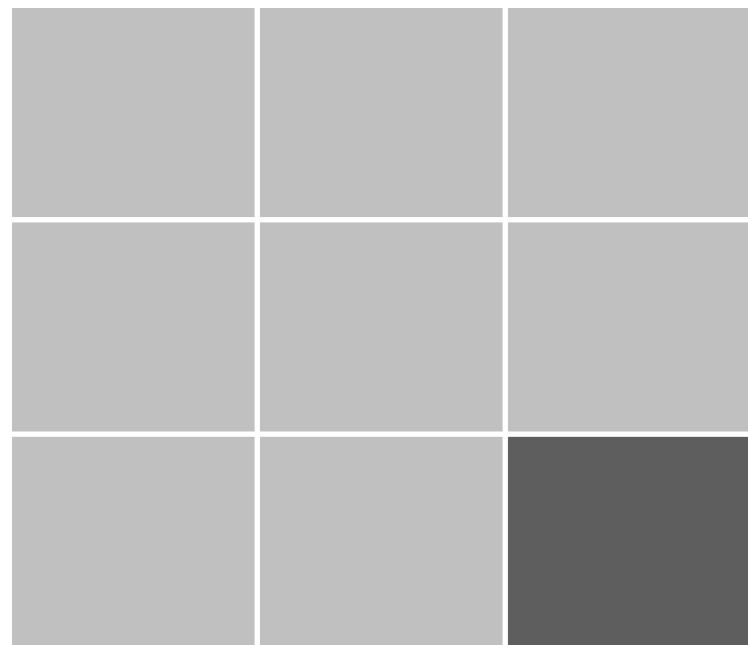
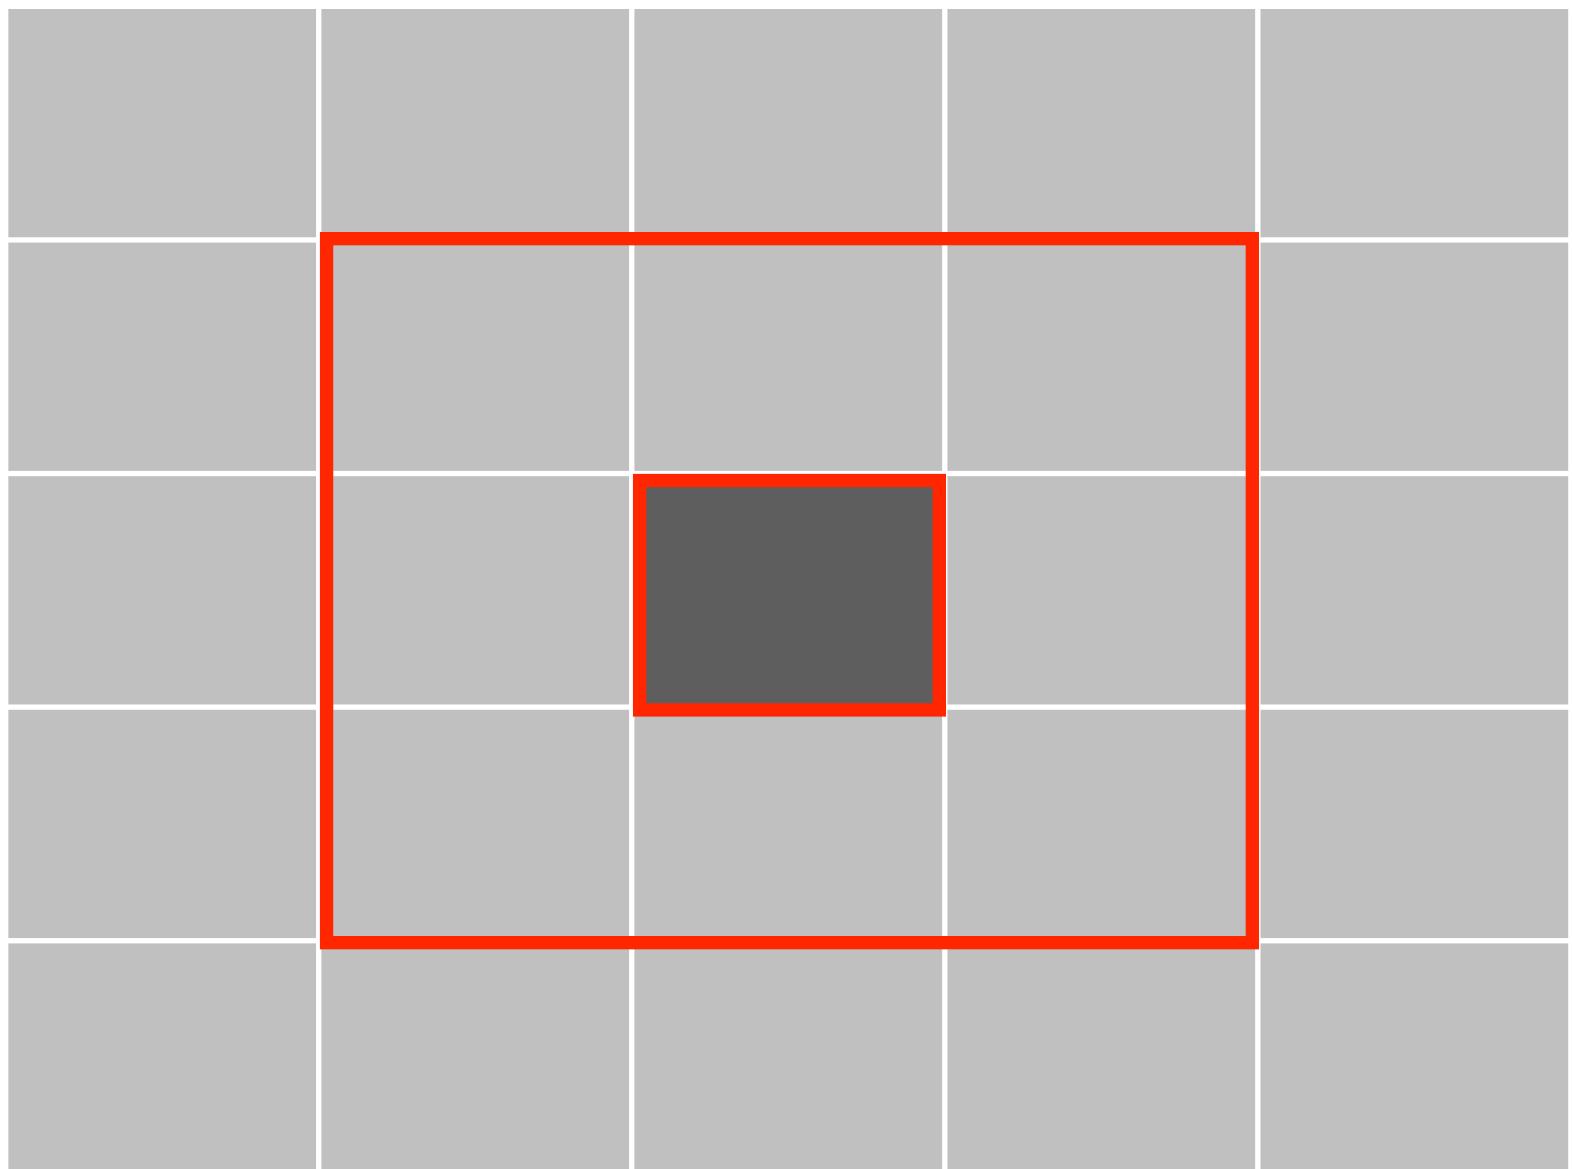
- Similar at 8 other locations
- ➡ not very unique

Corner detector intuition



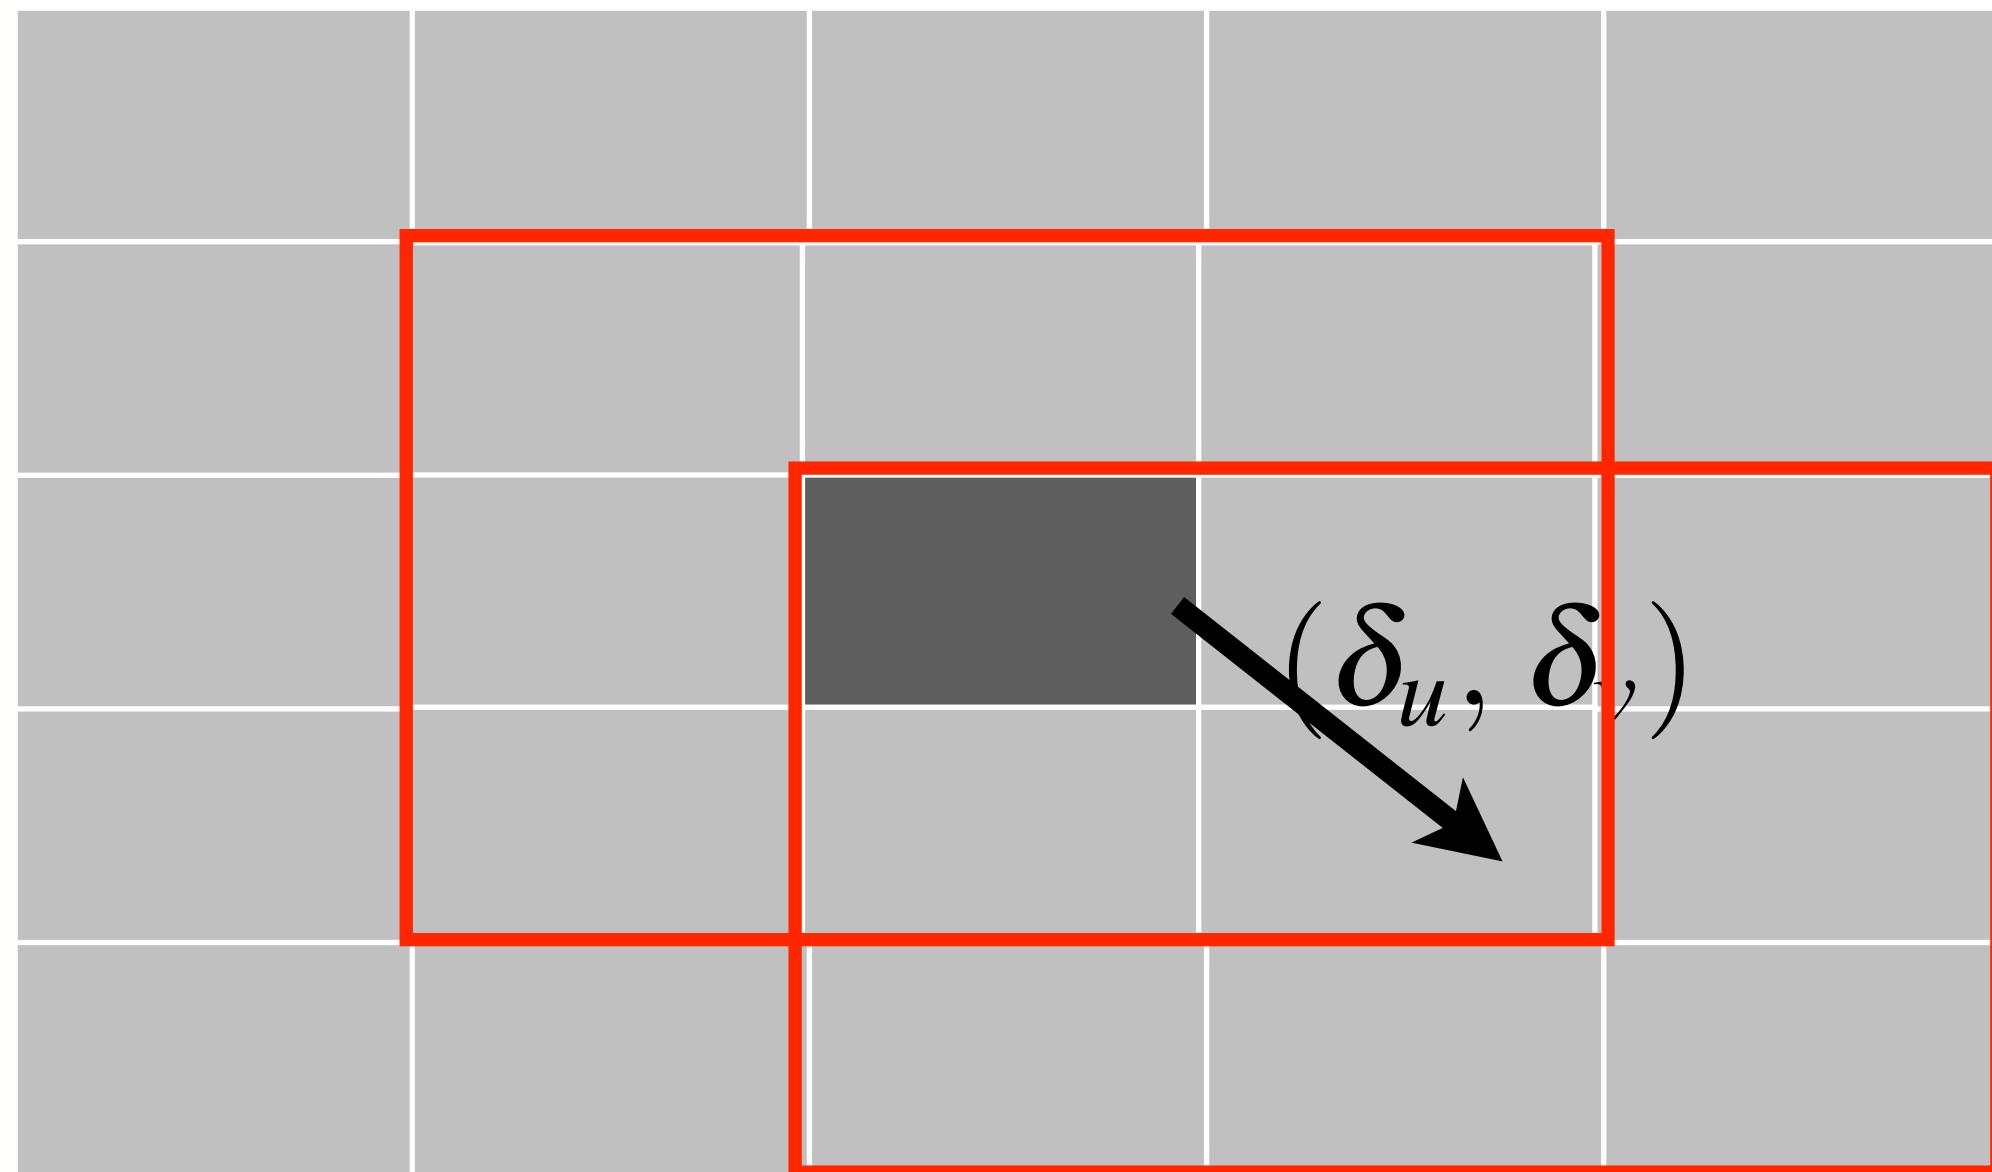
- Similar at 2 other locations
- vertical motion causes no change
- parallel to the line

Corner detector intuition



- Similar at no other locations
→ locally unique

...Corner detector intuition



- Similar at 0 other locations (locally unique)

$$s(u, v, \delta_u, \delta_v) = \sum_{(i,j) \in \mathcal{W}} (I[u + \delta_u + i, v + \delta_v + j] - I[u + i, v + j])^2$$

$$C_M(u, v) = \min_{(\delta_u, \delta_v) \in \mathcal{D}} s(u, v, \delta_u, \delta_v)$$

Corner detector

- More general approach
 - ➡ add a Gaussian weighting matrix

$$s(u, v, \delta_u, \delta_v) = \sum_{(i,j) \in \mathcal{W}} W[i, j] \left(\underbrace{I[u + \delta_u + i, v + \delta_v + j]} - I[u + i, v + j] \right)^2$$

$$s(u, v, \delta_u, \delta_v) = (\delta_u \quad \delta_v) A \begin{pmatrix} \delta_u \\ \delta_v \end{pmatrix}$$

structure tensor

$$A = \begin{pmatrix} G(\sigma_I) \otimes I_u^2 & G(\sigma_I) \otimes I_u I_v \\ G(\sigma_I) \otimes I_u I_v & G(\sigma_I) \otimes I_v^2 \end{pmatrix}$$

based on gradients
so intensity difference
is eliminated

- The eigenvalues of A contain important information

λ_2 small λ_2 large

λ_1 small	constant	edge
λ_1 large	edge	peak

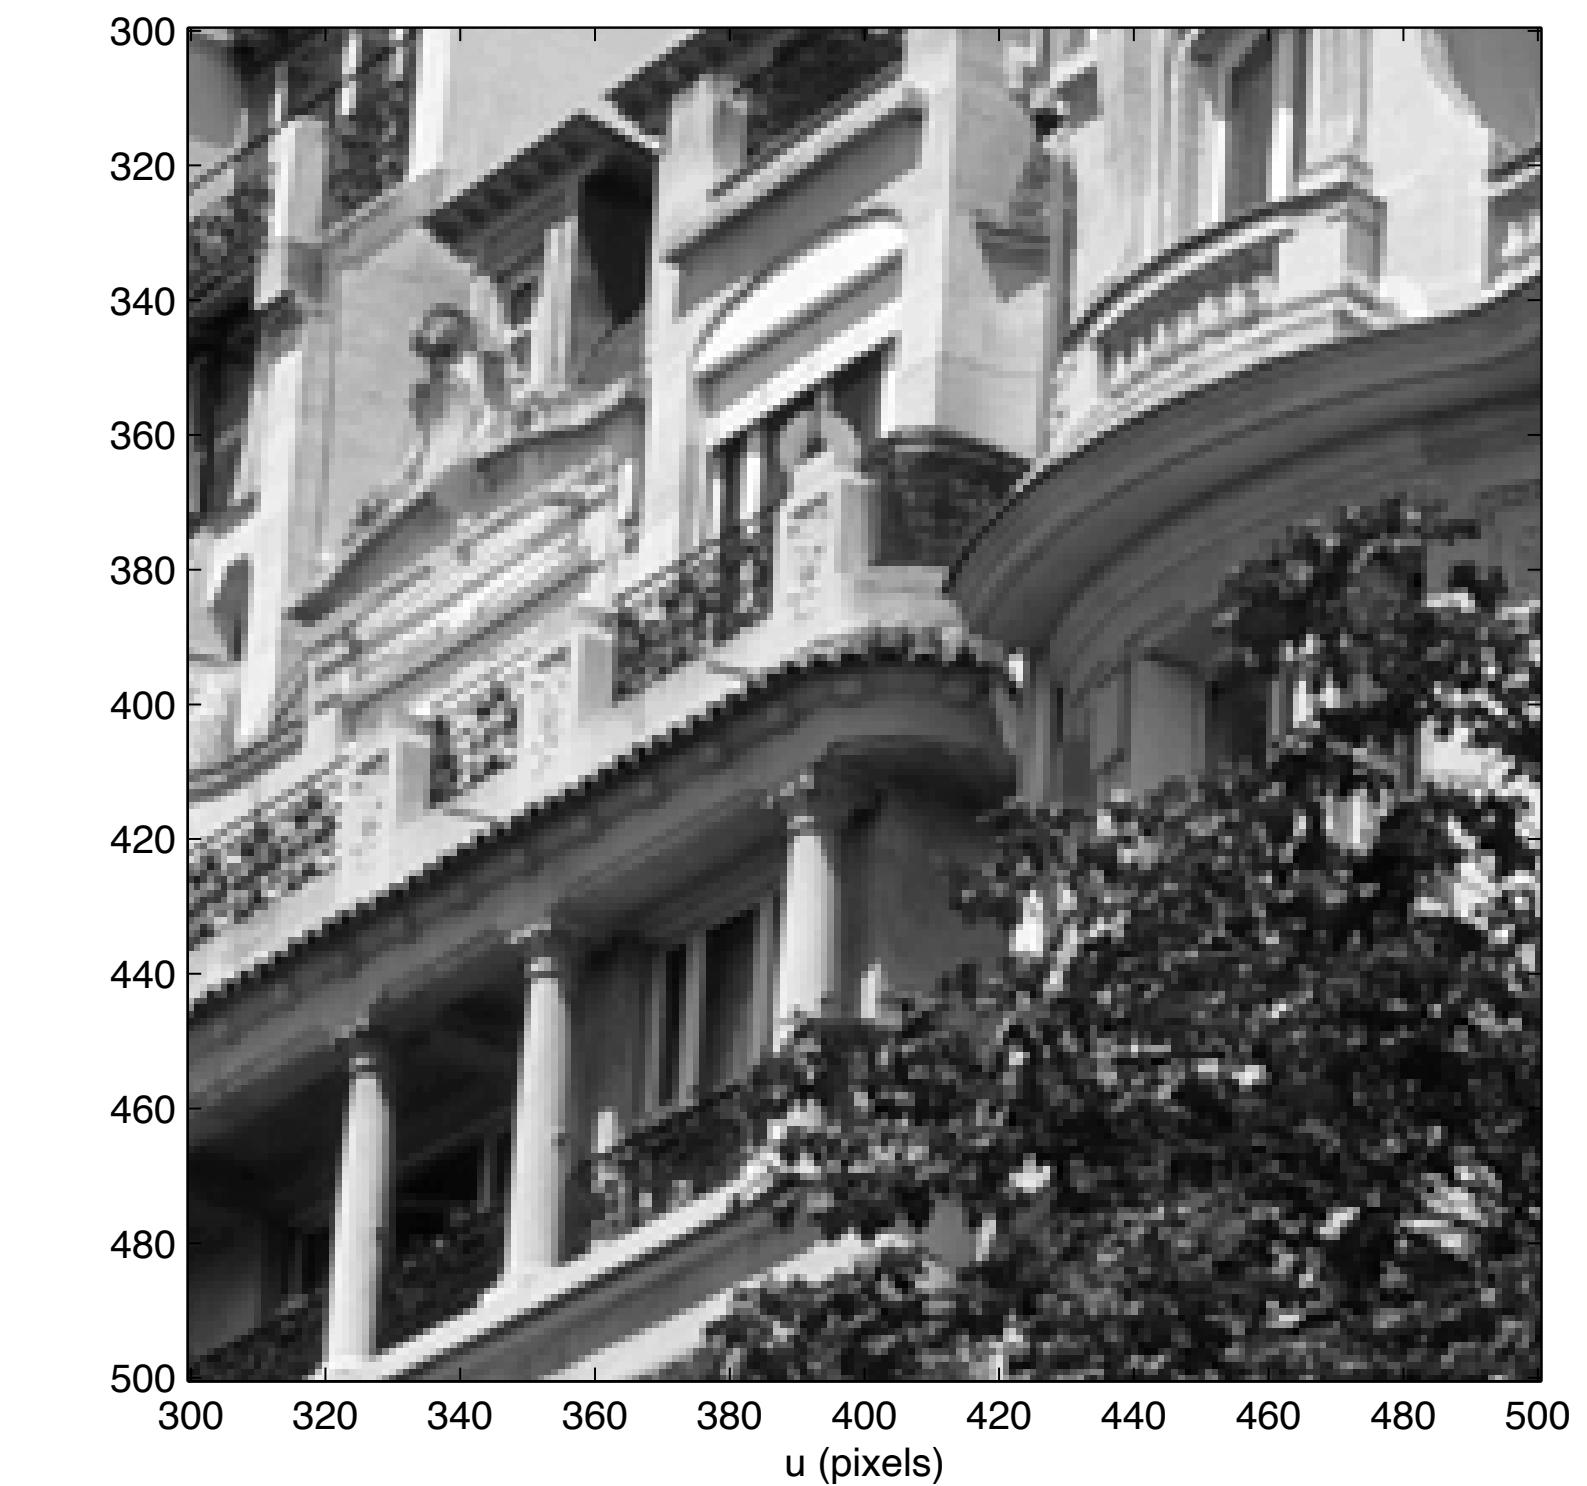
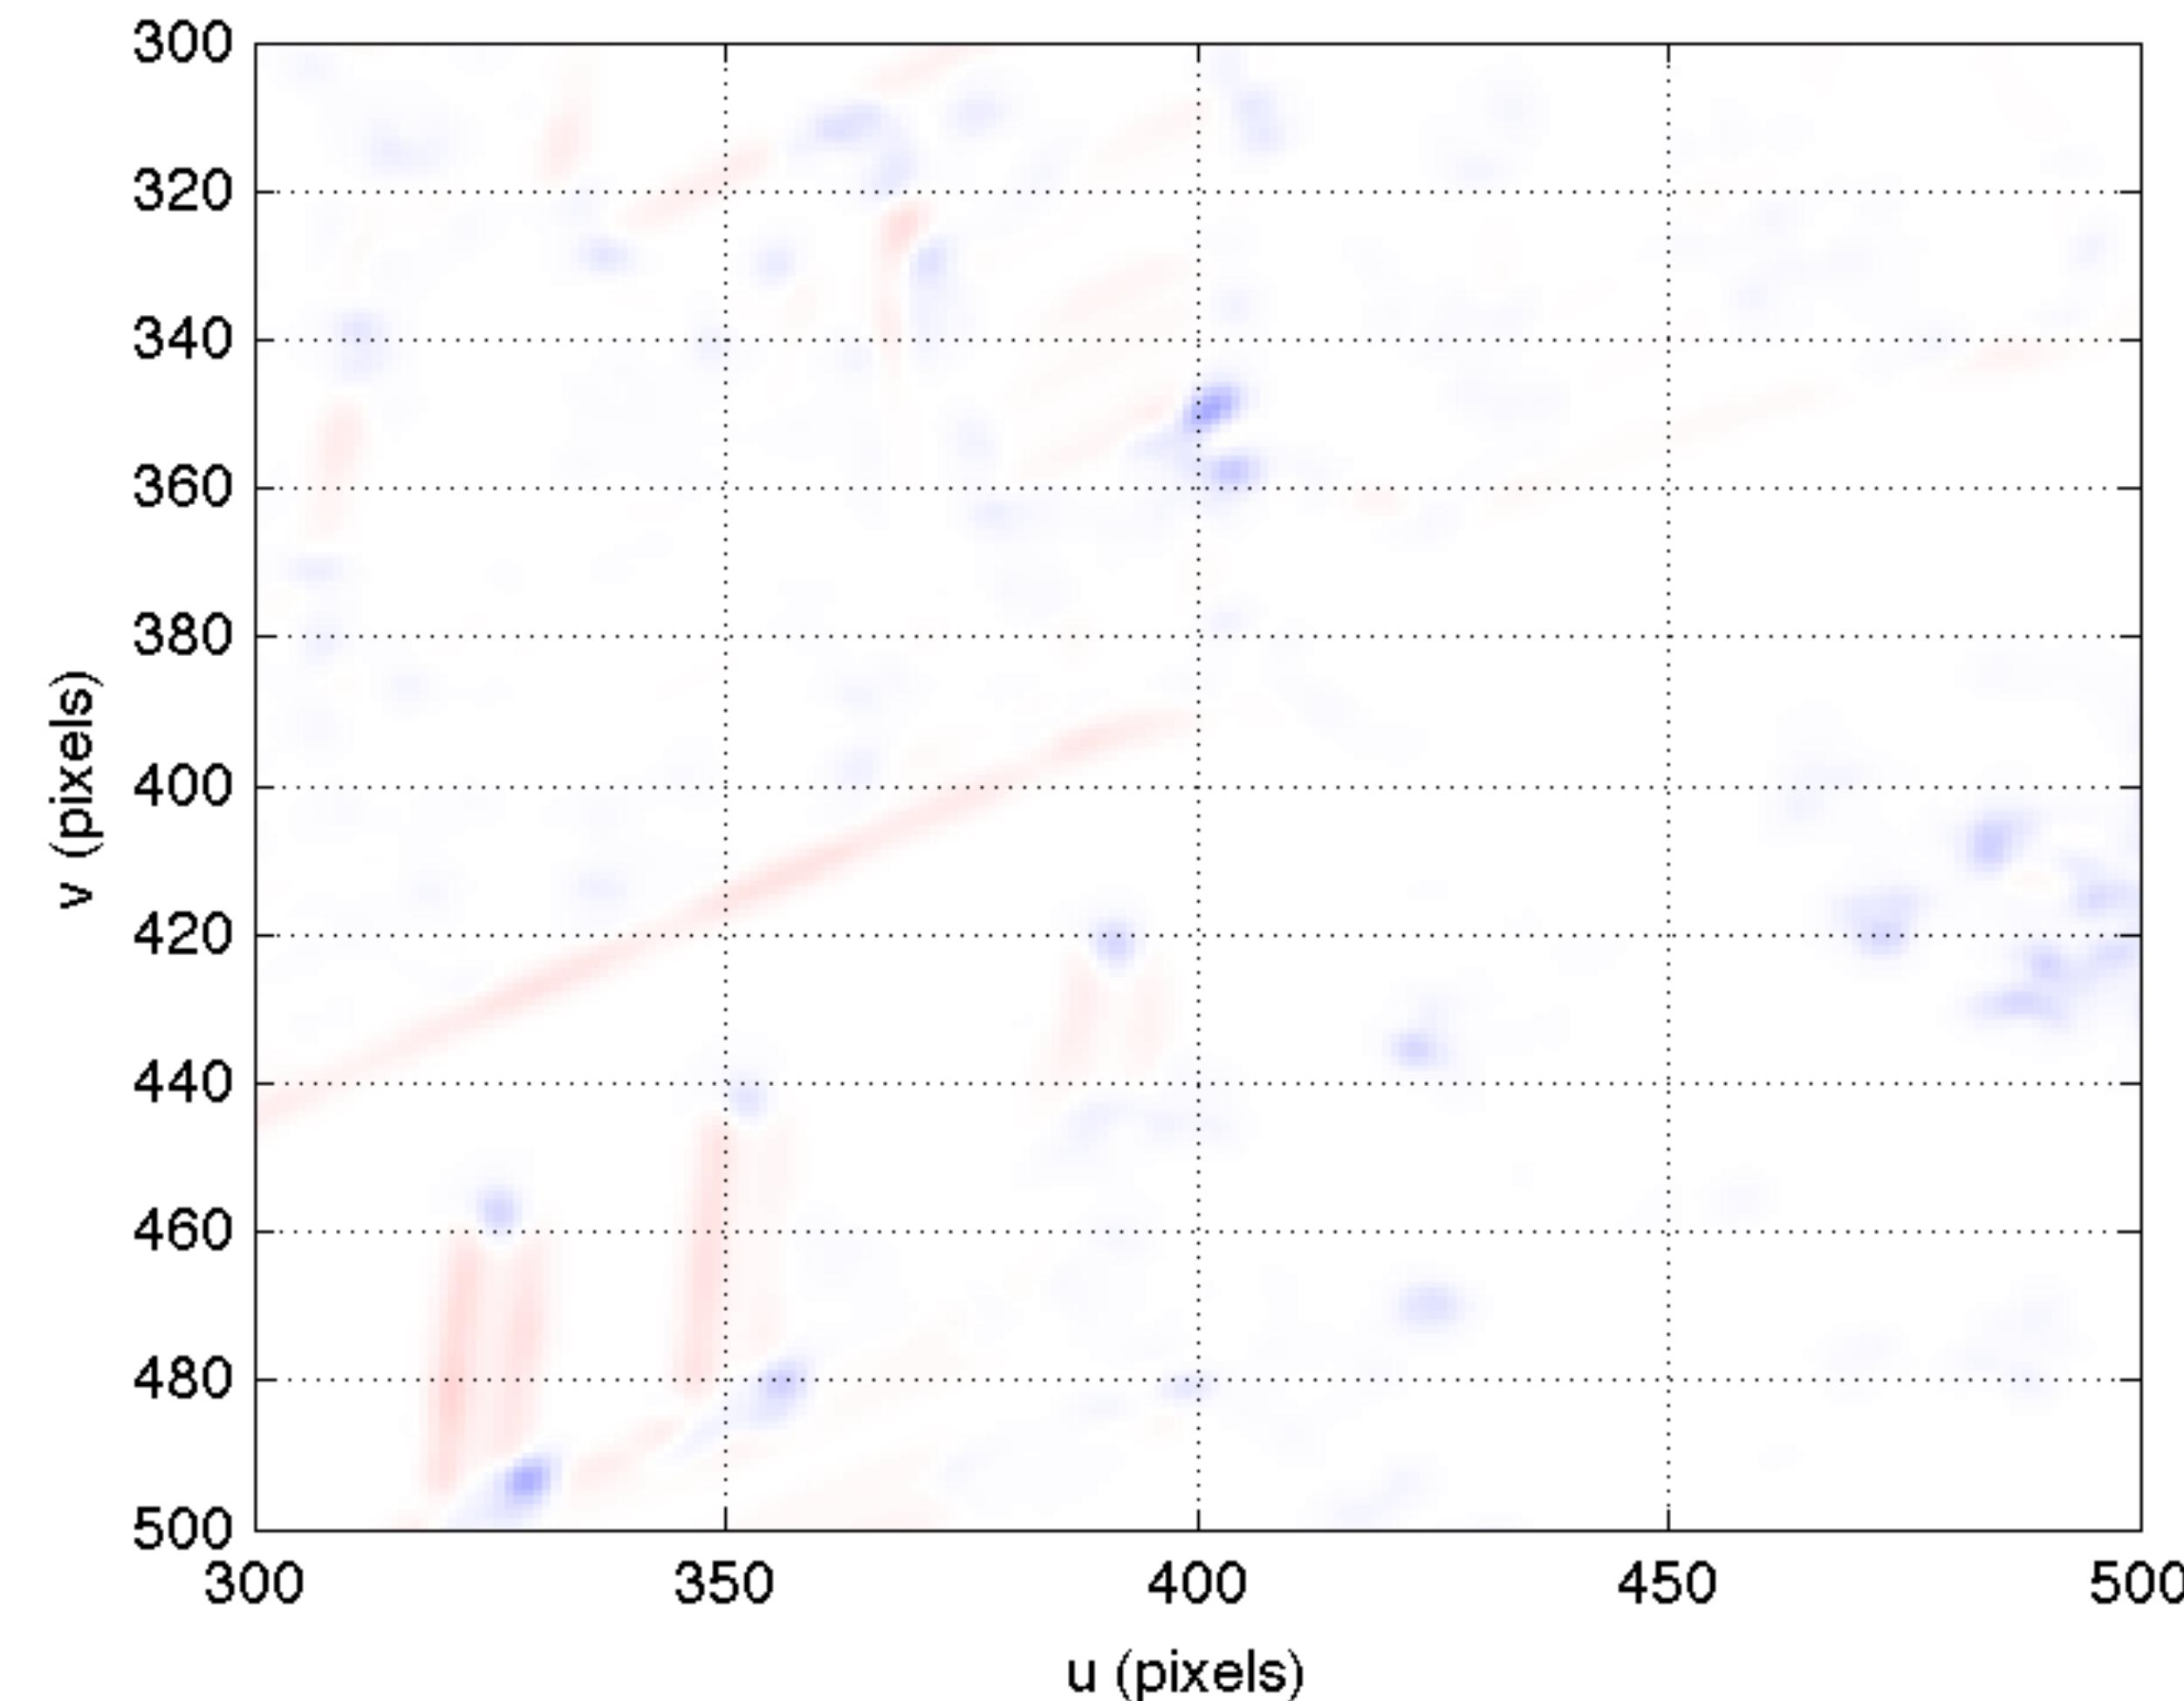
Shi-Tomasi detector

$$C_{ST}(u, v) = \min(\lambda_1, \lambda_2)$$

Harris detector

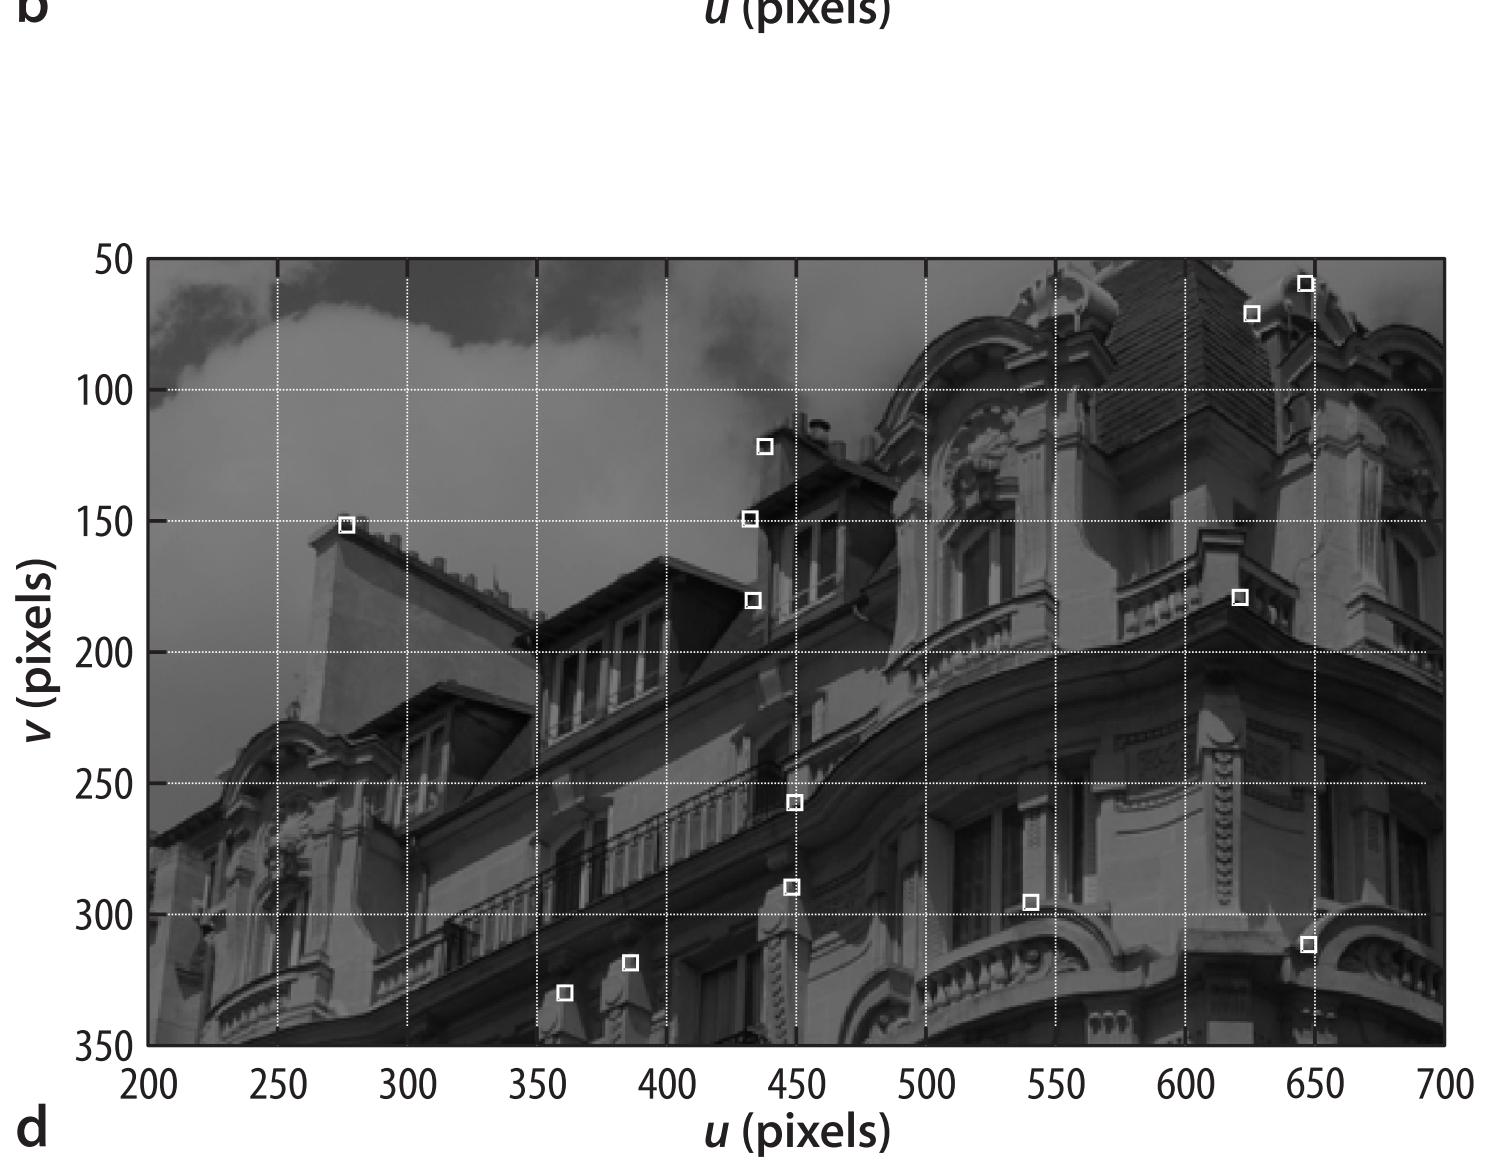
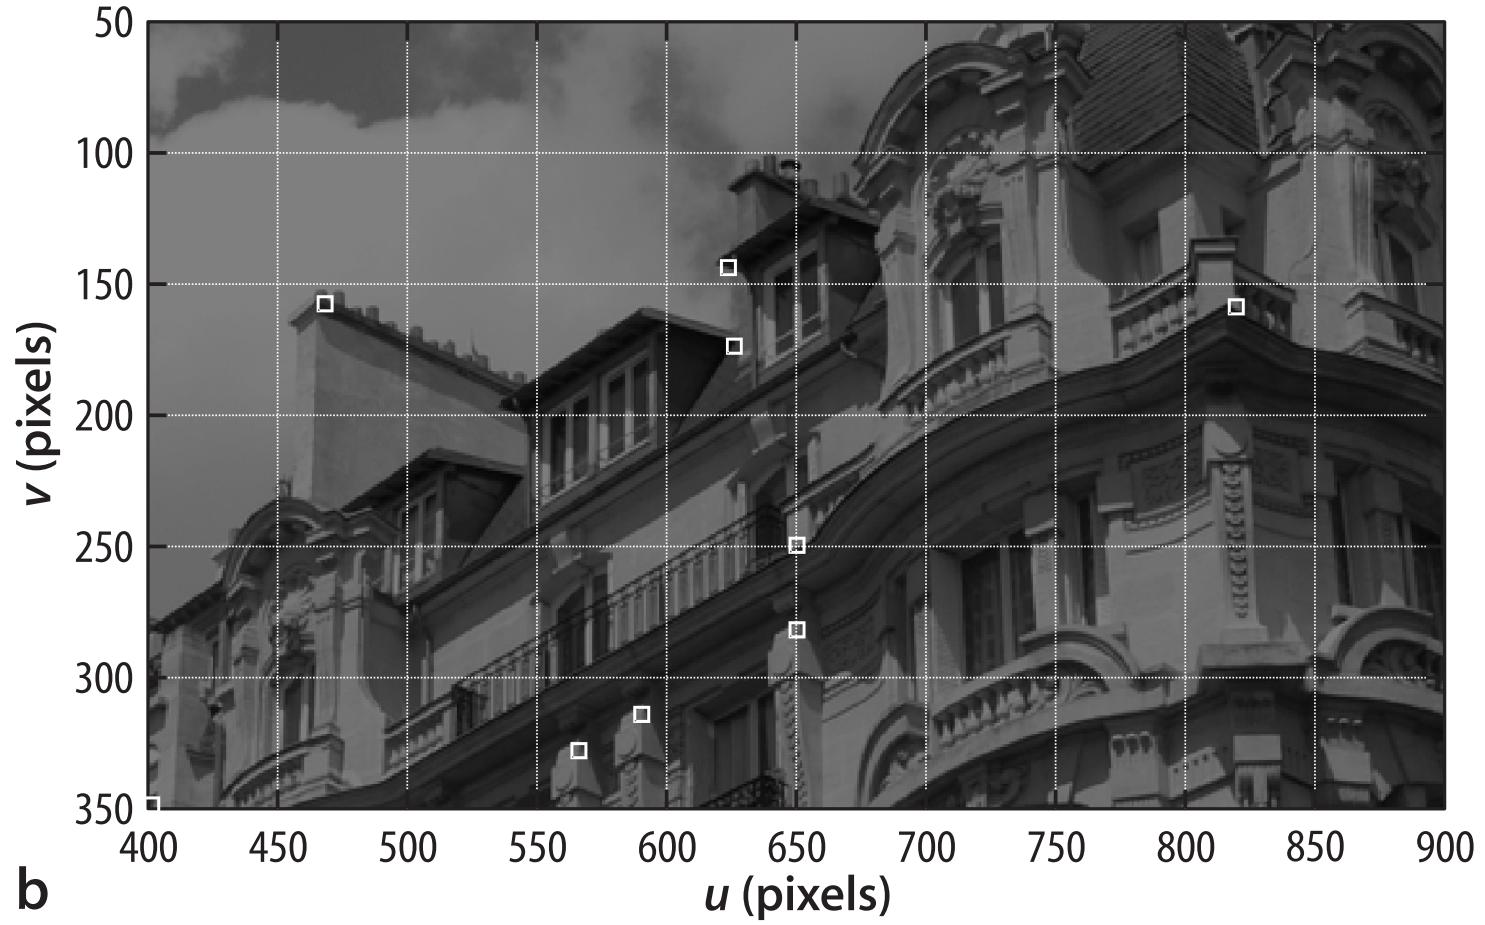
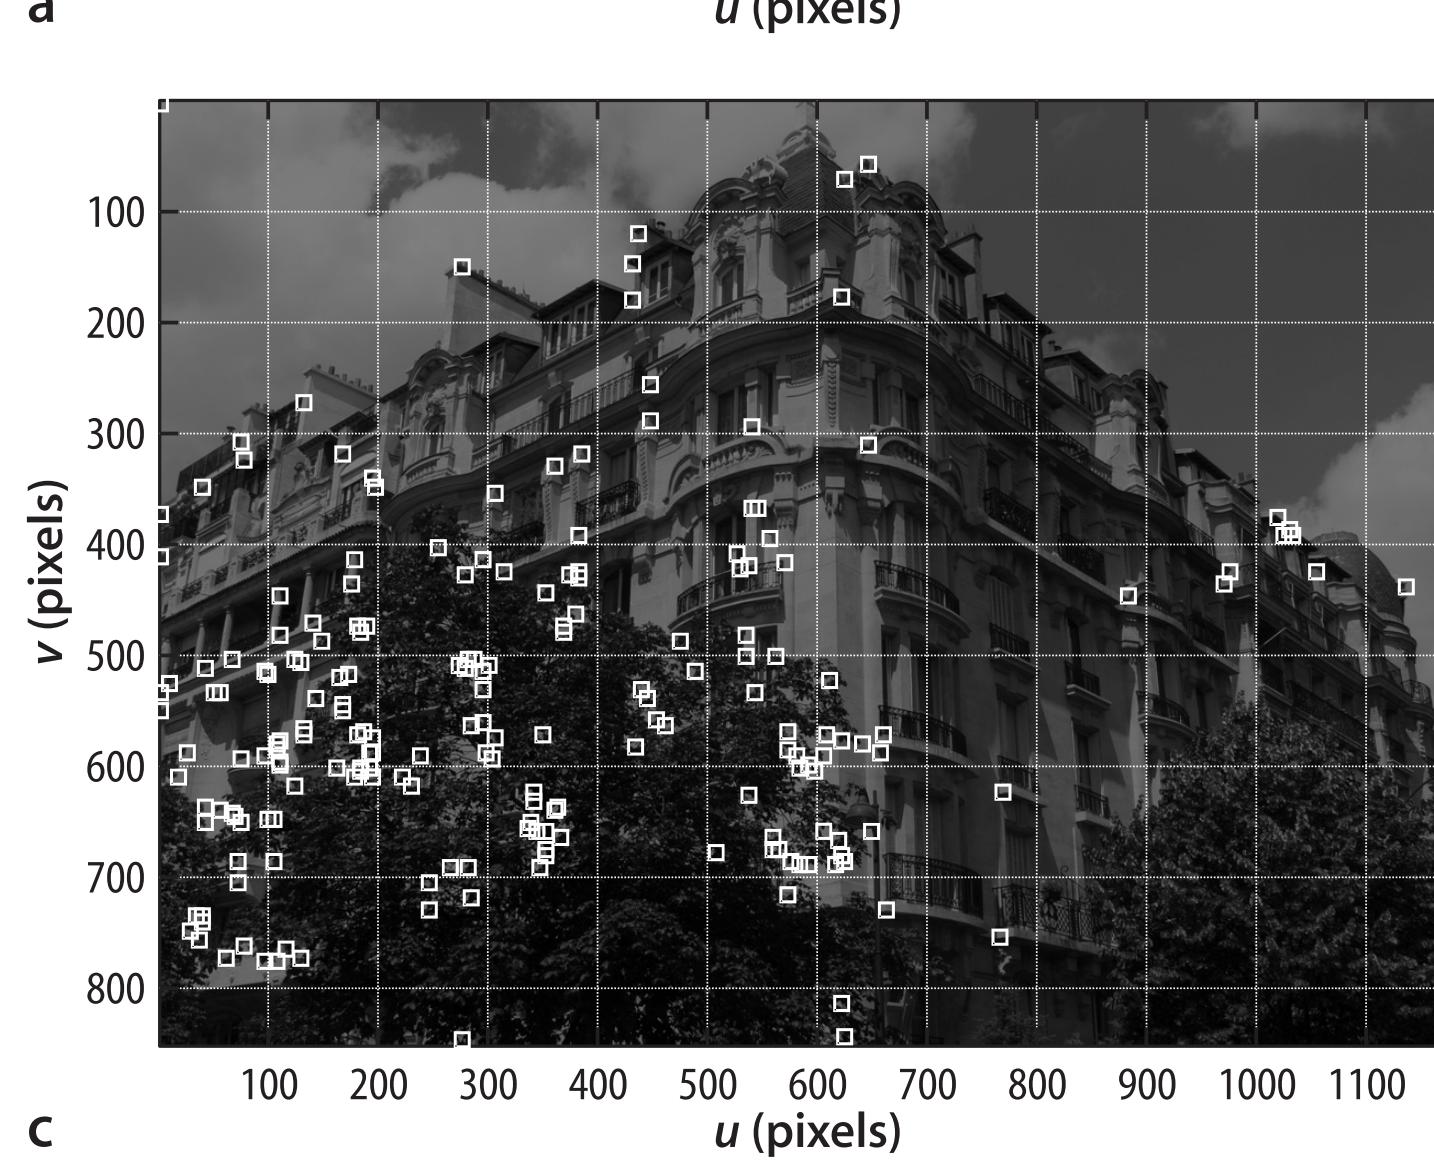
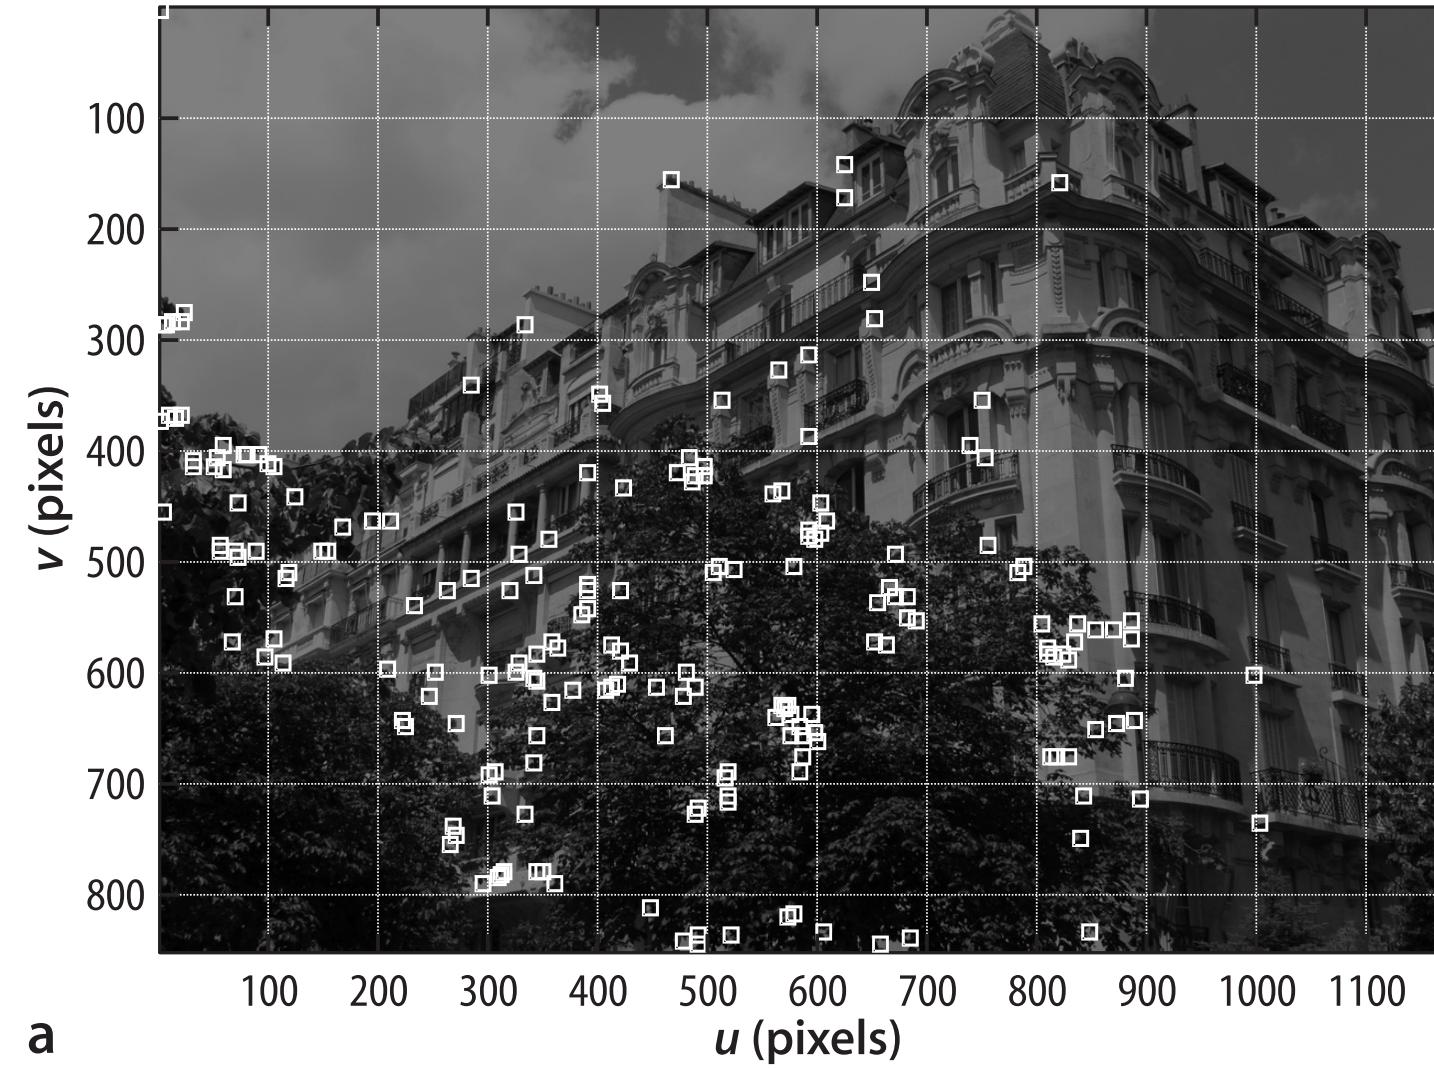
$$C_H(u, v) = \det(A) - k \text{tr}(A)$$

Harris corner value



blue is positive,
red is negative

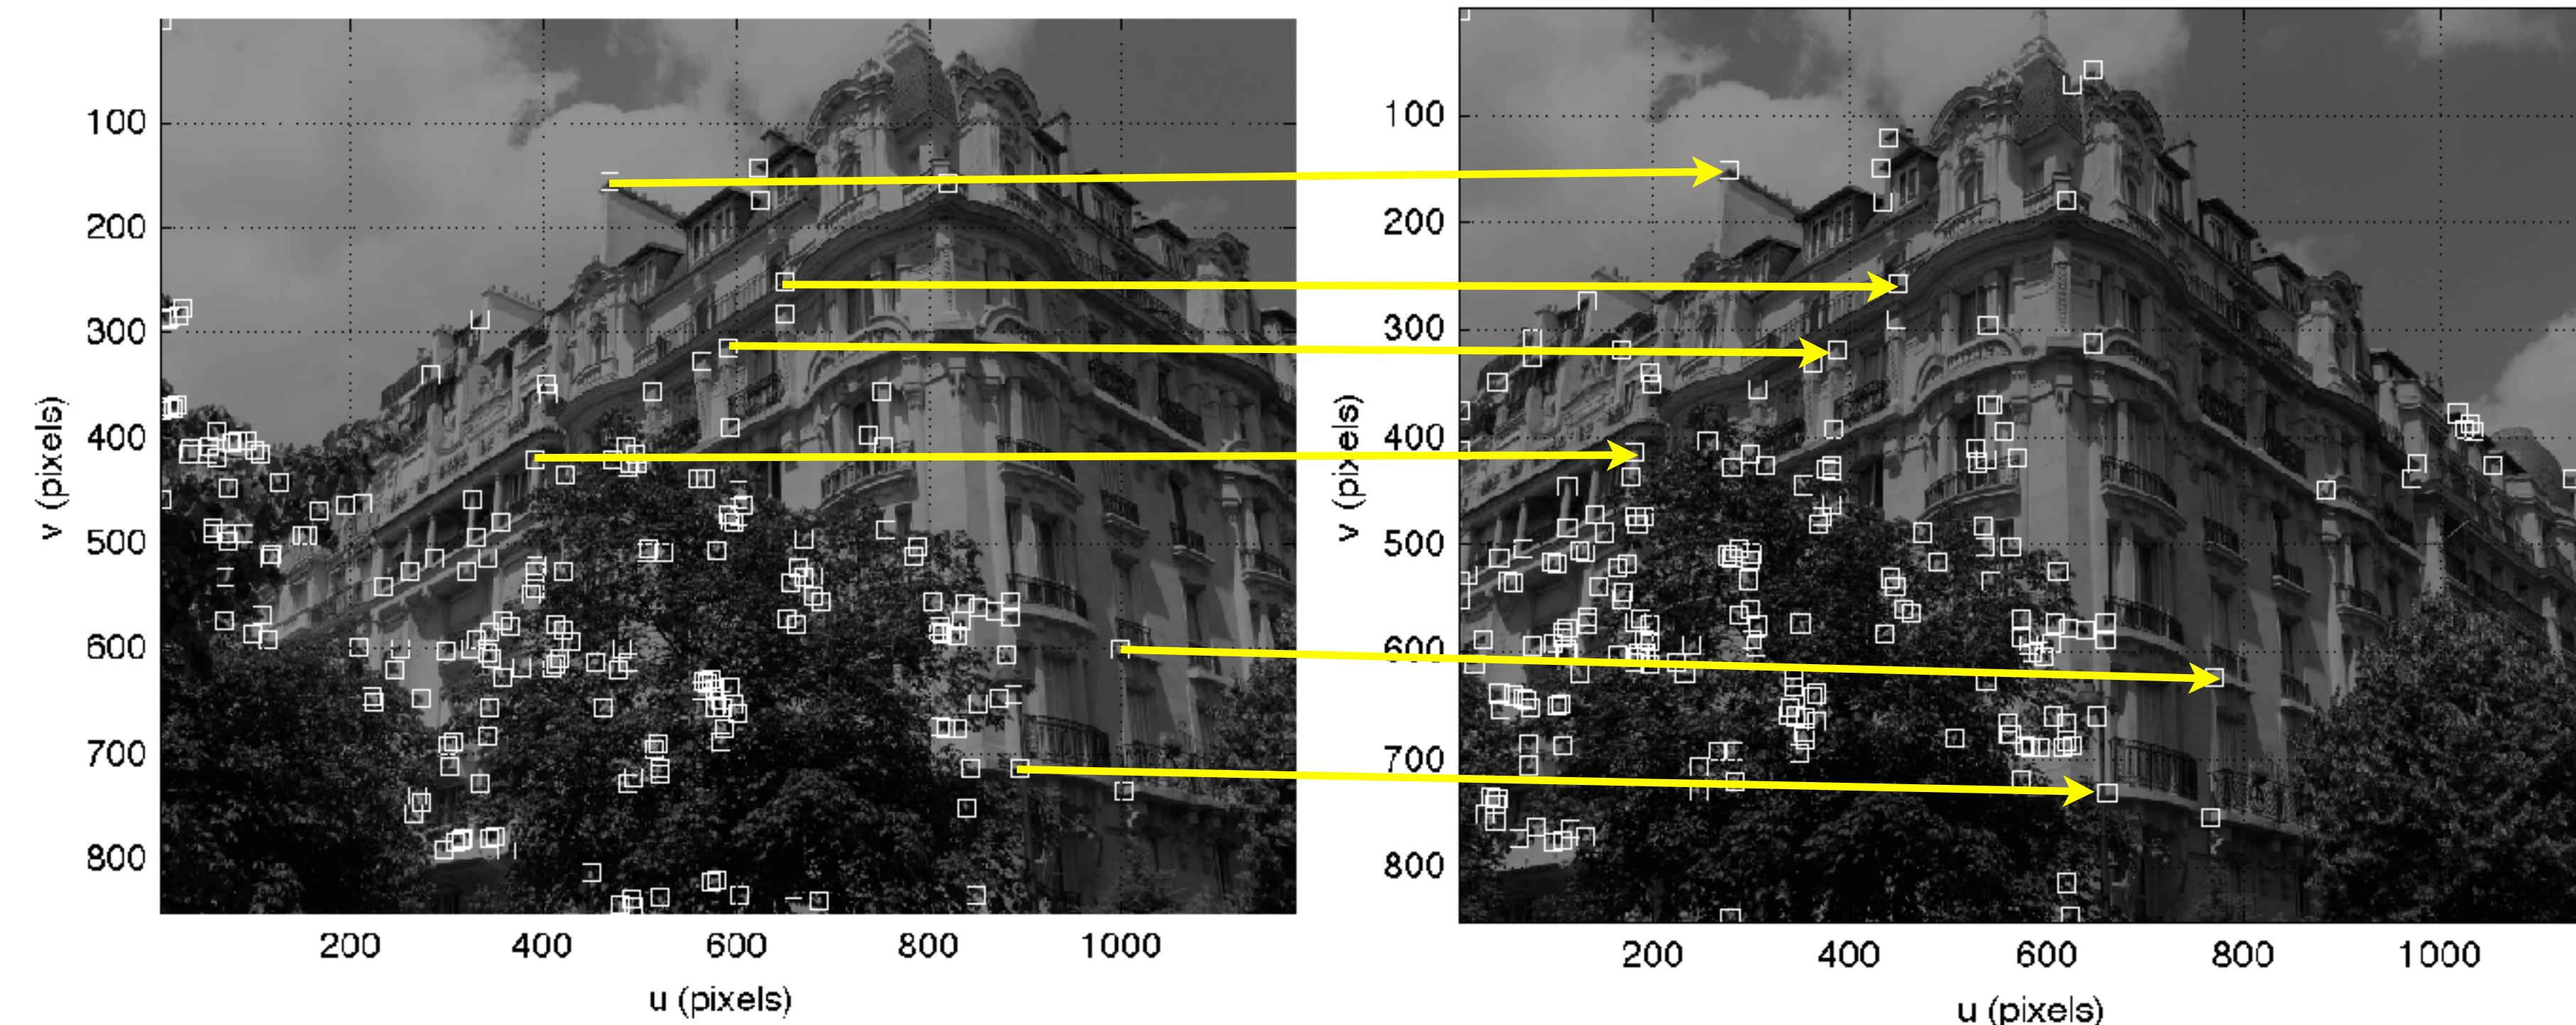
Harris features for 2 views



- Some features are clearly the same point in the scene (good)
- Unduly excited about the trees at the lower left (not good)

Comparing features

- We've found the coordinates of some interesting points in each image:
 $\{^1\mathbf{p}_i, i \in 1 \cdots N_1\}$ $\{^2\mathbf{p}_j, j \in 1 \cdots N_2\}$
- Now we need to determine the correspondence, which $^1\mathbf{p}_i \leftrightarrow ^2\mathbf{p}_j$



Feature matching

- We use a $W \times W$ window of pixels centred on each corner point (W is odd)
- We use a similarity metric to compare the windows
- For each point ${}^1\mathbf{p}_i$ we test the similarity against all the points $\{{}^2\mathbf{p}_j, j \in 1 \cdots N_2\}$ in the other image
- This is an $N_1 \times N_2$ search problem

Name	Measure	Toolbox function
SAD	$s = \sum_{(u,v) \in \mathbf{I}} \mathbf{I}_1[u, v] - \mathbf{I}_2[u, v] $	<code>sad()</code>
SSD	$s = \sum_{(u,v) \in \mathbf{I}} (\mathbf{I}_1[u, v] - \mathbf{I}_2[u, v])^2$	<code>ssd()</code>
ZNCC	$s = \frac{\sum_{(u,v) \in \mathbf{I}} \mathbf{I}_1[u, v] \cdot \mathbf{I}_2[u, v]}{\sqrt{\sum_{(u,v) \in \mathbf{I}} \mathbf{I}_1^2[u, v] \cdot \sum_{(u,v) \in \mathbf{I}} \mathbf{I}_2^2[u, v]}}$	<code>zncc()</code>

Tracking across frames



```
>> t = Tracker(im, c)
200 continuing tracks, 41 new tracks, 0 retired
241 continuing tracks, 46 new tracks, 0 retired
287 continuing tracks, 34 new tracks, 0 retired
.
.
```

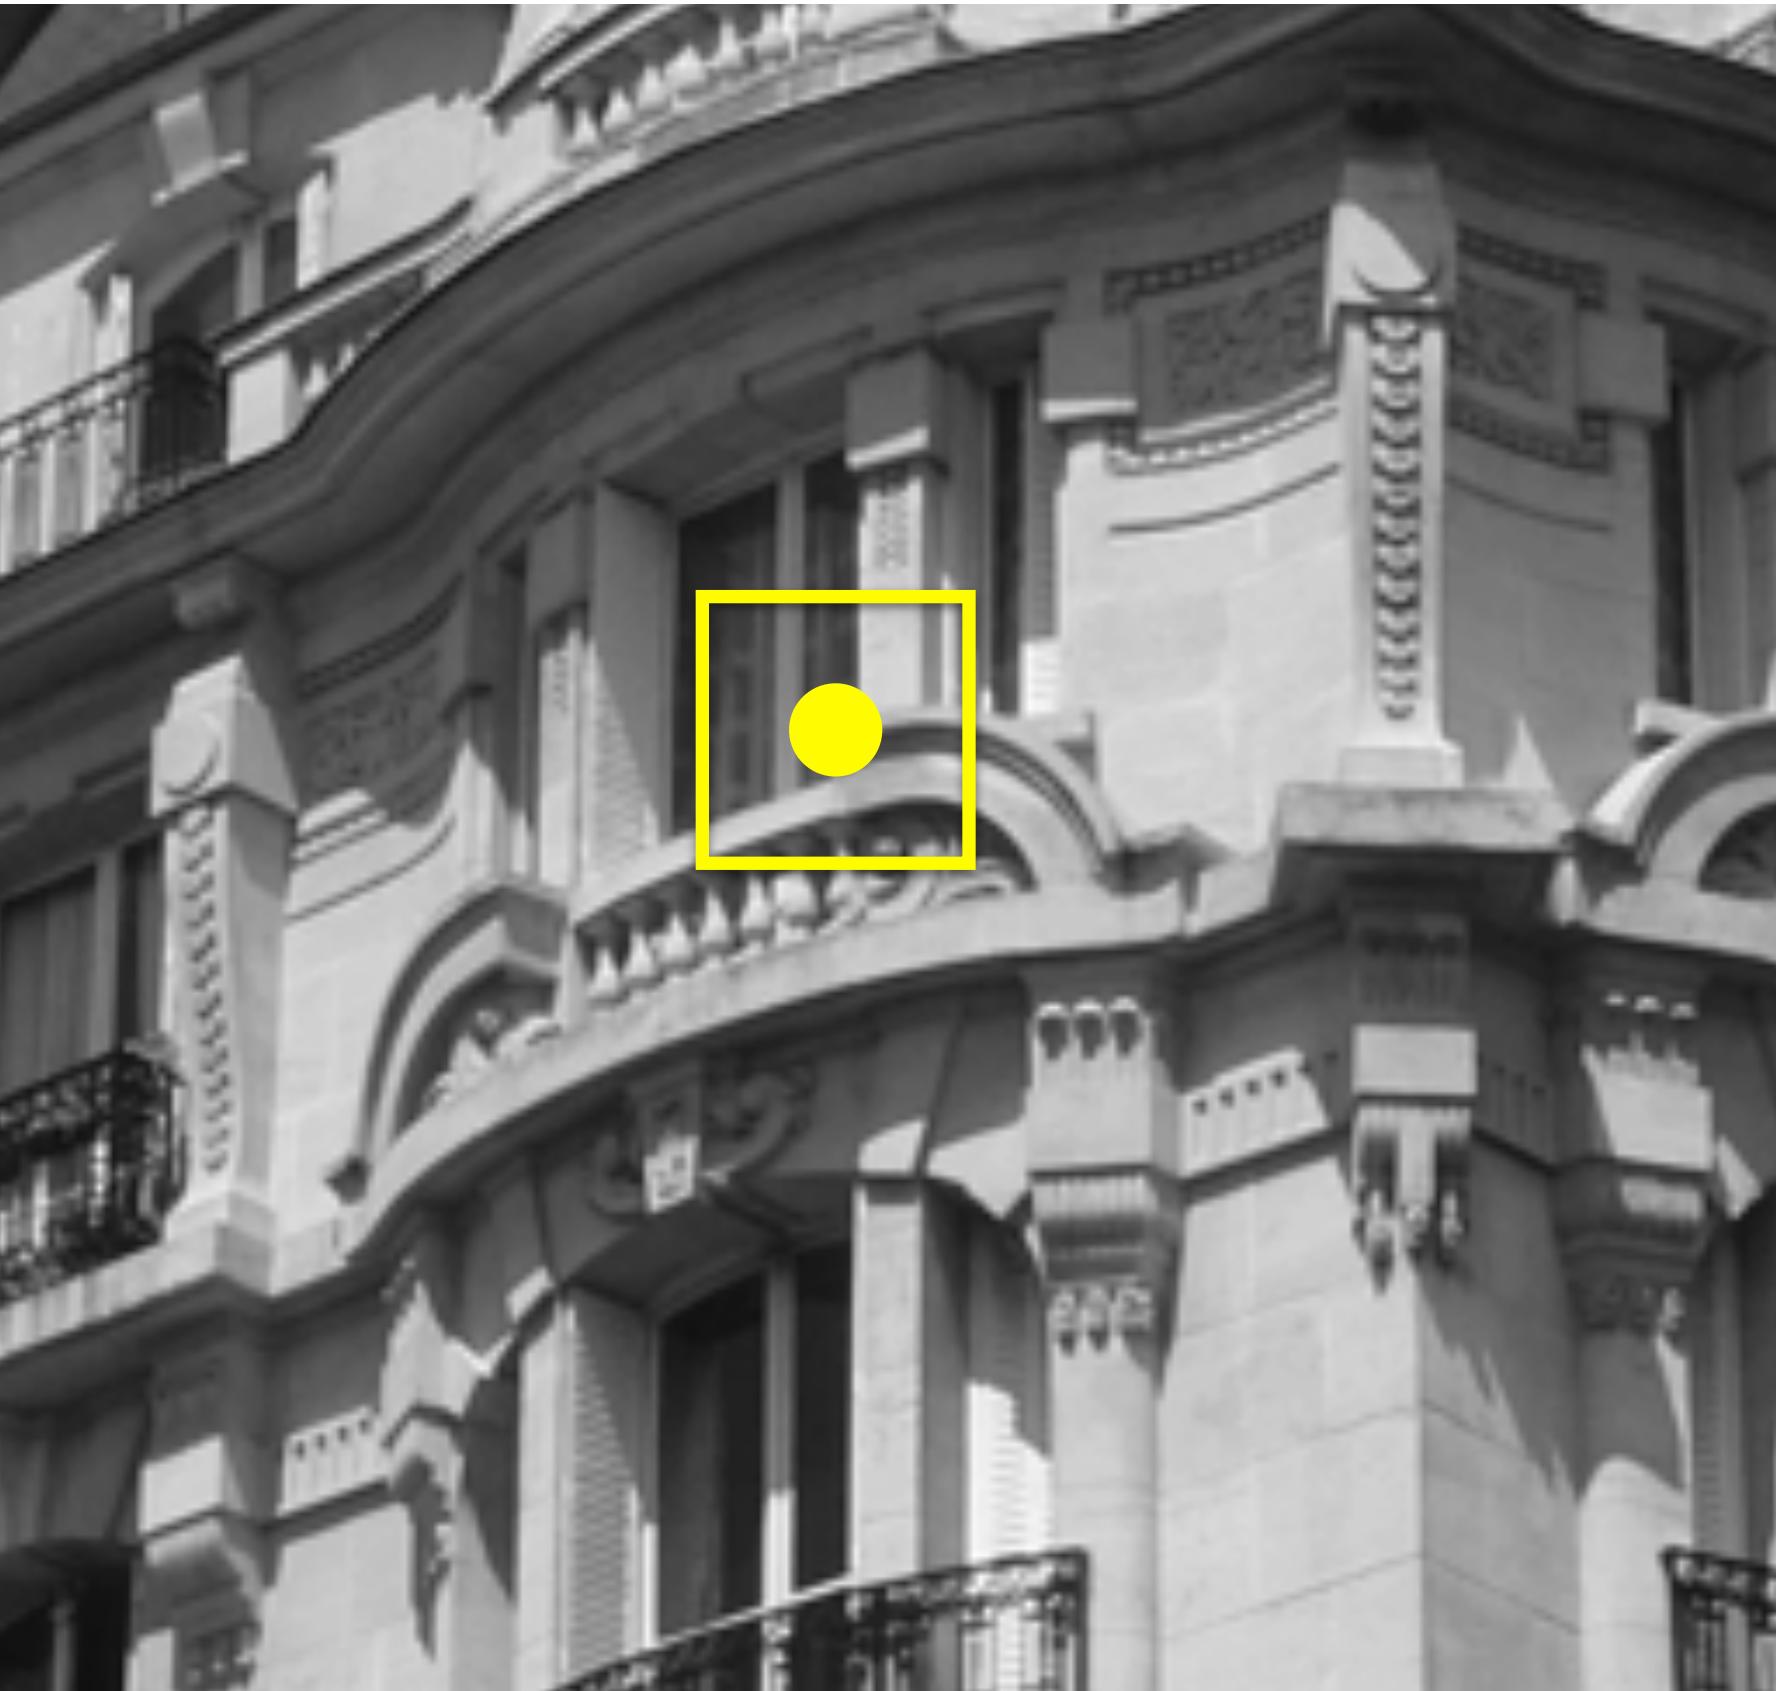
Harris feature recap

- Concise:
 - ➡ hundreds of features instead of millions of pixels
 - ➡ “*a description that is useful for the viewer and not cluttered with irrelevant information*” (Marr)
- Finds points that are distinct and easily located in a different view of the same scene
- Computationally efficient (good for real-time tracking)

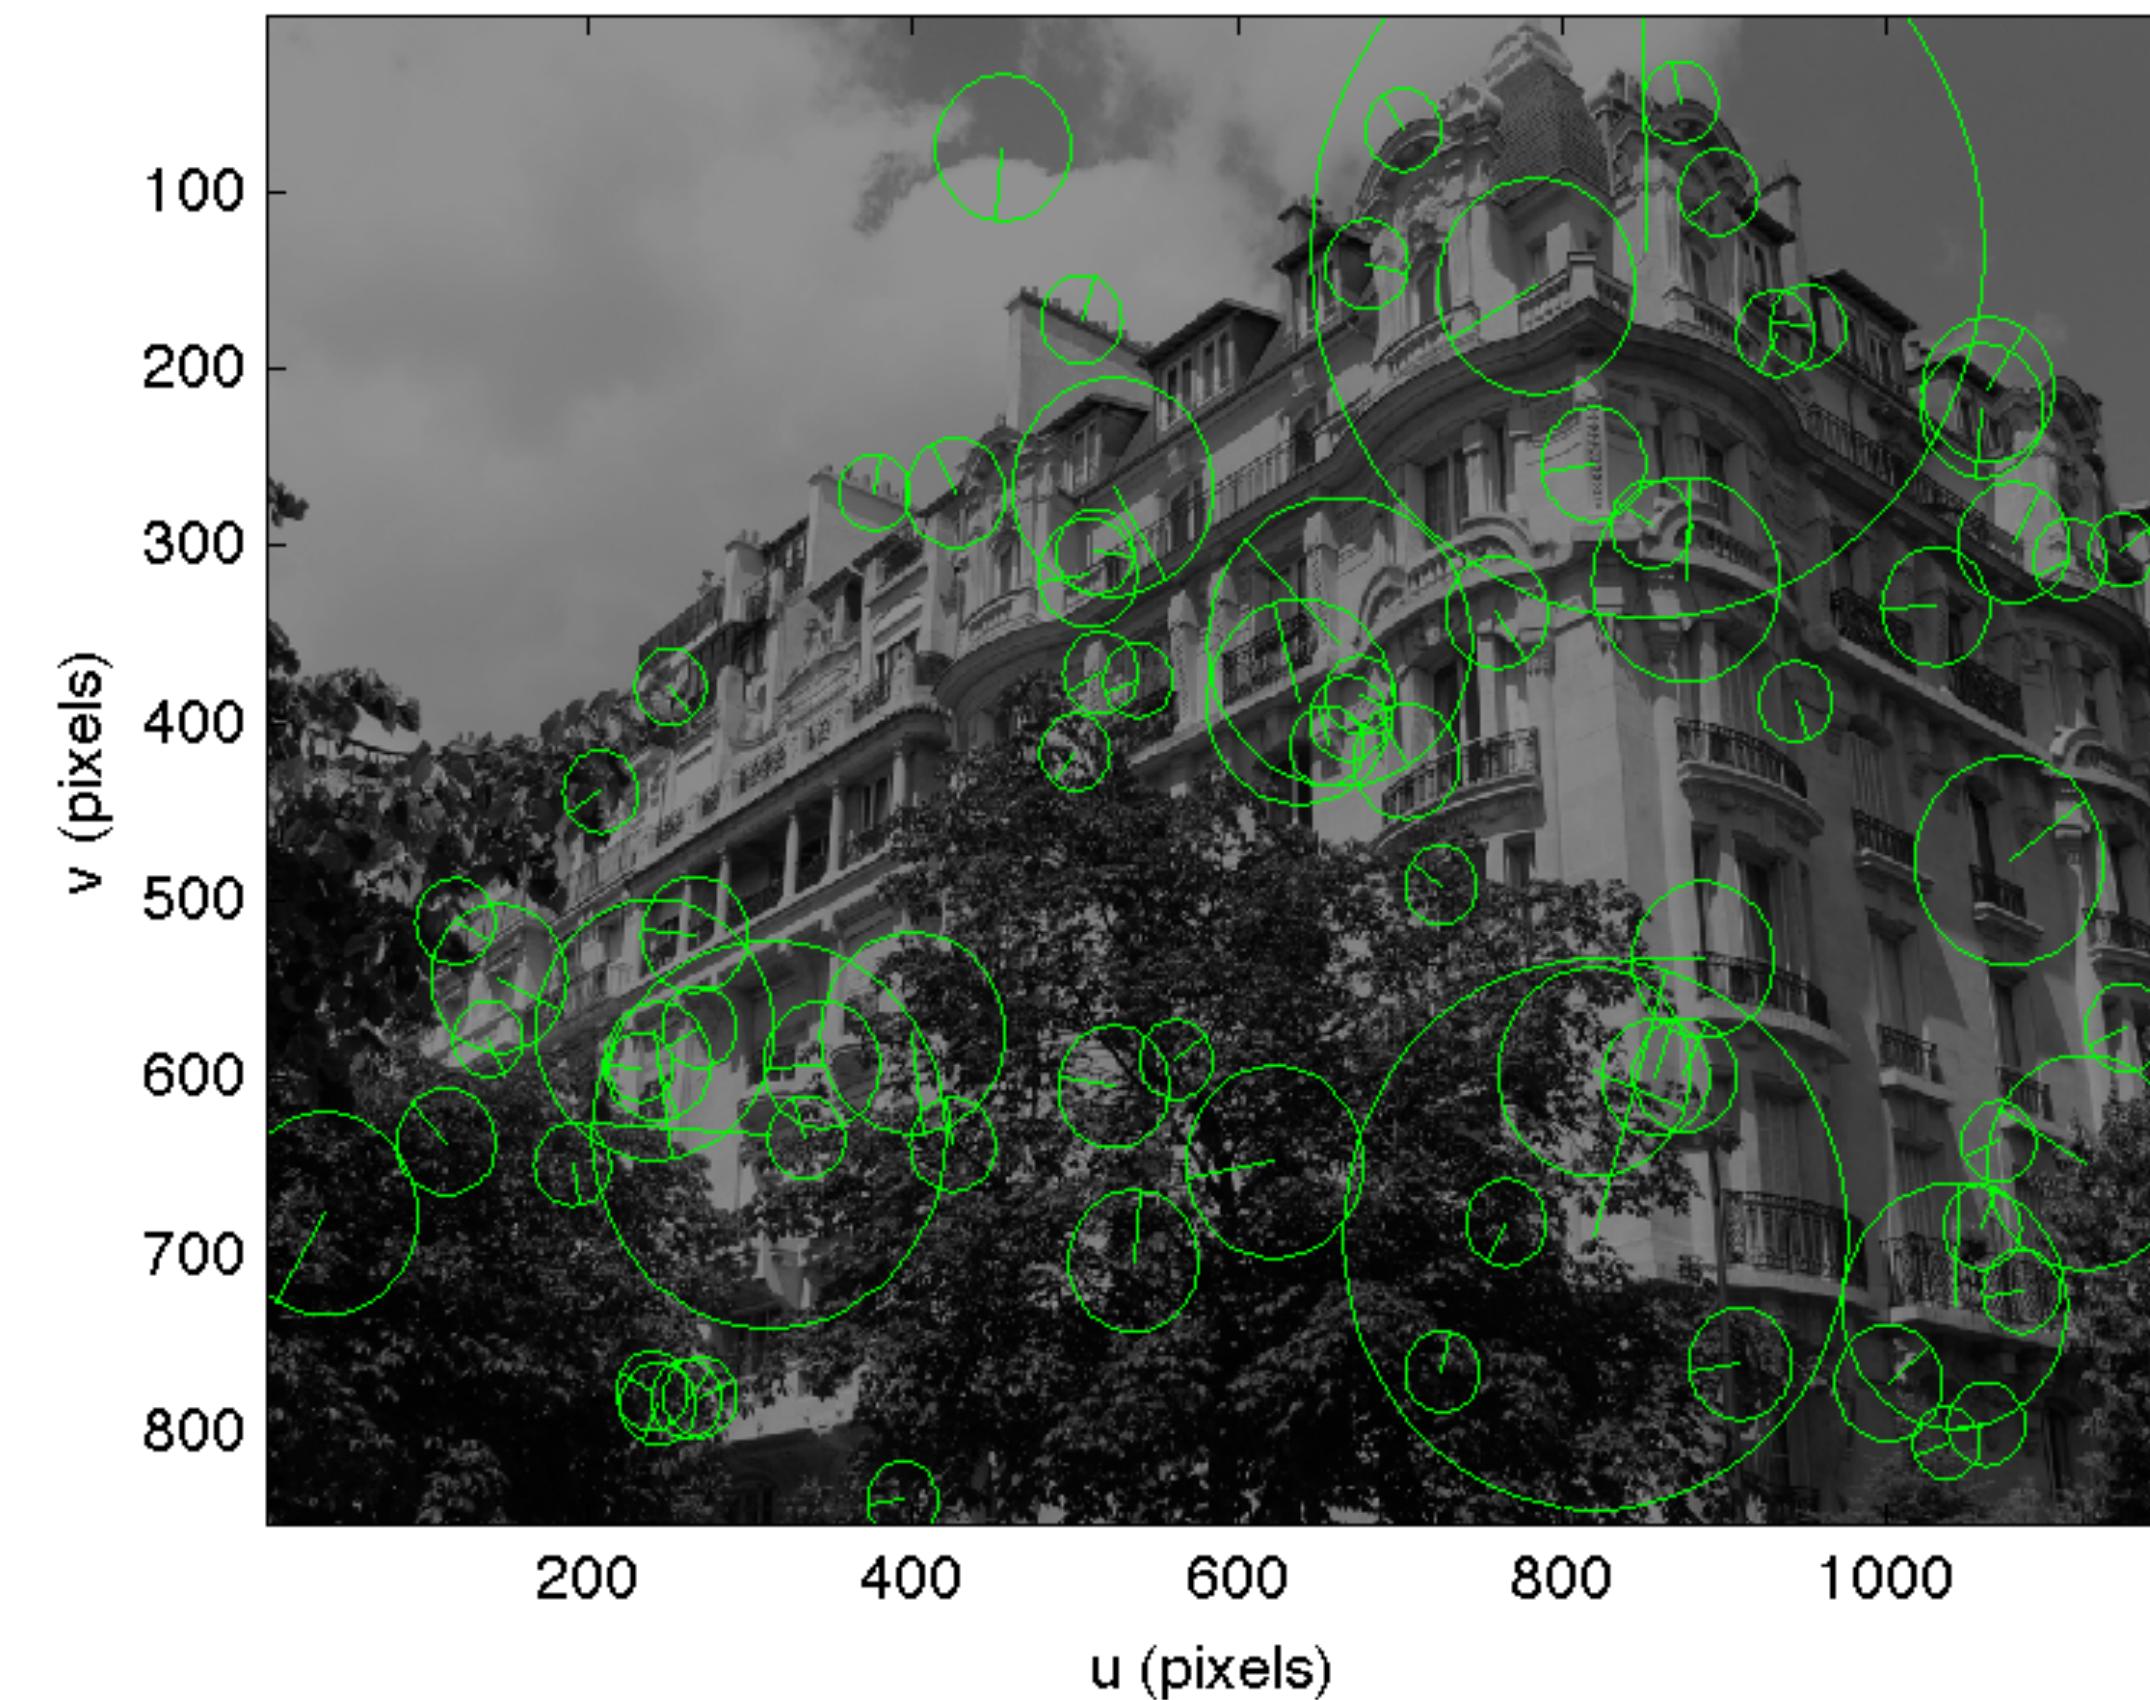
The problem of scale



The problem of orientation



Scale Invariant Feature Transform



SIFT detector (Lowe, 2004)

Harris vs SIFT

